

MATH5703M Statistical Theory

Jochen Voss

September 30, 2026

Contents

Introduction	vii
1 Statistical Models and Estimation	1
1.1 Statistical Models	1
1.2 Statistics and Estimators	3
1.3 The Road Ahead	4
2 Bias, Mean Squared Error and Consistency	6
2.1 Bias and Unbiasedness	6
2.2 Mean Squared Error	8
2.3 When Unbiasedness Is Not Enough	9
2.4 Consistency	10
3 R Basics and a First Simulation	13
3.1 First Steps in R	13
3.2 Simulating a Normal Sample	14
3.3 The Sample Mean as a Random Quantity	15
4 Likelihood and the Method of Moments	18
4.1 The Likelihood Function	18
4.2 The Score Function	19
4.3 The Method of Moments	20
4.4 Identifiability	22
5 Maximum Likelihood Estimation	25
5.1 The Maximum Likelihood Estimator	25
5.2 Worked Examples	26
5.3 When the Recipe Fails	27
5.4 Bias of the Maximum Likelihood Estimator	28
5.5 Invariance Under Reparametrisation	29
6 Estimators by Simulation	33
6.1 Functions and Loops	33
6.2 Empirical Bias and Mean Squared Error	34
6.3 Comparing Two Estimators	36
7 Exponential Families	39
7.1 Recognising the Shape	39
7.2 The Natural Parametrisation and the Cumulant Function	41
7.3 Several Parameters	42

7.4	What the Structure Buys	43
8	Sufficiency	47
8.1	Sufficient Statistics	47
8.2	The Factorisation Theorem	48
8.3	Sufficiency in Exponential Families	49
8.4	Minimal Sufficiency and Completeness	50
8.5	The Sample Maximum and Minimum	51
9	Consistency by Simulation	55
9.1	A Running Estimate	55
9.2	The Mean Squared Error as a Function of the Sample Size	56
9.3	A Sufficient Statistic in R	59
9.4	A Task to Take Home	60
10	Improving Estimators: Rao–Blackwell and Lehmann–Scheffe	62
10.1	The Rao–Blackwell Theorem	62
10.2	Completeness and the Lehmann–Scheffe Theorem	63
10.3	Ancillary Statistics and Basu’s Theorem	65
11	Fisher Information	69
11.1	Setting and Assumptions	69
11.2	The Score Has Mean Zero	70
11.3	Fisher Information	70
11.4	Examples	72
11.5	Exponential Families	72
11.6	Reparametrisation	73
11.7	When the Identities Fail	74
12	Sampling Distributions and Standard Errors	77
12.1	The Sampling Distribution of an Estimator	77
12.2	Standard Errors and a Wald Interval	78
12.3	Estimating a Probability by Simulation	79
12.4	Skewed Data	80
13	The Cramer–Rao Inequality	83
13.1	The Inequality	83
13.2	Estimating a Function of the Parameter	85
13.3	Efficiency and Attainment	85
13.4	When the Conditions Fail	87
14	Multiparameter Fisher Information	90
14.1	The Score Vector and the Information Matrix	90
14.2	Properties of the Information Matrix	91
14.3	The Cramer–Rao Bound for Vector Parameters	92
14.4	The Normal Distribution	93
14.5	Nuisance Parameters	94

15 Cramer–Rao and Efficiency in R	97
15.1 An Efficient Estimator Against the Bound	97
15.2 Attaining and Missing the Bound	98
15.3 Beating the Bound Without Breaking It	99
15.4 A Task to Take Home	101
16 Consistency of the MLE	103
16.1 The Regularity Conditions	103
16.2 The Likelihood Ratio Detects the Truth	104
16.3 Consistency	105
16.4 When the Assumptions Fail	106
17 Asymptotic Normality and Efficiency of the MLE	110
17.1 The Theorem	110
17.2 Proof of the Theorem	111
17.3 Standard Errors	112
17.4 Examples	113
17.5 Quality of the Approximation	114
18 Maximum Likelihood in R	117
18.1 The Log-Likelihood as an R Function	117
18.2 Standard Errors from the Observed Information	118
18.3 Starting Values	119
18.4 Coverage of the Wald Interval	120
18.5 A Task to Take Home	122
19 The EM Algorithm	124
19.1 Incomplete Data	124
19.2 The Algorithm	125
19.3 The Ascent Property	126
19.4 A Two-Component Normal Mixture	127
19.5 Practical Matters	129
20 Hypothesis Testing	131
20.1 Hypotheses	131
20.2 Tests, Critical Regions and Errors	131
20.3 The Power Function, Size and Level	132
20.4 Testing the Mean of a Normal Sample	133
20.5 The p -Value	134
21 EM in R	138
21.1 The E and M Steps as R Functions	138
21.2 A Larger Sample	139
21.3 Label Switching and a Bad Start	142
21.4 A Task to Take Home	143
22 The Neyman–Pearson Lemma	145
22.1 Simple Hypotheses and Most Powerful Tests	145
22.2 The Lemma	146
22.3 The Normal Mean	147

22.4 One-Sided Alternatives	148
22.5 Two-Sided Alternatives	149
23 Likelihood-Ratio Tests	152
23.1 The Generalised Likelihood Ratio	152
23.2 The Two-Sided Test for the Mean of a Normal Distribution . . .	153
23.3 Wilks' Theorem	154
23.4 Using the Theorem	155
24 Power and Size by Simulation	159
24.1 The Size of a Test by Simulation	159
24.2 The Power Curve	160
24.3 Wilks' Theorem in Action	161
24.4 A Task to Take Home	164
25 Confidence Intervals	166
25.1 Definition and Interpretation	166
25.2 Pivotal Quantities	167
25.3 Inverting a Test	168
25.4 Approximate Intervals from the MLE	169
25.5 One-Sided Intervals and the Choice of Interval	170
26 Bayesian Inference: Priors and Posteriors	173
26.1 Prior, Likelihood and Posterior	173
26.2 Conjugate Priors	174
26.3 Point Estimates from the Posterior	175
26.4 Flat, Improper and Jeffreys Priors	176
27 Confidence Versus Credibility	180
27.1 Coverage by Simulation	180
27.2 The Same Data, Two Intervals	182
27.3 From Prior to Posterior	183
28 Credible Intervals and Bayesian Testing	186
28.1 Credible Intervals	186
28.2 Bayesian Testing	187
28.3 Comparing Frequentist and Bayesian Answers	189
29 Standard Tests Revisited	193
29.1 Tests from Pivots	193
29.2 The One-Sample t -Test	194
29.3 The Chi-Squared Test for a Variance	195
29.4 Paired and Two-Sample Problems	196
29.5 Which Test?	197
30 Do the Standard Tests Work?	200
30.1 Real Data, Two Answers	200
30.2 The Replication Crisis in Miniature	201
30.3 A Stress Test	204
30.4 A Task to Take Home	205

31 The Exponential Model: Start to Finish	207
31.1 The Model and the Data	207
31.2 Point Estimation	208
31.3 Interval Estimation	208
31.4 Testing	209
31.5 The Bayesian Answer	209
31.6 Looking Back and Looking Forward	210
Preparing for the Exam	214
A Background from Probability	222
A.1 Expectation, Variance and Independence	222
A.2 Standard Distributions	223
A.3 Sampling from a Normal Population	224
A.4 Some Inequalities	225
A.5 Conditional Expectation	226
A.6 Convergence and Limit Theorems	227
A.7 Transformation of a Random Variable	229
A.8 Bayes' Theorem	229
B Introduction to R	230
B.1 Getting Started	230
B.2 Vectors and Arithmetic	231
B.3 Lists	233
B.4 Random Numbers and Distributions	234
B.5 Writing Your Own Code	235
B.6 Estimation by Simulation	237
B.7 Plotting, and Getting Results into a Report	239
B.8 Numerical Maximum Likelihood	240
C Solutions to the Exercises	242
C.1 Statistical Models and the Estimation Problem	242
C.2 Bias, Mean Squared Error and Consistency	243
C.3 Likelihood and the Method of Moments	245
C.4 Maximum Likelihood Estimation	247
C.5 Exponential Families	252
C.6 Sufficiency	256
C.7 Rao–Blackwell and Lehmann–Scheffe	260
C.8 Fisher Information	264
C.9 The Cramer–Rao Inequality	267
C.10 Multiparameter Fisher Information	271
C.11 Consistency of the MLE	274
C.12 Asymptotic Normality of the MLE	276
C.13 The EM Algorithm	277
C.14 Hypothesis Testing	280
C.15 The Neyman–Pearson Lemma	282
C.16 Likelihood-Ratio Tests	285
C.17 Confidence Intervals	288
C.18 Bayesian Inference: Priors and Posteriors	290
C.19 Credible Intervals and Bayesian Testing	294

<i>CONTENTS</i>	vi
C.20 Standard Tests Revisited	297
C.21 The Exponential Model: From Start to Finish	300
Bibliography	305
Index	307

Introduction

These are the lecture notes for the Statistical Theory (MATH5703M) module. The question we will consider in this module, in many different forms, is the following: What can we learn about the parameters of a model if we are given data from the model?

The first half of the module is concerned with estimation. We will start by considering statistical models and properties of an estimator, using the bias and mean squared error as measures of quality. We will then consider two different approaches to constructing estimators: maximum likelihood estimators and the method of moments. The unifying framework for this part of the module is the exponential family of distributions. Sufficiency, completeness and the Rao–Blackwell and Lehmann–Scheffe theorems will be used to study how to improve an estimator and to identify cases where an estimator cannot be improved. Finally, we will consider Fisher information and the Cramer–Rao bound to study limitations on the precision of unbiased estimators. We will conclude the study of estimation by considering large-sample theory for the maximum likelihood estimator and the EM algorithm.

The second half of the module is concerned with testing and interval estimation. We will develop the Neyman–Pearson theory of hypothesis tests and the likelihood-ratio test, and will then consider confidence intervals. A section on Bayesian inference will cover the choice of prior, credible intervals and Bayesian testing. We will conclude the study of testing by considering standard tests used in practice.

Alongside the theory, we will use simulation in R to illustrate the ideas of the module. Appendix B collects the R commands we use.

How to Read the Proofs

Many of the benefits of this module can be gained by studying the proofs, but a different approach is needed to read and understand a proof than is needed to understand ordinary text. The technique we will use here, called *self-explanation*, is the only technique for reading proofs which is supported by strong experimental evidence. In three experiments with undergraduate mathematics students, [Hodds et al. \(2014\)](#) showed that a short self-study training in self-explanation, consisting of around twenty minutes of slides and one worked example and one practice proof, significantly improved scores in proof-comprehension tests. The improvement in the laboratory experiment was close to one standard deviation. A fifteen-minute booklet version of the training, completed during a scheduled lecture, also showed a smaller but significant improvement, and the effect of this was still visible three weeks later. We will practise this technique together in

the first week of the module, and my hope is that you will use the technique throughout the module.

The idea of self-explanation is that you should pause after every line of the proof, and try to explain the line to yourself, before you read the next line. A good self-explanation answers the following three questions.

- What does this line state?
- Why is this statement true? Be specific about how you know: is it from an earlier line of the proof, or from a definition, or from an assumption of the theorem, or from a previous result?
- How does this line contribute to the proof? How does it move the argument towards the statement?

There are two bad habits which look similar to self-explanation, but are not: paraphrasing, *i.e.* restating a line in different words without understanding the reason, and ‘monitoring’, *i.e.* saying ‘yes, this makes sense’ and then carrying on. Both of these techniques help your understanding less than self-explanation does. It is not enough to just have a feeling that you understand a proof; you also need to be able to reconstruct the argument line by line. If you cannot name the reason for a statement, you have a gap in your understanding and you need to find it: write down exactly what is missing, look it up or work it out, and only then continue.

In practice this means reading with a pencil. Before you read the proof, read the statement carefully, and make sure you understand what is assumed and what is claimed. Then, as you read the proof, write down the justification for each step in the margin. If you cannot find a justification for a step, try to fill the gap yourself before you read the next line. This approach will make reading a proof slower, but it will also make sure that you understand the proof. The same technique can be used to understand the worked solutions in appendix C: try to solve the exercise yourself first, and only then look at the solution line by line to check your work.

What About AI?

Large language models are very good at the kind of material we cover in this module. Using one of these models, it is possible to get a solution to any of the exercises in these notes, to get answers for the examination questions, and even to get R code for most of the tasks required for the practical report.

The problem with this is that the aim of the module is for you to learn the material in these notes, not to solve the exercises (solutions for all of the exercises are in the appendix). While AI can explain every line of a proof to you, this is not the same as self-explanation. Similarly, the practical report requires you to show your judgement and to explain your work, but a report written by the model does not give you practice in either of these skills.

There is evidence of this problem in scientific studies: [Strömberg et al. \(2026\)](#) studied the effect of such tools on 26 811 secondary-school students over a period of 30 months: homework scores increased by 18%, homework time decreased by 30% and after six months scores in the monthly, closed-book examinations had

decreased by 20%. The losses were concentrated amongst the approximately 80% of users who completed their homework faster than any non-user and scored high marks for the homework, and the authors of the study interpreted these results as indicating that these users had “outsourced” their homework to the tool. Users who spent as much time on their homework as non-users did performed about as well in the examinations, so the loss in this group was not due to the tools themselves, but due to the way the tools were used.

A randomised trial in mathematics (Bastani et al., 2025) showed the same effect. Around 1000 secondary-school students either used a language model to learn the material or used only the textbook and notes. The students were then tested in a closed-book exam. The students who used the unrestricted language model scored 17% worse than the others, mainly because they used the model to find answers to the questions. When a version of the model was used which only gave hints instead of answers, the loss in performance disappeared. The students were not aware of the difference: students who used the unrestricted model did not believe that they had learned less than the other students.

I want to also mention that the exam for our module is closed book, pen and paper based. In the exam, you will not have access to AI support, and will need to rely on what you have personally learned.

There are ways to use these models to help understanding, but these require a slightly different approach: you ask the model for something other than the answer to the question at hand.

- Ask for more practice material. The model can easily generate more exercises for you. But remember that the only way to learn from these exercises is to solve them yourself!
- Ask for a second explanation. If a definition, a result or a step in a proof is not clear to you, ask the model to explain it in different words; sometimes a second explanation helps you to understand it. Afterwards, close the window and write it down from memory: a definition or result in full, a step in a proof by giving the reason why it holds in your own words. Reading the model’s explanation is only monitoring, while writing it down from memory shows that you have understood or at least memorised it.
- Ask which result a step relies on. If you are stuck on a step of a proof which you cannot yet justify, the model can sometimes name the result the step uses, for example Jensen’s inequality or the tower property of conditional expectation. Ask for the name only, look up the result in the index and work out the step yourself.
- Ask why, not what. A question like “Why do we divide by $n - 1$ in the sample variance?” is a question which you can first try to answer yourself, and then check your answer with the model’s answer. In contrast, a question like “solve this exercise” leaves nothing for you to do.
- Ask about R. What does an error message mean? What does a given argument of a command do? These are questions which can be answered by the help pages (see appendix B). It is not expected that you will know these details by heart from this module, so asking the model is fine.

One way to test whether you have understood an argument is to try to re-create the argument on paper, from memory, a week later. If you can do this, you have learned something! If you cannot do this, you only have a solution, not an understanding.

Reading

These notes are self-contained, but the following books can be used as additional reading. The links lead to the University of Leeds Library catalogue.

- G. Casella and R. L. Berger, *Statistical Inference* (2nd edition; available online). This is our main text. It covers most of the material of the module, from point estimation to sufficiency and the information inequality, and from hypothesis testing to confidence intervals.
- G. A. Young and R. L. Smith, *Essentials of Statistical Inference* (print only). A concise and rigorous graduate text at the level of this module. The book is particularly good on likelihood, decision theory and the large-sample behaviour of the maximum likelihood estimator.
- G. H. Givens and J. A. Hoeting, *Computational Statistics* (2nd edition; available online). This book is useful for learning about the EM algorithm and about numerical optimisation of the likelihood. It also introduces the computational ideas which we use in the R sessions.
- P. M. Lee, *Bayesian Statistics: An Introduction* (4th edition; available online). A self-contained introduction to the Bayesian material we cover, including the choice of prior, Jeffreys priors and Bayesian testing.
- W. J. Braun and D. J. Murdoch, *A First Course in Statistical Programming with R* (3rd edition; available online). This book is designed to support the R sessions and the practical report. It teaches programming and simulation in R from the ground up.

The following list has additional resources which either are more specialised and go deeper, or offer a gentler alternative to the more standard texts listed above.

- D. R. Cox and D. V. Hinkley, *Theoretical Statistics* (available online). The classic deeper treatment of the theory behind the module.
- E. L. Lehmann and G. Casella, *Theory of Point Estimation* (2nd edition; print only). The advanced reference for the estimation half: sufficiency, completeness and efficiency in full detail.
- A. C. Davison, *Statistical Models* (print only). Broad and likelihood-centred, with the EM algorithm and many worked applications.
- A. Gelman et al., *Bayesian Data Analysis* (3rd edition; available online). The standard modern reference for Bayesian statistics, well beyond what we need but the natural next step.

- M. H. DeGroot and M. J. Schervish, *Probability and Statistics* (4th edition; available online). A gentle and thorough account, useful for revising the prerequisite material from MATH2701.
- J. A. Rice, *Mathematical Statistics and Data Analysis* (3rd edition; print only). An accessible companion pitched a little below the core, good for a second explanation.

Lecture 1

Statistical Models and the Estimation Problem

The subjects of probability and statistics are in some sense the opposite of each other. In a probability course we are given a distribution $\mathbb{P}$ and we ask questions about the data generated by this distribution: how often should we expect ten tosses of a fair coin to result in seven heads? In contrast, in a statistics course we are given data and we ask questions about the distribution: what can we learn about the coin by observing seven heads?

probability: $\mathbb{P} \longrightarrow X_1, \dots, X_n$, statistics: $X_1, \dots, X_n \longrightarrow \mathbb{P}$.

The topic of this module is the second arrow in the above discussion. We will consider two questions: in *estimation* we consider how (and how well) we can determine an unknown quantity from data. In *hypothesis testing* we consider whether the data are compatible with a given claim. We will answer these questions first for the classical approach, and then for the Bayesian approach. In this first lecture we will lay the foundations for these questions by introducing the model and the estimators computed from the data.

1.1 Statistical Models

The data are given as $x_1, \dots, x_n$, a list of n numbers. Once we have computed an estimate from these numbers, we want to know how far it is from the truth, and this only makes sense if there is a true value to compare the estimate to. The data themselves do not contain a ‘true value’. Instead, this ground truth is provided by a model, a mathematical description of how the data were generated.

Our description of how the data were generated is called a model. We use a model to describe the data as the outcome of a random experiment. The data we have observed, $x_1, \dots, x_n$, are seen as the observed values of random variables $X_1, \dots, X_n$, where the distribution of the random variables depends on the unknown true value. Here we use the convention that upper case letters like X_i stand for random variables, and that lower case letters like x_i stand for the observed values for our dataset.

In the simplest and most common situation the observations are a *random sample*: the X_i are independent and identically distributed (i.i.d.), but we do not know the distribution. Instead, we assume that the distribution belongs to a known family of distributions, labelled by a parameter.

Definition 1.1. A *parametric model* for data $X_1, \dots, X_n$ consists of a family

$$(\mathbb{P}_\theta \mid \theta \in \Theta)$$

of distributions, and the assumption that the $X_1, \dots, X_n$ are i.i.d. with distribution $\mathbb{P}_\theta$ for an unknown value of the *parameter* θ . The set Θ of possible parameter values is called the *parameter space*.

Our assumption here is that the actually observed data, $x_1, \dots, x_n$, are a sample from $\mathbb{P}_\theta$ for one (unknown to us) θ . This value is called the *true parameter value* and is denoted by θ_0 . This is the value which our estimates are compared to: an estimate is good, if it is close to θ_0 , and bad if it is far away from θ_0 .

Since the distribution is determined by the parameter, learning about the distribution also means learning about the parameter. Thus, the diagram from the beginning of this lecture can be updated to

$$\begin{array}{lcl} \text{probability:} & \theta_0 & \longrightarrow \mathbb{P}_{\theta_0} \longrightarrow X_1, \dots, X_n, \\ \text{statistics:} & X_1, \dots, X_n & \longrightarrow \theta_0. \end{array}$$

The true parameter value can never be observed directly, and all information about the true parameter value comes from $X_1, \dots, X_n$.

The parameter θ does not necessarily have to be a single number. For example, if we assume that the data are normally distributed with unknown mean μ and unknown variance σ^2 , then $\mathbb{P}_\theta = N(\mu, \sigma^2)$ and the parameter is the pair $\theta = (\mu, \sigma^2)$. In this case, $\Theta = \mathbb{R} \times (0, \infty)$ is a subset of the plane. While we allow for more general parameters from the start, most of the examples in this text will involve a scalar parameter θ .

The distribution $\mathbb{P}_\theta$ describes, for every set A of possible values, the probability $\mathbb{P}_\theta(A)$ that an observation falls into A . In practice, a distribution is almost always given by a function $f(x; \theta)$ of the possible values x , and this can be done in two different ways. For discrete distributions, where the observations can take finitely or countably many values, *e.g.* for $\{0, 1\}$ or $\{0, 1, 2, \dots\}$, we have $f(x; \theta) = \mathbb{P}_\theta(\{x\})$, giving the probability of observing the value x . These numbers are called the *probability weights* of the distribution (sometimes also called the *probability mass function*). These weights are numbers in the interval $[0, 1]$, summing to one, and satisfying

$$\mathbb{P}_\theta(A) = \sum_{x \in A} f(x; \theta).$$

For continuous distributions, $f(x; \theta)$ is a *density*: a non-negative function, which is integrable and has integral one, such that

$$\mathbb{P}_\theta(A) = \int_A f(x; \theta) \, dx.$$

For continuous distributions, each value x has probability $\mathbb{P}_\theta(\{x\}) = 0$, and a density can take values larger than one. The function values $f(x; \theta)$ of a density are not probabilities, but integrals of the density over sets give probabilities.

In both cases, the family of functions $(f(x; \theta) \mid \theta \in \Theta)$ determines the model, and we will often write the model in this form. Much of the theory in this module is common to both cases. Where a statement is true for both cases, we

sometimes only state it for densities, and then the discrete case can be deduced by interpreting $f(x; \theta)$ as the probability weights and all integrals over x as sums.

We use the notation $\mathbb{P}_\theta$ also for the probabilities of events involving the data, and we write $\mathbb{E}_\theta$ for expectations, to indicate that the result depends on the value of the parameter. For example, $\mathbb{P}_\theta(X_1 \leq t) = \mathbb{P}_\theta((-\infty, t])$ is the probability that a single observation is at most t , when the data have distribution $\mathbb{P}_\theta$.

We will consider a small number of families of models throughout the module. The discrete families are listed in table 1.1 and the continuous ones in table 1.2. The corresponding means and variances, which we will need frequently in calculations, are shown in tables A.1 and A.2 in appendix A.

Family	Weights $f(x; \theta)$	Support	Θ
Bernoulli	$\theta^x(1 - \theta)^{1-x}$	$\{0, 1\}$	$(0, 1)$
Binomial	$\binom{m}{x} p^x(1 - p)^{m-x}$	$\{0, 1, \dots, m\} \times (0, 1)$	
Poisson	$\frac{\theta^x}{x!} e^{-\theta}$	$\{0, 1, 2, \dots\}$	$(0, \infty)$

Table 1.1: Discrete families of models considered in the module. For the binomial distribution the parameter is the pair $\theta = (m, p)$.

Family	Density $f(x; \theta)$	Support	Θ
Exponential	$\theta e^{-\theta x}$	$[0, \infty)$	$(0, \infty)$
Normal	$\frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$	$\mathbb{R}$	$\mathbb{R} \times (0, \infty)$
Uniform	$1/\theta$	$[0, \theta]$	$(0, \infty)$

Table 1.2: Continuous families of models considered in the module. For the normal distribution the parameter is the pair $\theta = (\mu, \sigma^2)$.

For families where θ is a vector, usually one component is assumed known and fixed, and only the remaining component is estimated. For example, for the binomial distribution the number m of trials is often assumed to be known, and for the normal distribution the variance σ^2 is often assumed to be known.

Example 1.2. Consider the model $X_1, \dots, X_n \sim \text{Uniform}(0, \theta)$ with $\theta > 0$, *i.e.* each observation is equally likely to take a value anywhere in the interval $[0, \theta]$ and $f(x; \theta) = 1/\theta$ for $0 \leq x \leq \theta$, and 0 otherwise. Since no observation can exceed θ , if θ increases, larger data values are possible. The range of the data can be used to draw inference about the parameter. This is in contrast to the normal distribution or the exponential distribution, where all values are possible for all θ . We will see later that the moving support of the uniform distribution leads to different behaviour than for the other models, and this example will return as a test case.

1.2 Statistics and Estimators

Once a model is fixed, we want to extract information about the unknown θ from the data. Since we do not know the value of θ , we need to consider quantities

which can be computed from the data.

Definition 1.3. A *statistic* is a function $T = T(X_1, \dots, X_n)$ of the data, which does not depend on the parameter θ .

Under this restriction the sample mean $\bar{X} = (X_1 + \dots + X_n)/n$ is a statistic, but $X_1 - \theta$ is not, since we need to know the parameter value to compute it.

Definition 1.4. An *estimator* for θ is a statistic $\hat{\theta} = \hat{\theta}(X_1, \dots, X_n)$, which can be used to approximate θ . The number $\hat{\theta}(x_1, \dots, x_n)$, computed from a given dataset $x_1, \dots, x_n$, is called an *estimate*.

Throughout the notes we will use hats to denote estimators: $\hat{\theta}$ is an estimator for θ , $\hat{\sigma}^2$ is an estimator for σ^2 , and so on. If two estimators for the same quantity are compared, the second one is denoted by a tilde: $\tilde{\theta}$ is a competitor of $\hat{\theta}$.

Most estimators are defined by a formula which works for all sample sizes, *e.g.* $\bar{X} = (X_1 + \dots + X_n)/n$. Technically, this formula defines a *family of estimators* $(\hat{\theta}_n)_{n \in \mathbb{N}}$, one for each sample size n . When we need to consider the dependence on n , and in particular when we study the behaviour of an estimator as the amount of data increases, we write $\hat{\theta}_n$ instead of $\hat{\theta}$.

An important note is that the definition does not imply anything about quality. Any statistic is an estimator, for example the constant function $\hat{\theta} = 0$ is a (bad) estimator for a Bernoulli success probability. We will discuss how to find and recognise good estimators in the following lectures.

The difference between an estimator and an estimate is a subtle but important one. An *estimate* is a single number, computed from the data, for example $\hat{\theta} = 0.7$. An *estimator* is a rule for computing such a number, typically a function of the random variables $X_1, \dots, X_n$, such that different outcomes of the experiment give rise to different values of the estimate.

We can only make statements about how good an estimator is, because $\hat{\theta}(X_1, \dots, X_n)$ is a random variable. The *sampling distribution* of an estimator is the distribution of the resulting estimates. We can ask questions like the following: If the data are generated by the true value θ_0 , how are the estimates distributed around θ_0 ? Is the distribution centred around the truth? If so, how large is the spread? A good estimator will have a sampling distribution which is concentrated around the true parameter value, for all possible values of the true parameter. In the next lecture we will learn how to quantify these ideas using the concepts of bias and mean squared error.

1.3 The Road Ahead

Now that we have established the model and the estimator, we can start to build the module. The topics we will cover include point estimation (both likelihood and method of moments), measurement and comparison of estimators, the best possible estimators (sufficiency, information in a sample), and the limits of estimation. We will then consider statistical tests and the related question of compatibility of data with a given hypothesis about θ . Finally, we will consider both estimation and testing from a Bayesian perspective, where the parameter is assumed to be random.

Summary.

- Statistics is the inverse of probability: we use data to learn about the model, rather than using the model to learn about the data.
- A parametric model assumes that the data are i.i.d. from a family $(\mathbb{P}_\theta \mid \theta \in \Theta)$ of distributions, for an unknown true parameter value θ_0 . The model provides the true value which an estimate is compared to.
- Discrete distributions are given by probability weights, which are probabilities, while continuous distributions are given by a density, which is not a probability. Probabilities are calculated as sums of weights or as integrals of the density.
- A statistic is a function of the data. An estimator is a statistic which is used to approximate θ , and an estimate is the value of the estimator for a given dataset.
- An estimator is a random variable, and we study its sampling distribution. A good estimator will have a sampling distribution concentrated around the true value θ_0 , for all possible values of θ_0 . In the next lecture we will learn how to quantify this idea.

Exercise 1.1. A coin with unknown probability θ of heads is tossed n times. The results of the individual tosses are given by $X_i = 1$, if the i -th toss is heads, and $X_i = 0$ otherwise.

1. Write down a parametric model for $X_1, \dots, X_n$, including the probability weights and the parameter space.
2. Determine an estimator $\hat{\theta}$ for θ .
3. In one sentence each, explain why $\hat{\theta}(X_1, \dots, X_n)$ is a random variable while $\hat{\theta}(x_1, \dots, x_n)$ is a number.

Exercise 1.2. Let $X_1, \dots, X_n \sim N(\mu, \sigma^2)$ be i.i.d., where both the mean μ and the variance σ^2 are unknown, so that the parameter is $\theta = (\mu, \sigma^2)$. Which of the following quantities are statistics?

$$(a) \max_{1 \leq i \leq n} X_i - \min_{1 \leq i \leq n} X_i, \quad (b) \frac{1}{n} \sum_{i=1}^n (X_i - \mu)^2, \quad (c) \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2.$$

Exercise 1.3. Let $X_1, \dots, X_n \sim \text{Uniform}(0, \theta)$ be i.i.d., as in example 1.2.

1. Write down the joint density $f(x_1, \dots, x_n; \theta)$, and state the values of the data for which it is non-zero.
2. Show that the estimator $\hat{\theta} = \max_i X_i$ always satisfies $\hat{\theta} \leq \theta$.
3. Show that in fact $\hat{\theta} < \theta$ with probability one, and deduce that $\mathbb{E}_\theta(\hat{\theta}) < \theta$, so that on average the estimator underestimates θ .

Lecture 2

Bias, Mean Squared Error and Consistency

We finished the previous lecture with the idea that a good estimator is one which has a sampling distribution concentrated around the true value θ_0 , for all possible true values. In this lecture we quantify this idea: the bias measures whether the estimator is on average right, the mean squared error measures how far away the estimator typically is from the truth, and consistency is the property that the estimator converges to the truth as the amount of data increases. The main idea to take away is the decomposition of the mean squared error into a variance and a squared bias.

2.1 Bias and Unbiasedness

The most basic question we can ask about an estimator is whether it gives the correct answer on average.

Definition 2.1. The *bias* of an estimator $\hat{\theta}$ for θ is given by

$$\text{bias}(\hat{\theta}) = \mathbb{E}_{\theta}(\hat{\theta}) - \theta.$$

The estimator is *unbiased*, if $\hat{\theta}$ has finite mean and $\text{bias}(\hat{\theta}) = 0$ for all $\theta \in \Theta$, *i.e.* if $\mathbb{E}_{\theta}(\hat{\theta}) = \theta$ for all θ .

The bias in general depends on the true value θ , as shown by the expectation with θ as a subscript.

The workhorse example is the sample mean. Since we will use the first two moments frequently, we state them here for reference.

Lemma 2.2. Let $X_1, \dots, X_n$ be *i.i.d.* with $\mathbb{E}(X_i) = \mu$ and $\text{Var}(X_i) = \sigma^2$. Define $\bar{X} = (X_1 + \dots + X_n)/n$. Then

$$\mathbb{E}(\bar{X}) = \mu \quad \text{and} \quad \text{Var}(\bar{X}) = \frac{\sigma^2}{n}.$$

Proof. Using the linearity of expectation we find

$$\mathbb{E}(\bar{X}) = \frac{1}{n} \sum_{i=1}^n \mathbb{E}(X_i) = \frac{1}{n} \cdot n\mu = \mu.$$

Since the X_i are independent, the variance of the sum can be found as the sum of the variances:

$$\text{Var}(\bar{X}) = \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) = \frac{1}{n^2} \cdot n\sigma^2 = \frac{\sigma^2}{n}.$$

This completes the proof. $\square$

Whenever the parameter of interest is the mean of the observations, as for the Bernoulli and Poisson models from table 1.1, lemma 2.2 shows that $\mathbb{E}_\theta(\bar{X}) = \theta$ and thus the sample mean is unbiased.

A more delicate example, which we will consider again in section 2.3, is the problem of estimating a variance. If a sample is observed with unknown mean, then one can estimate the variance σ^2 using the *sample variance*

$$S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2.$$

The unexpected divisor $n-1$ instead of the seemingly more obvious n is required to make this estimator unbiased.

Lemma 2.3. *Let $n \geq 2$ and let $X_1, \dots, X_n$ be i.i.d. with $\text{Var}(X_i) = \sigma^2$. Then $\mathbb{E}(S^2) = \sigma^2$ and thus the sample variance is an unbiased estimator for σ^2 .*

Proof. Let $\mu = \mathbb{E}(X_i)$. We prove the algebraic identity

$$\sum_{i=1}^n (X_i - \bar{X})^2 = \sum_{i=1}^n (X_i - \mu)^2 - n(\bar{X} - \mu)^2, \quad (2.1)$$

which holds for all numbers. Writing $X_i - \bar{X} = (X_i - \mu) - (\bar{X} - \mu)$, we can expand the square and use the relation $\sum_i (X_i - \mu) = n(\bar{X} - \mu)$ to simplify the expression of the cross term. Taking expectations on (2.1) and using the relations $\mathbb{E}((X_i - \mu)^2) = \sigma^2$ as well as $\mathbb{E}((\bar{X} - \mu)^2) = \text{Var}(\bar{X}) = \sigma^2/n$ from lemma 2.2 we find

$$\mathbb{E}\left(\sum_{i=1}^n (X_i - \bar{X})^2\right) = n\sigma^2 - n \cdot \frac{\sigma^2}{n} = (n-1)\sigma^2.$$

Dividing by $n-1$ gives the result $\mathbb{E}(S^2) = \sigma^2$, as claimed. $\square$

The calculation also shows what is wrong with the divisor n : the estimator $(1/n) \sum_i (X_i - \bar{X})^2$ has expectation $(n-1)\sigma^2/n$ and thus systematically underestimates the variance, since the observations are on average slightly closer to the sample mean $\bar{X}$ than the true mean μ would be. We will return to this comparison, and to the surprising fact that the biased estimator is in some sense better, in section 2.3.

The bias of this estimator is small if the sample is large. A property which is weaker than unbiasedness, but which holds for large sample sizes, is that the bias goes to zero in the limit.

Definition 2.4. Let $(\hat{\theta}_n)_{n \in \mathbb{N}}$ be a family of estimators for θ , depending on the sample size n . The estimator $\hat{\theta}_n$ is *asymptotically unbiased*, if

$$\lim_{n \rightarrow \infty} \text{bias}(\hat{\theta}_n) = 0$$

for all $\theta \in \Theta$.

Example 2.5. Let $X_1, \dots, X_n$ be i.i.d. with $\text{Var}(X_i) = \sigma^2$ and consider the family of estimators

$$\hat{\sigma}_n^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$$

for $n \geq 2$. Since $\hat{\sigma}_n^2 = (n-1)S^2/n$, we can use lemma 2.3 to conclude that $\mathbb{E}(\hat{\sigma}_n^2) = (n-1)\sigma^2/n$ and thus

$$\text{bias}(\hat{\sigma}_n^2) = \frac{n-1}{n}\sigma^2 - \sigma^2 = -\frac{\sigma^2}{n}.$$

While the bias for every n is non-zero, it converges to 0 as $n \rightarrow \infty$. Thus, $\hat{\sigma}_n^2$ is a biased but asymptotically unbiased estimator.

2.2 Mean Squared Error

If an estimator is unbiased, this does not mean that it is accurate. The mean squared error is a numerical measure of the typical size of the error $\hat{\theta} - \theta$.

Definition 2.6. The *mean squared error* (MSE) of an estimator $\hat{\theta}$ for θ is given by

$$\text{MSE}(\hat{\theta}) = \mathbb{E}_{\theta}((\hat{\theta} - \theta)^2).$$

We want the mean squared error to be small. It combines two different kinds of defects, a systematic offset and random scatter, and the following decomposition separates the two.

Theorem 2.7. For any estimator $\hat{\theta}$ for θ with finite variance,

$$\text{MSE}(\hat{\theta}) = \text{Var}(\hat{\theta}) + \text{bias}(\hat{\theta})^2.$$

Proof. Let $m = \mathbb{E}_{\theta}(\hat{\theta})$. Then we can write $\text{bias}(\hat{\theta}) = m - \theta$. Adding and subtracting m inside the square we get

$$\begin{aligned} \text{MSE}(\hat{\theta}) &= \mathbb{E}_{\theta}(((\hat{\theta} - m) + (m - \theta))^2) \\ &= \mathbb{E}_{\theta}((\hat{\theta} - m)^2) + 2(m - \theta) \mathbb{E}_{\theta}(\hat{\theta} - m) + (m - \theta)^2. \end{aligned}$$

Since $\mathbb{E}_{\theta}(\hat{\theta} - m) = m - m = 0$, the middle term disappears. The first term is $\text{Var}(\hat{\theta})$ and the last term is $\text{bias}(\hat{\theta})^2$. This is the required identity and the proof is complete. $\square$

The decomposition allows us to understand the rôle of unbiasedness: for an unbiased estimator, the bias term disappears and we have $\text{MSE}(\hat{\theta}) = \text{Var}(\hat{\theta})$, and thus, when comparing two unbiased estimators, we can only consider the variances. This is the basis of all optimality theory: among the unbiased estimators, the best one is the one with the smallest variance. In lecture 13 we will see that there is a lower bound for the variance of an unbiased estimator.

Using theorem 2.7 and lemma 2.2, we find that the mean squared error of the sample mean is given by $\text{MSE}(\bar{X}) = \sigma^2/n$. For Bernoulli samples this gives $\text{MSE}(\bar{X}) = \theta(1-\theta)/n$ and for Poisson samples $\text{MSE}(\bar{X}) = \theta/n$. In both cases the mean squared error decreases as $1/n$: more data give a proportionally smaller error.

2.3 When Unbiasedness Is Not Enough

Unbiasedness is an appealing property, but it is a constraint, not a goal, and insisting on it can cost us accuracy. We will show that for the variance of a normal sample the unbiased sample variance is beaten, in terms of mean squared error, by a biased competitor.

Let $X_1, \dots, X_n \sim N(\mu, \sigma^2)$ be i.i.d. and consider the family of estimators

$$\hat{\sigma}_c^2 = c \sum_{i=1}^n (X_i - \bar{X})^2,$$

for every constant $c > 0$. For $c = 1/(n-1)$ this is the unbiased sample variance S^2 and for $c = 1/n$ it is the underestimating version from the previous section. To compare these estimators, we need to know the distribution of the sum of squares. From proposition A.3 in appendix A we know that the scaled sum of squares satisfies

$$\frac{1}{\sigma^2} \sum_{i=1}^n (X_i - \bar{X})^2 \sim \chi_{n-1}^2.$$

If we write $W = \sum_i (X_i - \bar{X})^2$ and recall from definition A.2 that χ_{n-1}^2 has mean $n-1$ and variance $2(n-1)$, we get

$$\mathbb{E}(W) = (n-1)\sigma^2 \quad \text{and} \quad \text{Var}(W) = 2(n-1)\sigma^4.$$

Thus, the estimator $\hat{\sigma}_c^2 = cW$ has bias $\text{bias}(\hat{\sigma}_c^2) = (c(n-1) - 1)\sigma^2$ and, using theorem 2.7, we get

$$\text{MSE}(\hat{\sigma}_c^2) = c^2 \text{Var}(W) + \text{bias}(\hat{\sigma}_c^2)^2 = \left(2(n-1)c^2 + (c(n-1) - 1)^2\right)\sigma^4. \quad (2.2)$$

If we choose $c = 1/(n-1)$, the bias term disappears and we get

$$\text{MSE}(S^2) = \frac{2\sigma^4}{n-1}.$$

If we choose $c = 1/n$, then simplification of (2.2) gives

$$\text{MSE}(\hat{\sigma}_{1/n}^2) = \frac{2n-1}{n^2} \sigma^4.$$

The biased estimator wins, since $(2n-1)/n^2 < 2/(n-1)$. A little algebra shows that this inequality is equivalent to $(2n-1)(n-1) < 2n^2$, i.e. to $2n^2 - 3n + 1 < 2n^2$, which holds for all $n \geq 2$.

Remark. The value $c = 1/n$ is not optimal. The mean squared error (2.2) is a quadratic function of c and minimising this function over all $c > 0$ gives $c = 1/(n+1)$. For this value we get

$$\text{MSE}(\hat{\sigma}_{1/(n+1)}^2) = \frac{2\sigma^4}{n+1}.$$

Thus, neither the unbiased estimator nor the maximum-likelihood estimator $c = 1/n$ (which we will see in lecture 5) is optimal, but a third, different estimator exists.

By allowing a small, controlled bias, we have been able to reduce the variance and thus have strictly improved the mean squared error. The lesson to be learned here is not that bias is good, but that an estimator should be evaluated using its mean squared error instead of just the bias.

2.4 Consistency

The mean squared errors we have computed all shrink as the sample grows. Consistency is the property which describes this limiting behaviour: as the number of observations increases, a consistent estimator converges to the true value.

Definition 2.8. Consider a family $(\hat{\theta}_n)_{n \in \mathbb{N}}$ of estimators for θ , one for each sample size n . The estimator $\hat{\theta}_n$ is said to be *consistent* for θ , if $\hat{\theta}_n \xrightarrow{P} \theta$ as $n \rightarrow \infty$, i.e. if for every $\varepsilon > 0$ we have

$$\lim_{n \rightarrow \infty} \mathbb{P}_\theta(|\hat{\theta}_n - \theta| > \varepsilon) = 0,$$

for all $\theta \in \Theta$.

Here we write $\xrightarrow{P}$ for convergence in probability, as recalled in definition A.12 in appendix A: the probability of the estimate being off the truth by any fixed amount eventually becomes negligible. The most basic example of a consistent estimator is the sample mean, where consistency is a manifestation of the law of large numbers.

Theorem 2.9. Let $X_1, \dots, X_n$ be i.i.d. with $\mathbb{E}_\theta(X_i) = \theta$ and finite variance σ^2 . Then the sample mean $\bar{X}$ is a consistent estimator for θ .

Proof. Using Chebyshev's inequality, lemma A.6 and lemma 2.2, for any $\varepsilon > 0$ we have

$$\mathbb{P}_\theta(|\bar{X} - \theta| > \varepsilon) \leq \frac{\text{Var}(\bar{X})}{\varepsilon^2} = \frac{\sigma^2}{\varepsilon^2 n} \longrightarrow 0 \quad (n \rightarrow \infty).$$

Thus $\bar{X} \xrightarrow{P} \theta$ and the estimator is consistent. $\square$

Remark. The theorem requires the variance to be finite, as used in the Chebyshev argument; it is sufficient to have the mean finite, by the law of large numbers, theorem A.14 in appendix A. Some condition like this is needed for the result to hold. The standard counterexample is the Cauchy distribution, with density $1/(\pi(1+x^2))$ on $\mathbb{R}$ (see table A.2). This distribution has heavy tails and thus $\mathbb{E}(|X_1|)$ is infinite, i.e. the distribution has no mean. Similarly, the sample mean of a Cauchy distribution is itself Cauchy-distributed, with the same distribution as a single observation for every n (without proof). Thus, the sample mean does not concentrate as the sample size increases and hence is not a consistent estimator for the centre of the distribution. Consistency of an estimator is a property of the estimator and the model combined, and does not only depend on the form of the formula.

The argument used almost nothing about the sample mean except that its mean squared error converges to zero, and thus the same argument applies to any estimator with this property.

Theorem 2.10. If $\text{MSE}(\hat{\theta}_n) \rightarrow 0$ as $n \rightarrow \infty$, then $\hat{\theta}_n$ is consistent.

Proof. Using Markov's inequality, lemma A.5, for the non-negative random variable $(\hat{\theta}_n - \theta)^2$, for any $\varepsilon > 0$ we get

$$\mathbb{P}_\theta(|\hat{\theta}_n - \theta| > \varepsilon) = \mathbb{P}_\theta((\hat{\theta}_n - \theta)^2 > \varepsilon^2) \leq \frac{\mathbb{E}_\theta((\hat{\theta}_n - \theta)^2)}{\varepsilon^2} = \frac{\text{MSE}(\hat{\theta}_n)}{\varepsilon^2} \longrightarrow 0$$

as $n \rightarrow \infty$. Thus $\hat{\theta}_n \xrightarrow{P} \theta$ as required. $\square$

Combining this criterion with the bias-variance decomposition gives a sufficient condition in terms of the two quantities we know how to compute.

Corollary 2.11. *If $\hat{\theta}_n$ is asymptotically unbiased and $\text{Var}(\hat{\theta}_n) \rightarrow 0$ as $n \rightarrow \infty$, then $\hat{\theta}_n$ is consistent.*

Proof. From theorem 2.7 we know that $\text{MSE}(\hat{\theta}_n) = \text{Var}(\hat{\theta}_n) + \text{bias}(\hat{\theta}_n)^2$ and since both terms go to zero by assumption, we can conclude from theorem 2.10 that the result holds. $\square$

All estimators in this lecture where we have computed a mean squared error have $\text{MSE}(\hat{\theta}_n)$ of order $1/n$ and thus are consistent. Consistency is a coarse property: two consistent estimators can have very different rates of convergence and the mean squared error is a more detailed measure. We will see numerical evidence of consistency in the R session of lecture 9, and in lecture 16 we will see a proof that the maximum likelihood estimator is consistent for a wide class of cases.

Summary.

- The bias $\text{bias}(\hat{\theta}) = \mathbb{E}_\theta(\hat{\theta}) - \theta$ measures whether an estimator is right on average; an unbiased estimator has zero bias for all θ .
- The mean squared error $\text{MSE}(\hat{\theta}) = \mathbb{E}_\theta((\hat{\theta} - \theta)^2)$ is a measure for the typical squared error, given by $\text{MSE}(\hat{\theta}) = \text{Var}(\hat{\theta}) + \text{bias}(\hat{\theta})^2$.
- Unbiasedness is a constraint, not an aim: a biased estimator can have smaller mean squared error, as shown by the normal variance.
- An estimator is consistent, if $\hat{\theta}_n$ converges in probability to θ . A sufficient condition for this is $\text{MSE}(\hat{\theta}_n) \rightarrow 0$, i.e. being asymptotically unbiased with vanishing variance.

Exercise 2.1. Let $X_1, \dots, X_n$ be an i.i.d. sample from the Poisson distribution with parameter θ , i.e. $\mathbb{E}_\theta(X_i) = \text{Var}_\theta(X_i) = \theta$ and let $\hat{\theta} = \bar{X}$.

1. Show that $\hat{\theta}$ is an unbiased estimator for θ .
2. Determine the mean squared error of $\hat{\theta}$ and show that $\hat{\theta}$ is a consistent estimator.

Exercise 2.2. Let $X_1, \dots, X_n \sim \text{Uniform}(0, \theta)$ be i.i.d., as in example 1.2, and let $\hat{\theta} = \max_i X_i$. You may use that $\hat{\theta}$ has density $f(t) = n t^{n-1} / \theta^n$ for $0 \leq t \leq \theta$.

1. Show that $\mathbb{E}_\theta(\hat{\theta}) = n\theta/(n+1)$, and hence that the bias is $\text{bias}(\hat{\theta}) = -\theta/(n+1)$. Is $\hat{\theta}$ asymptotically unbiased?
2. Show that $\text{MSE}(\hat{\theta}) = 2\theta^2/((n+1)(n+2))$, and deduce that $\hat{\theta}$ is consistent.
3. Compare the rate at which this mean squared error shrinks with the $1/n$ rate found for the sample mean, and comment.

Exercise 2.3. Let $X_1, \dots, X_n$ be an i.i.d. sample from the exponential distribution with parameter θ , i.e. $f(x; \theta) = \theta e^{-\theta x}$ for $x \geq 0$ and $\mathbb{E}_\theta(X_i) = 1/\theta$. If we define an estimator by $\hat{\theta} = 1/\bar{X}$, and if we use the fact that $T = X_1 + \dots + X_n$ has the gamma density $f_T(x) = \theta^n x^{n-1} e^{-\theta x} / (n-1)!$ for $x \geq 0$, show that, for $n \geq 2$,

$$\mathbb{E}_\theta(\hat{\theta}) = \frac{n\theta}{n-1}.$$

From this result, find the bias of $\hat{\theta}$ and show that $\hat{\theta}$ is biased but asymptotically unbiased.

Lecture 3

R Basics and a First Simulation

This is the first of the R sessions, which illustrate the theory on the computer and build up the R skills needed for the practical report. The lecture has two aims: to get R up and running on the computer, and to learn how to simulate a random sample. The second part of the lecture uses simulation to illustrate the ideas from lectures 1 and 2: an estimator is a random quantity, and the bias and spread of an estimator can be seen by simulation. Appendix B covers the same ground at a slower pace and in more detail. This is the written back-stop for the R sessions, and is the place to look up R commands later. The commands we use in the lecture are shown in a framed box, and the R output is shown below, with each line of the output marked by `##`.

3.1 First Steps in R

Appendix B explains how to install R or how to use R on the university computers. Either way, we will end up at the *console* where we type in commands and R prints the results.

The commands we need to use today are as follows: To construct a list of numbers, called a *vector*, we can use `c`. To store a list of numbers under a name, we use the assignment arrow `<-`. Finally, we can use the built-in functions `mean`, `median` and `var` to summarise a vector.

```
x <- c(4, 7, 1, 5, 3)
x
mean(x)
median(x)
var(x)
```

```
## [1] 4 7 1 5 3
## [1] 4
## [1] 4
## [1] 5
```

The `[1]` at the start of each output line is a position marker, not part of the answer: it gives the index of the first value printed on that line. The assignment itself prints nothing; typing `x` on its own shows the stored values.

The function `var` does not use a general convention, but computes the sample variance S^2 from lecture 2, using the divisor $n - 1$. To verify this, we can compute the formula manually, using the fact that arithmetic in R operates on the individual elements of a vector, and that `length` can be used to determine the number of elements.

```
sum((x - mean(x))^2) / (length(x) - 1)
```

```
## [1] 5
```

The answer coincides with the result of `var(x)` above, and we now know that R uses the unbiased estimator from lemma 2.3 for variances. Further details about vectors, indexing and arithmetic can be found in appendix B. The appendix also explains how to understand the help page of an unfamiliar command; we will not repeat this information here.

3.2 Simulating a Normal Sample

The theory from the last two lectures considers data generated from a model with known true parameter. In this section we will see that we can simulate such data on a computer. The command `rnorm` can be used to generate a random sample from a normal distribution. The arguments of the command are the sample size and the values for the `mean` and `sd`.

Warning. R uses the standard deviation as a parameter of the normal distribution, not the variance. Thus, if we write `sd = 2`, we mean $\sigma = 2$ and thus $\sigma^2 = 4$. A common mistake is to forget this and to get simulations which are mysteriously too spread out.

Since random draws are different every time, it is often a nuisance to check our work and in a report it would be a disaster if the random numbers changed every time the report was printed. The solution to this problem is the `set.seed` command: if we call this command with a fixed number, the random-number generator will be set to a fixed state and the same seed will always return the same random numbers.

```
set.seed(201234567)
rnorm(3)
set.seed(201234567)
rnorm(3)
```

```
## [1] -1.0333701 -0.7724437  0.1556372
## [1] -1.0333701 -0.7724437  0.1556372
```

Any fixed number can be used as the seed. In the practical report, it is a good idea to use a personal number, like your student ID, to seed the random number generator. This way, your report is exactly reproducible, while your dataset is still different from the datasets of your classmates.

Finally, we can produce a proper sample: $n = 100$ observations from the normal distribution with mean $\mu = 10$ and variance $\sigma^2 = 4$. The command `summary` can be used to print the standard descriptive statistics for the sample, and `hist` can be used to produce a histogram. Finally, `abline` can be used to

add a straight line to an existing plot (in this case, a dashed vertical line at the sample mean).

```
set.seed(201234567)
sample <- rnorm(100, mean = 10, sd = 2)
summary(sample)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    4.710   9.045  10.166  10.006  11.350  13.877
```

The summary already tells a plausible story: the sample mean 10.006 is close to the true mean 10, and the quartiles are roughly symmetric around the mean. To get a better picture of the sample, we consider the histogram in figure 3.1. We adjust the horizontal range using the `xlim` argument so that the axis can be compared to the axis in a second figure, below.

```
hist(sample, col = "grey85", main = "", xlab = "x", xlim = c(4, 16))
abline(v = mean(sample), lty = 2)
```




Figure 3.1: A simulated sample of $n = 100$ observations from a normal distribution with mean $\mu = 10$ and variance $\sigma^2 = 4$. The dashed line gives the sample mean.

3.3 The Sample Mean as a Random Quantity

In lecture 1 we have discussed the difference between an estimate, a single number computed from data, and an estimator, which is a random variable. We can use simulation to visualise this difference: every call to `rnorm` below generates a fresh sample of size $n = 25$ from the same model as before, and the sample mean of each sample is a fresh realisation of the estimator $\bar{X}$. Before you execute the following three lines of R, make a prediction: will the three printed numbers be equal? If not, how much do you expect the three numbers to be away from the true value 10?

```
mu <- 10
n <- 25
set.seed(2026)
mean(rnorm(n, mean = mu, sd = 2))
mean(rnorm(n, mean = mu, sd = 2))
mean(rnorm(n, mean = mu, sd = 2))
```

```
## [1] 9.340722
## [1] 10.5942
## [1] 9.335904
```

Three samples, three different estimates, scattered around the true value 10. The scatter we see here is the sampling distribution, introduced in lecture 1. By repeating the experiment many times, we can map out this distribution. The `replicate` command can be used to repeat a command a given number of times, and to store the results in a vector. Loops and functions can be used to achieve the same effect (see appendix B). We repeat the experiment 10000 times, and compare the results to the predictions from lecture 2: From lemma 2.2 we know that the estimator $\bar{X}$ is unbiased, with standard deviation $\sigma/\sqrt{n} = 2/\sqrt{25} = 0.4$. From theorem 2.7 we know that the mean squared error of the estimator equals the variance $\sigma^2/n = 4/25 = 0.16$.

```
estimates <- replicate(10000, mean(rnorm(n, mean = mu, sd = 2)))
mean(estimates) - mu          # empirical bias
sd(estimates)                 # empirical spread
mean((estimates - mu)^2)      # empirical mean squared error
```

```
## [1] 0.004472127
## [1] 0.4006136
## [1] 0.1604952
```

The agreement is striking: the empirical bias is close to zero, the empirical spread 0.4006 matches the theoretical value 0.4, and the empirical mean squared error 0.1605 matches the theoretical value 0.16. This shows that the theory from lecture 2 is not an abstraction, but accurately predicts the output of the machine with two decimal digits of accuracy.

Finally we draw the histogram of all 10000 estimates, on the same horizontal axis as figure 3.1 so that the two plots can be compared directly. Before looking at the result, sketch the histogram you expect: where should it be centred, and how wide should it be compared with the data histogram in figure 3.1?

```
hist(estimates, breaks = 40, col = "grey85", main = "",
      xlab = "estimate", xlim = c(4, 16))
abline(v = mu, lty = 2)
```

While the data in figure 3.1 are spread out over most of the axis, the simulated estimates form a narrow spike around the truth. This is a consequence of the fact that, by averaging, the estimator has exchanged the spread $\sigma = 2$ of an individual observation for the much smaller spread 0.4 of the estimator. The picture of a histogram of an estimator which is close to the truth is the one we had in mind in lecture 2, when we discussed small bias and small mean squared




Figure 3.2: The sampling distribution of the sample mean $\bar{X}$ for $n = 25$: a histogram of 10000 simulated estimates, where the true mean $\mu = 10$ is marked by the dashed line. The axis is the same as in figure 3.1.

error. In the next R session, lecture 6, we will use the same simulation setup to compare different estimators. In lecture 9 we will watch the spike get narrower as n increases. This is an example of consistency.

Before the next session, please type in the commands of this session and try to understand the results.

Summary.

- The written reference for all R commands used in this session is appendix B.
- The command `rnorm` can be used to simulate samples from a normal distribution. The standard deviation of the distribution can be controlled using the `sd` argument of the command. The commands `mean`, `median`, `var`, `summary` and `hist` can be used to summarise and display the simulated samples.
- If `set.seed` is called with a fixed number, *e.g.* the student ID, then the output of all subsequent simulation commands is exactly reproducible.
- If an estimate is re-computed using many simulated samples, a histogram of the resulting sampling distribution can be plotted. The bias, spread and mean squared error of the sampling distribution can be computed and shown to agree with the theoretically predicted values from lecture 2.

Lecture 4

Likelihood and the Method of Moments

In the previous two lectures we have introduced the problem of estimation, but all the estimators we have considered so far were just guesses. In this lecture we introduce the *likelihood function*, the most important object in this module: the concepts of sufficiency (in lecture 8), Fisher information (in lecture 11), maximum likelihood (in lecture 5) and the Bayesian approach (in lecture 26) are all based on this function. We also discuss the *method of moments*, a simple method for constructing an estimator from a model. We will see the weaknesses of this approach which will lead to the introduction of the maximum likelihood estimator in the next lecture.

4.1 The Likelihood Function

From lecture 1 we know that a parametric model specifies a joint density for the data, for each $\theta \in \Theta$. For an i.i.d. sample this is the product of the marginal densities. The idea of this section is to consider the converse problem, *i.e.* to consider, once the data $x_1, \dots, x_n$ are fixed, how well a given value of θ explains the data.

Definition 4.1. Let $x_1, \dots, x_n$ be data from a parametric model $(f(x; \theta) \mid \theta \in \Theta)$. Then the *likelihood function* is the function $L: \Theta \rightarrow [0, \infty)$, given by

$$L(\theta) = \prod_{i=1}^n f(x_i; \theta),$$

and the *log-likelihood* is $\ell(\theta) = \log L(\theta)$, given by

$$\ell(\theta) = \sum_{i=1}^n \log f(x_i; \theta).$$

While the formula for L is the same as the joint density of the sample, the interpretation of this function has changed: while the joint density is a function of the data with the parameter fixed, the likelihood is a function of the parameter with the data fixed. To emphasise that the data are used to compute the likelihood, we write $L(\theta; x_1, \dots, x_n)$, and if we substitute the random variables

X_i for the numbers x_i , the value of $L(\theta)$ is a random quantity. Both the numerical values and the random quantities can be used freely. For discrete data, $L(\theta)$ equals $\mathbb{P}_\theta(X_1 = x_1, \dots, X_n = x_n)$, *i.e.* the probability that the model with parameter θ generates the observed data.

Example 4.2. Assume we have observed 7 heads in $n = 10$ tosses of a coin with unknown heads probability θ . Using the Bernoulli weights from table 1.1, we find the likelihood to be

$$L(\theta) = \prod_{i=1}^{10} \theta^{x_i} (1 - \theta)^{1-x_i} = \theta^7 (1 - \theta)^3, \quad \theta \in (0, 1).$$

For the fair coin case we have $L(0.5) = 0.5^{10} \approx 0.00098$ and $L(0.7) = 0.7^7 \cdot 0.3^3 \approx 0.00222$. Both values are very small, since any sequence of ten tosses is unlikely, but the ratio of these values is more informative: $\theta = 0.7$ is 2.3 times better at explaining the data than the fair coin model is, and this is the kind of information we are interested in.

Warning. The likelihood is not a probability distribution over θ : the parameter is an unknown constant, statements like “the probability that $\theta = 0.7$ is small” make no sense, and L does not integrate to one over Θ . We will learn more about this in lecture 26, where we will see how the Bayesian approach can be used to turn the parameter into a random variable.

In practice we almost always use the log-likelihood: the likelihood is a product of n factors, typically each less than one, so for realistic sample sizes it is astronomically small. In contrast, ℓ is a sum of moderate size and it is much easier to take derivatives of a sum than of a product. Since the logarithm is strictly increasing, the functions L and ℓ order the parameter values in the same way, so no information is lost when we use the logarithm.

Example 4.3. For an i.i.d. sample $x_1, \dots, x_n$ from the exponential model with density $f(x; \theta) = \theta e^{-\theta x}$ for $x \geq 0$, the likelihood is

$$L(\theta) = \prod_{i=1}^n \theta e^{-\theta x_i} = \theta^n \exp\left(-\theta \sum_{i=1}^n x_i\right).$$

Thus, the log-likelihood is

$$\ell(\theta) = n \log \theta - \theta \sum_{i=1}^n x_i, \quad \theta > 0.$$

The product in this expression has been replaced by a smooth function of θ , which uses only two summaries of the data, n and $\sum_i x_i$. We will learn more about this approach to summarising data in lecture 8.

4.2 The Score Function

If the log-likelihood function is differentiable, then the derivative is sometimes called the *score* function. In this section we only use the score function to find the maximum of the likelihood. (We will learn the deeper meaning of the score function, as the carrier of information in the sample, in lecture 11.)

Definition 4.4. Let $\ell(\theta)$ be the log-likelihood function and assume that it is differentiable at θ . Then the *score function* is given by

$$\ell'(\theta) = \frac{d}{d\theta} \ell(\theta) = \sum_{i=1}^n \frac{\partial}{\partial \theta} \log f(x_i; \theta).$$

The score is the sum of n terms, one for each observation, and if we assume that the data are random, these terms are i.i.d. This additive structure allows us to apply the law of large numbers and the central limit theorem to the likelihood function later in the module. Since the score is the derivative of the log-likelihood, a zero of the score is a good starting point to find the parameter value which gives the best explanation for the data.

Example 4.5. Continuing from example 4.3, the score for the exponential sample is given by

$$\ell'(\theta) = \frac{n}{\theta} - \sum_{i=1}^n x_i, \quad \theta > 0.$$

The score is positive for $\theta < 1/\bar{x}$ and negative for $\theta > 1/\bar{x}$, where $\bar{x} = (1/n) \sum_i x_i$. Thus the log-likelihood increases up to $1/\bar{x}$ and then decreases, so $\theta = 1/\bar{x}$ is the best explanation for the data. The idea of using the maximiser of the likelihood to estimate parameters is the basis of maximum likelihood estimation, as discussed in lecture 5.

4.3 The Method of Moments

Before we consider how to maximise the likelihood, we will discuss a much simpler method, the oldest general purpose method for constructing estimators. This method is based on the fact that the model often gives the moments of the observations as functions of the model parameters θ , and that the data can be used to compute the empirical moments. The idea of the method is then to find an estimator such that the two sets of moments coincide.

Let

$$\mu_j(\theta) = \mathbb{E}_\theta(X_1^j), \quad m_j = \frac{1}{n} \sum_{i=1}^n X_i^j$$

for $j = 1, 2, \dots$. Then $\mu_j(\theta)$ is the j -th *population moment*, a deterministic function of θ , and m_j is the corresponding *sample moment*, a statistic. By the law of large numbers (theorem A.14 in appendix A), we have $m_j \xrightarrow{P} \mu_j(\theta_0)$: for large sample size the sample moments are close to the population moments corresponding to the true parameter value.

Definition 4.6. Let $\theta = (\theta_1, \dots, \theta_k)$ be a parameter with k components. Then the *method of moments estimator* $\hat{\theta}$ is found by solving the system of k equations

$$\mu_j(\hat{\theta}) = m_j, \quad j = 1, \dots, k,$$

for $\hat{\theta} \in \Theta$, if this system of equations has a unique solution.

The recipe uses as many moment equations as there are unknown parameter components, and the required moments can usually be found in tables like tables A.1 and A.2. We illustrate the method for three of our running examples.

Example 4.7. For the exponential model with rate θ , we have $\mu_1(\theta) = 1/\theta$ and the only moment equation is $1/\hat{\theta} = \bar{X}$. Thus, the method of moments estimator is given by

$$\hat{\theta} = \frac{1}{\bar{X}}.$$

Coincidentally, this is the same as the zero of the score from example 4.5.

Example 4.8. For the normal model with $\theta = (\mu, \sigma^2)$ the parameter has two components. Thus we need to match the first two moments, $\mu_1(\theta) = \mu$ and $\mu_2(\theta) = \sigma^2 + \mu^2$. We can solve these two moment equations (see exercise 4.3) to find

$$\hat{\mu} = \bar{X} \quad \text{and} \quad \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2.$$

Note the divisor n : since we are matching moments without considering the bias, the method gives the biased plug-in variance estimator from lecture 2, instead of the unbiased sample variance S^2 . In this example, from the mean squared error comparison in lecture 2, this is not a problem.

Example 4.9. Let $X_1, \dots, X_n \sim \text{Gamma}(\alpha, \beta)$ be i.i.d. where the shape α and the rate β are unknown. Then $\theta = (\alpha, \beta)$. From table A.2 we know $\mathbb{E}_\theta(X_1) = \alpha/\beta$ and $\text{Var}_\theta(X_1) = \alpha/\beta^2$. Instead of matching μ_1 and μ_2 we can equivalently match the mean and the variance: if we write $V = m_2 - m_1^2 = (1/n) \sum_i (X_i - \bar{X})^2$ for the plug-in variance, we get the equations

$$\frac{\hat{\alpha}}{\hat{\beta}} = \bar{X} \quad \text{and} \quad \frac{\hat{\alpha}}{\hat{\beta}^2} = V.$$

Dividing the first equation by the second we find $\hat{\beta} = \bar{X}/V$ and thus

$$\hat{\alpha} = \frac{\bar{X}^2}{V} \quad \text{and} \quad \hat{\beta} = \frac{\bar{X}}{V}.$$

This example shows the method in a good light: for the gamma distribution the maximum likelihood equations from lecture 5 have no closed-form solution, whereas the method of moments finds explicit formulas in two lines.

We now weigh the merits of the method. On the credit side, it is quick, often gives explicit formulas where the likelihood does not, and is consistent under weak conditions: the sample moments converge by the law of large numbers, and if the solution of the moment equations depends continuously on the moments, the continuous mapping theorem (theorem A.17) carries the convergence over to the estimator.

The method has three main weaknesses. First, the estimates can be contradicted by the data: for the uniform model $X_1, \dots, X_n \sim \text{Uniform}(0, \theta)$ we have $\mu_1(\theta) = \theta/2$ and thus $\hat{\theta} = 2\bar{X}$, but for the dataset $x_1 = 0.1, x_2 = 0.2, x_3 = 0.9$ this gives $\hat{\theta} = 2 \cdot 0.4 = 0.8$, an estimate where the observation $x_3 = 0.9$ could never have occurred. This problem is immediately detected by the likelihood, since $L(0.8) = 0$, but the moment equations do not detect this problem. Secondly, the choice of moments is arbitrary: if we had used the second moment instead of the first, we would in general have obtained a different estimator. Finally, by taking only a few moments, much of the information in the data is lost. A method which makes use of all the information in the data should be more successful. This is the aim of the next few lectures.

Remark. In some cases, both approaches coincide: for the exponential model the moment estimator $1/\bar{X}$ is also the maximiser of the likelihood (from example 4.5), the same is true for the Poisson distribution (exercise 4.1) and for the normal distribution the two methods coincide (lecture 5). In contrast, for the uniform model considered above, the maximum likelihood estimator is $\max_i X_i$ instead of $2\bar{X}$ and has a very different mean squared error (as shown in exercises 2.2 and 4.4).

4.4 Identifiability

Both recipes of this lecture, and the whole estimation programme, assume that the parameter can be recovered from the distribution of the data: if two parameter values lead to the same distribution, no data can distinguish between them and no good estimator for θ can exist. This property of a model deserves a name.

Definition 4.10. A parametric model $(\mathbb{P}_\theta \mid \theta \in \Theta)$ is *identifiable*, if different parameter values lead to different distributions, *i.e.* if we have $\mathbb{P}_{\theta_1} \neq \mathbb{P}_{\theta_2}$ for all $\theta_1, \theta_2 \in \Theta$ with $\theta_1 \neq \theta_2$.

All the families of models listed in tables 1.1 and 1.2 are identifiable. Non-identifiability usually occurs in models which are over-parametrised, as the following example shows.

Example 4.11. Assume that we want to model a measurement which is subject to two additive effects. We can assume that $X_1, \dots, X_n \sim N(\alpha + \beta, 1)$ are i.i.d. with parameter $\theta = (\alpha, \beta) \in \mathbb{R}^2$. Then the distribution of the data depends on θ only via the sum $\alpha + \beta$. Thus, the parameters $(1, 2)$ and $(0, 3)$ give rise to the same distribution and the model is not identifiable. Both recipes from this lecture fail in this case: the likelihood is constant on the lines $\alpha + \beta = \text{const}$ and the moment equations only contain the unknowns in the combination $\alpha + \beta$. To solve the problem we can either reparametrise the model using $\mu = \alpha + \beta$ or we can restrict the parameter space, for example by setting $\beta = 0$.

Identifiability is a question about the parametrisation, not about the data or the sample size. A more subtle version of the problem is the case of the mixture models from lecture 19, where the models are only identifiable up to relabelling the mixture components. For the rest of the module, we will assume that the models are identifiable, unless we explicitly state otherwise.

Summary.

- The likelihood $L(\theta) = \prod_i f(x_i; \theta)$ is the joint density of the observed data, as a function of θ with the data fixed, and is not a probability distribution over θ .
- The score $\ell'(\theta)$, the derivative of the log-likelihood function $\ell(\theta) = \sum_i \log f(x_i; \theta)$, finds the parameter value which best explains the data. The score returns in lecture 11.
- The method of moments solves the equations $\mu_j(\hat{\theta}) = m_j$ for $j = 1, \dots, k$. The method is fast, explicit and consistent, but it can return impossible estimates and does not make good use of the data. The maximum likelihood method from lecture 5 is better.
- A model is identifiable, if different parameter values lead to different distributions. If the model is not identifiable, no estimator can work and we need to change the parametrisation.

Exercise 4.1. Let $X_1, \dots, X_n$ be an i.i.d. sample from a Poisson distribution with parameter $\theta > 0$ and let $x_1, \dots, x_n$ be the observed data.

1. Write down the likelihood $L(\theta)$ and show that the log-likelihood is $\ell(\theta) = \log \theta \sum_i x_i - n\theta - \sum_i \log(x_i!)$.
2. Compute the score $\ell'(\theta)$ and find the value of θ where the score equals zero.
3. Compute the method of moments estimator for θ and compare your result to your answer in the previous part.

Exercise 4.2. Let $X_1, \dots, X_n$ be an i.i.d. sample from the geometric distribution with parameter $p \in (0, 1)$, *i.e.* from the distribution with $f(x; p) = (1 - p)^{x-1}p$ for $x \in \{1, 2, \dots\}$ and $\mathbb{E}_p(X_1) = 1/p$. (There is some disagreement among authors whether the geometric distribution counts the number of trials until the first success, inclusive of this trial, starting at 1 as in the present exercise, or the number of failures before the first success, starting at 0. In this module we always count trials.)

1. Using the method of moments, determine an estimator $\hat{p}$ for p .
2. Show that the value of $\hat{p}$ always lies in the interval $(0, 1]$, *i.e.* that the estimate can never be contradicted by the data (this is in contrast to the example of the uniform distribution in the text).

Exercise 4.3. Let $X_1, \dots, X_n \sim N(\mu, \sigma^2)$ i.i.d. with parameter $\theta = (\mu, \sigma^2)$, as in example 4.8.

1. Using the relation $\text{Var}(X_1) = \mathbb{E}(X_1^2) - \mathbb{E}(X_1)^2$, show that the first two population moments are $\mu_1(\theta) = \mu$ and $\mu_2(\theta) = \sigma^2 + \mu^2$.
2. Write down the two moment equations and solve them to verify the estimators $\hat{\mu} = \bar{X}$ and $\hat{\sigma}^2 = (1/n) \sum_i (X_i - \bar{X})^2$ from example 4.8.

Exercise 4.4. Let $X_1, \dots, X_n \sim \text{Uniform}(0, \theta)$ i.i.d. and let $\hat{\theta} = 2\bar{X}$ be the method of moments estimator as given in the text.

1. Show that $\hat{\theta}$ is unbiased and that $\text{MSE}(\hat{\theta}) = \theta^2/(3n)$.
2. In exercise 2.2 we have seen that $\text{MSE}(\max_i X_i) = 2\theta^2/((n+1)(n+2))$ for the biased estimator $\max_i X_i$. Show that $\max_i X_i$ has strictly smaller mean squared error for all $n \geq 3$ and comment on your result.

Exercise 4.5. Consider the model $X_1, \dots, X_n \sim N(0, \theta^2)$ i.i.d. for the parameter $\theta \in \Theta = \mathbb{R} \setminus \{0\}$.

1. Show that the model is not identifiable.
2. Find a smaller parameter space $\Theta' \subseteq \Theta$, where the model is identifiable.
3. On your space Θ' , determine the method of moments estimator for θ , using the second moment.

Lecture 5

Maximum Likelihood Estimation

In the previous lecture we have seen that for the exponential model, the parameter value which maximises the likelihood is the one which ‘best’ explains the data. In this lecture we will make this idea into the main estimation method of the module: we define the maximum likelihood estimator, practise computing it for the families from tables 1.1 and 1.2, and will see problems where more thought is needed. We will see that the estimator is not unbiased in general, but is invariant under reparametrisation; the deeper justification of the method will be given in lectures 13, 16 and 17.

5.1 The Maximum Likelihood Estimator

Example 4.5 already contains the idea: for the exponential sample the log-likelihood increases up to $\theta = 1/\bar{x}$ and then decreases, so this value fits the data best. We now formalise this principle.

Definition 5.1. A *maximum likelihood estimator* (MLE) for θ is an estimator $\hat{\theta} = \hat{\theta}(X_1, \dots, X_n)$, which maximises the likelihood function over the parameter space, *i.e.* which satisfies

$$L(\hat{\theta}) \geq L(\theta) \quad \text{for all } \theta \in \Theta.$$

Since the logarithm is strictly increasing, $\hat{\theta}$ maximises L if and only if it maximises the log-likelihood ℓ and for the reasons given in lecture 4 we mostly work with ℓ . The definition is careful to avoid the problem of a maximum not existing or not being unique. For example, the likelihood could increase as the edge of the parameter space is approached, so that no maximum exists. Similarly, the maximum may not be unique. In such cases, any maximiser can be used. For the standard families of models and typical data, neither of these problems will occur and it will be safe to speak of “the” MLE.

If the set Θ is an interval and the log-likelihood function is differentiable, then calculus can be used to find the MLE: at a maximum inside the interval Θ , the derivative of ℓ equals zero and thus the MLE satisfies the *likelihood equation*

$$\ell'(\theta) = 0,$$

i.e. the MLE is a zero of the score function from definition 4.4. The general procedure is to find the log-likelihood function, to solve the likelihood equation and to verify that the solution is a maximum, for example by checking that $\ell''(\theta) < 0$ or by checking that the score changes from positive to negative.

Warning. The solution of the likelihood equation only gives *interior* critical points. These could be minima or inflection points and the checking step is required to avoid wrong answers. Also, the above procedure assumes that the maximum is interior and that ℓ is differentiable. We will see that these assumptions can be violated.

5.2 Worked Examples

We start our analysis by considering the “running families” of the module.

Example 5.2. In example 4.5, the score of the exponential model equals zero at $\theta = 1/\bar{x}$. Since $\ell''(\theta) = -n/\theta^2 < 0$ for all $\theta > 0$, the function ℓ is strictly concave and this root is the only global maximum. Thus, the MLE for the rate of the exponential distribution is

$$\hat{\theta} = \frac{1}{\bar{X}}.$$

This coincides with the method of moments estimator from example 4.7.

Example 5.3. Let $X_1, \dots, X_n$ be an i.i.d. Bernoulli sample with success probability θ , and let $s = \sum_i x_i$ be the observed number of successes. As in example 4.2 we have the likelihood $L(\theta) = \theta^s(1 - \theta)^{n-s}$ and thus

$$\ell(\theta) = s \log \theta + (n - s) \log(1 - \theta), \quad \theta \in (0, 1).$$

Assume first $0 < s < n$. Then the likelihood equation is

$$\ell'(\theta) = \frac{s}{\theta} - \frac{n - s}{1 - \theta} = 0.$$

This equation implies $s(1 - \theta) = (n - s)\theta$ and the only solution to this equation is $\theta = s/n$. Since $\ell''(\theta) = -s/\theta^2 - (n - s)/(1 - \theta)^2 < 0$ everywhere, this root is the global maximum and the MLE for the proportion of successes is

$$\hat{\theta} = \frac{s}{n} = \bar{X}.$$

For the dataset from example 4.2, seven heads in ten tosses, we get $\hat{\theta} = 0.7$, the value which we already recognised in the example.

Remark. The assumption $0 < s < n$ is made, because at the two extremes the MLE does not exist: for $s = n$ the function $\ell(\theta) = n \log \theta$ is strictly increasing on $(0, 1)$ and the supremum is approached at the boundary point $\theta = 1$, which is outside the open parameter space, and thus there is no maximiser. This is a simple case of the existence caveat from definition 5.1. In this case we could, for example, report the boundary value $s/n = 1$ or else extend the parameter space to the closed interval.

Example 5.4. Consider $X_1, \dots, X_n \sim N(\mu, \sigma^2)$, where both parameters are unknown. Then we have $\theta = (\mu, \sigma^2)$. Taking logarithms we get

$$\ell(\mu, \sigma^2) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2.$$

Since we have a two-dimensional parameter, the likelihood equation translates to a system of two equations: we set both partial derivatives equal to zero. Taking derivatives with respect to μ we find

$$\frac{\partial \ell}{\partial \mu} = \frac{1}{\sigma^2} \sum_{i=1}^n (x_i - \mu),$$

which is zero only for $\mu = \bar{x}$. Since the sum of squares $\sum_i (x_i - \mu)^2$ takes its minimum for $\mu = \bar{x}$, this value maximises ℓ over μ for every fixed σ^2 and we can substitute $\mu = \bar{x}$ to find the maximum of the remaining function of one variable. Let $v = \sigma^2$ and $W = \sum_i (x_i - \bar{x})^2$. Then we have

$$\ell(\bar{x}, v) = -\frac{n}{2} \log(2\pi v) - \frac{W}{2v},$$

with derivative

$$\frac{\partial \ell}{\partial v} = -\frac{n}{2v} + \frac{W}{2v^2} = \frac{W - nv}{2v^2},$$

which is positive for $v < W/n$ and negative for $v > W/n$. Thus, the maximum is at $v = W/n$. Hence, the MLE is the pair

$$\hat{\mu} = \bar{X} \quad \text{and} \quad \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2.$$

This is the plug-in estimator, obtained by using the method of moments in example 4.8, with divisor n : for the normal distribution, both methods give the same result for both coordinates. In particular, the MLE for the variance is the biased estimator from section 2.3, where we showed that the bias in this case was harmless. Whether the bias is a general property of maximum likelihood is a question we will consider later.

5.3 When the Recipe Fails

Both assumptions in the recipe, *i.e.* an interior maximum and differentiable log-likelihood, break down for the uniform model. In this case we have to return to definition 5.1 and maximise the expression directly.

Example 5.5. Let $X_1, \dots, X_n \sim \text{Uniform}(0, \theta)$ be i.i.d., as in example 1.2. From the joint density found in exercise 1.3, the likelihood is

$$L(\theta) = \begin{cases} \theta^{-n} & \text{if } \theta \geq \max_i x_i, \\ 0 & \text{otherwise.} \end{cases}$$

Calculus is useless here: on the region where L is positive its derivative $-n\theta^{-n-1}$ is negative everywhere, so the likelihood equation has no solution, and at $\theta = \max_i x_i$ the function is not even continuous. But maximising directly is easy:

the likelihood is zero for $\theta < \max_i x_i$ and strictly decreasing for $\theta \geq \max_i x_i$, and thus largest at the left edge of the positive region. The MLE is the sample maximum,

$$\hat{\theta} = \max_i X_i.$$

This identifies the estimator we have been tracking since exercise 1.3, which by exercise 4.4 beats the method of moments estimator $2\bar{X}$ in mean squared error for every $n \geq 3$. The likelihood also rules out impossible parameter values automatically: for the dataset $x_1 = 0.1$, $x_2 = 0.2$, $x_3 = 0.9$ of lecture 4, any $\theta < 0.9$ has likelihood zero, so the contradiction the method of moments overlooked cannot arise.

The uniform model fails the recipe in an obvious way; a failure which is harder to spot is a likelihood equation with several roots. Consider the t location model with $\nu \geq 1$ degrees of freedom, where ν is fixed and known and the density is given by

$$f(x; \theta) = c_\nu \left(1 + \frac{(x - \theta)^2}{\nu}\right)^{-(\nu+1)/2}, \quad x \in \mathbb{R},$$

with location $\theta \in \mathbb{R}$ and a normalising constant c_ν free of θ . The log-likelihood is differentiable everywhere, so the recipe seems to apply, but after clearing denominators the likelihood equation

$$\ell'(\theta) = (\nu + 1) \sum_{i=1}^n \frac{x_i - \theta}{\nu + (x_i - \theta)^2} = 0$$

turns into a polynomial equation of degree $2n - 1$ in θ and for $n \geq 2$ this equation can have several roots. Exercise 5.4 shows that for $n = 2$ there are three roots when the observations are more than $2\sqrt{\nu}$ apart: the mid-point of the interval between the observations, which is then a local *minimum* of the likelihood, and two global maxima of equal likelihood, one close to each of the observations. Thus, the model prefers to fit one data point and to consider the other as an outlier, and in this case the data cannot choose between the two maxima and the MLE will not be unique. The exercise also shows that a root of the likelihood equation can be a minimum and that the checking step in the recipe is not optional.

For larger samples the likelihood equation of the t location model cannot be solved in closed form, and similarly the likelihood equations for many other models, *e.g.* the two-parameter gamma distribution from example 4.9, cannot be solved in closed form. The solution to this problem, which is to maximise ℓ numerically, will be discussed, together with its pitfalls, in lecture 18.

5.4 Bias of the Maximum Likelihood Estimator

Nothing in definition 5.1 mentions bias and indeed the MLE is not always unbiased (we already know this from example 5.4 for the normal distribution). In this section we consider the exponential distribution with rate θ , where all moments can be determined exactly.

Let $X_1, \dots, X_n$ be i.i.d. exponential with rate θ and let $\hat{\theta} = 1/\bar{X}$ be the MLE from example 5.2. Setting $T = X_1 + \dots + X_n$, so that $\hat{\theta} = n/T$, we found in exercise 2.3, using the gamma density of T , that $\mathbb{E}_\theta(\hat{\theta}) = n\theta/(n - 1)$ for

$n \geq 2$, while for $n = 1$ the expectation is infinite. Thus the MLE is biased, with bias $\theta/(n-1)$, and multiplying by $(n-1)/n$ removes the bias: the estimator $\hat{\theta} = (n-1)\hat{\theta}/n = (n-1)/T$ satisfies $\mathbb{E}_{\theta}(\hat{\theta}) = \theta$.

From lecture 2 we know that removing the bias does not automatically give the estimator with the smaller mean squared error. In exercise 5.3 we compute the second moment of $\hat{\theta}$ and find $\text{MSE}(\tilde{\theta})/\text{MSE}(\hat{\theta}) = (n-1)/(n+2) < 1$ for all $n \geq 3$. Thus, after removing the bias, this estimator is strictly better. The same exercise shows that a third scaling of $\hat{\theta}$ would be even better, similar to how the divisor $n+1$ in the normal distribution variance was chosen.

For the normal variance the biased divisor- n estimator was better than the unbiased sample variance S^2 (section 2.3); for the exponential rate the bias-corrected estimator is better. Neither result is a theorem: the bias-variance trade-off can go either way, and the only way to know which is better is to compute both mean squared errors. In favour of the MLE we note that the bias decreases to zero as the sample size increases, and in lecture 17 we will see that for large sample size the MLE is in some sense the best estimator.

5.5 Invariance Under Reparametrisation

Often the quantity we want to estimate is not the parameter itself, but a function of the parameter, *e.g.* a probability, a mean or a median. The following theorem shows that under maximum likelihood such quantities transform together with the parameter.

Theorem 5.6 (Invariance of the MLE). *Let $g: \Theta \rightarrow \Psi$ be a bijection and let $\hat{\theta}$ be a maximum likelihood estimator for θ . Then $g(\hat{\theta})$ is a maximum likelihood estimator for the transformed parameter $\psi = g(\theta)$.*

The idea of the proof is that re-labelling the parameter does not change the family of distributions: in terms of ψ the likelihood is $\tilde{L}(\psi) = L(g^{-1}(\psi))$, so maximising $\tilde{L}$ over Ψ is the same maximisation as before. Writing this down is exercise 5.5.

Remark. If g is not a bijection, then the likelihood of ψ is defined to be the maximal likelihood of the values of θ which map to this value. While the statement $\hat{\psi} = g(\hat{\theta})$ is still true, we will not require the general case in this module.

As an application, exercise 5.8 finds the MLE for the *median* waiting time $m = (\log 2)/\theta$ of the exponential model: since $\theta \mapsto (\log 2)/\theta$ is a bijection, theorem 5.6 gives $\hat{m} = (\log 2) \bar{X}$ with no further maximisation. The exercise also contains a surprise: $\hat{m}$ is unbiased, despite the MLE for the rate being biased. In general, unbiasedness is lost under nonlinear reparametrisation, and thus there is no analogue of theorem 5.6 for unbiasedness.

Since we have seen that the MLE can be biased, non-unique and even non-existent, we need to do more work to understand the method. The arguments for the validity of the MLE can be found in lectures 13, 16 and 17: the main result is the Cramer–Rao bound, which no unbiased estimator can violate. For large sample sizes, the MLE is shown to be consistent and to asymptotically achieve this bound. In preparation for these arguments, lecture 8 summarises the observation from example 4.3 that the likelihood function can be collapsed to a small number of data summaries.

Summary.

- The maximum likelihood estimator $\hat{\theta}$ is the parameter value which maximises the likelihood (or log-likelihood) of the data.
- For smooth problems, where the maximum is an interior point of the parameter space, the MLE can be found by solving the likelihood equation $\ell'(\theta) = 0$ and verifying that the root is a maximum. For the standard families of models this gives $\bar{X}$ (Bernoulli and Poisson), $1/\bar{X}$ (exponential rate) and $(\bar{X}, (1/n) \sum_i (X_i - \bar{X})^2)$ (normal).
- If the recipe fails, the definition can still be used: for the uniform distribution the maximum sits at the boundary, at $\hat{\theta} = \max_i X_i$, and for the t location model with two observations far apart the MLE is not unique.
- The MLE is not unbiased in general. Whether removing the bias improves the mean squared error depends on the model: for the exponential rate the corrected estimator wins, reversing the normal-variance comparison of lecture 2.
- The MLE is invariant under reparametrisation: the MLE for $g(\theta)$ is $g(\hat{\theta})$. Unbiasedness is not invariant.
- The full justification of the method, including efficiency and large-sample optimality, follows in lectures 13, 16 and 17.

Exercise 5.1. Let $X_1, \dots, X_n$ be an i.i.d. sample from the geometric distribution with parameter $p \in (0, 1)$, i.e. $f(x; p) = (1 - p)^{x-1}p$ for $x \in \{1, 2, \dots\}$ as in exercise 4.2.

1. Show that the log-likelihood function is given by $\ell(p) = n \log p + (\sum_i x_i - n) \log(1 - p)$.
2. Solve the likelihood equation and show that your solution is a maximum. Deduce that the MLE is $\hat{p} = 1/\bar{X}$.
3. Discuss your results in the light of the method of moments estimator from exercise 4.2.

Exercise 5.2. Let $X_1, \dots, X_n$ be i.i.d. with density $f(x; \theta) = \theta^{-2} x e^{-x/\theta}$ for $x \geq 0$, where $\theta > 0$. This is the gamma distribution with known shape 2 and scale θ , i.e. $\text{Gamma}(2, 1/\theta)$ in the rate parametrisation from table A.2.

1. Find the MLE $\hat{\theta}$ for θ .
2. Show that your estimator $\hat{\theta}$ is unbiased.
3. Determine $\text{MSE}(\hat{\theta})$ and use this to show that $\hat{\theta}$ is consistent.

Exercise 5.3. Let $X_1, \dots, X_n$ be i.i.d. exponential with rate θ , let $\hat{\theta} = 1/\bar{X}$ be the MLE and assume $n \geq 3$. Consider the family of estimators $c\hat{\theta}$ for constants $c > 0$.

1. Using the gamma density of $T = X_1 + \dots + X_n$ as in exercise 2.3, show that $\mathbb{E}_\theta(\hat{\theta}^2) = n^2 \theta^2 / ((n-1)(n-2))$.

2. Show that

$$\text{MSE}(c\hat{\theta}) = \left(\frac{n^2}{(n-1)(n-2)} c^2 - \frac{2n}{n-1} c + 1 \right) \theta^2,$$

and deduce the values $\text{MSE}(\hat{\theta}) = (n+2)\theta^2/((n-1)(n-2))$ for the MLE and $\text{MSE}(\tilde{\theta}) = \theta^2/(n-2)$ for the unbiased version $\tilde{\theta} = (n-1)\hat{\theta}/n$ from the text, and the ratio $\text{MSE}(\tilde{\theta})/\text{MSE}(\hat{\theta}) = (n-1)/(n+2)$.

3. Determine the value of c which minimises the mean squared error, and show that the minimal mean squared error is $\theta^2/(n-1)$. Compare your result with the MLE ($c = 1$), the unbiased version $c = (n-1)/n$, and the result for the normal variance in section 2.3. Comment upon your result.

Exercise 5.4. Let $x_1 > x_2$ be any two observations from the t location model given in section 5.3. Let $m = (x_1 + x_2)/2$ be the mid-point and $d = (x_1 - x_2)/2$ be half the distance between the two observations.

1. Let $a = x_1 - \theta$ and $b = x_2 - \theta$. Show that the likelihood equation $\ell'(\theta) = 0$ is equivalent to $(a+b)(\nu + ab) = 0$.
2. Show that for $\theta = m + t$ the log-likelihood function is $\ell(m+t) = C - \frac{\nu+1}{2} \log P(t)$, where C does not depend on t and

$$P(t) = (\nu + d^2)^2 + 2(\nu - d^2)t^2 + t^4.$$

3. Show that for $d^2 < \nu$ the mid-point m is the unique MLE, but for $d^2 > \nu$ the likelihood equation has three roots $m, m \pm \sqrt{d^2 - \nu}$, where m is a local minimum of ℓ and the other two are maximum likelihood estimates with equal likelihood.

Exercise 5.5. Prove the statement of theorem 5.6: Let $g: \Theta \rightarrow \Psi$ be a bijection and let the likelihood in the new parametrisation be given by $\tilde{L}(\psi) = L(g^{-1}(\psi))$. Show that $g(\hat{\theta})$ is the maximum of $\tilde{L}$ over Ψ .

Exercise 5.6. Let X be an observation from the exponential distribution with rate $\theta > 0$, i.e. $\mathbb{P}_\theta(X > x) = e^{-\theta x}$ for $x \geq 0$. The following properties of the exponential distribution will be needed later in the module.

1. Show that the exponential distribution has no memory, i.e. that for all $s, t \geq 0$ we have

$$\mathbb{P}_\theta(X > s+t | X > s) = \mathbb{P}_\theta(X > t).$$

2. Show that for all $c \geq 0$ we have

$$\mathbb{E}_\theta(X | X > c) = c + \frac{1}{\theta}.$$

Comment upon your result in the light of the first part of the question.

Exercise 5.7. Let $X_1, \dots, X_n$ be an i.i.d. sample from the Poisson distribution with parameter $\theta > 0$, and assume that $\sum_i x_i \geq 1$.

1. In exercise 4.1 we have seen that the score equals zero at $\bar{x}$. Show that $\bar{X}$ is indeed the MLE for θ .
2. Determine the MLE for the probability $p_0 = \mathbb{P}_\theta(X_1 = 0)$ of observing zero.
3. Using the relation $\mathbb{E}_\theta(z^{X_1}) = e^{\theta(z-1)}$ for $z > 0$, show that your estimator from the previous part is biased, but asymptotically unbiased. Comment on your result in the light of the fact that $\bar{X}$ is unbiased.

Exercise 5.8. Assume that we have observed the i.i.d. exponential waiting times $X_1, \dots, X_n$ with rate $\theta > 0$, and that we want to estimate the *median* waiting time m , *i.e.* the value such that $\mathbb{P}_\theta(X_1 \leq m) = 1/2$.

1. Show that $m = g(\theta) = (\log 2)/\theta$ and that the map g maps $(0, \infty)$ to $(0, \infty)$ bijectively.
2. Using theorem 5.6 and example 5.2, show that the MLE for the median is given by $\hat{m} = (\log 2) \bar{X}$.
3. Show that $\hat{m}$ is an unbiased estimator for m , and comment on why this does not contradict the bias of the MLE $\hat{\theta} = 1/\bar{X}$ for the rate observed in this lecture.

Lecture 6

Estimators by Simulation

This is the R session for week 2. In lectures this week we introduced two estimators for the uniform model: the method of moments estimator $2\bar{X}$ in lecture 4 and the maximum likelihood estimator $\max_i X_i$ in lecture 5. In exercise 4.4 you used algebra to determine which of these two estimators is better. Today we will use the Monte Carlo methods from lecture 3 to answer the same question, and we will learn how to write our own function in R and how to repeat this function in a loop. (Appendix B shows both approaches, but at a slower pace.) The approach we follow here is a good one to get into: before you do any simulation, try to predict what the result will be. Only then can the output of your simulation either confirm or contradict your expectation.

6.1 Functions and Loops

The model we consider today is $X_1, \dots, X_n \sim \text{Uniform}(0, \theta)$, and to study the estimators for this model using simulation, we will need to repeat the same experiment thousands of times: we have to sample from the model, compute the estimate, and store the result. While we could just type in the commands to do this over and over again in the console, as we did in lecture 3, this approach is not scalable. Instead, we will encapsulate the experiment into a *function* and then use R to repeat this function many times.

To generate samples from the uniform distribution, we can use the function `runif` and to specify the boundaries of the distribution we use `min` and `max`, as explained in appendix B. To create a function, we use the keyword `function`, followed by a list of arguments and then the body of the function enclosed in braces. The return value of a function is the value of the last line of the body. The following function performs a complete experiment for the method of moments estimator: it simulates a fresh sample of size `n` from the model with true parameter `theta`, and then returns the computed value $2\bar{x}$ of the estimate.

```
est_mom <- function(n, theta) {  
  sample <- runif(n, min = 0, max = theta)  
  2 * mean(sample)  
}  
set.seed(2026)  
est_mom(10, 1)  
est_mom(10, 1)
```

```
## [1] 0.8844232
## [1] 0.6339234
```

Defining the function prints nothing; only the two calls do. Each call draws its own sample, so two calls give two different estimates, both scattered around the true value $\theta = 1$: once more the estimator is a random quantity, as seen in lecture 3 for the sample mean.

To repeat the experiment systematically, we can use a **for**-loop: we set aside a vector to store the results using `numeric(5)`, then use a counter `i` which takes the values 1 to 5 (generated by `seq_len(5)`), and then store the estimate for each iteration of the loop in the results vector. (Appendix B explains why we use `seq_len` instead of the shorter notation `1:5`.)

```
set.seed(2026)
estimates <- numeric(5)
for (i in seq_len(5)) {
  estimates[i] <- est_mom(10, 1)
}
estimates
```

```
## [1] 0.8844232 0.6339234 0.5469486 0.9469991 0.8432259
```

Since we reset the seed, the first two values are exactly the two estimates computed above: the loop repeats the same call and keeps the answers. For this everyday pattern of repeating an expression and collecting the results, R provides the shortcut `replicate` which can be used to achieve the same effect in one line.

```
set.seed(2026)
replicate(5, est_mom(10, 1))
```

```
## [1] 0.8844232 0.6339234 0.5469486 0.9469991 0.8432259
```

The five values agree with the loop, as they must. We use `replicate` for the rest of this lecture; the explicit loop returns in lecture 9, where the sample size changes from one run to the next.

6.2 Empirical Bias and Mean Squared Error

With `est_mom` in place we can now use the Monte Carlo approach from lecture 3: we generate many estimates and then compute summary statistics. If $\hat{\theta}^{(k)}$ is the estimate using the k -th of N simulated samples, then by the law of large numbers (theorem A.14), the average of the deviations $\hat{\theta}^{(k)} - \theta$ is approximately equal to the bias $\mathbb{E}_{\theta}(\hat{\theta}) - \theta$ from lecture 2, and the average of the squared deviations is approximately equal to the mean squared error $\mathbb{E}_{\theta}((\hat{\theta} - \theta)^2)$. In R, both summary statistics can be computed in a single line: the empirical bias is given by `mean(estimates) - theta` and the empirical mean squared error is given by `mean((estimates - theta)^2)`.

We now have all the ingredients for a complete Monte Carlo study: a function to perform the experiment, `replicate` to repeat it and two one-liners for the summaries. The box below shows such a study in full, but with the lines of the

R script intentionally shuffled. Before you continue, arrange the lines of this box in the correct order to form an R script which performs a Monte Carlo study with $\theta = 1$, $n = 10$ and 10000 repetitions.

```
estimates <- replicate(10000, est_mom(n, theta))
mean((estimates - theta)^2)    # empirical MSE
theta <- 1
n <- 10
est_mom <- function(n, theta) {
  sample <- runif(n, min = 0, max = theta)
  2 * mean(sample)
}
mean(estimates) - theta        # empirical bias
set.seed(2026)
```

The rule is that all names must be defined before they are used: the function and the values for `theta` and `n` must come before the `replicate` line, the seed must be set before any random draws are made, and the two summaries need the finished vector `estimates` to be computed, so they come last. There is more than one order which satisfies these conditions, so any of these orders is correct. The reconstructed study is shown below; for comparison the theoretical mean squared error is also shown. Before we run the program, we make a final prediction: from exercise 4.4 we know that $2\bar{X}$ is unbiased and has $\text{MSE}(2\bar{X}) = \theta^2/(3n)$, where $1/30 \approx 0.0333$, so should the two printed summaries be exactly 0 and 0.0333, or just close to these values? How close?

```
est_mom <- function(n, theta) {
  sample <- runif(n, min = 0, max = theta)
  2 * mean(sample)
}
theta <- 1
n <- 10
set.seed(2026)
estimates <- replicate(10000, est_mom(n, theta))

mean(estimates) - theta        # empirical bias
mean((estimates - theta)^2)    # empirical MSE
theta^2 / (3 * n)              # the theoretical value
```

```
## [1] 0.00226647
## [1] 0.03303851
## [1] 0.03333333
```

The numbers are close to the theory, but not identical, and this is expected: the empirical bias and mean squared error are averages over 10000 random repetitions, so are estimates of the true values, but not identical to the true values. Changing the seed of the random number generator would result in slightly different results. Similarly, if the number of repetitions was increased, the results would get closer to the theoretical values. So, agreement to two decimal places is a good result, and an empirical bias of exactly zero would be a cause for concern (not for celebration).

6.3 Comparing Two Estimators

The real use of the Monte Carlo method is to compare estimators in models where the algebra runs out. Before trusting it there, we should watch it succeed on a contest whose outcome we know: by exercises 2.2 and 4.4, the maximum likelihood estimator $\max_i X_i$ of example 5.5 is biased but has smaller mean squared error than the unbiased estimator $2\bar{X}$ for all $n \geq 3$. If the simulation reproduces these known answers, we may trust the same recipe in the practical report and in lecture 18, where no formula is available.

We can write the maximum likelihood estimator as a function, in analogy to `est_mom`. To allow for a fair comparison, we use the same seed 2026 as we used above, to ensure that both estimators use the same 10000 samples. Before you try the code, what should the two printed numbers be close to? From exercise 2.2 we know that the bias of the maximum likelihood estimator is $-\theta/(n+1) = -1/11$ and the mean squared error is $2\theta^2/((n+1)(n+2)) = 2/132 \approx 0.0152$.

```
est_mle <- function(n, theta) {
  sample <- runif(n, min = 0, max = theta)
  max(sample)
}
set.seed(2026)
mle <- replicate(10000, est_mle(n, theta))

mean(mle) - theta           # empirical bias
mean((mle - theta)^2)       # empirical MSE
```

```
## [1] -0.08960997
## [1] 0.01479181
```

Both predictions are correct: the empirical bias -0.0896 is close to the theoretical value $-1/11 \approx -0.0909$ and the empirical mean squared error 0.0148 is close to the theoretical value 0.0152 , which is less than half of the 0.0330 for $2\bar{X}$ shown above. Thus, in this example, the biased estimator performs better than the unbiased estimator, as shown in exercise 4.4.

The histograms make the comparison visible, and this is the time for the main prediction of the day. Before you run the code, sketch the two histograms you expect, using the theory above for the centres and the widths (the standard deviation of $2\bar{X}$ is $\sqrt{1/30} \approx 0.18$). Which of the two histograms will be more spread out? Where will the two histograms be centred? Will the second histogram be bell-shaped like the first one? The command `par(mfrow = c(2, 1))` can be used to arrange the following two plots as a column, and by using `xlim` for both plots we keep the axes directly comparable, as in lecture 3.

```
par(mfrow = c(2, 1))
hist(estimates, breaks = 40, col = "grey85", main = "method of moments",
     xlab = "estimate", xlim = c(0.4, 1.6))
abline(v = theta, lty = 2)
hist(mle, breaks = 40, col = "grey85", main = "maximum likelihood",
     xlab = "estimate", xlim = c(0.4, 1.6))
abline(v = theta, lty = 2)
```





Figure 6.1: Simulated sampling distributions of the two estimators for the uniform parameter, using the same axis: the method of moments estimator $2\bar{X}$ (top) and the maximum likelihood estimator $\max_i X_i$ (bottom). The true value $\theta = 1$ is given by the dashed line.

Figure 6.1 shows that the first two guesses were right, and that the third guess is wrong for most people. The top histogram is a symmetric, bell-shaped histogram, centred around the dashed line: the estimator $2\bar{X}$ is unbiased. The bell-shape is expected, since it is an average and by the central limit theorem (theorem A.15), averages are asymptotically normal. (We will learn more about this result in lecture 12.) The histogram for the maximum likelihood estimator is much narrower, and is shifted to the left of the dashed line. This indicates a bias $-1/11$. The histogram is not bell-shaped at all: no simulated estimates are above 1, since the maximum of a sample from $\text{Uniform}(0, \theta)$ can never be above θ and the estimates are converging to this boundary. This shows that not every estimator has a bell-shaped sampling distribution: the central limit theorem only applies to averages and does not apply to the maximum. We will see this histogram again in lecture 13, where the uniform model will be used as a counterexample to the general theory.

In lecture 9 we will vary the sample size and will see the histograms getting tighter around the truth.

Just type the commands you see in this solution in your own R session, to make sure you understand what they do.

Summary.

- Appendix B is the written reference for functions, loops and `replicate`.
- A function created with `function(n, theta) {...}` packages one simulation experiment; a `for`-loop or, more conveniently, `replicate` repeats it and collects the estimates.
- The one-liner `mean(estimates) - theta` gives the empirical bias, and `mean((estimates - theta)^2)` gives the empirical mean squared error; both approximate the true values up to a simulation error that shrinks as the number of repetitions grows.
- For the uniform model with $n = 10$, the simulation reproduced the analytic results of exercise 4.4: the biased estimator $\max_i X_i$ beats the unbiased $2\bar{X}$ in mean squared error, and its sampling distribution is a skewed pile below θ , not a bell.
- Predict before you run: a simulation only teaches when its output can confirm or contradict a stated expectation.

Lecture 7

Exponential Families

In lectures 4 and 5 we repeatedly collapsed products of n densities to expressions which summarised the data using only one or two sums. This lecture explains the reason for this coincidence: almost all families of models considered here have the same algebraic shape, and models of this shape are called *exponential families*. This shape allows us to define the data summary used in the likelihood, the moments of this summary, and the smoothness required for the later theory, and many of the results in the following weeks are almost automatic for this class.

7.1 Recognising the Shape

We start by considering a familiar distribution and then rewrite its probability weights so that the data and the parameter are as far apart as possible. For the Bernoulli distribution with $x \in \{0, 1\}$ we can write

$$\begin{aligned} f(x; \theta) &= \theta^x (1 - \theta)^{1-x} = \exp(x \log \theta + (1 - x) \log(1 - \theta)) \\ &= \exp\left(x \log \frac{\theta}{1 - \theta} + \log(1 - \theta)\right). \end{aligned}$$

Here we have written the weights as the exponential of the product of one function of θ and one function of x , minus a function of only θ . The Poisson weights and the exponential density can be written in a similar way: we have $\theta^x = e^{x \log \theta}$ and $\theta e^{-\theta x} = e^{-\theta x + \log \theta}$, up to a factor depending on x . We will give a name to this structure.

Definition 7.1. A parametric model $(f(x; \theta) \mid \theta \in \Theta)$ with $\Theta \subseteq \mathbb{R}$ is a *one-parameter exponential family*, if the density or the probability weights can be written as

$$f(x; \theta) = h(x) \exp(\eta(\theta) T(x) - B(\theta))$$

for all x and all $\theta \in \Theta$, where $h \geq 0$ and T are functions of x alone and η and B are functions of θ alone. In particular, the support $\{x \mid f(x; \theta) > 0\}$ is $\{x \mid h(x) > 0\}$ and does not depend on θ . The statistic $T(X)$ is called the *natural statistic* and $\eta(\theta)$ is called the *natural parameter* of the family.

Since the exponential is always positive, $f(x; \theta)$ is positive if and only if $h(x)$ is. This proves the statement about the support. The representation is not unique, since we can replace η by $c\eta$ and T by T/c for any constant $c \neq 0$.

Finally, since the density must integrate (or sum) to one, the function B is determined by the other three functions: we have

$$B(\theta) = \log \int h(x) \exp(\eta(\theta) T(x)) dx. \quad (7.1)$$

These rewritings give the following examples.

Example 7.2. The Bernoulli distribution with $\theta \in (0, 1)$ is a one-parameter exponential family with $h(x) = 1$ for $x \in \{0, 1\}$, natural statistic $T(x) = x$, natural parameter $\eta(\theta) = \log(\theta/(1 - \theta))$, the *log-odds* of success, and $B(\theta) = -\log(1 - \theta)$.

Example 7.3. The Poisson distribution with $\theta > 0$ is a one-parameter exponential family with $h(x) = 1/x!$ for $x \in \{0, 1, 2, \dots\}$, natural statistic $T(x) = x$, natural parameter $\eta(\theta) = \log \theta$ and $B(\theta) = \theta$.

Example 7.4. The exponential distribution with rate $\theta > 0$ is a one-parameter exponential family with $h(x) = 1$ for $x \geq 0$ and $h(x) = 0$ otherwise, natural statistic $T(x) = x$, natural parameter $\eta(\theta) = -\theta$ and $B(\theta) = -\log \theta$.

The normal distribution is also an exponential family, but since both parameters are unknown, we need to use the two-parameter version of the definition from section 7.3.

Remark. The uniform distribution on $[0, \theta]$ is not an exponential family, since the support of the distribution changes with the parameter. In contrast, the support of an exponential family is always fixed. This is the same reason why the likelihood equation was useless in example 5.5. Similarly, the t location model from section 5.3 is not an exponential family, despite the fact that its support is all of $\mathbb{R}$; we do not prove this.

The shape survives the passage from one observation to a sample.

Lemma 7.5. Let $X_1, \dots, X_n$ be i.i.d. from the one-parameter exponential family in definition 7.1. Then the joint density of the sample is

$$f(x_1, \dots, x_n; \theta) = \prod_{i=1}^n h(x_i) \exp\left(\eta(\theta) \sum_{i=1}^n T(x_i) - nB(\theta)\right),$$

which is of the form of definition 7.1 with the same natural parameter $\eta(\theta)$, natural statistic $T_n = \sum_{i=1}^n T(X_i)$, the function $\prod_i h(x_i)$ instead of h and $nB(\theta)$ instead of $B(\theta)$.

Proof. Since the observations are independent, the joint density is the product $\prod_{i=1}^n h(x_i) \exp(\eta(\theta) T(x_i) - B(\theta))$ of the marginal densities. Combining the n exponents using the rule $e^a e^b = e^{a+b}$ gives the required form. This completes the proof. $\square$

Read as a function of θ , the joint density is the likelihood, and thus the log-likelihood of an exponential family sample is

$$\ell(\theta) = \sum_{i=1}^n \log h(x_i) + \eta(\theta) T_n - nB(\theta). \quad (7.2)$$

The first term does not involve θ and thus, except for the term involving n , the likelihood depends on the data only via the single number $T_n = \sum_i T(x_i)$. This is an extension of example 4.3 to the whole class; in lecture 8 we will see in which sense the value T_n contains all the information about θ .

7.2 The Natural Parametrisation and the Cumulant Function

Since the parameter enters the density only via $\eta(\theta)$, it is often convenient to use η as the parameter. In this case, the normalising function (7.1) could be given a name, since its derivatives are the moments of the natural statistic.

Definition 7.6. Let h and T be as in definition 7.1. Then the *natural parametrisation* of the family is given by

$$f(x; \eta) = h(x) \exp(\eta T(x) - K(\eta)), \quad \eta \in \mathcal{N},$$

where the *cumulant function* (also called the *log-partition function*) is given by

$$K(\eta) = \log \int h(x) e^{\eta T(x)} dx,$$

where the sum replaces the integral for discrete models, and the *natural parameter space* is the set $\mathcal{N} = \{\eta \in \mathbb{R} \mid K(\eta) < \infty\}$ of all η such that the integral is finite.

Since K makes $f(x; \eta)$ integrate to one, the formula defines a probability distribution for every $\eta \in \mathcal{N}$, whether or not η is of the form $\eta(\theta)$; comparing with (7.1) shows that $\eta(\theta) \in \mathcal{N}$ and $B(\theta) = K(\eta(\theta))$, so that the original family sits inside this larger one. For the three examples of section 7.1 we find $K(\eta) = \log(1 + e^\eta)$ with $\mathcal{N} = \mathbb{R}$ for the Bernoulli distribution, $K(\eta) = e^\eta$ with $\mathcal{N} = \mathbb{R}$ for the Poisson distribution and $K(\eta) = -\log(-\eta)$ with $\mathcal{N} = (-\infty, 0)$ for the exponential distribution. In all three cases $\mathcal{N}$ is an open interval and $\theta \mapsto \eta(\theta)$ is a bijection onto $\mathcal{N}$, making the natural parametrisation a relabelling in the sense of theorem 5.6.

The first two derivatives of K give the mean and the variance of $T(X)$. The name of K is derived from the fact that $s \mapsto K(\eta + s) - K(\eta)$ is the cumulant generating function of $T(X)$, as given in exercise 7.6.

Theorem 7.7. Let η be an interior point of the natural parameter space $\mathcal{N}$. Then, in the natural parametrisation,

$$\mathbb{E}_\eta(T(X)) = K'(\eta) \quad \text{and} \quad \text{Var}_\eta(T(X)) = K''(\eta).$$

Proof. Set $Z(\eta) = \int h(x) e^{\eta T(x)} dx$, so that $K = \log Z$. Since $h \geq 0$ is not identically zero and $e^{\eta T(x)} > 0$, we have $Z(\eta) > 0$ and thus the division in the proof is justified. Since η is in the interior of $\mathcal{N}$, we can take derivatives inside the integral as many times as required; this is a result from analysis which we use without proof. Taking derivatives once, we get

$$Z'(\eta) = \int T(x) h(x) e^{\eta T(x)} dx.$$

Dividing by $Z(\eta) = e^{K(\eta)}$ changes the integrand to $T(x)$ times the density $f(x; \eta)$ and thus we get

$$K'(\eta) = \frac{Z'(\eta)}{Z(\eta)} = \int T(x) h(x) e^{\eta T(x) - K(\eta)} dx = \mathbb{E}_\eta(T(X)).$$

Taking derivatives again, we find

$$K''(\eta) = \frac{Z''(\eta)}{Z(\eta)} - \left(\frac{Z'(\eta)}{Z(\eta)} \right)^2,$$

and using the same argument as above, with $T(x)^2$ in the integrand, we find $Z''(\eta)/Z(\eta) = \mathbb{E}_\eta(T(X)^2)$. Thus we have $K''(\eta) = \mathbb{E}_\eta(T(X)^2) - \mathbb{E}_\eta(T(X))^2 = \text{Var}_\eta(T(X))$ and the proof is complete. $\square$

The theorem allows us to compute moments by differentiation: for the Poisson distribution $K(\eta) = e^\eta$ we find $K'(\eta) = K''(\eta) = e^\eta = \theta$, the mean and variance from table A.1, and in exercise 7.4 we consider the exponential and normal distributions. For T_n , being a sum of n i.i.d. copies of $T(X)$, we can use the theorem to find $\mathbb{E}_\eta(T_n) = nK'(\eta)$ and $\text{Var}_\eta(T_n) = nK''(\eta)$.

7.3 Several Parameters

The normal distribution where both μ and σ^2 are unknown does not fit definition 7.1, since there are two unknown quantities in the exponent. The extension here is to allow the sum of products in the exponent, one for each parameter.

Definition 7.8. A parametric model $(f(x; \theta) \mid \theta \in \Theta)$ with $\Theta \subseteq \mathbb{R}^k$ is a *k-parameter exponential family*, if the density or the probability weights can be written as

$$f(x; \theta) = h(x) \exp\left(\sum_{j=1}^k \eta_j(\theta) T_j(x) - B(\theta)\right)$$

for all x and all $\theta \in \Theta$, where $h \geq 0$ and $T_1, \dots, T_k$ are functions of x alone and $\eta_1, \dots, \eta_k$ and B are functions of θ alone. The vector $(T_1(X), \dots, T_k(X))$ is the *natural statistic* and $(\eta_1(\theta), \dots, \eta_k(\theta))$ is the *natural parameter* of the family.

The support of the distribution is again $\{x \mid h(x) > 0\}$ and the proof of lemma 7.5 applies without change: an i.i.d. sample is again a k -parameter exponential family, where the natural statistic is the vector of the k sums $\sum_i T_j(X_i)$. The cumulant function can be extended to this situation, but we do not need this result here.

Example 7.9. By expanding the square in the exponent of the density of the normal distribution $N(\mu, \sigma^2)$ from table 1.2, we can write

$$f(x; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{x^2}{2\sigma^2} + \frac{\mu}{\sigma^2} x - \frac{\mu^2}{2\sigma^2}\right).$$

If only μ is unknown, the factor $(2\pi\sigma^2)^{-1/2}$ and the term $-x^2/(2\sigma^2)$ contain no unknown quantities and we can define the function $h(x) = (2\pi\sigma^2)^{-1/2} e^{-x^2/(2\sigma^2)}$: we have found a one-parameter exponential family with natural statistic $T(x) = x$, natural parameter $\eta(\mu) = \mu/\sigma^2$ and $B(\mu) = \mu^2/(2\sigma^2)$. If both parameters are unknown, we have to interpret the term $-x^2/(2\sigma^2)$ as a second product and we get a two-parameter exponential family with $h(x) = 1$, natural statistic (x, x^2) , natural parameter $(\mu/\sigma^2, -1/(2\sigma^2))$ and $B(\theta) = \mu^2/(2\sigma^2) + \frac{1}{2} \log(2\pi\sigma^2)$. For a sample, the natural statistic $(\sum_i x_i, \sum_i x_i^2)$ is equivalent to the pair $(\bar{x}, \sum_i (x_i - \bar{x})^2)$ used to derive the maximum likelihood estimator in example 5.4.

The gamma distribution with both parameters unknown is a second example of a two-parameter exponential family. If we write $x^{\alpha-1} = e^{(\alpha-1)\log x}$, then the natural statistic is $(\log x, x)$ and the details can be worked out using exercise 7.3. For a sample of size n , the natural statistic is the pair $(\sum_i \log x_i, \sum_i x_i)$, which is different from the pair of moments used in example 4.9. This is the first sign that the moment estimator is missing some information. (We will see in lecture 8 that this is indeed the case.)

For the completeness statements in lecture 10 it is important that the natural parameter varies in as many directions as there are natural statistics. Mathematically, this can be expressed as the following condition.

Definition 7.10. A one-parameter exponential family is of *full rank*, if the set $\eta(\Theta) = \{\eta(\theta) \mid \theta \in \Theta\}$ of natural parameter values contains a non-degenerate open interval. A k -parameter exponential family is of *full rank*, if the set $\{(\eta_1(\theta), \dots, \eta_k(\theta)) \mid \theta \in \Theta\}$ contains a non-empty open subset of $\mathbb{R}^k$.

In the one-parameter case, the condition is almost always satisfied, since the image of an interval Θ under a continuous, non-constant η is an interval with more than one point; all the one-parameter examples above are of full rank. For several parameters, the condition can be more interesting: the natural parameter values of the normal distribution fill the open half-plane $\{(\eta_1, \eta_2) \mid \eta_2 < 0\}$ (example 7.9) and the natural parameter values of the gamma distribution fill an open quadrant (exercise 7.3), but the natural parameter values of $N(\theta, \theta^2)$ with $\theta > 0$, where the standard deviation equals the mean, form a curve in $\mathbb{R}^2$ which does not contain an open set and thus the two-parameter family in this case is not of full rank (exercise 7.7).

7.4 What the Structure Buys

The definitions in this lecture have applications in many areas. In lecture 8 we see that the natural statistic T_n of a sample is *sufficient* and lemma 7.5 is essentially the whole proof. In lecture 10 we see that full rank implies that T_n is *complete* and we use the Lehmann–Scheffe theorem to find the best unbiased estimators. In lecture 11 we find that the Fisher information for the natural parameter equals $K''(\eta)$. In lecture 13 we see that the estimators which achieve the Cramer–Rao bound are, up to affine transformations, natural statistics. Finally, in lecture 26 we find that conjugate priors arise because the likelihood (7.2) depends on the data only via T_n and n .

Exponential families also satisfy the regularity conditions from lectures 11, 16 and 17: the support of the distribution is fixed, the density is smooth in θ whenever η and B are, and we have differentiation under the integral sign as in the proof of theorem 7.7. The log-likelihood $\ell(\eta) = \sum_i \log h(x_i) + \eta T_n - nK(\eta)$ has second derivative $-nK''(\eta) = -n \text{Var}_\eta(T(X)) < 0$, since $T(X)$ is not constant, and thus is strictly concave in η . For this reason, the likelihood equation from exercise 7.5 has at most one root and this root is then the maximum likelihood estimator. This shows that the uniqueness hypothesis of the consistency theorem from lecture 16 holds for the whole class.

Summary.

- A one-parameter exponential family has density of the form $f(x; \theta) = h(x) \exp(\eta(\theta)T(x) - B(\theta))$, where the natural statistic is T , the natural parameter is η and the support of the distribution does not depend on θ . Examples of exponential families include the Bernoulli, Poisson, exponential, normal and gamma distributions. The uniform distribution is not an exponential family.
- An i.i.d. sample from an exponential family is again an exponential family. In this case, the natural statistic is $T_n = \sum_i T(X_i)$ and the likelihood depends on the data only through T_n .
- In the natural parametrisation $f(x; \eta) = h(x) \exp(\eta T(x) - K(\eta))$ the cumulant function K can be used to find the moments of the natural statistic: we have $\mathbb{E}_\eta(T(X)) = K'(\eta)$ and $\text{Var}_\eta(T(X)) = K''(\eta)$.
- A family is of full rank, if the values of the natural parameter corresponding to an open set of the correct dimension can be achieved. This condition is required for completeness in [lecture 10](#).
- By construction, exponential families satisfy the regularity conditions of the later theory and the log-likelihood function is strictly concave in the natural parameter.

Exercise 7.1. Let X be geometrically distributed with parameter $p \in (0, 1)$, *i.e.* $f(x; p) = (1 - p)^{x-1}p$ for $x \in \{1, 2, \dots\}$, as in [exercises 4.2](#) and [5.1](#).

1. Show that this distribution is a one-parameter exponential family and identify the functions h , T , η and B . What is the natural statistic of an i.i.d. sample $X_1, \dots, X_n$?
2. Determine the natural parameter space $\mathcal{N}$ and the cumulant function K . Check that $B(p) = K(\eta(p))$. Is this family of full rank?
3. Using [theorem 7.7](#), determine the mean and variance of X and compare the result with [table A.1](#).

Exercise 7.2. Let X be binomially distributed with a known number m of trials and unknown success probability $\theta \in (0, 1)$, *i.e.* $f(x; \theta) = \binom{m}{x} \theta^x (1 - \theta)^{m-x}$ for $x \in \{0, 1, \dots, m\}$.

1. Show that this distribution is a one-parameter exponential family, with the same natural parameter as the Bernoulli distribution from [example 7.2](#), and identify the functions h , T and B .
2. Determine the natural parameter space and the cumulant function.
3. Using [theorem 7.7](#), determine the mean and variance of X , and the mean and variance of the natural statistic T_n for an i.i.d. sample of size n .

Exercise 7.3. Let X be gamma distributed, $X \sim \text{Gamma}(\alpha, \beta)$ with shape $\alpha > 0$ and rate $\beta > 0$, both unknown, *i.e.* with density given by $f(x; \alpha, \beta) = \beta^\alpha x^{\alpha-1} e^{-\beta x} / \Gamma(\alpha)$ for $x > 0$, as given in [table A.2](#).

1. Show that this model is a two-parameter exponential family, as given in definition 7.8, with $h(x) = 1$ for $x > 0$ and $h(x) = 0$ otherwise, natural statistic $(\log x, x)$, natural parameter $(\alpha - 1, -\beta)$, and function $B(\alpha, \beta) = \log \Gamma(\alpha) - \alpha \log \beta$.
2. Show that the set of natural parameter values is the open quadrant

$$\{(\eta_1, \eta_2) \mid \eta_1 > -1, \eta_2 < 0\}$$

and thus the family has full rank as given in definition 7.10.

3. Assume that the shape α is known, as in exercise 5.2. Show that then the factor $x^{\alpha-1}$ can be incorporated into the function h to yield a one-parameter exponential family with natural statistic $T(x) = x$ and natural parameter $-\beta$.

Exercise 7.4.

1. For the exponential distribution with rate θ , use the cumulant function $K(\eta) = -\log(-\eta)$ from the text to find the mean and variance of an observation.
2. For the normal distribution with known variance σ^2 and unknown mean μ , where h , T and η are as in example 7.9, use the method of completing the square to show that $K(\eta) = \sigma^2 \eta^2 / 2$ with $\mathcal{N} = \mathbb{R}$ and use this to derive the mean and variance of an observation. Check that your result satisfies $B(\mu) = K(\eta(\mu))$.

Exercise 7.5. Let $X_1, \dots, X_n$ be an i.i.d. sample from a one-parameter exponential family in its natural parametrisation, with $\mathcal{N}$ an open interval and $\text{Var}_\eta(T(X_1)) > 0$ for all $\eta \in \mathcal{N}$.

1. Show that the likelihood equation $\ell'(\eta) = 0$ from lecture 5 can be written as

$$\frac{1}{n} \sum_{i=1}^n T(x_i) = \mathbb{E}_\eta(T(X_1)),$$

i.e. the maximum likelihood estimator is the parameter value under which the population mean of the natural statistic equals its sample mean.

2. Show that every root of the likelihood equation is the unique maximum likelihood estimator for η .
3. Apply the equation from the first part to the Poisson distribution and to the Bernoulli distribution, and use theorem 5.6 to recover the estimators found in exercise 5.7 and example 5.3.
4. Explain how the first part of this exercise accounts for the coincidences between the method of moments and maximum likelihood which we noted in lecture 4, and why no such coincidence is to be expected for the gamma distribution of exercise 7.3.

Exercise 7.6. Let X be a random variable with density

$$f(x; \eta) = h(x) \exp(\eta T(x) - K(\eta))$$

given by definition 7.6, and let s be such that $\eta + s \in \mathcal{N}$.

1. Show that $\mathbb{E}_\eta(e^{sT(X)}) = \exp(K(\eta + s) - K(\eta))$, i.e. that $M(s) = K(\eta + s) - K(\eta)$ is the logarithm of the moment generating function of $T(X)$.
2. Let η be an interior point of $\mathcal{N}$. Then, since M is differentiable at $s = 0$, we can find $M'(0)$ and $M''(0)$ from the first part of this exercise. Use these results to give an alternative proof of theorem 7.7.
3. Using the first part of the question, show that for the Poisson distribution with parameter θ we have $\mathbb{E}_\theta(z^X) = e^{\theta(z-1)}$ for all $z > 0$. This is the same identity as in exercise 5.7.

Exercise 7.7. Consider the normal distribution $N(\theta, \theta^2)$ with mean and standard deviation both equal to $\theta > 0$.

1. Using example 7.9, show that this model is a two-parameter exponential family with natural statistic and natural parameter given by

$$T(x) = (x, x^2) \quad \text{and} \quad \eta(\theta) = (1/\theta, -1/(2\theta^2)).$$

2. Show that as θ varies, the natural parameter values follow the curve $\eta_2 = -\eta_1^2/2$ with $\eta_1 > 0$. Conclude that this family is not of full rank in the sense of definition 7.10.

Lecture 8

Sufficiency

In example 4.2 we have observed ten coin tosses and got seven heads. The likelihood in this example is $L(\theta) = \theta^7(1 - \theta)^3$. This function depends on the data only via the number 7 of heads, but not on the order of heads and tails, *i.e.* for learning about θ we can replace the full record of observations by this single number. The idea of this example is made precise in this lecture: A *sufficient statistic* contains all information about the parameter. The factorisation theorem helps to find sufficient statistics. The concept of sufficiency forms the basis of improvements of estimators in lecture 10.

In this lecture we write $X = (X_1, \dots, X_n)$ for the data, $x = (x_1, \dots, x_n)$ for an observed value, and

$$f(x; \theta) = \prod_{i=1}^n f(x_i; \theta)$$

for the joint density or joint probability weights of the sample. The vector argument indicates whether we consider the joint function of the sample, or the function f of a single observation. This expression, as a function of θ for fixed data, is the likelihood $L(\theta)$ from lecture 4.

8.1 Sufficient Statistics

A statistic $T = T(X_1, \dots, X_n)$ in the sense of definition 1.3 reduces the data to a smaller summary, and usually some information is lost in the process. The following definition expresses that no information about θ is lost.

Definition 8.1. A statistic $T = T(X_1, \dots, X_n)$ is *sufficient* for θ , if the conditional distribution of $(X_1, \dots, X_n)$ given the value $T = t$ does not depend on θ , for all values t of T .

If T is sufficient, anyone who knows the value t of T can generate a new data set with the same distribution as the original data set, by sampling from the conditional distribution of X given $T = t$. Since this distribution does not depend on θ , all information about θ that we could obtain by studying the data set can now be obtained by studying t alone.

For the coin tosses in the introduction, exercise 8.1 shows that the definition is correct: if the number of heads $T = \sum_{i=1}^n X_i$ is given as t , each of the $\binom{n}{t}$ possible arrangements of heads and tails has conditional probability $1/\binom{n}{t}$, independent

of the value of θ . Thus, the positions of the heads do not contain any information about the coin and T is sufficient.

Two observations help to place the definition. First, the whole sample $T = (X_1, \dots, X_n)$ is always sufficient, since given the whole sample there is no variability left; sufficiency alone therefore says nothing about how much the data have been reduced, and we return to this in section 8.4. Secondly, a one-to-one function $\varphi(T)$ of a sufficient statistic carries the same information and is sufficient as well; for the coin tosses the sample mean $\bar{X} = T/n$ is thus sufficient too.

8.2 The Factorisation Theorem

Directly verifying definition 8.1 requires knowledge of the distribution of T and involves a conditional probability which, for most models, is tedious to compute. The following theorem replaces the computation with an inspection of the joint density.

Theorem 8.2 (Fisher–Neyman factorisation theorem). *The statistic T is sufficient for θ , if and only if there are functions g and h such that*

$$f(x; \theta) = g(T(x); \theta) h(x) \quad (8.1)$$

for all data vectors x and all $\theta \in \Theta$.

The function g may depend on the data only through $T(x)$, and the function h may not depend on θ at all: the parameter and the data interact only through T . Nothing in the statement requires T or θ to be one-dimensional. We give a complete proof for discrete data and then explain what changes in the continuous case.

Proof for discrete data. Since we have discrete observations, the joint probability weight $f(x; \theta) = \mathbb{P}_\theta(X = x)$ is the probability of observing the data vector x .

Assume first that T is sufficient and let x be a data vector with $t = T(x)$. Since the event $\{X = x\}$ is contained in the event $\{T = t\}$, we have

$$\mathbb{P}_\theta(X = x) = \mathbb{P}_\theta(X = x, T = t) = \mathbb{P}_\theta(T = t) \mathbb{P}_\theta(X = x | T = t)$$

whenever $\mathbb{P}_\theta(T = t) > 0$. Using the sufficiency of the statistic, the conditional probability on the right does not depend on θ and thus is a function $h(x)$ of the data. The first factor is a function $g(t; \theta)$ of t and θ . If $\mathbb{P}_\theta(T = t) = 0$ for the given θ , both sides in (8.1) equal zero since $g(t; \theta) = 0$. For the data vectors x with $\mathbb{P}_\theta(T = T(x)) = 0$ for every θ , which are never observed, we define $h(x) = 0$. Thus (8.1) holds for all x and θ .

Conversely, assume that (8.1) holds and let t be a value of T and θ a parameter value with $\mathbb{P}_\theta(T = t) > 0$. Summing the factorisations for all data vectors y with $T(y) = t$, we find

$$\mathbb{P}_\theta(T = t) = \sum_{y: T(y)=t} f(y; \theta) = g(t; \theta) \sum_{y: T(y)=t} h(y) = g(t; \theta) H(t),$$

where $H(t)$ is the last sum which does not depend on θ . Since the left-hand side is positive, we have $g(t; \theta) \neq 0$. For data x with $T(x) = t$ we get

$$\mathbb{P}_\theta(X = x | T = t) = \frac{\mathbb{P}_\theta(X = x)}{\mathbb{P}_\theta(T = t)} = \frac{g(t; \theta) h(x)}{g(t; \theta) H(t)} = \frac{h(x)}{H(t)},$$

and for x with $T(x) \neq t$ the conditional probability is zero. Thus, in both cases the value does not depend on θ and T is sufficient. This completes the proof. $\square$

For continuous data the event $\{T = t\}$ has probability zero and the conditional distribution of X given $T = t$ cannot be described as a quotient of probabilities anymore. If the statistic can be extended to a change of variables, *i.e.* if we can find coordinates U such that $x \mapsto (T(x), U(x))$ is a bijection, then the argument presented above can be applied with densities instead of probabilities: The joint density of (T, U) will factor in a similar way, and after integrating out u and dividing we will obtain a conditional density of U given $T = t$ which does not depend on θ . In general, no such completion of the proof is possible, and the proof will have to replace the family of measures by one dominating probability measure, for example a weighted sum of countably many of the $\mathbb{P}_\theta$, such that the density of $\mathbb{P}_\theta$ is a function of T . We will not go into detail here and will use theorem 8.2 for continuous models in analogy to the proof for discrete models, with the understanding that the factorisation, similar to the density, only needs to hold up to sets of probability zero.

The factorisation has a consequence for maximum likelihood. For fixed data the likelihood is $L(\theta) = g(T(x); \theta)h(x)$, and the factor $h(x)$ does not affect where the maximum over θ is attained: the maximum likelihood estimator, whenever it is unique, is a function of any sufficient statistic. The examples of lecture 5 confirm this: the Bernoulli MLE is T/n , and the uniform MLE is the sample maximum, which exercise 8.2 shows to be sufficient.

8.3 Sufficiency in Exponential Families

For the exponential families from lecture 7, no work is needed to find a sufficient statistic, since the densities are already given in factorised form.

Corollary 8.3. *Let $X_1, \dots, X_n$ be an i.i.d. sample from a one-parameter exponential family as given in definition 7.1 and let T be the natural statistic. Then $T_n = \sum_{i=1}^n T(X_i)$ is sufficient for θ .*

Proof. From lemma 7.5 we know that the joint density of the sample is

$$f(x; \theta) = \left(\prod_{i=1}^n h(x_i) \right) \exp \left(\eta(\theta) \sum_{i=1}^n T(x_i) - nB(\theta) \right).$$

Thus, we can write the joint density in the form (8.1) with $g(t; \theta) = \exp(\eta(\theta)t - nB(\theta))$ and $h(x) = \prod_i h(x_i)$. Consequently, T_n is sufficient by theorem 8.2. $\square$

The same argument shows that for the k -parameter exponential families of lecture 7 the vector $(\sum_i T_1(X_i), \dots, \sum_i T_k(X_i))$ of summed natural statistics is sufficient for the parameter vector. It is instructive to see the factorisation directly in the standard models.

Example 8.4. Let $X_1, \dots, X_n$ be i.i.d. Poisson with parameter $\theta > 0$. Then the joint probability weights are given by

$$f(x; \theta) = \prod_{i=1}^n \frac{\theta^{x_i}}{x_i!} e^{-\theta} = \theta^{\sum_i x_i} e^{-n\theta} \cdot \frac{1}{\prod_i x_i!},$$

which is of the form (8.1) with $g(t; \theta) = \theta^t e^{-n\theta}$ and $h(x) = 1/\prod_i x_i!$. Thus the sum $T = \sum_i X_i$ is sufficient for θ . Similarly, the sample mean $\bar{X}$ is sufficient. For data sets with the same sum, the likelihood functions are proportional: this was already shown in exercise 4.1 for the log-likelihood, and lecture 9 illustrates this in R.

For the normal distribution $N(\mu, \sigma^2)$, where both parameters are unknown, we can find the sufficient statistic by expanding the squares in the exponent of the density. It transpires that the joint density depends on the data only via the pair $T = (\sum_i X_i, \sum_i X_i^2)$. Thus, T is a sufficient statistic for (μ, σ^2) . Since $\sum_i (x_i - \bar{x})^2 = \sum_i x_i^2 - n\bar{x}^2$, the pair $(\bar{X}, S^2)$ maps to T one-to-one and thus is also a sufficient statistic. This is the pair of statistics on which the maximum likelihood estimator from example 5.4 is based. If σ^2 is known, the factor $\exp(-\sum_i x_i^2/(2\sigma^2))$ can be moved into $h(x)$ and then $\bar{X}$ alone suffices as a sufficient statistic for μ . The factorisation is worked out in exercise 8.4.

8.4 Minimal Sufficiency and Completeness

As we have seen in section 8.1, the whole sample is always sufficient. For an i.i.d. sample, the vector of ordered observations is also sufficient, since the joint density does not change when the x_i are interchanged. The aim of this section is to characterise the statistic which allows us to reduce the data as much as possible, while still retaining all the information.

Definition 8.5. A sufficient statistic T is *minimal sufficient* for θ , if for every sufficient statistic S there is a function φ such that $T = \varphi(S)$.

A minimal sufficient statistic can thus be computed from every other sufficient statistic. The definition is awkward to check directly, because it refers to all sufficient statistics at once; the following criterion instead compares the likelihood functions of two data sets.

Theorem 8.6. Let T be a statistic, such that every data vector has positive likelihood for some $\theta \in \Theta$, and such that for all data vectors x and y we have

$$\frac{f(x; \theta)}{f(y; \theta)} \text{ does not depend on } \theta \quad \text{if and only if} \quad T(x) = T(y).$$

Then T is minimal sufficient for θ .

The proof of this theorem is only sketched here. The statement of sufficiency follows from theorem 8.2: For every value t of T we can choose a fixed data vector x_t such that $T(x_t) = t$. Then the ratio $f(x; \theta)/f(x_{T(x)}; \theta)$, by assumption, does not depend on θ and can be used as $h(x)$ while $g(t; \theta) = f(x_t; \theta)$. To see minimality, let S be any sufficient statistic. Then we can write f as $f(x; \theta) = \tilde{g}(S(x); \theta) \tilde{h}(x)$. If $S(x) = S(y)$, then $f(x; \theta)/f(y; \theta) = \tilde{h}(x)/\tilde{h}(y)$ does not depend on θ and thus we have $T(x) = T(y)$ by assumption: the value of T is already determined by the value of S and T is a function of S .

The criterion states that $T(x)$ contains the likelihood function of the data up to a constant factor. We apply the criterion once, to the Poisson sum. The two-parameter normal case is worked out in exercise 8.4.

Example 8.7. Let $X_1, \dots, X_n$ be i.i.d. Poisson with parameter $\theta > 0$, and let $T = \sum_i X_i$. For two data vectors x and y , the joint probability weights of example 8.4 give

$$\frac{f(x; \theta)}{f(y; \theta)} = \frac{\theta^{\sum_i x_i} e^{-n\theta} / \prod_i x_i!}{\theta^{\sum_i y_i} e^{-n\theta} / \prod_i y_i!} = \theta^{\sum_i x_i - \sum_i y_i} \frac{\prod_i y_i!}{\prod_i x_i!}.$$

If $\sum_i x_i = \sum_i y_i$, the power of θ disappears and the ratio is a constant. If the sums differ by $k \neq 0$, the ratio contains the factor θ^k , which is not constant as θ ranges over $(0, \infty)$. Thus the ratio is free of θ if and only if $\sum_i x_i = \sum_i y_i$, and T is minimal sufficient by theorem 8.6.

For the results of lecture 10 we will need one more property of a statistic, concerning functions of the statistic which are unbiased estimators for zero.

Definition 8.8. A statistic T is *complete*, if for every function g with finite mean $g(T)$ under $\mathbb{P}_\theta$ for all $\theta \in \Theta$, the condition $\mathbb{E}_\theta(g(T)) = 0$ for all $\theta \in \Theta$ implies $\mathbb{P}_\theta(g(T) = 0) = 1$ for all $\theta \in \Theta$.

This means that the only function of a complete statistic which has expectation zero for all parameter values is the zero function. The condition is violated for statistics which do not reduce the data enough: in exercise 8.6 we will see that the whole Bernoulli sample is sufficient, but not complete, whereas the total $\sum_i X_i$ is complete, as we will see in lecture 10 (proposition 10.5). The property is important because any two unbiased estimators for the same quantity, which are both functions of a complete statistic T , differ by a function of T with expectation zero, and by completeness this function must be zero with probability one: for each quantity there is at most one unbiased estimator which is a function of a complete statistic. In lecture 10 we will use this uniqueness to turn the Rao–Blackwell improvement into an optimality result. In the lecture we will also discuss which of the standard statistics are complete.

8.5 The Sample Maximum and Minimum

So far we have considered sufficient statistics which are sums, and we have seen the distributions of these sums in appendix A. For the uniform distribution we will consider the sample maximum, and in order to determine the bias or the mean squared error of the maximum, we need to know the distribution of the maximum. The following result gives the distribution of the maximum and minimum of an i.i.d. sample.

Proposition 8.9. Let $X_1, \dots, X_n$ be i.i.d. with distribution function F and density f . Furthermore, let $X_{(n)} = \max_i X_i$ and $X_{(1)} = \min_i X_i$. Then the random variables $X_{(n)}$ and $X_{(1)}$ have densities

$$f_{(n)}(y) = n F(y)^{n-1} f(y) \quad \text{and} \quad f_{(1)}(y) = n (1 - F(y))^{n-1} f(y).$$

Proof. The maximum is at most y , if and only if all observations are at most y . Using independence we get

$$\mathbb{P}(X_{(n)} \leq y) = \mathbb{P}(X_1 \leq y, \dots, X_n \leq y) = \prod_{i=1}^n \mathbb{P}(X_i \leq y) = F(y)^n.$$

Using the chain rule for differentiation with respect to y , we find that the maximum has density $nF(y)^{n-1}f(y)$. Similarly, the minimum is strictly greater than y , if and only if all observations are strictly greater than y . Thus we get

$$\mathbb{P}(X_{(1)} \leq y) = 1 - \mathbb{P}(X_1 > y, \dots, X_n > y) = 1 - (1 - F(y))^n.$$

Using the chain rule for differentiation we find $n(1 - F(y))^{n-1}f(y)$. This completes the proof. $\square$

Example 8.10. Let $X_1, \dots, X_n \sim \text{Uniform}(0, \theta)$ be i.i.d., as in example 1.2. The distribution function is $F(y) = y/\theta$ for $0 \leq y \leq \theta$, and the density is $f(y) = 1/\theta$ on $[0, \theta]$. By proposition 8.9 the sample maximum has density

$$f_{(n)}(y) = n\left(\frac{y}{\theta}\right)^{n-1} \frac{1}{\theta} = \frac{ny^{n-1}}{\theta^n}, \quad 0 \leq y \leq \theta,$$

and $f_{(n)}(y) = 0$ outside this interval. This is the density supplied without proof in exercise 2.2, from which we computed $\mathbb{E}_\theta(X_{(n)}) = n\theta/(n+1)$ there. The support of the maximum moves with θ , as the support of a single observation does, and exercise 8.2 shows that the maximum is sufficient for θ , applying the factorisation theorem with some care about the support.

The minimum plays the corresponding role for models where the density is largest at the left-hand end of the support. In exercise 8.5 we see that the minimum of an exponential sample is itself exponentially distributed, but with rate multiplied by n .

Summary.

- A statistic T is sufficient for θ , if the conditional distribution of the data given T does not depend on θ : the data can be reduced to T without losing information about θ .
- The factorisation theorem states that T is sufficient if and only if $f(x; \theta) = g(T(x); \theta) h(x)$. Using this result, we can see that a unique MLE is always a function of any sufficient statistic.
- For exponential families, the summed natural statistic $\sum_i T(X_i)$ is sufficient. This includes the Bernoulli total, the Poisson sum and the normal pair $(\sum_i X_i, \sum_i X_i^2)$.
- A minimal sufficient statistic is a function of any sufficient statistic, and satisfies the condition $T(x) = T(y)$ if and only if the likelihood ratio $f(x; \theta)/f(y; \theta)$ does not depend on θ . A complete statistic is any statistic such that the only function of the statistic which has expectation zero for all θ is the zero function. In lecture 10 we will use this property to identify the unique best unbiased estimator.
- The densities of the sample maximum and minimum are $nF^{n-1}f$ and $n(1 - F)^{n-1}f$, obtained by differentiating F^n and $1 - (1 - F)^n$, respectively.

Exercise 8.1. Let $X_1, \dots, X_n$ be i.i.d. Bernoulli with success probability θ and let $T = \sum_{i=1}^n X_i$ be the number of successes. From the remark about sums in appendix A we know that $T \sim \text{Binomial}(n, \theta)$.

1. Let $t \in \{0, 1, \dots, n\}$ and let $x \in \{0, 1\}^n$ be a data vector with $\sum_i x_i = t$. Show that $\mathbb{P}_\theta(X = x | T = t) = 1/\binom{n}{t}$.
2. What is the conditional probability, if $\sum_i x_i \neq t$? Using definition 8.1, show that T is sufficient for θ .
3. Using the factorisation theorem 8.2, show that the result in the previous part is correct.

Exercise 8.2. Let $X_1, \dots, X_n \sim \text{Uniform}(0, \theta)$ be i.i.d. with $\theta > 0$.

1. Using the indicator notation $\mathbf{1}_{\{0 \leq x \leq \theta\}}$ for the density of a single observation, show that the joint density can be written as

$$f(x; \theta) = \theta^{-n} \mathbf{1}_{\{\max_i x_i \leq \theta\}} \mathbf{1}_{\{\min_i x_i \geq 0\}}.$$

2. Deduce from theorem 8.2 that the sample maximum $X_{(n)} = \max_i X_i$ is sufficient for θ . Identify the functions g and h .
3. Why can the factor with $\min_i x_i$ be included in h , but the factor with $\max_i x_i$ cannot?

Exercise 8.3. Let $X_1, \dots, X_n$ be i.i.d. exponential with rate $\theta > 0$, i.e. $f(x; \theta) = \theta e^{-\theta x}$ for $x \geq 0$.

1. Show that $T = \sum_i X_i$ is sufficient for θ .
2. Using theorem 8.6, show that T is minimal sufficient.
3. In example 5.2 we have found the MLE $\hat{\theta} = 1/\bar{X}$. Write $\hat{\theta}$ as a function of T and explain why the factorisation in the first part of the question guarantees that such a function exists.

Exercise 8.4. Let $X_1, \dots, X_n \sim N(\mu, \sigma^2)$ be i.i.d.

1. Assume that μ is known and σ^2 is unknown. Show that $\sum_i (X_i - \mu)^2$ is sufficient for σ^2 .
2. Assume that both parameters are unknown. By expanding the squares in the exponent of the joint density, show that the joint density depends on the data only via the pair $T = (\sum_i X_i, \sum_i X_i^2)$. Using theorem 8.2, show that T is sufficient for (μ, σ^2) .
3. Using theorem 8.6, show that the pair T from the previous part is minimal sufficient for (μ, σ^2) .
4. Is the sample mean $\bar{X}$ on its own sufficient for (μ, σ^2) when both parameters are unknown? Justify your answer using the previous part.

Exercise 8.5. Let $X_1, \dots, X_n$ be i.i.d. exponential with rate $\theta > 0$, and let $X_{(1)} = \min_i X_i$.

1. Show that $X_{(1)}$ is exponentially distributed with rate $n\theta$, first by proposition 8.9 and then directly from $\mathbb{P}_\theta(X_{(1)} > y)$.
2. Show that $nX_{(1)}$ is an unbiased estimator for the mean $1/\theta$, and compute its variance.
3. Compare the variance of $nX_{(1)}$ with that of the sample mean $\bar{X}$, which is also unbiased for $1/\theta$, and comment on the behaviour of the two estimators as n grows. Is $nX_{(1)}$ a function of the sufficient statistic $\sum_i X_i$?

Exercise 8.6. Let $X_1, \dots, X_n$ be i.i.d. Bernoulli with success probability $\theta \in (0, 1)$, where $n \geq 2$, and let $T = (X_1, \dots, X_n)$ be the whole sample.

1. Show that $g(T) = X_1 - X_2$ satisfies $\mathbb{E}_\theta(g(T)) = 0$ for all θ and $\mathbb{P}_\theta(g(T) = 0) < 1$.
2. Using definition 8.8, show that T is not complete, despite T being sufficient.

Lecture 9

Consistency by Simulation

This is the R session for week 3. In the previous two sessions we were promised that we would watch an estimator converge to the truth, *i.e.* the consistency from lecture 2. Today we will see this convergence in two different ways: we will watch an estimate converge to the truth as observations are added, and we will plot the mean squared error against n . On the way there we will extend the Monte Carlo skills we acquired in lecture 6 by introducing a loop which varies the sample size and a matrix to store the results. We will conclude by doing a computation which shows what lecture 8 means by a sufficient statistic. As before, appendix B contains a written reference for this topic. We will predict the outcome of the simulations before we start each one.

9.1 A Running Estimate

Consistency, definition 2.8, is a statement about the sequence of estimators $\hat{\theta}_1, \hat{\theta}_2, \dots$ as the sample size increases. The most direct way to visualise this sequence is to simulate a long sample and to re-compute the estimate after each observation. We will consider the exponential distribution with rate $\theta = 2$, where the maximum likelihood estimator for the rate is $\hat{\theta}_n = 1/\bar{X}_n$, as found in example 5.2, and we will see that this estimator is consistent for θ , by theorems 2.9 and A.17.

The command `cumsum` can be used to compute the running totals X_1, X_1+X_2 and so on in one step, and dividing the count n by these totals gives all 2000 running estimates at once. Before you look at the output, predict the five printed numbers. How far from 2 do you expect the estimate after one observation to be? How far after 2000?

```
theta <- 2
N      <- 2000
set.seed(2026)
x      <- rexp(N, rate = theta)
n      <- seq_len(N)
running <- n / cumsum(x)
running[c(1, 10, 100, 1000, 2000)]
```

```
## [1] 5.033385 1.446844 1.770152 1.921852 1.939830
```

The first estimate is $1/x_1$, which could be any number; in this example it is more than twice the truth. By $n = 1000$ the estimate is within 0.1 of the truth. To get a better understanding of the path taken by the estimate, we plot the estimate against n . The argument `log = "x"` for the `plot` command tells R to plot n on a logarithmic axis, so that we can still see the early part of the path. Before we look at the plot, try to sketch the path you expect the estimate to take. Where is the curve most volatile? Once the curve enters the band $\theta \pm 0.1$, does it stay there?

```
plot(n, running, type = "l", log = "x", ylim = c(1, 3.5),
     xlab = "n", ylab = "running estimate")
abline(h = theta, lty = 2)
abline(h = theta + c(-0.1, 0.1), lty = 3)
```




Figure 9.1: The running estimate $1/\bar{X}_n$ for the exponential rate $\theta = 2$, plotted against n on a logarithmic axis. The truth is given by the dashed line, the band $\theta \pm 0.1$ is given by the dotted lines.

Figure 9.1 shows the expected behaviour: the path is wild for small n and ends inside the band. It does not converge in the sense of a sequence of numbers where, after some point, no term is outside the band; the path leaves and re-enters the band at random times. What decreases is the probability of being outside the band at a given n , as given in definition 2.8.

9.2 The Mean Squared Error as a Function of the Sample Size

In lecture 6 we estimated the mean squared error of the two uniform estimators $2\bar{X}$ and $\max_i X_i$ for $n = 10$. To see how each decreases with n we repeat the study for several sample sizes, and this is the everyday shape of a simulation study in the practical report: an outer loop over a setting, an inner Monte Carlo run and a table of results. We first put both estimates into one function, so that both are computed from the same sample and returned as a vector of length two. The command `replicate` collects the 10000 repetitions into a *matrix*, a

rectangular table of numbers with two rows, one per estimator, and the command `rowMeans` can be used to average each row.

```
est_both <- function(n, theta) {
  sample <- runif(n, min = 0, max = theta)
  c(2 * mean(sample), max(sample))
}
```

The outer loop runs over a vector `ns` of sample sizes, using `seq_along(ns)` to count through its positions. The command `matrix(0, nrow = length(ns), ncol = 2)` can be used to set aside a matrix for the results, `mse[j, ]` refers to the j -th row of this matrix, and row and column names, set using `rownames` and `colnames`, can be used to make the matrix a readable table. As in lecture 6, the box below shows the complete script, with the lines of the loop kept together for clarity, but with the lines in a shuffled order. Before you read the script, you should try to work out the correct order of lines to make the script work.

```
signif(mse, 3)
for (j in seq_along(ns)) {
  ests <- replicate(10000, est_both(ns[j], theta))
  mse[j, ] <- rowMeans((ests - theta)^2)
}
mse <- matrix(0, nrow = length(ns), ncol = 2)
set.seed(2026)
rownames(mse) <- ns
theta <- 1
colnames(mse) <- c("2 * mean", "max")
ns <- c(5, 10, 20, 50, 100, 200, 500, 1000)
```

The rule for writing R scripts is still in force: all names must be defined before they are used, *i.e.* `theta` and `ns` must come first, the seed of the random number generator must be set before the loop starts, and the names and the display (using the `signif` function to round the numbers) require the matrix to be complete. Before you try the reconstructed script, try to predict the last row of the output. What will the two mean squared errors be for $n = 1000$? For $n = 10$ the second mean squared error was approximately half of the first one. Will the ratio between the two mean squared errors stay close to two, or will it change?

```
theta <- 1
ns <- c(5, 10, 20, 50, 100, 200, 500, 1000)
mse <- matrix(0, nrow = length(ns), ncol = 2)
set.seed(2026)
for (j in seq_along(ns)) {
  ests <- replicate(10000, est_both(ns[j], theta))
  mse[j, ] <- rowMeans((ests - theta)^2)
}
rownames(mse) <- ns
colnames(mse) <- c("2 * mean", "max")
signif(mse, 3)
```

```
##      2 * mean      max
```

```
## 5    0.065900 4.58e-02
## 10   0.033300 1.52e-02
## 20   0.016600 4.48e-03
## 50   0.006610 7.51e-04
## 100  0.003310 1.91e-04
## 200  0.001660 5.07e-05
## 500  0.000668 7.86e-06
## 1000 0.000331 1.99e-06
```

The row for $n = 10$ coincides with the two numbers from lecture 6 up to simulation error (the samples are not the same, since the loop has already sampled the $n = 5$ values first), and the values for each column decrease as n increases, as expected from the sufficient condition for consistency given in corollary 2.11. The ratio between the mean squared errors is not close to two anymore: for $n = 1000$ the mean squared error for the maximum is smaller by a factor of approximately 170. The exact values from exercises 2.2 and 4.4 show that this is expected, since $\theta^2/(3n)$ decreases like $1/n$ while $2\theta^2/((n+1)(n+2))$ decreases like $1/n^2$. If we plot the data on log-log axes, *i.e.* if we use the argument `log = "xy"` for the plot, then a quantity proportional to n^{-k} can be plotted as a straight line with slope $-k$. What will the two sets of points look like on a log-log plot? Will the two sets of points be parallel?

```
plot(ns, mse[, 1], log = "xy", pch = 16, xlab = "n", ylab = "MSE",
     ylim = range(mse))
points(ns, mse[, 2], pch = 1)
lines(ns, theta^2 / (3 * ns), lty = 2)
lines(ns, 2 * theta^2 / ((ns + 1) * (ns + 2)), lty = 3)
legend("bottomleft", legend = c("2 * mean", "max"), pch = c(16, 1))
```




Figure 9.2: Simulated mean squared errors for $2\bar{X}$ (filled circles) and $\max_i X_i$ (open circles), for the uniform parameter $\theta = 1$, on logarithmic axes. The exact values $\theta^2/(3n)$ (dashed line) and $2\theta^2/((n+1)(n+2))$ (dotted line) are also shown.

Both sets of points lie on straight lines: the slope is -1 for $2\bar{X}$ and the slope is -2 for the maximum. The usual rate of decrease for a mean squared error is like $1/n$. All the estimators we will consider in the R strand of this module will

have slope -1 . The faster rate for the maximum is a consequence of the support of the uniform distribution moving with θ . We will consider this example again in lecture 13.

We can also consider the full sampling distribution of the maximum at $n = 50$, since in lecture 8 we know the exact density: by proposition 8.9 and example 8.10, the maximum of n observations from $\text{Uniform}(0, 1)$ has density $f(y) = ny^{n-1}$ for $0 \leq y \leq 1$. To plot this density on top of a histogram, we need to plot the histogram using the density scale. This can be achieved by using the argument `freq = FALSE` of `hist`. Then the command `curve` can be used to plot any function of `x` on the horizontal axis, on top of the existing plot, if `add = TRUE` is used. Sketch the density $50y^{49}$ before the plot appears. Where is the density largest? What is the density at $y = 1$?

```
set.seed(2026)
mle50 <- replicate(10000, max(runif(50, min = 0, max = 1)))
hist(mle50, freq = FALSE, breaks = 40, col = "grey85", main = "",
     xlab = "estimate")
curve(50 * x^49, add = TRUE, lwd = 2)
```

The curve closely follows the tops of the bars, and reaches the value 50 at the boundary of the interval. The histogram is not a bell-shaped distribution as in figure 6.1: a sampling distribution can be concentrated around the truth, without being normal. In figure B.1 in appendix B you can see the same construction for a normal sample.

9.3 A Sufficient Statistic in R

To conclude these notes, we consider an example for use in lecture 8: for Poisson samples, the total number of counts $\sum_i X_i$ is sufficient by example 8.4. The factorisation theorem 8.2 states that for data sets with the same total count, the likelihoods are proportional, differing by a constant shift on the log scale. We consider three data sets of size $n = 5$, where the first two have the same total count. We compute the log-likelihood for these data sets using the `dpois` function in R, with the `log = TRUE` option to get the log-likelihoods, and without any manual simplification of the results. The `sapply` command can be used to evaluate a function on every element of a vector; here we use it to evaluate the log-likelihoods on a grid of θ -values. Will the difference between the first two curves be the same at every grid point? Will the same be true for the first and third curves? Do the three maxima coincide?

```
x1 <- c(2, 0, 3, 1, 4)
x2 <- c(2, 2, 2, 2, 2)
x3 <- c(0, 1, 0, 2, 1)
c(sum(x1), sum(x2), sum(x3))
loglik <- function(theta, data) {
  sum(dpois(data, lambda = theta, log = TRUE))
}
grid <- seq(0.5, 4, by = 0.01)
l11 <- sapply(grid, loglik, data = x1)
l12 <- sapply(grid, loglik, data = x2)
l13 <- sapply(grid, loglik, data = x3)
```

```
range(l11 - l12)
range(l11 - l13)
grid[c(which.max(l11), which.max(l12), which.max(l13))]
```

```
## [1] 10 10 4
## [1] -2.197225 -2.197225
## [1] -9.128696 3.347953
## [1] 2.0 2.0 0.8
```

The difference between the first two curves is constant at every grid point, taking the value -2.197 . This is equal to $\log(32/288)$, the logarithm of the ratio of the two factors $1/\prod_i x_i!$, which do not depend on θ . The difference to the third curve varies between -9.1 and 3.3 . The first two maxima are both at $\bar{x} = 2$ while the third maximum is at 0.8 . The likelihood does not allow us to infer anything about θ beyond what is already known from the total count: for $\mathbf{x}1$ and $\mathbf{x}2$, the information about θ is the same, once the total count is known.

9.4 A Task to Take Home

The following task is a review of the topics of this lecture, in the style of a practical report. Following the structure of section 9.2, for the exponential rate with $\theta = 2$ and sample sizes given in `ns`, estimate the bias and mean squared error of the maximum likelihood estimator $1/\bar{X}$ using 10000 repetitions each, in a table with named rows and columns. Create a plot of the mean squared error against n , using logarithmic axes. From the plot, determine the slope. Compare the empirical bias to the exact value from exercise 2.3. Write three sentences about what the table and plot show about the consistency of the estimator $1/\bar{X}$.

Just like it is important to understand the syntax of a programming language, it is also important to be able to read the output of R commands. Before the next session, type in the commands from this session, and check that the output matches what is shown in the text.

Summary.

- Appendix B is the written reference for `cumsum`, the matrices returned by `replicate` and `curve`.
- The command `cumsum` can be used to compute all running estimates of a growing sample at once. Plotted against n , the path keeps leaving and re-entering a band around the truth; what decreases is the probability of being outside it.
- A loop over a vector of sample sizes, with the results of each Monte Carlo run stored in one row of a matrix, gives the mean squared error as a function of n . On log-log axes a mean squared error proportional to n^{-k} is a straight line with slope $-k$: for the uniform parameter, $2\bar{X}$ shows slope -1 , the usual rate, and $\max_i X_i$ slope -2 .
- With the argument `freq = FALSE` a histogram is drawn on the density scale, and `curve(..., add = TRUE)` can be used to superimpose a theoretical density on top of it.
- Two Poisson data sets with the same total have log-likelihoods which differ by a constant: the total is sufficient.

Lecture 10

Improving Estimators: Rao–Blackwell and Lehmann–Scheffe

In lecture 8 we have seen that a sufficient statistic T can be used to summarise all the information in the data about the parameter. In this lecture we will use this idea to construct improved versions of estimators. The Rao–Blackwell theorem states that conditioning an unbiased estimator on the value of T never increases its variance. The Lehmann–Scheffe theorem states that, if T is a complete sufficient statistic, this procedure gives the unique best unbiased estimator. Finally, Basu’s theorem gives the independence of $\bar{X}$ and S^2 for normal samples and this result will be used in lectures 25 and 29. The proofs of these three theorems all use facts about conditional expectations, as shown in propositions A.9 and A.11 in appendix A.

10.1 The Rao–Blackwell Theorem

Assume that we want to estimate a function $g(\theta)$ of the parameter, *e.g.* we could be interested in the probability $\mathbb{P}_\theta(X_1 = 0)$ of zero counts. As in lecture 2, an estimator U is *unbiased* for $g(\theta)$, if $\mathbb{E}_\theta(U) = g(\theta)$ for all $\theta \in \Theta$. While unbiased estimators are often easy to find, they often turn out to be wasteful, since they only use a small amount of the available sample. The Rao–Blackwell theorem uses a sufficient statistic T to improve an existing estimator by replacing U with the conditional expectation $V = \mathbb{E}_\theta(U|T)$, *i.e.* the average of U for all samples which have the same value of T as the observed sample.

Theorem 10.1 (Rao–Blackwell). *Let T be a sufficient statistic for θ and let U be an unbiased estimator for $g(\theta)$ with $\text{Var}_\theta(U) < \infty$ for all θ . Then $V = \mathbb{E}_\theta(U|T)$ satisfies the following properties.*

1. *The random variable V does not depend on θ , but is a function of T alone and thus a statistic.*
2. *The statistic V is an unbiased estimator for $g(\theta)$.*
3. *For all $\theta \in \Theta$ we have $\text{Var}_\theta(V) \leq \text{Var}_\theta(U)$, with equality if and only if $\mathbb{P}_\theta(U = V) = 1$.*

Proof. From definition 8.1 we know that the conditional distribution of the sample $(X_1, \dots, X_n)$ given T does not depend on θ . Since U is a function of the sample, the conditional distribution and thus the conditional expectation given T do not depend on θ . Thus, V is a function of T and by definition 1.3 a statistic.

For the second statement, using the tower property from proposition A.9, we find

$$\mathbb{E}_\theta(V) = \mathbb{E}_\theta(\mathbb{E}_\theta(U|T)) = \mathbb{E}_\theta(U) = g(\theta)$$

for all θ . This shows that V is unbiased.

Finally, for the third statement, we can use the law of total variance, proposition A.11, with U instead of X and T instead of Y . Since $\mathbb{E}_\theta(U|T) = V$ we get

$$\text{Var}_\theta(U) = \mathbb{E}_\theta(\text{Var}_\theta(U|T)) + \text{Var}_\theta(V).$$

Since the conditional variance $\text{Var}_\theta(U|T)$ is non-negative, we find $\text{Var}_\theta(V) \leq \text{Var}_\theta(U)$. Equality holds, if and only if $\mathbb{E}_\theta(\text{Var}_\theta(U|T)) = 0$ and since a non-negative random variable with expectation zero is zero with probability one, this is the case if and only if $\text{Var}_\theta(U|T) = 0$ with probability one, *i.e.* if and only if U coincides with its conditional expectation V with probability one. This completes the proof. $\square$

Since both estimators are unbiased, the mean squared errors equal the variances, as given in theorem 2.7, and thus V is at least as good as U , and strictly better unless U was already a function of T . The procedure of passing from U to V is called *Rao-Blackwellisation*. The necessity of sufficiency is highlighted by the fact that conditioning on a statistic which is not sufficient in general will result in a quantity which depends on θ but is not an estimator.

Example 10.2. Let $X_1, \dots, X_n$ be i.i.d. Poisson with parameter θ and assume that we want to estimate $g(\theta) = \mathbb{P}_\theta(X_1 = 0) = e^{-\theta}$. The indicator $U = \mathbf{1}_{\{X_1=0\}}$ is unbiased for $g(\theta)$, since $\mathbb{E}_\theta(U) = \mathbb{P}_\theta(X_1 = 0)$. By corollary 8.3 the sum $T = \sum_i X_i$ is sufficient. Given $T = t$, the first observation is distributed as Binomial($t, 1/n$) and thus we have

$$V = \mathbb{E}_\theta(U|T) = \mathbb{P}_\theta(X_1 = 0|T) = \left(1 - \frac{1}{n}\right)^T.$$

The Rao-Blackwellised estimator now uses all n observations. Exercises 10.1 and 10.2 show how to find the conditional distribution and how to apply the same procedure to a continuous model.

10.2 Completeness and the Lehmann-Scheffe Theorem

The Rao-Blackwell theorem improves an existing estimator, but two different unbiased estimators could in principle be improved to two different functions of T . The concept of completeness, definition 8.8, solves this problem: if T is complete, there can only be one unbiased estimator for $g(\theta)$ among the functions of T and thus all Rao-Blackwellisations will converge to this estimator. We start our discussion by naming the target.

Definition 10.3. An unbiased estimator V for $g(\theta)$ with $\text{Var}_\theta(V) < \infty$ for all $\theta \in \Theta$ is a *uniformly minimum variance unbiased estimator* (UMVUE) for $g(\theta)$, if $\text{Var}_\theta(V) \leq \text{Var}_\theta(W)$ for all $\theta \in \Theta$ and for every unbiased estimator W for $g(\theta)$ with finite variance.

The word “uniformly” refers to the parameter: the inequality has to hold for every θ at once, and there is no reason a priori why such an estimator should exist. The following theorem shows that, whenever a complete sufficient statistic exists, a UMVUE exists as well and is easy to identify.

Theorem 10.4 (Lehmann–Scheffé). *Let T be a complete sufficient statistic for θ and let $V = \varphi(T)$ be a function of T , an unbiased estimator for $g(\theta)$ with $\text{Var}_\theta(V) < \infty$ for all θ . Then V is a UMVUE for $g(\theta)$. Furthermore, V is the unique UMVUE: any other unbiased estimator W for $g(\theta)$ with $\text{Var}_\theta(W) = \text{Var}_\theta(V)$ for all θ satisfies $\mathbb{P}_\theta(W = V) = 1$ for all θ .*

Proof. To show the uniqueness of V as an unbiased estimator for $g(\theta)$ which is a function of T , let $\psi(T)$ be another such function. Then we have $\mathbb{E}_\theta(\psi(T)) = g(\theta) = \mathbb{E}_\theta(\varphi(T))$ for all θ . The function $h = \psi - \varphi$ satisfies $\mathbb{E}_\theta(h(T)) = 0$ for all θ and, since T is complete, by definition 8.8, we have $\mathbb{P}_\theta(h(T) = 0) = 1$ for all θ , i.e. $\psi(T) = \varphi(T)$ with probability one.

Now let W be any unbiased estimator for $g(\theta)$ with finite variance. Then, by the Rao–Blackwell theorem, theorem 10.1, the random variable $\mathbb{E}_\theta(W|T)$ is a function of T which is unbiased for $g(\theta)$. Thus, by the first step of the proof, it equals V with probability one. By the Rao–Blackwell theorem we have then

$$\text{Var}_\theta(V) = \text{Var}_\theta(\mathbb{E}_\theta(W|T)) \leq \text{Var}_\theta(W)$$

for all θ , i.e. V is a UMVUE. If $\text{Var}_\theta(W) = \text{Var}_\theta(V)$ for all θ , then the equality case of the Rao–Blackwell theorem implies $W = \mathbb{E}_\theta(W|T) = V$ with probability one. This completes the proof. $\square$

The theorem turns the search for a best unbiased estimator into a two-step recipe: find a complete sufficient statistic T , then find any function of T which is unbiased for $g(\theta)$, by adjusting a constant or by Rao–Blackwellising a crude unbiased estimator. For the models of this module the following two facts settle the completeness.

Proposition 10.5.

1. Let $X_1, \dots, X_n$ be an i.i.d. sample from an exponential family of full rank, as given in definitions 7.1 and 7.10. Then the corresponding natural statistic $T_n = \sum_{i=1}^n T(X_i)$ is complete.
2. Let $X_1, \dots, X_n$ be i.i.d. and uniformly distributed on the set $(0, \theta)$ where $\theta > 0$. Then the sample maximum $\max_i X_i$ is complete.

This proposition is given without proof. The first statement rests on the uniqueness theorem for Laplace transforms. The second statement rests on taking derivatives of the relation $\int_0^\theta h(t) t^{n-1} dt = 0$ with respect to θ . Both statements are allowed to be quoted. Together with corollary 8.3 and exercise 8.2 they form a complete sufficient statistic for all families of the form given in tables 1.1 and 1.2 (with m known for the binomial). If we use a one-to-one function of a complete sufficient statistic, the new function is again complete and sufficient. Thus we can use the function $\bar{X}$ instead of $\sum_i X_i$ if required.

Example 10.6. For the Poisson sample from example 10.2, the natural parameter $\eta = \log \theta$ takes values on all of $\mathbb{R}$ and thus the sum $T = \sum_i X_i$ is a complete and sufficient statistic. Since $\bar{X} = T/n$ is a function of T and is unbiased for θ , it is the UMVUE for θ by theorem 10.4. Similarly, the Rao-Blackwellised indicator $V = (1 - 1/n)^T$ is the UMVUE for $e^{-\theta}$.

Example 10.7. Let $X_1, \dots, X_n$ be i.i.d. $N(\mu, \sigma^2)$ with both parameters unknown. This is a two-parameter exponential family of full rank with natural statistic $T(x) = (x, x^2)$, and thus the pair $(\bar{X}, S^2)$ is complete and sufficient; exercise 10.4 carries out the verification. Since $\bar{X}$ and S^2 are unbiased for μ and σ^2 by lemmas 2.2 and 2.3, they are the UMVUEs for these parameters.

The second example puts section 2.3 into perspective: there the biased estimator $\frac{1}{n+1} \sum_i (X_i - \bar{X})^2$ had smaller mean squared error than S^2 . No contradiction, since the Lehmann-Scheffe theorem compares S^2 to unbiased estimators only.

10.3 Ancillary Statistics and Basu's Theorem

A sufficient statistic contains all the information about θ . At the opposite extreme are statistics which contain none.

Definition 10.8. A statistic $A = A(X_1, \dots, X_n)$ is *ancillary* for θ , if the distribution of A under $\mathbb{P}_\theta$ does not depend on $\theta \in \Theta$.

The running example is a normal sample $X_1, \dots, X_n \sim N(\mu, \sigma^2)$ with known variance σ^2 and unknown mean μ . Writing $Z_i = X_i - \mu$ we have $X_i - \bar{X} = Z_i - \bar{Z}$ and the Z_i are i.i.d. $N(0, \sigma^2)$, for every value of μ . Thus, the vector of residuals $(X_1 - \bar{X}, \dots, X_n - \bar{X})$ is ancillary for μ and every function of the residuals is ancillary, *e.g.* the sample variance S^2 and the range $\max_i X_i - \min_i X_i$; we will consider the range in exercise 10.5.

One might expect an ancillary statistic to have nothing to do with a sufficient one. Under completeness this becomes a theorem, and it is a surprisingly effective way of proving independence.

Theorem 10.9 (Basu). *Let T be a complete sufficient statistic for θ and let A be an ancillary statistic. Then T and A are independent of each other under $\mathbb{P}_\theta$ for all $\theta \in \Theta$.*

Proof. Let B be a set of possible values of A . Since A is ancillary, the probability $p_B = \mathbb{P}_\theta(A \in B)$ does not depend on θ . We compare this probability to the conditional probability

$$h(T) = \mathbb{P}_\theta(A \in B | T) = \mathbb{E}_\theta(\mathbf{1}_{\{A \in B\}} | T).$$

Since A is a function of the sample and since T is sufficient, the conditional distribution of A given T does not depend on θ . Thus, $h(T)$ is a function of T alone. Using the tower property from proposition A.9 we find

$$\mathbb{E}_\theta(h(T)) = \mathbb{P}_\theta(A \in B) = p_B$$

for all θ . Thus the function $h - p_B$ satisfies $\mathbb{E}_\theta(h(T) - p_B) = 0$ for all θ . Since T is complete, we have $h(T) = p_B$ with probability one for every $\mathbb{P}_\theta$.

It remains to show that this implies independence. For any set C of possible values of T , the indicator function $\mathbf{1}_{\{T \in C\}}$ is a function of T and can be taken out of the conditional expectation given T (see proposition A.10 in appendix A). Using the tower property we find

$$\begin{aligned}\mathbb{P}_\theta(T \in C, A \in B) &= \mathbb{E}_\theta(\mathbf{1}_{\{T \in C\}} \mathbf{1}_{\{A \in B\}}) = \mathbb{E}_\theta(\mathbf{1}_{\{T \in C\}} \mathbb{E}_\theta(\mathbf{1}_{\{A \in B\}} | T)) \\ &= \mathbb{E}_\theta(\mathbf{1}_{\{T \in C\}} p_B) = \mathbb{P}_\theta(T \in C) \mathbb{P}_\theta(A \in B).\end{aligned}$$

Since B and C were arbitrary, we have shown the independence of T and A and the proof is complete. $\square$

The proof uses each hypothesis exactly once: ancillarity makes the unconditional probability constant in θ , sufficiency makes the conditional probability a statistic, and completeness forces these two to be the same. The most important application of this result in this module is the following.

Corollary 10.10. *Let $X_1, \dots, X_n$ be i.i.d. $N(\mu, \sigma^2)$. Then $\bar{X}$ and S^2 are independent.*

Proof. Fixing the value of σ^2 , we can consider $\mu \in \mathbb{R}$ to be the only parameter. The density of a single observation is

$$f(x; \mu) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-x^2/(2\sigma^2)} \exp\left(\frac{\mu}{\sigma^2} x - \frac{\mu^2}{2\sigma^2}\right).$$

This is of the form given in definition 7.1 with natural statistic $T(x) = x$ and natural parameter $\eta(\mu) = \mu/\sigma^2$. As μ ranges over $\mathbb{R}$, so does $\eta(\mu) = \mu/\sigma^2$, that is, the natural parameter takes all values in $\mathbb{R}$, and thus the family is of full rank. By corollary 8.3 and proposition 10.5, the sum $\sum_i X_i$, and thus $\bar{X}$, is a complete and sufficient statistic for μ . Furthermore, S^2 is a function of the residuals $X_i - \bar{X}$ and thus, as shown above, is ancillary for μ . By Basu's theorem, theorem 10.9, the statistics $\bar{X}$ and S^2 are independent under $\mathbb{P}_{\mu, \sigma^2}$ for every μ and, since σ^2 was arbitrary, the claim is proved. $\square$

This completes the proof of the independence statement of proposition A.3. The result from lectures 25 and 29 on the ratio of $\bar{X} - \mu$ to S uses this proposition. The proof contains a useful trick: when both parameters are unknown, neither $\bar{X}$ nor S^2 can be ancillary and we cannot apply Basu's theorem directly. The solution to this problem is to fix the value of σ^2 , *i.e.* to consider the independence statement for one distribution at a time.

In lecture 13 we will see a second way to derive optimality, using the Cramer-Rao bound, which does not require conditioning.

Summary.

- Rao-Blackwell: If an unbiased estimator is conditioned on a sufficient statistic T , the conditioned estimator is unbiased and has no larger variance.
- A UMVUE is an unbiased estimator which has minimal variance among all unbiased estimators, for all θ .
- Lehmann-Scheffé: If T is complete and sufficient, then the only function of T which is unbiased for $g(\theta)$ is the UMVUE for $g(\theta)$.
- The natural statistic for a full-rank exponential family and the maximum of a uniform sample are complete, but without proof.
- Basu: A complete sufficient statistic is independent of all ancillary statistics. Using this result, one can show that $\bar{X}$ and S^2 of a normal sample are independent.

Exercise 10.1. Let $X_1, \dots, X_n$ be i.i.d. Poisson with parameter $\theta > 0$, let $T = \sum_{i=1}^n X_i$ and let $U = \mathbf{1}_{\{X_1=0\}}$, as in example 10.2. You may use the fact that a sum of independent Poisson random variables is again Poisson, with the parameters added.

1. Show that U is an unbiased estimator for $e^{-\theta}$ and determine the variance of this estimator.
2. Show that for $t \in \{0, 1, 2, \dots\}$ and $x \in \{0, 1, \dots, t\}$ we have

$$\mathbb{P}_\theta(X_1 = x | T = t) = \binom{t}{x} \left(\frac{1}{n}\right)^x \left(1 - \frac{1}{n}\right)^{t-x},$$

i.e. that the conditional distribution of X_1 given $T = t$ is Binomial($t, 1/n$) and that this distribution does not depend on θ .

3. Deduce that $V = \mathbb{E}_\theta(U|T) = (1 - 1/n)^T$.
4. Show directly that V is an unbiased estimator for $e^{-\theta}$, using the probability generating function $\mathbb{E}_\theta(s^T) = \exp(n\theta(s - 1))$ of the Poisson distribution with parameter $n\theta$.
5. Explain why V is the UMVUE for $e^{-\theta}$ and discuss the differences between this UMVUE and the MLE $e^{-\bar{X}}$, introduced in exercise 5.7.

Exercise 10.2. Let $n \geq 2$ and let $X_1, \dots, X_n$ be i.i.d. exponential with rate $\theta > 0$, *i.e.* $f(x; \theta) = \theta e^{-\theta x}$ for $x \geq 0$ and let $T = \sum_{i=1}^n X_i$. You may use the fact that the sum of k independent random variables of this type has density $f_k(t) = \theta^k t^{k-1} e^{-\theta t} / (k-1)!$ for $t \geq 0$.

1. Explain why X_1 is an unbiased estimator for the mean $1/\theta$ and why T is sufficient for θ .

2. Let $R = \sum_{i=2}^n X_i$. Use the independence of X_1 and R to find the joint density of (X_1, T) . Deduce the conditional density of X_1 given $T = s$. The result is

$$f(x|s) = \frac{(n-1)(s-x)^{n-2}}{s^{n-1}} \quad \text{for } 0 < x < s,$$

and zero elsewhere. Check that this density does not depend on θ .

3. Show that $\mathbb{E}_\theta(X_1|T=s) = s/n$ and conclude that Rao-Blackwellising X_1 leads to the sample mean $\bar{X}$, again unbiased for $1/\theta$.
4. Find a second argument for $\mathbb{E}_\theta(X_1|T) = T/n$, based on the symmetry of the sample, without any integration.

Exercise 10.3. Let $X_1, \dots, X_n$ be i.i.d. uniformly distributed on the set $(0, \theta)$ where $\theta > 0$, and let $M = \max_i X_i$.

1. Using the mean of M from exercise 2.2, show that $\hat{\theta} = (n+1)M/n$ is an unbiased estimator for θ .
2. Show that $\hat{\theta}$ is the UMVUE for θ .
3. The estimator $\tilde{\theta} = 2\bar{X}$ is also unbiased for θ . Determine the variances of $\hat{\theta}$ and of $\tilde{\theta}$ and verify that they are ordered as in theorem 10.4.
4. Without computing any conditional expectations, find $\mathbb{E}_\theta(\tilde{\theta}|M)$.

Exercise 10.4. Let $X_1, \dots, X_n$ be i.i.d. $N(\mu, \sigma^2)$, where both parameters are unknown. This is the situation of example 10.7.

1. Write the density of an observation in the form from definition 7.8 and identify the natural statistic $T(x) = (x, x^2)$ and the natural parameter $\eta(\mu, \sigma^2)$.
2. Show that the family of densities has full rank, in the sense of definition 7.10, by determining the range of η .
3. Deduce that $(\sum_i X_i, \sum_i X_i^2)$ is complete and sufficient and argue that the same holds for the pair $(\bar{X}, S^2)$.

Exercise 10.5. Let $X_1, \dots, X_n$ be i.i.d. $N(\mu, \sigma^2)$ where σ^2 is known and $\mu \in \mathbb{R}$ is unknown. Furthermore, let $R = \max_i X_i - \min_i X_i$ be the sample range.

1. Show that R is ancillary for μ .
2. Show that $\bar{X}$ and R are independent.
3. Now assume that μ is known and $\sigma^2 > 0$ is unknown. Is R ancillary for σ^2 ? Is R/σ a statistic?

Lecture 11

Fisher Information

Since the data are random, the score function from lecture 4 is a random variable, and the Fisher information, given by the variance of the score, measures how much information the sample contains about the parameter. In lecture 13 we will use the Fisher information to obtain a lower bound for the variance of any unbiased estimator, and in lecture 17 we will see that the Fisher information equals the asymptotic variance of the maximum likelihood estimator. In this lecture we will define the Fisher information, prove two identities which allow us to compute the Fisher information, and we will compute the Fisher information for the standard families of models.

11.1 Setting and Assumptions

Throughout this lecture, $X_1, \dots, X_n$ are i.i.d. with density or probability weights $f(x; \theta)$, where the parameter space $\Theta \subseteq \mathbb{R}$ is an open interval. As in lecture 4, the log-likelihood is the sum $\ell(\theta) = \sum_{i=1}^n \ell_i(\theta)$ of the terms $\ell_i(\theta) = \log f(X_i; \theta)$ and the score of definition 4.4 is its derivative $\ell'(\theta) = \sum_i \ell'_i(\theta)$. In this section we study the distribution of the score under $\mathbb{P}_\theta$, for a fixed value of θ .

The results of this lecture are obtained by differentiating, with respect to θ , the identity $\int f(x; \theta) dx = 1$, and we need this differentiation to be legitimate. We assume the following throughout. The support $\{x \mid f(x; \theta) > 0\}$ of the distribution does not depend on θ . For every x the function $\theta \mapsto f(x; \theta)$ is twice continuously differentiable on Θ . Finally, the integral of $f(x; \theta)$ over the support may be differentiated under the integral sign twice, so that

$$\frac{d^k}{d\theta^k} \int f(x; \theta) dx = \int \frac{\partial^k}{\partial \theta^k} f(x; \theta) dx \quad \text{for } k = 1, 2, \quad (11.1)$$

with sums in place of integrals for a discrete model. These three assumptions are referred to as the conditions of this section. Most standard families of models from tables 1.1 and 1.2 satisfy these conditions, except for the uniform distribution (in section 11.7). The exponential families in the sense of definition 7.1 with twice continuously differentiable natural parameter $\eta(\theta)$ also satisfy these conditions (since the exchange is the step in the proof of theorem 7.7). These conditions are included in the full list of regularity conditions in definition 16.1 of lecture 16.

11.2 The Score Has Mean Zero

The first identity is that, if θ is the correct parameter value, the data do not systematically deviate from this value: the score at the true parameter value is zero on average.

Lemma 11.1. *Let the situation be as described in section 11.1. Then*

$$\mathbb{E}_\theta(\ell'_i(\theta)) = 0 \quad \text{for } i = 1, \dots, n, \quad \text{and thus} \quad \mathbb{E}_\theta(\ell'(\theta)) = 0$$

for all $\theta \in \Theta$.

Proof. Using the chain rule, we find the derivative of $\ell_i(\theta) = \log f(X_i; \theta)$ as

$$\ell'_i(\theta) = \frac{\frac{\partial}{\partial \theta} f(X_i; \theta)}{f(X_i; \theta)}. \quad (11.2)$$

Since the denominator is positive (because X_i is in the support), we can cancel the density with the denominator to find

$$\begin{aligned} \mathbb{E}_\theta(\ell'_i(\theta)) &= \int \frac{\frac{\partial}{\partial \theta} f(x; \theta)}{f(x; \theta)} f(x; \theta) dx = \int \frac{\partial}{\partial \theta} f(x; \theta) dx \\ &= \frac{d}{d\theta} \int f(x; \theta) dx = \frac{d}{d\theta} 1 = 0, \end{aligned}$$

where the integral is over the support of the distribution, which does not depend on θ , and the third equality comes from the case $k = 1$ of (11.1). Summing over i gives the statement for the whole sample. This completes the proof. $\square$

If the parameter value is wrong, the score is not mean zero: for the exponential model from example 4.5, the score has expectation $n/\theta - n/\theta_0$ if the true parameter value is θ_0 , positive for $\theta < \theta_0$ and negative for $\theta > \theta_0$. The score on average points towards the correct parameter value from both sides. This is the mechanism behind the consistency proof of lecture 16.

11.3 Fisher Information

The next question we could ask is about the fluctuations of the score around zero. If the score is typically large in magnitude, this indicates that the data allow us to distinguish well between nearby parameter values. In contrast, if the score is typically close to zero, this indicates that the log-likelihood is nearly flat at the truth. The variance of the score is a measure for how informative the sample is.

Definition 11.2. Under the conditions of section 11.1, the *Fisher information* of the sample $X_1, \dots, X_n$ about θ is given by

$$\mathcal{I}_n(\theta) = \text{Var}_\theta(\ell'(\theta)),$$

where the variance of the score is taken at θ . The Fisher information of a single observation is denoted by $\mathcal{I}(\theta) = \mathcal{I}_1(\theta)$ and is given by the variance of $\ell'_1(\theta)$.

The variance of the score is often awkward to compute directly, and the second expression of the following theorem is nearly always the quickest route to the information.

Theorem 11.3. *Suppose we are in the situation of section 11.1. Then*

$$\mathcal{I}_n(\theta) = \mathbb{E}_\theta(\ell'(\theta)^2) = -\mathbb{E}_\theta(\ell''(\theta))$$

for all $\theta \in \Theta$.

Proof. The first equality follows from the fact that $\ell'(\theta)$ has mean zero by lemma 11.1, and thus the variance equals the second moment. The second equality is obtained by taking the derivative of the ratio (11.2) using the quotient rule:

$$\ell_i''(\theta) = \frac{\frac{\partial^2}{\partial \theta^2} f(X_i; \theta)}{f(X_i; \theta)} - \left(\frac{\frac{\partial}{\partial \theta} f(X_i; \theta)}{f(X_i; \theta)} \right)^2 = \frac{\frac{\partial^2}{\partial \theta^2} f(X_i; \theta)}{f(X_i; \theta)} - \ell_i'(\theta)^2.$$

Similar to the proof of lemma 11.1, the expectation of the first term cancels out in this case, using the case $k = 2$ of (11.1):

$$\mathbb{E}_\theta \left(\frac{\frac{\partial^2}{\partial \theta^2} f(X_i; \theta)}{f(X_i; \theta)} \right) = \int \frac{\partial^2}{\partial \theta^2} f(x; \theta) dx = \frac{d^2}{d\theta^2} \int f(x; \theta) dx = 0.$$

Thus we have $\mathbb{E}_\theta(\ell_i''(\theta)) = -\mathbb{E}_\theta(\ell_i'(\theta)^2)$ for all i . The terms $\ell_1'(\theta), \dots, \ell_n'(\theta)$ are independent of each other, since they are functions of different observations, and since they have mean zero, the variance of the sum is the sum of the variances:

$$\begin{aligned} \text{Var}_\theta(\ell'(\theta)) &= \sum_{i=1}^n \text{Var}_\theta(\ell_i'(\theta)) = \sum_{i=1}^n \mathbb{E}_\theta(\ell_i'(\theta)^2) \\ &= -\sum_{i=1}^n \mathbb{E}_\theta(\ell_i''(\theta)) = -\mathbb{E}_\theta(\ell''(\theta)), \end{aligned}$$

where we used $\ell''(\theta) = \sum_i \ell_i''(\theta)$ in the last step. This completes the proof. $\square$

The second expression gives the name of the quantity: it is the expected curvature of the log-likelihood at the true parameter value. Large information means that the log-likelihood has a sharp maximum, well-defined by the data. Small information means the log-likelihood is flat. In lecture 17 this idea will be turned into a theorem: the approximate variance of the maximum likelihood estimator is $1/\mathcal{I}_n(\theta)$ for large sample size, and the numerical value for the curvature $-\ell''(\hat{\theta})$ at the maximum likelihood estimate can be obtained by a numerical optimiser (lecture 18).

The proof also shows how the information of the sample relates to the information of one observation.

Proposition 11.4. *With the same assumptions as in section 11.1 we have*

$$\mathcal{I}_n(\theta) = n\mathcal{I}(\theta)$$

for all $\theta \in \Theta$.

Proof. The score is the sum of independent terms and thus the variance can be found as the sum $\sum_i \text{Var}_\theta(\ell_i'(\theta))$, as in the proof of theorem 11.3, and since the observations are identically distributed, each term in the sum equals

$$\text{Var}_\theta(\ell_1'(\theta)) = \mathcal{I}(\theta).$$

This completes the proof. $\square$

11.4 Examples

We now compute the information for the standard families of tables 1.1 and 1.2. By proposition 11.4 it suffices to work with a single observation $X = X_1$.

Example 11.5. For the Bernoulli distribution with success probability $\theta \in (0, 1)$ we have $\ell_1(\theta) = X \log \theta + (1 - X) \log(1 - \theta)$ and thus

$$\ell'_1(\theta) = \frac{X}{\theta} - \frac{1 - X}{1 - \theta}, \quad \ell''_1(\theta) = -\frac{X}{\theta^2} - \frac{1 - X}{(1 - \theta)^2}.$$

Using $\mathbb{E}_\theta(X) = \theta$ we find

$$\mathcal{I}(\theta) = -\mathbb{E}_\theta(\ell''_1(\theta)) = \frac{\theta}{\theta^2} + \frac{1 - \theta}{(1 - \theta)^2} = \frac{1}{\theta(1 - \theta)}.$$

Thus, the information of the sample is $\mathcal{I}_n(\theta) = n/(\theta(1 - \theta))$. Exercise 11.1 checks these identities for this model directly.

Example 11.6. For the Poisson distribution with mean $\theta > 0$ we have $\ell_1(\theta) = X \log \theta - \theta - \log X!$ and thus

$$\ell'_1(\theta) = \frac{X}{\theta} - 1, \quad \ell''_1(\theta) = -\frac{X}{\theta^2}, \quad \mathcal{I}(\theta) = \frac{\mathbb{E}_\theta(X)}{\theta^2} = \frac{1}{\theta}.$$

Example 11.7. In the case of the normal distribution with unknown mean μ and known variance σ^2 , the parameter is μ and $\ell_1(\mu) = -\frac{1}{2} \log(2\pi\sigma^2) - (X - \mu)^2/(2\sigma^2)$. Thus we have

$$\ell'_1(\mu) = \frac{X - \mu}{\sigma^2}, \quad \ell''_1(\mu) = -\frac{1}{\sigma^2}, \quad \mathcal{I}(\mu) = \frac{1}{\sigma^2}.$$

Since the second derivative does not depend on the data, we do not need to consider the expectation. Note that $\text{Var}_\mu(\bar{X}) = \sigma^2/n = 1/\mathcal{I}_n(\mu)$ by lemma 2.2 and from lecture 13 we know that no unbiased estimator can be better.

Example 11.8. For the exponential distribution with rate $\theta > 0$, we have $\ell_1(\theta) = \log \theta - \theta X$ and thus

$$\ell'_1(\theta) = \frac{1}{\theta} - X, \quad \ell''_1(\theta) = -\frac{1}{\theta^2}, \quad \mathcal{I}(\theta) = \frac{1}{\theta^2}.$$

Again, the second derivative is non-random. In fact, section 11.6 shows that the dependence of the information on the scale of the parameter is a general property.

11.5 Exponential Families

All four models from section 11.4 are exponential families. For such models, the information is the second derivative of the cumulant function K , given by definition 7.6.

Proposition 11.9. *Let $f(x; \eta) = h(x) \exp(\eta T(x) - K(\eta))$ be an exponential family in its natural parametrisation. Then the Fisher information of a single observation about η is given by*

$$\mathcal{I}(\eta) = K''(\eta) = \text{Var}_\eta(T(X)).$$

On the other hand, if the family is written as $f(x; \theta) = h(x) \exp(\eta(\theta)T(x) - B(\theta))$ as in definition 7.1 and if η is two times continuously differentiable, then the information about θ is given by

$$\mathcal{I}(\theta) = \eta'(\theta)^2 K''(\eta(\theta)) = \eta'(\theta)^2 \text{Var}_\theta(T(X)).$$

Proof. In the natural parametrisation we have $\ell_1(\eta) = \log h(X) + \eta T(X) - K(\eta)$ and thus $\ell'_1(\eta) = T(X) - K'(\eta)$ and $\ell''_1(\eta) = -K''(\eta)$. By theorem 7.7 we have $\mathbb{E}_\eta(T(X)) = K'(\eta)$ and $\text{Var}_\eta(T(X)) = K''(\eta)$ and thus, since the second derivative is non-random, we find $\mathcal{I}(\eta) = K''(\eta) = \text{Var}_\eta(T(X))$. For the general parametrisation we have $B(\theta) = K(\eta(\theta))$, since both quantities are determined by the requirement that the density must integrate to one, and thus, using the chain rule, we find

$$\ell'_1(\theta) = \eta'(\theta) (T(X) - K'(\eta(\theta))).$$

The factor $\eta'(\theta)$ is a constant and thus the variance of the score is

$$\mathcal{I}(\theta) = \eta'(\theta)^2 \text{Var}_\theta(T(X)) = \eta'(\theta)^2 K''(\eta(\theta)),$$

as claimed. $\square$

The proposition turns the computation of the information into that of a variance, which for the standard families can be read off tables A.1 and A.2; exercises 11.2 and 11.4 use it in this way.

11.6 Reparametrisation

The information depends on which parameter we use to describe a model, and the factor $\eta'(\theta)^2$ in proposition 11.9 is an instance of the general rule.

Proposition 11.10. *Let $\psi = g(\theta)$ where $g: \Theta \rightarrow g(\Theta)$ is a differentiable bijection with $g'(\theta) \neq 0$ for all $\theta \in \Theta$. Let $\mathcal{I}_\theta$ denote the Fisher information for a single observation when the model is parametrised by θ and let $\mathcal{I}_\psi$ be the information when the model is parametrised by ψ . Then,*

$$\mathcal{I}_\psi(\psi) = \frac{\mathcal{I}_\theta(\theta)}{g'(\theta)^2} \quad \text{for } \psi = g(\theta).$$

Proof. With the new parametrisation, the density is $\tilde{f}(x; \psi) = f(x; g^{-1}(\psi))$ and the log-likelihood for a single observation is $\tilde{\ell}_1(\psi) = \ell_1(g^{-1}(\psi))$. Since g' never vanishes on the interval Θ , by Darboux's theorem g is strictly monotonic and g^{-1} is differentiable. Using the chain rule and the derivative of the inverse function we get $(g^{-1})'(\psi) = 1/g'(\theta)$ for $\theta = g^{-1}(\psi)$. This gives

$$\tilde{\ell}'_1(\psi) = \ell'_1(\theta) (g^{-1})'(\psi) = \frac{\ell'_1(\theta)}{g'(\theta)}.$$

Since θ is fixed for fixed ψ , $1/g'(\theta)$ is a constant and taking variances under $\mathbb{P}_\theta$ gives $\mathcal{I}_\psi(\psi) = \mathcal{I}_\theta(\theta)/g'(\theta)^2$, as claimed. $\square$

For the exponential distribution, reparametrising from the rate θ to the mean $\mu = 1/\theta$ changes the information $1/\theta^2$ from example 11.8 to $\mathcal{I}_\mu(\mu) = 1/\mu^2$; see exercise 11.5. Thus, the information is not a property of the model in isolation, but depends on the parametrisation. What does not change under $\psi = g(\theta)$ is the combination $\mathcal{I}_\theta(\theta) (d\theta)^2$, since $d\psi = g'(\theta) d\theta$, and this invariance is the basis of the Jeffreys prior in lecture 26.

11.7 When the Identities Fail

The uniform distribution from example 1.2 shows what can go wrong if the conditions from section 11.1 are violated.

Example 11.11. Let X be uniformly distributed on $[0, \theta]$, *i.e.* $f(x; \theta) = 1/\theta$ for $0 \leq x \leq \theta$. The support of the distribution depends on θ , but we can still compute the derivatives on the support to see whether the identities hold. For $0 < X < \theta$ we have $\ell_1(\theta) = -\log \theta$ and thus $\ell'_1(\theta) = -1/\theta$ and $\ell''_1(\theta) = 1/\theta^2$. The score is a negative constant and thus we have $\mathbb{E}_\theta(\ell'_1(\theta)) = -1/\theta \neq 0$, in contradiction to lemma 11.1. The variance is zero, whereas $\mathbb{E}_\theta(\ell'_1(\theta)^2) = 1/\theta^2$ and $-\mathbb{E}_\theta(\ell''_1(\theta)) = -1/\theta^2$: the three expressions for theorem 11.3 take three different values, one of which is negative. The proof breaks down at the interchange between differentiation and integration: for this model, the two sides of (11.1) with $k = 1$ differ by the term $f(\theta; \theta) = 1/\theta$ from the moving upper limit, which is missed by the differentiation under the integral sign. Exercise 11.6 can be used to work out the two computations.

For this model, there is no Fisher information and the two results based on it are not applicable: the sample maximum does not satisfy the lower bound from lecture 13, the mean squared error decreases like $1/n^2$ per exercise 2.2 instead of like $1/n$, and the asymptotic normality from lecture 17 also fails.

Summary.

- Under the conditions of section 11.1 (a support free of θ , twice continuously differentiable density and differentiation under the integral sign), the score at the true parameter value has mean zero.
- The Fisher information $\mathcal{I}_n(\theta)$ is the variance of the score, and it equals $\mathbb{E}_\theta(\ell'(\theta)^2)$ and $-\mathbb{E}_\theta(\ell''(\theta))$, the expected curvature of the log-likelihood at the truth.
- Information from independent observations can be combined, $\mathcal{I}_n(\theta) = n\mathcal{I}(\theta)$; for an exponential family we have $\eta'(\theta)^2 \text{Var}_\theta(T)$; under a reparametrisation $\psi = g(\theta)$ we get $\mathcal{I}_\psi(\psi) = \mathcal{I}_\theta(\theta)/g'(\theta)^2$.
- We have learned the following standard values: $1/(\theta(1-\theta))$ for the Bernoulli distribution, $1/\theta$ for the Poisson distribution, $1/\sigma^2$ for the mean of the normal distribution and $1/\theta^2$ for the rate parameter in the exponential distribution. For the uniform distribution these relations do not hold, since the support of the distribution moves with the parameter.
- The information is the currency of the two results presented later: it gives a bound on the variance of unbiased estimators in lecture 13 and it determines the asymptotic variance of the MLE in lecture 17.

Exercise 11.1. Let $X_1, \dots, X_n$ be i.i.d. Bernoulli with success probability $\theta \in (0, 1)$, and let $\ell'_1(\theta)$ be the score of a single observation from example 11.5.

1. Verify lemma 11.1 for this model by computing $\mathbb{E}_\theta(\ell'_1(\theta))$ directly from the two possible values of X_1 .
2. Compute $\text{Var}_\theta(\ell'_1(\theta))$ directly, and check that it agrees with the value of $\mathcal{I}(\theta)$ found in example 11.5.
3. Compare $1/\mathcal{I}_n(\theta)$ with the variance of the estimator $\bar{X}$ for θ .

Exercise 11.2. Let $X_1, \dots, X_n$ be i.i.d. with density $f(x; \theta) = \theta^{-2} x e^{-x/\theta}$ for $x \geq 0$, where $\theta > 0$, the gamma distribution with known shape 2 and scale θ from exercise 5.2.

1. Compute the score $\ell'_1(\theta)$ of a single observation and verify that it has mean zero.
2. Show that $\mathcal{I}(\theta) = 2/\theta^2$, once via $-\mathbb{E}_\theta(\ell''_1(\theta))$ and once via the variance of the score, using $\text{Var}_\theta(X_1) = 2\theta^2$.
3. Write the density in the form of definition 7.1 with $T(x) = x$, identify $\eta(\theta)$, and recover $\mathcal{I}(\theta)$ from proposition 11.9.
4. In exercise 5.2 we found that the MLE $\hat{\theta} = \bar{X}/2$ is unbiased with $\text{MSE}(\hat{\theta}) = \theta^2/(2n)$. Compare this with $1/\mathcal{I}_n(\theta)$.

Exercise 11.3. Let $X_1, \dots, X_n$ be i.i.d. from the geometric distribution with parameter $p \in (0, 1)$, i.e. $f(x; p) = (1-p)^{x-1}p$ for $x \in \{1, 2, \dots\}$, where $\mathbb{E}_p(X_1) = 1/p$ and $\text{Var}_p(X_1) = (1-p)/p^2$, as in exercises 4.2 and 5.1.

1. Compute the score $\ell'_1(p)$ of a single observation and verify that it has mean zero.
2. Determine $\mathcal{I}(p) = 1/(p^2(1-p))$ and write down $\mathcal{I}_n(p)$.
3. Verify the value from the previous part by taking the variance of the score.

Exercise 11.4. Let $X_1, \dots, X_n$ be i.i.d. Poisson with mean $\theta > 0$.

1. From lecture 7 we know that the natural parameter is $\eta = \log \theta$ and the cumulant function is $K(\eta) = e^\eta$. Using proposition 11.9, find the information $\mathcal{I}_\eta(\eta)$ about η in terms of θ .
2. Recover $\mathcal{I}_\theta(\theta) = 1/\theta$ from the previous part using proposition 11.10.
3. In exercise 5.7 we have estimated the probability $p_0 = \mathbb{P}_\theta(X_1 = 0) = e^{-\theta}$. Determine the information $\mathcal{I}_{p_0}(p_0)$ about p_0 and describe its behaviour as $\theta \rightarrow \infty$. Give a one-sentence answer as to why this behaviour is reasonable.

Exercise 11.5. Let $X_1, \dots, X_n$ be i.i.d. exponential with rate $\theta > 0$ and let $\mu = 1/\theta$ be the mean.

1. Using proposition 11.10 and example 11.8, verify the value $\mathcal{I}_\mu(\mu) = 1/\mu^2$ from section 11.6.
2. Show that the information of the sample about μ is n/μ^2 and compare the reciprocal of this quantity to the variance of the estimator $\bar{X}$ for μ .

Exercise 11.6. Let X be uniformly distributed on the interval $[0, \theta]$ as described in example 11.11. Compute both sides of (11.1) for this model, using $k = 1$, and verify that the difference between the two quantities corresponds to the contribution $f(\theta; \theta) = 1/\theta$ of the (moving) upper limit of the integral.

Lecture 12

Sampling Distributions and Standard Errors

This is the R session for week 4. We will watch the sampling distribution of a maximum likelihood estimator converge to a normal distribution, we will use the Fisher information from this week's lectures to compute a standard error and a confidence interval, and we will learn how to estimate a probability using simulation. (The latter skill will be needed for the coverage checks in the practical report, and for the size and power calculations in lecture 24.)

12.1 The Sampling Distribution of an Estimator

The model of the day is $X_1, \dots, X_n$ i.i.d. from the exponential distribution with rate θ , where the maximum likelihood estimator is $\hat{\theta} = 1/\bar{X}$ by example 5.2. The command `rexp` can be used to simulate such a sample, and as in lecture 6 we package one run as a function. Before you run the code below, predict the two histograms: for $n = 10$, will the estimates form a symmetric bell around the true value $\theta = 1$? Will the peak sit above, below or at 1? How will the picture change for $n = 100$?

```
est_rate <- function(n, theta) {  
  sample <- rexp(n, rate = theta)  
  1 / mean(sample)  
}  
theta <- 1  
set.seed(2026)  
small <- replicate(10000, est_rate(10, theta))  
large <- replicate(10000, est_rate(100, theta))
```

As in lecture 9, we require the histograms to be generated with `freq = FALSE` and we overlay the density by using `curve`. The density we use is the large-sample approximation from lecture 17, which is normal with mean θ and standard deviation $\theta/\sqrt{n}$. The command `qqnorm` can be used to plot the sorted estimates against the quantiles of a standard normal distribution, and `qqline` can be used to add the line for the quantiles of a normal distribution.

```
par(mfrow = c(2, 2))  
hist(small, breaks = 40, freq = FALSE, col = "grey85", main = "n = 10",
```

```

xlab = "estimate")
curve(dnorm(x, mean = theta, sd = theta / sqrt(10)),
      add = TRUE, lwd = 2)
qqnorm(small, main = "n = 10")
qqline(small)
hist(large, breaks = 40, freq = FALSE, col = "grey85", main = "n = 100",
      xlab = "estimate")
curve(dnorm(x, mean = theta, sd = theta / sqrt(100)),
      add = TRUE, lwd = 2)
qqnorm(large, main = "n = 100")
qqline(large)

```




Figure 12.1: The sampling distribution of $\hat{\theta} = 1/\bar{X}$ for the exponential model with $\theta = 1$: histograms with the normal approximation overlaid (left column) and normal quantile plots (right column), for 10000 samples of size $n = 10$ (top row) and $n = 100$ (bottom row).

Figure 12.1 shows that for $n = 10$ the estimates are not normally distributed: the histogram has a long right tail, the peak is to the left of 1 and the quantile plot shows quantiles diverging from the line. Since $1/\bar{X}$ is not an average, the central limit theorem (theorem A.15) does not apply directly. For $n = 100$ the distortion is gone: this is asymptotic normality of the maximum likelihood estimator, proved in lecture 17.

12.2 Standard Errors and a Wald Interval

The standard deviation $\theta/\sqrt{n}$ of the normal curves in figure 12.1 is given by the Fisher information from lecture 11: By example 11.8 and definition 11.2 the sample contains information $\mathcal{I}_n(\theta) = n/\theta^2$ and lecture 17 will show that the spread of the estimator is approximately $1/\sqrt{\mathcal{I}_n(\theta)} = \theta/\sqrt{n}$. Normally, the value of θ is unknown and we have to evaluate the information at the estimate

instead. The *standard error* of the estimate is given by

$$\text{se}(\hat{\theta}) = \frac{1}{\sqrt{\mathcal{I}_n(\hat{\theta})}} = \frac{\hat{\theta}}{\sqrt{n}}.$$

This is a measure of the likely size of the estimation error, computed from the data. Since the estimator is approximately normally distributed with this spread, the *Wald interval* $\hat{\theta} \pm 1.96 \text{ se}(\hat{\theta})$ covers the true value with probability approximately 95%. This is a confidence interval of the kind we have seen in MATH2701, and is used informally here and more formally in lecture 25. The theorem which describes the resulting recipe is corollary 17.3. Before you try the code, make a prediction: will the interval cover the true value 1?

```
n <- 30
set.seed(2026)
x <- rexp(n, rate = theta)
theta_hat <- 1 / mean(x)
se <- theta_hat / sqrt(n)
theta_hat
se
theta_hat + c(-1, 1) * 1.96 * se
```

```
## [1] 0.6337667
## [1] 0.1157094
## [1] 0.4069762 0.8605572
```

The interval $[0.407, 0.861]$ does not contain the true value 1. This is not a programming error: the sample happened to contain unusually large observations, and a 95% interval misses the truth in about one sample out of twenty. From a single sample it is not possible to know whether this is one of the misses; how often the recipe is wrong in the long run can be answered by simulation.

12.3 Estimating a Probability by Simulation

The probability of an event A is the expected value of the indicator of A , and by the law of large numbers (theorem A.14) this expectation can be approximated by the average of the indicator over many independent repetitions. In R, the `mean` of the logical vector `abs(large - theta) <= 0.2` can be used to estimate $\mathbb{P}(|\hat{\theta} - \theta| \leq 0.2)$, since `mean` interprets `TRUE` as 1 and `FALSE` as 0. The same principle applies to all probabilities we estimate by simulation: for estimating the coverage probability of an interval, we take the average of the indicator that the interval contains the truth. Similarly, the rejection rate of a test (see lecture 24) is the average of the indicator that the test rejects.

We apply this to the coverage of the Wald interval: the box below contains a complete script, with its lines shuffled. Before you continue, arrange the lines on paper into a working script which estimates, from 10000 samples of size $n = 10$, the probability that the Wald interval contains θ .

```
mean(ci[1, ] <= theta & theta <= ci[2, ]) # empirical coverage
theta <- 1
wald <- function(n, theta) {
```

```
x <- rexp(n, rate = theta)
theta_hat <- 1 / mean(x)
se <- theta_hat / sqrt(n)
theta_hat + c(-1, 1) * 1.96 * se
}
ci <- replicate(10000, wald(10, theta))
set.seed(2026)
```

Names must be defined before they are used and the seed set before the first random draw (lecture 6): the function and `theta` come first, then the seed, then `replicate`, then the summary. One detail is new: the function returns the two endpoints, which `replicate` collects as the columns of a two-row matrix, so that `ci[1, ]` and `ci[2, ]` hold the lower and upper endpoints and `&` combines two logical vectors elementwise. Below, the reconstructed script is extended by a loop over $n = 10, 30$ and 100 . Before running it, predict the coverage at $n = 10$: given the skewness in figure 12.1, above or below 0.95?

```
wald <- function(n, theta) {
  x <- rexp(n, rate = theta)
  theta_hat <- 1 / mean(x)
  se <- theta_hat / sqrt(n)
  theta_hat + c(-1, 1) * 1.96 * se
}
theta <- 1
ns <- c(10, 30, 100)
coverage <- numeric(3)
set.seed(2026)
for (i in seq_along(ns)) {
  ci <- replicate(10000, wald(ns[i], theta))
  coverage[i] <- mean(ci[1, ] <= theta & theta <= ci[2, ])
}
data.frame(n = ns, coverage = coverage)
```

```
##      n coverage
## 1  10   0.9535
## 2  30   0.9539
## 3 100   0.9504
```

The coverage is close to the nominal 0.95 for all three sample sizes, even for $n = 10$ where the sampling distribution was far from normal. The total hides an imbalance: the standard error $\hat{\theta}/\sqrt{n}$ is proportional to the estimate, so small estimates get short intervals and at $n = 10$ most misses are below the truth (check this yourself!).

12.4 Skewed Data

Tables A.1 and A.2 list three distributions which the lectures do not otherwise use, the negative binomial, the Pareto and the Rayleigh distributions. The command `rnbinom` can be used to draw from the negative binomial distribution, counting failures before the `size`-th success as in the table. The other two distributions cannot be directly generated in base R. Instead, if U is uniformly

distributed on the interval $(0, 1)$ and F is a continuous, strictly increasing distribution function, then $F^{-1}(U)$ has distribution function F . For the Pareto distribution with $\alpha = 3$, this gives $X = (1 - U)^{-1/3}$. For the Rayleigh distribution with $\theta = 2$, the square X^2 is exponentially distributed with mean θ , *i.e.* $X = \sqrt{2E}$ where E is standard exponentially distributed. Before you try the code, try to sketch the three histograms you would expect: which distribution is discrete? Which distribution starts at 1? Which distribution has the longest tail? Finally, once you have done this, look up the means and variances in the table.

```
N <- 10000
set.seed(2026)
nb <- rbinom(N, size = 3, prob = 0.4)
pa <- (1 - runif(N))^(1 / 3)
ra <- sqrt(2 * rexp(N))
c(mean(nb), var(nb))
c(mean(pa), var(pa))
c(mean(ra), var(ra))
```

```
## [1] 4.5308 11.5188
## [1] 1.4993207 0.6484568
## [1] 1.2450801 0.4328587
```

The empirical means and variances match the values from the table, except for the Pareto distribution where the variance 0.648 is much smaller than the table value 0.75: This discrepancy is caused by the fact that the 10000 samples have not yet observed the large, rare events which contribute significantly to the variance of this heavy-tailed distribution. The script `R/S12-sheet.R` continues by computing 10000 sample means of $n = 30$ draws from each distribution and by superimposing the normal density from the central limit theorem over the histograms. Which of the three histograms of means do you expect will be furthest away from the normal curve?

In figure 12.2 thirty observations are enough for the central limit theorem to work well in the first and third columns, but not in the middle column where the heavy tail causes the histogram of means to be still visibly skewed. The heavier the tail, the more n is required and the same care is needed when using the Wald intervals above.

As a task to try at home, in the style of a short practical report, choose a rate θ , let $n = 50$, set the random number generator seed to your student ID, and simulate one sample from an exponential distribution. Determine the estimate $\hat{\theta} = 1/\bar{x}$, the standard error, and the Wald interval for this sample. Write down your results in one sentence each. Finally, using 10000 samples, estimate the coverage of the interval at $n = 50$. Also, in separate analyses, determine how often the interval is entirely below the truth and how often it is entirely above the truth. In two sentences, discuss whether you would trust the interval for this sample size.

Before the next session, type the commands used in this session into R, and check that your answers match the text.




Figure 12.2: Top: histograms of 10000 draws from the negative binomial distribution with $r = 3$ and $p = 0.4$, the Pareto distribution with $\alpha = 3$ and the Rayleigh distribution with $\theta = 2$. Bottom: histograms of 10000 sample means of $n = 30$ draws, with the normal density from the central limit theorem overlaid; the two middle panels are cut off on the right.

Summary.

- The written reference for the commands used in this session is appendix B.
- The normal approximation to the sampling distribution of $\hat{\theta} = 1/\bar{X}$ is poor for $n = 10$ and good for $n = 100$. The standard error $\text{se}(\hat{\theta}) = 1/\sqrt{\mathcal{I}_n(\hat{\theta})}$ and the Wald interval $\hat{\theta} \pm 1.96 \text{ se}(\hat{\theta})$ can be computed from data.
- A probability can be estimated by simulation as the mean of a logical vector, *i.e.* by taking the average of an indicator function over replicates. This method can be used to compute both coverage probabilities and rejection rates.
- The negative binomial, Pareto and Rayleigh distributions can be simulated using `rnbinom`, the inverse distribution function and the square root of an exponential variable, respectively. The heavier the tail of a distribution, the longer the central limit theorem takes to apply.

Lecture 13

The Cramer–Rao Inequality

In lecture 10 we have found the best unbiased estimators by conditioning on a complete sufficient statistic. In contrast, in this lecture we will discuss a question which could be asked about any unbiased estimator: What are the limitations of unbiased estimators? In the situation of lecture 11, we will see that the variance of an unbiased estimator cannot be smaller than $1/\mathcal{I}_n(\theta)$. This is the Cramer–Rao inequality, which gives a precise meaning to the word “optimal” and which we will use in lecture 17 to characterise the maximum likelihood estimator. In this lecture we will prove the inequality, generalise it to functions of the parameter, and then see which estimators achieve the bound. We will conclude by considering the uniform distribution, where the conditions for the bound to hold are violated and where it is possible to beat the bound.

13.1 The Inequality

Throughout this lecture, $X_1, \dots, X_n$ are i.i.d. with density or probability weights $f(x; \theta)$, the parameter space $\Theta \subseteq \mathbb{R}$ is an open interval, and the conditions of section 11.1 hold. The score $\ell'(\theta)$ and the Fisher information $\mathcal{I}_n(\theta) = \text{Var}_\theta(\ell'(\theta))$ are as in lecture 11; we assume in addition that $0 < \mathcal{I}_n(\theta) < \infty$ for all $\theta \in \Theta$, and we write $f(x; \theta) = \prod_{i=1}^n f(x_i; \theta)$ for the joint density of a value $x = (x_1, \dots, x_n)$ of the sample.

The idea of the proof can be stated in advance: the covariance between an unbiased estimator W and the score equals one, for any model, and then the Cauchy–Schwarz inequality implies that small Fisher information implies large variance. We state the result for the covariance as a separate lemma, for use later for functions of the parameter.

Lemma 13.1. *Let $W = W(X_1, \dots, X_n)$ be a statistic such that $\text{Var}_\theta(W) < \infty$ for all $\theta \in \Theta$, and assume that the expectation of W can be differentiated under the integral sign, i.e. that*

$$\frac{d}{d\theta} \int W(x) f(x; \theta) dx = \int W(x) \frac{\partial}{\partial \theta} f(x; \theta) dx, \quad (13.1)$$

where the sum replaces the integral for discrete models, holds. Then we have

$$\text{Cov}_\theta(W, \ell'(\theta)) = \frac{d}{d\theta} \mathbb{E}_\theta(W)$$

for all $\theta \in \Theta$.

Proof. Since the score has mean zero (lemma 11.1), the covariance is given by the expectation of the product $\text{Cov}_\theta(W, \ell'(\theta)) = \mathbb{E}_\theta(W \ell'(\theta))$. As in equation (11.2), the score of the sample is the ratio $\ell'(\theta) = \frac{\partial}{\partial \theta} f(X; \theta) / f(X; \theta)$ of the joint density and thus the density cancels when we take the expectation:

$$\begin{aligned} \mathbb{E}_\theta(W \ell'(\theta)) &= \int W(x) \frac{\frac{\partial}{\partial \theta} f(x; \theta)}{f(x; \theta)} f(x; \theta) \, dx = \int W(x) \frac{\partial}{\partial \theta} f(x; \theta) \, dx \\ &= \frac{d}{d\theta} \int W(x) f(x; \theta) \, dx, \end{aligned}$$

where the integral is over the support of the distribution, which does not depend on θ , and where we used the assumption (13.1) in the last step. The final integral can be evaluated as $\mathbb{E}_\theta(W)$. This completes the proof. $\square$

The assumption (13.1) is an additional exchange of differentiation and integration. For exponential families, this relation holds for all statistics with finite variance, by a similar argument from analysis as given in the proof of theorem 7.7, but we will not verify this relation for individual examples.

Theorem 13.2 (Cramer–Rao inequality). *Assume the conditions of section 11.1 and $0 < \mathcal{I}_n(\theta) < \infty$. Let W be an unbiased estimator for θ with $\text{Var}_\theta(W) < \infty$ which satisfies (13.1). Then we have*

$$\text{Var}_\theta(W) \geq \frac{1}{\mathcal{I}_n(\theta)}$$

for all $\theta \in \Theta$.

Proof. Since W is unbiased, we have $\mathbb{E}_\theta(W) = \theta$ and using lemma 13.1 we find that the covariance $\text{Cov}_\theta(W, \ell'(\theta))$ is the derivative of θ with respect to θ , i.e. it equals 1. Using the Cauchy–Schwarz inequality, lemma A.8, for W and $\ell'(\theta)$, we find

$$1 = \text{Cov}_\theta(W, \ell'(\theta))^2 \leq \text{Var}_\theta(W) \text{Var}_\theta(\ell'(\theta)) = \text{Var}_\theta(W) \mathcal{I}_n(\theta),$$

where we used definition 11.2 for the last step. Dividing by the positive number $\mathcal{I}_n(\theta)$ gives the claim. This completes the proof. $\square$

Three comments help to read the theorem. First, by proposition 11.4 the bound is $1/(n\mathcal{I}(\theta))$: no unbiased estimator can have a standard deviation which decreases faster than the $1/\sqrt{n}$ of the sample mean (lemma 2.2). Second, the bound is independent of the estimator used, and only depends on the model: the bound states how much the data can reveal about θ , regardless of the unbiased method used to estimate it. For the normal mean with known variance the bound is $\sigma^2/n = \text{Var}_\mu(\bar{X})$, as shown in example 11.7, and thus $\bar{X}$ achieves this bound. Third, by theorem 2.7, the bound also acts as a lower bound for the mean squared error, but only amongst unbiased estimators: in section 2.3 we have seen a biased estimator for the normal variance, with smaller mean squared error than S^2 .

13.2 Estimating a Function of the Parameter

Often the quantity of interest is a function $g(\theta)$ of the parameter, for example the probability $e^{-\theta}$ of zero counts in example 10.2. If W is an unbiased estimator for $g(\theta)$, the only change in the proof is the value of the covariance.

Theorem 13.3. *Suppose the conditions of theorem 13.2 hold and that $g: \Theta \rightarrow \mathbb{R}$ is differentiable. Let W be an unbiased estimator for $g(\theta)$ with $\text{Var}_\theta(W) < \infty$ which satisfies (13.1). Then,*

$$\text{Var}_\theta(W) \geq \frac{g'(\theta)^2}{\mathcal{I}_n(\theta)}$$

for all $\theta \in \Theta$.

Proof. We have $\mathbb{E}_\theta(W) = g(\theta)$ and from lemma 13.1 we know that the covariance satisfies $\text{Cov}_\theta(W, \ell'(\theta)) = g'(\theta)$. Using the Cauchy–Schwarz inequality we find $g'(\theta)^2 \leq \text{Var}_\theta(W) \mathcal{I}_n(\theta)$ and dividing by $\mathcal{I}_n(\theta)$ completes the proof. $\square$

Theorem 13.2 is the case $g(\theta) = \theta$; it is the general form which is worth memorising: identify the function g the estimator is unbiased for, differentiate, square and divide by the information. Both numerator and denominator depend on how the model is parametrised, but by proposition 11.10 the bound itself does not: it is a property of the estimator and the model, not of the name we give to the parameter. Exercise 13.6 verifies this.

13.3 Efficiency and Attainment

We give a name to the estimators which attain the bound.

Definition 13.4. Let W be an unbiased estimator for $g(\theta)$ that satisfies the assumptions of theorem 13.3. Then the *efficiency* of W is given by

$$\text{eff}_\theta(W) = \frac{g'(\theta)^2 / \mathcal{I}_n(\theta)}{\text{Var}_\theta(W)},$$

where the efficiency is the ratio between the Cramer–Rao bound and the actual variance. The estimator W is called *efficient*, if $\text{eff}_\theta(W) = 1$, i.e. if $\text{Var}_\theta(W) = g'(\theta)^2 / \mathcal{I}_n(\theta)$, for all $\theta \in \Theta$.

From theorem 13.3 we know that the efficiency is contained in the interval $(0, 1]$, and since the bound is proportional to $1/n$, an efficiency of $1/2$ means that the estimator uses the information from n observations as well as an efficient estimator would use the information from $n/2$. An efficient estimator is a UMVUE in the sense of definition 10.3, because no unbiased estimator can be better than the bound, but the converse statement is not true, since the bound is not always attainable: in exercise 13.4, the UMVUE for the rate of the exponential distribution has efficiency $(n-2)/n$. These two results complement each other: the Lehmann–Scheffe theorem allows us to identify the best unbiased estimator, whereas the Cramer–Rao bound provides an absolute scale for comparing estimators.

The proof tells us when the bound is attained: equality in the Cramer–Rao inequality is equality in the Cauchy–Schwarz step, and this only holds for random variables which are proportional to each other.

Proposition 13.5. *With the same assumptions as in theorem 13.3, the estimator W is efficient if and only if there is a function $c(\theta)$ such that*

$$W - g(\theta) = c(\theta) \ell'(\theta) \quad \text{with probability one, for every } \theta \in \Theta, \quad (13.2)$$

and in this case $c(\theta) = g'(\theta)/\mathcal{I}_n(\theta)$. In particular, if the model is a one-parameter exponential family as in definition 7.1 and the density is

$$f(x; \theta) = h(x) \exp(\eta(\theta)T(x) - B(\theta)),$$

where η is continuously differentiable with $\eta'(\theta) \neq 0$ for all θ , and where $T_n = \sum_{i=1}^n T(X_i)$ is the natural statistic, then T_n/n is an efficient estimator for the mean $\mu(\theta) = \mathbb{E}_\theta(T(X_1))$.

Proof. Let $U = W - g(\theta)$ and $V = \ell'(\theta)$ be two random variables with mean zero. Then $\text{Cov}_\theta(W, V) = \mathbb{E}_\theta(UV)$, $\text{Var}_\theta(W) = \mathbb{E}_\theta(U^2)$ and $\mathcal{I}_n(\theta) = \mathbb{E}_\theta(V^2) > 0$. The proof of lemma A.8 in appendix A considers the quadratic $q(t) = \mathbb{E}_\theta((U - tV)^2) = \mathbb{E}_\theta(U^2) - 2t\mathbb{E}_\theta(UV) + t^2\mathbb{E}_\theta(V^2)$, which is non-negative for all real t , and the inequality states that the discriminant is non-positive. Efficiency of W implies that the inequality is an equality, i.e. that the discriminant is zero and thus that q has a real root t_0 . Since $q(t_0) = 0$ is the expectation of the non-negative random variable $(U - t_0V)^2$, we have $U = t_0V$ with probability one, which is (13.2) with $c(\theta) = t_0$. Conversely, if $U = c(\theta)V$ with probability one, we have $\mathbb{E}_\theta(UV)^2 = c(\theta)^2\mathbb{E}_\theta(V^2)^2 = \mathbb{E}_\theta(U^2)\mathbb{E}_\theta(V^2)$. This is an equality in the Cauchy–Schwarz inequality and thus implies efficiency. To identify $c(\theta)$ we take the covariance of both sides of (13.2) with $\ell'(\theta)$. The left-hand side gives $g'(\theta)$ by lemma 13.1. The right-hand side gives $c(\theta)\mathcal{I}_n(\theta)$.

For the exponential family, the proof of proposition 11.9 gives the score of a single observation as $\ell'_i(\theta) = \eta'(\theta)(T(X_i) - K'(\eta(\theta)))$, and theorem 7.7 identifies $K'(\eta(\theta))$ with $\mu(\theta)$. Summing over i we find $\ell'(\theta) = \eta'(\theta)(T_n - n\mu(\theta))$, and thus

$$\frac{T_n}{n} - \mu(\theta) = \frac{1}{n\eta'(\theta)} \ell'(\theta),$$

which is (13.2) for $W = T_n/n$, $g = \mu$ and $c(\theta) = 1/(n\eta'(\theta))$. The estimator T_n/n is unbiased for $\mu(\theta)$ by construction and has finite variance $K''(\eta(\theta))/n$, and the conditions of section 11.1 hold for the family as noted there. This completes the proof. $\square$

The condition (13.2) is restrictive. In an exponential family, an efficient estimator must be an affine function of the natural statistic T_n , and the coefficients of any such affine function cannot depend on θ (since W is fixed while θ varies over Θ). Thus, the only functions of θ which can lead to an efficient estimator are affine functions $a + b\mu(\theta)$ of the mean of the natural statistic, estimated by $a + bT_n/n$. For the Poisson distribution, an exponential family with $T(x) = x$ and $\mu(\theta) = \theta$, the sample mean is an efficient estimator for θ . Exercise 13.3 shows that the attainment condition is violated in this case, by showing that $e^{-\theta}$ is not an affine function of θ . Thus, the UMVUE $(1 - 1/n)^{T_n}$ from example 10.6 is the best one among the unbiased estimators, but does not achieve the bound.

13.4 When the Conditions Fail

As we have seen, the proof of the inequality in this lecture involves two instances of differentiating under the integral sign, once in lemma 11.1 and once in the assumption (13.1). If the support of the distribution moves as θ does, both of these exchanges will fail, since an additional term is added to the integral which is not included in the exchange.

Example 13.6. Let $X_1, \dots, X_n$ be i.i.d. with a uniform distribution on $(0, \theta)$ and $\theta > 0$. The support of the distribution depends on θ and the conditions from section 11.1 are violated. Thus, theorem 13.2 does not apply. We can still see what the formula would be for this model: in example 11.11 the three expressions of theorem 11.3 had three different values, and taking the positive of these three values, $\mathbb{E}_\theta(\ell'_1(\theta)^2) = 1/\theta^2$, and multiplying by n as in proposition 11.4, we find the formal expression n/θ^2 and the formal bound θ^2/n .

Now consider the estimators for θ based on the maximum $M = \max_i X_i$. The MLE M is biased by exercise 2.2, and thus the mean squared error alone does not violate the inequality, since the inequality is a statement about unbiased estimators. The bias-corrected estimator $W = (n+1)M/n$ is unbiased and by exercise 10.3 has variance

$$\text{Var}_\theta(W) = \frac{\theta^2}{n(n+2)} < \frac{\theta^2}{n} \quad \text{for every } n \geq 1,$$

which is of order $1/n^2$ instead of $1/n$. An unbiased estimator with variance smaller than the formal bound in this way is said to be *super-efficient*. The proof fails at the covariance. The score of the sample on the support is the constant $\ell'(\theta) = -n/\theta$, and thus we have $\text{Cov}_\theta(W, \ell'(\theta)) = 0$ instead of 1: the two sides of (13.1) are $\frac{d}{d\theta} \mathbb{E}_\theta(W) = 1$ on the left and $\int W(x) \frac{\partial}{\partial \theta} f(x; \theta) dx = -(n/\theta) \mathbb{E}_\theta(W) = -n$ on the right, and the difference is the contribution of the moving boundary of the support.

There is no contradiction: the inequality has hypotheses, and the uniform model violates these hypotheses. Super-efficiency is a property of models where the parameter is an endpoint of the support, where a single observation near the endpoint can be used to estimate θ much more accurately than for a regular model. In lecture 16 we will prove the consistency of M using a direct argument. In lecture 17 we will see that the limiting distribution of M is not normal. In lecture 15 we will use simulation to demonstrate the $1/n^2$ decay of the variance.

Summary.

- Under the conditions of section 11.1, all unbiased estimators W for $g(\theta)$ satisfy $\text{Var}_\theta(W) \geq g'(\theta)^2 / \mathcal{I}_n(\theta)$. For $g(\theta) = \theta$ we get the bound $1 / \mathcal{I}_n(\theta)$. The proof uses $\text{Cov}_\theta(W, \ell'(\theta)) = g'(\theta)$ and the Cauchy–Schwarz inequality.
- The efficiency of an unbiased estimator is the ratio of the bound to the variance, with values in $(0, 1]$. An efficient estimator is a UMVUE, but the converse is not true.
- The bound is saturated if and only if $W - g(\theta) = c(\theta)\ell'(\theta)$ with probability one. For the one-parameter exponential family, the natural statistic T_n/n is efficient for the mean, and the only functions of θ which can be used to construct an efficient estimator are affine functions of the mean.
- The bound depends on regularity conditions. For example, for the uniform distribution on $(0, \theta)$, the support of the distribution depends on the parameter. The unbiased estimator $(n+1)M/n$ has variance $\theta^2/(n(n+2))$, which is smaller than the formal value θ^2/n . Thus, the estimator is super-efficient.

Exercise 13.1. Let $X_1, \dots, X_n$ be i.i.d. exponentially distributed with rate $\theta > 0$. Define $f(x; \theta) = \theta e^{-\theta x}$ for $x \geq 0$ and $T = \sum_{i=1}^n X_i$. Use the fact that the chi-squared distribution χ_k^2 from definition A.2 coincides with the Gamma($k/2, 1/2$) distribution from table A.2 and that T is Gamma(n, θ) distributed.

1. Show that $Y_1 = 2\theta X_1$ has density $\frac{1}{2}e^{-y/2}$ for $y \geq 0$ and consequently $Y_1 \sim \chi_2^2$.
2. Using definition A.2, show that $2\theta T \sim \chi_{2n}^2$.
3. Using the gamma density of T and the change of variables formula from proposition A.19, show that you get the same result as in the previous part.
4. Using the mean and variance of the chi-squared distribution, find $\mathbb{E}_\theta(T)$ and $\text{Var}_\theta(T)$. Check your values by computing the moments of a single observation.

Exercise 13.2. Let $X_1, \dots, X_n$ be i.i.d. Bernoulli with success probability $\theta \in (0, 1)$, and let $T_n = \sum_{i=1}^n X_i$.

1. Write down the Cramer–Rao bound for θ , using example 11.5.
2. Show that $\bar{X}$ is an efficient estimator for θ .
3. Verify condition (13.2) for $W = \bar{X}$ directly, by writing the score $\ell'(\theta)$ in terms of $\bar{X} - \theta$, and check that the resulting $c(\theta)$ equals $g'(\theta)/\mathcal{I}_n(\theta)$.
4. Write down the Cramer–Rao bound for the odds $g(\theta) = \theta/(1 - \theta)$, and explain why no efficient estimator for the odds exists.
5. Show that in fact no unbiased estimator for the odds exists at all.

Exercise 13.3. Let $X_1, \dots, X_n$ be i.i.d. Poisson with mean $\theta > 0$ and let $T_n = \sum_{i=1}^n X_i$.

1. Write down the Cramer–Rao bound for θ , using example 11.6, and show that $\bar{X}$ achieves this bound.
2. Verify condition (13.2) for $W = \bar{X}$ directly, by writing the score $\ell'(\theta)$ in terms of $\bar{X} - \theta$, and by checking that the resulting $c(\theta)$ equals $g'(\theta)/\mathcal{I}_n(\theta)$.
3. Explain why no efficient estimator for $g(\theta) = e^{-\theta}$ exists, and why this does not contradict the fact that $(1 - 1/n)^{T_n}$ is the UMVUE for $e^{-\theta}$, using example 10.6.

Exercise 13.4. Let $X_1, \dots, X_n$, where $n \geq 3$, be i.i.d. exponential with rate $\theta > 0$, let $T = \sum_{i=1}^n X_i$ and let $\tilde{\theta} = (n - 1)/T$ be the unbiased version of the MLE from lecture 5.

1. Write down the Cramer–Rao bound for θ , using example 11.8.
2. In exercise 5.3 we found $\text{MSE}(\tilde{\theta}) = \theta^2/(n - 2)$. Determine the efficiency of $\tilde{\theta}$ and describe how it behaves as $n \rightarrow \infty$.
3. Show that no unbiased estimator for θ can be efficient.
4. Show that $\tilde{\theta}$ is the UMVUE for θ .
5. Show that $\bar{X}$ is an efficient estimator for the mean $\mu = 1/\theta$, and explain how this is compatible with the previous results.

Exercise 13.5. Let $X_1, \dots, X_n$ be i.i.d. $N(\mu, \sigma^2)$. Assume first that μ is known and $\sigma^2 > 0$ is the parameter.

1. Show that $\mathcal{I}_n(\sigma^2) = n/(2\sigma^4)$ and find the Cramer–Rao bound for σ^2 .
2. Show that $\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \mu)^2$ is an efficient estimator for σ^2 , using definition A.2.
3. Now assume that μ is unknown. From lemma 2.3 we know that the sample variance S^2 is an unbiased estimator for σ^2 . Using proposition A.3, find $\text{Var}(S^2)$ and compare the result with the bound from the first part of the question. Lecture 14 shows that this bound still holds when μ is unknown. What can you say about S^2 , in light of example 10.7?

Exercise 13.6. Let $\psi = g(\theta)$ where $g: \Theta \rightarrow g(\Theta)$ is a continuously differentiable bijection with $g'(\theta) \neq 0$ for all $\theta \in \Theta$. Furthermore, let $\mathcal{I}_\psi$ and $\mathcal{I}_\theta$ be the Fisher information for the sample with model parametrised by ψ and by θ , respectively.

1. Using propositions 11.4 and 11.10, show that $\mathcal{I}_\psi(\psi) = \mathcal{I}_\theta(\theta)/g'(\theta)^2$.
2. Let W be an unbiased estimator for $g(\theta)$ which satisfies the assumptions of theorem 13.3. Show that the bound from theorem 13.2, for W as an unbiased estimator for the parameter ψ , coincides with the bound from theorem 13.3.

Lecture 14

Multiparameter Fisher Information

In lectures 11 and 13 the parameter was a single number, but most models we actually fit, *e.g.* the normal distribution with unknown mean and variance, have several parameters. In this lecture we will extend the concepts of Fisher information and the Cramer–Rao bound to a vector parameter $\theta = (\theta_1, \dots, \theta_k)$: the score will be a vector, the information will be a $k \times k$ matrix, and the bound will be an inequality of covariance matrices. The new feature here is that we will need to invert the matrix, and the inverse will give the cost of estimating one parameter in the presence of the other parameters.

14.1 The Score Vector and the Information Matrix

Setting the stage as in lecture 11, we now assume that the parameter space $\Theta \subseteq \mathbb{R}^k$ is an open set and that the derivatives with respect to θ are partial derivatives, denoted by ∂_j for the derivative with respect to θ_j . The data $X_1, \dots, X_n$ are now i.i.d. with density or probability weights $f(x; \theta)$ and the log-likelihood function is $\ell(\theta) = \sum_{i=1}^n \ell_i(\theta)$, where $\ell_i(\theta) = \log f(X_i; \theta)$, and we assume that the conditions of section 11.1 hold for each coordinate, allowing for differentiation under the integral sign for any pair of coordinates, and we assume that each coordinate of the score has finite variance under $\mathbb{P}_\theta$, so that the entries of the information matrix below are finite.

Definition 14.1. The *score vector* of the sample is given by the gradient of the log-likelihood function

$$\nabla \ell(\theta) = (\partial_1 \ell(\theta), \dots, \partial_k \ell(\theta)),$$

as a column vector in $\mathbb{R}^k$.

For $k = 1$ the score vector is the score $\ell'(\theta)$ from definition 4.4. Setting the score vector to zero leads to the system of likelihood equations for the normal distribution considered in example 5.4. As in lecture 11, we will consider the score vector to be a random vector under $\mathbb{P}_\theta$ and we will determine the covariance of this random vector. The covariance matrix $\text{Cov}(Y)$ of a random vector $Y = (Y_1, \dots, Y_k)$ has entries $\text{Cov}(Y_j, Y_l)$. The diagonal of this matrix contains the variances.

Definition 14.2. Given the conditions of section 11.1, the *Fisher information matrix* of the sample about θ is defined as

$$\mathcal{I}_n(\theta) = \text{Cov}_\theta(\nabla \ell(\theta)), \quad \text{i.e.} \quad \mathcal{I}_n(\theta)_{jl} = \text{Cov}_\theta(\partial_j \ell(\theta), \partial_l \ell(\theta)),$$

where this is the $k \times k$ covariance matrix of the score vector. The information matrix for a single observation is denoted by $\mathcal{I}(\theta) = \mathcal{I}_1(\theta)$.

The diagonal element $\mathcal{I}_n(\theta)_{jj}$ is the Fisher information about θ_j in the sense of definition 11.2 with all other coordinates fixed to their true values; the off-diagonal elements are new and measure the strength of correlations between the components of the score. Both identities from lecture 11 carry over to this situation, separately for each coordinate.

Proposition 14.3. *With the assumptions of section 11.1 in place, we have $\mathbb{E}_\theta(\nabla \ell(\theta)) = 0$ and*

$$\mathcal{I}_n(\theta)_{jl} = \mathbb{E}_\theta(\partial_j \ell(\theta) \partial_l \ell(\theta)) = -\mathbb{E}_\theta(\partial_j \partial_l \ell(\theta)) \quad \text{for } j, l = 1, \dots, k.$$

Thus, $\mathcal{I}_n(\theta)$ is minus the expected Hessian matrix of the log-likelihood.

The proof of this result is similar to the proof of lemma 11.1 and theorem 11.3, but with partial derivatives instead of $d/d\theta$. We omit the details of the proof here, but we can mention that the ratio $\partial_j f/f$ has mean zero since the density in the numerator and denominator cancels. Using the quotient rule we find $\partial_j \partial_l \ell_i = \partial_j \partial_l f(X_i; \theta)/f(X_i; \theta) - \partial_j \ell_i \partial_l \ell_i$ where the first term has mean zero. Finally, the independence of the observations allows us to carry over the resulting identity from one observation to the sample. As for the scalar case, the second expression is the more practical one, at least for $k = 2$, because it requires only three second derivatives (one mixed derivative, occurring twice by symmetry).

14.2 Properties of the Information Matrix

The three properties of the scalar information which we used most often in lecture 11 have matrix counterparts. Recall that a symmetric matrix A is *positive semi-definite*, if $a^\top A a \geq 0$ for all $a \in \mathbb{R}^k$, and *positive definite*, if $a^\top A a > 0$ for all $a \neq 0$; a positive definite matrix has a positive definite inverse.

Proposition 14.4. *Suppose the conditions from section 11.1 hold. Then the following statements hold.*

1. *The matrix $\mathcal{I}_n(\theta)$ is symmetric and positive semi-definite. In particular we have $a^\top \mathcal{I}_n(\theta) a = \text{Var}_\theta(a^\top \nabla \ell(\theta))$ for all $a \in \mathbb{R}^k$.*
2. *Information is additive: $\mathcal{I}_n(\theta) = n \mathcal{I}(\theta)$.*
3. *Let $\varphi \mapsto \theta = h(\varphi)$ be a continuously differentiable bijection between open subsets of $\mathbb{R}^k$ — here φ is a parameter vector, the several-parameter counterpart of the $\psi = g(\theta)$ in lecture 5 — and let $J(\varphi)$ be the invertible Jacobian matrix with entries $J_{jl} = \partial \theta_j / \partial \varphi_l$. Then the information matrix for φ is given by*

$$\mathcal{I}_n^\varphi(\varphi) = J(\varphi)^\top \mathcal{I}_n(\theta) J(\varphi) \quad \text{for } \theta = h(\varphi).$$

Proof. For the first statement, the symmetry of the matrix follows from the fact that $\text{Cov}(Y_j, Y_l) = \text{Cov}(Y_l, Y_j)$ for all j and l . For $a \in \mathbb{R}^k$, the variance of the linear combination $a^\top \nabla \ell(\theta)$ is given by

$$\text{Var}_\theta \left(\sum_{j=1}^k a_j \partial_j \ell(\theta) \right) = \sum_{j=1}^k \sum_{l=1}^k a_j a_l \mathcal{I}_n(\theta)_{jl} = a^\top \mathcal{I}_n(\theta) a,$$

and thus is non-negative. The second statement is proved by using the fact that the score vector of the sample is the sum of n i.i.d. score vectors of the individual observations and that the covariance matrix of the sum of independent random vectors equals the sum of the covariance matrices. Finally, for the third statement, we note that the log-likelihood in the new parametrisation is $\tilde{\ell}(\varphi) = \ell(h(\varphi))$ and thus we find

$$\frac{\partial \tilde{\ell}}{\partial \varphi_l}(\varphi) = \sum_{j=1}^k \partial_j \ell(\theta) \frac{\partial \theta_j}{\partial \varphi_l} = \sum_{j=1}^k J_{jl} \partial_j \ell(\theta),$$

i.e. $\nabla \tilde{\ell}(\varphi) = J^\top \nabla \ell(\theta)$. Since J is a fixed matrix for fixed φ , the covariance matrix of $J^\top \nabla \ell(\theta)$ is given by $J^\top \mathcal{I}_n(\theta) J$. This completes the proof. $\square$

For $k = 1$ the third statement is a consequence of proposition 11.10: If $\psi = g(\theta)$ we have $h = g^{-1}$ and $J = 1/g'(\theta)$. Positive definiteness is violated in cases where the log-likelihood does not respond to changes of θ in some direction, for example when the model is not identifiable in the sense of definition 4.10. From now on we assume that $\mathcal{I}_n(\theta)$ is positive definite so that the inverse exists.

14.3 The Cramer–Rao Bound for Vector Parameters

Theorem 13.2 bounds the variance of an unbiased estimator for a scalar parameter from below. For a vector parameter the natural object to bound is the covariance matrix, and the natural reading of “at least” between symmetric matrices is that the difference is positive semi-definite.

Theorem 14.5. *Suppose we have the conditions of section 11.1 and a positive definite matrix $\mathcal{I}_n(\theta)$. Let $T = (T_1, \dots, T_k)$ be an unbiased estimator for θ , so that $\mathbb{E}_\theta(T_j) = \theta_j$ for all j and all θ , where the variances are finite. Then the matrix*

$$\text{Cov}_\theta(T) - \mathcal{I}_n(\theta)^{-1}$$

is positive semi-definite for all $\theta \in \Theta$. In particular we have

$$\text{Var}_\theta(T_j) \geq (\mathcal{I}_n(\theta)^{-1})_{jj} \quad \text{and} \quad \text{Var}_\theta(a^\top T) \geq a^\top \mathcal{I}_n(\theta)^{-1} a$$

for $j = 1, \dots, k$ and for all $a \in \mathbb{R}^k$.

Proof. We only sketch the argument, which is the proof of theorem 13.2 for linear combinations. Differentiating $\mathbb{E}_\theta(T_j) = \theta_j$ under the integral sign, as in that proof, gives $\text{Cov}_\theta(T_j, \partial_l \ell(\theta)) = 1$ for $l = j$ and 0 otherwise, and thus by bilinearity $\text{Cov}_\theta(a^\top T, b^\top \nabla \ell(\theta)) = a^\top b$ for all $a, b \in \mathbb{R}^k$. Using the Cauchy–Schwarz inequality, lemma A.8, and the first statement of proposition 14.4, we find

$$(a^\top b)^2 \leq (a^\top \text{Cov}_\theta(T) a) (b^\top \mathcal{I}_n(\theta) b).$$

Choosing $b = \mathcal{I}_n(\theta)^{-1}a$, both $a^\top b$ and $b^\top \mathcal{I}_n(\theta)b$ equal $a^\top \mathcal{I}_n(\theta)^{-1}a$, which is positive for $a \neq 0$, and dividing by this value gives $a^\top \text{Cov}_\theta(T) a \geq a^\top \mathcal{I}_n(\theta)^{-1}a$. This is the claim, since $a^\top \text{Cov}_\theta(T) a = \text{Var}_\theta(a^\top T)$, and for the diagonal case we have $a = e_j$, with the j th unit vector. This completes the sketch of the proof. $\square$

The bound is almost always applied in its diagonal form: The variance of any unbiased estimator for θ_j , when the other components are unknown, is at least the j th diagonal entry of the inverse information matrix. Note that this is the diagonal of the inverse, not the reciprocal of the diagonal; we will see the difference in section 14.5. An estimator which achieves the bound for θ_j is said to be efficient for θ_j , as in lecture 13. The bound for a function $g(\theta)$ with gradient ∇g is $\nabla g(\theta)^\top \mathcal{I}_n(\theta)^{-1} \nabla g(\theta)$, which includes the statement of theorem 13.3 for $k = 1$, as well as the linear case $g(\theta) = a^\top \theta$, which the theorem already covers. The result is also used in exercise 14.3.

14.4 The Normal Distribution

The only example of a distribution which all students should be able to reproduce without notes is the normal distribution with unknown parameters. We will consider this example in detail.

Example 14.6. Let $X_1, \dots, X_n$ be i.i.d. $N(\mu, \sigma^2)$ where $\theta = (\mu, \sigma^2)$, both unknown. From example 5.4 we know that the log-likelihood function is

$$\ell(\mu, \sigma^2) = -\frac{n}{2} \log(2\pi) - \frac{n}{2} \log \sigma^2 - \frac{1}{2\sigma^2} \sum_{i=1}^n (X_i - \mu)^2,$$

and the first partial derivatives are given by

$$\frac{\partial \ell}{\partial \mu} = \frac{1}{\sigma^2} \sum_{i=1}^n (X_i - \mu), \quad \frac{\partial \ell}{\partial \sigma^2} = -\frac{n}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{i=1}^n (X_i - \mu)^2.$$

Taking second derivatives, where we treat σ^2 as a single variable, we find the three second derivatives as

$$\begin{aligned} \frac{\partial^2 \ell}{\partial \mu^2} &= -\frac{n}{\sigma^2}, & \frac{\partial^2 \ell}{\partial \mu \partial \sigma^2} &= -\frac{1}{\sigma^4} \sum_{i=1}^n (X_i - \mu), \\ \frac{\partial^2 \ell}{\partial (\sigma^2)^2} &= \frac{n}{2\sigma^4} - \frac{1}{\sigma^6} \sum_{i=1}^n (X_i - \mu)^2. \end{aligned}$$

To compute the expectations, we use the fact that $\mathbb{E}_\theta(X_i - \mu) = 0$ and $\mathbb{E}_\theta((X_i - \mu)^2) = \sigma^2$. The mixed derivative has expectation zero. The last term has expectation $n/(2\sigma^4) - n\sigma^2/\sigma^6 = -n/(2\sigma^4)$. Thus, by proposition 14.3, we find the information matrix and its inverse to be

$$\mathcal{I}_n(\mu, \sigma^2) = \begin{pmatrix} \frac{n}{\sigma^2} & 0 \\ 0 & \frac{n}{2\sigma^4} \end{pmatrix} \quad \text{and} \quad \mathcal{I}_n(\mu, \sigma^2)^{-1} = \begin{pmatrix} \frac{\sigma^2}{n} & 0 \\ 0 & \frac{2\sigma^4}{n} \end{pmatrix}. \quad (14.1)$$

We should commit these matrices to memory. The upper left entry is the information n/σ^2 about the mean, where the variance is known from example 11.7.

The sample mean, with variance σ^2/n , achieves the bound for μ . For the variance the bound is $2\sigma^4/n$. The UMVUE S^2 from example 10.7 has variance $2\sigma^4/(n-1)$, as shown in the computation of section 2.3, and thus does not achieve the bound. Since S^2 is the UMVUE, no unbiased estimator for σ^2 achieves the bound when μ is unknown. If we had parametrised the model by the standard deviation instead, *i.e.* if we had $\varphi = (\mu, \sigma)$, then the third statement of proposition 14.4, with $J = \text{diag}(1, 2\sigma)$, gives $\mathcal{I}_n^\varphi(\mu, \sigma) = \text{diag}(n/\sigma^2, 2n/\sigma^2)$. See exercise 14.1 for details.

The information matrix for the normal distribution is diagonal. The following section explains the significance of this fact and what happens when this is not the case.

14.5 Nuisance Parameters

Assume that we are interested in θ_1 only, and that the remaining components $\theta_2, \dots, \theta_k$ are *nuisance parameters*, *i.e.* unknown but not of interest. If the values of the nuisance parameters were known, the Cramer–Rao bound for θ_1 would be $1/\mathcal{I}_n(\theta)_{11}$. Since the values are unknown, the bound is $(\mathcal{I}_n(\theta)^{-1})_{11}$ by theorem 14.5. For $k = 2$ we can compare the two bounds explicitly. For this we write

$$\mathcal{I}_n(\theta) = \begin{pmatrix} a & b \\ b & c \end{pmatrix} \quad \text{and find} \quad \mathcal{I}_n(\theta)^{-1} = \frac{1}{ac - b^2} \begin{pmatrix} c & -b \\ -b & a \end{pmatrix},$$

where $ac - b^2 > 0$ since the matrix is positive definite, and

$$(\mathcal{I}_n(\theta)^{-1})_{11} = \frac{c}{ac - b^2} = \frac{1}{a} \cdot \frac{1}{1 - \rho^2}, \quad \text{where } \rho = \frac{b}{\sqrt{ac}} \quad (14.2)$$

is the correlation between the two components of the score vector. Since $0 \leq \rho^2 < 1$, the bound where the nuisance parameter is unknown is $1/(1 - \rho^2)$ times larger than the bound where the nuisance parameter is known, and both bounds are equal if and only if $b = 0$. Two parameters with $\mathcal{I}_n(\theta)_{12} = 0$ are said to be *orthogonal*; in this case, not knowing θ_2 does not incur any additional cost for the bound when estimating θ_1 . The same inequality $(\mathcal{I}_n(\theta)^{-1})_{jj} \geq 1/\mathcal{I}_n(\theta)_{jj}$ holds for all j and all k : not knowing a parameter can never be beneficial.

The normal distribution is the orthogonal case: in example 14.6 the mixed derivative has mean zero, and the bound for μ is σ^2/n whether or not σ^2 is known. This is a special feature of the normal distribution, not the rule. For the gamma distribution with shape $\alpha > 0$ and rate $\beta > 0$ both unknown, as in exercise 7.3, all three second derivatives of the log-likelihood $\ell_1(\alpha, \beta)$ of one observation are non-random, and the information matrix of one observation is

$$\mathcal{I}(\alpha, \beta) = \begin{pmatrix} \psi'(\alpha) & -\frac{1}{\beta} \\ -\frac{1}{\beta} & \frac{\alpha}{\beta^2} \end{pmatrix},$$

where ψ' , the second derivative of $\log \Gamma$, is the *trigamma function*; exercise 14.2 derives this matrix. The off-diagonal entry is non-zero and by (14.2) we have

$\rho^2 = 1/(\alpha\psi'(\alpha))$; for $\alpha = 2$, with $\psi'(2) = \pi^2/6 - 1 \approx 0.645$, this gives $1/(1-\rho^2) \approx 4.45$, so that the bound with the other parameter unknown is about four and a half times the bound with the other parameter known.

The inverse information matrix is needed when we estimate several parameters simultaneously. In lecture 17 we will see that, for large sample size, the maximum likelihood estimator is approximately normally distributed with covariance matrix $\mathcal{I}_n(\theta_0)^{-1}$. In lecture 18 we will see how we can numerically find the standard errors of the estimator by inverting the observed Hessian matrix, and in lecture 23 we will see that the number of parameters k can be seen as the number of degrees of freedom in this system.

Summary.

- For a vector parameter the score is the gradient $\nabla\ell(\theta)$, and the Fisher information matrix $\mathcal{I}_n(\theta)$ is its covariance matrix, equal to minus the expected Hessian of the log-likelihood.
- The information matrix is symmetric and positive semi-definite, adds over independent observations, $\mathcal{I}_n = n\mathcal{I}$, and transforms as $J^\top \mathcal{I}_n J$ under a change of parametrisation with Jacobian J .
- Cramer–Rao: for an unbiased estimator T for θ the matrix $\text{Cov}_\theta(T) - \mathcal{I}_n(\theta)^{-1}$ is positive semi-definite; in particular $\text{Var}_\theta(T_j) \geq (\mathcal{I}_n(\theta)^{-1})_{jj}$.
- For $N(\mu, \sigma^2)$ with both parameters unknown, $\mathcal{I}_n = \text{diag}(n/\sigma^2, n/(2\sigma^4))$ with inverse $\text{diag}(\sigma^2/n, 2\sigma^4/n)$.
- The diagonal of the inverse dominates the reciprocal of the diagonal, with equality for orthogonal parameters: nuisance parameters inflate the bound by $1/(1-\rho^2)$ in the two-parameter case, which is one for the normal distribution and about 4.45 for the gamma distribution with shape 2.

Exercise 14.1. Let $X_1, \dots, X_n$ be i.i.d. $N(\mu, \sigma^2)$ and consider the parametrisation $\varphi = (\mu, \sigma)$, where the standard deviation is used.

1. Compute the Jacobian matrix J of $(\mu, \sigma) \mapsto (\mu, \sigma^2)$ and use proposition 14.4 and (14.1) to show that

$$\mathcal{I}_n^\varphi(\mu, \sigma) = \text{diag}(n/\sigma^2, 2n/\sigma^2).$$

2. By taking derivatives of the log-likelihood $\ell(\mu, \sigma)$ twice and taking expectations, verify this result directly.
3. Write down the Cramer–Rao bound for an unbiased estimator for σ . Using Jensen’s inequality, lemma A.7, show that the sample standard deviation $S = \sqrt{S^2}$ is not unbiased for σ and thus cannot be compared to $\text{Var}(S)$ directly.

Exercise 14.2. Consider the gamma distribution with shape $\alpha > 0$ and rate $\beta > 0$. The log-likelihood of observing one value is $\ell_1(\alpha, \beta) = \alpha \log \beta - \log \Gamma(\alpha) + (\alpha - 1) \log X - \beta X$ from exercise 7.3, and the information matrix $\mathcal{I}(\alpha, \beta)$ is given in section 14.5.

1. Compute the first and second partial derivatives of $\ell_1(\alpha, \beta)$ and check the matrix.
2. Assume that the shape parameter is known to be $\alpha = 2$. In this case, β is the only parameter. From the matrix derive the Fisher information about β . Why is this a diagonal element of $\mathcal{I}(\alpha, \beta)$ and not of the inverse matrix?
3. Exercise 11.2 considers parametrisation of the same model in terms of the scale parameter $\theta = 1/\beta$. Use proposition 11.10 to derive the value $\mathcal{I}(\theta) = 2/\theta^2$ found in that exercise.
4. Now assume that both parameters are unknown. Show that the Cramer–Rao bound for an unbiased estimator for β from a sample of size n is given by

$$\frac{\beta^2 \psi'(\alpha)}{n(\alpha \psi'(\alpha) - 1)},$$

and determine the ratio of this bound to the bound for known shape parameter $\alpha = 2$.

Exercise 14.3. Let $X_1, \dots, X_m$ be i.i.d. with distribution $N(\mu_1, \sigma^2)$ and, independently of these, let $Y_1, \dots, Y_n$ be i.i.d. with distribution $N(\mu_2, \sigma^2)$, where the common variance σ^2 is known and $\theta = (\mu_1, \mu_2)$ is unknown.

1. Write down the log-likelihood of all $m + n$ observations and show that the information matrix is $\mathcal{I}(\mu_1, \mu_2) = \text{diag}(m/\sigma^2, n/\sigma^2)$.
2. Using theorem 14.5, find the Cramer–Rao bound for an unbiased estimator for the difference $\mu_1 - \mu_2$. Show that $\bar{X} - \bar{Y}$ achieves this bound.
3. Assume now that σ^2 is also unknown, *i.e.* when we have $\theta = (\mu_1, \mu_2, \sigma^2)$. Show that the 3×3 information matrix is still diagonal and discuss the implications for the estimation of $\mu_1 - \mu_2$ when σ^2 is unknown.

Lecture 15

Cramer–Rao and Efficiency in R

This is the R session for week 5. Last week in lectures we proved the Cramer–Rao inequality (theorem 13.2). This week we will try to translate this bound into numbers: we will consider an efficient estimator with variance on the bound, an inefficient estimator and the reason for its inefficiency, and the uniform model where the bound is beaten because the theorem does not apply. As usual, appendix B has a written version of the material we will discuss, and we will try to predict the outcome of each simulation before we perform the simulation.

15.1 An Efficient Estimator Against the Bound

Our model of choice for this section is $X_1, \dots, X_n$ i.i.d. Poisson with mean $\theta = 3$. From example 11.6 we know that the Cramer–Rao bound for an unbiased estimator for θ is θ/n . The sample mean, $\bar{X}$, is an unbiased estimator for θ with $\text{Var}_\theta(\bar{X}) = \theta/n$ (see lemma 2.2) and thus, by definition 13.4, is an efficient estimator. We use simulation to verify this claim, using the loop over `ns` from lecture 9. The command `var` can be used to compute the sample variance of a vector of estimates, and dividing the bound by this value allows us to estimate the efficiency of the estimator. The full script for this section, with lines intentionally shuffled (except for the function body and the loop which are left as one piece), is given in the following box. Before you can run the script, you will need to arrange the lines yourself.

```
bound <- theta / ns
data.frame(n = ns, variance = signif(v, 3), bound = signif(bound, 3),
           efficiency = round(bound / v, 3))
est_pois <- function(n, theta) {
  mean(rpois(n, lambda = theta))
}
set.seed(2026)
for (j in seq_along(ns)) {
  v[j] <- var(replicate(10000, est_pois(ns[j], theta)))
}
theta <- 3
v <- numeric(length(ns))
ns <- c(5, 10, 20, 50, 100, 200)
```

As in lectures 6 and 9, all names must be defined before they are used: the function, `theta` and `ns` come first, `v` needs `length(ns)`, the seed must be set before the loop, and the table must be created last. Before you try the reconstructed script, make a prediction for the efficiency column. Will all entries be exactly 1? Can an entry be larger than 1, despite the fact that theorem 13.2 states that no unbiased estimator can have a variance smaller than the bound?

```
est_pois <- function(n, theta) {
  mean(rpois(n, lambda = theta))
}
theta <- 3
ns <- c(5, 10, 20, 50, 100, 200)
v <- numeric(length(ns))
set.seed(2026)
for (j in seq_along(ns)) {
  v[j] <- var(replicate(10000, est_pois(ns[j], theta)))
}
bound <- theta / ns
data.frame(n = ns, variance = signif(v, 3), bound = signif(bound, 3),
           efficiency = round(bound / v, 3))
```

```
##      n variance bound efficiency
## 1    5   0.5970 0.600      1.005
## 2   10   0.3000 0.300      1.001
## 3   20   0.1510 0.150      0.995
## 4   50   0.0590 0.060      1.017
## 5  100   0.0297 0.030      1.010
## 6  200   0.0150 0.015      0.999
```

The empirical variances in the table follow the bound for all sample sizes, and the estimated efficiencies are clustered around 1 with half of the entries above this value. The following argument shows that this is no contradiction: the statement of the theorem only bounds the true variance $\text{Var}_\theta(\bar{X})$, whereas the table shows the variance of 10000 random estimates. The random sample is off the true value by one or two percent in either direction, as shown in lecture 6. An estimated efficiency of 1.017 means that, within the limits of error of the simulation, the estimator is efficient.

15.2 Attaining and Missing the Bound

In this section we show that not every unbiased estimator is efficient. As an example, from exercise 13.4 we know that the bias-corrected maximum likelihood estimator $\tilde{\theta} = (n-1)/\sum_i X_i$ for the exponential rate θ is unbiased and has variance $\theta^2/(n-2)$. On the other hand, the bound from example 11.8 is θ^2/n and thus the efficiency of this estimator is $(n-2)/n$. We can perform a similar analysis as in the previous section, this time for $\tilde{\theta}$ with $\theta = 2$: To do this we write a new function (below `est_pois` in the code) to compute the estimator, keep the loop as before, adjust the range of `ns` to 5 to 100, and replace the bound by `theta^2 / ns`. The resulting script is `R/S15-efficiency.R`. Before you can run the script, you need to write down $(n-2)/n$ for $n = 5, 10, 20, 50$ and 100. Considering the simulation error we have observed, for which sample sizes could the simulation be expected to distinguish $\tilde{\theta}$ from an efficient estimator?

```
est_exp <- function(n, theta) {
  (n - 1) / sum(rexp(n, rate = theta))
}
data.frame(n = ns, efficiency = round(bound / v, 3), exact = (ns - 2) /
  ns)
```

```
##      n efficiency exact
## 1    5      0.607 0.60
## 2   10      0.803 0.80
## 3   20      0.903 0.90
## 4   50      0.929 0.96
## 5  100      1.008 0.98
```

For $n = 5, 10$ and 20 the simulation coincides with the exact efficiencies up to two decimal digits. For $n = 50$ and 100 the difference to 1 is 4% and 2%, comparable to the size of the simulation error. Thus, the last two rows do not give useful results and we would need significantly more than 10000 repetitions to resolve these results. The loss of efficiency is a constant factor which goes to zero as n increases: $\tilde{\theta}$ is asymptotically efficient in the sense of lecture 17, but not efficient.

Where does the difference come from? The proof of theorem 13.2 in lecture 13 uses the Cauchy–Schwarz inequality between the estimator T and the score $\ell'(\theta)$. Dividing through shows that the efficiency of T is given by the squared correlation between T and the score. This equals 1, if and only if $T - \theta$ is a multiple of the score. For the Poisson sample $\bar{X} - \theta$ equals θ/n times the score, and thus the correlation is 1 without any calculation. For the exponential sample $\tilde{\theta} - \theta$ is a function of the score, but not a linear one. The following function computes the estimation error and the score for a given sample of size $n = 5$. The command `cor` can be used to compute the correlation of two vectors. Predict the value of the squared correlation before you run the code.

```
n <- 5
err_score <- function(n, theta) {
  x <- rexp(n, rate = theta)
  c((n - 1) / sum(x) - theta, n / theta - sum(x))
}
set.seed(2026)
expo <- replicate(2000, err_score(n, 2))
cor(expo[1, ], expo[2, ])^2
```

```
## [1] 0.6147813
```

The squared correlation 0.615 coincides with the efficiency $3/5$ from the table, up to simulation error. A plot of `expo[1, ]` against `expo[2, ]` shows that the 2000 points lie on a curve instead of on a straight line. The curvature of this curve is evidence that the bound does not account for all of the variance.

15.3 Beating the Bound Without Breaking It

Example 13.6 is this week's counterexample: for $X_1, \dots, X_n$ i.i.d. uniformly distributed on the set $(0, \theta)$, the estimator $(n + 1) \max_i X_i / n$ is unbiased and

has variance $\theta^2/(n(n+2))$, whereas the Cramer–Rao formula, using the value $\mathbb{E}_\theta(\ell'(\theta)^2) = 1/\theta^2$ from example 11.11, would give θ^2/n . The following function computes, for the same sample, the bias-corrected maximum and the method of moments estimator $2\bar{X}$, both of which are unbiased. Thus, the variances are what the theorem would be about, if it held. The command `apply` can be used to apply a function to every row (second argument 1) or column (2) of a matrix, in this case to compute the variance `var` for each row of estimates. Before you try the code, make a prediction: which of the two columns will be below the column θ^2/n ? Will the other column be above it?

```
est_both <- function(n, theta) {
  x <- runif(n, min = 0, max = theta)
  c((n + 1) * max(x) / n, 2 * mean(x))
}
theta <- 1
ns <- c(5, 10, 20, 50, 100, 200, 500, 1000)
vars <- matrix(0, nrow = length(ns), ncol = 2)
set.seed(2026)
for (j in seq_along(ns)) {
  ests <- replicate(10000, est_both(ns[j], theta))
  vars[j, ] <- apply(ests, 1, var)
}
rownames(vars) <- ns
colnames(vars) <- c("(n + 1) max / n", "2 * mean")
signif(cbind(vars, "theta^2 / n" = theta^2 / ns), 3)
```

```
##      (n + 1) max / n  2 * mean  theta^2 / n
## 5          2.76e-02 0.065900    0.200
## 10         8.35e-03 0.033300    0.100
## 20         2.37e-03 0.016600    0.050
## 50         3.79e-04 0.006610    0.020
## 100        9.59e-05 0.003310    0.010
## 200        2.59e-05 0.001660    0.005
## 500        3.93e-06 0.000668    0.002
## 1000       1.00e-06 0.000331    0.001
```

The variance of $2\bar{X}$ is $\theta^2/(3n)$ by exercise 2.2, one third of the formal bound at every n . The first column is below the formal bound throughout and falls away from it: at $n = 1000$ the variance of the corrected maximum is 10^{-6} , one thousandth of θ^2/n . As in lecture 9, a plot on log-log axes shows the rates. Predict the slopes of the two sets of points and of the dashed line θ^2/n .

```
plot(ns, vars[, 1], log = "xy", pch = 16, xlab = "n", ylab = "variance",
     ylim = range(vars, theta^2 / ns))
points(ns, vars[, 2], pch = 1)
lines(ns, theta^2 / ns, lty = 2)
lines(ns, theta^2 / (ns * (ns + 2)), lty = 3)
legend("bottomleft", legend = c("(n + 1) max / n", "2 * mean"),
     pch = c(16, 1))
```

In figure 15.1 the open circles are parallel to the dashed line with slope -1 , corresponding to the usual rate, and the filled circles follow the dotted line with slope -2 : figure 9.2 again, with the value θ^2/n added and filled circles




Figure 15.1: Simulated variances for $(n+1) \max_i X_i/n$ (filled circles) and $2\bar{X}$ (open circles), for the uniform parameter $\theta = 1$, on logarithmic axes. The dashed line shows the value θ^2/n from the Cramer–Rao formula. The dotted line shows the exact variance $\theta^2/(n(n+2))$ from example 13.6.

placed below it from the first sample size on. Nothing is broken. The regularity conditions of theorem 13.2 are violated by the uniform distribution, since the support of the distribution moves with θ : example 11.11 shows that the score does not have mean zero and that the differentiation under the integral sign used in the proof is not possible. Thus, the value θ^2/n cannot be a bound and indeed it is a formula used outside its domain of validity, and an estimator with variance smaller than this value is called super-efficient. The same failure of regularity is shown in exercise 17.3 to cause the rate $1/n$ of the maximum.

15.4 A Task to Take Home

The following task allows you to practice the methods we have learned in this lecture, in a style similar to the practical report. First, you should repeat the study of section 15.1 for the Bernoulli distribution with $\theta = 0.3$, where $\bar{X}$ is unbiased and the bound is $\theta(1-\theta)/n$ as shown in example 11.5. Create a table which shows the estimated efficiencies for the sample sizes in `ns`. Next, consider the sample variance S^2 , computed using `var`, as an estimator for σ^2 for a normal sample with $\mu = 0$ and $\sigma^2 = 1$. The bound $2\sigma^4/n$ for this estimator is given in example 14.6. Compute the exact variance of S^2 using the chi-squared distribution of $(n-1)S^2/\sigma^2$ from proposition A.3. Compare the exact efficiency of your estimator to the values in the table. Finally, in three sentences, discuss which of the two estimators is efficient, how much less efficient the other one is, and how many repetitions of your simulation would be required to see the less efficient estimator at $n = 100$.

Before the next session, you should type the commands used in this session, and should check that you get the same results as the text.

Summary.

- The written reference for the functions `var` and `apply` is appendix [B](#).
- The estimated efficiency, *i.e.* the ratio of the Cramer–Rao bound to the sample variance of the simulated estimates, fluctuates around the true efficiency by one or two percent for 10000 repetitions and can be greater than 1. A simulation is only able to detect an efficiency loss, if the loss is bigger than this error.
- The Poisson mean is efficient for all n . The bias-corrected exponential-rate estimator has efficiency $(n - 2)/n$. The efficiency is the square of the correlation between the estimator and the score, and equals 1, if and only if the estimator is a linear function of the score.
- For the uniform parameter, the variance of $(n + 1) \max_i X_i/n$ decreases as $1/n^2$, and is smaller than the value θ^2/n from the Cramer–Rao formula, since the regularity conditions are violated. The formula is not a bound in this case.

Lecture 16

Consistency of the MLE

In lecture 5 we have seen that the maximum likelihood estimator can be biased, that it may not be unique and that it may not even exist. The following two lectures present the results which, despite these problems, allow us to justify the use of the MLE: for large sample size the maximum likelihood estimator converges to the true parameter value, and does so with minimal spread. In the current lecture we will prove the first of these two results, the consistency of the MLE, from a set of assumptions which we will explicitly state. The proof of the asymptotic normality in lecture 17 builds on these same assumptions and on the consistency result. The numerical aspects, concerning estimates and standard errors, will be discussed in lecture 18.

16.1 The Regularity Conditions

Throughout this and the following lectures, $X_1, \dots, X_n$ will be i.i.d. with density or probability weights $f(x; \theta)$, where the parameter θ is a one-dimensional parameter and $\Theta \subseteq \mathbb{R}$ is an open interval. The true parameter value is denoted by θ_0 and $\mathbb{P}_{\theta_0}$ and $\mathbb{E}_{\theta_0}$ are probabilities and expectations w.r.t. this value. As in lecture 4, the log-likelihood is given by

$$\ell(\theta) = \sum_{i=1}^n \ell_i(\theta), \quad \text{where } \ell_i(\theta) = \log f(X_i; \theta),$$

and for random data the $\ell_i(\theta)$ are i.i.d. random variables. The proofs in these two lectures will make use of the additive structure of the log-likelihood, allowing for applications of the law of large numbers and the central limit theorem for the likelihood.

The large-sample results are true for most models of interest, but there are exceptions; the uniform distribution from example 5.5 is a case in which the results do not hold. We start our discussion by stating the assumptions required for the theorems to hold.

Definition 16.1. The model $(f(x; \theta) \mid \theta \in \Theta)$ satisfies the *regularity conditions*, if the following four statements hold.

- (R1) The parameter is identifiable in the sense of definition 4.10, *i.e.* different values of θ lead to different distributions.
- (R2) The support $\{x \mid f(x; \theta) > 0\}$ does not depend on θ .

- (R3) For every x the function $\theta \mapsto f(x; \theta)$ is three times continuously differentiable on Θ , the order of differentiation with respect to θ and integration (or summation) with respect to x can be interchanged twice, and the Fisher information $\mathcal{I}(\theta)$ of a single observation is a continuous function of θ with $0 < \mathcal{I}(\theta) < \infty$ for all $\theta \in \Theta$.
- (R4) For every $\theta_0 \in \Theta$ there are a $\delta > 0$ and a function $M(x)$ such that $\mathbb{E}_{\theta_0}(M(X_1)) < \infty$ and

$$\left| \frac{\partial^3}{\partial \theta^3} \log f(x; \theta) \right| \leq M(x) \quad \text{for all } x \text{ and all } \theta \text{ with } |\theta - \theta_0| < \delta.$$

The four conditions play different roles. Condition (R1) is needed for the question to make sense: if two parameter values give the same distribution, no estimator can converge to both. Condition (R2) excludes models like the uniform distribution, where the parameter moves the edge of the support and the likelihood jumps to zero; the Cramer–Rao bound fails for such models, as we have seen in lecture 13, and the arguments of this lecture fail for them too. Condition (R3) is the smoothness needed to differentiate the log-likelihood, and under (R2) and (R3) it makes the two identities of lecture 11 available (lemma 11.1 and theorem 11.3),

$$\mathbb{E}_{\theta}(\ell'_i(\theta)) = 0 \quad \text{and} \quad \mathcal{I}(\theta) = \text{Var}_{\theta}(\ell'_i(\theta)) = -\mathbb{E}_{\theta}(\ell''_i(\theta)). \quad (16.1)$$

Condition (R4) is not needed in this lecture at all: it controls the remainder term in the Taylor expansion in the proof of asymptotic normality in lecture 17, and it is listed here so that both lectures work from the same definition. Although the condition looks technical, it is usually easy to verify, because the third derivative of the log-likelihood is an explicit function of x and θ ; exercise 16.1 does this for the exponential distribution.

The usual families of distributions from tables 1.1 and 1.2, with the exception of the uniform distribution, satisfy the regularity conditions (with m known for the binomial distribution). The same is true for the t location model from section 5.3, despite the fact that the likelihood equation may have more than one solution. We will return to this problem in section 16.4.

16.2 The Likelihood Ratio Detects the Truth

If θ_0 is the true value, it seems natural that for a large sample the likelihood $L(\theta_0)$ should be larger than $L(\theta)$ for any wrong value θ . This fact is the engine of the consistency proof, and thus we prove it first.

Theorem 16.2. *Suppose (R1) and (R2). Let $\theta \neq \theta_0$ and assume that $Y_1 = \log(f(X_1; \theta)/f(X_1; \theta_0))$ has finite expectation under $\mathbb{P}_{\theta_0}$. Then*

$$\lim_{n \rightarrow \infty} \mathbb{P}_{\theta_0}(L(\theta_0) > L(\theta)) = 1.$$

Proof. Using logarithms and dividing by n , we can rewrite the event $L(\theta_0) > L(\theta)$ as

$$\frac{1}{n} \sum_{i=1}^n \log \frac{f(X_i; \theta)}{f(X_i; \theta_0)} < 0,$$

where, by (R2), the ratios are well defined and positive for all observations. The left-hand side is the sample mean of the i.i.d. random variables $Y_i = \log(f(X_i; \theta)/f(X_i; \theta_0))$ and by the law of large numbers, theorem A.14, converges in probability to $\mathbb{E}_{\theta_0}(Y_1)$. Therefore, we only need to show that this expectation is strictly negative, so that the sample mean is negative with probability approaching one.

Since the logarithm is strictly concave, the function $-\log$ is convex and we can use Jensen's inequality, lemma A.7, to get

$$\mathbb{E}_{\theta_0}(Y_1) = \mathbb{E}_{\theta_0}\left(\log \frac{f(X_1; \theta)}{f(X_1; \theta_0)}\right) \leq \log \mathbb{E}_{\theta_0}\left(\frac{f(X_1; \theta)}{f(X_1; \theta_0)}\right),$$

and the inequality is strict unless the ratio $f(X_1; \theta)/f(X_1; \theta_0)$ is constant with probability one. If the ratio was constant, we would have $f(x; \theta) = c f(x; \theta_0)$ on the common support and, since both functions have integral one, this would imply $c = 1$, *i.e.* the two distributions would be the same which (R1) excludes for $\theta \neq \theta_0$. Thus the inequality is strict. Writing the expectation on the right-hand side as an integral over the common support, we can cancel the density $f(x; \theta_0)$ and get

$$\mathbb{E}_{\theta_0}\left(\frac{f(X_1; \theta)}{f(X_1; \theta_0)}\right) = \int \frac{f(x; \theta)}{f(x; \theta_0)} f(x; \theta_0) dx = \int f(x; \theta) dx = 1,$$

with a sum instead of an integral for a discrete model. Thus we find $\mathbb{E}_{\theta_0}(Y_1) < \log 1 = 0$ and the proof is complete.

The assumption that $\mathbb{E}_{\theta_0}(Y_1)$ is finite excludes nothing of interest: the only alternative is $\mathbb{E}_{\theta_0}(Y_1) = -\infty$, and in that case the conclusion holds all the more easily. Replacing Y_i by $\max(Y_i, -M)$ for a constant $M > 0$ gives i.i.d. variables with finite expectation, to which the law of large numbers applies, and that expectation is negative once M is large enough; since the sample mean of the Y_i is at most that of the truncated variables, it is again negative with probability approaching one. $\square$

The quantity $-\mathbb{E}_{\theta_0}(Y_1)$ from the proof is the *Kullback–Leibler divergence* of the wrong model from the true model, and the theorem states that on average each observation increases the log-likelihood of the truth by this fixed positive amount, compared to the log-likelihood of any of the competitors. It is important to note that the theorem only compares fixed values and does not state that $L(\theta_0)$ will eventually be bigger than $L(\theta)$ for all wrong θ , or even that it is a statement about the maximiser. In the next section we will see a proof that allows us to conclude that the statement holds for the maximiser.

16.3 Consistency

From definition 2.8 we know that an estimator $\hat{\theta}_n$ is consistent for θ , if $\hat{\theta}_n \xrightarrow{P} \theta$ for all possible true values. In lecture 2 we have already seen how to prove consistency using the mean squared error. For the maximum likelihood estimator this approach is not feasible in general, because there is no general formula for the mean squared error. Instead, we will use the likelihood equation from lecture 5 to show that with probability tending to one there is a root of the equation which is close to θ_0 and, if the root is unique, this root coincides with the MLE.

Theorem 16.3. *Suppose that the regularity conditions (R1), (R2) and (R3) hold, that $\log(f(X_1; \theta)/f(X_1; \theta'))$ has finite mean under $\mathbb{P}_{\theta'}$ for all $\theta, \theta' \in \Theta$, and that for every sample size $n \geq 1$ and every sample with coordinates in the support of the model, which by (R2) is the same for every θ , the likelihood equation $\ell'(\theta) = 0$ has a unique root which coincides with the maximum likelihood estimator $\hat{\theta}_n$. Then $\hat{\theta}_n$ is a consistent estimator for θ .*

Proof. Let the true value be θ_0 and let $\varepsilon > 0$ be small enough that the interval $(\theta_0 - \varepsilon, \theta_0 + \varepsilon)$ is contained in the open set Θ . Let $\theta_1 = \theta_0 - \varepsilon$ and $\theta_2 = \theta_0 + \varepsilon$ and consider the event

$$S_\varepsilon = \{L(\theta_1) < L(\theta_0) \text{ and } L(\theta_2) < L(\theta_0)\}.$$

Since the event S_ε is the intersection of two events of the form considered in theorem 16.2, both of which have probability tending to one, we have $\mathbb{P}_{\theta_0}(S_\varepsilon) \rightarrow 1$ as $n \rightarrow \infty$.

We now show that on the event S_ε the MLE is within ε of θ_0 . Assume that the sample lies in S_ε and, to be specific, that $L(\theta_1) \leq L(\theta_2)$; the other case is symmetric. By (R3) the likelihood is a continuous function of θ and since $L(\theta_1) \leq L(\theta_2) < L(\theta_0)$ the intermediate value theorem guarantees a point $\tilde{\theta}_1 \in [\theta_1, \theta_0]$ with $L(\tilde{\theta}_1) = L(\theta_2)$. The function L is differentiable and has the same value at the two ends of the interval $[\tilde{\theta}_1, \theta_2]$, and thus by Rolle's theorem there is a point $\theta^* \in (\tilde{\theta}_1, \theta_2)$ with $L'(\theta^*) = 0$. Since $L > 0$ on this interval, the point θ^* is also a root of the likelihood equation $\ell'(\theta) = 0$. By assumption, this root is unique and equals the MLE, and thus $\hat{\theta}_n = \theta^*$. By construction θ^* is in (θ_1, θ_2) and we have shown that

$$S_\varepsilon \subseteq \{|\hat{\theta}_n - \theta_0| < \varepsilon\}.$$

Taking probabilities we find $\mathbb{P}_{\theta_0}(|\hat{\theta}_n - \theta_0| < \varepsilon) \geq \mathbb{P}_{\theta_0}(S_\varepsilon) \rightarrow 1$, i.e. $\mathbb{P}_{\theta_0}(|\hat{\theta}_n - \theta_0| \geq \varepsilon) \rightarrow 0$ as $n \rightarrow \infty$. Since $\varepsilon > 0$ was arbitrary and θ_0 was an arbitrary point of Θ , this is the statement $\hat{\theta}_n \xrightarrow{P} \theta$ for every true value, and the proof is complete. $\square$

The structure of this proof will be re-used in lecture 17. The probabilistic input comes from theorem 16.2 which gives an event of probability tending to one; on this event the argument is then a standard application of calculus, and the uniqueness of the MLE allows us to pass from “some root of the likelihood equation” to “the MLE”. In the standard examples, uniqueness is typically proved by showing that the log-likelihood is strictly concave, as for the exponential distribution in example 5.2; exercise 16.1 completes the proof of this model.

Remark. The assumption that the likelihood equation has a unique root for every sample is slightly stronger than required. For example, for the Bernoulli distribution in example 5.3, the likelihood equation has no root in the case where all observations are equal, but this only happens with probability $\theta_0^n + (1 - \theta_0)^n$ which tends to zero. The proof is unchanged if we assume that the condition holds on an event with probability tending to one, since we can intersect S_ε with such an event.

16.4 When the Assumptions Fail

Each hypothesis of theorem 16.3 can fail, and the failures are the cases where more care is needed in practice.

If the parameter is not identifiable, *i.e.* if (R1) fails, the question is ill posed: in example 4.11 the same distribution of the data could be generated by two different parameter values and thus no function of the data could converge to one of these parameter values rather than the other. In this case we would need to reparametrise the model, as shown in the example.

If the support of the model depends on the parameter, (R2), the likelihood is not a smooth function of θ and the MLE is typically not a root of the likelihood equation. For example, for the uniform distribution on $(0, \theta)$, the MLE is the sample maximum, as shown in example 5.5, and the theorem does not apply in this case. Nevertheless, the estimator is consistent: the mean squared error of the estimator goes to zero as the sample size increases, as shown in exercise 2.2, and thus theorem 2.10 applies; exercise 16.3 gives a direct proof of the same fact. For the same model, in lecture 13, we have seen that the Cramer–Rao bound can be violated if (R2) fails and in lecture 17 we will see that the limiting distribution of the MLE is not normal in this case.

If the likelihood equation has several roots, the regularity conditions may all hold and the theorem still does not apply, because we cannot identify the root found by Rolle’s theorem with the MLE; the t location model of section 5.3 is of this type. What survives is this: under (R1) to (R3) there is always a sequence of roots of the likelihood equation which is consistent, since the root θ^* found above lies within ε of θ_0 . What is lost is the guarantee that the root we find, or the global maximiser, is this consistent one; we do not prove the general statement in this module. The practical consequence of this is that numerical maximisation of the likelihood should be started with a simple consistent estimator, such as the sample median for a location parameter, so that the iteration converges to the nearby consistent root. This is a lesson we will take up in lecture 18.

The following questions, about the rate of convergence of the MLE to θ_0 and the distribution of the error for large n , are answered in the asymptotic normality theorem of lecture 17. The proof of this theorem uses the consistency shown here, and condition (R4).

Summary.

- The regularity conditions (R1) to (R4) from definition 16.1 state that the parameters must be identifiable, the support of the model must not depend on θ , the log-likelihood must be three times differentiable with finite and positive Fisher information, and a bound must be placed on the third derivative of the log-likelihood.
- For fixed parameter values, the likelihood of the true value eventually exceeds the likelihood of any wrong value, with probability approaching one. The proof uses the law of large numbers and Jensen's inequality.
- If the likelihood equation has only one root, and this root is the MLE, then the MLE is consistent. The proof finds a root of the likelihood equation near θ_0 , using Rolle's theorem and the fact that this root coincides with the MLE.
- The theorem does not apply if the support of the model depends on the parameter (since then the MLE is typically not a root of the likelihood equation, as for the uniform distribution, even though it may still be consistent), or if the likelihood equation has multiple roots (since then only the existence of a consistent root can be proved).

Exercise 16.1. Let $X_1, \dots, X_n$ be i.i.d. with exponential rate $\theta > 0$, i.e. $f(x; \theta) = \theta e^{-\theta x}$ for all $x \geq 0$.

1. Show that the regularity conditions (R1) to (R4) hold for this model. Hint: For (R3) you can use the value $\mathcal{I}(\theta) = 1/\theta^2$ from example 11.8. For (R4) you should take derivatives of $\log f(x; \theta)$ with respect to θ and then find a bound $M(x)$.
2. Show that the likelihood equation has only one root for each sample and that this root is the MLE. Use theorem 16.3 to conclude that $\hat{\theta}_n = 1/\bar{X}$ is consistent for θ .
3. Find an alternative proof of the consistency of $1/\bar{X}$, using the law of large numbers (theorem A.14), the continuous mapping theorem (theorem A.17), and any other arguments you need.

Exercise 16.2. Let X_1 be the only observation of a Bernoulli random variable with success probability $\theta_0 \in (0, 1)$, and let $\theta \in (0, 1)$ be another parameter value.

1. Compute the expectation $\mathbb{E}_{\theta_0}(\log(f(X_1; \theta)/f(X_1; \theta_0)))$ from the proof of theorem 16.2, as an explicit function of θ and θ_0 .
2. Using the inequality $\log y \leq y - 1$ for $y > 0$, with equality only for $y = 1$, show that this expectation is negative for all $\theta \neq \theta_0$.

Exercise 16.3. Let $X_1, \dots, X_n$ be i.i.d. uniformly distributed on the interval $(0, \theta)$ and let $\hat{\theta}_n = \max_i X_i$ be the MLE from example 5.5.

1. Show that for $0 < \varepsilon < \theta$ we have $\mathbb{P}_\theta(|\hat{\theta}_n - \theta| > \varepsilon) = (1 - \varepsilon/\theta)^n$ and thus $\hat{\theta}_n$ is consistent for θ .

2. Which of the assumptions of theorem 16.3 are violated for this model? Why cannot the theorem be applied in this case, despite the fact that the conclusion is still true?

Lecture 17

Asymptotic Normality and Efficiency of the MLE

In lecture 16 we have seen that under regularity conditions the maximum likelihood estimator converges to the true parameter value. In this lecture we will discuss the error $\hat{\theta}_n - \theta_0$ in more detail: It transpires that for large sample size the error is approximately normally distributed with mean zero and variance given by the Cramer–Rao bound from lecture 13. Thus, the MLE is unbiased in the limit, and more importantly, is as accurate as any unbiased estimator can be. The proof of this result is a challenging one, but it is based on tools we have already learned in this module. The argument is based on Taylor expansion of the score function, and the resulting theorem allows us to compute standard errors and to derive approximate confidence statements. We will see the resulting numbers in lecture 18 and will make the statements precise in lecture 25.

17.1 The Theorem

Using the notation from lecture 16, we write $\mathcal{I}(\theta)$ for the Fisher information of a single observation and $\mathcal{I}_n(\theta) = n\mathcal{I}(\theta)$ for the Fisher information of the sample, as given in definition 11.2 and proposition 11.4. The result is formulated in terms of the scaled error $\sqrt{n}(\hat{\theta}_n - \theta_0)$, as this is the only error term which has a non-degenerate limit: by consistency the error goes to zero and the factor $\sqrt{n}$ is required to balance this effect.

Theorem 17.1. *Suppose that the regularity conditions (R1) to (R4) from definition 16.1 are satisfied, that $\log(f(X_1; \theta)/f(X_1; \theta'))$ has finite mean under $\mathbb{P}_{\theta'}$ for all $\theta, \theta' \in \Theta$, and that for every sample size $n \geq 1$ and every sample where all coordinates are in the support of the model, the likelihood equation $\ell'(\theta) = 0$ has a unique solution $\hat{\theta}_n$ which coincides with the maximum likelihood estimator. Then*

$$\sqrt{n}(\hat{\theta}_n - \theta_0) \xrightarrow{d} N\left(0, \frac{1}{\mathcal{I}(\theta_0)}\right) \quad (n \rightarrow \infty).$$

Before we give the proof, we unpack the statement of the theorem. The theorem informally states that, as n gets large,

$$\hat{\theta}_n \approx N\left(\theta_0, \frac{1}{n\mathcal{I}(\theta_0)}\right) = N\left(\theta_0, \frac{1}{\mathcal{I}_n(\theta_0)}\right).$$

This statement contains three different aspects: (1) The centre of the approximating distribution equals the true value, *i.e.* the MLE is asymptotically unbiased (this does not mean that the MLE is unbiased for finite n , only that it is asymptotically unbiased). (2) The spread of the distribution decreases as $1/\sqrt{n}$, as for the sample mean. (3) The variance of the approximating distribution equals $1/\mathcal{I}_n(\theta_0)$. This is the Cramer–Rao lower bound from lecture 13: no unbiased estimator can have smaller variance for any n and the MLE achieves this bound in the limit. The last property has a name.

Definition 17.2. An estimator $\tilde{\theta}_n$ for θ , which does not have to be the maximum likelihood estimator, is *asymptotically efficient*, if $\sqrt{n}(\tilde{\theta}_n - \theta) \xrightarrow{d} N(0, 1/\mathcal{I}(\theta))$ for all $\theta \in \Theta$.

In this language theorem 17.1 states that under the regularity conditions the MLE is asymptotically efficient. In lecture 13 we have seen that an efficient estimator, one which achieves the Cramer–Rao bound exactly, only exists for special models and parametrisations, namely the natural statistic of an exponential family. The asymptotic statement is much wider in applicability, as it holds for every regular model and every parametrisation, since by theorem 5.6 the MLE transforms together with the parameter.

17.2 Proof of the Theorem

We will use a Taylor expansion of the score function ℓ' around the true value. Since the MLE is a root of the score, we can use Taylor expansion of $\ell'(\hat{\theta}_n)$ around the point θ_0 to write the error term $\hat{\theta}_n - \theta_0$ in terms of $\ell'(\theta_0)$ and $\ell''(\theta_0)$, which are both sums of i.i.d. terms and thus can be analysed using the limit theorems for sample means.

Proof of theorem 17.1. Let the true value be θ_0 . Then, by (R3), the function ℓ' is two times continuously differentiable. Thus, by Taylor’s theorem with the Lagrange form of the remainder, for ℓ' around the point θ_0 we get

$$\ell'(\hat{\theta}_n) = \ell'(\theta_0) + \ell''(\theta_0)(\hat{\theta}_n - \theta_0) + \frac{1}{2} \ell'''(\theta_n^*)(\hat{\theta}_n - \theta_0)^2,$$

where θ_n^* is a point between θ_0 and $\hat{\theta}_n$. Since $\hat{\theta}_n$ solves the likelihood equation, the left-hand side is zero. Solving for $\hat{\theta}_n - \theta_0$ and multiplying by $\sqrt{n}$ we can write the scaled error as a ratio:

$$\sqrt{n}(\hat{\theta}_n - \theta_0) = \frac{A_n}{B_n - C_n}, \quad (17.1)$$

where

$$A_n = \frac{1}{\sqrt{n}} \ell'(\theta_0), \quad B_n = -\frac{1}{n} \ell''(\theta_0), \quad C_n = \frac{1}{2n} \ell'''(\theta_n^*)(\hat{\theta}_n - \theta_0).$$

We now take the limit of each of the three terms separately.

The numerator A_n is a scaled sum of i.i.d. terms: by definition of the score we have $\ell'(\theta_0) = \sum_i \ell'_i(\theta_0)$ and using the identities (16.1) we find that each term

has mean zero and finite, positive variance $\mathcal{I}(\theta_0)$. Thus, by the central limit theorem, theorem A.15, for the sample mean of the $\ell'_i(\theta_0)$ we get

$$A_n = \sqrt{n} \left(\frac{1}{n} \sum_{i=1}^n \ell'_i(\theta_0) - 0 \right) \xrightarrow{d} N(0, \mathcal{I}(\theta_0)).$$

The first term in the denominator is also a sample mean: $B_n = -\frac{1}{n} \sum_i \ell''_i(\theta_0)$ and by the law of large numbers, theorem A.14, and the second identity in (16.1) we find

$$B_n \xrightarrow{P} -\mathbb{E}_{\theta_0}(\ell''_1(\theta_0)) = \mathcal{I}(\theta_0).$$

The remainder term C_n is where the consistency from lecture 16 and condition (R4) come into play. Let δ and M be as in (R4). On the event $\{|\hat{\theta}_n - \theta_0| < \delta\}$, the intermediate point θ_n^* also lies within δ of θ_0 . Thus, writing $\ell'''(\theta_n^*) = \sum_i \ell'''_i(\theta_n^*)$ as a sum of the observations and using (R4) for each term, we find

$$\left| \frac{1}{n} \ell'''(\theta_n^*) \right| \leq \frac{1}{n} \sum_{i=1}^n |\ell'''_i(\theta_n^*)| \leq \frac{1}{n} \sum_{i=1}^n M(X_i)$$

on this event. By theorem 16.3, the probability of the event goes to one, and by the law of large numbers the right-hand side converges in probability to the finite constant $\mathbb{E}_{\theta_0}(M(X_1))$. Thus, for all sufficiently large K we have $\mathbb{P}_{\theta_0}(|\ell'''(\theta_n^*)|/n > K) \rightarrow 0$. The other factor of C_n is $\hat{\theta}_n - \theta_0$, which converges in probability to zero by theorem 16.3. For any $\eta > 0$, the event $|C_n| > \eta$ is contained in the union of the events $|\ell'''(\theta_n^*)|/n > K$ and $|\hat{\theta}_n - \theta_0| > 2\eta/K$, both of which have probability that goes to zero, and thus $C_n \xrightarrow{P} 0$.

It remains to combine these three limits, and the appropriate tool for this is Slutsky's theorem, theorem A.16. Since $B_n \xrightarrow{P} \mathcal{I}(\theta_0)$ and $C_n \xrightarrow{P} 0$, the denominator satisfies $B_n - C_n \xrightarrow{P} \mathcal{I}(\theta_0)$, where the value is constant and non-zero. By the quotient part of Slutsky's theorem, for the representation (17.1), we have

$$\sqrt{n}(\hat{\theta}_n - \theta_0) = \frac{A_n}{B_n - C_n} \xrightarrow{d} \frac{N(0, \mathcal{I}(\theta_0))}{\mathcal{I}(\theta_0)} = N\left(0, \frac{1}{\mathcal{I}(\theta_0)}\right),$$

where in the last step we used the fact that dividing a $N(0, \sigma^2)$ random variable by the constant $c > 0$ gives a $N(0, \sigma^2/c^2)$ random variable, in this case with $\sigma^2 = \mathcal{I}(\theta_0)$ and $c = \mathcal{I}(\theta_0)$. This completes the proof. $\square$

The structure of this argument is typical of most large-sample results in statistics: Taylor expansion is used to transform a question about an implicitly defined estimator into a question about sums of i.i.d. terms, the central limit theorem is used to analyse the leading term, the law of large numbers is used to analyse the coefficient of the leading term, the remainder is shown to be negligible using consistency of the estimator, and Slutsky's theorem is used to combine the individual results. The only tricky step is the analysis of the remainder term, and condition (R4) is the reason that this term is straightforward to analyse.

17.3 Standard Errors

The variance $1/\mathcal{I}_n(\theta_0)$ in theorem 17.1 depends on the unknown θ_0 , and so cannot be reported as it stands. In practice we replace θ_0 by its estimate, and the next result says that this does not disturb the limit.

Corollary 17.3. *With the same assumptions as in theorem 17.1, we have*

$$\sqrt{n\mathcal{I}(\hat{\theta}_n)}(\hat{\theta}_n - \theta_0) \xrightarrow{d} N(0, 1) \quad (n \rightarrow \infty).$$

Proof. By (R3) the function $\mathcal{I}$ is continuous and positive on Θ and thus the function $\theta \mapsto \sqrt{\mathcal{I}(\theta)/\mathcal{I}(\theta_0)}$ is continuous at θ_0 . By $\hat{\theta}_n \xrightarrow{p} \theta_0$ from theorem 16.3, the continuous mapping theorem from theorem A.17 implies that $\sqrt{\mathcal{I}(\hat{\theta}_n)/\mathcal{I}(\theta_0)} \xrightarrow{p} 1$. By multiplying the result of theorem 17.1 by this factor and using the product part of Slutsky's theorem, theorem A.16, we find

$$\sqrt{n\mathcal{I}(\hat{\theta}_n)}(\hat{\theta}_n - \theta_0) = \sqrt{\frac{\mathcal{I}(\hat{\theta}_n)}{\mathcal{I}(\theta_0)}} \cdot \sqrt{n\mathcal{I}(\theta_0)}(\hat{\theta}_n - \theta_0) \xrightarrow{d} 1 \cdot N(0, 1),$$

and the claim is proved. $\square$

The quantity

$$\text{se}(\hat{\theta}_n) = \frac{1}{\sqrt{\mathcal{I}_n(\hat{\theta}_n)}} = \frac{1}{\sqrt{n\mathcal{I}(\hat{\theta}_n)}}$$

is called the *standard error* of the MLE. As a statistic, this quantity from the corollary is an approximation to the standard deviation of $\hat{\theta}_n$ for large n . The corollary also allows us to justify approximate confidence statements: Since a standard normal variable takes values in $(-1.96, 1.96)$ with probability 0.95, the random interval $\hat{\theta}_n \pm 1.96 \text{ se}(\hat{\theta}_n)$ covers θ_0 with probability which tends to 0.95. Intervals of this form were used in an informal way in lecture 12 and we will make these intervals precise in lecture 25. In lecture 18 we will use simulation to check how close the real coverage of these intervals is to 0.95 for moderate n . If we do not have a closed formula for $\mathcal{I}$, the denominator $\mathcal{I}_n(\hat{\theta}_n)$ can be replaced by the observed curvature $-\ell''(\hat{\theta}_n)$ of the log-likelihood at the maximum. This approach is taken by a numerical optimiser and we will use the resulting values in lecture 18.

17.4 Examples

For the MLEs which are sample means, the Bernoulli and Poisson MLEs from lecture 5, theorem 17.1 does not give any new results: the central limit theorem allows us to find the asymptotic normality of these estimators directly and from exercises 17.1 and 17.2 we know that the two variances are the same. For the normal mean with known variance, the MLE $\bar{X}$ is exactly normally distributed for every n , with variance $\sigma^2/n = 1/\mathcal{I}_n(\bar{X})$. Thus, for this case the theorem gives an exact result. The first new case is an MLE which is not a sample mean.

Example 17.4. Let $X_1, \dots, X_n$ be i.i.d. with exponential rate θ . Then the MLE for the parameter value is $\hat{\theta}_n = 1/\bar{X}$ from example 5.2 and by exercise 16.1 the regularity conditions and the uniqueness of the root hold. From example 11.8 we know that the Fisher information is given by $\mathcal{I}(\theta) = 1/\theta^2$. Thus, by theorem 17.1,

$$\sqrt{n} \left(\frac{1}{\bar{X}} - \theta \right) \xrightarrow{d} N(0, \theta^2),$$

with standard error $\text{se}(\hat{\theta}_n) = \hat{\theta}_n/\sqrt{n}$. We can compare the result to exact results as follows: In lecture 5 we found $\mathbb{E}_\theta(\hat{\theta}_n) = n\theta/(n-1)$, which is a bias of $\theta/(n-1)$.

In exercise 5.3 we found the variance $\text{Var}_\theta(\hat{\theta}_n) = n^2\theta^2/((n-1)^2(n-2))$. If we multiply this variance by n , it converges to θ^2 . If we multiply the bias by $\sqrt{n}$, it converges to zero, as predicted by the theorem: for large n , the bias is of a smaller order than the spread of the distribution, and thus disappears in the limit.

The theorem can also be applied to functions of the parameter. The mean of the exponential distribution is $\mu = g(\theta) = 1/\theta$ and by theorem 5.6 the MLE for this parameter is $\hat{\mu}_n = 1/\hat{\theta}_n = \bar{X}$. Using the delta method, theorem A.18 with $g'(\theta) = -1/\theta^2$, we get

$$\sqrt{n}(\bar{X} - \mu) \xrightarrow{d} N(0, g'(\theta)^2 \theta^2) = N\left(0, \frac{1}{\theta^2}\right) = N(0, \mu^2).$$

This is the central limit theorem for $\bar{X}$, since the exponential distribution has variance $1/\theta^2 = \mu^2$. The standard error of $\bar{X}$, $\bar{X}/\sqrt{n}$, is the standard error we would obtain by computing the standard error from first principles.

The example illustrates a general point: for any smooth function $g(\theta)$ of the parameter, theorem 5.6 gives the MLE $g(\hat{\theta}_n)$, and the delta method gives the asymptotic distribution and standard error $|g'(\hat{\theta}_n)| \text{se}(\hat{\theta}_n)$, without any further maximisation.

17.5 Quality of the Approximation

Theorem 17.1 is a limit statement and does not give information about the required sample size for the normal approximation to be applicable. The required value of n for the normal approximation to be usable depends on the model and there are three situations where care is needed. First, for small samples from a model with a skewed distribution: for the exponential rate, $\hat{\theta}_n = 1/\bar{X}$ has a right-skewed distribution and a bias of $\theta/(n-1)$ which is not captured in the limit but which is visible for n around 10. Secondly, when a parameter value is close to the boundary of the parameter space: for the Bernoulli distribution with θ_0 close to zero or one, the standard error $\sqrt{\hat{\theta}_n(1-\hat{\theta}_n)/n}$ is zero whenever all observations have the same value, and thus the interval $\hat{\theta}_n \pm 1.96 \text{se}(\hat{\theta}_n)$ has zero width in this case. Finally, if the regularity conditions are violated: for the uniform distribution on $(0, \theta)$, exercise 17.3 shows that the error of the MLE is of order $1/n$ instead of $1/\sqrt{n}$, and that the limiting distribution is exponential instead of normal, as expected from the super-efficiency of the parameter estimates as shown in lecture 13. The first case is investigated by simulation in lecture 12, by checking the coverage of the interval $\hat{\theta}_n \pm 1.96 \text{se}(\hat{\theta}_n)$ for moderate n .

Finally, the theorem can be extended to a vector parameter $\theta = (\theta_1, \dots, \theta_k)$: Under appropriate regularity conditions, the MLE is asymptotically normal with mean θ_0 and covariance matrix $\mathcal{I}_n(\theta_0)^{-1}$, the inverse of the Fisher information matrix from lecture 14. Without proof, the argument is similar to the one above, but uses Taylor expansion for several variables. For the normal distribution with unknown parameters, where the information matrix is diagonal, this gives the asymptotic normality of $\bar{X}$ and of the variance estimator separately.

Summary.

- Under regularity conditions and assuming the uniqueness of the likelihood root, we have $\sqrt{n}(\hat{\theta}_n - \theta_0) \xrightarrow{d} N(0, 1/\mathcal{I}(\theta_0))$: the MLE is asymptotically unbiased, the error is of order $1/\sqrt{n}$, and the asymptotic variance equals the Cramer–Rao bound, *i.e.* the MLE is asymptotically efficient.
- The proof uses Taylor expansion of the score around θ_0 , central limit theorem for the score, law of large numbers for the second derivative, consistency and (R4) for the remainder, and Slutsky’s theorem to combine these results.
- Replacing θ_0 by $\hat{\theta}_n$ in the expression for the variance does not change the limit and thus we get the standard error $\text{se}(\hat{\theta}_n) = 1/\sqrt{\mathcal{I}_n(\hat{\theta}_n)}$ and the approximate 95% confidence interval $\hat{\theta}_n \pm 1.96 \text{ se}(\hat{\theta}_n)$ for the parameter value.
- Using the delta method, we can apply the result to any smooth function of the parameter. The standard error for such an application is $|g'(\hat{\theta}_n)| \text{ se}(\hat{\theta}_n)$.
- The approximation can be poor for small n in models with skewed distributions and at the boundary of the parameter space, and it breaks down completely when the regularity conditions are violated, *e.g.* for the uniform distribution.

Exercise 17.1. Let $X_1, \dots, X_n$ be i.i.d. Bernoulli with success probability $\theta \in (0, 1)$, where the MLE $\hat{\theta}_n = \bar{X}$ and Fisher information $\mathcal{I}(\theta) = 1/(\theta(1 - \theta))$ are given in example 11.5.

1. State the result of theorem 17.1 for this model and verify that it coincides with the central limit theorem for $\bar{X}$.
2. The *log-odds* are given by $\psi = g(\theta) = \log(\theta/(1 - \theta))$. Find the MLE for ψ and use the delta method to show that $\sqrt{n}(\hat{\psi}_n - \psi) \xrightarrow{d} N(0, 1/(\theta(1 - \theta)))$.
3. Determine the standard error of $\hat{\psi}_n$ and an approximate 95% interval for ψ , in terms of $\hat{\theta}_n$ and n .

Exercise 17.2. Let $X_1, \dots, X_n$ be i.i.d. Poisson with parameter $\theta > 0$, where MLE $\hat{\theta}_n = \bar{X}$ and Fisher information $\mathcal{I}(\theta) = 1/\theta$ are given in example 11.6.

1. State the result of theorem 17.1 for this model and verify that the result coincides with the central limit theorem.
2. In exercise 5.7 we have found the MLE $\hat{p}_0 = e^{-\bar{X}}$ for the probability $p_0 = \mathbb{P}_\theta(X_1 = 0) = e^{-\theta}$ and we have seen that this MLE is biased. Using the delta method, determine the asymptotic distribution of $\sqrt{n}(\hat{p}_0 - p_0)$. Why does the bias, found in the previous step, not contradict the result from this question?

Exercise 17.3. Let $X_1, \dots, X_n$ be i.i.d. uniformly distributed on the interval $(0, \theta)$. The MLE for this model is $\hat{\theta}_n = \max_i X_i$.

1. For every $x \geq 0$, show that $\mathbb{P}_\theta(n(\theta - \hat{\theta}_n) > x) \rightarrow e^{-x/\theta}$ as $n \rightarrow \infty$, *i.e.* that $n(\theta - \hat{\theta}_n)$ converges in distribution to the exponential distribution with mean θ , *i.e.* with rate $1/\theta$.
2. Show that $\sqrt{n}(\hat{\theta}_n - \theta) \xrightarrow{P} 0$. Why is the MLE not asymptotically normal in the sense of theorem 17.1? Why does this make sense in the context of the super-efficiency of the uniform model from lecture 13?

Lecture 18

Maximum Likelihood in R

This is the R session for week 6. In lectures 16 and 17 we have been concentrating on some mathematical proofs, and in this week we will be considering the resulting theorems in a numerical context: we will find the maximum likelihood estimate for a model where the likelihood equation cannot be solved explicitly, we will learn how to read off the standard errors from the output of the optimiser, and we will use simulation to determine how often the interval $\hat{\theta} \pm 1.96 \text{ se}(\hat{\theta})$ from corollary 17.3 contains the truth. The written reference for this session is again appendix B.

18.1 The Log-Likelihood as an R Function

The model of the day is $X_1, \dots, X_n$ i.i.d. from the gamma distribution with shape α and rate β , both unknown. Example 4.9 gave the method of moments estimates $\tilde{\alpha} = \bar{X}^2/V$ and $\tilde{\beta} = \bar{X}/V$, with V the plug-in variance, and section 5.3 noted that the likelihood equations have no closed-form solution. We simulate one sample of size $n = 40$ with $\alpha = 2.5$ and $\beta = 0.5$. The command `dgamma` with `log = TRUE` can be used to evaluate the logarithm of the density at every observation, and the sum of these values is $\ell(\alpha, \beta)$. Since the R optimisers *minimise*, the function returns the negative log-likelihood, and since they hand over all parameters as one vector, it takes the vector `par`, shape first and rate second.

```
alpha <- 2.5
beta  <- 0.5
n     <- 40
set.seed(2026)
x <- rgamma(n, shape = alpha, rate = beta)
nll <- function(par, data) {
  -sum(dgamma(data, shape = par[1], rate = par[2], log = TRUE))
}
nll(c(2.5, 0.5), x)
nll(c(1, 1), x)
```

```
## [1] 102.0691
## [1] 217.3799
```

The true parameter values lead to a much smaller negative log-likelihood than the pair $(1, 1)$. Here we assume that we know the rate $\beta = 0.5$ and that we

want to estimate the shape parameter α . The command `optimise` can be used to find the minimum of a function of one variable in a given interval, without requiring a starting value. Before you try the command, make a sketch of the negative log-likelihood as a function of α on the interval from 0.1 to 20. How many dips will there be? Will the minimum lie above or below 2.5?

```
optimise(function(a) nll(c(a, 0.5), x), interval = c(0.1, 20))
```

```
## $minimum
## [1] 2.624051
##
## $objective
## [1] 101.924
```

The estimate $\hat{\alpha} = 2.62$ is a little above the true value, and the minimum 101.9 lies below the value 102.1 where the truth is located. Since both parameters are now unknown, we have to use `optim`, which minimises a function of a vector, and which requires a starting value. From section 16.4 we know that a good way to find a starting value is to use a simple, consistent estimator. One possible choice for this is the method of moments estimate. Before you try the code, predict whether the maximum likelihood estimates will agree with the method of moments estimates to one decimal place.

```
V      <- mean((x - mean(x))^2)
start <- c(mean(x)^2 / V, mean(x) / V)
start
fit <- optim(start, nll, data = x)
fit$par
fit$convergence
fit$counts
```

```
## [1] 2.2942540 0.4221649
## [1] 2.258769 0.415672
## [1] 0
## function gradient
##      43      NA
```

The extra argument `data = x` is passed on to our function. The result `fit` is a list: `fit$par` holds the estimates, $\hat{\alpha} = 2.26$ and $\hat{\beta} = 0.42$, which agree with the start to one decimal place; `fit$convergence` is 0 if the search finished properly, and any other value means the estimates cannot be trusted; `fit$counts` records that 43 evaluations were needed. By default `optim` uses the Nelder–Mead method, which needs no derivatives; appendix B shows the alternative "BFGS".

18.2 Standard Errors from the Observed Information

From theorem 17.1 and the vector version of this result from lecture 17 we know that the MLE is approximately normally distributed with covariance matrix $\mathcal{I}_n(\theta_0)^{-1}$. As in section 17.3, we can replace $\mathcal{I}_n(\hat{\theta})$ by the observed information, *i.e.* the matrix of second derivatives of the negative log-likelihood at the maximum. The `optim` function can compute this matrix for us if we ask for it

explicitly, using `hessian = TRUE`. The matrix can be inverted using the `solve` command, the `diag` command gives the diagonal of this matrix, and the square roots of the diagonal elements are the standard errors. Before you look at the output, predict which of the two standard errors will be larger and by how much.

```
fit <- optim(start, nll, data = x, hessian = TRUE)
fit$hessian
covmat <- solve(fit$hessian)
se <- sqrt(diag(covmat))
result <- cbind(estimate = fit$par, se = se,
                lower = fit$par - 1.96 * se, upper = fit$par + 1.96 * se
                )
rownames(result) <- c("shape", "rate")
round(result, 3)
cov2cor(covmat)[1, 2]
```

```
##           [,1]      [,2]
## [1,]  22.18712 -96.2299
## [2,] -96.22990 522.9200
##      estimate      se lower upper
## shape    2.259 0.473  1.333  3.185
## rate      0.416 0.097  0.225  0.606
## [1] 0.8933916
```

The table prints the results of a maximum likelihood estimate in an easy to read format. Each row corresponds to one parameter, and shows the parameter estimate, the standard error and the interval $\hat{\theta} \pm 1.96 \text{ se}(\hat{\theta})$ from corollary 17.3. Both intervals capture the true parameter values. The standard error for the shape parameter is approximately five times larger than the standard error for the rate parameter, but, since the estimates differ by a similar factor, the two standard errors are comparable relative to the estimates. The final row of the table, computed using the `cov2cor` function, shows that the two estimates have correlation 0.89. It is easy to show that, if we increase the shape and the rate simultaneously, while keeping the mean α/β fixed, the log-likelihood changes only slightly. Thus the data constrains the ratio α/β well, but poorly constrains each parameter individually. This is the problem of a nuisance parameter as discussed in lecture 14.

18.3 Starting Values

The gamma log-likelihood function has only one peak, so the method is forgiving. Before you try the code, make a prediction: what happens if we start with very different values, *e.g.* with shape 30 and rate 30? Will `optim` still find the maximum? Will it take more than 43 evaluations to find the maximum, when we used the “good” start?

```
fit_far <- optim(c(30, 30), nll, data = x)
fit_far$par
fit_far$counts
```

```
## [1] 2.2571696 0.4154238
```

```
## function gradient
##      97      NA
```

The search reaches the same maximum, up to two decimal places, but it takes more than twice as long and R complains about several warnings and the presence of `NaNs` (not a number) in the output: the search tried to use negative parameter values where `dgamma` is not defined. We can prevent this by adding a line at the top of the function which returns `Inf` if `any(par <= 0)`. This will prevent the optimiser from trying these values. For a “good” start this is not required. The real problem of a poor start is when the likelihood has several local maxima, as in section 5.3. We illustrate this effect using a toy function which is not a model: the logarithm of a two-component normal mixture density, evaluated at θ , has a local maximum at $\theta = 2$ and a lower local maximum at $\theta = -2$ (see figure 18.1). Which of these two peaks will `optimise` report for the interval from -6 to 4 , which contains both maxima? Which `optim` result will you get if you start at -3 and at 3 , respectively? To find the maximum, `optim` needs to use a different method than the default method, and we can use the `BFGS` method. Also, we need to tell `optimise` that we want to find a maximum by using the argument `maximum = TRUE`.

```
toy <- function(theta) {
  log(0.7 * dnorm(theta, mean = 2) + 0.3 * dnorm(theta, mean = -2, sd =
    0.7))
}
curve(toy(x), from = -5, to = 5)
optimise(toy, interval = c(-6, 4), maximum = TRUE)$maximum
optim(-3, function(theta) -toy(theta), method = "BFGS")$par
optim(3, function(theta) -toy(theta), method = "BFGS")$par
```

```
## [1] -1.99892
## [1] -1.998923
## [1] 2
```

Both optimisers find a local maximum, but neither of them can determine whether this is the global maximum: `optim` starts at and ends at the first peak it encounters, and `optimise` assumes that there is only one peak. There is no command which allows us to fix this problem. The only way to avoid this problem is to use some theory and to follow certain habits: always start with a consistent estimator (as we did for the gamma distribution earlier in this lecture), always plot the log-likelihood function (if the number of parameters is one), and always start the `optim` procedure again to see whether the solution is different.

18.4 Coverage of the Wald Interval

While corollary 17.3 shows that the interval $\hat{\theta} \pm 1.96 \text{se}(\hat{\theta})$ covers the truth with probability approaching 0.95, in section 17.5 we have seen that sometimes, for moderate n , this is not the case. To verify the claim for the gamma shape and rate, we use simulation in this section, following the approach to coverage introduced in lecture 12. The complete script for our simulation is given in the following box, with the lines shuffled (except for the bodies of the function and of the loop). Before you read the rest of this section, try to do the following:




Figure 18.1: The toy function with the two peaks. The peak found by `optim` depends on the starting point. If `optimise` is used, it returns the lower peak, despite the interval containing both peaks.

(1) write down the lines in a way which forms a working script, and (2) predict the coverages for $n = 20$: will they be above or below 0.95? Which of the two, shape or rate, will be worse?

```
data.frame(n = ns, shape = coverage[, 1], rate = coverage[, 2])
fit_gamma <- function(n, alpha, beta) {
  x      <- rgamma(n, shape = alpha, rate = beta)
  V      <- mean((x - mean(x))^2)
  start  <- c(mean(x)^2 / V, mean(x) / V)
  fit    <- optim(start, nll, data = x, hessian = TRUE)
  c(fit$par, sqrt(diag(solve(fit$hessian))))
}
set.seed(2026)
for (i in seq_along(ns)) {
  res    <- replicate(5000, fit_gamma(ns[i], alpha, beta))
  lower  <- res[1:2, ] - 1.96 * res[3:4, ]
  upper  <- res[1:2, ] + 1.96 * res[3:4, ]
  coverage[i, 1] <- mean(lower[1, ] <= alpha & alpha <= upper[1, ])
  coverage[i, 2] <- mean(lower[2, ] <= beta & beta <= upper[2, ])
}
alpha <- 2.5
coverage <- matrix(0, nrow = length(ns), ncol = 2)
ns      <- c(20, 100)
beta    <- 0.5
```

Every name must exist before it is used: the function and the four assignments come first, with `coverage` after `ns`; then the seed, then the loop, and the table last. The one new detail is the return value, the two estimates followed by the two standard errors, which `replicate` stacks as the columns of a matrix; thus `res[1:2, ]` holds the estimates and `res[3:4, ]` the standard errors. The script in working order is `R/S18-coverage.R`; it fits the model 5000 times for each sample size, takes a few seconds, and prints the following table.

```
##      n  shape  rate
## 1   20 0.9624 0.9604
## 2  100 0.9508 0.9526
```

At $n = 100$ both coverages are within simulation error of the nominal 0.95, as expected from corollary 17.3. At $n = 20$ both coverages are *above* the nominal level, around 0.96: the interval is slightly too wide rather than too narrow. The total hides an imbalance again, and the script also counts for the shape at $n = 20$ how many misses are below the truth and how many are above.

```
## below above
## 0.0296 0.0116
```

Almost three quarters of the misses are below the true shape. Since the standard error increases with the estimate, an underestimated shape will result in a short interval which does not reach the truth, whereas an overestimated shape will result in a long interval which usually still covers the truth. The same mechanism was at work for the exponential rate in lecture 12.

18.5 A Task to Take Home

The following task is in the style of a practical report. The published scripts of this lecture can be used as a starting point. Use your student ID to seed the random number generator, and choose your own shape and rate. Simulate one gamma sample of size $n = 30$. Use `optim` to fit the model by maximum likelihood, starting with estimates from the method of moments. For each parameter, report the estimate, the standard error, and the interval $\hat{\theta} \pm 1.96 \text{ se}(\hat{\theta})$, in three separate sentences. For each interval, state whether the interval contains the value you chose. Estimate the coverage of both intervals when $n = 30$, using 2000 replicates. Write a short paragraph about whether you would trust the intervals for this sample size. Finally, fit the normal model with unknown mean and variance, using `optim`, to a normal sample of size 30. Check that the estimates match the closed forms from example 5.4, and that the standard error of the mean matches $\hat{\sigma}/\sqrt{n}$.

Before the next session, you should type the commands used in this session, and should check that you get the same results as the text.

Summary.

- The written documentation for `optim`, `optimise`, `solve` and the Hessian matrix can be found in appendix B.
- The negative log-likelihood can be represented as an R function for one parameter vector, using the density with `log = TRUE`. `optimise` can then be used to minimise this function over an interval for one parameter, and `optim` can be used to find the minimum for several parameters, starting from a given point. It is important to check `fit$convergence` to see whether the optimisation was successful.
- If `hessian = TRUE`, the optimiser returns the observed information matrix. The inverse of this matrix can be used to approximately determine the covariance matrix of the MLE, and the square roots of the diagonal elements can be used to determine the standard errors reported in section 17.3. The results of a fit can be printed in a table which shows the estimates, standard errors and intervals for each parameter.
- A good starting point for the optimisation procedure is to use a consistent estimator, *e.g.* based on the method of moments. Also, since an optimiser finds a local maximum, it is important to plot the log-likelihood and to re-run the optimisation procedure with a different starting point.
- For the gamma distribution, with shape and rate parameters, the coverage of the Wald interval is slightly above 95% for $n = 20$ and close to 95% for $n = 100$, with most of the misses being below the true value. This is caused by small estimates having small standard errors.

Lecture 19

The EM Algorithm

For the standard families of tables 1.1 and 1.2 the likelihood equation from lecture 5 could be solved in closed form, but for models with latent variables or missing data, as considered in this lecture, this is no longer the case. The expectation-maximisation (EM) algorithm, introduced in this lecture, alternates between filling in the missing information and re-estimating the parameter values. We prove that each round of the EM algorithm cannot decrease the likelihood. We illustrate the EM algorithm with the help of a worked example, considering a mixture of two normal distributions. This example also gives the non-identifiable model from lecture 4. In lecture 21 we will discuss how to implement the EM algorithm in R.

19.1 Incomplete Data

Assume that we are given data consisting of i.i.d. pairs (X_i, Z_i) , $i = 1, \dots, n$, with joint density or joint probability weights $f_c(x, z; \theta)$, but we have observed only $X = (X_1, \dots, X_n)$. The pair (X, Z) with $Z = (Z_1, \dots, Z_n)$ is called the *complete data*, the values X are called the *observed data* and the unobserved values Z_i are called *latent variables*. The *complete-data log-likelihood* is given by

$$\ell_c(\theta) = \sum_{i=1}^n \log f_c(X_i, Z_i; \theta),$$

and the *observed-data log-likelihood* is given by the marginal density of a single observation,

$$f(x; \theta) = \sum_z f_c(x, z; \theta), \quad \ell(\theta) = \sum_{i=1}^n \log f(X_i; \theta),$$

where the integral is used instead of a sum when Z_i is continuous. The MLE for θ is the maximiser of ℓ , but the algebra from lecture 5 is not applicable, since the sum inside the logarithm prevents any simplifications. The function ℓ_c , being a sum of standard log-likelihoods, would be easy to maximise, but it cannot even be evaluated as it depends on Z . The EM algorithm is based on this observation.

In the example we consider in the lecture, the population consists of two groups, with proportion π for group 1 and $1 - \pi$ for group 2. The measurement for group 1 is $N(\mu_1, \sigma^2)$ and for group 2 is $N(\mu_2, \sigma^2)$, where we assume that

we know σ^2 . We are given observations of the measurement X_i , but we do not know which group $Z_i \in \{1, 2\}$ the observation belongs to. If we write φ_σ for the density of $N(0, \sigma^2)$, then the observations are i.i.d. with density

$$f(x; \theta) = \pi \varphi_\sigma(x - \mu_1) + (1 - \pi) \varphi_\sigma(x - \mu_2), \quad \theta = (\pi, \mu_1, \mu_2) \in (0, 1) \times \mathbb{R} \times \mathbb{R}.$$

This is an example of a *two-component normal mixture*. If we knew the groups, the MLEs for this example would be easy to find: we would get the proportion of group 1 and the two group means by using examples 5.3 and 5.4, but without the groups we cannot find the MLEs in closed form.

19.2 The Algorithm

If we knew the parameter, we could compute the conditional distribution of the latent variables given the data; if we knew the latent variables, we could maximise ℓ_c . Neither of these quantities is known, so we start with a guess θ_0 and alternate between filling in the missing values, replacing ℓ_c by its conditional expectation given the observed data, and maximising over θ .

Definition 19.1. Let $\theta_0 \in \Theta$ be given. Then the *EM algorithm* for finding the parameter values $\theta_1, \theta_2, \dots$ consists of the following iterative steps for $k = 0, 1, 2, \dots$:

(E) Expectation step: compute

$$Q(\theta | \theta_k) = \mathbb{E}_{\theta_k}(\ell_c(\theta) | X),$$

the conditional expectation of the complete-data log-likelihood given the observed data, where the conditional distribution of Z is computed using the parameter value θ_k .

(M) Maximisation step: set

$$\theta_{k+1} = \arg \max_{\theta \in \Theta} Q(\theta | \theta_k).$$

The second argument θ_k of Q enters only through the conditional distribution of Z given X . By Bayes' rule, proposition A.20 in appendix A, this conditional distribution has density or probability weights

$$k(z | x; \theta) = \frac{f_c(x, z; \theta)}{f(x; \theta)}.$$

Since the pairs (X_i, Z_i) are independent, once the data are observed $Q(\theta | \theta_k)$ is a sum over the observations, given by

$$Q(\theta | \theta_k) = \sum_{i=1}^n \sum_z k(z | x_i; \theta_k) \log f_c(x_i, z; \theta).$$

In cases where ℓ_c is linear in the indicator functions of the latent variables, *e.g.* for the mixture distribution considered in this section, the E-step is to replace the indicator functions by the corresponding conditional probabilities given the data.

19.3 The Ascent Property

The algorithm only ever maximises Q , and it is not obvious that this does anything for the observed-data log-likelihood ℓ . The central result of this lecture is that it does. We assume throughout that the set of values z with $f_c(x, z; \theta) > 0$ does not depend on θ , the analogue for the latent variables of condition (R2) in definition 16.1. We also assume that $f(x; \theta) > 0$ for every data point x and every $\theta \in \Theta$, and that $\mathbb{E}_{\theta'}(\log f_c(x, Z; \theta) \mid X = x)$ is finite for all $\theta, \theta' \in \Theta$ and every data point x , so that $Q(\theta|\theta_k)$ and $H(\theta|\theta_k)$ in the decomposition (19.1) below are well defined. Rearranging the formula for $k(z|x; \theta)$ and taking logarithms, we have

$$\log f(x; \theta) = \log f_c(x, z; \theta) - \log k(z|x; \theta)$$

for every such z . The left-hand side does not involve z and thus taking conditional expectations given $X = x$ under θ_k , with the random variable Z in place of z , does not change this. Summing over the sample gives the decomposition

$$\ell(\theta) = Q(\theta|\theta_k) - H(\theta|\theta_k), \quad \text{where } H(\theta|\theta_k) = \sum_{i=1}^n \mathbb{E}_{\theta_k}(\log k(Z_i|X_i; \theta) \mid X_i), \quad (19.1)$$

valid for all $\theta, \theta_k \in \Theta$.

Theorem 19.2. *Let $\theta_0, \theta_1, \theta_2, \dots$ be generated by the EM algorithm from definition 19.1. Then*

$$\ell(\theta_{k+1}) \geq \ell(\theta_k) \quad \text{for all } k = 0, 1, 2, \dots$$

Proof. We first show that $H(\theta|\theta_k) \leq H(\theta_k|\theta_k)$ for all θ . Since H is a sum over the observations, we only need to consider a single observation with observed value x . Let $Y = k(Z|x; \theta)/k(Z|x; \theta_k)$. This value is well defined and positive by our assumption on the supports. Since $-\log$ is convex, we can use Jensen's inequality, lemma A.7, for the conditional distribution of Z given $X = x$ under θ_k to get

$$\begin{aligned} & \mathbb{E}_{\theta_k}(\log k(Z|x; \theta) \mid X = x) - \mathbb{E}_{\theta_k}(\log k(Z|x; \theta_k) \mid X = x) \\ &= \mathbb{E}_{\theta_k}(\log Y \mid X = x) \leq \log \mathbb{E}_{\theta_k}(Y \mid X = x) \\ &= \log \sum_z \frac{k(z|x; \theta)}{k(z|x; \theta_k)} k(z|x; \theta_k) = \log \sum_z k(z|x; \theta) = \log 1 = 0, \end{aligned}$$

since $z \mapsto k(z|x; \theta)$ is a probability distribution over z (with sums replaced by integrals for a continuous latent variable). This is the same argument as in the proof of theorem 16.2.

Now we can use the decomposition (19.1) at $\theta = \theta_{k+1}$ and at $\theta = \theta_k$, both with the same second argument θ_k , and taking the difference we get

$$\ell(\theta_{k+1}) - \ell(\theta_k) = (Q(\theta_{k+1}|\theta_k) - Q(\theta_k|\theta_k)) - (H(\theta_{k+1}|\theta_k) - H(\theta_k|\theta_k)).$$

The first bracket is non-negative, since θ_{k+1} is the maximiser of $\theta \mapsto Q(\theta|\theta_k)$ in the M step, and the second bracket is non-positive by the first step of the proof. Thus the difference is non-negative as claimed. $\square$

The statement of the theorem does not imply that EM finds the MLE: the sequence $\ell(\theta_k)$ converges whenever ℓ is bounded above, but the iteration may stop at a local maximum or, in principle, at a saddle point. In general we can only say that the points where the EM algorithm converges are stationary points of ℓ .

Proposition 19.3. *Let Θ be open and ℓ as well as $\theta \mapsto Q(\theta|\theta')$ differentiable for all $\theta' \in \Theta$. Then*

$$\ell'(\theta') = \left. \frac{\partial}{\partial \theta} Q(\theta|\theta') \right|_{\theta=\theta'} \quad \text{for all } \theta' \in \Theta.$$

In particular, if $\theta_{k+1} = \theta_k$ for some k , the point θ_k is a root of the likelihood equation $\ell'(\theta) = 0$.

Proof. From (19.1) we know that the function $\theta \mapsto H(\theta|\theta') = Q(\theta|\theta') - \ell(\theta)$ is differentiable. Furthermore, from the first step of the proof of theorem 19.2 we know that this function has its maximum at an interior point θ' . Thus, the derivative at this point must be zero. Taking derivatives of the decomposition at $\theta = \theta'$ gives the identity. If $\theta_{k+1} = \theta_k$, then θ_k is an interior maximiser of $\theta \mapsto Q(\theta|\theta_k)$ and the right-hand side equals zero at $\theta' = \theta_k$. Thus we have $\ell'(\theta_k) = 0$. This completes the proof. $\square$

For a vector parameter the same argument applies, but now using gradients instead of derivatives. Conversely, if the function $\theta \mapsto Q(\theta|\theta')$ is strictly concave for all θ' , then all roots of the likelihood equation are fixed points: The derivative of $\theta \mapsto Q(\theta|\theta^*)$ is zero at a root θ^* and the only maxima of a strictly concave function are the points where the derivative is zero; exercises 19.2 and 19.4 show that this argument applies for two mixtures.

19.4 A Two-Component Normal Mixture

We can apply the algorithm for finding the parameters of a mixture from section 19.1. The E and M steps for this mixture model are as follows.

Example 19.4. Let $Z_1, \dots, Z_n$ be i.i.d. with $\mathbb{P}(Z_i = 1) = \pi$ and $\mathbb{P}(Z_i = 2) = 1 - \pi$ and given $Z_i = j$ let $X_i \sim N(\mu_j, \sigma^2)$, where $\sigma^2 > 0$ is known and $\theta = (\pi, \mu_1, \mu_2)$ is unknown. Then the complete-data density is $f_c(x, 1; \theta) = \pi \varphi_\sigma(x - \mu_1)$ and $f_c(x, 2; \theta) = (1 - \pi) \varphi_\sigma(x - \mu_2)$. Thus, the complete-data log-likelihood is

$$\ell_c(\theta) = \sum_{i=1}^n \left[\mathbf{1}_{\{Z_i=1\}} (\log \pi + \log \varphi_\sigma(x_i - \mu_1)) + \mathbf{1}_{\{Z_i=2\}} (\log(1 - \pi) + \log \varphi_\sigma(x_i - \mu_2)) \right],$$

which is linear in the indicators. For the E step, using Bayes' rule with the current value $\theta_k = (\pi_k, \mu_{1,k}, \mu_{2,k})$, we find

$$\gamma_i = \mathbb{P}_{\theta_k}(Z_i = 1 | X_i = x_i) = \frac{\pi_k \varphi_\sigma(x_i - \mu_{1,k})}{\pi_k \varphi_\sigma(x_i - \mu_{1,k}) + (1 - \pi_k) \varphi_\sigma(x_i - \mu_{2,k})},$$

This is called the *responsibility* of group 1 for the observation i , and $Q(\theta|\theta_k)$ is $\ell_c(\theta)$ with $\mathbf{1}_{\{Z_i=1\}}$ replaced by γ_i and $\mathbf{1}_{\{Z_i=2\}}$ replaced by $1 - \gamma_i$.

For the M step we maximise over $\theta = (\pi, \mu_1, \mu_2)$ with the γ_i fixed. Since $\log \varphi_\sigma(x - \mu) = -(x - \mu)^2/(2\sigma^2) + \text{const}$, the partial derivatives of Q are

$$\begin{aligned}\frac{\partial Q}{\partial \pi} &= \frac{\sum_i \gamma_i}{\pi} - \frac{\sum_i (1 - \gamma_i)}{1 - \pi}, \\ \frac{\partial Q}{\partial \mu_1} &= \frac{1}{\sigma^2} \sum_{i=1}^n \gamma_i (x_i - \mu_1), \quad \frac{\partial Q}{\partial \mu_2} = \frac{1}{\sigma^2} \sum_{i=1}^n (1 - \gamma_i) (x_i - \mu_2),\end{aligned}$$

and setting these to zero gives the updates

$$\pi_{k+1} = \frac{1}{n} \sum_{i=1}^n \gamma_i, \quad \mu_{1,k+1} = \frac{\sum_i \gamma_i x_i}{\sum_i \gamma_i}, \quad \mu_{2,k+1} = \frac{\sum_i (1 - \gamma_i) x_i}{\sum_i (1 - \gamma_i)}. \quad (19.2)$$

The second derivatives are negative and thus these are maxima. They are the complete-data MLEs of examples 5.3 and 5.4 with indicators replaced by responsibilities.

For a numerical example we can choose $\sigma = 1$, together with the observations $x = (0, 1, 5, 6, 7)$ and initial value $\theta_0 = (1/2, 2, 4)$. In this case we have $\pi_0 = 1/2$ and to cancel the weights we divide the numerator and denominator by $\varphi_1(x_i - 2)$ to get

$$\gamma_i = \frac{1}{1 + \exp(\frac{1}{2}(x_i - 2)^2 - \frac{1}{2}(x_i - 4)^2)} = \frac{1}{1 + e^{2x_i - 6}}.$$

So, with four decimal digits we have

$$\gamma = (0.9975, 0.9820, 0.0180, 0.0025, 0.0003).$$

This shows that the two small observations are assigned nearly completely to group 1, while the three large observations are assigned to group 2. Now, with $\sum_i \gamma_i = 2.0003$, $\sum_i \gamma_i x_i = 1.0891$ and $\sum_i (1 - \gamma_i) x_i = 17.9109$, we can use the update (19.2) to get

$$\theta_1 = \left(\frac{2.0003}{5}, \frac{1.0891}{2.0003}, \frac{17.9109}{2.9997} \right) = (0.4001, 0.5445, 5.9709).$$

The log-likelihood of the observed data increases from $\ell(\theta_0) = -17.52$ to $\ell(\theta_1) = -9.21$, as required by theorem 19.2 (see exercise 19.3). For another iteration we find

$$\theta_2 = (0.4000, 0.5001, 6.0000).$$

At this point the iteration no longer changes by more than four decimal digits: the responsibilities are so close to 0 or 1 that the mean of each group coincides with the mean of the observations for that group. The algorithm converged quickly in this example, because the groups were well separated. In cases where the groups have more overlap, more iterations will be required for the algorithm to converge.

In the example we assumed that the variance was known and common to both groups. If, on the other hand, we had assumed separate, unknown variances for the two groups, then the estimate for one component could be centred on an observation and the variance of this component could be shrunk to zero. In this case, the likelihood would become infinite at this point and the MLE would not exist.

19.5 Practical Matters

Three issues arise whenever EM is used, and the mixture shows all three. The first is the promised failure of identifiability: the parameter values (π, μ_1, μ_2) and $(1 - \pi, \mu_2, \mu_1)$ give the same density $f(x; \theta)$, and thus the model is not identifiable in the sense of definition 4.10, the difficulty foreshadowed after example 4.11. This is called *label switching*: the likelihood has two global maxima which are mirror images of each other, and EM started from $(1/2, 4, 2)$ instead of $(1/2, 2, 4)$ in example 19.4 converges to $(0.6, 6.0, 0.5)$ instead of $(0.4, 0.5, 6.0)$. Either answer is fine on its own, but averaging different runs is not, and the remedy of lecture 4 is to restrict the parameter space by a convention such as $\mu_1 < \mu_2$. The second issue is local maxima: theorem 19.2 guarantees ascent and nothing more, and a poor starting value can lead the iteration to a local maximum of ℓ without warning. The standard precaution is to run EM from several starting values and to keep the run with the largest observed-data log-likelihood, which we can always evaluate. The third issue is when to stop: since $\ell(\theta_k)$ is non-decreasing, we stop when its increase falls below a small tolerance; convergence is geometric. Finally, EM produces an estimate but no standard error; that is obtained afterwards from the observed information at the fixed point, as in section 17.3. Lecture 21 shows all three issues in R.

Summary.

- In an incomplete-data model, the complete data (X, Z) has an easy log-likelihood ℓ_c , but only X is observed and the observed-data log-likelihood ℓ has a sum inside the logarithm.
- EM iteratively computes the expectation $Q(\theta|\theta_k) = \mathbb{E}_{\theta_k}(\ell_c(\theta) | X)$ and then maximises the expectation Q over θ .
- We have shown that $\ell(\theta_{k+1}) \geq \ell(\theta_k)$ for all k . The proof uses $\ell = Q - H$ and Jensen's inequality to show that the maximum of $\theta \mapsto H(\theta|\theta_k)$ is attained at θ_k . A fixed point of the iteration is a root of the likelihood equation.
- For a two-component normal mixture, the E step computes the responsibilities $\gamma_i = \mathbb{P}_{\theta_k}(Z_i = 1 | X_i = x_i)$ and the M step computes the weight as the average responsibility and the means as the responsibility-weighted mean of the data.
- Mixture parameters are only identifiable up to relabelling the components (label switching). Local maxima of the likelihood function can be avoided by running EM with different starting values and by comparing runs of the algorithm using the observed-data log-likelihood.

Exercise 19.1. Let X be geometrically distributed with parameter $p \in (0, 1)$, i.e. $f(x; p) = (1 - p)^{x-1}p$ for $x \in \{1, 2, \dots\}$, as in exercise 5.1.

1. Show that $\mathbb{P}_p(X \geq m) = (1 - p)^{m-1}$ for $m \in \{1, 2, \dots\}$ and use this result to show that $\mathbb{P}_p(X \geq m + k | X \geq m) = \mathbb{P}_p(X \geq k + 1)$ for all $k \geq 0$. Comment upon your result in the light of the memorylessness of the exponential distribution from exercise 5.6.
2. Show that

$$\mathbb{E}_p(X | X \geq m) = m - 1 + \mathbb{E}_p(X) = m - 1 + \frac{1}{p} \quad \text{for all } m \in \{1, 2, \dots\}.$$

Exercise 19.2. Suppose f_1 and f_2 are two known, distinct densities, and that $X_1, \dots, X_n$ are i.i.d. with density

$$f(x; \pi) = \pi f_1(x) + (1 - \pi) f_2(x),$$

where $\pi \in (0, 1)$ is the only unknown parameter. Let $Z_i \in \{1, 2\}$ be the component observation i belongs to, as in example 19.4.

1. Write down the observed-data log-likelihood $\ell(\pi)$ and the likelihood equation $\ell'(\pi) = 0$.
2. Write down the complete-data log-likelihood and show that the E step determines the responsibilities $\gamma_i(\pi_k) = \pi_k f_1(x_i) / (\pi_k f_1(x_i) + (1 - \pi_k) f_2(x_i))$ and the M step determines the update $\pi_{k+1} = \frac{1}{n} \sum_i \gamma_i(\pi_k)$.
3. Show that $\gamma_i(\pi) - \pi = \pi(1 - \pi) (f_1(x_i) - f_2(x_i)) / f(x_i; \pi)$ and thus that $\frac{1}{n} \sum_i \gamma_i(\pi) - \pi = \pi(1 - \pi) \ell'(\pi) / n$. Deduce that $\pi^* \in (0, 1)$ is a fixed point of the EM iteration if and only if it solves the likelihood equation, as in proposition 19.3.

Exercise 19.3. For the data, starting value and first iterate of the numerical illustration in example 19.4, compute the observed-data log-likelihood $\ell(\theta) = \sum_{i=1}^5 \log(\pi \varphi_1(x_i - \mu_1) + (1 - \pi) \varphi_1(x_i - \mu_2))$ for $\theta_0 = (1/2, 2, 4)$ and $\theta_1 = (0.4001, 0.5445, 5.9709)$. Show that the likelihood has increased. Which of the five observations contributes most to the increase in likelihood? Why?

Exercise 19.4. Consider the two-component normal mixture from example 19.4, but with known weight $\pi \in (0, 1)$ and known variance σ^2 . The only unknown quantities are the means $\theta = (\mu_1, \mu_2)$.

1. Show that the EM updates for this case are given by the two mean updates from (19.2), where the responsibilities are computed using the known π .
2. Compute the derivative of the observed-data log-likelihood and show that

$$\frac{\partial \ell}{\partial \mu_1}(\theta) = \frac{1}{\sigma^2} \sum_{i=1}^n \gamma_i(\theta) (x_i - \mu_1),$$

where $\gamma_i(\theta)$ is the responsibility at θ . Deduce that the likelihood equations coincide with the fixed-point equations of the EM iteration.

3. What are the updates in the limit $\pi \rightarrow 1$? Comment upon your result in the context of example 5.4.

Lecture 20

Hypothesis Testing

So far we have asked questions about the value of θ . In this lecture we will consider the second question from lecture 1: if there is a claim about θ , is it supported by the data? We will spend some time to carefully set up the language of hypothesis testing, which will be used in lectures 22, 23, 25 and 29. We will conclude the lecture by considering two examples, both about the mean of a normal sample.

20.1 Hypotheses

Throughout, $X_1, \dots, X_n$ follow a parametric model in the sense of definition 1.1. Questions like the one about fairness of a coin or about the effect of a treatment ask whether the true parameter value θ_0 belongs to a subset of the parameter space Θ .

Definition 20.1. A *hypothesis* is a statement of the form $\theta \in \Theta'$ for a subset $\Theta' \subseteq \Theta$. In a testing problem, the *null hypothesis* $H_0: \theta \in \Theta_0$ and the *alternative hypothesis* $H_1: \theta \in \Theta_1$ are given, where Θ_0 and Θ_1 are disjoint subsets of Θ . A hypothesis is called *simple*, if the corresponding subset consists of a single point, and *composite* otherwise.

As an abbreviation, from this lecture on we write θ_0 for the value chosen to test a null hypothesis. In lectures 16 and 17, the same symbol was used to denote the unknown, true value of the parameter. The context should make it clear which meaning is intended.

In most problems we have $\Theta_1 = \Theta \setminus \Theta_0$, *i.e.* we can assume that exactly one of the two hypotheses is true. For example, for the mean θ of a normal sample, the hypothesis $H_0: \theta = \theta_0$ is simple, whereas the one-sided alternative $H_1: \theta > \theta_0$, the two-sided alternative $H_1: \theta \neq \theta_0$ and the one-sided null hypothesis $H_0: \theta \leq \theta_0$ are all composite.

The two hypotheses have different roles: the null hypothesis is the default position, and the alternative is the claim which requires evidence before we can act on it. Section 20.3 will make this distinction precise.

20.2 Tests, Critical Regions and Errors

A test decides between the two hypotheses under consideration, and can be described by the set of data values where it rejects H_0 .

Definition 20.2. A *test* of H_0 against H_1 is a rule which for every possible data vector $x = (x_1, \dots, x_n)$ decides whether to “reject H_0 ” or “do not reject H_0 ”. The set

$$C = \{x \mid \text{the test rejects } H_0 \text{ for the data } x\}$$

is called the *critical region* (or rejection region) of the test. Alternatively, a test can be described by its *test function* $\varphi(x) = 1$ for $x \in C$ and $\varphi(x) = 0$ otherwise, where 1 indicates rejection.

In practice, the critical region is almost always of the form $C = \{x \mid T(x) > c\}$ or $C = \{x \mid T(x) \geq c\}$ for a statistic T given by definition 1.3, called the *test statistic*, and a constant c called the *critical value*. The two forms of the critical region differ if $T(x) = c$ has positive probability (exercise 20.2 shows an example of this case, for discrete distributions).

A test applied to random data will sometimes wrongly decide. This can happen in two different ways.

Definition 20.3. A *type I error* occurs, if the test wrongly rejects H_0 when H_0 is true. A *type II error* occurs, if the test wrongly does not reject H_0 when H_1 is true.

Table 20.1 summarises the possible outcomes of statistical tests, in relation to the truth of the hypothesis being tested.

	do not reject H_0	reject H_0
H_0 true	correct decision	type I error
H_1 true	type II error	correct decision (power)

Table 20.1: The four possible outcomes of a test.

To understand the interpretation of this table, we can use the following example: in a criminal trial, the null hypothesis is the presumption of innocence of the defendant, and the alternative is guilt. If an innocent person is convicted, this is a type I error. If a guilty person is acquitted, this is a type II error. The standard of proof “beyond reasonable doubt” is such that many acquittals of the guilty are acceptable in order to avoid too many wrongful convictions. The reason for this is that the two types of error have different consequences: the consequences of type I errors are more serious. Furthermore, since an acquittal does not prove innocence, we write “do not reject H_0 ” instead of “accept H_0 ” in the table.

20.3 The Power Function, Size and Level

The probabilities of the two errors depend on the true parameter value, and both can be found using a single function of θ .

Definition 20.4. The *power function* of a test with critical region C is given by

$$\beta(\theta) = \mathbb{P}_\theta(X \in C) = \mathbb{E}_\theta(\varphi(X)), \quad \theta \in \Theta,$$

i.e. the probability of rejecting H_0 when the true parameter value is θ . For $\theta \in \Theta_1$ the value $\beta(\theta)$ is called the *power* of the test at the alternative θ .

For $\theta \in \Theta_0$ the value $\beta(\theta)$ is the probability of a type I error, and for $\theta \in \Theta_1$ the value $1 - \beta(\theta)$ is the probability of a type II error. The two probabilities cannot both be reduced by making the critical region larger or smaller. If the size of the critical region is increased, the value $\beta(\theta)$ increases for all θ , *i.e.* type I errors become more likely and type II errors become less likely. If the size of the critical region is decreased, the opposite happens. For this reason, we have to choose one of the two types of error to control and it transpires that we can decide which type of error to control by considering the asymmetry between the two types of hypotheses: type I errors are errors where the default position is abandoned without good reason, and thus we should restrict the probability of type I errors first and only then should we consider the power of the test.

Definition 20.5. The *size* of a test with power function β is given by

$$\sup_{\theta \in \Theta_0} \beta(\theta),$$

where the supremum is the maximum probability of type I error among all possible null hypotheses. For $\alpha \in (0, 1)$, the test is a *level α test*, if the size of the test is less than or equal to α .

For a simple null hypothesis $H_0: \theta = \theta_0$, the size of the test equals $\beta(\theta_0)$. The number α is called the *significance level* and the values 0.05 and 0.01 are only conventions. Among all level α tests, we prefer the test with the largest power at the alternatives. In lecture 22 we will see how to find the test with optimal power, when both the null and the alternative hypothesis are simple.

20.4 Testing the Mean of a Normal Sample

We now construct a test from first principles for $X_1, \dots, X_n$ i.i.d. $N(\theta, \sigma^2)$, where the variance σ^2 is known. Here we write Φ for the standard normal distribution function and z_q for the q -quantile of the standard normal distribution, *i.e.* we have $\Phi(z_q) = q$. The most commonly used values of this function are $z_{0.95} = 1.645$ and $z_{0.975} = 1.96$.

Example 20.6. We want to test $H_0: \theta = \theta_0$ against $H_1: \theta > \theta_0$ at significance level α . If $\bar{X}$ is significantly larger than θ_0 , we can conclude that the alternative is more likely to be true. Thus, we can choose $\bar{X}$ as the test statistic and reject if $\bar{X} > c$. From proposition A.3 we know that under H_0 the standardised mean $\sqrt{n}(\bar{X} - \theta_0)/\sigma$ is standard normally distributed. Thus, the probability of type I errors is

$$\beta(\theta_0) = \mathbb{P}_{\theta_0}(\bar{X} > c) = \mathbb{P}_{\theta_0}\left(\frac{\sqrt{n}(\bar{X} - \theta_0)}{\sigma} > \frac{\sqrt{n}(c - \theta_0)}{\sigma}\right) = 1 - \Phi\left(\frac{\sqrt{n}(c - \theta_0)}{\sigma}\right).$$

Setting this probability equal to α , we find the critical value as $\sqrt{n}(c - \theta_0)/\sigma = z_{1-\alpha}$ and thus the level α test rejects H_0 if

$$\bar{X} > \theta_0 + z_{1-\alpha} \frac{\sigma}{\sqrt{n}}, \quad \text{equivalently if} \quad Z = \frac{\sqrt{n}(\bar{X} - \theta_0)}{\sigma} > z_{1-\alpha}. \quad (20.1)$$

This is the one-sided z -test. Since $\sqrt{n}(\bar{X} - \theta)/\sigma$ is standard normally distributed under $\mathbb{P}_\theta$, we can use the same standardisation to determine the power of the

test. The power function is

$$\beta(\theta) = \mathbb{P}_\theta\left(\frac{\sqrt{n}(\bar{X} - \theta)}{\sigma} > z_{1-\alpha} - \frac{\sqrt{n}(\theta - \theta_0)}{\sigma}\right) = \Phi\left(\frac{\sqrt{n}(\theta - \theta_0)}{\sigma} - z_{1-\alpha}\right). \quad (20.2)$$

Since the function Φ is increasing, the graph of β is an S-shaped curve which starts at 0 and converges to 1, passing through the point (θ_0, α) . Close to θ_0 the power is only slightly larger than α , but as n increases, the curve gets steeper. This fact is used in exercise 20.1 to determine the required sample size.

For $n = 25$, $\sigma = 2$, $\theta_0 = 10$ and $\alpha = 0.05$ the test rejects if $\bar{X} > 10 + 1.645 \cdot 2/5 = 10.658$. At the alternative $\theta = 11$ the power is $\Phi(2.5 - 1.645) = \Phi(0.855) = 0.804$, at $\theta = 10.5$ it is only $\Phi(1.25 - 1.645) = \Phi(-0.395) = 0.346$: a shift of half a standard deviation is detected in barely a third of all samples.

Since the power function (20.2) is increasing, the same test also has size $\sup_{\theta \leq \theta_0} \beta(\theta) = \beta(\theta_0) = \alpha$ for the composite null hypothesis $H_0: \theta \leq \theta_0$, attained at the boundary point; this is why definition 20.5 takes a supremum.

Example 20.7. Assume that we want to test for both directions, *i.e.* we test $H_0: \theta = \theta_0$ against $H_1: \theta \neq \theta_0$ and we reject for $|\bar{X} - \theta_0| > c$. Under H_0 the test statistic $Z = \sqrt{n}(\bar{X} - \theta_0)/\sigma$ is standard normally distributed and, since the normal distribution is symmetric, we have

$$\mathbb{P}_{\theta_0}(|\bar{X} - \theta_0| > c) = \mathbb{P}_{\theta_0}\left(|Z| > \frac{\sqrt{n}c}{\sigma}\right) = 2\left(1 - \Phi\left(\frac{\sqrt{n}c}{\sigma}\right)\right).$$

Setting this equal to α , we find $\sqrt{n}c/\sigma = z_{1-\alpha/2}$ and thus the level α two-sided z -test rejects H_0 if

$$|\bar{X} - \theta_0| > z_{1-\alpha/2} \frac{\sigma}{\sqrt{n}}, \quad \text{equivalently if} \quad |Z| > z_{1-\alpha/2}. \quad (20.3)$$

Let $\delta = \sqrt{n}(\theta - \theta_0)/\sigma$. Then, similar to the result in example 20.6, we find the power function

$$\beta(\theta) = \Phi(-z_{1-\alpha/2} - \delta) + 1 - \Phi(z_{1-\alpha/2} - \delta). \quad (20.4)$$

This function is symmetric about θ_0 and has its minimum value α at this point. Thus, the graph of this power function is U-shaped instead of S-shaped. For $n = 25$, $\sigma = 2$, $\theta_0 = 10$ and $\alpha = 0.05$ the test rejects if $|\bar{X} - 10| > 1.96 \cdot 2/5 = 0.784$ and at $\theta = 11$ we have $\delta = 2.5$ and the power $\Phi(-4.46) + 1 - \Phi(-0.54) = 0.705$.

At $\theta = 11$ the one-sided test has power 0.804 against 0.705, since it rejects all values with probability α from the upper tail. Thus, the one-sided test is superior to the two-sided test for all alternatives above θ_0 , but is useless for alternatives below. Before we consider the data, we need to decide which departures from H_0 are of interest. In lecture 22 we will see that there is no single test which is best for both directions, and in lecture 24 we will consider how size and power behave in simulation.

20.5 The p -Value

A test at a fixed level only states whether H_0 is rejected or not. The p -value, introduced in this section, allows us to quantify the strength of the evidence in a more detailed way.

Definition 20.8. Consider a family of tests with critical regions C_α for all levels $\alpha \in (0, 1)$, such that $C_\alpha \subseteq C_{\alpha'}$ whenever $\alpha < \alpha'$. Then the *p-value* of the data x is given by

$$p(x) = \inf \{ \alpha \in (0, 1) \mid x \in C_\alpha \},$$

the smallest level at which the data lead to rejection of H_0 .

The tests from section 20.4 are nested in this way, since $z_{1-\alpha}$ decreases as α increases. For a test which rejects for large values of a statistic T and which has a simple null hypothesis, the smallest level at which the data can lead to rejection is

$$p(x) = \mathbb{P}_{\theta_0}(T(X) \geq T(x)). \quad (20.5)$$

This is the probability under H_0 of the test statistic taking a value at least as extreme as the observed value. For the one-sided z -test this probability is $p(x) = 1 - \Phi(z)$, for the observed value $z = \sqrt{n}(\bar{x} - \theta_0)/\sigma$. For the two-sided test we get $p(x) = 2(1 - \Phi(|z|))$. For the situation of example 20.6, the observed mean $\bar{x} = 10.9$ leads to $z = 2.25$ and thus to $p = 1 - \Phi(2.25) = 0.0122$ for the one-sided test and $p = 0.0244$ for the two-sided test. We can see that we can reject H_0 at level α if and only if $p(x) \leq \alpha$, and thus the *p-value* is a compact summary of a test. Under H_0 , the *p-value* of a continuous test statistic is uniformly distributed on $(0, 1)$ (see exercise 20.4). This result is used in lecture 30.

Both statements assume that the test statistic is continuous. For discrete test statistics, *e.g.* the number of successes in a Bernoulli sample, the tail probability (20.5) can only take the countably many values $\mathbb{P}_{\theta_0}(T(X) \geq t)$ and the *p-value* is discontinuous. Most levels α are not attained by any critical region. We can still perform the test for $p(x) \leq \alpha$ and the resulting test is *conservative*. The size of the test is the largest attainable value which does not exceed the level α . The test will reject a true H_0 less often than the nominal level suggests and the test will lose some power. The *p-value* is no longer uniformly distributed either, but still satisfies $\mathbb{P}_{\theta_0}(p(X) \leq \alpha) \leq \alpha$. This is what allows the level to still be valid.

Since the *p-value* is a probability, it is tempting to interpret the results in wrong ways; the following statements about a *p-value* of 0.0122 are all wrong.

- “The probability that H_0 is true is 0.0122.” Since θ is a fixed, unknown number, H_0 is either true or false; the probability in (20.5) refers to the data. To get a probability for a hypothesis, we have to introduce a prior distribution for θ , *i.e.* we have to follow the Bayesian approach from lecture 28.
- “The probability that the observed data could have been produced by chance is 0.0122.” The *p-value* corresponds to extreme values of T , not to the probability of the data.
- “If we had observed a *p-value* of 0.4, we would have known that H_0 is true.” Data which is compatible with H_0 may be compatible with many different alternatives: the test from example 20.6 has power only 0.346 at $\theta = 10.5$.
- “The effect is important, because *p* is small.” If n is large enough, any difference between θ and θ_0 will result in a tiny *p-value*. Reports should include both the estimate and the standard error.

None of these mistakes invalidates the p -value as a tool, but they do show that the p -value answers a very specific question, and that the same value should not be used to answer different questions.

Summary.

- A testing problem contrasts a null hypothesis $H_0: \theta \in \Theta_0$ against an alternative $H_1: \theta \in \Theta_1$. The hypothesis is simple, if the corresponding subset of Θ is a single point, and composite otherwise.
- A test is given by its critical region C , typically of the form $T(x) > c$ for a test statistic T and critical value c .
- Type I errors are cases where H_0 is true but is wrongly rejected by the test, type II errors are cases where H_0 is false but is not (correctly) rejected. The power function $\beta(\theta) = \mathbb{P}_\theta(X \in C)$ gives the probability of type I errors for Θ_0 and the power of the test for Θ_1 .
- The size of a test is $\sup_{\Theta_0} \beta$. A level α test has size at most α . The level is chosen first, and then the power of the test can be considered. For example, for the normal mean with known variance, the one-sided and two-sided z -tests can be derived.
- The p -value is the smallest level at which the data lead to rejection of the hypothesis. Note that this is not the probability that H_0 is true, and a large p -value does not establish H_0 .

Exercise 20.1. Consider the one-sided z -test from example 20.6, at level α , and let $\theta_1 > \theta_0$ be the alternative and $1 - \gamma \in (\alpha, 1)$ be the target power.

1. Show that the test has power at least $1 - \gamma$ at θ_1 if and only if

$$n \geq \left(\frac{(z_{1-\alpha} + z_{1-\gamma}) \sigma}{\theta_1 - \theta_0} \right)^2.$$

2. For $\sigma = 2$, $\theta_1 - \theta_0 = 0.5$, $\alpha = 0.05$ and target power 0.9 (so that $z_{0.9} = 1.282$), determine the smallest sample size which achieves the target power.
3. How does the required sample size change, when the difference $\theta_1 - \theta_0$ is halved? What happens when σ is doubled?

Exercise 20.2. Let $X_1, \dots, X_{10}$ be i.i.d. Bernoulli with success probability p and consider testing $H_0: p = 1/2$ against $H_1: p > 1/2$ using the test which rejects if $T = \sum_i X_i \geq c$ for an integer c .

1. Write down the size of the test as a function of c , using the binomial distribution of T under H_0 .
2. Determine the smallest c such that the test has level 0.05. Compute the size of the test for this case. Justify your answer by showing that no choice of c can lead to a size of exactly 0.05.

3. Determine the power of the test from the previous part when $p = 0.7$, $p = 0.8$ and $p = 0.9$. Comment on the capabilities and limitations of the test with ten observations.
4. Nine successes are observed. Compute the p -value and state the decision at level 0.05.

Exercise 20.3. Let $X_1, \dots, X_n \sim \text{Uniform}(0, \theta)$ be i.i.d. as in example 1.2 and let $M = \max_i X_i$, where the density was found in example 8.10. We test $H_0: \theta = \theta_0$ against $H_1: \theta > \theta_0$.

1. Show that $\mathbb{P}_\theta(M \leq m) = (m/\theta)^n$ for $0 \leq m \leq \theta$.
2. Consider the test which rejects if $M > c$, for a constant $c \leq \theta_0$. Determine the size of this test and the value of c which makes the size equal to α .
3. Determine the power function of the test from the previous part, and evaluate this function at $n = 5$, $\theta_0 = 1$, $\alpha = 0.05$ and $\theta = 1.2$.
4. Consider the second test, which rejects if $M > \theta_0$. Show that this test has size 0. Determine the power function of this test and compare the power functions of the two tests. Which of the two tests would you choose and why?

Exercise 20.4. A sample of size $n = 16$ from $N(\theta, \sigma^2)$ with known $\sigma = 4$ has mean $\bar{x} = 52.3$. We want to test $\theta_0 = 50$.

1. Determine the p -value of the one-sided z -test with $H_0: \theta = 50$ against $H_1: \theta > 50$ and state the decisions for 0.05 and 0.01.
2. Determine the p -value of the two-sided test against $H_1: \theta \neq 50$ and state the same two decisions.
3. Your colleague summarises the first part of the question as “there is a 1.1% chance that the mean is 50”. What is wrong with this sentence? Write one correct sentence to summarise the first part of the question.
4. Show that under H_0 the p -value of the one-sided z -test, $p(X) = 1 - \Phi(Z)$ with $Z = \sqrt{n}(\bar{X} - \theta_0)/\sigma$, satisfies $\mathbb{P}_{\theta_0}(p(X) \leq u) = u$ for all $u \in (0, 1)$.

Lecture 21

EM in R

This is the R session for week 7. In lecture 19 the EM algorithm was introduced and an ascent property was proved. In example 19.4 the E and M steps were illustrated for a two-component normal mixture. Today we will type the two steps into R functions, check the functions against the hand computation from the example, and then run the functions on a simulated sample of size 200. We will conclude by considering the problem of label switching and the effect of a bad start. The written reference for this session is appendix B.

21.1 The E and M Steps as R Functions

Throughout, the parameter is the vector $\theta = (\pi, \mu_1, \mu_2)$ and the common standard deviation σ is known. The E step of example 19.4 computes the responsibilities γ_i , the M step computes the updates (19.2), and both involve only vector arithmetic. The function `dnorm` is also used to compute the normal density. We also need to compute the observed-data log-likelihood ℓ . By theorem 19.2, this must increase.

```
resp <- function(x, theta, sigma) {
  a <- theta[1] * dnorm(x, mean = theta[2], sd = sigma)
  b <- (1 - theta[1]) * dnorm(x, mean = theta[3], sd = sigma)
  a / (a + b)
}
mstep <- function(x, gamma) {
  c(mean(gamma),
    sum(gamma * x) / sum(gamma),
    sum((1 - gamma) * x) / sum(1 - gamma))
}
loglik <- function(x, theta, sigma) {
  sum(log(theta[1] * dnorm(x, mean = theta[2], sd = sigma) +
    (1 - theta[1]) * dnorm(x, mean = theta[3], sd = sigma)))
}
```

We first test our function on the data from example 19.4. For this we use the five observations $x = (0, 1, 5, 6, 7)$ where $\sigma = 1$ and initial value $\theta_0 = (1/2, 2, 4)$. The results of one iteration of the algorithm are then compared to the responsibilities and iterate θ_1 and to the two values for the log-likelihood computed by hand.

```
x      <- c(0, 1, 5, 6, 7)
theta  <- c(0.5, 2, 4)
gamma  <- resp(x, theta, 1)
round(gamma, 4)
theta1 <- mstep(x, gamma)
theta1
c(loglik(x, theta, 1), loglik(x, theta1, 1))
```

```
## [1] 0.9975 0.9820 0.0180 0.0025 0.0003
## [1] 0.4000671 0.5444727 5.9709581
## [1] -17.518842 -9.212934
```

Everything agrees with the hand computation ($\mu_{2,1}$ differs in the last digit because the example rounded the intermediate sums), and the last line is the check for exercise 19.3. We now iterate, keeping the parameter and the log-likelihood after each step in a matrix. Before you run the loop, predict the log-likelihood after the second iteration: far above -9.21 , slightly above it, or equal to it?

```
theta <- c(0.5, 2, 4)
path  <- matrix(0, nrow = 4, ncol = 4,
                dimnames = list(1:4, c("pi", "mu1", "mu2", "loglik")))
for (k in 1:4) {
  theta <- mstep(x, resp(x, theta, 1))
  path[k, ] <- c(theta, loglik(x, theta, 1))
}
round(path, 4)
```

```
##      pi    mu1    mu2 loglik
## 1 0.4001 0.5445 5.971 -9.2129
## 2 0.4000 0.5001 6.000 -9.2097
## 3 0.4000 0.5001 6.000 -9.2097
## 4 0.4000 0.5001 6.000 -9.2097
```

The second iteration gains only 0.003, and after it nothing moves: the responsibilities are so close to 0 and 1 that each mean is the mean of its own group, and the iteration has reached its fixed point, a root of the likelihood equation by proposition 19.3.

21.2 A Larger Sample

Convergence in two steps is a consequence of the large difference between the two groups. Here we simulate a sample where the groups overlap. We can simulate such a sample by first randomly choosing a group, using `rbinom` to get $Z_i = 1$ with probability $\pi = 0.3$. We then simulate the measurements from the normal distribution corresponding to the group, *i.e.* with means 0 and 3 and $\sigma = 1$. The command `ifelse` can be used to choose the correct mean for each observation. The box below contains the complete script for simulation and EM iteration, with lines shuffled (except for the body of the loop). Before you can run the script, you have to arrange these lines into a working script. The loop finishes when the log-likelihood increases by less than 10^{-8} , as described in section 19.5.

The vector `ll` stores the log-likelihood after each iteration, with the initial value at position 1.

```
theta <- c(0.5, q[1], q[2])
for (k in 1:1000) {
  theta <- mstep(x, resp(x, theta, 1))
  ll[k + 1] <- loglik(x, theta, 1)
  if (ll[k + 1] - ll[k] < 1e-8) break
}
z <- rbinom(n, size = 1, prob = 0.3)
n <- 200
q <- quantile(x, c(0.25, 0.75), names = FALSE)
set.seed(2026)
ll <- loglik(x, theta, 1)
x <- rnorm(n, mean = ifelse(z == 1, 0, 3), sd = 1)
```

The rule is the one from lecture 9: each line must come after all lines on which it depends, so the order of lines must be `seed`, `n`, `z` and `x`, quartiles `q`, start value, first entry of `ll`, and then the loop. The start values are chosen such that the weight is at 1/2 and the two means are at the two quartiles. This is a cheap way to split the sample into two groups. Before you try the script, try to answer the following two questions: How many iterations will the loop take? (In the hand example from lecture, two iterations were needed.) Will the values for the log-likelihood increase by the same amount in each iteration, or differently?

```
table(z)
k
round(theta, 4)
round(ll[1:6], 3)
round(ll[k + 1], 3)
```

```
## z
##   0   1
## 152  48
## [1] 28
## [1] 0.2397 0.2202 3.0761
## [1] -375.257 -370.274 -365.880 -362.046 -359.040 -356.973
## [1] -354.505
```

The fixed point $\hat{\theta} = (0.240, 0.220, 3.076)$ recovers the true group-1 proportion $48/200 = 0.24$ and the two means to within about two tenths. The loop required 28 iterations, because the responsibilities of the two means were far away from 0 and 1 and because in each iteration the means were only slightly changed. The recorded values are increasing with each iteration, as theorem 19.2 guarantees, but at decreasing rates: 5.0, then 4.4, 3.8, 3.0, 2.1, as predicted by the geometric convergence of section 19.5. By proposition 19.3 the fixed point is a root of the likelihood equation. To test this, we can use `optim`, as shown in lecture 18, by maximising `loglik` with the same starting values as we used in the EM algorithm. If we use the command `hessian = TRUE`, we can get the standard errors from section 17.3. (The EM algorithm does not provide the standard errors.)


```
fit <- optim(c(0.5, q[1], q[2]), function(par) -loglik(x, par, 1),
            hessian = TRUE)
round(fit$par, 4)
round(sqrt(diag(solve(fit$hessian))), 4)
```

```
## [1] 0.2397 0.2202 3.0761
## [1] 0.0383 0.1965 0.0961
```

The two methods agree to four decimal places, and the standard errors say that the fitted weight is 0.240 ± 0.038 , with the smaller group's mean determined only half as well as the larger group's. At the fixed point the responsibility of group 1 is a decreasing function of x which equals $1/2$ where $\pi\varphi_\sigma(x - \mu_1) = (1 - \pi)\varphi_\sigma(x - \mu_2)$. Before you look at figure 21.1, predict where the crossing lies: at the midpoint 1.65 of the two fitted means, or above or below this point?

```
gamma <- resp(x, theta, 1)
half <- (theta[3]^2 - theta[2]^2) / (2 * (theta[3] - theta[2])) +
        log(theta[1] / (1 - theta[1])) / (theta[3] - theta[2])
round(half, 3)
plot(x, gamma, pch = ifelse(z == 1, 16, 1), xlab = "x",
     ylab = "responsibility of group 1")
```

```
## [1] 1.244
```




Figure 21.1: The fitted mixture density with its two weighted components (top panel) and the responsibilities of group 1 at the fixed point as a function of the observations (bottom panel). The filled circles give the observations which were assigned to group 1. The dotted lines give the responsibility $1/2$ and the crossing point 1.244.

The crossing happens at 1.244, well below the mid-point 1.648, since the weights are unequal: an observation which is half-way between the means has three times the probability of being from the larger group. The filled and open circles give the true group of each observation, which cannot be known by EM: the responsibilities agree with the truth when the observations are far away from the crossing point, but between 0.7 and 1.6 both types of observations can be found and the group cannot be decided. The top panel is also plotted using the script `R/S21-mixture.R`; the three functions are given in `R/S21-hand.R`.

21.3 Label Switching and a Bad Start

Section 19.5 lists two ways in which a start can lead EM astray; both cases are represented in the sample. The script `R/S21-starts.R` encapsulates the loop as a function `em` and then runs this function with four different starts: the quartile start, the start with the two means swapped, the start $(1/2, 3, 3)$ with both means equal, and the start $(1/2, 2.39, 2.41)$ with nearly equal means. Before you look at the table, predict the outcome of each row. Where does the start with the swapped means end up? What is the log-likelihood there? What happens to the responsibilities when $\mu_1 = \mu_2$ and thus to the next update?

```
good    <- em(x, c(0.5, q[1], q[2]), 1)
swapped <- em(x, c(0.5, q[2], q[1]), 1)
equal   <- em(x, c(0.5, 3, 3), 1)
near    <- em(x, c(0.5, 2.39, 2.41), 1)
```

##		pi	mu1	mu2	loglik	iterations
##	good	0.2397	0.2202	3.0761	-354.5046	28
##	swapped	0.7603	3.0761	0.2202	-354.5046	28
##	equal	0.5000	2.3915	2.3915	-421.8067	2
##	near	0.2397	0.2202	3.0761	-354.5046	33

The start with swapped means ends in the mirror image $(0.760, 3.076, 0.220)$ after the same 28 iterations and with the same log-likelihood: both parameter values lead to the same density, the model is not identifiable in the sense of definition 4.10 and the choice of labelling in the start determines which of the two maxima we reach. The start with equal parameters is the trap. In this case, with $\mu_1 = \mu_2$, the two normal densities in the responsibility cancel, all γ_i are equal to π and the M step sets both means equal to the sample mean while leaving the weight unchanged; another step does not change anything and the loop stops after two iterations at $(1/2, 2.3915, 2.3915)$, effectively a single normal distribution, with a log-likelihood 67 below the maximum. Theorem 19.2 is not violated: the log-likelihood did not decrease but failed to increase. The resulting resting point is a root of the likelihood equation, by proposition 19.3, but a saddle point instead of a maximum and thus EM will not be able to distinguish between the two. The nearly equal start shows that the trap is narrow: the iteration moves away from the saddle point and then climbs to the maximum, in 33 iterations, as shown in figure 21.2 for the three traces.

For this model, with known common variance, a grid of hundreds of starting values finds only the two mirror-image maxima and the saddle line $\mu_1 = \mu_2$ as stable solutions. In contrast, if the variances were unknown or if there were




Figure 21.2: Log-likelihood of the given data as a function of the iteration number for the EM algorithm, for the start with quartile-spaced means (solid line), for the start with nearly equal means (dotted line) and for the start with equal means (dash-dotted line with crosses). The trace of the start with interchanged means coincides with the solid line.

three or more components, additional local maxima would occur and the procedure from section 19.5 of starting the EM algorithm with several runs and then choosing the run with the highest observed-data log-likelihood becomes essential. For a two-component mixture, the twelve lines of `resp`, `mstep` and `loglik` are all we need. For serious work there are packages: `mclust` fits mixtures of this kind, and `requireNamespace("mclust", quietly = TRUE)` reports whether it is installed. Nothing in this session relies on it.

21.4 A Task to Take Home

The following task is in the style of a practical report. The published scripts of this lecture can be used as a starting point. Using your student ID as the seed, choose a weight π between 0.2 and 0.5, and choose two means which are two to three units apart. Simulate one sample of size $n = 300$ from the mixture with $\sigma = 1$. Run the EM algorithm with three different starting values (the quartile start and two of your choice) and show the fixed point, log-likelihood and number of iterations for each run in a table. Justify which run you would keep and mention if label switching occurred. Finally, use `optim` with `hessian = TRUE` to estimate the standard errors for the kept run. For each parameter, show the estimate, the standard error and the interval $\hat{\theta} \pm 1.96 \text{ se}(\hat{\theta})$ from corollary 17.3, in three sentences. For each sentence, state whether the interval contains the value you chose. Plot the responsibilities against the data. Write a short paragraph about how many observations have responsibility values between 0.2 and 0.8. Comment on how well-separated the two groups are in your plot. Finally, repeat the whole exercise with the two means only one unit apart, and say in one sentence what changes.

Typing the commands yourself is an important step towards understanding how the R scripts work. Before the next session, please type in the commands

from this session, and check that you get the same results as shown in the text.

Summary.

- The written reference for the R code in this session is appendix B.
- The E-step and M-step from example 19.4 involve only a few lines of vector arithmetic. For five observations, these lines can reproduce the manual computation, and the algorithm converges to the fixed point in two steps.
- For a sample of size 200 with overlapping groups, it takes 28 iterations for EM to converge. The log-likelihood increases in every step, by increasingly small amounts as required by theorem 19.2. The `optim` function finds the same fixed point as the manual calculation, as predicted by proposition 19.3, and `optim` can be used to compute the standard errors.
- The responsibility of group 1 is a decreasing S-shaped function of x . The function crosses $1/2$ below the mid-point of the means when group 1 is the smaller group. Close to the crossing point it is no longer possible to decide which group a sample belongs to.
- If the labels are swapped before starting the EM algorithm, the algorithm converges to the mirror-image maximum with the same log-likelihood. If the two groups have the same mean, the EM algorithm converges to a saddle point immediately, *i.e.* a root of the likelihood equation which is not a maximum. To avoid this, the EM algorithm should be started with different values and the best of the resulting runs should be chosen, based on the log-likelihood of the observed data.

Lecture 22

The Neyman–Pearson Lemma

In lecture 20 it was clear that for the mean of a normal sample we could not do better than to reject for large values of $\bar{X}$. This lecture proves that this is indeed correct: for the most basic testing problem, where one distribution is being tested against a different distribution, the Neyman–Pearson lemma allows us to identify the most powerful test. This is in analogy to the role of the Cramer–Rao bound from lecture 13 for estimators. The same test is optimal for all one-sided alternatives, but there is no single test which is optimal for a two-sided alternative. In fact, in lecture 23 we will see the required tool for this case.

22.1 Simple Hypotheses and Most Powerful Tests

We consider the problem of testing a simple null hypothesis $H_0: \theta = \theta_0$ against a simple alternative $H_1: \theta = \theta_1$, where $\theta_0 \neq \theta_1$ are two fixed points of the parameter space, for data $X = (X_1, \dots, X_n)$ with joint density or probability weights $f(x; \theta)$ as in lecture 8. Under either hypothesis the distribution of the data is completely specified; section 22.4 shows that the solution of this problem also settles more realistic ones.

We will compare different tests, so we write $\beta_C(\theta) = \mathbb{P}_\theta(X \in C)$ for the power function of definition 20.4 of the test with critical region C . For simple hypotheses, the size from definition 20.5 and the power are the single numbers $\beta_C(\theta_0)$ and $\beta_C(\theta_1)$. Among all tests with level α , *i.e.* with $\beta_C(\theta_0) \leq \alpha$, we want to find the test which makes type II errors least often.

Definition 22.1. Let $0 < \alpha < 1$. A test with critical region C is said to be *most powerful at level α* for testing $H_0: \theta = \theta_0$ against $H_1: \theta = \theta_1$, if $\beta_C(\theta_0) \leq \alpha$ and

$$\beta_C(\theta_1) \geq \beta_{C'}(\theta_1)$$

for all critical regions C' with $\beta_{C'}(\theta_0) \leq \alpha$.

As with the uniformly minimum variance unbiased estimator from definition 10.3, the definition does not give a method to find such a test or even whether such a test exists; we will answer both of these questions in the next section.

22.2 The Lemma

If $f(x; \theta_1)$ is much larger than $f(x; \theta_0)$, the data favour the alternative, and we should reject. The natural test statistic is the *likelihood ratio*

$$R(x) = \frac{f(x; \theta_1)}{f(x; \theta_0)} = \frac{L(\theta_1)}{L(\theta_0)},$$

and the natural test rejects if $R(X)$ exceeds some threshold k . The lemma states that no test with the same level can have more power. We write the inequality $f(x; \theta_1) > k f(x; \theta_0)$ instead of $R(x) > k$ so that points with $f(x; \theta_0) = 0$ do not require separate treatment and so that the boundary of the critical region can include any points where equality holds, as required for discrete models and in exercise 22.3.

Theorem 22.2 (Neyman–Pearson lemma). *Let $k \geq 0$ and let C be a critical region with*

$$\{x \mid f(x; \theta_1) > k f(x; \theta_0)\} \subseteq C \subseteq \{x \mid f(x; \theta_1) \geq k f(x; \theta_0)\},$$

where $\alpha = \beta_C(\theta_0)$ is the size of the critical region. Then the following statements hold.

1. *The test with critical region C is the most powerful test at level α for the hypothesis $H_0: \theta = \theta_0$ against $H_1: \theta = \theta_1$.*
2. *If C' is another critical region, which is most powerful at level α for the same problem, then C and C' differ at most on the boundary set $\{x \mid f(x; \theta_1) = k f(x; \theta_0)\}$, except for a set of observations with probability zero for every $\theta \in \Theta$.*

Proof. Let C' be a critical region with $\beta_{C'}(\theta_0) \leq \alpha$ and let $\mathbf{1}_C$ and $\mathbf{1}_{C'}$ be the indicator functions for the two regions. Then the proof is based on the observation that

$$h(x) = (\mathbf{1}_C(x) - \mathbf{1}_{C'}(x))(f(x; \theta_1) - k f(x; \theta_0)) \geq 0 \quad \text{for all } x. \quad (22.1)$$

If $x \in C$, the first factor is either 0 or 1 and the second factor is non-negative by the choice of C . If $x \notin C$, the first factor is either 0 or -1 and the second factor is zero or negative by the choice of C .

Integrating (22.1) over all x , with a sum in place of the integral for a discrete model, and using $\int \mathbf{1}_A(x) f(x; \theta) dx = \mathbb{P}_\theta(X \in A) = \beta_A(\theta)$ for each of the four terms, we find

$$0 \leq \int h(x) dx = (\beta_C(\theta_1) - \beta_{C'}(\theta_1)) - k(\beta_C(\theta_0) - \beta_{C'}(\theta_0)). \quad (22.2)$$

Since $\beta_C(\theta_0) = \alpha \geq \beta_{C'}(\theta_0)$ and $k \geq 0$, the subtracted term is non-negative and thus we have

$$\beta_C(\theta_1) - \beta_{C'}(\theta_1) \geq k(\alpha - \beta_{C'}(\theta_0)) \geq 0,$$

and this is the first statement.

For the second statement assume that C' is also most powerful at level α . Then $\beta_{C'}(\theta_1) \geq \beta_C(\theta_1)$ by definition 22.1, and $\beta_C(\theta_1) \geq \beta_{C'}(\theta_1)$ by the first

statement, so that the two powers are equal and the first bracket in (22.2) vanishes. The right-hand side of (22.2) is then $-k(\alpha - \beta_{C'}(\theta_0)) \leq 0$, and together with the inequality on the left this gives $\int h(x) dx = 0$. Now $h(x) > 0$ holds exactly for those x which lie in one of the regions C and C' but not in the other and at which $f(x; \theta_1) \neq k f(x; \theta_0)$, since at such points both factors of h are non-zero and, by the first paragraph, of the same sign. For a discrete model a sum of non-negative terms is zero only if every term is zero, and thus this set is empty. For a model with densities, a non-negative function with integral zero vanishes outside a set of measure zero, and a set of measure zero has probability zero under every density $f(x; \theta)$. This completes the proof. $\square$

Two comments help in using the lemma. First, the constant k is rarely computed: whenever R is a strictly monotone function of a simpler statistic T , we rewrite $R(x) > k$ as $T(x) > c$ or $T(x) < c$ and choose c from the distribution of $T(X)$ under θ_0 to give size α . Second, the lemma shows that the test is most powerful at the level equal to its own size, whatever value that size takes. For models with densities, $R(X)$ usually has a continuous distribution, and thus in theory any size can be achieved and the boundary set in the second statement has probability zero. Together these results show that the second statement implies that the most powerful test is essentially unique. In addition, it can be shown (for example in exercise 22.4) that the most powerful test can be chosen to depend on the data only via a sufficient statistic.

Remark. For discrete models the size $\mathbb{P}_{\theta_0}(R(X) > k)$ of tests can only take certain values as k changes and it is in general not possible to obtain a test with a given size α . The usual solution in textbooks, to randomise the test on the boundary set to achieve the required size, is not taken here. Instead, we consider the test with the largest possible size which is less than or equal to α . This test, by the result of exercise 22.1, will be the most powerful test for the smaller level.

22.3 The Normal Mean

We now apply the lemma to the normal-mean example from lecture 20: the test constructed there using common sense is the most powerful one.

Example 22.3. Let $X_1, \dots, X_n$ be i.i.d. $N(\theta, \sigma^2)$ where σ^2 is known, and consider the hypothesis $H_0: \theta = \theta_0$ against the alternative $H_1: \theta = \theta_1$, where $\theta_1 > \theta_0$. The joint density of the observations is $f(x; \theta) = (2\pi\sigma^2)^{-n/2} \exp(-\sum_i (x_i - \theta)^2 / (2\sigma^2))$. In the ratio of two densities, the constant factors cancel. After expanding the squares in the definition of the test statistic, we find

$$\sum_{i=1}^n (x_i - \theta_1)^2 - \sum_{i=1}^n (x_i - \theta_0)^2 = -2(\theta_1 - \theta_0) \sum_{i=1}^n x_i + n(\theta_1^2 - \theta_0^2),$$

and thus

$$\log R(x) = \frac{n(\theta_1 - \theta_0)}{\sigma^2} \bar{x} - \frac{n(\theta_1^2 - \theta_0^2)}{2\sigma^2}.$$

Since $\theta_1 - \theta_0 > 0$, the right-hand side is a strictly increasing function of $\bar{x}$ and the inequality $R(x) > k$ is equivalent to $\bar{x} > c$ where $c = \sigma^2 \log k / (n(\theta_1 - \theta_0)) + (\theta_1 + \theta_0)/2$. Thus, the critical region is of the form $\{\bar{x} > c\}$ and we can choose the

value of c by considering the size of the test: under H_0 we have $\bar{X} \sim N(\theta_0, \sigma^2/n)$ by proposition A.3 and thus

$$\mathbb{P}_{\theta_0}(\bar{X} > c) = 1 - \Phi\left(\frac{\sqrt{n}(c - \theta_0)}{\sigma}\right) = \alpha$$

if and only if

$$c = c_\alpha = \theta_0 + z_{1-\alpha} \frac{\sigma}{\sqrt{n}},$$

where $z_{1-\alpha}$ is the $(1 - \alpha)$ -quantile of the standard normal distribution. By theorem 22.2 the test which rejects for $\bar{X} > c_\alpha$ is most powerful at level α and, since the boundary set $\{\bar{x} = c_\alpha\}$ has probability zero, this is the only test up to sets of probability zero. This is the test from example 20.6. The power of this test for the alternative is

$$\beta(\theta_1) = \mathbb{P}_{\theta_1}(\bar{X} > c_\alpha) = 1 - \Phi\left(z_{1-\alpha} - \frac{\sqrt{n}(\theta_1 - \theta_0)}{\sigma}\right).$$

For $n = 25$, $\sigma = 2$, $\theta_0 = 10$ and $\alpha = 0.05$ we have $z_{0.95} = 1.645$ and thus $c_\alpha = 10 + 1.645 \cdot 2/5 = 10.658$, and the power against $\theta_1 = 11$ is $1 - \Phi(1.645 - 2.5) = \Phi(0.855) = 0.804$, the value we found in lecture 20. The lemma allows us to turn this number into a statement about all possible procedures: no test of level 0.05, based on the given 25 observations, can be better.

22.4 One-Sided Alternatives

The threshold $c_\alpha = \theta_0 + z_{1-\alpha}\sigma/\sqrt{n}$ in example 22.3 does not involve θ_1 : the alternative entered the computation only through the sign of $\theta_1 - \theta_0$, and thus every $\theta_1 > \theta_0$ leads to the same critical region. This has a consequence for the composite problem of testing $H_0: \theta = \theta_0$ against $H_1: \theta > \theta_0$, the problem of interest in lecture 20.

Proposition 22.4. *Consider testing $H_0: \theta = \theta_0$ against $H_1: \theta > \theta_0$. Assume that C is a critical region with $\beta_C(\theta_0) = \alpha$ which, for every $\theta_1 > \theta_0$, satisfies the condition from theorem 22.2 for testing θ_0 against θ_1 , for some constant $k = k(\theta_1)$ which may depend on θ_1 . Then, for every $\theta_1 > \theta_0$ and every critical region C' with $\beta_{C'}(\theta_0) \leq \alpha$, we have $\beta_C(\theta_1) \geq \beta_{C'}(\theta_1)$.*

Proof. Let $\theta_1 > \theta_0$ and let C' be a critical region with $\beta_{C'}(\theta_0) \leq \alpha$. Then, since the null hypothesis is simple, C' is also a test of level α for the simple hypothesis θ_0 against θ_1 and by assumption C satisfies the condition from theorem 22.2 for this problem. Thus, by the first statement of the theorem, we have $\beta_C(\theta_1) \geq \beta_{C'}(\theta_1)$ and since θ_1 was arbitrary this completes the proof. $\square$

A test which has this property is most powerful against every alternative in H_1 simultaneously and is said to be *uniformly most powerful* at level α . We only use this name as a description and do not consider the general theory of such tests.

For the normal mean, example 22.3 shows that the critical region $\{\bar{x} > \theta_0 + z_{1-\alpha}\sigma/\sqrt{n}\}$ is the Neyman–Pearson region for every $\theta_1 > \theta_0$, and thus the

one-sided test of example 20.6 is uniformly most powerful against $H_1: \theta > \theta_0$, with power function

$$\beta_+(\theta) = 1 - \Phi\left(z_{1-\alpha} - \frac{\sqrt{n}(\theta - \theta_0)}{\sigma}\right).$$

For alternatives $\theta_1 < \theta_0$ the same computation with the inequality reversed gives the region $\{\bar{x} < \theta_0 - z_{1-\alpha}\sigma/\sqrt{n}\}$, uniformly most powerful against $H_1: \theta < \theta_0$, with power function

$$\beta_-(\theta) = 1 - \Phi\left(z_{1-\alpha} + \frac{\sqrt{n}(\theta - \theta_0)}{\sigma}\right).$$

22.5 Two-Sided Alternatives

We now consider the alternative $H_1: \theta \neq \theta_0$. The two one-sided tests disagree: one rejects for large values of $\bar{X}$, the other for small values, and each one is poor on the other side of θ_0 , since for $\theta < \theta_0$ we have $\beta_+(\theta) < \alpha$. Thus, no single test can agree with both tests simultaneously, and the uniqueness statement in the lemma implies that this argument constitutes a proof.

Proposition 22.5. *Let $X_1, \dots, X_n$ be i.i.d. $N(\theta, \sigma^2)$ with known σ^2 and let $0 < \alpha < 1$. Then there is no critical region C^* with $\beta_{C^*}(\theta_0) \leq \alpha$ such that $\beta_{C^*}(\theta) \geq \beta_{C'}(\theta)$ for all $\theta \neq \theta_0$ and any critical region C' with $\beta_{C'}(\theta_0) \leq \alpha$.*

Proof. Suppose that C^* were such a region, and fix some $\theta_1 > \theta_0$. Then C^* is most powerful at level α for testing θ_0 against θ_1 , and so is the region $C_+ = \{\bar{x} > c_\alpha\}$ of example 22.3. The boundary set $\{\bar{x} = c_\alpha\}$ has probability zero, and thus by the second statement of theorem 22.2 the regions C^* and C_+ differ only on a set which has probability zero for every θ . Consequently the two tests have the same power function, $\beta_{C^*}(\theta) = \beta_+(\theta)$ for all θ . Now take any $\theta_2 < \theta_0$. From the formulas of section 22.4 we get

$$\beta_{C^*}(\theta_2) = \beta_+(\theta_2) < 1 - \Phi(z_{1-\alpha}) = \alpha < \beta_-(\theta_2),$$

since Φ is strictly increasing and $\sqrt{n}(\theta_2 - \theta_0)/\sigma < 0$. The left-sided test $C_- = \{\bar{x} < \theta_0 - z_{1-\alpha}\sigma/\sqrt{n}\}$ has size α and thus is an admissible competitor C' , and it has strictly larger power at θ_2 than C^* . This contradicts the assumed property of C^* , which completes the proof. $\square$

The argument is not particular to the normal distribution: whenever the one-sided most powerful regions disagree, a uniformly most powerful two-sided test would have to coincide with both, which is impossible, and so we have to settle for a compromise. The two-sided test in example 20.7 is such a compromise: the power of this test is less than $\beta_+(\theta)$ for $\theta > \theta_0$, but it is strictly larger than α for all $\theta \neq \theta_0$. In lecture 23 we will see that this test is the result of applying the likelihood principle to a composite alternative. The generalised likelihood-ratio test obtained in this way is a general-purpose tool which can be used when no uniformly most powerful test exists.

Summary.

- A test is most powerful at level α for a simple null hypothesis against a simple alternative, if it has level α and no other test of level α has larger power.
- The Neyman–Pearson lemma: The test which rejects when $f(x; \theta_1) > k f(x; \theta_0)$, where k is chosen so that the size is α , is most powerful at level α . Any other most powerful test coincides with this test outside the boundary set. The proof involves integration of $(\mathbf{1}_C - \mathbf{1}_{C'})(f(x; \theta_1) - k f(x; \theta_0)) \geq 0$.
- In practice, $R(x) > k$ can be rewritten as $T(x) > c$ for a simple statistic T and c can be found by considering the distribution of T under H_0 ; for the normal mean this gives $\bar{X} > \theta_0 + z_{1-\alpha}\sigma/\sqrt{n}$.
- If the region of acceptance does not depend on the alternative θ_1 , the same test will be most powerful against all $\theta_1 > \theta_0$: the test is then uniformly most powerful for the one-sided problem.
- For the two-sided alternative, no uniformly most powerful test can exist, since the two one-sided tests disagree; instead, the likelihood-ratio tests from lecture 23 can be used.

Exercise 22.1. Let $X_1, \dots, X_n$ be i.i.d. Poisson with mean $\theta > 0$ and let $S = \sum_{i=1}^n X_i$, Poisson with mean $n\theta$. We test $H_0: \theta = \theta_0$ against $H_1: \theta = \theta_1$ where $\theta_1 > \theta_0$.

1. Show that the likelihood ratio $R(x)$ is a strictly increasing function of $s = \sum_i x_i$ and that the most powerful test rejects when $S > c$. Why does c not depend on θ_1 ? What are the consequences for testing against the composite alternative $H_1: \theta > \theta_0$?
2. Let $n = 5$, $\theta_0 = 1$ and $\alpha = 0.05$. Determine the smallest integer c such that the test has size at most 0.05. Compute the actual size of the test for this value. You may use $\mathbb{P}(S \leq 8) = 0.9319$ and $\mathbb{P}(S \leq 9) = 0.9682$ for a Poisson random variable S with mean 5.
3. Determine the power of the test at $\theta_1 = 2$, using $\mathbb{P}(S \leq 9) = 0.4579$ for a Poisson random variable with mean 10.
4. Why cannot a test of the form $\{S > c\}$ have size exactly 0.05? At which level is the test from part (b) most powerful?

Exercise 22.2. Let $X_1, \dots, X_n$ be i.i.d. Bernoulli with success probability $p \in (0, 1)$, and let $S = \sum_{i=1}^n X_i$. In exercise 20.2 we tested $H_0: p = 1/2$ against $H_1: p > 1/2$ by rejecting for large values of S .

1. For a fixed alternative $p_1 > 1/2$, show that the likelihood ratio is given by

$$R(x) = \left(\frac{p_1}{1-p_1} \right)^s (2(1-p_1))^n, \quad s = \sum_{i=1}^n x_i,$$

and that it is strictly increasing in s . Deduce that the most powerful test rejects when $S > c$, and that the test of exercise 20.2 is uniformly most powerful against $H_1: p > 1/2$.

2. Let $n = 20$. Using $\mathbb{P}_{1/2}(S \geq 14) = 0.0577$ and $\mathbb{P}_{1/2}(S \geq 15) = 0.0207$, find the critical region of the largest size not exceeding 0.05, and compute its power at $p_1 = 0.7$ using $\mathbb{P}_{0.7}(S \leq 14) = 0.5836$. Compare the size and the power with those found for $n = 10$ in exercise 20.2.
3. For an alternative $p_1 < 1/2$, find the most powerful critical region, and explain why the two one-sided problems lead to different tests.

Exercise 22.3. Let $X_1, \dots, X_n$ be i.i.d. uniformly distributed on the set $(0, \theta)$, with joint density $f(x; \theta) = \theta^{-n} \mathbf{1}_{\{m \leq \theta\}}$ for $x_1, \dots, x_n \geq 0$, where $m = \max_i x_i$, and let $M = \max_i X_i$. From example 8.10 we know that $\mathbb{P}_\theta(M \leq y) = (y/\theta)^n$ for $0 \leq y \leq \theta$. We want to test $H_0: \theta = \theta_0$ against $H_1: \theta = \theta_1$, where $\theta_1 > \theta_0$, at significance level $\alpha \in (0, 1)$. In exercise 20.3 we have considered tests based on M for this model, and in the following we will see that these tests are optimal.

1. Let $k = (\theta_0/\theta_1)^n$. Determine the sets $\{x \mid f(x; \theta_1) > k f(x; \theta_0)\}$ and $\{x \mid f(x; \theta_1) = k f(x; \theta_0)\}$ in terms of m .
2. Show that the critical region $C = \{m > c\}$ for $c = \theta_0(1 - \alpha)^{1/n}$ has size α and satisfies the condition from theorem 22.2 for this k . Conclude that this is the most powerful test at level α . Compute the power of the test and explain why the test is uniformly most powerful against $H_1: \theta > \theta_0$.
3. Show that the region $C'' = \{m > \theta_0\} \cup \{m \leq \theta_0 \alpha^{1/n}\}$ also has size α and the same power as C . Explain, using the second statement of theorem 22.2, why this does not contradict the theorem, even though the model has densities.

Exercise 22.4. Let T be a sufficient statistic for the model $f(x; \theta)$ and let $\theta_0 \neq \theta_1$ be two parameter values. Using the factorisation theorem 8.2, show that the likelihood ratio $R(x) = f(x; \theta_1)/f(x; \theta_0)$ depends on the data only through $T(x)$, for all x where the denominator is positive.

Lecture 23

Likelihood-Ratio Tests

Lecture 22 ended with a gap: the Neyman–Pearson lemma does not consider composite null hypotheses or two-sided alternatives, but these are the most common situations in statistical tests. The purpose of this lecture is to describe a generalisation of the likelihood-ratio test, which can be used for such cases. The generalised likelihood-ratio test, introduced in this section, is based on the same principles as the Neyman–Pearson ratio, but is modified to account for the more general case. In the second part of this lecture, we will see that the large-sample behaviour of the generalised likelihood-ratio test, described by Wilks’ theorem, is based on the asymptotic theory from lecture 17. The generalised likelihood-ratio test is the test computed by most statistical software packages when a likelihood-ratio test is requested.

23.1 The Generalised Likelihood Ratio

As in lecture 22, we write $L(\theta)$ for the likelihood of the data $X = (X_1, \dots, X_n)$ and $\ell(\theta) = \log L(\theta)$ for the log-likelihood, where the parameter space Θ is now allowed to be a subset of $\mathbb{R}^k$. We consider the following testing problem:

$$H_0: \theta \in \Theta_0 \quad \text{against} \quad H_1: \theta \in \Theta \setminus \Theta_0$$

for a subset $\Theta_0 \subseteq \Theta$, where both hypotheses can be composite. Since the alternative is the complement of the null hypothesis within Θ , the null hypothesis is a restriction of the full model and we have *nested* hypotheses.

If the hypothesis is composite, the likelihood principle suggests representing it in the Neyman–Pearson comparison of the two likelihoods by the parameter value which gives the best fit to the data, *i.e.* by the maximum of the likelihood over the hypothesis.

Definition 23.1. The *generalised likelihood ratio* for testing $H_0: \theta \in \Theta_0$ against $H_1: \theta \in \Theta \setminus \Theta_0$ is given by

$$\Lambda(x) = \frac{\sup_{\theta \in \Theta_0} L(\theta)}{\sup_{\theta \in \Theta} L(\theta)}.$$

The *generalised likelihood-ratio test* at level α rejects H_0 , if $\Lambda(X) \leq c$, where the constant c is chosen such that $\sup_{\theta \in \Theta_0} \mathbb{P}_\theta(\Lambda(X) \leq c) \leq \alpha$. The test can equivalently be written as rejecting H_0 for large values of the statistic

$$W = -2 \log \Lambda(X).$$

The denominator of Λ is the likelihood at the maximum likelihood estimator $\hat{\theta}$ from lecture 5, the numerator is the likelihood at the *restricted* maximum likelihood estimator $\hat{\theta}_0$, which is the best value within H_0 ; if H_0 is simple, we have $\Theta_0 = \{\theta_0\}$ and the numerator is just $L(\theta_0)$. Since $\Theta_0 \subseteq \Theta$ we always have $0 \leq \Lambda \leq 1$: values close to one mean that the restriction to Θ_0 costs little in terms of fit, values close to zero mean that H_0 is a much worse fit for the data than the full model is, and this is the evidence on which we base the rejection. The factor -2 in W is the scaling which, as we will see in section 23.3, gives the limiting distribution of W its standard form. In terms of the log-likelihood we have

$$W = 2(\ell(\hat{\theta}) - \ell(\hat{\theta}_0)), \quad (23.1)$$

which is twice the amount the maximised log-likelihood decreases when the null hypothesis is imposed.

The test is an extension of the Neyman–Pearson test, not a replacement. For example, for a two-point parameter space $\Theta = \{\theta_0, \theta_1\}$ where $\Theta_0 = \{\theta_0\}$, and using the notation $R = L(\theta_1)/L(\theta_0)$ for the likelihood ratio from lecture 22, we have

$$\Lambda = \frac{L(\theta_0)}{\max(L(\theta_0), L(\theta_1))} = \min\left(1, \frac{1}{R}\right),$$

and for $c < 1$ the event $\Lambda \leq c$ is the same as the event $R \geq 1/c$; thus, for this case, the test is the most powerful test from theorem 22.2 with $k = 1/c$. For composite hypotheses no optimality is guaranteed and in fact proposition 22.5 shows that for the two-sided normal problem no such test can exist.

Two more properties follow directly from the definition. First, if T is a sufficient statistic, then we can use the factorisation theorem 8.2 to write $L(\theta) = g(T(x); \theta)h(x)$, where the factor $h(x)$ can be cancelled between numerator and denominator. This shows that Λ can be written as a function of the data which depends on the data only through T . Secondly, if $\psi = g(\theta)$ is a one-to-one reparametrisation, the suprema in the definition of Λ do not change and thus Λ and W do not depend on the choice of parametrisation of the model. This is the same invariance as in theorem 5.6, and exercise 23.3 verifies it in an example.

23.2 The Two-Sided Test for the Mean of a Normal Distribution

We will now consider the question which was left open at the end of lecture 22, *i.e.* we will consider the two-sided alternative for the mean of a normal sample with known variance. This example will carry over as an example for all the distributions considered later, since for this case the distribution of W under H_0 can be found exactly.

Example 23.2. Let $X_1, \dots, X_n$ be i.i.d. $N(\theta, \sigma^2)$ with known σ^2 and consider $H_0: \theta = \theta_0$ against the alternative $H_1: \theta \neq \theta_0$. Thus we have $\Theta = \mathbb{R}$ and $\Theta_0 = \{\theta_0\}$. From example 5.4 we know that the log-likelihood

$$\ell(\theta) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (X_i - \theta)^2$$

attains its maximum for $\mathbb{R}$ at $\hat{\theta} = \bar{X}$. Since the null hypothesis is simple, the restricted maximiser is θ_0 and using formula (23.1) we find

$$W = 2(\ell(\bar{X}) - \ell(\theta_0)) = \frac{1}{\sigma^2} \left(\sum_{i=1}^n (X_i - \theta_0)^2 - \sum_{i=1}^n (X_i - \bar{X})^2 \right).$$

Using the identity (2.1) from lecture 2, where θ_0 is substituted for μ , we find that the difference of the two sums of squares is $n(\bar{X} - \theta_0)^2$. Thus we have

$$W = \frac{n(\bar{X} - \theta_0)^2}{\sigma^2} = Z^2, \quad \text{where } Z = \frac{\sqrt{n}(\bar{X} - \theta_0)}{\sigma}. \quad (23.2)$$

The test rejects for large values of W , *i.e.* for large values of $|Z|$. Thus we have identified the two-sided z -test from example 20.7: the likelihood principle allowed us to derive the test which in lecture 22 was derived by common sense.

Under $\mathbb{P}_{\theta_0}$, the statistic Z is standard normally distributed by proposition A.3 and thus, by definition A.2, its square $W = Z^2$ satisfies $W \sim \chi_1^2$ for every sample size n . Let $\chi_{r,p}^2$ be the p -quantile of the χ_r^2 distribution. Then the level α test rejects if $W \geq \chi_{1,1-\alpha}^2$. Since $W = Z^2$, the events $W \geq \chi_{1,1-\alpha}^2$ and $|Z| \geq z_{1-\alpha/2}$ coincide and thus we have $\chi_{1,1-\alpha}^2 = z_{1-\alpha/2}^2$; from the tables we find $\chi_{1,0.95}^2 = 3.841 = 1.96^2$. For $n = 25$, $\sigma = 2$, $\theta_0 = 10$ and observed mean $\bar{x} = 10.9$, for the data from lecture 20, we find $w = 25 \cdot 0.81/4 = 5.06$, which is larger than 3.841 and thus H_0 is rejected at level 0.05. The p -value is $\mathbb{P}(\chi_1^2 \geq 5.06) = 2(1 - \Phi(\sqrt{5.06})) = 2(1 - \Phi(2.25)) = 0.0244$, which is the same value as obtained for the two-sided test in section 20.5.

A familiar test has been recovered from the general recipe, and the next section shows that the chi-squared distribution of W is the general case for large samples, even in models where normality is not exactly satisfied.

23.3 Wilks' Theorem

In example 23.2 the distribution of W under H_0 was known, because $\bar{X}$ is exactly normal; in most models this will not be the case and Wilks' theorem provides the large-sample approximation which allows the test to be used practically.

The theorem is stated for hypotheses which fix some coordinates of the parameter. Let $\Theta \subseteq \mathbb{R}^k$ be open, write $\theta = (\theta_1, \dots, \theta_k)$, and let the null hypothesis fix the first r coordinates at given values,

$$\Theta_0 = \{ \theta \in \Theta \mid \theta_j = \theta_j^0 \text{ for all } j \leq r \}, \quad (23.3)$$

where $1 \leq r \leq k$. For $r = k$ the null hypothesis is simple, and for $r < k$ the free coordinates $\theta_{r+1}, \dots, \theta_k$ are *nuisance parameters*, unknown parameters about which the hypothesis says nothing. The number r of constraints is the difference between the dimensions of Θ and Θ_0 .

Theorem 23.3 (Wilks). *Let $X_1, \dots, X_n$ be i.i.d. from a model with parameter $\theta \in \Theta \subseteq \mathbb{R}^k$ which satisfies the regularity conditions of definition 16.1, extended to a k -dimensional parameter as at the end of lecture 17, and suppose that, for every sample where all coordinates are in the support of the model, the maximum likelihood estimator is the unique root of the likelihood equations, both over Θ and*

over Θ_0 . Let Θ_0 be of the form (23.3) and let $W_n = -2 \log \Lambda_n$ be the generalised likelihood-ratio statistic for a sample of size n . Then for every $\theta \in \Theta_0$ we have

$$W_n \xrightarrow{d} \chi_r^2 \quad (n \rightarrow \infty)$$

under $\mathbb{P}_\theta$.

Whatever the model, the only feature of the problem which enters the limit is the number r of constraints. In example 23.2 we have $k = r = 1$ and there the statement holds exactly for every n . We do not prove the theorem: together with the vector case of theorem 17.1, on which its general form rests, it is one of the few results of this module which we state without proof. The following argument shows where the chi-squared distribution comes from.

Consider the case $k = r = 1$ of a one-dimensional parameter and a simple null hypothesis $H_0: \theta = \theta_0$, so that $W_n = 2(\ell(\hat{\theta}_n) - \ell(\theta_0))$. We will expand the log-likelihood around the maximum likelihood estimator (MLE) instead of around θ_0 as in lecture 17: using Taylor's theorem with the Lagrange form of the remainder we find

$$\ell(\theta_0) = \ell(\hat{\theta}_n) + \ell'(\hat{\theta}_n)(\theta_0 - \hat{\theta}_n) + \frac{1}{2} \ell''(\theta_n^*)(\theta_0 - \hat{\theta}_n)^2$$

for some θ_n^* between θ_0 and $\hat{\theta}_n$. Since $\hat{\theta}_n$ solves the likelihood equation, the first derivative at this point is zero and rearranging this gives

$$W_n = -\ell''(\theta_n^*)(\hat{\theta}_n - \theta_0)^2 = \left(-\frac{1}{n} \ell''(\theta_n^*)\right) \left(\sqrt{n}(\hat{\theta}_n - \theta_0)\right)^2. \quad (23.4)$$

Now assume that H_0 is true. Then, by theorem 17.1, the second term is the square of a quantity which converges in distribution to $N(0, 1/\mathcal{I}(\theta_0))$, and by the continuous mapping theorem, theorem A.17, this shows that the second term converges in distribution to $Y^2/\mathcal{I}(\theta_0)$ where Y is standard normally distributed. The first term is the term B_n from the proof of theorem 17.1 with θ_n^* instead of θ_0 and since θ_n^* is between θ_0 and the consistent estimator $\hat{\theta}_n$, the bound (R4) on the third derivative shows that this term still converges in probability to $\mathcal{I}(\theta_0)$. Finally, by Slutsky's theorem, theorem A.16, we find

$$W_n \xrightarrow{d} \mathcal{I}(\theta_0) \cdot \frac{Y^2}{\mathcal{I}(\theta_0)} = Y^2 \sim \chi_1^2,$$

by definition A.2. Since the Fisher information cancels, the limit is independent of the details of the model.

For a simple null hypothesis in one dimension, this argument is close to a complete proof. The proof for the general case is more challenging: the expansion will be in k variables and will use the vector version of theorem 17.1, and the quadratic form which replaces Y^2 will consist of r independent, squared standard normal variables, one for each constraint. This is the origin of the number r of degrees of freedom in the argument.

23.4 Using the Theorem

Wilks' theorem turns the generalised likelihood-ratio test into a recipe: compute the statistic $W = 2(\ell(\hat{\theta}) - \ell(\hat{\theta}_0))$ from the two maximised log-likelihoods, found

numerically if need be as in lecture 18; count the degrees of freedom $r = \dim \Theta - \dim \Theta_0$; reject H_0 at level α , if $W \geq \chi_{r,1-\alpha}^2$; and report the approximate p -value $\mathbb{P}(\chi_r^2 \geq w)$ for the observed value w , which for $r = 1$ equals $2(1 - \Phi(\sqrt{w}))$. Only the counting of degrees of freedom requires some care. For example, in the case of the mean of a normal sample with unknown variance, and $H_0: \mu = \mu_0$, we have $k = 2$ and $r = 1$, since the variance is unrestricted under both hypotheses. From exercise 23.2 we know that $r = m - 1$ for a multinomial distribution with m categories and a fully specified null hypothesis.

Remark. The hypothesis (23.3) fixes some of the coordinates, but since W is invariant under reparametrisation, the theorem applies to all smooth, nested hypotheses of dimension $k - r$. For example, the line $\{\mu_1 = \mu_2\}$ in the plane of two normal means can be transformed to $\theta_1 = \mu_1 - \mu_2 = 0$ with $\theta_2 = \mu_2$ free by a change of coordinates. Thus, the rule $r = \dim \Theta - \dim \Theta_0$ can be applied directly.

The theorem is a limit statement and there are three cases where problems can occur when applying the theorem to a fixed n . First, the approximation may be poor for small sample size: for the mean of a normal sample with unknown variance, exercise 23.1 shows that the level 0.05 test based on $\chi_{1,0.95}^2$ rejects too often at $n = 10$. Lecture 24 watches the approximation improve as n increases, by simulation. Whenever the exact distribution of W or of an equivalent test statistic is known, as in the situation of the exercise, this distribution should be used. Lecture 29 takes this approach, using the exact distribution, for the standard tests. Secondly, if the regularity conditions are violated, for example if the support of the distribution depends on the parameter or if the null value coincides with the boundary of the parameter space, the expansion (23.4) will break down and the limit will not be χ_r^2 . Finally, the theorem only describes W under H_0 . The theorem gives the size of the test, but the power of the test must be determined separately.

Software output often lists two relatives of the likelihood-ratio test: the *Wald test*, using $n\mathcal{I}(\hat{\theta})(\hat{\theta} - \theta_0)^2$ as in corollary 17.3, and the *score test*, using the squared score $\ell'(\theta_0)^2/(n\mathcal{I}(\theta_0))$. Both tests have the same chi-squared limit under H_0 as W , but here we use the likelihood-ratio test throughout.

Summary.

- For nested hypotheses $H_0: \theta \in \Theta_0$ against $H_1: \theta \in \Theta \setminus \Theta_0$, the generalised likelihood ratio is $\Lambda = \sup_{\Theta_0} L / \sup_{\Theta} L \in [0, 1]$, and the generalised likelihood-ratio test rejects for small Λ , *i.e.* for large $W = -2 \log \Lambda = 2(\ell(\hat{\theta}) - \ell(\hat{\theta}_0))$.
- For two simple hypotheses the test is the Neyman–Pearson test; in general it depends on the data only through a sufficient statistic and is unchanged by reparametrisation.
- For a normal mean with known variance and $H_0: \theta = \theta_0$, the statistic is $W = n(\bar{X} - \theta_0)^2 / \sigma^2 = Z^2$, which is exactly χ_1^2 under H_0 and the test is the two-sided z -test.
- Wilks' theorem: under regularity conditions, if H_0 fixes r of the k parameters, then $W \xrightarrow{d} \chi_r^2$ under H_0 . The theorem is given without proof, but from the arguments for $k = r = 1$, using Taylor expansion of ℓ around $\hat{\theta}_n$ and theorem 17.1, it is clear why the limit is χ_1^2 .
- In use: compute W from the two maximised log-likelihoods, count $r = \dim \Theta - \dim \Theta_0$, reject if $W \geq \chi_{r, 1-\alpha}^2$. The approximation can be poor for small n , the test fails if the regularity conditions are violated, and no information about power is available.

Exercise 23.1. Let $X_1, \dots, X_n$ be i.i.d. $N(\mu, \sigma^2)$, where both parameters are unknown. Consider the hypothesis $H_0: \mu = \mu_0$ against the alternative $H_1: \mu \neq \mu_0$, where σ^2 is a nuisance parameter. Let $Q = \sum_i (X_i - \bar{X})^2$, $S^2 = Q/(n-1)$ and $t = \sqrt{n}(\bar{X} - \mu_0)/S$.

1. Following example 5.4, show that the maximised log-likelihood function over Θ is given by

$$-\frac{n}{2} \log(2\pi Q/n) - \frac{n}{2}.$$

2. Show that, restricting ourselves to the region $\mu = \mu_0$, the maximum of the likelihood function is attained for $\hat{\sigma}_0^2 = \frac{1}{n} \sum_i (X_i - \mu_0)^2$ and find expressions for $\hat{\sigma}_0^2$ in terms of Q and $\bar{X} - \mu_0$ using the identity (2.1).

3. Deduce that

$$W = n \log \left(1 + \frac{t^2}{n-1} \right),$$

and that the generalised likelihood-ratio test rejects for large values of $|t|$.

4. State the approximate distribution of W under H_0 given by theorem 23.3 and justify your answer for the degrees of freedom. Recall that, under H_0 , the statistic t follows the t_{n-1} distribution by definition A.4 and corollary 10.10. For $n = 10$, determine the value of $|t|$ above which the approximate test with critical value $\chi_{1, 0.95}^2 = 3.841$ rejects and compare your result with the 0.975-quantile 2.262 of the t_9 distribution. Comment upon your result.

Exercise 23.2. An experiment with m possible outcomes is repeated n times independently. The number of times each outcome occurs is given by N_j for outcome j . Thus, $(N_1, \dots, N_m)$ is multinomially distributed with probability weights

$$\mathbb{P}_p(N_1 = n_1, \dots, N_m = n_m) = \frac{n!}{n_1! \cdots n_m!} p_1^{n_1} \cdots p_m^{n_m}, \quad n_1 + \cdots + n_m = n,$$

where $p = (p_1, \dots, p_m)$ satisfies $p_j > 0$ and $\sum_j p_j = 1$. We test $H_0: p = p^0$ for a given vector $p^0 = (p_1^0, \dots, p_m^0)$ against the alternative $p \neq p^0$.

1. Show that the log-likelihood function is $\sum_j n_j \log p_j$ up to an additive constant and use the Lagrange multiplier method, with constraint $\sum_j p_j = 1$, to show that the maximum of the likelihood function is at $\hat{p}_j = n_j/n$.
2. Show that

$$W = 2 \sum_{j=1}^m n_j \log \frac{n_j}{n p_j^0},$$

where $0 \log 0 = 0$.

3. Why does theorem 23.3 state that $W \xrightarrow{d} \chi_{m-1}^2$ under H_0 ?
4. A die is thrown 60 times and the six faces appear 13, 8, 11, 6, 12 and 10 times, respectively. Test the hypothesis that the die is fair at level 0.05, using $\chi_{5,0.95}^2 = 11.07$.
5. Let $n_j = n p_j^0 + d_j$ so that $\sum_j d_j = 0$. Show that, by expanding the logarithm to second order in $d_j/(n p_j^0)$, we find

$$W \approx \sum_{j=1}^m \frac{(n_j - n p_j^0)^2}{n p_j^0}.$$

This is Pearson's chi-squared statistic. Compute the value for the die in the previous part of the question.

Exercise 23.3. Let $X_1, \dots, X_n$ be i.i.d. with Bernoulli distribution, *i.e.* with success probability $p \in (0, 1)$. Furthermore, let $S = \sum_i X_i$ and $\hat{p} = S/n$ and consider the hypothesis $H_0: p = p_0$ against the alternative $H_1: p \neq p_0$.

1. Show that

$$W = 2 \left(S \log \frac{\hat{p}}{p_0} + (n - S) \log \frac{1 - \hat{p}}{1 - p_0} \right),$$

where we use the convention $0 \log 0 = 0$.

2. Approximately determine the distribution of W under H_0 for large n and use the resulting distribution to derive a test at level α .
3. In $n = 50$ trials, we observed 32 successes. Test the hypothesis $H_0: p = 1/2$ at level 0.05 and determine the approximate p -value.
4. In exercise 17.1, the model was reparametrised using the log-odds $\psi = \log(p/(1-p))$. Show that the generalised likelihood-ratio statistic for testing $H_0: \psi = \psi_0$ coincides with the one in part (a) for $p_0 = e^{\psi_0}/(1 + e^{\psi_0})$, the value of the null proportion.

Lecture 24

Power and Size by Simulation

This is the R session for week 8. The lectures this week have been about the Neyman–Pearson lemma and the likelihood-ratio tests. In this session we will see how these tests perform in practice. The R technique we will use here is the indicator average from lecture 12, where the indicator now stands for the event of being rejected. The written reference for this is in appendix B.

24.1 The Size of a Test by Simulation

The test we consider today is the one-sided z -test from example 20.6, in the numerical setting of the example: we have $n = 25$ observations from $N(\theta, \sigma^2)$ with known $\sigma = 2$ and we want to test $H_0: \theta = 10$ against $H_1: \theta > 10$ at level $\alpha = 0.05$, *i.e.* we reject if and only if $\bar{X} > 10.658$. By definition 20.4, the power function $\beta(\theta) = \mathbb{P}_\theta(\bar{X} > 10.658)$ is a probability, and by theorem A.14 we can approximate this probability by the proportion of samples where the test rejects. As in lecture 12, we write a function which, for one run of the test, returns TRUE if the test rejects and FALSE otherwise. We can then use `mean` to get the proportion of rejections. Before you try the code, predict the value of `mean(rejected)`. Will it be exactly 0.05? If not, how far away from 0.05 would it have to be before you would suspect that there might be a mistake in the code?

```
n      <- 25
sigma  <- 2
theta0 <- 10
crit   <- theta0 + qnorm(0.95) * sigma / sqrt(n)
crit
z_reject <- function(theta) {
  x <- rnorm(n, mean = theta, sd = sigma)
  mean(x) > crit
}
set.seed(2026)
rejected <- replicate(10000, z_reject(theta0))
head(rejected)
mean(rejected)
sqrt(0.05 * 0.95 / 10000)
```

```
## [1] 10.65794
```

```
## [1] FALSE FALSE FALSE FALSE FALSE
## [1] 0.0523
## [1] 0.002179449
```

The rejection rate for H_0 is 0.0523. This is the size of the test in the sense of definition 20.5. The value is not exactly 0.05, and it should not be: as in lecture 12, the estimate has standard deviation $\sqrt{0.05 \cdot 0.95/10000} = 0.0022$. The last line of the output shows this value. A value between 0.046 and 0.054 is what we expect for a correct test. If we get 0.07, we should suspect that there is an error in the code. We can use the same function to get the rejection rate for any other θ . Predict the three numbers which follow, for the first value of θ which is less than θ_0 , and the other two values are the alternatives considered in example 20.6.

```
mean(replicate(10000, z_reject(9.5)))
mean(replicate(10000, z_reject(10.5)))
mean(replicate(10000, z_reject(11)))
```

```
## [1] 0.0016
## [1] 0.344
## [1] 0.7953
```

At $\theta = 9.5$ the test rejects 0.2% of samples: for the composite null hypothesis $H_0: \theta \leq 10$ the size is attained at the boundary point $\theta_0 = 10$, as remarked after example 20.6. The values 0.344 and 0.795 agree with the exact powers 0.346 and 0.804 from that example to within simulation error.

24.2 The Power Curve

A single number is a poor summary of a test, so we consider the power function of a grid of alternatives. We compare the one-sided z -test to the two-sided z -test from example 20.7, and the one-sided z -test for the first 12 observations. The final test is wasteful, but a legal level 0.05 test. What is the cost of the waste? The following box contains the full script, with lines in random order (except for the function body). Before you read the rest of this section, try to arrange these lines into a working script.

```
colnames(power) <- thetas
decide <- function(theta) {
  x      <- rnorm(n, mean = theta, sd = sigma)
  z      <- sqrt(n) * (mean(x) - theta0) / sigma
  z_half <- sqrt(12) * (mean(x[1:12]) - theta0) / sigma
  c(one = z > qnorm(0.95), two = abs(z) > qnorm(0.975),
    half = z_half > qnorm(0.95))
}
set.seed(2026)
theta0 <- 10
round(power[, thetas %in% c(9.5, 10, 10.5, 11, 11.5)], 3)
power <- sapply(thetas, function(theta) rowMeans(replicate(5000, decide(
  theta))))
thetas <- seq(9, 12, by = 0.1)
n      <- 25
sigma  <- 2
```

The rule is as in lectures 9 and 12: all names must be defined before they are used, and the seed must be set before the first random draw. The new function returns a named logical vector of length three, giving the decisions for the three tests. The command `replicate` then returns a three-row matrix, and `rowMeans` returns the three rejection rates. The command `sapply` can then be used to apply this recipe to every value of the grid. Before you try the reconstructed script, predict the row for the two-sided test. Is the rejection rate of the two-sided test at $\theta = 9.5$ above or below the rejection rate of the one-sided test? Is the rejection rate at $\theta = 11$ above or below? Predict the position of the half-sample test at $\theta = 11$, closer to 0.8 or to 0.5?

```
n      <- 25
sigma  <- 2
theta0 <- 10
decide <- function(theta) {
  x      <- rnorm(n, mean = theta, sd = sigma)
  z      <- sqrt(n) * (mean(x) - theta0) / sigma
  z_half <- sqrt(12) * (mean(x[1:12]) - theta0) / sigma
  c(one = z > qnorm(0.95), two = abs(z) > qnorm(0.975),
    half = z_half > qnorm(0.95))
}
thetas <- seq(9, 12, by = 0.1)
set.seed(2026)
power <- sapply(thetas, function(theta) rowMeans(replicate(5000, decide(
  theta))))
colnames(power) <- thetas
round(power[, thetas %in% c(9.5, 10, 10.5, 11, 11.5)], 3)
```

```
##      9.5    10   10.5    11   11.5
## one  0.002 0.052 0.359 0.803 0.982
## two  0.227 0.051 0.248 0.701 0.963
## half 0.007 0.053 0.230 0.526 0.827
```

Figure 24.1, produced by the script `R/S24-power.R`, plots the data for the three rows and the grid. The exact power function (20.2) of the one-sided test is given by the line. All three curves start at 0.05 for $\theta_0 = 10$ as required for level 0.05 tests, and the filled circles are on the exact curve at this point. The rest of the plot shows the power of the tests. The two-sided test is always smaller than the one-sided test for $\theta > 10$ and is much larger than the one-sided test for $\theta < 10$. Thus, neither curve is always larger than the other curve, and this illustrates proposition 22.5. From example 22.3 and proposition 22.4 we know that no level 0.05 test based on these 25 observations can reject $\theta = 11$ more than 80% of the time, and the crosses indicate what power a poor test would have for this bound: the power at $\theta = 11$ decreases from 0.80 to 0.53, and this is the cost of having discarded thirteen of the observations.

24.3 Wilks' Theorem in Action

Assume now that the variance is not known. In this case we cannot perform the z -test and we have two parameters, where $H_0: \theta = 10$ is fixed. From exercise 23.1




Figure 24.1: Estimated power functions for three level 0.05 tests for $H_0: \theta = 10$, using 5000 samples of size $n = 25$ for each grid point: one-sided z -test (filled circles), two-sided z -test (open circles) and one-sided z -test on the first 12 observations only (crosses). For comparison, the line gives the exact power function of the one-sided test from lecture 20.

we know that the likelihood-ratio statistic from lecture 23 is given by

$$W = -2 \log \Lambda = n \log \left(1 + \frac{t^2}{n-1} \right), \quad t = \frac{\sqrt{n}(\bar{X} - \theta_0)}{S},$$

where S is the sample standard deviation, and by Wilks' theorem (theorem 23.3) the distribution of W under H_0 converges to a chi-squared distribution with one degree of freedom as $n \rightarrow \infty$. In contrast, in the case with known variance from lecture 23, the chi-squared distribution was exact. Here, the approximation will not be exact and we need to determine how well the approximation holds for the n we have. We can simulate W under H_0 for $n = 10$ and $n = 100$, using $\sigma = 2$ again, and we can reject the hypothesis if W is greater than the critical value 3.841 of the chi-squared distribution, as if the theorem were exact. Before you try the code, predict the two rejection rates. Will the rate for $n = 10$ be above or below 0.05? Note that the value S in the denominator of t is variable, since n can be small.

```
theta0 <- 10
sigma <- 2
wilks <- function(n) {
  x <- rnorm(n, mean = theta0, sd = sigma)
  t <- sqrt(n) * (mean(x) - theta0) / sd(x)
  n * log(1 + t^2 / (n - 1))
}
set.seed(2026)
w10 <- replicate(10000, wilks(10))
w100 <- replicate(10000, wilks(100))
crit <- qchisq(0.95, df = 1)
crit
c(n10 = mean(w10 > crit), n100 = mean(w100 > crit))
```

```
## [1] 3.841459
##      n10    n100
## 0.0745 0.0516
```

The rejection rate for $n = 100$ is 0.052, which is within simulation error of the nominal level. The rejection rate for $n = 10$ is 0.075: the test with nominal level 0.05 has type I errors half again as often as it should, since S is a noisy estimator for σ using only ten observations. The top row of figure 24.2 shows the two histograms, together with the density of the chi-squared distribution. The plot below shows the sorted values of W together with the quantiles of the chi-squared distribution. This is similar to the plot we used in lecture 12 using `qqnorm`. The command `qchisq` with `ppoints(10000)` can be used to get the quantiles. Before you look at figure 24.2, predict the shape of the quantile plot for $n = 10$: is the quantile plot on the line, or is it curved above or below the line?




Figure 24.2: The likelihood-ratio statistic $W = n \log(1 + t^2/(n - 1))$ simulated 10000 times under H_0 for $n = 10$ (left) and $n = 100$ (right). Top: histograms with the density of the chi-squared distribution with one degree of freedom (truncated at 8). Bottom: the sorted values of the test statistic against the quantiles of the chi-squared distribution with one degree of freedom, with the identity line.

The histograms do not do a good job of showing the difference in sample size, because the density of the chi-squared distribution is unbounded at zero and because most of the mass of the distribution is concentrated in the first few bars. In contrast, the quantile plots show better: for $n = 10$ the points start deviating from the line (above) at around 2, indicating that the upper tail of W is heavier than expected from the limit. For $n = 100$ the points stay on the line throughout the plot. The solution to the problem of small n is to use the exact distribution: under H_0 the test statistic t follows the t -distribution with $n - 1$ degrees of freedom, and we can reject if $|t|$ exceeds the 0.975-quantile of this distribution. This is the one-sample t -test from lecture 29. Since W is increasing as a function of t^2 , this is equivalent to rejecting if W exceeds the value $n \log(1 + t_{0.975, n-1}^2/(n - 1))$.

```

ns <- c(10, 100)
exact <- ns * log(1 + qt(0.975, df = ns - 1)^2 / (ns - 1))
exact
c(n10 = mean(w10 > exact[1]), n100 = mean(w100 > exact[2]))

```

```

## [1] 4.501803 3.899844
##      n10      n100
## 0.0532 0.0501

```

Using the exact threshold 4.50 instead of 3.84, the rejection rate for $n = 10$ is back to 0.05 and for $n = 100$ the two thresholds are close enough that it does not matter which one is used. An asymptotic level can only be guaranteed in the limit; otherwise, an exact test should be used if available, and if not, a simulation under H_0 can be used to determine the actual size at a given n .

24.4 A Task to Take Home

The following task is in the style of a practical report. Set the seed of the random number generator to your student ID and consider the setting from section 24.2, but with sample size $n = 10$ instead of $n = 25$. Estimate the size of the one-sided z -test for 10000 samples and state whether the estimated size is consistent with 0.05, within simulation error. Next, estimate the power of the test on the grid `seq(9, 13, by = 0.1)` and create a plot of the estimated power curve, together with the exact power function from lecture 20. From the plot, determine the smallest θ at which the power is greater than 0.9. Finally, consider the test which rejects for large values of the sample *median*. The critical value must be found first: simulate the median of 10 observations under $\theta_0 = 10$ and use the 0.95-quantile of the simulated values to get a test with size 0.05. (Using the critical value of the z -test would lead to a much larger size and comparing the power of two tests with different size is not meaningful.) Estimate the power curve of the calibrated median test and add it to your plot. Write a short paragraph comparing the two power curves, and use example 22.3 to explain why the median test cannot win.

All of this session concerns a single test at a time. The question of what happens when many tests are performed on the same data set and when only the significant tests are reported will be discussed in lecture 30.

Before the next session, type the commands of this session into R, and check that what you get is what the text says you should get.

Summary.

- The written reference for the commands used in this session is appendix [B](#).
- We can estimate the size and power of a test by simulation, by taking the average of the indicator function for the rejection. An estimated rate for 10000 samples has standard deviation of approximately 0.002 near 0.05.
- A power curve can be generated by applying the rejection-rate recipe to a grid of alternative values, using `sapply`. For the one-sided z -test, the estimated points lie on the exact curve from lecture [20](#).
- For every level 0.05 test, the point 0.05 for θ_0 is passed through. The tests differ in power only. The two-sided test is weaker than the one-sided test on one side and stronger on the other, and a test which excludes observations is weaker for all alternatives, as shown in lecture [22](#).
- For the normal mean with unknown variance, the likelihood-ratio statistic is only approximately chi-squared: for $n = 10$ the chi-squared critical value corresponds to size 0.075 instead of 0.05, and for $n = 100$ the approximation is good. If an exact test is available, this should be used. If not, simulation under H_0 should be used.

Lecture 25

Confidence Intervals

So far we have considered estimators which output a single number. In this lecture we will consider intervals of parameter values, constructed such that the true parameter value lies inside the interval with a given probability. We will see two different ways to construct such an interval, one using the concept of a pivotal quantity and the other one based on the inversion of a family of hypothesis tests. This latter approach also allows us to relate the testing theory from lectures 20 to 23 to the topic of estimation. The Bayesian equivalents of confidence intervals are called credible intervals and we will learn about these in lecture 28. In lecture 27 we will compare the two approaches to confidence intervals using simulation.

25.1 Definition and Interpretation

Here, $X = (X_1, \dots, X_n)$ is a sample from a parametric model with parameter $\theta \in \Theta \subseteq \mathbb{R}$, and $\mathbb{P}_\theta$ is the distribution of the sample under θ . For a fixed small $\alpha \in (0, 1)$, typically $\alpha = 0.05$, the aim is to find an interval which misses the truth with probability at most α .

Definition 25.1. Let $L(X)$ and $U(X)$ be statistics such that $L(X) \leq U(X)$. The random interval $[L(X), U(X)]$ is a *confidence interval* for θ with *confidence level* $1 - \alpha$, if

$$\mathbb{P}_\theta(L(X) \leq \theta \leq U(X)) \geq 1 - \alpha \quad \text{for all } \theta \in \Theta.$$

The probability on the left-hand side is called the *coverage probability* of the interval at θ .

The endpoints of the interval are statistics, *i.e.* functions of the data, while the parameter θ is an unknown constant which does not change between samples. For continuous data we can usually achieve equality in the definition, but for discrete data the coverage probability of the interval can jump as θ changes. For this reason, the definition only requires the left-hand side to satisfy the inequality. More generally, any random subset $C(X) \subseteq \Theta$ of the parameter space such that $\mathbb{P}_\theta(\theta \in C(X)) \geq 1 - \alpha$ for all θ is called a *confidence set*.

Some care is needed when interpreting a confidence interval. For example, if the data lead to the 95% interval $[9.32, 11.28]$, both endpoints are now fixed numbers, and so is the unknown θ ; the statement $9.32 \leq \theta \leq 11.28$ is either

true or false. The 95% describes the procedure: 95% of the computed intervals would cover the truth in repeated samples, but we do not know whether the interval in front of us is one of these. The statement “ θ is in $[9.32, 11.28]$ with probability 0.95” is wrong, since it assumes that the parameter is random (this only holds in the Bayesian model from lecture 28). The statement “95% of the observations are in the interval” is wrong, since this would be a confusion of parameter with data. The correct interpretation of such an interval can be illustrated by simulating many intervals and seeing how many of these cover the truth. This is done in lecture 27.

25.2 Pivotal Quantities

The most direct approach to constructing a confidence interval is to construct an interval based on a quantity whose distribution is known exactly, regardless of the parameter value. Such a quantity is called a *pivotal quantity*.

Definition 25.2. A function $Q(X, \theta)$ of the sample and the parameter is a *pivotal quantity* or *pivot*, if the distribution of $Q(X, \theta)$ under $\mathbb{P}_\theta$ does not depend on $\theta \in \Theta$.

A pivot is not a statistic, because it contains the unknown θ , but we know the quantiles of the pivot distribution. To get a confidence interval, we can choose constants $a < b$ such that $\mathbb{P}_\theta(a \leq Q(X, \theta) \leq b) = 1 - \alpha$ and then solve the inequalities $a \leq Q(X, \theta) \leq b$ for θ . If Q is monotone in θ , the solution of these inequalities forms an interval which, with probability exactly $1 - \alpha$, contains the unknown parameter value. Using the quantile notation z_p from lecture 20, we have $z_{0.975} = 1.96$ and, using symmetry, $z_{1-p} = -z_p$. Similarly, we write $t_{k,p}$ and $\chi_{k,p}^2$ for the p -quantiles of the t_k and the χ_k^2 distribution, respectively. These quantiles are all tabulated.

Example 25.3. Let $X_1, \dots, X_n$ be i.i.d. $N(\theta, \sigma^2)$ where σ^2 is known. Then, by proposition A.3, $\bar{X} \sim N(\theta, \sigma^2/n)$ and thus

$$Q(X, \theta) = \frac{\sqrt{n}(\bar{X} - \theta)}{\sigma} \sim N(0, 1)$$

for all θ . Thus, Q is a pivot. We can find a confidence interval for θ by bracketing Q between $-z_{1-\alpha/2}$ and $z_{1-\alpha/2}$ and then solving the resulting inequalities for θ . The resulting confidence interval is

$$\left[\bar{X} - z_{1-\alpha/2} \frac{\sigma}{\sqrt{n}}, \bar{X} + z_{1-\alpha/2} \frac{\sigma}{\sqrt{n}} \right]$$

with level $1 - \alpha$. For $n = 16$ observations, with $\sigma = 2$ and $\bar{x} = 10.3$, the 95% interval is $10.3 \pm 1.96 \cdot 0.5 = [9.32, 11.28]$. This is the same interval as we discussed above.

If σ^2 is unknown, this pivot is not available. If we replace σ by the sample standard deviation S , we get a new pivot and we have to consider the t distribution from appendix A.

Example 25.4. Let $X_1, \dots, X_n$ be i.i.d. $N(\mu, \sigma^2)$, where both parameters are unknown, and let S^2 be the sample variance. By proposition A.3, the two values

$$Z = \frac{\sqrt{n}(\bar{X} - \mu)}{\sigma} \sim N(0, 1) \quad \text{and} \quad V = \frac{(n-1)S^2}{\sigma^2} \sim \chi_{n-1}^2$$

are independent. By definition A.4, the ratio

$$T = \frac{Z}{\sqrt{V/(n-1)}} = \frac{\sqrt{n}(\bar{X} - \mu)}{S}$$

is distributed as t_{n-1} and, since the unknown σ cancels, T is a pivot for μ , independently of the nuisance parameter σ^2 . The t -interval

$$\left[\bar{X} - t_{n-1, 1-\alpha/2} \frac{S}{\sqrt{n}}, \bar{X} + t_{n-1, 1-\alpha/2} \frac{S}{\sqrt{n}} \right],$$

with level $1 - \alpha$, can be obtained by bracketing the pivot as above. For the data from example 25.3, with sample standard deviation $s = 2.1$, we can use the table to find $t_{15, 0.975} = 2.131$ and the 95% interval $10.3 \pm 2.131 \cdot 2.1/4 = [9.18, 11.42]$. The wider interval is the price we pay for not knowing σ . As n increases, this difference decreases, since t_{n-1} converges to the standard normal distribution.

Example 25.5. In the same setting, $V = (n-1)S^2/\sigma^2 \sim \chi_{n-1}^2$ is a pivot for σ^2 , independently of the unknown μ . If we bracket this value between $\chi_{n-1, \alpha/2}^2$ and $\chi_{n-1, 1-\alpha/2}^2$ and then solve for σ^2 , the order of the endpoints is reversed, and we obtain the confidence interval

$$\left[\frac{(n-1)S^2}{\chi_{n-1, 1-\alpha/2}^2}, \frac{(n-1)S^2}{\chi_{n-1, \alpha/2}^2} \right]$$

for σ^2 with level $1 - \alpha$. For $n = 16$ and $s^2 = 4.41$ the table gives $\chi_{15, 0.025}^2 = 6.262$ and $\chi_{15, 0.975}^2 = 27.488$, and thus the 95% interval is given by

$$[66.15/27.488, 66.15/6.262] = [2.41, 10.56].$$

This interval is far from symmetric around the estimate 4.41.

If we take the square root of both endpoints, the interval $[1.55, 3.25]$ for σ has the same coverage probability: applying a strictly increasing function to both endpoints does not change the coverage probability, whereas applying a decreasing function swaps the endpoints.

25.3 Inverting a Test

A level- α test of $H_0: \theta = \theta_0$ decides whether a single value θ_0 is compatible with the data; a confidence interval collects all the values which are. The following theorem, using the notions of critical region and level from definitions 20.2 and 20.5, makes this correspondence exact, in both directions.

Theorem 25.6. For every $\theta_0 \in \Theta$, let $A(\theta_0)$ be the acceptance region of a level- α test of $H_0: \theta = \theta_0$, that is the set of samples for which H_0 is not rejected. Then

$$C(X) = \{\theta_0 \in \Theta \mid X \in A(\theta_0)\},$$

is a confidence set for θ with level $1 - \alpha$. Conversely, if $C(X)$ is a confidence set with level $1 - \alpha$, then for each $\theta_0 \in \Theta$ the test which rejects $H_0: \theta = \theta_0$ if and only if $\theta_0 \notin C(X)$ has level α .

Proof. By the definition of $C(X)$, for every θ_0 and every sample we have $\theta_0 \in C(X)$ if and only if $X \in A(\theta_0)$. Let θ be the true value. Then, for $\theta_0 = \theta$, under $\mathbb{P}_\theta$ the null hypothesis $H_0: \theta_0 = \theta$ is true and the test with acceptance region $A(\theta)$ rejects with probability at most α . We have

$$\mathbb{P}_\theta(\theta \in C(X)) = \mathbb{P}_\theta(X \in A(\theta)) = 1 - \mathbb{P}_\theta(X \notin A(\theta)) \geq 1 - \alpha.$$

Since θ was arbitrary, $C(X)$ is a confidence set with level $1 - \alpha$. For the converse statement, the test for $H_0: \theta = \theta_0$ rejects if and only if $\theta_0 \notin C(X)$. Thus, under $\mathbb{P}_{\theta_0}$, this event has probability $1 - \mathbb{P}_{\theta_0}(\theta_0 \in C(X)) \leq \alpha$ and the test has level α . This completes the proof. $\square$

The theorem is a change of viewpoint, rather than a new result: the event “ θ_0 is not rejected by X ” can be read both as a test (with θ_0 fixed and X random) and as a confidence set (with X fixed and θ_0 varying). As an illustration, the size- α two-sided test from example 20.7 for the mean of an i.i.d. $N(\theta, \sigma^2)$ sample with known σ^2 rejects $H_0: \theta = \theta_0$ if $|\bar{X} - \theta_0| > z_{1-\alpha/2} \sigma / \sqrt{n}$ and the values not rejected by the test form the interval from example 25.3. Similarly, by inverting the one-sided test from example 20.6, we find a one-sided confidence interval. We will consider this type of interval again in section 25.5. Finally, in exercise 25.4, the same approach is repeated with unknown σ . In practice, the duality is most often used in the opposite direction: a value θ_0 is rejected at level α if it does not belong to the $1 - \alpha$ interval.

25.4 Approximate Intervals from the MLE

Exact pivots only exist for special models. For every regular model, however, corollary 17.3 provides an approximate pivot: the standardised error $(\hat{\theta}_n - \theta) / \text{se}(\hat{\theta}_n)$ of the maximum likelihood estimator, with $\text{se}(\hat{\theta}_n) = 1 / \sqrt{\mathcal{I}_n(\hat{\theta}_n)}$, converges in distribution to $N(0, 1)$ whatever the value of θ .

Proposition 25.7. *With the same assumptions as in corollary 17.3, the Wald interval*

$$[\hat{\theta}_n - z_{1-\alpha/2} \text{se}(\hat{\theta}_n), \hat{\theta}_n + z_{1-\alpha/2} \text{se}(\hat{\theta}_n)]$$

has coverage probability that converges to $1 - \alpha$ as $n \rightarrow \infty$, for all $\theta \in \Theta$.

Proof. Let $Z_n = (\hat{\theta}_n - \theta) / \text{se}(\hat{\theta}_n)$ and $z = z_{1-\alpha/2}$. The event that the interval contains θ is the event $|\hat{\theta}_n - \theta| \leq z \text{se}(\hat{\theta}_n)$. This event is the same as $-z \leq Z_n \leq z$. From corollary 17.3 we know that $Z_n \xrightarrow{d} Z$ with $Z \sim N(0, 1)$. Writing F_n for the distribution function of Z_n , for every $\varepsilon > 0$ we have

$$\begin{aligned} \mathbb{P}_\theta(-z \leq Z_n \leq z) &\geq \mathbb{P}_\theta(-z < Z_n \leq z) = F_n(z) - F_n(-z), \\ \mathbb{P}_\theta(-z \leq Z_n \leq z) &\leq \mathbb{P}_\theta(-z - \varepsilon < Z_n \leq z) = F_n(z) - F_n(-z - \varepsilon). \end{aligned}$$

Since Φ is continuous everywhere, definition A.13 gives $F_n(t) \rightarrow \Phi(t)$ for every t , and thus the lower bound converges to $\Phi(z) - \Phi(-z) = 1 - \alpha$ and the upper bound converges to $\Phi(z) - \Phi(-z - \varepsilon)$ as $n \rightarrow \infty$. Letting $\varepsilon \downarrow 0$ and using the continuity of Φ again, we find that $\mathbb{P}_\theta(-z \leq Z_n \leq z) \rightarrow 1 - \alpha$ as $n \rightarrow \infty$ and thus the coverage probability converges to $1 - \alpha$. This is the required result. $\square$

The Wald interval, for 95%, is the interval $\hat{\theta}_n \pm 1.96 \operatorname{se}(\hat{\theta}_n)$. This is the interval introduced in lectures 12 and 18. Using the delta method, theorem A.18, the same approach can be used to derive the interval $g(\hat{\theta}_n) \pm z_{1-\alpha/2} |g'(\hat{\theta}_n)| \operatorname{se}(\hat{\theta}_n)$ for any smooth function $g(\theta)$ (see exercise 25.3). However, the coverage of this interval is only guaranteed in the limit, and section 17.5 lists the cases to be aware of.

Example 25.8. Let $X_1, \dots, X_n$ be i.i.d. Bernoulli with success probability $\theta \in (0, 1)$. The MLE is $\hat{\theta}_n = \bar{X}$ by example 5.3 and the Fisher information is $\mathcal{I}(\theta) = 1/(\theta(1 - \theta))$ by example 11.5. Thus the Wald interval is $\bar{X} \pm z_{1-\alpha/2} \sqrt{\bar{X}(1 - \bar{X})/n}$. For 37 successes in $n = 100$ trials we have $\bar{x} = 0.37$, the standard error is $\sqrt{0.37 \cdot 0.63/100} = 0.048$, and the 95% interval is $0.37 \pm 1.96 \cdot 0.048 = [0.275, 0.465]$. If θ is close to 0 or 1, or if n is small, the coverage of this interval may be less than $1 - \alpha$, since the normal approximation to $\bar{X}$ is poor in these cases and the standard error equals zero if all observations are the same.

25.5 One-Sided Intervals and the Choice of Interval

Sometimes it is only important to consider one direction of a confidence interval, for example when we want to show that a failure rate is less than or equal to some value. A *lower confidence bound* for θ with level $1 - \alpha$ is a statistic $L(X)$ such that $\mathbb{P}_\theta(L(X) \leq \theta) \geq 1 - \alpha$ for all θ . The corresponding *one-sided confidence interval* is then $[L(X), \infty)$. An upper confidence bound is defined symmetrically. In the pivot method, a one-sided interval can be obtained by considering the single inequality $Q(X, \theta) \leq b$, where b is the $(1 - \alpha)$ -quantile of the pivot. For example, for the normal distribution with known variance, this gives the lower bound $\bar{X} - z_{1-\alpha} \sigma/\sqrt{n}$ as promised in section 25.3, where $z_{1-\alpha}$ is used instead of $z_{1-\alpha/2}$ since all of the error probability is allocated to one side. In the context of the pivot method, any pair of values $a < b$ with $\mathbb{P}(a \leq Q \leq b) = 1 - \alpha$ defines a valid two-sided interval. Here we use the *equal-tailed* choice, i.e. a and b are the $(\alpha/2)$ - and $(1 - \alpha/2)$ -quantiles, respectively. This choice is the shortest interval for a symmetric, unimodal pivot density, but may not be shortest for skewed densities like the chi-squared distribution; the interval boundaries can be found in the tables. Exercise 25.2 shows an example where the shortest interval is one-sided. In lecture 28 we will see the same choice of interval again, when we consider credible intervals.

Summary.

- A confidence interval covers the true parameter with probability at least $1 - \alpha$ for all θ . The interval is random, the parameter is fixed, and the level describes the procedure, not the individual interval.
- An exact interval can be obtained by bracketing the pivot between two quantiles and solving for θ . For normal samples this can be used to derive the z -, t - and chi-squared intervals.
- Tests and confidence sets are dual: the values θ_0 not rejected by a family of level- α tests form a $1 - \alpha$ confidence set, and conversely.
- The Wald interval $\hat{\theta}_n \pm z_{1-\alpha/2} \text{se}(\hat{\theta}_n)$ asymptotically covers $1 - \alpha$ for all regular models, but the finite-sample coverage may be lower.
- One-sided intervals use $z_{1-\alpha}$ instead of $z_{1-\alpha/2}$. For two-sided intervals the equal-tailed choice is the norm.

Exercise 25.1. Let $X_1, \dots, X_n$ be i.i.d. exponential with rate $\theta > 0$ and let $T = \sum_{i=1}^n X_i$. In exercise 13.1 we have seen that $2\theta T \sim \chi_{2n}^2$.

1. Explain why $Q = 2\theta T$ is a pivot for θ and derive the equal-tailed confidence interval for θ with level $1 - \alpha$.
2. For $n = 10$ observations $\sum_i x_i = 25.0$, compute the 95% interval, using the values $\chi_{20,0.025}^2 = 9.591$ and $\chi_{20,0.975}^2 = 34.170$.
3. The MLE for the rate is $\hat{\theta} = 1/\bar{X}$ and the standard error is $\hat{\theta}/\sqrt{n}$, as shown in example 17.4. Compute the Wald interval from proposition 25.7 for the given data and compare the two intervals.

Exercise 25.2. Let $X_1, \dots, X_n$ be i.i.d. uniformly distributed on the set $[0, \theta]$ and let $M = \max_i X_i$.

1. Using the density of M from example 8.10, show that $\mathbb{P}_\theta(M/\theta \leq u) = u^n$ for $0 \leq u \leq 1$ and conclude that M/θ is a pivot.
2. Show that $[M, M/\alpha^{1/n}]$ is a confidence interval for θ with level $1 - \alpha$ and compute the interval for $n = 5$, $\max_i x_i = 3.2$ and $\alpha = 0.05$.
3. Among all intervals $[M/b, M/a]$ with $0 < a < b \leq 1$ and $\mathbb{P}_\theta(a \leq M/\theta \leq b) = 1 - \alpha$, show that the interval from the previous part is the shortest. (Hint: think about how the length of the interval changes when b is moved and a is adjusted to maintain the same coverage.)

Exercise 25.3. Let $X_1, \dots, X_n$ be i.i.d. Poisson with mean $\theta > 0$. Then the MLE is $\hat{\theta}_n = \bar{X}$ and the Fisher information is $\mathcal{I}(\theta) = 1/\theta$ from example 11.6.

1. Find the Wald interval for θ with level $1 - \alpha$ and determine the interval for level 95% when $n = 20$ is the number of counts with $\sum_i x_i = 30$.
2. In exercise 17.2 we have found the asymptotic distribution of the MLE $\hat{p}_0 = e^{-\bar{X}}$ for $p_0 = \mathbb{P}_\theta(X_1 = 0) = e^{-\theta}$. Use this distribution to find an approximate confidence interval for p_0 and work out the interval for the given data.

3. Another interval for p_0 can be found by applying the decreasing function $g(\theta) = e^{-\theta}$ to the boundaries of the interval from the first part. Show that this interval also has asymptotic coverage $1 - \alpha$, compute the interval for the given data and comment on the differences between the two intervals.

Exercise 25.4. Let $X_1, \dots, X_n$ be i.i.d. $N(\mu, \sigma^2)$, where both parameters are unknown. For $\mu_0 \in \mathbb{R}$ consider the test which rejects $H_0: \mu = \mu_0$ if $\sqrt{n}|\bar{X} - \mu_0|/S > t_{n-1, 1-\alpha/2}$.

1. Show that the test has size α , independent of the value of σ^2 .
2. Using theorem 25.6, determine the confidence set of values μ_0 which are not rejected. Check that your confidence set coincides with the t -interval from example 25.4.
3. For the data given in example 25.4, can you decide whether $H_0: \mu = 9$ and $H_0: \mu = 10$ should be rejected at level 5%, without having to compute the test statistic?
4. Now consider the one-sided test which rejects $H_0: \mu = \mu_0$ if $\sqrt{n}(\bar{X} - \mu_0)/S > t_{n-1, 1-\alpha}$. Invert this test and describe the resulting confidence set.

Lecture 26

Bayesian Inference: Priors and Posteriors

So far the parameter θ has been a fixed but unknown number, and a sentence such as “ θ is probably close to 0.3” had no meaning. In this lecture and in lecture 28 we will see that the parameter itself can be given a probability distribution: the *prior*, chosen before the data are seen, and the *posterior*, obtained after the data are seen. (We will see an update in pictures in lecture 27.)

26.1 Prior, Likelihood and Posterior

Since θ is now a random variable, the model density is a conditional density given the parameter, and in the Bayesian lectures we write it as $f(x|\theta)$, where $x = (x_1, \dots, x_n)$ is the whole data set, so that $f(x|\theta) = \prod_i f(x_i|\theta)$ is the likelihood $L(\theta)$ of lecture 4.

Definition 26.1. Let θ be a random variable with values in Θ and density $\pi(\theta)$, and assume that given θ the data $X = (X_1, \dots, X_n)$ have joint density $f(x|\theta)$. Then the density π is called the *prior density* (or prior) of θ and the conditional density $\pi(\theta|x)$ of θ given $X = x$ is called the *posterior density* (or posterior). The parameters in the prior are called *hyperparameters*.

From now on, unless otherwise specified, π will denote a prior or posterior density. In lecture 19, we used the same letter to denote the mixture weight for a two-component model.

The posterior density can be found using Bayes’ theorem as given in proposition A.20 in appendix A, where the parameter plays the role of A and the data play the role of B .

Theorem 26.2. *In the situation of definition 26.1, the posterior density is given by*

$$\pi(\theta|x) = \frac{f(x|\theta) \pi(\theta)}{m(x)}, \quad \text{where } m(x) = \int_{\Theta} f(x|\vartheta) \pi(\vartheta) \, d\vartheta \quad (26.1)$$

is the marginal density of the data. The formula holds for all x with $m(x) > 0$.

Proof. The joint density of (θ, X) is given by $f(x|\theta) \pi(\theta)$. Integrating out θ gives the marginal density $m(x)$ of X . Dividing the joint density by $m(x)$ gives the conditional density of θ given $X = x$, as claimed. $\square$

The denominator $m(x)$ makes the posterior integrate to one and does not involve θ ; in practice we never compute it directly, but write

$$\pi(\theta|x) \propto f(x|\theta) \pi(\theta), \quad (26.2)$$

where $\propto$ means “equal up to a factor free of θ ”, and we try to recognise the right-hand side as the kernel of a known density, whose normalising constant then comes from table A.2. In words, the posterior is proportional to the likelihood times the prior.

The posterior is the complete Bayesian answer: a statement like “ $\mathbb{P}(\theta > 0.3|x) = 0.95$ ” now has a meaning. Point estimates are summaries of the posterior (section 26.3), and intervals and tests will be considered in lecture 28.

26.2 Conjugate Priors

The rule (26.2) works best if the posterior distribution belongs to the same family of distributions as the prior distribution, so that the update amounts to no more than a change of the hyperparameters.

Definition 26.3. A family $\mathcal{P}$ of prior distributions is *conjugate* for a model $(f(x|\theta) | \theta \in \Theta)$, if for every prior in $\mathcal{P}$ and every data set x the posterior $\pi(\theta|x)$ also belongs to $\mathcal{P}$.

For exponential families, conjugate priors can be written down by inspection: the likelihood depends on the natural parameter η only through the exponent $\eta T_n - nK(\eta)$ (lemma 7.5), and a prior with an exponent of the same form reproduces itself.

Proposition 26.4. Let $f(x|\eta) = h(x) \exp(\eta T(x) - K(\eta))$ with $\eta \in \mathcal{N}$ be a one-parameter exponential family in its natural parametrisation, and for real numbers a and b let

$$\pi_{a,b}(\eta) \propto \exp(a\eta - bK(\eta)), \quad \eta \in \mathcal{N},$$

whenever the right-hand side has a finite integral over $\mathcal{N}$. Then these priors are conjugate for the model: if the prior is $\pi_{a,b}$, the data $x_1, \dots, x_n$ have natural statistic $T_n = \sum_{i=1}^n T(x_i)$ and the marginal density satisfies $0 < m(x) < \infty$, then the posterior is $\pi_{a+T_n, b+n}$.

Proof. Using (26.2) and lemma 7.5 we find

$$\begin{aligned} \pi(\eta|x) &\propto \prod_{i=1}^n h(x_i) \exp(\eta T_n - nK(\eta)) \exp(a\eta - bK(\eta)) \\ &\propto \exp((a + T_n)\eta - (b + n)K(\eta)), \end{aligned}$$

since $\prod_i h(x_i)$ does not depend on η . The right-hand side is the kernel of $\pi_{a+T_n, b+n}$ and by assumption $m(x) < \infty$ this kernel has finite integral and thus is a density. This shows that the posterior is $\pi_{a+T_n, b+n}$. This completes the proof. $\square$

The update adds the natural statistic to a and the sample size to b : the prior acts like b earlier observations with statistic total a . We now work the recipe for our simplest model in the original parameter.

Example 26.5. Let $X_1, \dots, X_n$ be i.i.d. Bernoulli with success probability $\theta \in (0, 1)$ and let the prior be the Beta(a, b) distribution with $a, b > 0$. From table A.2 we know that this distribution has density

$$\pi(\theta) = \frac{\theta^{a-1}(1-\theta)^{b-1}}{B(a, b)}, \quad 0 < \theta < 1,$$

and the mean of this distribution is $a/(a+b)$. If $T = \sum_i x_i$ is the number of successes in the sample, then the likelihood is $f(x|\theta) = \theta^T(1-\theta)^{n-T}$, as in example 4.2, and thus we have

$$\pi(\theta|x) \propto \theta^T(1-\theta)^{n-T} \cdot \theta^{a-1}(1-\theta)^{b-1} = \theta^{a+T-1}(1-\theta)^{b+n-T-1},$$

where we have omitted the constant $1/B(a, b)$ for simplicity. The right-hand side is the kernel of a Beta($a+T, b+n-T$) density and thus we have found the posterior distribution: the Beta distributions are conjugate for the Bernoulli model, and the update adds the successes to a and the failures to b . The posterior mean is given by

$$\mathbb{E}(\theta|x) = \frac{a+T}{a+b+n} = \frac{a+b}{a+b+n} \cdot \frac{a}{a+b} + \frac{n}{a+b+n} \cdot \bar{x}, \quad (26.3)$$

which is a weighted average of the prior mean $a/(a+b)$ and the sample proportion $\bar{x} = T/n$, *i.e.* the maximum likelihood estimate from example 5.3. The prior can be interpreted as having observed $a+b$ additional samples, of which a were successes. This *equivalent sample size* is a measure for how the hyperparameters should be chosen in practice (exercise 26.1). As $n \rightarrow \infty$, the weight of the prior decreases and with enough data the prior will be completely forgotten. For $a = b = 2$, $n = 10$ and $T = 7$, for example, the posterior distribution is Beta(9, 5) with mean $9/14 \approx 0.643$, between the prior mean 0.5 and the sample proportion 0.7.

In terms of the log-odds η , the example is an instance of proposition 26.4. Using the same arguments, we can show that normal priors are conjugate for the normal mean with known variance. The cumulant function in this case is a quadratic function of η (exercise 7.4). The result is given here without proof, and the proof, based on completing the square, is given as exercise 26.3.

Proposition 26.6. Let $X_1, \dots, X_n$ be i.i.d. $N(\theta, \sigma^2)$ with known variance σ^2 and let the prior be $\theta \sim N(m, \tau^2)$. Then the posterior distribution is $N(m_n, \tau_n^2)$ where

$$\frac{1}{\tau_n^2} = \frac{1}{\tau^2} + \frac{n}{\sigma^2} \quad \text{and} \quad m_n = \frac{1/\tau^2}{1/\tau^2 + n/\sigma^2} m + \frac{n/\sigma^2}{1/\tau^2 + n/\sigma^2} \bar{x}.$$

The reciprocal of a variance is called a *precision*, and in these terms the proposition is easy to remember: precisions add, and the posterior mean is the precision-weighted average of the prior mean and $\bar{x}$. Again, the prior washes out as $n \rightarrow \infty$.

26.3 Point Estimates from the Posterior

If a single number is required for θ , the posterior mean, median and mode can all be used. If the errors are measured in terms of squared error, as in lecture 2, the posterior mean is the optimal choice.

Theorem 26.7. *Assume that the posterior distribution has finite variance. Then the posterior mean $\hat{\theta}_B = \mathbb{E}(\theta|x)$ minimises the posterior expected squared error, i.e. for every number d we have*

$$\mathbb{E}((\theta - d)^2 | x) \geq \mathbb{E}((\theta - \hat{\theta}_B)^2 | x) = \text{Var}(\theta|x),$$

and the equality holds only for $d = \hat{\theta}_B$.

Proof. Adding and subtracting $\hat{\theta}_B$ inside the square, as in the proof of theorem 2.7, we get

$$\mathbb{E}((\theta - d)^2 | x) = \mathbb{E}((\theta - \hat{\theta}_B)^2 | x) + 2(\hat{\theta}_B - d) \mathbb{E}(\theta - \hat{\theta}_B | x) + (\hat{\theta}_B - d)^2,$$

where all expectations are taken with respect to the posterior. The middle term vanishes, because $\mathbb{E}(\theta|x) = \hat{\theta}_B$, and the last term is non-negative and zero only for $d = \hat{\theta}_B$. This completes the proof. $\square$

The estimator $\hat{\theta}_B$ is called the *Bayes estimator* (under squared error) for θ . The theorem is a statement about fixed data. For fixed θ , and varying data, exercise 26.4 shows that the Bernoulli $\hat{\theta}_B$ is biased but consistent, and that neither the Bernoulli estimator nor $\bar{X}$ has smaller mean squared error for all θ .

26.4 Flat, Improper and Jeffreys Priors

Without prior knowledge to encode, the obvious choice is the *flat prior* which is a constant density on Θ . This prior does not prefer any value of θ . For the Bernoulli model this prior is $\text{Beta}(1, 1)$ and from example 26.5 we know that the posterior mean is $(T + 1)/(n + 2)$, which is close to $\bar{x}$ but is shifted towards $1/2$. Since a constant density on an unbounded parameter space has infinite integral, the flat prior is *improper*. Nevertheless, this prior is still in use today: if $\int f(x|\theta) \pi(\theta)$ has finite integral, it is normalised and used as the posterior. The finiteness of the integral must be checked in each case (exercise 26.2). For the normal mean, the flat prior on $\mathbb{R}$ is the limit $\tau \rightarrow \infty$ in proposition 26.6 and in this case leads to the proper posterior $N(\bar{x}, \sigma^2/n)$.

The flat prior has a more basic defect: flatness depends on the parametrisation. If θ is uniform on $(0, 1)$, then $\psi = \theta^2$ has density $1/(2\sqrt{\psi})$ by proposition A.19, and a prior which claims to know nothing about θ claims to know something about θ^2 . The Fisher information provides a parametrisation-free rule, since proposition 11.10 shows how it transforms.

Definition 26.8. In the situation of section 11.1, the *Jeffreys prior* for θ is given by

$$\pi_J(\theta) \propto \sqrt{\mathcal{I}(\theta)}, \quad \theta \in \Theta,$$

where $\mathcal{I}(\theta)$ is the Fisher information of an observation.

This prior is justified by the following invariance property, which the flat prior lacks.

Proposition 26.9. *Let $\psi = g(\theta)$ where $g: \Theta \rightarrow g(\Theta)$ is a continuously differentiable bijection with $g'(\theta) \neq 0$ for all θ and let $h = g^{-1}$.*

1. If θ has the flat prior on Θ , then ψ has the density $|h'(\psi)|$ on $g(\Theta)$. The density is flat only if g is affine.
2. If θ has the Jeffreys prior $\sqrt{\mathcal{I}_\theta(\theta)}$, then ψ has the density $\sqrt{\mathcal{I}_\psi(\psi)}$. This is the Jeffreys prior computed using the ψ -parametrisation.

Proof. Using the change-of-variable formula from proposition A.19, if a density π_θ for θ is given, then a density of $\pi_\psi(\psi) = \pi_\theta(h(\psi)) |h'(\psi)|$ for ψ can be derived. The same substitution can be applied to improper densities. In the case of the flat prior $\pi_\theta \equiv 1$, we have $|h'(\psi)|$ and the derivative of h is constant only if h , i.e. if g , is affine. For the Jeffreys prior we have $\pi_\theta(\theta) = \sqrt{\mathcal{I}_\theta(\theta)}$ and using proposition 11.10 and $h'(\psi) = 1/g'(\theta)$ for $\theta = h(\psi)$ we get

$$\pi_\psi(\psi) = \sqrt{\mathcal{I}_\theta(\theta)} |h'(\psi)| = \frac{\sqrt{\mathcal{I}_\theta(\theta)}}{|g'(\theta)|} = \sqrt{\frac{\mathcal{I}_\theta(\theta)}{g'(\theta)^2}} = \sqrt{\mathcal{I}_\psi(\psi)}.$$

This is the Jeffreys prior in the ψ -parametrisation. This completes the proof. $\square$

The Jeffreys prior is a property of the model, rather than of the way it is labelled, and as such it can be proper or improper.

Example 26.10. For the Bernoulli model, from example 11.5 we know $\mathcal{I}(\theta) = 1/(\theta(1-\theta))$ and thus the Jeffreys prior is given by

$$\pi_J(\theta) \propto \theta^{-1/2}(1-\theta)^{-1/2}, \quad 0 < \theta < 1.$$

This is the density of the Beta(1/2, 1/2) distribution. The Jeffreys prior is proper and has more weight near the boundaries of the interval than the flat prior does. From example 26.5 we know that the posterior distribution is Beta($T+1/2$, $n-T+1/2$) and the mean of the posterior is $(T+1/2)/(n+1)$. This value is between the maximum likelihood estimate T/n and the estimate $(T+1)/(n+2)$ obtained using the flat prior. For large n , all three values coincide.

For location and scale models the Jeffreys prior is the same for every base density f_0 .

Proposition 26.11. Let f_0 be a fixed density such that the following models satisfy the conditions of section 11.1.

1. Consider the location model $f(x|\theta) = f_0(x-\theta)$ with $\theta \in \mathbb{R}$. The Fisher information in this model is constant as a function of θ and thus the Jeffreys prior is flat, $\pi_J(\theta) \propto 1$ on $\mathbb{R}$.
2. Consider the scale model $f(x|\sigma) = f_0(x/\sigma)/\sigma$ with $\sigma > 0$. The Fisher information in this model is $\mathcal{I}(\sigma) = c/\sigma^2$ for a constant $c > 0$. Thus, the Jeffreys prior is $\pi_J(\sigma) \propto 1/\sigma$ on $(0, \infty)$.

Both priors are improper.

Exercise 26.5 verifies both statements by writing the score as a function of the standardised observation. The normal mean with known variance is a location model, so its Jeffreys prior is flat. Similarly, the normal standard deviation with known mean is a scale model with $\mathcal{I}(\sigma) = 2/\sigma^2$.

Summary.

- The data can be used to update the prior density $\pi(\theta)$ to the posterior $\pi(\theta|x) \propto f(x|\theta)\pi(\theta)$. The normalising constant can be found by recognising the kernel of a known density.
- Priors are conjugate for a family if the posterior always belongs to the same family. For a one-parameter exponential family in its natural parametrisation, the priors $\exp(a\eta - bK(\eta))$ are conjugate. Beta priors are conjugate for the Bernoulli model, and normal priors are conjugate for the normal mean with known variance.
- The posterior mean is a weighted average of the prior mean and the maximum likelihood estimate. As n increases, the prior weight decreases. The posterior mean is the Bayes estimator for θ under squared error.
- A flat prior is not invariant under reparametrisation, but the Jeffreys prior $\pi_J(\theta) \propto \sqrt{I(\theta)}$ is. The Jeffreys prior is $\text{Beta}(1/2, 1/2)$ for the Bernoulli model, is flat for a location parameter, and is proportional to $1/\sigma$ for a scale parameter. An improper prior can only be used if the posterior is proper.

Exercise 26.1. A Bernoulli model is to be fitted with a $\text{Beta}(a, b)$ prior for the success probability θ .

1. Find a and b such that the prior mean is 0.3 and the equivalent sample size $a + b$ is 10. You can find the prior standard deviation in table A.2.
2. Find a and b such that the prior has mean 0.3 and standard deviation 0.1. What is the equivalent sample size for this prior?
3. For the prior in the first part of the question, a sample of size $n = 20$ has $T = 9$ successes. Determine the posterior distribution and the posterior mean. Write the posterior mean as a weighted average of the prior mean and $\bar{x}$.

Exercise 26.2. The *Haldane prior* for the Bernoulli success probability is the limit $a, b \rightarrow 0$ of the $\text{Beta}(a, b)$ prior, *i.e.* the function $\pi(\theta) \propto \theta^{-1}(1 - \theta)^{-1}$ for $0 < \theta < 1$.

1. Show that this prior is improper.
2. Show that the formal posterior $\pi(\theta|x) \propto \theta^{T-1}(1 - \theta)^{n-T-1}$ is a proper distribution if and only if $0 < T < n$. Identify the resulting distribution.
3. For $0 < T < n$, determine the posterior mean and compare it to the maximum likelihood estimate.
4. Describe what happens for $T = 0$.

Exercise 26.3. Using the following steps, prove proposition 26.6.

1. Show that the likelihood, as a function of θ , satisfies $f(x|\theta) \propto \exp(-n(\theta - \bar{x})^2/(2\sigma^2))$.
2. Multiply the prior density with this function and show, by completing the square in θ , that

$$\pi(\theta|x) \propto \exp\left(-\frac{(\theta - m_n)^2}{2\tau_n^2}\right)$$

where m_n and τ_n^2 are as in the proposition.

3. What happens to m_n and τ_n^2 when $\tau \rightarrow \infty$ with n fixed? And when $n \rightarrow \infty$ with τ fixed?

Exercise 26.4. Let $X_1, \dots, X_n$ be i.i.d. Bernoulli with success probability θ and let $\hat{\theta}_B = (a + T)/(a + b + n)$ with $T = \sum_i X_i$ be the Bayes estimator from example 26.5. Now fix θ and let $\hat{\theta}_B$ be an estimator as in lecture 2.

1. What are the bias and variance of $\hat{\theta}_B$? Show that the estimator is biased for every $\theta \neq a/(a + b)$, but consistent.
2. For $a = b = 1$ and $n = 10$, compare the estimators $\text{MSE}(\hat{\theta}_B)$ and $\text{MSE}(\bar{X})$ for $\theta = 1/2$ and $\theta = 0.1$.

Exercise 26.5. Let f_0 be a fixed density as in proposition 26.11 and let $\ell_0(z) = \log f_0(z)$.

1. For the location model $f(x|\theta) = f_0(x - \theta)$, show that the score is $-\ell'_0(x - \theta)$ and thus that the Fisher information does not depend on θ .
2. For the scale model $f(x|\sigma) = f_0(x/\sigma)/\sigma$, show that the score is $-(1 + z\ell'_0(z))/\sigma$ where $z = x/\sigma$ and that $\mathcal{I}(\sigma) = c/\sigma^2$ where c does not depend on σ .
3. For the $N(0, \sigma^2)$ model with parameter σ , show that $\mathcal{I}(\sigma) = 2/\sigma^2$.

Lecture 27

Confidence Versus Credibility

This is the R session for week 9 of the course. In lectures this week we have introduced the topics of confidence intervals and Bayesian priors and posteriors. The aim of this R session is to visualise these concepts: we will consider a hundred confidence intervals, together with a count of how often these intervals “miss” the target, we will consider a confidence interval and a Bayesian interval for the same data, and we will consider how a prior can be transformed into a posterior as observations are added one by one. Today we will use the Bayesian interval in an informal way, in preparation for the fact that in one week (in lecture 28) we will introduce the concept of a credible interval and will then be able to make a proper comparison between confidence intervals and credible intervals.

27.1 Coverage by Simulation

The interval of the day is the z -interval from example 25.3, using the numerical values from the example ($n = 16$, $\sigma = 2$ are known), so that the 95% interval is $\bar{X} \pm 1.96 \cdot 0.5$. By definition 25.1 of the level 0.95, the random interval contains the fixed value θ in 95% of all samples. To verify this, we consider $\theta = 10$ and compute the interval for 100 samples. Before you execute the code, answer the following questions: (1) How many of the 100 intervals will not cover the value 10? (2) Will all intervals have the same width? (3) Will the misses be clustered or evenly distributed among the 100 samples?

```
theta <- 10
sigma <- 2
n <- 16
z_interval <- function() {
  x <- rnorm(n, mean = theta, sd = sigma)
  mean(x) + c(-1, 1) * qnorm(0.975) * sigma / sqrt(n)
}
set.seed(2026)
ci <- replicate(100, z_interval())
covers <- ci[1, ] <= theta & theta <= ci[2, ]
sum(covers)
which(!covers)
```

```
## [1] 95
## [1] 1 29 34 63 83
```


As in lecture 12, `covers` is a vector of `TRUE` and `FALSE`, one for each interval, indicating whether the interval covers the truth. The command `mean` counts the proportion of `TRUE` values, which are interpreted as ones and zeros by R. For the example, 95 out of 100 intervals cover the truth. The misses are for samples 1, 29, 34, 63 and 83. The exact number 95 is a coincidence, determined by the seed of the random number generator: the number of intervals which cover the truth is binomially distributed with standard deviation 2.2, so 92 or 98 would have been just as unremarkable. The command `segments` can be used to plot each interval as a horizontal line, and the samples where the interval covers the truth are plotted in grey while the misses are plotted in black.

```
plot(NULL, xlim = range(ci), ylim = c(1, 100),
     xlab = expression(theta), ylab = "sample number")
segments(ci[1, ], 1:100, ci[2, ], 1:100,
        col = ifelse(covers, "grey60", "black"), lwd = 2)
abline(v = theta, lty = 2)
```




Figure 27.1: One hundred 95% confidence intervals for the mean of a $N(\theta, 4)$ sample of size $n = 16$, with true value $\theta = 10$ (dashed line). The grey intervals cover the truth, the black intervals do not.

Figure 27.1 illustrates the confidence level of the procedure in a graphical way. The intervals in the figure all have the same width, since the value σ is known: the intervals move, while the dashed line stays fixed. The misses are not clustered but are scattered throughout the figure, since the samples are independent. If we had only observed sample 29, the interval $[10.43, 12.39]$ would have been indistinguishable from the others, and there would have been no indication in the data that this was a miss. For this reason, the statement “ θ is in the interval $[10.43, 12.39]$ with probability 0.95” is wrong: the value 10 is either in the interval or it is not, but in this instance it is not. The probability refers to the procedure, not to the individual values. Before the code is executed, predict the long-run proportion of intervals which cover the truth, and determine the distance from 0.95 where an error is no longer considered significant.

```
ci <- replicate(10000, z_interval())
```

```
mean(ci[1, ] <= theta & theta <= ci[2, ])
```

```
## [1] 0.9473
```

The long-run coverage is 0.947. The standard error of this estimate is $\sqrt{0.95 \cdot 0.05 / 10000} = 0.002$, so the estimate is about one and a half standard errors below the nominal value 0.95: this is good agreement, considering the simulation error. The pivot argument in example 25.3 guarantees that this result holds for every value of θ .

27.2 The Same Data, Two Intervals

We now take the data of example 25.8, $T = 37$ successes in $n = 100$ Bernoulli trials, and compute two intervals for the success probability θ . The first is the Wald interval $\bar{x} \pm 1.96\sqrt{\bar{x}(1-\bar{x})/n}$ of that example. For the second we treat θ as random, as in lecture 26: the flat-prior posterior is $\text{Beta}(1+T, 1+n-T)$ by example 26.5, and the interval between its 2.5% and 97.5% quantiles contains θ with posterior probability 0.95; this is the equal-tailed credible interval of lecture 28, and the command `qbeta` can be used to find the quantiles. Before you run the code, predict whether the credible interval will be wider or narrower than the Wald interval, and whether it will sit to the left or to the right of it.

```
n <- 100
total <- 37
xbar <- total / n
se <- sqrt(xbar * (1 - xbar) / n)
xbar + c(-1, 1) * qnorm(0.975) * se
a <- 1 + total
b <- 1 + n - total
qbeta(c(0.025, 0.975), a, b)
```

```
## [1] 0.2753721 0.4646279
```

```
## [1] 0.2817646 0.4680688
```

The two intervals, $[0.275, 0.465]$ and $[0.282, 0.468]$, coincide to two decimal digits; the credible interval is slightly shifted towards $1/2$, due to the fact that the flat prior adds one success and one failure to the data. The Wald interval states that the recipe used to compute it covers the truth in approximately 95% of samples. It does not make a statement about this particular interval. The credible interval states that, given the observations and the flat prior, the probability that θ falls between 0.282 and 0.468 is 0.95: the sentence forbidden in the previous section is the correct interpretation in this context, since θ is a random variable in the Bayesian model. The posterior distribution can be used to answer other questions; the command `pbeta` can be used to evaluate its distribution function.

```
lower <- xbar - qnorm(0.975) * se
upper <- xbar + qnorm(0.975) * se
pbeta(upper, a, b) - pbeta(lower, a, b)
1 - pbeta(0.5, a, b)
qbeta(c(0.025, 0.975), 12 + total, 3 + n - total)
```

```
## [1] 0.9531666
## [1] 0.004667524
## [1] 0.3374894 0.5171204
```

The Wald interval itself has posterior probability 0.953 under the flat prior, and the posterior probability that θ exceeds $1/2$ is 0.005. The frequentist interval cannot make a statement about this probability. The final line shows the price of departing from the flat prior: under the informative prior $\text{Beta}(12, 3)$, with mean 0.8 and equivalent sample size 15, the same data give the credible interval $[0.337, 0.517]$. A credible interval always depends on the prior, and numerical agreement with a confidence interval is a property of the flat prior and of the sample size. We will consider this comparison in more detail in lecture 28.

27.3 From Prior to Posterior

The following picture illustrates the update from lecture 26, where we process the observations one by one. In the example, we simulate 100 Bernoulli observations with $\theta = 0.3$. After k observations, with s successes, the prior distribution with $\text{Beta}(a, b)$ has evolved into the posterior distribution $\text{Beta}(a + s, b + k - s)$ with mean $(a + s)/(a + b + k)$. This result was obtained in example 26.5. The box below contains a script for computing the posterior mean of the flat prior after $k = 20$ observations. The lines in the box are shuffled, so you will need to sort them into a working script before you can try the code.

```
successes <- cumsum(x)
theta <- 0.3
(a + successes[k]) / (a + b + k)    # posterior mean after k
  observations
x <- rbinom(100, size = 1, prob = theta)
k <- 20
set.seed(2026)
a <- 1
b <- 1
```

As always, names must be defined before they are used, and the seed must be set before the first random draw: `theta` first, followed by the seed, then `rbinom` with `size = 1` to generate the Bernoulli draws, then `cumsum` to generate the running total of successes; `a`, `b` and `k` can be placed anywhere before the last line. The script `R/S27-posterior.R` extends this to handle three priors simultaneously, stored as the rows of a matrix: the flat prior $\text{Beta}(1, 1)$, the Jeffreys prior $\text{Beta}(1/2, 1/2)$ from example 26.10 and the misplaced prior $\text{Beta}(12, 3)$ from the previous section. Before you look at the output of the following code, predict the result: which of the three posterior means will be closest to 0.3 after five observations? Will you still be able to distinguish between the three posterior means after one hundred observations?

```
priors <- rbind(flat = c(1, 1), jeffreys = c(0.5, 0.5),
               informative = c(12, 3))
post_mean <- function(k) {
  s <- if (k == 0) 0 else successes[k]
  (priors[, 1] + s) / (priors[, 1] + priors[, 2] + k)
}
```

```
head(x, 10)
successes[c(5, 20, 100)]
round(sapply(c(0, 5, 20, 100), post_mean), 3)
```

```
## [1] 0 0 0 0 0 0 0 1 0 0
## [1] 0 2 28
##           [,1] [,2] [,3] [,4]
## flat      0.5 0.143 0.136 0.284
## jeffreys   0.5 0.083 0.119 0.282
## informative 0.8 0.600 0.400 0.348
```

The start of the sequence is bad: the first five observations are all failures and after twenty observations there have been only two successes. The probability of at most two successes in twenty observations is 0.035 for $\theta = 0.3$. For this case, the prediction at $k = 5$ is dominated by the flat prior: the posterior mean 0.143 is closest to 0.3, since five failures have driven both weak priors below the truth, and the Jeffreys prior is driven lower than the flat prior is. The misplaced prior is still at 0.6. The two weak priors follow the data and end at 0.14 and 0.12, both far below the truth. The misplaced prior, moving from 0.8 to 0.4, is the closest of the three at $k = 20$: with twenty observations, the data alone is not enough to avoid being misled, but a prior in the correct region helps and a prior in the wrong region does harm. The data does not allow us to infer the correct region. By $k = 100$ the two weak priors are indistinguishable at 0.28 and even the misplaced prior, with an equivalent sample size of 15, is now mostly aligned with the sample proportion. The script plots the three posterior densities at each stage using `curve` and `dbeta`. Before you look at figure 27.2, sketch the panel for $k = 5$: which curves are steep, where do they take their peaks and how wide are the curves?

Figure 27.2 summarises the update procedure. At $k = 0$ the three curves correspond to the prior distributions from lecture 26. One can see that the Jeffreys prior is not flat. At $k = 5$ the two weak posteriors have the largest values at 0, since no successes have been observed so far, and the informative posterior has not yet significantly changed. At $k = 20$ all three posteriors have the form of humps, and the posteriors are narrower than in the previous step: as data are added, the posterior gets sharper, irrespective of the prior distribution. At $k = 100$ the two weak posteriors coincide and the informative posterior overlaps these two: this is a consequence of the prior being washed out in example 26.5, and the panel for $k = 20$ serves as a reminder that data are needed to achieve this effect.

Repeat the picture in figure 27.1 for the t -interval from example 25.4, but with $n = 5$, `sigma` replaced with `sd(x)` and `qnorm(0.975)` replaced with `qt(0.975, df = 4)`. The widths of the intervals will change. Estimate how often the intervals miss the target value, and estimate the long-run coverage from 10000 samples. Repeat the experiment, but this time leaving `qnorm(0.975)` instead of the t quantile, and comment on your result. Finally, choose a prior mean and an equivalent sample size as in exercise 26.1, use your student ID as the seed for the random number generator, and repeat the posterior sequence of the last section, both with and without the prior. For both cases, determine the first k where the posterior mean is within 0.05 of the truth.

Type the commands used in this session into R, and check that you get the




Figure 27.2: Posterior densities for the probability of success in a Bernoulli trial after $k = 0, 5, 20$ and 100 observations of one simulated sequence with $\theta = 0.3$ (grey vertical line), for the priors Beta(1, 1) (solid), Beta(1/2, 1/2) (dashed) and Beta(12, 3) (dotted).

same results as the text.

Summary.

- The written reference for the commands used in this session is appendix B.
- The confidence level of a statistical procedure is a property of the procedure: in 100 simulated trials of 95% intervals, 95 intervals contained the truth, the misses were indistinguishable from the successes, and in the long run the coverage was 0.947. A computed confidence interval either contains the parameter or it does not.
- For the same Bernoulli data, the Wald interval and the credible interval with flat prior coincide to two decimal digits, but they are different statements: the credible interval is a probability statement about θ , conditioned on the data and on the prior, and thus depends on the prior.
- The Beta posterior gets more and more sharp as observations are added, and in the limit it converges to the sample proportion for every prior. For 20 observations the prior is still noticeable, but for 100 the prior is largely washed out.

Lecture 28

Credible Intervals and Bayesian Testing

Lecture 26 ended with the posterior distribution as the complete Bayesian answer. In this lecture we answer the interval question of lecture 25 and the testing question of lectures 20 to 23 in posterior terms, and then set the two approaches side by side on the mean of a normal sample; lecture 27 shows the same comparison in pictures.

We use the two conjugate pairs from lecture 26: the Beta posterior from example 26.5 for the Bernoulli distribution, and for the normal distribution $X_1, \dots, X_n \sim N(\theta, \sigma^2)$ with known variance and prior $\theta \sim N(m, \tau^2)$, we have the posterior $N(m_n, \tau_n^2)$ from proposition 26.6, with

$$\frac{1}{\tau_n^2} = \frac{1}{\tau^2} + \frac{n}{\sigma^2} \quad \text{and} \quad m_n = \tau_n^2 \left(\frac{m}{\tau^2} + \frac{n\bar{x}}{\sigma^2} \right). \quad (28.1)$$

We consider the data set from example 25.3: $n = 16$ observations with $\sigma = 2$ and $\bar{x} = 10.3$. Assume that before the data were collected we expected the mean to be around 12 and we chose the prior $N(12, 1)$. Then the posterior precision is $1/1 + 16/4 = 5$ and the posterior mean is $m_n = 0.2 \cdot (12 + 4 \cdot 10.3) = 10.64$, i.e. the posterior distribution is $N(10.64, 0.2)$.

28.1 Credible Intervals

A confidence interval, as given in definition 25.1, is a random interval which covers the fixed parameter with a given probability. In contrast, in a Bayesian model the data are fixed after observation, the parameter is random and an interval estimate is a set which contains most of the posterior probability.

Definition 28.1. Let $\pi(\theta|x)$ be the posterior density of θ given the data x and let $\alpha \in (0, 1)$. An interval $[l, u]$ is a *credible interval* for θ with *credibility* $1 - \alpha$, if

$$\mathbb{P}(l \leq \theta \leq u|x) = \int_l^u \pi(\vartheta|x) d\vartheta = 1 - \alpha.$$

The *equal-tailed* credible interval is the interval $[q_{\alpha/2}, q_{1-\alpha/2}]$, between the $\alpha/2$ - and $(1 - \alpha/2)$ -quantiles of the posterior distribution.

The probability in the definition is a probability about θ , for the given data set: the sentence “ θ is in $[l, u]$ with probability 0.95” (which we had to forbid in

lecture 25) is the exact statement a 95% credible interval makes. The trade-off is that this probability depends on the prior distribution.

Any interval which has posterior probability $1 - \alpha$ works, and the equal-tailed interval is the usual choice, because the edges of this interval are quantiles. The alternative would be to find the shortest such set.

Definition 28.2. A *highest posterior density* (HPD) set with credibility $1 - \alpha$ is the set

$$H_c = \{ \theta \in \Theta \mid \pi(\theta|x) \geq c \},$$

where $c > 0$ is chosen such that $\mathbb{P}(\theta \in H_c|x) = 1 - \alpha$.

The set H_c is the smallest set with posterior probability $1 - \alpha$. Any time we want to exchange part of this set with a region where the density is below c , we need to include more length in the interval to make up for the lost probability. For a unimodal posterior distribution, H_c forms an interval, called the *HPD interval*, where the two boundaries have the same posterior density. For a symmetric, unimodal posterior distribution, this interval will be an equal-tailed interval (see exercise 28.1).

Example 28.3. If the posterior distribution is equal to the normal distribution $N(m_n, \tau_n^2)$, then the density is symmetric around m_n and decreases as it gets further away from this point. Thus, the HPD interval is the equal-tailed interval

$$[m_n - z_{1-\alpha/2} \tau_n, m_n + z_{1-\alpha/2} \tau_n]. \quad (28.2)$$

For the data set and prior used in the section, we have the prior $N(12, 1)$ and $\tau_n = \sqrt{0.2} = 0.447$, and the 95% credible interval is $10.64 \pm 1.96 \cdot 0.447$, *i.e.* the interval $[9.76, 11.52]$. The credible interval is shifted towards the prior mean and is shorter than the 95% confidence interval $[9.32, 11.28]$ from example 25.3, because the prior provides additional information. In the case of the flat prior, $\tau \rightarrow \infty$ in section 26.4, the posterior distribution is $N(\bar{x}, \sigma^2/n)$ and the credible interval (28.2) coincides with the confidence interval $\bar{x} \pm z_{1-\alpha/2} \sigma/\sqrt{n}$. This coincidence is shown in section 28.3.

For a skewed posterior the HPD interval is shifted towards the mode, slightly shorter than the equal-tailed interval, and the numerical determination of this interval is the topic of exercise 28.2, which shows the difference between the two intervals for the skewed posterior $\text{Beta}(9, 5)$ in example 26.5. As n increases, the difference between the intervals decreases, because the posterior distribution approaches a symmetric normal distribution and the equal-tailed interval is then appropriate.

28.2 Bayesian Testing

In the Bayesian model a hypothesis $H_0: \theta \in \Theta_0$, in the sense of definition 20.1, is an event concerning the random variable θ and thus has a probability before and after the data are seen. Testing does not require new theory.

Definition 28.4. Let $H_0: \theta \in \Theta_0$ and $H_1: \theta \in \Theta_1$ be hypotheses with positive prior probability. Then the *posterior probability* of H_0 is given by

$$\mathbb{P}(H_0|x) = \mathbb{P}(\theta \in \Theta_0|x) = \int_{\Theta_0} \pi(\vartheta|x) d\vartheta.$$

Similarly, we can find the posterior probability of H_1 . A *Bayesian test* rejects H_0 if $\mathbb{P}(H_0|x)$ falls below a threshold or, if both hypotheses are treated equally, chooses the hypothesis with the higher posterior probability.

The threshold plays the role of the level α of definition 20.5, but the two numbers are not comparable: α bounds the probability of a type I error over repeated samples, whereas the threshold bounds the posterior probability of H_0 for the data in hand. Any asymmetry between the hypotheses can be expressed by the prior.

Example 28.5. We consider the running data set, testing $H_0: \theta \leq 10$ against $H_1: \theta > 10$. Using the prior $N(12, 1)$, we find the posterior distribution to be $N(10.64, 0.2)$. Thus we have

$$\mathbb{P}(H_0|x) = \mathbb{P}(\theta \leq 10|x) = \Phi\left(\frac{10 - 10.64}{0.447}\right) = \Phi(-1.43) = 0.076,$$

and consequently $\mathbb{P}(H_1|x) = 0.924$. A symmetric decision between the hypotheses favours H_1 . Using the rule with threshold 0.05, we would not have rejected H_0 . For the flat prior case, the posterior is $N(10.3, 0.25)$ and we get $\mathbb{P}(H_0|x) = \Phi(-0.6) = 0.274$. This is a much less certain result than in the example above: the prior $N(12, 1)$, which had only 2.3% of its mass below 10, contributed significantly to the evidence against H_0 . The same approach can be used to find the posterior probability of a hypothesis about a proportion: for the posterior $\text{Beta}(9, 5)$ in example 26.5, the posterior probability of $H_0: \theta \leq 1/2$ is the value of the Beta distribution function at $1/2$. This can be computed using any statistical package (in R, the function is `pbeta`), and the result is 0.133. Exercise 28.3 shows how to compute this integral manually.

A simple hypothesis $\theta = \theta_0$ has probability zero under every prior with a density, and thus for two simple hypotheses the prior must be a two-point distribution. The posterior in this case is then simple.

Definition 28.6. Let $H_0: \theta = \theta_0$ and $H_1: \theta = \theta_1$ be simple hypotheses and let the prior have probability p_0 for θ_0 and $p_1 = 1 - p_0$ for θ_1 . Then the *prior odds* of H_0 against H_1 are p_0/p_1 and the *posterior odds* are $\mathbb{P}(H_0|x)/\mathbb{P}(H_1|x)$.

Using the result of Bayes' theorem, proposition A.20, we find $\mathbb{P}(H_j|x) = p_j f(x|\theta_j)/m(x)$ for $j = 0, 1$ and dividing these two expressions we find that the marginal density $m(x)$ cancels out. This gives

$$\frac{\mathbb{P}(H_0|x)}{\mathbb{P}(H_1|x)} = \frac{p_0}{p_1} \cdot \frac{f(x|\theta_0)}{f(x|\theta_1)}, \quad (28.3)$$

i.e. the posterior odds are the prior odds times the likelihood ratio. In this context, the likelihood ratio is also called the *Bayes factor* in favour of H_0 . This is the inverse of the ratio considered in the Neyman–Pearson lemma, theorem 22.2. From the posterior odds r we can derive the posterior probabilities as $\mathbb{P}(H_0|x) = r/(1 + r)$.

Example 28.7. We consider the running data set again, this time to test the hypothesis $H_0: \theta = 10$ against the hypothesis $H_1: \theta = 11$. Since the likelihood of

observing a normal sample is $f(x|\theta) \propto \exp(-n(\bar{x}-\theta)^2/(2\sigma^2))$ from exercise 26.3, we find the Bayes factor to be

$$\frac{f(x|\theta_0)}{f(x|\theta_1)} = \exp\left(\frac{n}{2\sigma^2}((\bar{x}-\theta_1)^2 - (\bar{x}-\theta_0)^2)\right) = \exp\left(\frac{n(\theta_1-\theta_0)(\theta_0+\theta_1-2\bar{x})}{2\sigma^2}\right).$$

Thus we have $n = 16$, $\sigma = 2$ and $\bar{x} = 10.3$ and the exponent is $16 \cdot 1 \cdot 0.4/8 = 0.8$ and the Bayes factor is $e^{0.8} = 2.23$: the data are closer to the value 10 than to 11 and the data favour H_0 by a factor of approximately two. If the prior odds are 1, the posterior odds are 2.23 and $\mathbb{P}(H_0|x) = 2.23/3.23 = 0.69$. We see that the Bayes factor multiplies all prior opinions by the same factor, and in cases where the prior is very strong, it can be difficult to overturn the prior opinion. The situation in exercise 28.4 is more complicated, because the prior odds are not equal.

28.3 Comparing Frequentist and Bayesian Answers

Now we have two different sets of answers, for the same data, for the normal mean. We conclude by comparing the two sets.

The first difference is in the nature of the statements. A 95% confidence interval is a statement about a procedure: it says that for repeated samples, the interval covers the fixed truth in 95% of cases. A 95% credible interval is a statement about the parameter, given the data and the prior. Similarly, a p -value is computed assuming H_0 is true, whereas $\mathbb{P}(H_0|x)$ is the probability of the hypothesis given the data.

The second observation is that for a location parameter with a flat prior, both numbers coincide. Example 28.3 shows the two intervals coinciding, and the tests coincide as well: for $H_0: \theta \leq \theta_0$ the flat-prior posterior probability $\Phi(\sqrt{n}(\theta_0 - \bar{x})/\sigma)$ equals, by the symmetry of Φ , the p -value $1 - \Phi(\sqrt{n}(\bar{x} - \theta_0)/\sigma)$ of the one-sided z -test from example 20.6; for the running data set both values equal 0.274. This is the case because the flat-prior posterior density in θ has the same shape as the sampling density of $\bar{X}$ in $\bar{x}$. This is the reason why the misreadings in lectures 20 and 25 were so tempting: the forbidden sentence about the p -value is, for the simplest model, a true sentence about the flat-prior posterior, and only the interpretation needs to be changed.

The coincidence is only true for special models and priors. A genuine prior will move the Bayesian answer closer to the prior mean and will concentrate the answer around this value, as in the case of $N(12, 1)$ above. If the prior is correct, this is an advantage, but if the prior is wrong this effect will work against us. The frequentist guarantee from definition 25.1 applies for all θ and thus cannot be violated in this way; exercise 28.5, however, shows that the frequentist coverage of a credible interval can be worse than 95% when θ is far from the prior mean.

As the sample size increases, the influence of the prior decreases. For the normal mean, equation (28.1) shows that the weight of the prior mean in m_n goes to zero and that $\tau_n^2 \leq \sigma^2/n$ so that both the posterior mean and the posterior standard deviation converge to the values they would have with flat prior. The same effect holds for all regular models. The general version of this result is the Bayesian analogue of theorem 17.1: under the regularity conditions from definition 16.1, and for any continuous, positive prior density at the true

value, the posterior distribution is approximately

$$N\left(\hat{\theta}_n, \frac{1}{\mathcal{I}_n(\hat{\theta}_n)}\right) \quad \text{for large } n, \quad (28.4)$$

where $\hat{\theta}_n$ is the maximum likelihood estimator. Without proof, we state that this result holds. The proof is similar to the one for theorem 17.1: the log-posterior is $\ell(\theta) + \log \pi(\theta)$ up to a constant, the log-likelihood is approximately $\ell(\hat{\theta}_n) - \frac{1}{2}\mathcal{I}_n(\hat{\theta}_n)(\theta - \hat{\theta}_n)^2$ at its maximum, and this quadratic term grows like n while $\log \pi(\theta)$ does not. From (28.4) we know that the equal-tailed credible interval is approximately the Wald interval $\hat{\theta}_n \pm z_{1-\alpha/2} \text{se}(\hat{\theta}_n)$ of proposition 25.7, independently of the prior: for large sample size both approaches will determine the same interval with approximately frequentist coverage $1 - \alpha$.

The two approaches differ most when testing for a point null hypothesis $H_0: \theta = \theta_0$ against a two-sided alternative. In this case a Bayesian test requires a prior which puts a lump of probability on θ_0 but which spreads the remaining probability over the alternative. The exact distribution of the remaining probability over the alternative matters: if the prior is diffuse on the alternative, H_0 can have posterior probability of more than one half for a data set with p -value 0.01. We will not consider this case in detail, but we can conclude that small p -value and small posterior probability of H_0 are different quantities and that in this case the difference can be substantial.

Summary.

- A credible interval with credibility $1 - \alpha$ has posterior probability $1 - \alpha$: this is a probability statement about θ for the given data, taking the prior into account. The equal-tailed interval is based on posterior quantiles, the HPD interval is the shortest one, and for symmetric, unimodal posteriors both intervals coincide.
- For the normal mean with known variance and normal prior, the $1 - \alpha$ credible interval is $m_n \pm z_{1-\alpha/2} \tau_n$; for the flat prior this interval coincides with the confidence interval $\bar{x} \pm z_{1-\alpha/2} \sigma / \sqrt{n}$.
- A Bayesian test compares posterior probabilities of the hypotheses, or rejects H_0 if $\mathbb{P}(H_0|x)$ falls below a threshold. For two simple hypotheses, the posterior odds are the prior odds times the likelihood ratio, known as the Bayes factor.
- The level of a test and the confidence level of an interval are statements about the procedure applied to repeated samples, whereas the credibility of an interval is a statement about the parameter given one data set. For a location parameter with flat prior, the intervals, and the one-sided test and p -value, agree numerically while the statements do not have the same meaning.
- An informative prior will shift and shorten the Bayesian answer; as $n \rightarrow \infty$ the posterior is approximately $N(\hat{\theta}_n, 1/\mathcal{I}_n(\hat{\theta}_n))$ and then the credible and Wald intervals coincide. For point-null hypothesis testing, the two approaches can differ significantly.

Exercise 28.1. Let the posterior density $\pi(\theta|x)$ be continuous, strictly increasing on $(-\infty, \mu]$ and strictly decreasing on $[\mu, \infty)$.

1. Show that the HPD set H_c from definition 28.2 is an interval $[l, u]$ with $l < \mu < u$ and $\pi(l|x) = \pi(u|x)$.
2. Assume that the density is symmetric about μ , *i.e.* that $\pi(\mu + t|x) = \pi(\mu - t|x)$ for all t . Show that the HPD interval with credibility $1 - \alpha$ coincides with the equal-tailed interval $[q_{\alpha/2}, q_{1-\alpha/2}]$.
3. Now assume that the posterior density is continuous and strictly decreasing on $[0, \infty)$ and zero elsewhere. Show that the HPD set with credibility $1 - \alpha$ is the one-sided interval $[0, q_{1-\alpha}]$ and that this interval is shorter than the equal-tailed interval.

Exercise 28.2. In example 26.5, where a Bernoulli sample with $n = 10$ and $T = 7$ successes was combined with the prior $\text{Beta}(2, 2)$ to obtain the posterior distribution $\text{Beta}(9, 5)$ with mean $9/14 = 0.643$, the quantiles of the Beta distribution had to be found numerically. Using any statistical package, and using the `qbeta` function in R, we find the 0.025- and 0.975-quantiles of $\text{Beta}(9, 5)$ to be 0.386 and 0.861.

1. Write down the equal-tailed 95% credible interval for θ and compare the mid-point of this interval to the posterior mean.
2. Show that the posterior density has its maximum at $\theta = 2/3$, *i.e.* that the posterior distribution is skewed with the mode above the mean.
3. The HPD interval for $\text{Beta}(9, 5)$, found numerically, is $[0.401, 0.874]$. Using definition 28.2, explain why this interval is shifted towards the mode, compared to the equal-tailed interval. Also explain why the difference between the two intervals decreases as n increases.

Exercise 28.3. Assume that we have observed two successes in three Bernoulli trials with success probability θ , and the prior distribution of θ is flat on the interval $(0, 1)$.

1. Determine the posterior distribution of θ .
2. Calculate the posterior probability of $H_0: \theta \leq 1/2$ by integrating the posterior density. Also, calculate the posterior odds of H_0 against $H_1: \theta > 1/2$.
3. The prior odds of H_0 against H_1 were 1. By what factor have the data changed these odds? How is this factor related to the likelihood? Why is it not simply the ratio of the likelihoods at the two parameter values, as in (28.3)?
4. Write down the equations for the boundaries of the 95% equal-tailed credible interval and explain why these equations can only be solved numerically.

Exercise 28.4. In the situation of example 20.6, a sample of size $n = 25$ from $N(\theta, 4)$ has mean $\bar{x} = 10.9$. We want to compare the simple hypotheses $H_0: \theta = 10$ and $H_1: \theta = 11$.

1. Compute the Bayes factor in favour of H_0 .

2. Before the data were available, H_0 was three times as likely as H_1 . Compute the posterior odds and the posterior probability of H_0 .
3. In lecture 20 we have seen that the one-sided p -value 0.0122 can be found for the data. Compare your two numbers and explain why they can disagree, despite the fact that both values are against H_0 .
4. What value must the prior odds in favour of H_0 have for the posterior probability of H_0 to be greater than one half?

Exercise 28.5. Let $X_1, \dots, X_n$ be i.i.d. with $N(\theta, \sigma^2)$ where σ^2 is known, and let the prior be $N(m, \tau^2)$. Let $w = (1/\tau^2)/(1/\tau^2 + n/\sigma^2)$ be the weight of the prior mean in m_n . Consider the credible interval (28.2) to be a random interval, as introduced in lecture 25, where θ is fixed and the data are random.

1. Show that $\tau_n^2 = (1-w)\sigma^2/n$ and that under $\mathbb{P}_\theta$ the posterior mean satisfies $m_n - \theta \sim N(w(m - \theta), (1-w)^2\sigma^2/n)$.
2. Determine the coverage probability $\mathbb{P}_\theta(m_n - z_{1-\alpha/2}\tau_n \leq \theta \leq m_n + z_{1-\alpha/2}\tau_n)$ as a function of θ .
3. For $n = 16$, $\sigma = 2$, $\tau = 1$, $m = 12$ and $\alpha = 0.05$, determine the coverage probability for $\theta = 12$ and for $\theta = 10$ and discuss the result as $|\theta - m| \rightarrow \infty$.
4. Is the credible interval a confidence interval with level 0.95 in the sense of definition 25.1? What happens to the coverage probability as $n \rightarrow \infty$ with θ fixed?

Lecture 29

Standard Tests Revisited

The most commonly used statistical tests, the z -test, the t -test, the chi-squared test for a variance and the F -test, are often taught as recipes. In this lecture we will consider these tests again, from the perspective of the module: each of these tests is based on a pivot in the sense of lecture 25, and the test is performed by comparing a statistic to a quantile from a table. We will work through the one-sample t -test in detail, including how to read the tables, and will state the variance, paired, two-sample and F -tests in a similar way. We will conclude by giving a table which shows which test is appropriate for which question.

29.1 Tests from Pivots

Every test in this lecture is based on the following principle: the test statistic is known to follow a certain distribution under the null hypothesis, and the test rejects if the statistic exceeds a quantile of this distribution, taken from a table. Using the language of lecture 25, the test statistic is a pivot $Q(X, \theta)$ for the hypothesised value $\theta = \theta_0$ and the test is the one which, when inverted by theorem 25.6, gives the confidence interval based on this pivot.

We use the notation z_p , $t_{k,p}$ and $\chi_{k,p}^2$ for the quantiles from lecture 25, so that $\mathbb{P}(T \leq t_{k,p}) = p$ for $T \sim t_k$. A table shows these quantiles for a few values of p . For a one-sided test at level α , we need the $(1 - \alpha)$ -quantile, and for a two-sided test at the same level we need the $(1 - \alpha/2)$ -quantile. A p -value in the sense of definition 20.8 cannot be found exactly in a table, but can be bracketed between two tabulated levels; we will see an example of this approach in example 29.2, below.

The most basic example is the z -test from lecture 20: if $X_1, \dots, X_n \sim N(\mu, \sigma^2)$ are i.i.d. with known σ^2 , then $Z = \sqrt{n}(\bar{X} - \mu_0)/\sigma$ is standard normally distributed under $H_0: \mu = \mu_0$ and in examples 20.6 and 20.7 we can reject for $Z > z_{1-\alpha}$ or $|Z| > z_{1-\alpha/2}$ for the one-sided and two-sided alternative, respectively. The simplicity of this test is countered by the fact that the variance of the data is known. The rest of the lecture is concerned with how to replace Z when σ^2 is unknown.

29.2 The One-Sample t -Test

Let $X_1, \dots, X_n$ be i.i.d. $N(\mu, \sigma^2)$, where both parameters are unknown, and let $\bar{X}$ and S^2 be the sample mean and sample variance, respectively. If we replace σ in the definition of Z by the estimate S , we get the test statistic

$$T = \frac{\sqrt{n}(\bar{X} - \mu_0)}{S}. \quad (29.1)$$

In example 25.4 we have seen that for $\mu = \mu_0$ this is the ratio of the standard normal variable $\sqrt{n}(\bar{X} - \mu_0)/\sigma$ to the square root of $V/(n-1)$, where $V = (n-1)S^2/\sigma^2 \sim \chi_{n-1}^2$ by proposition A.3, and that the two values are independent of each other by corollary 10.10. Consequently, $T \sim t_{n-1}$ by definition A.4. This leads to the following tests.

Proposition 29.1. *Let $n \geq 2$ and let $X_1, \dots, X_n$ be i.i.d. $N(\mu, \sigma^2)$, where the parameters μ and σ^2 are unknown, $\mu_0 \in \mathbb{R}$ and let T be given by (29.1). Then the three tests of $H_0: \mu = \mu_0$, which reject if*

$$T > t_{n-1, 1-\alpha}, \quad T < -t_{n-1, 1-\alpha}, \quad \text{or} \quad |T| > t_{n-1, 1-\alpha/2},$$

for the alternatives $\mu > \mu_0$, $\mu < \mu_0$ and $\mu \neq \mu_0$, respectively, have size α , independent of the value of σ^2 .

Proof. Under H_0 we have $T \sim t_{n-1}$ for every value of σ^2 , as shown above. By the definition of the quantile we have $\mathbb{P}(T > t_{n-1, 1-\alpha}) = 1 - (1 - \alpha) = \alpha$. This is the size of the first test. The density of the t distribution is symmetric about zero and thus we have $\mathbb{P}(T < -t_{n-1, 1-\alpha}) = \mathbb{P}(T > t_{n-1, 1-\alpha}) = \alpha$ for the second test and for the third test we have $\mathbb{P}(|T| > t_{n-1, 1-\alpha/2}) = 2 \cdot \alpha/2 = \alpha$. Since these probabilities do not depend on σ^2 , the proof is complete. $\square$

These are the one-sided and two-sided *one-sample t -tests*. Since $t_{n-1, p} > z_p$ for every $p > 1/2$, the t -test requires a larger standardised deviation than the z -test before it can reject the hypothesis, and this is the price we pay for estimating σ . As n increases, this price decreases as the t_{n-1} distribution gets closer to the standard normal distribution. As in lecture 20, the one-sided tests have level α for the composite null hypotheses $\mu \leq \mu_0$ and $\mu \geq \mu_0$, since the probability of rejection increases as μ moves into the alternative.

Example 29.2. A quantity has nominal value 10, and ten measurements give

10.2, 9.7, 11.1, 10.8, 10.4, 9.9, 11.3, 10.6, 10.1, 10.9.

Assume that the measurements are i.i.d. $N(\mu, \sigma^2)$. We want to test $H_0: \mu = 10$ against $H_1: \mu \neq 10$ at level $\alpha = 0.05$. The sum of the observations is 105, so that $\bar{x} = 10.5$; the sum of the squared deviations from $\bar{x}$ is 2.52, so that $s^2 = 2.52/9 = 0.28$ and $s = 0.529$. The observed value of the test statistic is

$$t = \frac{\sqrt{10}(10.5 - 10)}{0.529} = 2.99.$$

The row corresponding to $n - 1 = 9$ degrees of freedom in the t -table shows that we have $t_{9, 0.95} = 1.833$, $t_{9, 0.975} = 2.262$, $t_{9, 0.99} = 2.821$ and $t_{9, 0.995} = 3.250$. For a two-sided test with level 0.05 we need $t_{9, 0.975} = 2.262$ and since $|t| = 2.99 > 2.262$

we can reject H_0 : the data are not compatible with a mean of 10. The same row also brackets the p -value. The observed value 2.99 falls between $t_{9,0.99}$ and $t_{9,0.995}$ and thus the two-sided p -value, $2\mathbb{P}(T \geq 2.99)$ for $T \sim t_9$, is between 0.01 and 0.02; the two-sided test also rejects at level 0.02 but not at level 0.01. If instead the question had been whether the measurements were high, the one-sided test for $H_0: \mu = 10$ against $H_1: \mu > 10$ at level 0.05 would have compared $t = 2.99$ to $t_{9,0.95} = 1.833$ and would have rejected H_0 , with a p -value between 0.005 and 0.01.

The t -test is the test obtained by applying the likelihood-ratio principle from lecture 23. In exercise 23.1 the generalised likelihood-ratio statistic for $H_0: \mu = \mu_0$, where σ^2 is unknown for both hypotheses, is found to be $-2 \log \Lambda = n \log(1 + T^2/(n-1))$, which is a strictly increasing function of T^2 . Thus, rejecting large values of $-2 \log \Lambda$ corresponds to the two-sided t -test. Following the general method of lecture 23, we would use the critical value from the χ_1^2 approximation of Wilks' theorem, theorem 23.3; however, in this case the null distribution of T is known exactly and the t_{n-1} quantile is the exact critical value for every n . As n gets large, the two approaches will be equal, since $n \log(1 + t^2/(n-1)) \rightarrow t^2$ and $t_{n-1,1-\alpha/2} \rightarrow z_{1-\alpha/2}$ and then $z_{1-\alpha/2}^2 = \chi_{1,1-\alpha}^2$.

29.3 The Chi-Squared Test for a Variance

The second question we can ask a normal sample, about the spread of the sample, is whether the variance equals a given value σ_0^2 . In proposition A.3 we have seen that the pivot $(n-1)S^2/\sigma^2 \sim \chi_{n-1}^2$ does not depend on the mean, and the test will follow the structure of the test from the previous section.

Proposition 29.3. *Let $n \geq 2$ and let $X_1, \dots, X_n$ be i.i.d. $N(\mu, \sigma^2)$, where μ and σ^2 are unknown, $\sigma_0^2 > 0$ and let*

$$V = \frac{(n-1)S^2}{\sigma_0^2}. \quad (29.2)$$

Each of the three tests for $H_0: \sigma^2 = \sigma_0^2$, which reject if

$$V > \chi_{n-1,1-\alpha}^2, \quad V < \chi_{n-1,\alpha}^2, \quad \text{or} \quad V \notin [\chi_{n-1,\alpha/2}^2, \chi_{n-1,1-\alpha/2}^2],$$

for the alternatives $\sigma^2 > \sigma_0^2$, $\sigma^2 < \sigma_0^2$ and $\sigma^2 \neq \sigma_0^2$, respectively, has size α , whatever the value of μ .

Proof. Under H_0 we have $V = (n-1)S^2/\sigma^2 \sim \chi_{n-1}^2$ by proposition A.3, for every value of μ . The first two rejection probabilities are $1 - (1 - \alpha) = \alpha$ and α by the definition of the quantiles. For the third, the events $V < \chi_{n-1,\alpha/2}^2$ and $V > \chi_{n-1,1-\alpha/2}^2$ are disjoint and each have probability $\alpha/2$, so the rejection probability is α . This completes the proof. $\square$

Unlike the t distribution, the chi-squared distribution is not symmetric and so we need different values from the table for the two tails of a two-sided test. This is the inverse of the confidence interval from example 25.5. Similarly, the one-sided test against $\sigma^2 > \sigma_0^2$ has level α for the composite null hypothesis $\sigma^2 \leq \sigma_0^2$, since the probability of the test to reject the null hypothesis increases with σ^2 . This is exercise 29.2.

The numerical carrying-out of the test is also part of that exercise, and it follows the pattern of example 29.2 with the chi-squared table in place of the t -table. Note that tests about a variance have low power for small samples, because S^2 is a much more variable estimator than $\bar{X}$.

29.4 Paired and Two-Sample Problems

The tests so far have considered a single sample. Here we will consider three more standard tests, which compare two sets of measurements each, and are again pivot tests.

The first of these three tests is the *paired t -test*. Assume that we have observed data in pairs (X_i, Y_i) , $i = 1, \dots, n$, e.g. before and after measurements for the same subject. Since the two measurements for one subject are not independent, we can reduce the data for each pair to the difference $D_i = X_i - Y_i$, model the differences as i.i.d. $N(\mu_D, \sigma_D^2)$ and then apply proposition 29.1 with $\mu_0 = 0$: the test statistic is $T = \sqrt{n} \bar{D} / S_D$, constructed from the sample mean and variance of the differences, and is t_{n-1} distributed under $H_0: \mu_D = 0$. The model is only for the differences, as the example in exercise 29.3 shows.

The second test is the *two-sample t -test*. Here we assume that $X_1, \dots, X_m$ are i.i.d. $N(\mu_X, \sigma^2)$ and that $Y_1, \dots, Y_n$ are i.i.d. $N(\mu_Y, \sigma^2)$, where the two samples are independent of each other, and where we assume that the unknown variances of the two populations are the same. We consider $H_0: \mu_X = \mu_Y$. The estimator for the common variance is now

$$S_p^2 = \frac{(m-1)S_X^2 + (n-1)S_Y^2}{m+n-2},$$

and the test statistic is

$$T = \frac{\bar{X} - \bar{Y}}{S_p \sqrt{1/m + 1/n}} \sim t_{m+n-2} \quad \text{under } H_0. \quad (29.3)$$

The distribution of this test statistic is similar to the distribution of the test statistic for the one-sample case: under H_0 the numerator is $N(0, \sigma^2(1/m+1/n))$, $(m+n-2)S_p^2/\sigma^2$ is the sum of two independent χ_{m-1}^2 and χ_{n-1}^2 variables and thus is χ_{m+n-2}^2 by definition A.2, and numerator and denominator are independent of each other by corollary 10.10, applied to each of the two samples, and the independence of the samples. Thus, the test can be used to test H_0 against $\mu_X \neq \mu_Y$, rejecting if $|T| > t_{m+n-2, 1-\alpha/2}$, and the one-sided variants of the test use $t_{m+n-2, 1-\alpha}$.

The third test is the *F -test for two variances*. In the same two-sample setting, but now with variances σ_X^2 and σ_Y^2 which are allowed to be different, the ratio $(S_X^2/\sigma_X^2)/(\sigma_Y^2/S_Y^2)$ is the ratio of two independent chi-squared variables, divided by the degrees of freedom. Thus, it is $F_{m-1, n-1}$ distributed by definition A.4. Under $H_0: \sigma_X^2 = \sigma_Y^2$ the unknown variances cancel and the statistic $F = S_X^2/S_Y^2$ is $F_{m-1, n-1}$ distributed. The two-sided test rejects if F does not fall into the interval $[F_{m-1, n-1, \alpha/2}, F_{m-1, n-1, 1-\alpha/2}]$, where $F_{k, l, p}$ is the p -quantile of $F_{k, l}$. Normally, tables only give the upper quantiles and the lower quantiles can be found by using $F_{k, l, p} = 1/F_{l, k, 1-p}$. This test is also used to check the assumption of equal variances in (29.3).

29.5 Which Test?

Table 29.1 summarises the six tests. The choice between them is determined by the following three questions: Is the parameter in question a mean or a variance? Is there one sample or two? If two samples, are the observations paired or independent? All six tests assume that the data are normally distributed. In lecture 30 we will use simulation to study what happens when this assumption is violated.

Question	H_0	Statistic	Null distribution
one mean, σ^2 known	$\mu = \mu_0$	$\sqrt{n}(\bar{X} - \mu_0)/\sigma$	$N(0, 1)$
one mean, σ^2 unknown	$\mu = \mu_0$	$\sqrt{n}(\bar{X} - \mu_0)/S$	t_{n-1}
paired differences	$\mu_D = 0$	$\sqrt{n}\bar{D}/S_D$	t_{n-1}
one variance, μ unknown	$\sigma^2 = \sigma_0^2$	$(n-1)S^2/\sigma_0^2$	χ_{n-1}^2
two means, common σ^2	$\mu_X = \mu_Y$	$\frac{\bar{X} - \bar{Y}}{S_p \sqrt{1/m + 1/n}}$	t_{m+n-2}
two variances	$\sigma_X^2 = \sigma_Y^2$	S_X^2/S_Y^2	$F_{m-1, n-1}$

Table 29.1: The standard tests for normal data. The pivot for each test is evaluated under H_0 , and the final column gives the table from which the critical value can be found (using the $(1 - \alpha)$ -quantile for a one-sided and the $(1 - \alpha/2)$ -quantile (together with the $(\alpha/2)$ -quantile for the chi-squared and F tests) for a two-sided alternative).

Summary.

- Each of the standard tests for normal data is based on the pivot from lecture 25, evaluated for the null value, and compared to a quantile of the known null distribution. The test is the inverse of the corresponding confidence interval.
- The one-sample t -test uses $T = \sqrt{n}(\bar{X} - \mu_0)/S \sim t_{n-1}$ under $H_0: \mu = \mu_0$ and rejects for $T > t_{n-1, 1-\alpha}$ (one-sided) or $|T| > t_{n-1, 1-\alpha/2}$ (two-sided). It is the likelihood-ratio test with an exact critical value.
- The chi-squared test for a variance uses $V = (n-1)S^2/\sigma_0^2 \sim \chi_{n-1}^2$ under $H_0: \sigma^2 = \sigma_0^2$. A two-sided test requires both lower and upper quantiles, since the distribution is not symmetric.
- The paired t -test is a one-sample t -test for the differences. Similarly, the two-sample t -test with pooled variance and the F -test for two variances follow this pattern, using the t_{m+n-2} and $F_{m-1, n-1}$ distributions.
- The table lists only a few quantiles and often the p -value is bracketed between two tabulated levels.

Exercise 29.1. Eight independent measurements of a quantity with nominal value 20 are given by

19.1, 18.6, 20.3, 19.4, 18.9, 19.8, 18.2, 19.7.

Assume that the measurements are i.i.d. $N(\mu, \sigma^2)$, where both parameters are unknown. From the t -table we find $t_{7,0.95} = 1.895$, $t_{7,0.975} = 2.365$, $t_{7,0.99} = 2.998$ and $t_{7,0.995} = 3.499$.

1. Determine $\bar{x}$, s^2 and the observed value of the test statistic T from (29.1) for $\mu_0 = 20$.
2. For $H_0: \mu = 20$, test $H_1: \mu < 20$ at levels 0.05 and 0.01. Bracket the p -value.
3. For $H_0: \mu = 20$, test $H_1: \mu \neq 20$ at the same two levels and bracket the p -value.
4. Calculate the 99% t -interval for μ from example 25.4 and discuss how it relates to your answer in the previous part.

Exercise 29.2. A sample of $n = 15$ i.i.d. $N(\mu, \sigma^2)$ observations has sample variance $s^2 = 2.9$, and the process generating the sample is known to have variance at most 1.5. From the chi-squared table we find $\chi_{14,0.025}^2 = 5.629$, $\chi_{14,0.05}^2 = 6.571$, $\chi_{14,0.95}^2 = 23.685$, $\chi_{14,0.975}^2 = 26.119$ and $\chi_{14,0.99}^2 = 29.141$.

1. Test $H_0: \sigma^2 \leq 1.5$ against $H_1: \sigma^2 > 1.5$ at levels 0.05 and 0.01. Bracket the p -value.
2. Test $H_0: \sigma^2 = 1.5$ against $H_1: \sigma^2 \neq 1.5$ at level 0.05.
3. Show that the power function of the one-sided test from proposition 29.3 is given by $\beta(\sigma^2) = \mathbb{P}(W > \chi_{n-1,1-\alpha}^2 \sigma_0^2 / \sigma^2)$ where $W \sim \chi_{n-1}^2$, and that this power function is increasing in σ^2 . Deduce that the test has size α for the composite null hypothesis $\sigma^2 \leq \sigma_0^2$.
4. Using the power function, explain how it can happen that the one-sided test from part (a) and the two-sided test from part (b) can give different results at the same level.

Exercise 29.3. The following data give the reaction times of six subjects before and after a training period.

subject	1	2	3	4	5	6
before (x_i)	72	65	80	58	69	75
after (y_i)	68	66	74	55	63	70

We want to know whether the training reduced the mean reaction time. From the t -table we find $t_{5,0.95} = 2.015$, $t_{5,0.975} = 2.571$, $t_{5,0.99} = 3.365$ and $t_{10,0.95} = 1.812$.

1. Write down the model and hypotheses for the paired t -test and carry out the test at levels 0.05 and 0.01.
2. Determine a 95% confidence interval for the mean reduction in reaction time.
3. Your colleague, not knowing that the two rows in the data frame give paired data, performs a two-sample t -test (29.3). The sample variances of the two rows are $s_X^2 = 59.77$ and $s_Y^2 = 42.80$, respectively. What is the test statistic and the decision of your colleague at level 0.05? Justify your answer by showing that the paired analysis is appropriate, and that the two analyses are different.

Exercise 29.4. In example 25.4 we have observed a sample of $n = 16$ values, where $\bar{x} = 10.3$ and $s = 2.1$, and we have found the 95% t -interval $[9.18, 11.42]$ for μ . In example 25.5, for the same data, we have found the 95% interval $[2.41, 10.56]$ for σ^2 . In exercise 25.4 we have seen that the hypothesis $H_0: \mu = 9$ was rejected at level 0.05, using theorem 25.6 alone. Now use the values $t_{15,0.95} = 1.753$, $t_{15,0.975} = 2.131$, $\chi_{15,0.025}^2 = 6.262$ and $\chi_{15,0.975}^2 = 27.488$.

1. Perform a two-sided t -test for $H_0: \mu = 9$ and state the decision. Repeat the test for $H_0: \mu = 11.5$ and check that the result is consistent with the interval.
2. For $H_0: \sigma^2 = \sigma_0^2$, perform a two-sided test at level 0.05 for $\sigma_0^2 = 2$ and $\sigma_0^2 = 9$, either from the interval alone or by computing the test statistic (29.2).
3. Perform a one-sided test for $H_0: \mu = 9.5$ against $H_1: \mu > 9.5$ at level 0.05. Why cannot the two-sided interval be used to answer this test? Which confidence set does this correspond to?

Lecture 30

Do the Standard Tests Work?

This is the R session for week 10. We perform a t -test for a real data set and derive the Bayesian answer for the same question. Then we consider three cases where a correct test can still be misleading: selective reporting of results, repeated testing and when the model assumption is violated. As before, appendix B contains a written reference for the material covered in the session.

30.1 Real Data, Two Answers

The data set `sleep`, built into R, contains the amount of extra sleep (in hours) for ten patients, when treated with two different drugs. The two measurements for each patient are paired, and thus we can combine the two measurements into one pair difference as shown in the third row of table 29.1 and then perform a one-sample t -test. The command `t.test` can be used to perform the test which example 29.2 performed manually. Before you try the command, look at the ten differences and consider whether the mean of these differences is positive, and whether the two-sided p -value is likely to be less than 0.05.

```
head(sleep, 3)
d <- with(sleep, extra[group == 2] - extra[group == 1])
d
t.test(d)
```

```
##   extra group ID
## 1   0.7      1  1
## 2  -1.6      1  2
## 3  -0.2      1  3
## [1] 1.2 2.4 1.3 1.3 0.0 1.0 1.8 0.8 4.6 1.4
##
## One Sample t-test
##
## data:  d
## t = 4.0621, df = 9, p-value = 0.002833
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
##  0.7001142 2.4598858
## sample estimates:
## mean of x
##      1.58
```

Every number in the output is a number we have seen before: the statistic $t = \sqrt{10} \cdot 1.58/1.23 = 4.06$ is compared with the t_9 distribution, the p -value 0.0028 is $2\mathbb{P}(T \geq 4.06)$ as in definition 20.8, and the interval $[0.70, 2.46]$ is the t -interval from example 25.4.

The Bayesian answer. We can model the differences as $N(\theta, \sigma^2)$. To apply proposition 26.6 we pretend that σ is known and equal to $s = 1.23$, *i.e.* we ignore the uncertainty in σ as in a z -test; a prior distribution on σ is beyond the scope of this module. Here we assume that the prior distribution is $N(0, 10^2)$. This prior is large enough to be overridden by ten observations. Before we execute the code, predict the posterior probability of $H_0: \theta \leq 0$: is it larger or smaller than the one-sided p -value 0.0014, output by `t.test` with `alternative = "greater"` (see the script `R/S30-sleep.R` to compute the probability)?

```
n      <- length(d)
dbar   <- mean(d)
s      <- sd(d)
m      <- 0
tau    <- 10
prec   <- 1 / tau^2 + n / s^2
m_n    <- (m / tau^2 + n * dbar / s^2) / prec
tau_n  <- sqrt(1 / prec)
c(mean = m_n, sd = tau_n)
pnorm(0, mean = m_n, sd = tau_n)
1 - pnorm(1, mean = m_n, sd = tau_n)
1 - pnorm(sqrt(n) * dbar / s)
```

```
##      mean      sd
## 1.5776132 0.3886648
## [1] 2.46355e-05
## [1] 0.9313799
## [1] 2.431373e-05
```

The posterior distribution is $N(1.578, 0.389^2)$, and the probability $\mathbb{P}(\theta \leq 0|d)$ is $2.5 \cdot 10^{-5}$, as found in example 28.5. The final line, the one-sided p -value of the z -test with $\sigma = s$ assumed known, coincides with the posterior probability $\mathbb{P}(\theta \leq 0|d)$ to two significant figures, as discussed in section 28.3 for the case of nearly flat prior. The t -test's p -value is sixty times larger, reflecting the fact that the t_9 tail is heavier than the normal tail. This is the cost of not knowing σ . The middle line of the output answers a question which cannot be answered by the p -value: what is the posterior probability that the effect is larger than one hour of sleep? The answer is 0.93.

30.2 The Replication Crisis in Miniature

A test at level 0.05 has a 5% chance of rejecting a true null hypothesis, which is harmless for one test but not when many tests are run and only the significant ones are reported. We simulate twenty t -tests, each on a sample of size 20 from $N(0, 1)$, so that every null hypothesis is true. By exercise 20.4 each p -value is uniform on $(0, 1)$ under H_0 . Before you run the code, predict how many of the twenty p -values will be less than or equal to 0.05.

```
n <- 20
m <- 20
set.seed(2026)
p <- replicate(m, t.test(rnorm(n))$p.value)
round(p, 3)
sum(p < 0.05)
```

```
## [1] 0.083 0.151 0.194 0.959 0.594 0.601 0.070 0.447 0.656 0.275
      0.362 0.820
## [13] 0.352 0.735 0.489 0.308 0.511 0.740 0.845 0.301
## [1] 0
```

The expected count is one, and this study happened to produce none. The number which matters is the *family-wise error rate*, the probability that a study of twenty tests reports at least one false discovery. Since the tests are independent, the exact value is $1 - 0.95^{20}$. Predict it, then check by simulation over 2000 studies.

```
study <- function() any(replicate(m, t.test(rnorm(n))$p.value) < 0.05)
mean(replicate(2000, study()))
1 - 0.95^m
```

```
## [1] 0.6315
## [1] 0.6415141
```

Nearly two studies in three would find a significant result for nothing. This can happen either by testing twenty hypotheses at once, or by testing one hypothesis twenty times. The following function simulates a study where a sample of size 100 is tested after every ten observations, and the procedure is continued until $p < 0.05$: this is the case of the researcher who “collects a few more data points and looks again”. Before you try the function, predict the rejection rate for this procedure under H_0 .

```
peek <- function() {
  x <- rnorm(100)
  for (k in seq(10, 100, by = 10)) {
    if (t.test(x[1:k])$p.value < 0.05) return(TRUE)
  }
  FALSE
}
mean(replicate(2000, peek()))
```

```
## [1] 0.206
```

Ten looks at the same data raise the type I error rate from 5% to 21%: each single test has level 0.05, but the procedure formed by choosing when to stop does not, and the reported p -value does not know how it was chosen.

The solution to this problem is to adjust the significance level. The two most common corrections are named here for reference, but I won't test your knowledge of these corrections in the exam. The *Bonferroni correction* is very simple: it just rejects $p < \alpha/m$, so that the family-wise error rate is at most α . This method is very conservative. The *Benjamini–Hochberg procedure* controls

the *false discovery rate*, *i.e.* the expected proportion of false rejections, by sorting the p -values in order $p_{(1)} \leq \dots \leq p_{(m)}$, and then finding the largest i such that $p_{(i)} \leq \alpha i/m$. The i smallest p -values are then rejected. The `p.adjust` function can be used to perform either correction, and returns the adjusted p -values for comparison with α . The following box contains a script for a study with 200 tests, 180 with a true null hypothesis and 20 with a real effect $\theta = 0.8$. The lines are in random order, so you will need to arrange them to form a working script before you can try the code.

```
table(rejected = p.adjust(p, "BH") < 0.05, effect)
theta  <- c(rep(0, 180), rep(0.8, 20))
p <- sapply(theta, function(th) t.test(rnorm(n, mean = th))$p.value)
table(rejected = p < 0.05, effect)
n      <- 20
m      <- 200
set.seed(2026)
effect <- theta > 0
table(rejected = p.adjust(p, "bonferroni") < 0.05, effect)
```

As in lectures 12 and 24, all names must be defined before they are used: `n`, `m` and `theta` must come first, `effect` depends on `theta`, the seed for the random draw in the `sapply` line must be set before the line itself, and the three tables can be given in any order. The command `table` can be used to cross-classify the decision with the truth. This will give the 2×2 table from lecture 20. Before you try the reconstructed script, predict the number of false discoveries among the 180 true null hypotheses in the first table, and predict which of the two methods, Bonferroni or Benjamini–Hochberg, will result in more of the 20 real effects being retained.

```
##          effect
## rejected FALSE TRUE
##   FALSE   170    0
##   TRUE    10   20
##          effect
## rejected FALSE TRUE
##   FALSE   180   14
##   TRUE     0    6
##          effect
## rejected FALSE TRUE
##   FALSE   180    3
##   TRUE     0   17
```

Without correction there are 30 discoveries, of which 10 are false: one in three of the reported findings is wrong, despite all tests being performed correctly. The Bonferroni method eliminates all ten false discoveries, but also 14 of the 20 real effects. The Benjamini–Hochberg method also eliminates all ten false discoveries but manages to keep 17 of the real effects: the guarantee is weaker, but the method finds many more results. Figure 30.1, generated by the script `R/S30-forking.R`, shows the 40 smallest p -values and their rank.

In the figure only six of the real effects are below the Bonferroni line, which does not move with the rank; the Benjamini–Hochberg line rises with i , and the last point which is below, the seventeenth, is just before the open circles start at rank 19.




Figure 30.1: The 40 smallest of 200 p -values against their rank on a logarithmic scale: real effects (filled circles) and true null hypotheses (open circles), with the raw level 0.05 (dashed), the Bonferroni level 0.05/200 (dotted) and the Benjamini–Hochberg line 0.05 i /200 (solid).

30.3 A Stress Test

Every test in table 29.1 assumes that the data are normally distributed. In lecture 12 we have seen that for distributions with heavy tails the central limit theorem takes effect only slowly, and section 17.5 lists further situations in which large-sample approximations are unreliable. We now test the t -test for *contaminated normal* data, *i.e.* data which consist of $N(\theta, 1)$ with probability 0.95 and $N(\theta, 100)$ with probability 0.05. This could, for example, happen if one in twenty of the measurements is made using a faulty instrument. The resulting distribution is still symmetric around θ and has variance 5.95. Before you try the code, consider the following questions: at $n = 20$ and level 0.05, does the size of the t -test stay close to 0.05? What happens to the power of the test at $\theta = 0.5$?

```
rcontam <- function(n, theta, eps = 0.05) {
  big <- rbinom(n, size = 1, prob = eps) == 1
  theta + rnorm(n, sd = ifelse(big, 10, 1))
}
n <- 20
decide <- function(theta, contaminated) {
  x <- if (contaminated) rcontam(n, theta) else rnorm(n, mean = theta)
  c(t = t.test(x)$p.value < 0.05, wilcoxon = wilcox.test(x)$p.value <
    0.05)
}
set.seed(2026)
rates <- sapply(c(0, 0.5, 1), function(theta) {
  c(clean = rowMeans(replicate(5000, decide(theta, FALSE))),
    contaminated = rowMeans(replicate(5000, decide(theta, TRUE))))
})
colnames(rates) <- c("theta = 0", "theta = 0.5", "theta = 1")
round(rates, 3)
```


##	theta = 0	theta = 0.5	theta = 1
## clean.t	0.051	0.561	0.989
## clean.wilcoxon	0.049	0.536	0.987
## contaminated.t	0.033	0.330	0.698
## contaminated.wilcoxon	0.048	0.460	0.940

The function `decide` also performs the Wilcoxon signed-rank test, which uses only signs and ranks of the data and assumes symmetry instead of normality. This test is not part of the module under consideration here, but is included for comparison. The first column of the table shows that the size of the t -test is preserved under contamination, and in fact decreases to 0.033: a single ‘wild’ observation has the effect of inflating S more than $\bar{X}$ would, and thus the test statistic is reduced. The power columns tell a different story: for $\theta = 1$ the t -test on clean data rejects 99% of the samples, whereas the t -test on contaminated data rejects only 70% of them. In contrast, the Wilcoxon test, which does not lose much power for clean data, still retains 94%. The t -test is not lying about the null hypothesis, but is failing to detect the effects present in the data. Figure 30.2, generated by the script `R/S30-stress.R`, shows the three power curves as a function of the grid of alternatives.




Figure 30.2: Estimated power functions at level 0.05 for 2000 samples of size $n = 20$, for the t -test on clean, normally distributed data (filled circles), the t -test on contaminated data (open circles), and the Wilcoxon signed-rank test on contaminated data (crosses).

The contaminated t -test has the correct size but the wrong power curve; if the contamination was asymmetric, the size would change as well. The moral of the story, as in lecture 24, is that when we are not sure about the model, we should simulate using a plausible alternative model, and should consider both columns of the table.

30.4 A Task to Take Home

Seed the generator with your student ID. First repeat the Bayesian analysis of the sleep data with the priors $N(0, 1)$ and $N(2, 0.5^2)$ and in two sentences state

the change in the posterior probability of $\theta \leq 0$. Then modify the function `peek` to test every five observations and estimate the rejection rate for H_0 from 2000 runs. Compare the estimated rate with 0.206. Finally, repeat the stress test with $n = 50$ and a contamination probability of 0.1 and create a table of sizes and powers for the test. Write a short paragraph about whether a larger sample size would improve the t -test in this example.

Before the next session, you must first understand the commands used in this session. Type in the commands yourself, and make sure that you get the same results as shown in the text.

Summary.

- The written reference for the commands used in this session is appendix B.
- The command `t.test` can be used to get the test statistic, the p -value and the t -interval for the data from lectures 25 and 29. The Bayesian answer from proposition 26.6, using a weak prior, gives the same result. In addition, the Bayesian answer can be used to get the probability of an effect of a given size.
- Under H_0 the p -value is uniformly distributed, *i.e.* twenty independent tests will produce a false discovery with probability 0.64. Similarly, repeatedly testing the same growing data set at ten points has a rejection probability of 0.21. The Bonferroni method controls the family-wise error rate by reducing the significance level of each test, but this method will miss most of the real effects. The Benjamini–Hochberg method controls the false discovery rate and thus preserves most of the real effects. Both methods can be implemented using the function `p.adjust`.
- For contaminated normal data the t -test maintains its size but loses a lot of power. A test which is not appropriate for the given data can fail in either of the two columns of the test output. Simulation under the alternative can be used to determine which of the two columns is more strongly affected.

Lecture 31

The Exponential Model: From Start to Finish

In this final lecture we will consider the exponential distribution with unknown rate. We will follow the approach a statistician would take when considering this model: we will consider the model and the data, how to estimate the parameter, how well the estimate is performing, what intervals and tests we can use, and how things change when we consider Bayesian inference. The individual steps of this approach are based on results we have seen in previous lectures, and the aim of this lecture is to see how these results can be combined to consider an example. We choose the exponential distribution because this is the running example of the module, and because most of the individual steps in the analysis should be already known to you from lectures and exercises.

31.1 The Model and the Data

Assume that the lifetime of a component of a machine is exponentially distributed with rate $\theta > 0$. Then the lifetimes $X_1, \dots, X_n$ of n components are i.i.d. with density $f(x; \theta) = \theta e^{-\theta x}$ for $x \geq 0$. The mean lifetime is $\mathbb{E}_\theta(X_i) = 1/\theta$, the variance is $1/\theta^2$ (table A.2). In example 7.4 we have seen that this is a one-parameter exponential family with natural statistic $T(x) = x$ and natural parameter $\eta(\theta) = -\theta$. Thus, by lemma 7.5, the joint density of the sample is

$$f(x_1, \dots, x_n; \theta) = \theta^n \exp\left(-\theta \sum_{i=1}^n x_i\right)$$

for $x_1, \dots, x_n \geq 0$. Since this density depends on the data only through $T = \sum_{i=1}^n X_i$, we can use the factorisation theorem 8.2 to conclude that T is sufficient for θ . Since the family has full rank in the sense of definition 7.10 (the natural parameter space is the open half-line), we can use proposition 10.5 to conclude that T is also complete. Thus, from lectures 7 to 10 we know that all our results can be based on T alone, and that any unbiased function of T will be the best unbiased estimator for the corresponding quantity.

As data we have observed $n = 10$ lifetimes (in years) of the components, and the sum of these values is $\sum_i x_i = 25.0$, *i.e.* we have $\bar{x} = 2.5$. The manufacturer claims that the mean lifetime is two years, *i.e.* $\theta_0 = 0.5$, and we want to test whether the components of the machine last longer than the manufacturer claims.

31.2 Point Estimation

The log-likelihood is $\ell(\theta) = n \log \theta - \theta T$, with derivative $\ell'(\theta) = n/\theta - T$ and second derivative $-n/\theta^2 < 0$, so the likelihood equation has the unique root $\hat{\theta} = n/T = 1/\bar{X}$ and this root is the maximum: $\hat{\theta}$ is the MLE for θ (example 5.2). For our data we find $\hat{\theta} = 10/25 = 0.4$. By the invariance of the MLE, theorem 5.6, the MLE for the mean lifetime $1/\theta$ is $\bar{X} = 2.5$ years.

The MLE for the rate is biased: exercise 2.3 shows $\mathbb{E}_\theta(\hat{\theta}) = n\theta/(n-1)$, so that $\hat{\theta}$ overestimates θ by a factor of $n/(n-1)$. The corrected estimator $\tilde{\theta} = (n-1)/T$ is unbiased. Since $\tilde{\theta}$ is an unbiased function of the complete sufficient statistic T , the Lehmann–Scheffe theorem 10.4 shows that this is the UMVUE for θ . The variance of the UMVUE can be found as $\theta^2/(n-2)$ (exercise 5.3). For the data considered here, we find $\tilde{\theta} = 9/25 = 0.36$.

What is the optimal performance of an estimator? From example 11.8 we know that the Fisher information of one observation is $\mathcal{I}(\theta) = 1/\theta^2$ and thus we have $\mathcal{I}_n(\theta) = n/\theta^2$. Since the model is regular (exercise 16.1), the Cramer–Rao inequality (theorem 13.2) gives the bound θ^2/n for the variance of any unbiased estimator. The variance of the UMVUE is $\theta^2/(n-2)$, which is larger than the bound and by proposition 13.5 no unbiased estimator can attain the bound, since the bound is only attained for an affine function of the natural statistic T and no such function is unbiased for the rate. The efficiency $(n-2)/n$ of $\tilde{\theta}$ tends to 1 as $n \rightarrow \infty$. The situation for the mean $1/\theta$ is different: $\bar{X}$ is unbiased with variance $1/(n\theta^2)$. This variance equals the bound $1/\mathcal{I}_n(\mu)$ for the mean (exercise 11.5) and thus $\bar{X}$ is efficient.

For large samples, theorems 16.3 and 17.1 apply, because the regularity conditions are satisfied: $\hat{\theta}_n$ is consistent and $\sqrt{n}(\hat{\theta}_n - \theta) \xrightarrow{d} N(0, \theta^2)$. From corollary 17.3 we know that the standard error is

$$\text{se}(\hat{\theta}) = \frac{1}{\sqrt{\mathcal{I}_n(\hat{\theta})}} = \frac{\hat{\theta}}{\sqrt{n}}.$$

For the data at hand we get $\text{se}(\hat{\theta}) = 0.4/\sqrt{10} = 0.126$. This is the number we would expect to see next to the estimate.

31.3 Interval Estimation

Proposition 25.7 turns the standard error into the approximate 95% confidence interval $\hat{\theta} \pm 1.96 \text{ se}(\hat{\theta})$. For the data considered here, this confidence interval is 0.4 ± 0.248 , *i.e.* $[0.152, 0.648]$. The interval is symmetric about the estimate, and the probability of the interval covering the correct value is only approximately 95%, since the interval is based on the normal limit at $n = 10$.

For the exponential distribution we can do better, because an exact pivot is available. In exercise 13.1, written there in terms of the mean $1/\theta$, we have seen that $2\theta T \sim \chi_{2n}^2$ for every value of θ , so $Q = 2\theta T$ is a pivot in the sense of lecture 25. Solving $\chi_{2n, \alpha/2}^2 \leq 2\theta T \leq \chi_{2n, 1-\alpha/2}^2$ for θ gives the exact interval

$$\left[\frac{\chi_{2n, \alpha/2}^2}{2T}, \frac{\chi_{2n, 1-\alpha/2}^2}{2T} \right]$$

with coverage exactly $1 - \alpha$ for every n . With $\chi_{20,0.025}^2 = 9.591$ and $\chi_{20,0.975}^2 = 34.170$ we find the interval $[9.591/50, 34.170/50] = [0.192, 0.683]$. The exact interval is shifted to the right relative to the Wald interval and is not symmetric about 0.4: it makes use of the fact that the sampling distribution of $\hat{\theta}$ is skewed, which the normal approximation does not. Both intervals are 95% confidence intervals in the sense of definition 25.1: the statement is about the procedure, and for the interval in front of us the parameter is either inside the interval or it is not.

31.4 Testing

We now test the manufacturer's claim. We take $H_0: \theta = \theta_0 = 0.5$ and start, as in lecture 22, with a simple alternative $H_1: \theta = \theta_1$ where $\theta_1 < \theta_0$ (a longer mean lifetime). The likelihood ratio is

$$\frac{f(x; \theta_1)}{f(x; \theta_0)} = \left(\frac{\theta_1}{\theta_0}\right)^n \exp((\theta_0 - \theta_1)T),$$

which is an increasing function of T , since $\theta_0 - \theta_1 > 0$. By the Neyman–Pearson lemma, theorem 22.2, the most powerful test at level α rejects H_0 if $T > c$, where c is chosen such that $\mathbb{P}_{\theta_0}(T > c) = \alpha$. Under H_0 the pivot gives $2\theta_0 T \sim \chi_{2n}^2$, and thus $c = \chi_{2n,1-\alpha}^2 / (2\theta_0)$. For $\alpha = 0.05$ we find $c = 31.410/1 = 31.41$. The critical region $\{T > 31.41\}$ does not depend on θ_1 and thus, by proposition 22.4, the same test is most powerful against all alternatives $\theta < \theta_0$. Since the observed value $T = 25.0$ is less than the critical value, we do not reject H_0 at the 5% level: the data do not show that the components have a longer lifetime than claimed. The p -value from definition 20.8 is $\mathbb{P}_{\theta_0}(T \geq 25.0) = \mathbb{P}(\chi_{20}^2 \geq 25.0) = 0.20$.

The power function of definition 20.4 can be found by evaluating the same pivot at the alternative: $\beta(\theta) = \mathbb{P}_{\theta}(T > c) = \mathbb{P}(\chi_{2n}^2 > 2\theta c)$, so for example $\beta(0.25) = \mathbb{P}(\chi_{20}^2 > 15.71) = 0.73$. If the true mean lifetime of the components is four years, then a sample of ten components has probability 0.73 of detecting this.

For the two-sided alternative $\theta \neq \theta_0$ there is no uniformly most powerful test, by the argument of proposition 22.5: the two one-sided tests reject on opposite tails of T . We use the generalised likelihood-ratio test of definition 23.1 instead. The statistic is

$$W = 2(\ell(\hat{\theta}) - \ell(\theta_0)) = 2n \left(\log \frac{\hat{\theta}}{\theta_0} - 1 + \frac{\theta_0}{\hat{\theta}} \right),$$

and by Wilks' theorem 23.3 it is approximately χ_1^2 -distributed under H_0 , since one parameter is fixed. For our data, $W = 20(\log 0.8 - 1 + 1.25) = 0.54$, well below $\chi_{1,0.95}^2 = 3.84$, so again H_0 is not rejected. The exercises at the end of this lecture verify the formula for W .

31.5 The Bayesian Answer

Since we can assume θ to be random with prior density $\pi(\theta)$, and since the likelihood function is $\theta^n e^{-\theta T}$ as a function of θ , a prior of the same shape as the

likelihood, namely the $\text{Gamma}(a, b)$ density $\pi(\theta) \propto \theta^{a-1}e^{-b\theta}$, is conjugate. This is a special case of theorem 26.2: the posterior distribution is given by

$$\pi(\theta|x) \propto \theta^{a-1}e^{-b\theta} \theta^n e^{-\theta T} = \theta^{a+n-1}e^{-(b+T)\theta},$$

which is again the $\text{Gamma}(a + n, b + T)$ distribution. The situation in the example fits the result of proposition 26.4: the update process increases a by n and b by T , so that the effect of the prior is as if we had observed a additional values with total lifetime b . In the example of the manufacturer's claim about the lifetime of the components, we can assume that the claim is equivalent to having observed two more prior observations with mean lifetime two years, *i.e.* $a = 2$ and $b = 4$. The prior mean is then $a/b = 0.5$. The posterior distribution is then $\text{Gamma}(12, 29)$ and, using theorem 26.7, the Bayes estimator for θ , using squared error loss, is given by

$$\mathbb{E}(\theta|x) = \frac{a+n}{b+T} = \frac{b}{b+T} \cdot \frac{a}{b} + \frac{T}{b+T} \cdot \frac{n}{T} = 0.138 \cdot 0.5 + 0.862 \cdot 0.4 = 0.414.$$

This is a weighted average of the prior mean and the MLE, where the weight of the data increases as T increases.

The credible interval of definition 28.1 comes from the same chi-squared fact as before: if $\theta \sim \text{Gamma}(a + n, b + T)$ then $2(b + T)\theta \sim \chi^2_{2(a+n)}$, so the equal-tailed 95% credible interval is $[\chi^2_{24,0.025}, \chi^2_{24,0.975}]/(2 \cdot 29) = [0.214, 0.679]$. For the manufacturer's claim, the posterior probability of long lifetimes is $\mathbb{P}(\theta < 0.5|x) = \mathbb{P}(\chi^2_{24} < 29) = 0.78$, up from the prior probability 0.59. Unlike the p -value, this is a probability about the parameter, computed after the data have been seen.

The Jeffreys prior from definition 26.8 is $\pi_J(\theta) \propto \sqrt{\mathcal{I}(\theta)} = 1/\theta$. This is the improper limit $a, b \rightarrow 0$ of the Gamma prior. Using this prior, the posterior distribution is $\text{Gamma}(n, T)$. The credible interval is $[\chi^2_{2n, \alpha/2}, \chi^2_{2n, 1-\alpha/2}]/(2T)$. This is the same as the exact confidence interval introduced in section 31.3. Also, $\mathbb{P}(\theta \geq \theta_0|x) = \mathbb{P}(\chi^2_{2n} \geq 2\theta_0 T)$. This is the p -value of the one-sided test. While the numbers coincide, the statements are different: the confidence interval covers the fixed θ in 95% of repetitions, whereas the credible interval covers the random θ with probability 0.95 given the data; and the p -value is a probability about the data under H_0 , whereas $\mathbb{P}(\theta \geq \theta_0|x)$ is a probability about the hypothesis. This illustrates the comparison from section 28.3, for a concrete model, and it is special that the numbers coincide for the Jeffreys prior and one-sided hypotheses.

31.6 Looking Back and Looking Forward

Every question about the exponential model we have been able to answer so far has used a small number of different tools: the exponential-family structure of the model allowed us to find a complete sufficient statistic, the likelihood function allowed us to find the estimator, the Fisher information allowed us to find the standard error and the bound, a pivot or the normal limit allowed us to find the interval, the Neyman–Pearson lemma and Wilks' theorem allowed us to find tests, and Bayes' theorem allowed us to find the posterior distribution. The same list of tools applies for any model, and it is this list which you should use when you are given a model not considered in these notes: Write down the

likelihood function, find the sufficient statistic, solve the likelihood equation and check that the root is a maximum, compute the information, and use the results to compute the standard error, the interval, and the tests. The exact pivot is a luxury of the exponential distribution; in general you will only have the large-sample results from lectures 16, 17 and 23, and some care is needed when these results are used.

This is also the material we will study after this module. The standard errors computed by all regression packages are the inverse Fisher information from lecture 14, the tests which compare fitted models from two different analyses are the likelihood-ratio tests from lecture 23, mixture models and hidden Markov models are fitted using the EM algorithm from lecture 19, and Bayesian computation, from the simple conjugate updating to the simulation methods of later modules, is based on the posterior distribution from lecture 26. The models under consideration here will be much larger, but the questions we ask of these models will be similar to the ones we considered in this module.

Summary.

- For the exponential distribution with rate θ , the sum T is complete and sufficient, the MLE is $1/\bar{X}$, the UMVUE is $(n-1)/T$, and no unbiased estimator can achieve the Cramer–Rao bound θ^2/n , although $\bar{X}$ is efficient for the mean.
- The standard error $\hat{\theta}/\sqrt{n}$ can be used to compute the Wald interval. The pivot $2\theta T \sim \chi_{2n}^2$ can be used to compute an exact interval, an exact one-sided test with known power function, and the p -value.
- A two-sided test can be based on $W = 2n(\log(\hat{\theta}/\theta_0) - 1 + \theta_0/\hat{\theta})$ and Wilks’ theorem.
- For the Gamma(a, b) prior, the posterior distribution is Gamma($a + n, b + T$), the posterior mean is a weighted average of the prior mean and the MLE, and the Jeffreys prior $1/\theta$ gives the exact interval and p -value numerically, but with a different meaning.
- The same short list of tools can be used to answer all questions about all regular models.

Exercise 31.1. Let $X_1, \dots, X_n$ be i.i.d. with exponential rate $\theta > 0$ and let $\hat{\theta} = 1/\bar{X}$ be the MLE.

1. Show that the generalised likelihood-ratio statistic for $H_0: \theta = \theta_0$ against $H_1: \theta \neq \theta_0$ is given by $W = 2n(\log(\hat{\theta}/\theta_0) - 1 + \theta_0/\hat{\theta})$, and that W depends on the data only through $u = \theta_0\bar{X}$, where $W = 2n(u - 1 - \log u)$.
2. Show that $u \mapsto u - 1 - \log u$ is zero at $u = 1$ and positive elsewhere, *i.e.* that $W \geq 0$ with equality if and only if $\hat{\theta} = \theta_0$.
3. For $n = 10$ and $\bar{x} = 2.5$, determine the value of W for $\theta_0 = 0.5$ and for $\theta_0 = 1$ and state the conclusion of the test at level 0.05 for both cases.

Exercise 31.2. Let $X_1, \dots, X_n$ be i.i.d. exponential with rate $\theta > 0$ and let $T = \sum_{i=1}^n X_i$.

1. Show that the Jeffreys prior for θ is $\pi_J(\theta) \propto 1/\theta$. Show that this prior is improper.
2. Show that the posterior distribution under the Jeffreys prior is proper for every sample with $T > 0$. Identify the posterior distribution as a Gamma distribution.
3. Using proposition 26.9 or otherwise, find the Jeffreys prior for the mean $\mu = 1/\theta$. Show that this prior is the image of π_J under the map $\theta \mapsto 1/\theta$.
4. Show that the equal-tailed 95% credible interval for θ , under the Jeffreys prior, coincides with the exact confidence interval from section 31.3, but in one sentence explain why the two intervals make different statements.

Exercise 31.3. Let $X_1, \dots, X_n$ with $n \geq 2$ be i.i.d. uniform on $(0, \theta)$, where $\theta > 0$ is unknown, and let $X_{(1)} = \min_i X_i$ and $X_{(n)} = \max_i X_i$. Without proof, we state that, given $X_{(n)} = m$, the remaining $n-1$ observations are i.i.d. uniform on $(0, m)$ and $\text{Var}_\theta((n+1)X_{(n)}/n) = \theta^2/(n(n+2))$.

- (a) Find the density of $X_{(1)}$ and show that $\mathbb{E}_\theta(X_{(1)}) = \theta/(n+1)$.
- (b) Show that $U = (n+1)X_{(1)}$ is an unbiased estimator for θ . Determine $\text{Var}_\theta(U)$ and show that U is not a consistent estimator for θ .
- (c) Show that $X_{(n)}$ is sufficient for θ .
- (d) Determine $V = \mathbb{E}_\theta(U|X_{(n)})$ and show that $V = (n+1)X_{(n)}/n$. Which theorem guarantees, without computation, that V is a statistic, unbiased for θ and with a variance of at most that of U ?
- (e) Explain why V is the UMVUE for θ . The Cramer–Rao formula, for this model, gives the value θ^2/n , which is larger than the variance of V . Explain why there is no contradiction.

Exercise 31.4. Seed trays are bought from two suppliers. Each tray contains $m = 10$ seeds, and each seed has probability $p_A = 0.2$ of germinating if the tray comes from supplier A, and probability $p_B = 0.7$ if the tray comes from supplier B, independently of the other seeds. Both probabilities are known. It is unknown what proportion $\pi \in (0, 1)$ of trays come from supplier A, but the supplier of individual trays is not recorded. For tray i , the number X_i of germinated seeds is observed, and we write $Z_i = 1$ if the tray came from supplier A and $Z_i = 2$ otherwise. Let f_A and f_B be the probability weights of the Binomial(10, 0.2) and Binomial(10, 0.7) distributions.

- (a) Write down the probability weights of the number X_i of germinated seeds and the observed-data log-likelihood $\ell(\pi)$ for n trays. Explain why the likelihood equation cannot be solved analytically and write down the complete-data log-likelihood $\ell_c(\pi)$.

- (b) Derive the E step of the EM algorithm, showing that $Q(\pi|\pi_k)$ can be found from $\ell_c(\pi)$ by replacing the indicator $\mathbf{1}_{\{Z_i=1\}}$ by the responsibility

$$\gamma_i = \frac{\pi_k f_A(x_i)}{\pi_k f_A(x_i) + (1 - \pi_k) f_B(x_i)}.$$

- (c) Derive the M step, $\pi_{k+1} = \frac{1}{n} \sum_{i=1}^n \gamma_i$, and show that this is a maximum of $\pi \mapsto Q(\pi|\pi_k)$.
- (d) For five trays, the number of germinated seeds was 2, 7, 1, 8 and 3. The probability weights of the two distributions are given by

x	1	2	3	7	8
$f_A(x)$	0.2684	0.3020	0.2013	0.000786	0.0000737
$f_B(x)$	0.000138	0.00145	0.00900	0.2668	0.2335

Given $\pi_0 = 1/2$, perform one iteration of the EM algorithm and comment on your responsibilities.

- (e) Assume now that p_A is unknown, and p_B is still known. Derive the M step for p_A . State the ascent property of the EM algorithm, and explain what the property does and does not guarantee for the sequence $(\pi_k, p_{A,k})$.

Preparing for the Exam

This chapter summarises what you need to do to prepare for the examination: it includes a map of the module, a summary of important results, information about what you were required to prove and what you were only required to know, advice on how to write your answers, details of common mistakes to avoid, and information on how to use the practice paper. None of this is new material, but it is gathered together here from different parts of the module for revision purposes.

The Module in One Page

The module consists of one question, split into two parts. The first part, covering lectures 1 to 19, deals with estimation: the question is what an estimator is, and how one can measure quality of an estimator. In this part of the module we also learn how estimators can be constructed using the method of moments and the maximum likelihood method, how a sufficient statistic can be used to compress the data and how conditioning on the data can be used to improve an estimator. In this part of the module we also learn about the Fisher information and how it can be used to give a bound on the variance of all unbiased estimators. We learn how the maximum likelihood estimator behaves for large sample size and how to compute the maximum likelihood estimator when some of the data are missing. Throughout this part of the module we will consider the exponential families, introduced in lecture 7. These families of distributions provide the sufficient statistic, the completeness of the sufficient statistic, the Fisher information and, in lecture 26, the conjugate priors. The second part of the module, covering lectures 20 to 29, deals with testing and interval estimation. In this part of the module we learn about the language of tests, the best test for simple hypotheses, the likelihood ratio for composite hypotheses, the duality of tests and confidence intervals, the Bayesian approach to the same questions and the standard tests for normal data. Table 31.1 shows the lectures for each week of the module and the results of each lecture and the later lectures which depend on the results of these lectures.

The Results You Need to Know

For each result in this section, you must know the statement and conditions, and you must be able to reproduce the proof. A proof is reconstructed from the idea, rather than being memorised, so for each result we have written down the idea in one line. If this line is not enough for you to work out the argument, then you need to spend some time revising the proof.

Week	Lectures and main results	Used later in
1	1: model, statistic, estimator (definitions 1.1, 1.3 and 1.4). 2: bias and mean squared error, theorem 2.7; consistency, theorem 2.10.	every later lecture; the bias corrections of lecture 5; the consistency proof of lecture 16
2	4: likelihood, score and method of moments (definitions 4.1, 4.4 and 4.6). 5: the MLE (definition 5.1), the likelihood equation and its check, invariance, theorem 5.6; failures, section 5.3.	lectures 11, 16, 17, 19 and 23
3	7: exponential families (definition 7.1), the cumulant function, theorem 7.7, full rank (definition 7.10). 8: sufficiency (definition 8.1), the factorisation theorem 8.2, minimal sufficiency (theorem 8.6), completeness (definition 8.8), maximum and minimum (proposition 8.9).	lectures 10, 11, 13 and 26
4	10: Rao–Blackwell, theorem 10.1; Lehmann–Scheffe, theorem 10.4; the completeness facts, proposition 10.5; Basu, theorem 10.9, and corollary 10.10. 11: Fisher information (definition 11.2), the identities, theorem 11.3, additivity and reparametrisation (propositions 11.4 and 11.10).	lectures 13, 16, 17, 25, 26 and 29
5	13: the Cramer–Rao inequality, theorems 13.2 and 13.3; efficiency (definition 13.4) and attainment, proposition 13.5; the uniform counterexample, example 13.6. 14: the information matrix (definition 14.2), theorem 14.5, the normal matrix of example 14.6.	lectures 17 and 23
6	16: the regularity conditions (definition 16.1), theorems 16.2 and 16.3. 17: asymptotic normality, theorem 17.1; standard errors, corollary 17.3; the delta method, theorem A.18.	lectures 23, 25 and 28
7	19: the EM algorithm (definition 19.1), the ascent property, theorem 19.2, fixed points, proposition 19.3, the normal mixture of example 19.4. 20: tests, errors, power, size and level (definitions 20.2 to 20.5), the p -value (definition 20.8), the z -tests of examples 20.6 and 20.7.	lectures 22, 23, 25 and 29
8	22: the Neyman–Pearson lemma, theorem 22.2; one-sided optimality, proposition 22.4; no two-sided UMP test, proposition 22.5. 23: the generalised likelihood ratio (definition 23.1), Wilks’ theorem 23.3, the exact case of example 23.2.	lectures 25 and 29
9	25: confidence intervals and pivots (definitions 25.1 and 25.2), duality, theorem 25.6, the Wald interval, proposition 25.7. 26: posterior, theorem 26.2; conjugacy, proposition 26.4 and example 26.5; the normal–normal model, proposition 26.6; the posterior mean, theorem 26.7; Jeffreys priors, definition 26.8 and proposition 26.9.	lectures 28 and 29
10	28: credible and HPD intervals (definitions 28.1 and 28.2), Bayesian tests (definition 28.6, example 28.5), the comparison of section 28.3. 29: the t -test and the variance test (propositions 29.1 and 29.3), table 29.1.	the examination
11	31: the exponential model worked from start to finish, a synthesis of the whole module.	

Table 31.1: The lectures week by week, with the results each contributes and the lectures which build on them. The R sessions are omitted.

We begin with estimation.

- Theorem 2.7, the decomposition $\text{MSE} = \text{Var} + \text{bias}^2$: add and subtract the mean inside the square, and the cross term vanishes. Theorem 2.10, a vanishing mean squared error gives consistency: Markov's inequality, lemma A.5, applied to $(\hat{\theta}_n - \theta)^2$.
- Theorem 5.6, the MLE for $g(\theta)$ is $g(\hat{\theta})$: a relabelling of the parameter does not change which distribution has the largest likelihood.
- Theorem 7.7, the moments of the natural statistic are derivatives of K : differentiate the normalising integral under the integral sign, once for the mean and twice for the variance.
- Theorem 8.2, in the discrete case: write the joint probability as $\mathbb{P}_\theta(T = t) \mathbb{P}_\theta(X = x | T = t)$, and observe that sufficiency is the statement that the second factor is free of θ ; for the converse, sum the factorisation over the data vectors with the same value of T .
- Theorem 10.1: the conditional expectation given a sufficient statistic is a statistic, the tower property keeps it unbiased, and the law of total variance, proposition A.11, shows that the variance can only decrease.
- Theorem 10.4: two unbiased functions of a complete statistic have a difference with mean zero, and thus are equal, so that the unbiased function of T is unique; Rao–Blackwell shows that no other unbiased estimator beats it.
- Theorem 10.9: the conditional probability $\mathbb{P}_\theta(A \in B | T)$ minus the unconditional one has mean zero for every θ , and completeness forces it to be zero.
- Theorem 11.3: differentiate $\int f(x; \theta) dx = 1$ once to see that the score has mean zero, lemma 11.1, and a second time, using the quotient rule, to obtain $\mathcal{I}_n = -\mathbb{E}_\theta(\ell'')$.
- Theorems 13.2 and 13.3, the Cramer–Rao inequality: the covariance of an unbiased estimator for $g(\theta)$ with the score equals $g'(\theta)$, by differentiating $\mathbb{E}_\theta(W) = g(\theta)$ under the integral sign, and the Cauchy–Schwarz inequality, lemma A.8, does the rest. Proposition 13.5 is the equality case of Cauchy–Schwarz: the estimator must be an affine function of the score.
- Theorem 16.2: by the law of large numbers the average log-likelihood ratio converges to

$$\mathbb{E}_{\theta_0}(\log(f(X; \theta)/f(X; \theta_0))),$$

and Jensen's inequality, lemma A.7, makes this expectation negative.

- Theorem 16.3: on the event that $L(\theta_0)$ exceeds both $L(\theta_0 - \varepsilon)$ and $L(\theta_0 + \varepsilon)$, whose probability tends to one by the previous theorem, the log-likelihood has an interior maximum and thus, by Rolle's theorem, a root of the likelihood equation within ε of θ_0 ; uniqueness of the root makes it the MLE.

- Theorem 17.1: expand the score about θ_0 , solve for $\sqrt{n}(\hat{\theta}_n - \theta_0)$, and apply the central limit theorem to the score, the law of large numbers to the second derivative, consistency and (R4) to the remainder and Slutsky's theorem, theorem A.16, to combine the three. Corollary 17.3 follows by replacing θ_0 by $\hat{\theta}_n$ in the variance, again by Slutsky's theorem.
- Theorem 19.2: write $\ell = Q - H$ as in (19.1); the M step increases Q , and Jensen's inequality shows that $H(\theta|\theta_k)$ is largest at $\theta = \theta_k$, so that H cannot increase.

We now turn our attention to testing, intervals and the Bayesian approach.

- Theorem 22.2: The integrand $(\mathbf{1}_C - \mathbf{1}_{C'})(f(x; \theta_1) - k f(x; \theta_0))$ is non-negative everywhere, and the theorem follows by integrating this non-negative function to get $\beta_C(\theta_1) - \beta_{C'}(\theta_1) \geq k(\beta_C(\theta_0) - \beta_{C'}(\theta_0)) \geq 0$. From proposition 22.4 we know that the region of the test is the same for all alternatives. From proposition 22.5 we know that there is no two-sided uniformly most powerful test, because such a test would have to agree with both one-sided tests, which reject on opposite tails.
- Theorem 25.6: The parameter θ_0 is in the interval $C(X)$ if and only if the test for θ_0 does not reject the value, and thus the probability of the interval covering θ_0 is one minus the size of the test.
- Proposition 25.7: From corollary 17.3 we know that the standardised error is approximately standard normally distributed, and so bracketing it between $\pm z_{1-\alpha/2}$ gives the interval.
- Theorem 26.2: This is an application of Bayes' theorem, proposition A.20, where the parameter is the event and the data is the evidence. Proposition 26.4: $\exp(a\eta - bK(\eta))$ can be multiplied by a likelihood of the form $\exp(\eta T_n - nK(\eta))$ to obtain an exponent of the same shape. Theorem 26.7: The same add-and-subtract argument as in theorem 2.7, but now under the posterior distribution.
- Proposition 26.9: The change of variables for densities, proposition A.19, provides the factor $|h'(\psi)|$ and the reparametrisation rule for the information, proposition 11.10, provides the same factor under the square root.
- Propositions 29.1 and 29.3: The statistics are pivots from lecture 25, evaluated at the null value, and the null distribution is given by proposition A.3 and definition A.4.

In addition to the theorems, the examination requires the recall of definitions and standard values, requiring only memory. The definitions are given in table 31.1 and should be given in the form in the notes, including the quantifier “for all $\theta \in \Theta$ ” wherever required. The standard values are the MLEs for the families in tables 1.1 and 1.2, as summarised in lecture 5, the Fisher information for one observation for the same families, as summarised in lecture 11, the normal information matrix (14.1) and its inverse, the one-sided and two-sided z -tests and their power functions from lecture 20, and the Beta–Bernoulli and normal–normal updates from lecture 26.

What Is (Not) Proved

The following list summarises the results the module assumes to be known without proof. (No one should try to improve a non-existent proof!) For each result, the examinable material is the statement, including any conditions, and how to use the result in a calculation.

- Wilks' theorem 23.3. Know the setting (23.3), the conclusion $W_n \xrightarrow{d} \chi_r^2$ under H_0 , and how to use the result: compute W from the two maximised log-likelihoods, count r as the number of constraints, and read off $\chi_{r,1-\alpha}^2$ from the table. The argument in lecture 23 for the case $k = r = 1$ is an explanation, not a proof. Also note that in example 23.2 the statement is exact for every n .
- The continuous case of the factorisation theorem 8.2. The proof of the discrete case is examinable, but for continuous data the theorem is used in both directions, with the indicator of a parameter-dependent support included in g or h as appropriate. The minimal-sufficiency criterion, theorem 8.6, is in a similar position: the proof was only sketched, but use of the result is examinable.
- The completeness results, proposition 10.5. Both results can be quoted as needed, and in the case of the exponential family the full-rank condition from definition 7.10 must be mentioned. Any other completeness results will be provided in the question.
- The vector Cramer–Rao bound, theorem 14.5. Know the statement, that the difference of the covariance matrix and the inverse information is positive semi-definite, and the consequence for single coordinates, $\text{Var}_\theta(T_j) \geq (\mathcal{I}_n^{-1})_{jj}$, together with the example of the normal distribution. The proof was sketched in lectures.
- The vector case of theorem 17.1, as given at the end of lecture 17: the MLE of a vector parameter is asymptotically normal with covariance matrix $\mathcal{I}_n(\theta_0)^{-1}$. Know the statement and what it means for the normal distribution, where the diagonal matrix shows the separation of the two coordinates.
- The large-sample normal approximation (28.4) to the posterior distribution. Know the statement, that under the regularity conditions and for a prior with continuous positive density the posterior is approximately $N(\hat{\theta}_n, 1/\mathcal{I}_n(\hat{\theta}_n))$, and its consequence, that the equal-tailed credible interval is approximately the Wald interval whatever the prior. The argument given in lecture 28 is an explanation, not a proof.

The material in appendix A is background knowledge which is used without proof, but the examination expects you to name the results when they are used. All the other material from the content lectures, including the MLE asymptotics from lectures 16 and 17, and the Lehmann–Scheffe theorem, was proved in full and is examinable as such.

Writing an Answer

The model answers in appendix C show the level of detail expected. Every statistical argument in this module has the same five steps, and an answer which makes each step visible is a complete answer.

1. State the model and the parameter. Write down the density or the probability weights, including the support and the parameter space, before doing anything with it. Many likelihood mistakes, in particular for models where the support depends on the parameter, are caused by leaving out the indicator of the support.
2. Name the result used: “by the factorisation theorem”, “by Wilks’ theorem”. A calculation which gives the correct result without explaining how it was obtained is not a complete answer.
3. Check the conditions. If the result requires a sufficient statistic to be complete, as for the Lehmann–Scheffe theorem 10.4, you should mention where the completeness comes from. If the result requires regularity conditions, as for the Cramer–Rao inequality, theorem 13.2, you should say that the support does not depend on the parameter. For a Wald interval or for Wilks’ theorem 23.3, you should say that the statement is an approximation for large samples. If a question asks why a result does not apply, this step is the whole answer.
4. Compute the result. Keep the computation on the same page where you write down the argument, to allow the reader to follow your argument.
5. Give the conclusion in words, in the language of the original question: “the test rejects H_0 at the 5% level”. Use the correct kind of statement: a confidence interval is a range of values which cover the parameter with a given probability, as repeated samples are taken. A credible interval is a range of values which have a given posterior probability, given the data. A p -value is not the probability that H_0 is true.

For a question like “show that” something is true, the answer is given in the question itself, and the argument is what counts. If one part of a question cannot be answered, later parts can still be answered: the questions are written so that later parts only use the statement of the previous parts, not the proofs of the previous parts.

Three habits help in every question. Distinguish the information of one observation from that of the sample, $\mathcal{I}_n = n\mathcal{I}$, and say which one a formula uses. Verify that a root of the likelihood equation is a maximum. Count degrees of freedom by counting constraints, and remember that a two-sided chi-squared test needs two quantiles, because the distribution is not symmetric. The next section lists the mistakes made most often.

Common Mistakes

The following mistakes are committed in scripts year after year, and they are all easy to avoid once the mistakes are known.

- Treating the parameter as random in a frequentist statement. A confidence interval is a random interval which covers the fixed value θ with probability 0.95 before the data are seen; the sentence “ θ lies in $[l, u]$ with probability 0.95” is a statement about a credible interval, definition 28.1, and is false for a confidence interval, definition 25.1.
- Using the Cramer–Rao inequality, or the asymptotic normality of the MLE, for a model whose support depends on the parameter. The uniform distribution on $(0, \theta)$ is a standard counterexample (examples 11.11 and 13.6); for such a model the estimator based on the maximum or the minimum has to be analysed directly.
- Claiming that an estimator is a UMVUE because it is unbiased and a function of a sufficient statistic. The Lehmann–Scheffe theorem requires the statistic to be complete, and this completeness needs to be shown, typically using proposition 10.5.
- Finding a root of the likelihood equation and calling it the MLE without checking that it is a maximum. The second derivative, or the concavity of the log-likelihood, is part of the answer, and for a model like the uniform distribution the maximum is not a root of the likelihood equation (section 5.3).
- Interpreting a non-rejection as a proof of H_0 , or interpreting a p -value as the probability that H_0 is true. A test controls the probability of type I errors, and the p -value is a probability, computed under H_0 , about the data, definition 20.8.
- Mixing up the one-observation and the sample information. The Cramer–Rao bound and the standard error use $\mathcal{I}_n(\theta) = n\mathcal{I}(\theta)$, and the limit in theorem 17.1 uses $\mathcal{I}(\theta)$ together with the factor $\sqrt{n}$; both forms are equivalent, but the n needs to be in one place or the other.
- Leaving out the indicator of the support when factorising a density. For the uniform distribution the factor $\mathbf{1}_{\{x_{(n)} \leq \theta\}}$ is the whole point of the factorisation, and without it the maximum would not be sufficient (exercise 8.2).
- Stopping at a number. An estimate of 0.111, a statistic of 3.61 or a posterior probability of 0.57 is not an answer until it has been compared to the relevant threshold and turned into a sentence about the question.

Using the Practice Paper

The exercises in appendix C are the best rehearsal for the examination you will get, and it is advisable to try them yourself before looking at the solutions.

A practice paper is most useful when it is sat before the answers are seen. If you have not yet done so, sit the paper under examination conditions, with the answers written out in full. A paper which has been read through, with a mental note that each part looks manageable, tells you very little. A paper which has been sat tells you where the time went, which parts you could start but not finish, and which results you thought you knew but could not write down.

Two habits from doing the paper carry over to the examination. If you find yourself spending more than a quarter of an hour on one part of the paper, write down what you have done so far and then move on to the next question. It is better to have a partial answer for all parts of a question than to have a complete answer for only some parts.

Thank you for taking part in this module, and good luck in the examination.

Appendix A

Background from Probability

This appendix summarises some important results from probability which we use in the module. The prerequisite for understanding the contents of this appendix is the knowledge of the MATH2701 Statistical Methods module; all the results summarised here should be known from that module (or from an equivalent background). The purpose of this appendix is to fix notation and to provide a point of reference for the main text. None of the material in this appendix is examinable on its own. The results included here are mostly statements about fully specified distributions (known models). The opposite direction, where we try to infer an unknown parameter from data, is the one we are concerned with in the main text. We include proofs of the basic tools, because these are short and instructive. We only state the deeper limit theorems, because proofs of these results would take up valuable space and time in this text.

A.1 Expectation, Variance and Independence

We write $\mathbb{E}(X)$ for the expectation of a random variable X and $\text{Var}(X)$ for its variance. We assume that the basic rules of algebra for expectations and variances are known. For reference, we state the following rules for expectations and variances.

Proposition A.1. *Let X and Y be random variables with finite variance, and let a and b be constants. Then*

$$\mathbb{E}(aX + bY) = a \mathbb{E}(X) + b \mathbb{E}(Y) \quad \text{and} \quad \text{Var}(aX + b) = a^2 \text{Var}(X).$$

The covariance is given by $\text{Cov}(X, Y) = \mathbb{E}(XY) - \mathbb{E}(X) \mathbb{E}(Y)$ and thus

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y) + 2 \text{Cov}(X, Y).$$

If X and Y are independent, then $\mathbb{E}(XY) = \mathbb{E}(X) \mathbb{E}(Y)$ and thus the covariance is zero. Similarly, the variance of the sum can be computed as the sum of the variances.

More generally, for independent $X_1, \dots, X_n$ the variance of the sum is the sum of the individual variances. We use this fact, for example, when we average a random sample. Since these identities are standard, we do not give proofs of these results here.

A.2 Standard Distributions

We will consider a small number of distributions in the module. Tables A.1 and A.2 summarise the probability weights or density, support, mean and variance of these distributions, for the discrete and for the continuous distributions respectively. The families of distributions considered in tables 1.1 and 1.2 are all included, together with their moments. The negative binomial, Pareto and Rayleigh distributions do not appear in tables 1.1 and 1.2; these distributions are used in the R sessions as a source of simulated data. The tables use the usual parameter names, but in the main text these families of distributions will be considered as parametric models where the unknown parameter is denoted by θ . For example, the rate parameter λ of the exponential distribution below is the θ of table 1.2. Throughout the module the exponential and gamma distributions are parametrised by their *rate*, as in the table and as in R's `dexp` and `dgamma`; some textbooks use the mean $1/\lambda$, or equivalently the scale, so a formula from elsewhere may need to be translated.

Family	$f(x)$	Support	Mean	Variance
Bernoulli(p)	$p^x(1-p)^{1-x}$	$\{0, 1\}$	p	$p(1-p)$
Binomial(m, p)	$\binom{m}{x} p^x(1-p)^{m-x}$	$\{0, \dots, m\}$	mp	$mp(1-p)$
Geometric(p)	$(1-p)^{x-1}p$	$\{1, 2, \dots\}$	$1/p$	$(1-p)/p^2$
Poisson(λ)	$\frac{\lambda^x}{x!} e^{-\lambda}$	$\{0, 1, 2, \dots\}$	λ	λ
Negative binomial(r, p)	$\binom{x+r-1}{x} p^r(1-p)^x$	$\{0, 1, 2, \dots\}$	$\frac{r(1-p)}{p}$	$\frac{r(1-p)}{p^2}$

Table A.1: The standard discrete distributions and their moments. Here $f(x)$ gives the probability weights. For the geometric distribution, the number of trials until the first success is considered, so the support starts at 1. In contrast, some texts and the statistical software R consider the number of failures before the first success, so the support starts at 0 (see appendix B).

Family	$f(x)$	Support	Mean	Variance
Uniform($0, \theta$)	$1/\theta$	$[0, \theta]$	$\theta/2$	$\theta^2/12$
Exponential(λ)	$\lambda e^{-\lambda x}$	$[0, \infty)$	$1/\lambda$	$1/\lambda^2$
Gamma(α, β)	$\frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x}$	$(0, \infty)$	α/β	α/β^2
Beta(α, β)	$\frac{x^{\alpha-1}(1-x)^{\beta-1}}{B(\alpha, \beta)}$	$[0, 1]$	$\frac{\alpha}{\alpha+\beta}$	$\frac{\alpha\beta}{(\alpha+\beta)^2(\alpha+\beta+1)}$
Normal(μ, σ^2)	$\frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$	$\mathbb{R}$	μ	σ^2
Cauchy	$\frac{1}{\pi(1+x^2)}$	$\mathbb{R}$	none	none
Pareto(α)	$\alpha x^{-(\alpha+1)}$	$[1, \infty)$	$\frac{\alpha}{\alpha-1}$	$\frac{\alpha}{(\alpha-1)^2(\alpha-2)}$
Rayleigh(θ)	$\frac{2x}{\theta} e^{-x^2/\theta}$	$[0, \infty)$	$\sqrt{\pi\theta}/2$	$(1-\pi/4)\theta$

Table A.2: The standard continuous distributions and their moments. Here $f(x)$ is the density, Γ is the gamma function and $B(\alpha, \beta) = \Gamma(\alpha)\Gamma(\beta)/\Gamma(\alpha+\beta)$ the beta function. The exponential distribution is the special case Gamma(1, λ). The mean of the Pareto distribution only exists for $\alpha > 1$ and the variance only for $\alpha > 2$.

In the module we often need the distribution of a sum of independent random

variables, for example when we compute the distribution of a total or of a sample mean. For several of the standard distributions the sum of independent copies again has one of the standard distributions, with only the parameters changed.

Remark. If $X_1, \dots, X_n$ are independent, then the sum of Bernoulli(p) variables is Binomial(n, p), the sum of Poisson variables is Poisson with the sum of the means, the sum of Exponential(λ) variables (or, more generally, of Gamma variables with the same rate) is Gamma with the sum of the shapes, and the sum of independent normals is normal with the sum of the means and the sum of the variances.

Another family, introduced in exercise 23.2 of lecture 23, is the multinomial distribution. If each of n independent trials has one of k possible categories, with probabilities $p_1, \dots, p_k$, where the probabilities must sum to one, then the vector of counts $(N_1, \dots, N_k)$ is *multinomial* distributed with $\mathbb{E}(N_j) = np_j$ and $\text{Var}(N_j) = np_j(1 - p_j)$. The individual counts are Binomial(n, p_j) and the multinomial distribution is a generalisation of the binomial distribution to the case where there are more than two categories.

A.3 Sampling from a Normal Population

The following distributions are needed to describe the sampling distributions of statistics computed from a normal sample. They form the basis of the confidence intervals in lecture 25 and of the standard tests in lecture 29. The first of these distributions also arose in the comparison of estimators for the variance in lecture 2. We start our discussion by considering the chi-squared distribution.

Definition A.2. Let $Z_1, \dots, Z_k$ be independent, standard normally distributed random variables. Then the distribution of $Z_1^2 + \dots + Z_k^2$ is called the *chi-squared distribution with k degrees of freedom*. This distribution is denoted by χ_k^2 . The mean and variance of the chi-squared distribution are k and $2k$, respectively.

The chi-squared distribution is a special case of the gamma distribution: the χ_k^2 distribution corresponds to the Gamma($k/2, 1/2$) distribution from table A.2, i.e. it has density $x^{k/2-1}e^{-x/2}/(2^{k/2}\Gamma(k/2))$ for $x > 0$. In statistics, the parametrisation in terms of degrees of freedom is more commonly used. The most important property of the chi-squared distribution for applications in statistics is that the sample mean and sample variance, being the two most common statistics for normally distributed samples, have known and independent distributions.

Proposition A.3. Let $X_1, \dots, X_n$ be i.i.d. $N(\mu, \sigma^2)$ variables. Then the sample mean $\bar{X}$ and the sample variance

$$S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$$

satisfy

$$\bar{X} \sim N\left(\mu, \frac{\sigma^2}{n}\right), \quad \frac{(n-1)S^2}{\sigma^2} \sim \chi_{n-1}^2,$$

and the two statistics $\bar{X}$ and S^2 are independent.

The independence of $\bar{X}$ and S^2 is the least obvious of these statements; we will prove this statement in lecture 10, as an application of Basu's theorem.

We now construct two more distributions, using the normal distribution and the chi-squared distribution as building blocks. These distributions describe what happens when trying to estimate an unknown variance.

Definition A.4. Let $Z \sim N(0, 1)$, $V \sim \chi_k^2$ and $W \sim \chi_l^2$ be independent. Then the t distribution with k degrees of freedom is the distribution of

$$T = \frac{Z}{\sqrt{V/k}}$$

denoted by t_k , and the F distribution with k and l degrees of freedom is the distribution of

$$F = \frac{V/k}{W/l},$$

denoted by $F_{k,l}$.

Using definition A.4 and proposition A.3 one can understand the origin of the t statistic for a normal sample: by dividing the standardised mean by the square root of the scaled sample variance, the unknown σ is replaced with its estimate and a t_{n-1} variable is obtained. We will see this argument in lecture 25.

A.4 Some Inequalities

A small number of basic inequalities are used throughout the module, mostly to estimate the probability of an estimator being away from its target. The most basic of these inequalities is an upper tail bound for non-negative variables, in terms of the mean.

Lemma A.5 (Markov's inequality). *Let X be a non-negative random variable and $a > 0$. Then*

$$\mathbb{P}(X \geq a) \leq \frac{\mathbb{E}(X)}{a}.$$

Proof. By definition, X is non-negative and thus at least a if $\{X \geq a\}$ and at least zero otherwise. Thus we can conclude $X \geq a \mathbf{1}_{\{X \geq a\}}$. Taking expectations we get

$$\mathbb{E}(X) \geq a \mathbb{E}(\mathbf{1}_{\{X \geq a\}}) = a \mathbb{P}(X \geq a).$$

Dividing by the positive number a completes the proof. $\square$

The following application of Markov's inequality, showing that a bound on the mean can be converted into a bound on the spread, completes this section.

Lemma A.6 (Chebyshev's inequality). *Let X have finite variance and $a > 0$. Then*

$$\mathbb{P}(|X - \mathbb{E}(X)| \geq a) \leq \frac{\text{Var}(X)}{a^2}.$$

Proof. The event $|X - \mathbb{E}(X)| \geq a$ coincides with the event $(X - \mathbb{E}(X))^2 \geq a^2$. By Markov's inequality, lemma A.5, for the non-negative variable $(X - \mathbb{E}(X))^2$ with threshold a^2 , we get

$$\mathbb{P}(|X - \mathbb{E}(X)| \geq a) = \mathbb{P}((X - \mathbb{E}(X))^2 \geq a^2) \leq \frac{\mathbb{E}((X - \mathbb{E}(X))^2)}{a^2} = \frac{\text{Var}(X)}{a^2}.$$

This is the required inequality. $\square$

The following inequality, which establishes a relation between the expectation of a convex function and the function of the expectation, is the basis of the information inequality for maximum likelihood and of the ascent property in the EM algorithm.

Lemma A.7 (Jensen's inequality). *Let g be convex and X be a random variable with finite mean. Then*

$$g(\mathbb{E}(X)) \leq \mathbb{E}(g(X)).$$

If g is strictly convex, then the inequality is strict unless X is constant with probability one.

Proof. Since g is convex, at the point $m = \mathbb{E}(X)$ there is a supporting line with a constant slope c , i.e. a function $g(x) \geq g(m) + c(x - m)$ for all x . Substituting X for x and taking expectations, the linear term has expectation zero since $\mathbb{E}(X - m) = 0$, and thus

$$\mathbb{E}(g(X)) \geq g(m) + c\mathbb{E}(X - m) = g(m) = g(\mathbb{E}(X)).$$

If g is strictly convex, the supporting line at m is unique and thus $g(x) > g(m) + c(x - m)$ for all $x \neq m$. In this case, if we have equality $\mathbb{E}(g(X)) = g(m)$, the non-negative random variable $g(X) - g(m) - c(X - m)$ has expectation zero and thus is zero with probability one, i.e. $X = m$ with probability one. This proves the claimed inequality and completes the proof. $\square$

Finally, the Cauchy–Schwarz inequality can be used to estimate a covariance in terms of variances. This result is, for example, used in the derivation of the Cramer–Rao lower bound in lecture 13.

Lemma A.8 (Cauchy–Schwarz inequality). *Let X and Y be random variables with finite variances. Then we have*

$$\text{Cov}(X, Y)^2 \leq \text{Var}(X) \text{Var}(Y).$$

Proof. We can assume that both variances are positive, since otherwise one of the random variables is constant and both sides of the inequality equal zero. Let $U = X - \mathbb{E}(X)$ and $V = Y - \mathbb{E}(Y)$, so that $\text{Cov}(X, Y) = \mathbb{E}(UV)$, $\text{Var}(X) = \mathbb{E}(U^2)$ and $\text{Var}(Y) = \mathbb{E}(V^2)$, and we want to show $\mathbb{E}(UV)^2 \leq \mathbb{E}(U^2) \mathbb{E}(V^2)$. For any real number t , the quantity $\mathbb{E}((U - tV)^2)$ is an expectation of a square and thus is non-negative. This gives

$$\mathbb{E}(U^2) - 2t\mathbb{E}(UV) + t^2\mathbb{E}(V^2) \geq 0$$

for all t . A quadratic polynomial which never takes a negative value has non-positive discriminant and thus we have $4\mathbb{E}(UV)^2 - 4\mathbb{E}(U^2)\mathbb{E}(V^2) \leq 0$ and thus we get the required inequality. This completes the proof. $\square$

A.5 Conditional Expectation

When one random variable carries information about another, we summarise it through conditional expectation. Given random variables X and Y , the conditional expectation $\mathbb{E}(X|Y)$ is the expectation of X computed as if the value of Y were known; because that value is itself random, $\mathbb{E}(X|Y)$ is a function of Y , and hence a random variable in its own right. Its own average recovers the ordinary expectation of X .

Proposition A.9 (Tower property). *Let X and Y be random variables, where X has finite mean. Then*

$$\mathbb{E}(\mathbb{E}(X|Y)) = \mathbb{E}(X).$$

Since the value of $\mathbb{E}(X|Y)$ is computed from the value of Y , we can pull out all factors which are functions of Y from the conditional expectation.

Proposition A.10 (Taking out what is known). *Let X and Y be random variables and let h be a function such that $h(Y)X$ has finite mean. Then*

$$\mathbb{E}(h(Y)X|Y) = h(Y)\mathbb{E}(X|Y).$$

In particular $\mathbb{E}(h(Y)|Y) = h(Y)$.

We already know the tower property and this rule. The only effect of conditioning can be to reduce the variance on average. The exact argument is given by the following decomposition, which we will use in lecture 10 when we improve an estimator by conditioning.

Proposition A.11 (Law of total variance). *Let X be a random variable with finite variance and let Y be any random variable. Then*

$$\text{Var}(X) = \mathbb{E}(\text{Var}(X|Y)) + \text{Var}(\mathbb{E}(X|Y)),$$

where $\text{Var}(X|Y) = \mathbb{E}(X^2|Y) - \mathbb{E}(X|Y)^2$ denotes the conditional variance.

Proof. Let $m = \mathbb{E}(X|Y)$ for the conditional mean. Using the tower property, proposition A.9, for X^2 we find $\mathbb{E}(X^2) = \mathbb{E}(\mathbb{E}(X^2|Y))$ and using the definition of the conditional variance we have $\mathbb{E}(X^2|Y) = \text{Var}(X|Y) + m^2$. Taking expectations we get

$$\mathbb{E}(X^2) = \mathbb{E}(\text{Var}(X|Y)) + \mathbb{E}(m^2).$$

Using the tower property again, we know $\mathbb{E}(m) = \mathbb{E}(X)$ and thus, after subtracting $\mathbb{E}(X)^2 = \mathbb{E}(m)^2$ on both sides, we can write $\mathbb{E}(X^2) - \mathbb{E}(X)^2$ as $\text{Var}(X)$ on the left-hand side and $\mathbb{E}(m^2) - \mathbb{E}(m)^2$ as $\text{Var}(m)$ on the right-hand side. This is the required identity and completes the proof. $\square$

A.6 Convergence and Limit Theorems

The large-sample theory of the module is based on two different concepts of a sequence of random variables converging to a limit. The weaker of the two concepts requires that the random variables are unlikely to differ from the limit by more than a fixed amount.

Definition A.12. A sequence of random variables X_n converges in probability to a limit X , written as $X_n \xrightarrow{\text{P}} X$, if for every $\varepsilon > 0$

$$\mathbb{P}(|X_n - X| > \varepsilon) \longrightarrow 0 \quad (n \rightarrow \infty).$$

The weaker (but still useful) convergence of distribution functions is used to approximate a whole distribution:

Definition A.13. A sequence of random variables X_n with distribution functions F_n converges *in distribution* to a limit X with distribution function F , written as $X_n \xrightarrow{d} X$, if $F_n(x) \rightarrow F(x)$ for all points x where F is continuous.

The two great theorems of large-sample statistics describe the behaviour of a sample mean of independent, identically distributed observations. The first theorem gives the mean of the sample, the second describes the size and shape of the fluctuations around the mean.

Theorem A.14 (Weak law of large numbers). *Let $X_1, X_2, \dots$ be i.i.d. with finite mean μ . Then the sample mean $\bar{X}_n$ satisfies $\bar{X}_n \xrightarrow{p} \mu$ as $n \rightarrow \infty$.*

The assumption of finite mean is important. The Cauchy distribution from table A.2 has no finite mean and the distribution of the sample mean is the same as the distribution of a single observation for all n . Thus, in this case nothing converges and this example is used in lecture 2 to argue that the sample mean is not consistent.

Theorem A.15 (Central limit theorem). *Let $X_1, X_2, \dots$ be i.i.d. with mean μ and finite variance σ^2 . Then*

$$\sqrt{n}(\bar{X}_n - \mu) \xrightarrow{d} N(0, \sigma^2) \quad (n \rightarrow \infty).$$

In order to combine these limits we will need rules which allow us to manipulate limits. Slutsky's theorem allows us to treat a term which converges to a constant as the constant itself, and the continuous mapping theorem allows us to apply a continuous function to a limit.

Theorem A.16 (Slutsky's theorem). *Let $X_n \xrightarrow{d} X$ and $Y_n \xrightarrow{p} c$ for some constant c . Then $X_n + Y_n \xrightarrow{d} X + c$ and $X_n Y_n \xrightarrow{d} cX$. If $c \neq 0$ we have $X_n/Y_n \xrightarrow{d} X/c$.*

Theorem A.17 (Continuous mapping theorem). *Let g be continuous at the possible values of X . Then $X_n \xrightarrow{d} X$ implies $g(X_n) \xrightarrow{d} g(X)$ and $X_n \xrightarrow{p} X$ implies $g(X_n) \xrightarrow{p} g(X)$.*

An important application of this theorem is to apply the central limit theorem after applying a smooth transformation to the data, in order to determine the limiting distribution of a function of an asymptotically normal statistic. This method is known as the delta method. We will use the delta method in lecture 17 to compute standard errors for functions of the maximum likelihood estimator.

Theorem A.18 (Delta method). *Let $\sqrt{n}(T_n - \theta) \xrightarrow{d} N(0, \tau^2)$ and let g be a differentiable function at θ with $g'(\theta) \neq 0$. Then*

$$\sqrt{n}(g(T_n) - g(\theta)) \xrightarrow{d} N(0, g'(\theta)^2 \tau^2).$$

These five theorems are given without proof, and are from the material covered in a first course in probability. We use the theorems as a ready-made toolkit.

A.7 Transformation of a Random Variable

We sometimes know the distribution of a random variable and want to find the distribution of a smooth, one-to-one transformation of this random variable. The change-of-variable formula is a mathematical tool which allows us to find the density of the transformed random variable. This formula is, for example, used when we reparametrise a model, as discussed in the context of the Jeffreys prior in lecture 26.

Proposition A.19 (Change of variable). *Let X be a continuous random variable with density f_X and let $Y = g(X)$ for a differentiable, strictly monotone function g with inverse h and with $g'(x) \neq 0$ for all x . Then Y has density*

$$f_Y(y) = f_X(h(y)) |h'(y)|$$

for all y in the range of g .

Proof. We only consider the case where g is increasing, the proof for the decreasing case is similar with changes in sign. The distribution function of Y is given by

$$F_Y(y) = \mathbb{P}(g(X) \leq y) = \mathbb{P}(X \leq h(y)) = F_X(h(y)),$$

where we used that g is increasing and has inverse h . Taking derivatives with respect to y , using the chain rule, we find $f_Y(y) = f_X(h(y)) h'(y)$ and thus $h'(y) > 0$ equals its own absolute value as given in the formula. This completes the proof. $\square$

The same argument applies to a vector of random variables: if X has a joint density f_X on $\mathbb{R}^k$ and $Y = g(X)$ for a differentiable, one-to-one map g with inverse h , then Y has joint density $f_Y(y) = f_X(h(y)) |\det h'(y)|$, where $h'(y)$ is the Jacobian matrix of h . For a linear map with determinant one, as in the solution to exercise 10.2 in appendix C, the factor is 1.

A.8 Bayes' Theorem

The final tool reverses the direction of conditioning, expressing the conditional density of one variable given another in terms of the reverse conditional. It is the mechanism behind every posterior distribution in the Bayesian lectures.

Proposition A.20 (Bayes' theorem). *Let A and B be random variables with joint density $f(a, b)$ and let B have strictly positive marginal density $f(b)$. Then the conditional density of A given $B = b$ is*

$$f(a|b) = \frac{f(b|a) f(a)}{f(b)},$$

where $f(b) = \int f(b|a) f(a) da$ is the marginal density of B .

If one of the two variables is discrete, the same formula holds with its probability weights in place of the density, and the integral in $f(b)$ becomes a sum over the values of A when A is discrete.

In the Bayesian lectures of the module, the unknown parameter plays the role of A with prior density $f(a)$ and the data are given by B . In this situation, proposition A.20 states that the posterior density is proportional to the product of the likelihood and the prior density. The denominator $f(B)$ does not depend on the parameter. This viewpoint is further developed in lectures 26 and 28.

Appendix B

Introduction to R

This appendix contains a short, self-contained reference to the R programming language as used in this module. We assume that the reader can execute commands in R and understand the output. We collect all R commands used in this module in one place, for the convenience of the reader. The aim is to provide enough information for the reader to not be held back by unfamiliarity with the R programming language.

Commands to be typed by the user are enclosed in a framed box, and the R output is shown below. Each line of the R output is marked with `##` so that it can be distinguished from the input.

B.1 Getting Started

At Leeds the easiest way to run R is via the AppsAnywhere service, which provides access to RStudio on university machines, either in the cloud or using a remote desktop. No installation is required to use this method. Alternatively, to run R on your own computer, you can install the two free programs R from the Comprehensive R Archive Network at <https://cran.r-project.org/> and RStudio Desktop from <https://posit.co/download/rstudio-desktop/>. RStudio provides a much more comfortable interface for using R. You need to install R first, before you can install RStudio.

Once RStudio is open, the large pane on the left is the *console*. This is where we type a command and see the result immediately. For simple calculations, the console is enough, but for anything we want to store or re-run later, we have to type the commands into a *script*, a plain text file, which we can open by clicking **File > New File > R Script**. We can then send individual lines of this file to the console, and we can save the whole file to disk.

Using the console is like using a calculator: we type an expression, and R prints the result.

```
1 + 1
```

```
## [1] 2
```

The `[1]` at the start of the output is a position marker, not part of the answer: it says that the first value printed on that line, here 2, is the first element of the result. For a single number this looks redundant, but it helps

once a result consists of many values spread out over several lines, because then each line starts with the index of the first value on that line.

To store a value for later use, we *assign* a name to it using the arrow `<-`, and we can later type the name to see the value. A list of numbers can be built using `c`, which combines values to form a vector.

```
x <- c(3, 1, 4, 1, 5)
x
mean(x)
```

```
## [1] 3 1 4 1 5
## [1] 2.8
```

Assignments to `x` do not print anything (this is intentional). If we type `x` on its own, R prints the five values. If we type `mean(x)`, R prints the average of the five values.

You will learn many commands, but it is impossible to remember them all. Every built-in function has a help page, which can be accessed by typing a question mark before the function name.

```
?mean
```

The page always has the same structure: the *Usage* block shows how to call the function, the *Arguments* block shows what each input argument does, and the *Examples* block at the bottom shows lines which can be copied into the console to see the function in action. The structure of a help page is the key to understanding an unfamiliar command. For example, if you are asked to use a function which you have never seen before, *e.g.* the exponential random-number generator, you can type `?rexp` to get the line `rexp(n, rate = 1)`. This line explains that the function is called `rexp`, takes a count `n` as its first argument, and has a ‘rate’ argument with a default value. Once you have learned how to read a help page, you will find that this approach is more useful than trying to memorise individual commands.

B.2 Vectors and Arithmetic

Most quantities we consider here are *vectors*, *i.e.* ordered sequences of numbers, constructed using the R command `c`. We can access elements of a vector by specifying either a single index or a range of indices. We can also select elements based on a condition. In the following code, all the text after the single `#` on each line is a *comment*, ignored by R. We use comments to explain what each line of code does.

```
x <- c(3, 1, 4, 1, 5, 9, 2, 6)
x[1]           # the first element
x[2:4]         # elements 2 to 4
x[x > 3]       # the elements bigger than 3
```

```
## [1] 3
## [1] 1 4 1
## [1] 4 5 9 6
```

Arithmetic on vectors is an operation which acts on all elements simultaneously. As a result, loops over the elements of a vector are rarely needed.

```
x + 10      # add 10 to every element
x^2        # square every element
```

```
## [1] 13 11 14 11 15 19 12 16
## [1]  9  1 16  1 25 81  4 36
```

Other functions operate on a vector to compute a single value. The most useful of these are `length` and `sum`, which return the count and the total, respectively, and the standard descriptive statistics.

```
length(x)
sum(x)
mean(x)
median(x)
var(x)
sd(x)
```

```
## [1] 8
## [1] 31
## [1] 3.875
## [1] 3.5
## [1] 7.553571
## [1] 2.748376
```

The functions `var` and `sd` compute the sample variance, using the divisor $n - 1$, and the square root of the sample variance, known as the sample standard deviation. The command `summary` can be used to print several of these values.

```
summary(x)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    1.000   1.750   3.500   3.875   5.250   9.000
```

The command `quantile` can be used to compute sample quantiles. By default the result includes percentage labels, as shown in the output of `summary`, but if the argument `names = FALSE` is used, the labels are omitted. This is useful if the values are to be used in later computations.

```
quantile(x, c(0.25, 0.75))
quantile(x, c(0.25, 0.75), names = FALSE)
```

```
## 25% 75%
## 1.75 5.25
## [1] 1.75 5.25
```

Results for different settings are often best read in tabular form. The commands `cbind` and `rbind` can be used to construct tables from vectors, by specifying the vectors as columns and rows, respectively. Names given in the call to either of these commands are used as column or row labels, respectively.

```
ns <- c(10, 20, 40)
vars <- c(0.1, 0.05, 0.025)
cbind(n = ns, variance = vars)
rbind(flat = c(1, 1), jeffreys = c(0.5, 0.5))
```

```
##          n variance
## [1,] 10      0.100
## [2,] 20      0.050
## [3,] 40      0.025
##          [,1] [,2]
## flat        1.0 1.0
## jeffreys    0.5 0.5
```

The command `cor` can be used to compute the correlation of two vectors. A covariance matrix can be converted into the corresponding matrix of correlations by `cov2cor`, which is how we read the correlation of two estimates off a numerical fit.

```
covmat <- matrix(c(4, 2, 2, 9), nrow = 2)
cov2cor(covmat)
```

```
##          [,1]      [,2]
## [1,] 1.0000000 0.3333333
## [2,] 0.3333333 1.0000000
```

Some of the data sets which come with R are stored as a *data frame*, a table whose columns have names. The command `with` can be used to compute with those columns without naming the table each time.

```
with(sleep, mean(extra))
```

```
## [1] 1.54
```

B.3 Lists

A vector holds values of one kind, and arithmetic can be performed on all values at once. A *list* is a second basic container which can hold objects of different kinds and different lengths, *e.g.* a number, a vector, a matrix, or even a piece of text. Usually a list is provided with names for its parts. An entry of a list is accessed by name, using a dollar sign followed by the name.

```
fit <- list(estimate = 2.5, se = 0.4, data = c(2.1, 2.9, 2.5))
fit$estimate
fit$estimate + 1.96 * fit$se
```

```
## [1] 2.5
## [1] 3.284
```

The command `names` can be used to see which parts a list has, and `str` prints a compact summary of the whole list, giving the kind and the size of each part.

```
names(fit)
str(fit)
```

```
## [1] "estimate" "se"      "data"
## List of 3
## $ estimate: num 2.5
## $ se      : num 0.4
## $ data    : num [1:3] 2.1 2.9 2.5
```

A part of a list can be used like an ordinary value of its type: `fit$data` is an ordinary vector, so we can use `length(fit$data)` to get the length of the vector as 3, and all the rules from the previous section apply to this vector. Arithmetic is not possible on the list as a whole, *e.g.* `fit * 2` is an error, since a list can contain parts which do not make sense when multiplied.

Lists are useful because many R commands return a list. For example, a command which fits a model returns an estimate, an uncertainty measure, and some information about the computation, all in a list. We can then take the relevant part of the list. In the section about numerical maximum likelihood, below, `optim` returns the results in a list and the estimate can be found as `fit$par`.

B.4 Random Numbers and Distributions

The basis of simulation is the use of random numbers. The names of the distribution functions in R are constructed using the following scheme: a one-letter prefix, followed by the name of the distribution. For the normal distribution (**norm**) the corresponding prefixes are **r** for random samples, **d** for the density, **p** for the cumulative distribution function and **q** for the inverse of the cumulative distribution function, *i.e.* the quantile function.

Since the simulation should be repeatable, we first have to fix the state of the random-number generator. This can be done by calling `set.seed` with a fixed seed.

```
set.seed(2026)
rnorm(5)      # five draws from the standard normal
set.seed(2026)
rnorm(5)      # the same seed gives the same draws
```

```
## [1] 0.52058907 -1.07969076 0.13923812 -0.08474878 -0.66663962
## [1] 0.52058907 -1.07969076 0.13923812 -0.08474878 -0.66663962
```

The following three prefixes answer the usual kind of questions about the same distribution.

```
dnorm(0)      # the N(0,1) density at 0
pnorm(1.96)   # the probability of being at most 1.96
qnorm(0.975)  # the 0.975 quantile
```

```
## [1] 0.3989423
## [1] 0.9750021
## [1] 1.959964
```

The same four prefixes can be used for most standard distributions: instead of `norm` we can write `binom`, `pois`, `exp`, `gamma` and so on. Thus, for a task which requires binomial or geometric draws, it is easy to write `rbinom` or `rgeom` and then just look at the parameters listed on the help page. The geometric distribution illustrates the need to read the help page: authors disagree on whether the geometric distribution counts the number of trials until the first success, or the number of failures before the first success. These notes count the number of trials, so that the support of the distribution in table A.1 starts at 1. In contrast, the help page for `rgeom` states that R counts the number of failures and thus starts the output at 0 and is one below our sequence. To compensate for this, we can use the command `rgeom` to simulate the geometric distribution from these notes by adding one to each of the generated values.

```
rbinom(5, size = 10, prob = 0.3) # five Binomial(10, 0.3) draws
rgeom(5, prob = 0.3) + 1        # five Geometric(0.3) draws, starting
                                at 1
```

```
## [1] 0 4 2 4 2
## [1] 4 14 3 3 2
```

B.5 Writing Your Own Code

A simulation repeats the same steps many times, so we can combine these steps into a *function* and then call this function as needed. The inputs of a function are given in brackets, and the output is the value of the last line of the function.

```
simulate_mean <- function(n, mu) {
  obs <- rnorm(n, mean = mu)
  mean(obs)
}
set.seed(1)
simulate_mean(10, 2)      # one estimate from a sample of size 10
```

```
## [1] 2.132203
```

To repeat a command several times, we can use a `for`-loop, where we store the results in a vector.

```
set.seed(1)
estimates <- numeric(5)
for (i in seq_len(5)) {
  estimates[i] <- simulate_mean(10, 2)
}
estimates
```

```
## [1] 2.132203 2.248845 1.866327 2.120730 2.134137
```

The function `replicate` can be used to do the same thing in one line. This is the most commonly used tool for this purpose.

```
set.seed(1)
replicate(5, simulate_mean(10, 2))
```

```
## [1] 2.132203 2.248845 1.866327 2.120730 2.134137
```

We used `seq_len(5)` instead of `1:5` above. While both forms work for positive numbers, the difference becomes apparent when the count may be zero: `1:0` counts down and returns two values, whereas `seq_len(0)` is empty and a loop over it does nothing (the correct behaviour). Using `seq_len` in this way avoids a subtle bug.

```
1:0
seq_len(0)
```

```
## [1] 1 0
## integer(0)
```

A function often has to be applied to each element of a vector, for example to compute an estimate for each of several parameter values. The command `sapply` can be used for this: it applies the function to every element and collects the results in a vector. The function can be written directly inside the call, without being given a name.

```
set.seed(1)
thetas <- c(0.5, 1, 2)
sapply(thetas, function(theta) mean(rexp(1000, rate = theta)))
```

```
## [1] 2.0626054 0.9835434 0.5409070
```

The three values estimate the means $1/\theta$, which are 2, 1 and 0.5.

Some computations depend on the input and may take either of two forms. We can use the construction `if (condition) a else b` to select one of two values depending on the condition. The result of this construction is a value, so we can assign it to a variable or return it from a function.

```
half_width <- function(n) {
  if (n < 2) NA else qt(0.975, df = n - 1) / sqrt(n)
}
half_width(1)
half_width(10)
```

```
## [1] NA
## [1] 0.7153569
```

The value returned by a function is normally the value of the last line, but the command `return` can be used to exit the function earlier. A common use of this command is to put a ‘guard’ at the top of a function which takes care of special cases not handled by the rest of the function. For example, in [lecture 18](#), we will need a function which is to be minimised, but which refuses negative parameter values.


```
nll <- function(theta) {
  if (theta <= 0) return(Inf)
  -sum(dexp(obs, rate = theta, log = TRUE))
}
obs <- c(0.7, 1.4, 0.2)
nll(-1)
nll(1)
```

```
## [1] Inf
## [1] 2.3
```

B.6 Estimation by Simulation

We can study an estimator without any algebra, by simulating a large number of samples, computing the estimate for each sample, and then summarising the results. This is the Monte Carlo method. We will consider an example where we estimate the mean μ of a normal distribution using the sample mean, where $\mu = 2$ and samples of size $n = 10$.

```
mu <- 2
n <- 10
set.seed(2026)
estimates <- replicate(2000, mean(rnorm(n, mean = mu)))
```

The vector `estimates` now contains 2000 sample means. We can take the average of these values to get an estimate of the truth, and the difference between the average and the truth is a measure of the bias of our estimator. Another measure of the quality of our estimator is the mean squared error, which we can estimate by taking the average of the squared differences between the estimates and the truth. For this estimator we know the exact value $\sigma^2/n = 1/10$, so we can use this value to check our estimate of the mean squared error.

```
mean(estimates) - mu          # empirical bias
mean((estimates - mu)^2)      # empirical mean squared error
1 / n                        # the theoretical value
```

```
## [1] 0.0105696
## [1] 0.101612
## [1] 0.1
```

A similar approach can be used to estimate a probability. A logical condition, *e.g.* `abs(estimates - mu) <= 0.5`, results in a vector of true and false values. Since R interprets these values as ones and zeros, the average of these values is the proportion of true values.

```
mean(abs(estimates - mu) <= 0.5)
```

```
## [1] 0.8765
```

This idiom, the mean of a condition, can be used to estimate any probability by simulation, *e.g.* how often a confidence interval covers the truth.

To watch an estimate settle as data accumulate, `cumsum` can be used to compute the running totals, and the running mean is then the quotient of the running total and the running count.

```
set.seed(1)
obs <- rnorm(200, mean = mu)
running <- cumsum(obs) / seq_len(200)
running[c(1, 10, 50, 200)]
```

```
## [1] 1.373546 2.132203 2.100448 2.035540
```

The running mean of the samples converges to 2 as the number of samples increases.

A simulated answer is only as good as the code behind it, and the usual failure is a plausible but wrong number instead of an error message. Two cheap habits can help to avoid this: try a small case first, and check the result against something known. For example, the mean of many standard normals should be close to zero.

```
mean(rnorm(1e5))
```

```
## [1] -0.002290606
```

There are two more idioms which come up repeatedly in the R sessions. First, if the expression to `replicate` returns a vector instead of a single number, *e.g.* the two ends of an interval, the result is a matrix where there is one column for each repetition, and one can select a row of this matrix using `ci[1, ]`. Secondly, if a loop is iterating through different settings, *e.g.* a vector `ns` of sample sizes, one can use `seq_along(ns)` to get the counter values for the positions in the vector, and one can use `data.frame` to summarise the results in tabular form.

```
set.seed(2026)
ci <- replicate(4, mean(rnorm(10, mean = mu)) + c(-1, 1) * 1.96 / sqrt(10))
ci
ci[1, ]
ns <- c(10, 40, 160)
width <- numeric(length(ns))
for (i in seq_along(ns)) width[i] <- 2 * 1.96 / sqrt(ns[i])
data.frame(n = ns, width = width)
```

```
##           [,1]      [,2]      [,3]      [,4]
## [1,] 0.7999089 1.173516 1.949872 1.431851
## [2,] 2.0395218 2.413128 3.189484 2.671464
## [3,] 0.7999089 1.173515 1.949871 1.431850
##      n      width
## 1  10 1.2396128
## 2  40 0.6198064
## 3 160 0.3099032
```

To summarise the contents of such a matrix row by row, one can use `apply`: the call `apply(ci, 1, mean)` applies the `mean` function to each row of the matrix (the 1 indicates that we want to operate on the rows, 2 would indicate the columns), *i.e.* this computes the average of the lower and upper ends of the four intervals. For means and sums specifically, R provides `rowMeans`, `colMeans`, `rowSums` and `colSums`, which do the same job more directly and are used by us when the summary is one of these four.

```
apply(ci, 1, mean)
```

```
## [1] 1.338787 2.578400
```

B.7 Plotting, and Getting Results into a Report

We can use plots to inspect data and results. The main functions used to create plots are `hist` (for histograms), `plot` (for points or lines) and `curve` (for plotting a function). The functions `abline` (to add a straight line) and `lines` (to add a curve) can be used to add to an existing plot. As an example, we consider the histogram of a sample, together with the density of the sample and the sample mean.

```
set.seed(2026)
obs <- rnorm(500, mean = 2, sd = 1)
hist(obs, freq = FALSE, breaks = 20, col = "grey85", main = "", xlab = "
  x")
curve(dnorm(x, mean = 2, sd = 1), add = TRUE, lwd = 2)
abline(v = mean(obs), lty = 2)
```

To compare the density to the histogram, we set `freq = FALSE` to put the histogram on a density scale and `add = TRUE` to plot the curve on top of the histogram. The result is shown in figure B.1.




Figure B.1: A histogram on the density scale, with the true density overlaid and the sample mean marked by a dashed line.

To assess whether data are normally distributed, the function `qqnorm` can be used to plot the sorted values against normal quantiles. If the points lie close to a straight line, the data are approximately normally distributed.

```
qqnorm(obs)
```

To include the plot in a report, we have to save it to a file. We can open a file for graphics output using `pdf` (or `png`) and then use `dev.off` to close the file and indicate that all the output should go to the file instead of to the screen.

```
pdf("myplot.pdf", width = 6, height = 4)
hist(obs, freq = FALSE, breaks = 20)
curve(dnorm(x, mean = 2, sd = 1), add = TRUE)
dev.off()
```

In RStudio we can use the **Export** button above the plot pane to save the current figure. Numbers from the console can be copied and pasted into the report as they appear, and if we store the commands in a script, we can re-run the script to get the exact same result.

B.8 Numerical Maximum Likelihood

For many models the likelihood equation cannot be solved by hand, and we have to find the maximum likelihood estimate numerically. The tool we use here is the `optim` function from R, which *minimises* a function. This means that we have to give the negative of the log-likelihood function as an argument. As an example we consider here a gamma sample with known rate 1 and unknown shape parameter. For this example, the maximum likelihood estimate cannot be found analytically.

```
set.seed(2026)
obs <- rgamma(200, shape = 3, rate = 1) # the true shape is 3
nll <- function(k) -sum(dgamma(obs, shape = k, rate = 1, log = TRUE))
fit <- optim(par = 1, fn = nll, method = "BFGS", hessian = TRUE)
fit$par
```

```
## [1] 3.132089
```

A command like `optim` returns its results in a *list*. The individual results can be accessed by name, using a dollar sign as in `fit$par` above. In our experiment we use the starting value `par`, the method "BFGS" and `hessian = TRUE`, so that `optim` also returns the matrix of second derivatives of the negative log-likelihood at the optimum. This matrix is called the observed information. In the R command below, we use `solve` to invert this matrix and the square root of the diagonal elements to compute approximate standard errors. These can be used to compute a Wald confidence interval.

```
se <- sqrt(diag(solve(fit$hessian)))
se
fit$par + c(-1, 1) * 1.96 * se
```

```
## [1] 0.1153831  
## [1] 2.905938 3.358240
```

The estimate 3.13 and its interval are close to the true shape 3, as expected. For a one-dimensional parameter, **optimise** is a simpler alternative to **optim**, which finds the minimum in an interval, but does not return the information.

Appendix C

Solutions to the Exercises

This appendix contains worked solutions to the exercises found at the end of each lecture. The solutions are grouped by lecture and refer to the exercises by number.

C.1 Statistical Models and the Estimation Problem

Here we solve the exercises at the end of lecture 1.

Solution to exercise 1.1.

1. Since the tosses are independent and since each toss has probability θ of being heads, we can describe $X_1, \dots, X_n$ as a sequence of independent and identically distributed Bernoulli trials with probability weights $f(x; \theta) = \theta^x(1 - \theta)^{1-x}$ for $x \in \{0, 1\}$ and parameter space $\Theta = (0, 1)$.
2. We can use the proportion of heads in the sample, $\hat{\theta} = \bar{X} = (1/n) \sum_{i=1}^n X_i$, as an estimator.
3. The value $\hat{\theta}(X_1, \dots, X_n)$ is a function of the random variables X_i before the coin tosses occur, so the value is still indeterminate. After the tosses have been made and the numbers $x_1, \dots, x_n$ are known, $\hat{\theta}(x_1, \dots, x_n)$ is a function of numbers, *i.e.* a single number.

□

Solution to exercise 1.2. A statistic is a function of the data which does not use the unknown parameter $\theta = (\mu, \sigma^2)$.

1. The range $\max_i X_i - \min_i X_i$ is computed from the data alone, so it is a statistic.
2. The value $(1/n) \sum_i (X_i - \mu)^2$ contains the unknown mean μ and thus cannot be computed from the data, so it is not a statistic. If μ was known, then it would be a statistic.
3. The sample variance $(1/(n-1)) \sum_i (X_i - \bar{X})^2$ uses the sample mean $\bar{X}$, which is computed from the data, in place of μ . Thus, it is a statistic.

□

Solution to exercise 1.3.

1. Since the observations are independent, the joint density is the product of the marginal densities:

$$f(x_1, \dots, x_n; \theta) = \prod_{i=1}^n \frac{1}{\theta} = \frac{1}{\theta^n},$$

whenever all observations are in the support, *i.e.* when $0 \leq x_i \leq \theta$ for all i . Otherwise, the density is zero. This can also be written as the density being non-zero only when $\min_i x_i \geq 0$ and $\max_i x_i \leq \theta$.

2. Each observation satisfies $X_i \leq \theta$ with probability one, since the support is $[0, \theta]$. The maximum of numbers, none of which are larger than θ , cannot exceed θ . Thus, $\hat{\theta} = \max_i X_i \leq \theta$.
3. Since the distribution of each X_i is continuous, we have $\mathbb{P}_\theta(X_i = \theta) = 0$ for all i . The maximum equals θ only if at least one observation equals this value, and thus we get

$$\mathbb{P}_\theta(\hat{\theta} = \theta) \leq \sum_{i=1}^n \mathbb{P}_\theta(X_i = \theta) = 0.$$

Together with the first part of the solution this shows that $\hat{\theta} < \theta$ with probability one. Thus, the random variable $\theta - \hat{\theta}$ is non-negative and is zero with probability zero. A non-negative random variable with expectation zero equals zero with probability one, and thus we find $\mathbb{E}_\theta(\theta - \hat{\theta}) > 0$, which is the same as $\mathbb{E}_\theta(\hat{\theta}) < \theta$. Thus, the estimator never overestimates θ and underestimates θ on average.

□

C.2 Bias, Mean Squared Error and Consistency

Here we solve the exercises at the end of lecture 2.

Solution to exercise 2.1.

1. Since $\mathbb{E}_\theta(X_i) = \theta$, we can use lemma 2.2 to conclude $\mathbb{E}_\theta(\bar{X}) = \theta$. Thus we have $\text{bias}(\hat{\theta}) = \mathbb{E}_\theta(\bar{X}) - \theta = 0$ for all θ and the estimator $\hat{\theta}$ is unbiased.
2. Since the estimator is unbiased, the mean squared error equals the variance. Using lemma 2.2 and the fact that $\text{Var}_\theta(X_i) = \theta$ for the Poisson distribution, we find

$$\text{MSE}(\hat{\theta}) = \text{Var}_\theta(\bar{X}) = \frac{\theta}{n}.$$

Since this converges to zero as $n \rightarrow \infty$, we can use theorem 2.10 to conclude that the estimator is consistent.

□

Solution to exercise 2.2.

1. Using the given density

$$\mathbb{E}_\theta(\hat{\theta}) = \int_0^\theta t \cdot \frac{n t^{n-1}}{\theta^n} dt = \frac{n}{\theta^n} \cdot \frac{\theta^{n+1}}{n+1} = \frac{n\theta}{n+1},$$

we find the bias to be $\text{bias}(\hat{\theta}) = n\theta/(n+1) - \theta = -\theta/(n+1)$. Since this converges to zero as $n \rightarrow \infty$, the estimator is asymptotically unbiased. Since the bias is negative, the maximum systematically underestimates θ , as expected from exercise 1.3.

2. We have

$$\mathbb{E}_\theta(\hat{\theta}^2) = \int_0^\theta t^2 \cdot \frac{n t^{n-1}}{\theta^n} dt = \frac{n\theta^2}{n+2},$$

and thus

$$\text{Var}(\hat{\theta}) = \frac{n\theta^2}{n+2} - \left(\frac{n\theta}{n+1}\right)^2 = \frac{n\theta^2}{(n+2)(n+1)^2},$$

where we used $(n+1)^2 - n(n+2) = 1$ in the last step. By theorem 2.7 we find

$$\begin{aligned} \text{MSE}(\hat{\theta}) &= \text{Var}(\hat{\theta}) + \text{bias}(\hat{\theta})^2 \\ &= \frac{n\theta^2}{(n+2)(n+1)^2} + \frac{\theta^2}{(n+1)^2} \\ &= \frac{2\theta^2}{(n+1)(n+2)}. \end{aligned}$$

Since this converges to zero as $n \rightarrow \infty$, we can conclude that $\hat{\theta}$ is consistent by theorem 2.10.

3. The mean squared error is of order $1/n^2$, while the mean squared error of the sample mean is of order $1/n$. Thus, the maximum converges to the target faster than the average does. The very fast convergence is caused by the support of the uniform model changing as a function of the parameter. This effect, mentioned in example 1.2, will be seen again in lecture 13.

□

Solution to exercise 2.3. Since $\hat{\theta} = n/T$, we have $\mathbb{E}_\theta(\hat{\theta}) = n \mathbb{E}_\theta(1/T)$ and using the gamma density of T we get

$$\mathbb{E}_\theta(1/T) = \int_0^\infty \frac{1}{x} \cdot \frac{\theta^n x^{n-1} e^{-\theta x}}{(n-1)!} dx = \frac{\theta^n}{(n-1)!} \int_0^\infty x^{n-2} e^{-\theta x} dx.$$

For $n \geq 2$, the remaining integral can be evaluated as the gamma integral $\int_0^\infty x^{n-2} e^{-\theta x} dx = (n-2)!/\theta^{n-1}$. Thus we get

$$\mathbb{E}_\theta(\hat{\theta}) = \frac{n\theta^n}{(n-1)!} \cdot \frac{(n-2)!}{\theta^{n-1}} = \frac{n\theta}{n-1}.$$

The bias is given by

$$\text{bias}(\hat{\theta}) = \frac{n\theta}{n-1} - \theta = \frac{\theta}{n-1}.$$

Since this is strictly positive, we see that $\hat{\theta}$ systematically overestimates θ on average and thus is biased. As $n \rightarrow \infty$, the bias converges to zero and thus $\hat{\theta}$ is asymptotically unbiased. This completes the solution. □

C.3 Likelihood and the Method of Moments

Here we solve the exercises at the end of lecture 4.

Solution to exercise 4.1.

1. Using the Poisson weights from table 1.1, we find the likelihood to be

$$L(\theta) = \prod_{i=1}^n \frac{\theta^{x_i}}{x_i!} e^{-\theta} = \frac{\theta^{\sum_i x_i}}{\prod_i x_i!} e^{-n\theta}.$$

Taking logarithms we get

$$\ell(\theta) = \log \theta \sum_{i=1}^n x_i - n\theta - \sum_{i=1}^n \log(x_i!),$$

as claimed.

2. Taking derivatives with respect to θ , and noting that the last term does not depend on θ , we find the score function to be

$$\ell'(\theta) = \frac{1}{\theta} \sum_{i=1}^n x_i - n.$$

Setting this equal to zero and solving gives $\theta = (1/n) \sum_i x_i = \bar{x}$. Thus, the score function is zero at the sample mean.

3. For the Poisson distribution we have $\mu_1(\theta) = \mathbb{E}_\theta(X_1) = \theta$. Thus, the only moment equation we have is $\hat{\theta} = \bar{X}$: the method of moments estimator is the sample mean. This is the value at which the score function is zero, *i.e.* for the Poisson distribution the method of moments coincides with the maximum likelihood estimator from lecture 5.

□

Solution to exercise 4.2.

1. Since we have $\mu_1(p) = \mathbb{E}_p(X_1) = 1/p$ and the parameter is scalar, we get the moment equation $1/\hat{p} = \bar{X}$. Solving this equation we find

$$\hat{p} = \frac{1}{\bar{X}}.$$

2. Since each observation satisfies $X_i \geq 1$, we have $\bar{X} \geq 1$, with equality if and only if all observations equal one. Thus, the interval of $\hat{p} = 1/\bar{X}$ values where the estimate is possible is $(0, 1]$. Since for every $p \in (0, 1)$ the dataset with values in $\{1, 2, \dots\}$ has positive probability, no estimate in this range is incompatible with the data (this is in contrast to the example of the uniform distribution in the text).

□

Solution to exercise 4.3.

1. From table A.2 we know that $\mathbb{E}_\theta(X_1) = \mu$ and $\text{Var}_\theta(X_1) = \sigma^2$. Thus, the first population moment is $\mu_1(\theta) = \mu$. Rewriting the relation $\text{Var}(X_1) = \mathbb{E}(X_1^2) - \mathbb{E}(X_1)^2$ we find

$$\mu_2(\theta) = \mathbb{E}_\theta(X_1^2) = \text{Var}_\theta(X_1) + \mathbb{E}_\theta(X_1)^2 = \sigma^2 + \mu^2,$$

as claimed.

2. Using the moments from the first part, the moment equations from definition 4.6 are

$$\hat{\mu} = m_1 = \bar{X} \quad \text{and} \quad \hat{\sigma}^2 + \hat{\mu}^2 = m_2 = \frac{1}{n} \sum_{i=1}^n X_i^2.$$

From the first equation we find $\hat{\mu} = \bar{X}$ directly. Substituting this value into the second equation we find

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n X_i^2 - \bar{X}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2,$$

where the last equality sign comes from expanding the square $(X_i - \bar{X})^2 = X_i^2 - 2X_i\bar{X} + \bar{X}^2$ and using the relation $\sum_i X_i = n\bar{X}$. These are the estimators given in example 4.8 and thus the solution is complete.

□

Solution to exercise 4.4.

1. From table A.2 we know that the uniform model has $\mathbb{E}_\theta(X_1) = \theta/2$ and $\text{Var}_\theta(X_1) = \theta^2/12$. Using lemma 2.2 we find $\mathbb{E}_\theta(\hat{\theta}) = 2\mathbb{E}_\theta(\bar{X}) = 2 \cdot \theta/2 = \theta$ and thus the estimator $\hat{\theta}$ is unbiased. Since the estimator is unbiased, the mean squared error equals the variance by theorem 2.7, and using lemma 2.2 again we get

$$\text{MSE}(\hat{\theta}) = \text{Var}(2\bar{X}) = 4 \cdot \frac{\theta^2/12}{n} = \frac{\theta^2}{3n}.$$

2. The maximum has smaller mean squared error, if and only if

$$\frac{2\theta^2}{(n+1)(n+2)} < \frac{\theta^2}{3n}.$$

This inequality is equivalent to $6n < (n+1)(n+2)$, i.e. to $0 < n^2 - 3n + 2 = (n-1)(n-2)$ for all $n \geq 3$. (The equality signs hold for $n = 1$ and $n = 2$.) Thus, for three or more observations the biased maximum is better than the unbiased method of moments estimator, and the $1/n^2$ rate of the maximum makes the difference increase quickly with n . This is another example of the lesson from lecture 2 that unbiasedness is not sufficient: the moment estimator squashes the data into the sample mean, whereas the maximum uses the fact that the support of the distribution moves with the parameter.

□

Solution to exercise 4.5.

1. The distribution $N(0, \theta^2)$ depends on θ only via θ^2 . Thus, for different parameter values θ and $-\theta$, the distribution is the same, *e.g.* $\theta_1 = 1$ and $\theta_2 = -1$ both lead to $N(0, 1)$. Thus, the model is not identifiable.
2. To make the model identifiable, we can restrict ourselves to $\Theta' = (0, \infty)$: for $0 < \theta_1 < \theta_2$ the distributions $N(0, \theta_1^2)$ and $N(0, \theta_2^2)$ have different variances and thus are different. (The other half line $(-\infty, 0)$ would work equally well.)
3. Since the first moment $\mathbb{E}_\theta(X_1) = 0$ for all θ , we cannot use this moment to derive any information about the parameter. We can use the second moment instead. Since $\mu_2(\theta) = \mathbb{E}_\theta(X_1^2) = \theta^2$, the moment equation is $\hat{\theta}^2 = (1/n) \sum_i X_i^2$ and taking the positive root, as required for $\Theta' = (0, \infty)$, we get

$$\hat{\theta} = \sqrt{\frac{1}{n} \sum_{i=1}^n X_i^2}.$$

This completes the solution.

□

C.4 Maximum Likelihood Estimation

Here we solve the exercises at the end of lecture 5.

Solution to exercise 5.1.

1. Using the geometric weights, the likelihood function is

$$L(p) = \prod_{i=1}^n (1-p)^{x_i-1} p = p^n (1-p)^{\sum_i x_i - n}.$$

Taking logarithms we get $\ell(p) = n \log p + (\sum_i x_i - n) \log(1-p)$, as claimed.

2. Taking derivatives we find the score as

$$\ell'(p) = \frac{n}{p} - \frac{\sum_i x_i - n}{1-p}.$$

Setting this equal to zero and clearing the denominators we find $n(1-p) = p(\sum_i x_i - n)$, *i.e.* $n = p \sum_i x_i$. The only solution to this equation is $p = n / \sum_i x_i = 1/\bar{x}$. Since all observations satisfy $x_i \geq 1$, we have $\sum_i x_i \geq n$ and thus

$$\ell''(p) = -\frac{n}{p^2} - \frac{\sum_i x_i - n}{(1-p)^2} < 0$$

for all $p \in (0, 1)$. Thus, ℓ is strictly concave and the root is the global maximum. Thus, the MLE is $\hat{p} = 1/\bar{X}$. (If all observations equal one, *i.e.* if $\sum_i x_i = n$, then $\ell(p) = n \log p$ is strictly increasing and no maximiser exists in the open interval $(0, 1)$. The same boundary phenomenon as for the Bernoulli distribution in the text occurs. In this case, the formula $\hat{p} = 1/\bar{X}$ gives the boundary value 1.)

3. In exercise 4.2 we have found the method of moments estimator to be $\hat{p} = 1/\bar{X}$. Thus, for the geometric distribution, as for the Bernoulli, Poisson and exponential distributions, the two different construction methods coincide.

□

Solution to exercise 5.2.

1. The log-likelihood function is

$$\ell(\theta) = -2n \log \theta + \sum_{i=1}^n \log x_i - \frac{1}{\theta} \sum_{i=1}^n x_i,$$

and the score function is

$$\ell'(\theta) = -\frac{2n}{\theta} + \frac{1}{\theta^2} \sum_{i=1}^n x_i.$$

Setting the score equal to zero, we find $\sum_i x_i = 2n\theta$ and thus the only root is $\theta = \bar{x}/2$. Since the score is positive for $\theta < \bar{x}/2$ and negative for $\theta > \bar{x}/2$, this root is the global maximum and the MLE is $\hat{\theta} = \bar{X}/2$.

2. Since the model is described by $\text{Gamma}(2, 1/\theta)$ in the rate parametrisation, we can use table A.2 to find $\mathbb{E}_\theta(X_1) = 2\theta$. Using lemma 2.2 we find $\mathbb{E}_\theta(\bar{X}) = 2\theta$ and thus $\mathbb{E}_\theta(\hat{\theta}) = (1/2) \cdot 2\theta = \theta$ for all θ . Thus, $\hat{\theta}$ is unbiased.
3. From table A.2 we find the variance to be $\text{Var}_\theta(X_1) = 2\theta^2$. Since the estimator is unbiased, the mean squared error equals the variance and using lemma 2.2 again, we find

$$\text{MSE}(\hat{\theta}) = \text{Var}\left(\frac{\bar{X}}{2}\right) = \frac{1}{4} \cdot \frac{2\theta^2}{n} = \frac{\theta^2}{2n}.$$

Since this converges to zero as $n \rightarrow \infty$, we can conclude that the estimator is consistent by theorem 2.10.

□

Solution to exercise 5.3.

1. Since $\hat{\theta} = n/T$ and since T has gamma density $\theta^n x^{n-1} e^{-\theta x} / (n-1)!$ on $(0, \infty)$, we get the same integral as in exercise 2.3, but with one fewer power of x , i.e. for $n \geq 3$ we have

$$\begin{aligned} \mathbb{E}_\theta(\hat{\theta}^2) &= n^2 \int_0^\infty \frac{1}{x^2} \cdot \frac{\theta^n x^{n-1} e^{-\theta x}}{(n-1)!} dx \\ &= \frac{n^2 \theta^n}{(n-1)!} \cdot \frac{(n-3)!}{\theta^{n-2}} = \frac{n^2}{(n-1)(n-2)} \theta^2, \end{aligned}$$

as claimed.

2. Expanding the square and using the linearity of the expectation, we find

$$\text{MSE}(c\hat{\theta}) = \mathbb{E}_\theta((c\hat{\theta} - \theta)^2) = c^2 \mathbb{E}_\theta(\hat{\theta}^2) - 2c\theta \mathbb{E}_\theta(\hat{\theta}) + \theta^2.$$

In the text we have seen that $\mathbb{E}_\theta(\hat{\theta}) = n\theta/(n-1)$, and using this value and the second moment from the previous part of the question, we get

$$\text{MSE}(c\hat{\theta}) = \left(\frac{n^2}{(n-1)(n-2)} c^2 - \frac{2n}{n-1} c + 1 \right) \theta^2.$$

For the MLE we set $c = 1$ and thus we get

$$\text{MSE}(\hat{\theta}) = \left(\frac{n^2}{(n-1)(n-2)} - \frac{2n}{n-1} + 1 \right) \theta^2 = \frac{n+2}{(n-1)(n-2)} \theta^2.$$

For the unbiased estimator we set $c = (n-1)/n$ and thus we get the mean squared error as the variance

$$\text{MSE}(\tilde{\theta}) = \left(\frac{n-1}{n-2} - 2 + 1 \right) \theta^2 = \frac{\theta^2}{n-2}.$$

Dividing the two expressions we get $\text{MSE}(\tilde{\theta})/\text{MSE}(\hat{\theta}) = (n-1)/(n+2)$, which is always smaller than 1 for $n \geq 3$.

3. The expression inside the brackets is a quadratic function of c with positive leading coefficient, so the minimum is at the point where the derivative is zero. This point is

$$c = \frac{n}{n-1} \cdot \frac{(n-1)(n-2)}{n^2} = \frac{n-2}{n}.$$

Substituting this value back into the expression inside the brackets we find

$$\text{MSE}\left(\frac{n-2}{n} \hat{\theta}\right) = \left(1 - \frac{n-2}{n-1}\right) \theta^2 = \frac{\theta^2}{n-1}.$$

Since $\theta^2/(n-1) < \theta^2/(n-2)$, this value is better than the unbiased estimator ($c = (n-1)/n$) and the MLE ($c = 1$), as shown in the previous part of the question. Since the winning estimator is smaller than the MLE, we know that the estimator is biased downwards. This is the same effect as we have seen for the normal variance in section 2.3, where the optimal divisor was $n+1$ instead of the unbiased value $n-1$ or the value found by maximum likelihood, *i.e.* n : the mean squared error can be minimised by a deliberately biased estimator, and the amount of bias depends on the family of distributions under consideration. There is no one-size-fits-all solution for this kind of problem.

□

Solution to exercise 5.4.

1. Up to the factor $\nu+1$, which we can cancel, the score is $a/(\nu+a^2) + b/(\nu+b^2)$. Multiplying through by the positive quantity $(\nu+a^2)(\nu+b^2)$, the likelihood equation becomes

$$a(\nu+b^2) + b(\nu+a^2) = \nu(a+b) + ab(a+b) = (a+b)(\nu+ab) = 0,$$

as claimed.

2. For $\theta = m + t$ we have $a = d - t$ and $b = -d - t$, and up to an additive constant the log-likelihood is $-\frac{\nu+1}{2}(\log(\nu + a^2) + \log(\nu + b^2))$. The product inside the logarithms is

$$\begin{aligned}(\nu + a^2)(\nu + b^2) &= (\nu + d^2 + t^2 - 2dt)(\nu + d^2 + t^2 + 2dt) \\ &= (\nu + d^2 + t^2)^2 - 4d^2t^2 = (\nu + d^2)^2 + 2(\nu - d^2)t^2 + t^4,\end{aligned}$$

which is the polynomial $P(t)$ and thus $\ell(m + t) = C - \frac{\nu+1}{2} \log P(t)$.

3. Since the logarithm is increasing, $\ell(m + t)$ takes its largest value when $P(t)$ is smallest. Let $u = t^2 \geq 0$. Then we have $P = (\nu + d^2)^2 + 2(\nu - d^2)u + u^2$, *i.e.* we get a quadratic function of u with minimum at $u = d^2 - \nu$. If $d^2 < \nu$, this minimum is at a negative value of u and thus P is increasing on $u \geq 0$ and takes its minimum at $t = 0$. The midpoint m is then the only maximiser of ℓ and, since the factor $\nu + ab = \nu - d^2 + t^2$ from the first part of the solution has no real zero, m is the only root of the likelihood equation. If $d^2 > \nu$, then P takes its minimum at $u = d^2 - \nu$, *i.e.* at $t = \pm\sqrt{d^2 - \nu}$ which are the zeros of $\nu + ab$. These two values are the maximum likelihood estimates, and since P is even in t , they have equal likelihood. Between these two values, at $t = 0$, the polynomial P has a local maximum, because the coefficient $2(\nu - d^2)$ of t^2 is negative. Thus, ℓ has a local minimum at the midpoint. This completes the proof.

□

Solution to exercise 5.5. Since g is a bijection, there exists an inverse function $g^{-1}: \Psi \rightarrow \Theta$ and we can re-index the model, using ψ instead of θ , *i.e.* if we let $\tilde{f}(x; \psi) = f(x; g^{-1}(\psi))$ the family of distributions is unchanged, but the likelihood function is $\tilde{L}(\psi) = L(g^{-1}(\psi))$. For any $\psi \in \Psi$, the value $g^{-1}(\psi)$ falls into the set Θ and by the definition of $\hat{\theta}$ we have

$$\tilde{L}(\psi) = L(g^{-1}(\psi)) \leq L(\hat{\theta}) = \tilde{L}(g(\hat{\theta})).$$

Thus, $g(\hat{\theta})$ is a maximum of the function $\tilde{L}$ on the set Ψ and consequently is a maximum likelihood estimator for ψ . This completes the proof. □

Solution to exercise 5.6.

1. We have $s + t \geq s$ and thus the event $\{X > s + t\}$ is contained in the event $\{X > s\}$. We get

$$\mathbb{P}_\theta(X > s + t | X > s) = \frac{\mathbb{P}_\theta(X > s + t)}{\mathbb{P}_\theta(X > s)} = \frac{e^{-\theta(s+t)}}{e^{-\theta s}} = e^{-\theta t} = \mathbb{P}_\theta(X > t),$$

i.e. the unconditional probability to survive beyond time t . We can conclude from this that, if a lifetime has already passed time s , the remaining lifetime has the same distribution as a fresh lifetime does. In other words, the time already elapsed does not change the distribution of the remaining lifetime.

2. The conditional density of X given the event $\{X > c\}$ is the density restricted to the set (c, ∞) , normalised to be a density again. We get $f(x; \theta)/\mathbb{P}_\theta(X > c)$ for $x > c$. Thus we have

$$\mathbb{E}_\theta(X|X > c) = \frac{1}{e^{-\theta c}} \int_c^\infty x \theta e^{-\theta x} dx.$$

Using integration by parts, with $u = x$ and $v' = \theta e^{-\theta x}$, we find $v = -e^{-\theta x}$ and thus

$$\int_c^\infty x \theta e^{-\theta x} dx = \left[-x e^{-\theta x} \right]_c^\infty + \int_c^\infty e^{-\theta x} dx = c e^{-\theta c} + \frac{e^{-\theta c}}{\theta}.$$

Dividing by $e^{-\theta c}$ gives $\mathbb{E}_\theta(X|X > c) = c + 1/\theta$ as required. This result is consistent with the first part of the solution: given survival until time c , the remaining lifetime $X - c$ is distributed like a fresh exponential lifetime. Thus, the conditional expectation is the unconditional mean $1/\theta$, plus the time c already elapsed, giving $c + 1/\theta$. This completes the solution. $\square$

Solution to exercise 5.7.

1. From the solution to exercise 4.1 we know that the score is given by $\ell'(\theta) = (1/\theta) \sum_i x_i - n$ and this score equals zero at $\theta = \bar{x}$. Let $S = \sum_i x_i \geq 1$. Then the second derivative is $\ell''(\theta) = -S/\theta^2 < 0$ for all $\theta > 0$ and thus ℓ is strictly concave. This shows that the root is the unique global maximum and thus the MLE is given by $\hat{\theta} = \bar{X}$.
2. The map $g(\theta) = e^{-\theta}$ is continuous and strictly decreasing and therefore maps $(0, \infty)$ onto $(0, 1)$ bijectively. Thus, by theorem 5.6, the MLE for $p_0 = g(\theta)$ is given by

$$\hat{p}_0 = e^{-\hat{\theta}} = e^{-\bar{X}}.$$

3. Using the independence of the observations and the given identity with $z = e^{-1/n}$ we get

$$\mathbb{E}_\theta(e^{-\bar{X}}) = \mathbb{E}_\theta\left(\prod_{i=1}^n e^{-X_i/n}\right) = \prod_{i=1}^n \mathbb{E}_\theta\left((e^{-1/n})^{X_i}\right) = \exp(n\theta(e^{-1/n} - 1)).$$

Since $e^u > 1 + u$ for all $u \neq 0$, we can choose $u = -1/n$ to get $e^{-1/n} - 1 > -1/n$ and thus $n(e^{-1/n} - 1) > -1$. This gives

$$\mathbb{E}_\theta(\hat{p}_0) = \exp(n\theta(e^{-1/n} - 1)) > e^{-\theta} = p_0,$$

and thus the estimator systematically overestimates p_0 on average and is biased for all n . On the other hand, expanding the exponential we find $n(e^{-1/n} - 1) = -1 + 1/(2n) - \dots \rightarrow -1$ as $n \rightarrow \infty$ and thus $\mathbb{E}_\theta(\hat{p}_0) \rightarrow e^{-\theta}$ and the estimator is asymptotically unbiased. There is no contradiction to the unbiasedness of $\bar{X}$: unbiasedness is not preserved under the nonlinear map g , as discussed after theorem 5.6, whereas the maximum likelihood property is. $\square$

Solution to exercise 5.8.

1. Solving the equation $\mathbb{P}_\theta(X_1 \leq m) = 1 - e^{-\theta m} = 1/2$ for m we find $e^{-\theta m} = 1/2$ and taking logarithms we find $m = (\log 2)/\theta$. The map $g(\theta) = (\log 2)/\theta$ is continuous and strictly decreasing on $(0, \infty)$ and since $g(\theta) \rightarrow \infty$ as $\theta \rightarrow 0$ and $g(\theta) \rightarrow 0$ as $\theta \rightarrow \infty$, it maps $(0, \infty)$ onto itself in a bijective way.
2. From example 5.2 we know that the MLE for the rate is given by $\hat{\theta} = 1/\bar{X}$. Since g is a bijection, theorem 5.6 shows that the MLE for the median $m = g(\theta)$ is given by

$$\hat{m} = g(\hat{\theta}) = \frac{\log 2}{\hat{\theta}} = (\log 2) \bar{X}.$$

3. From lemma 2.2 we know $\mathbb{E}_\theta(\bar{X}) = \mathbb{E}_\theta(X_1) = 1/\theta$ and thus $\mathbb{E}_\theta(\hat{m}) = (\log 2)/\theta = m$ for all θ and thus $\hat{m}$ is unbiased. There is no contradiction: unbiasedness is a property of a parametrisation and is in general lost under nonlinear transformations of the parameter. Thus, the bias of the MLE $\hat{\theta} = 1/\bar{X}$ for the rate is not relevant for the bias of the estimator $\hat{m} = (\log 2) \bar{X}$ for the median, which is a linear function of $\bar{X}$.

□

C.5 Exponential Families

Solution to exercise 7.1.

1. For $x \in \{1, 2, \dots\}$ we can write

$$\begin{aligned} f(x; p) &= (1-p)^{x-1}p = \exp((x-1)\log(1-p) + \log p) \\ &= \exp\left(x\log(1-p) - \log \frac{1-p}{p}\right). \end{aligned}$$

This is the form of definition 7.1 with $h(x) = 1$ for $x \in \{1, 2, \dots\}$ and $h(x) = 0$ otherwise, natural statistic $T(x) = x$, natural parameter $\eta(p) = \log(1-p)$ and $B(p) = \log((1-p)/p)$. By lemma 7.5 the natural statistic of a sample is $T_n = \sum_{i=1}^n X_i$.

2. For $\eta < 0$ the geometric series gives

$$\sum_{x=1}^{\infty} e^{\eta x} = \frac{e^\eta}{1 - e^\eta},$$

but for $\eta \geq 0$ the series diverges. Thus we have $\mathcal{N} = (-\infty, 0)$ and

$$K(\eta) = \eta - \log(1 - e^\eta), \quad \eta < 0.$$

Substituting $\eta(p) = \log(1-p)$, so that $e^\eta = 1-p$, we get $K(\eta(p)) = \log(1-p) - \log p = B(p)$ as required. As p ranges over $(0, 1)$, the natural parameter $\log(1-p)$ ranges over the entire open interval $(-\infty, 0)$, and thus the family has full rank.

3. We have

$$K'(\eta) = 1 + \frac{e^\eta}{1 - e^\eta} = \frac{1}{1 - e^\eta} \quad \text{and} \quad K''(\eta) = \frac{e^\eta}{(1 - e^\eta)^2}.$$

With $e^\eta = 1 - p$, by theorem 7.7, we find $\mathbb{E}_p(X) = 1/p$ and $\text{Var}_p(X) = (1 - p)/p^2$, the values given in table A.1. The mean $1/p$ is the moment considered in exercise 4.2.

□

Solution to exercise 7.2.

1. For $x \in \{0, 1, \dots, m\}$ we can write

$$f(x; \theta) = \binom{m}{x} \theta^x (1 - \theta)^{m-x} = \binom{m}{x} \exp\left(x \log \frac{\theta}{1 - \theta} + m \log(1 - \theta)\right).$$

This is the form of definition 7.1 with $h(x) = \binom{m}{x}$ for $x \in \{0, 1, \dots, m\}$ and $h(x) = 0$ otherwise, natural statistic $T(x) = x$, natural parameter $\eta(\theta) = \log(\theta/(1 - \theta))$, the same log-odds as in example 7.2, and $B(\theta) = -m \log(1 - \theta)$.

2. Using the binomial theorem, we find

$$\sum_{x=0}^m \binom{m}{x} e^{\eta x} = (1 + e^\eta)^m$$

for all $\eta \in \mathbb{R}$. Thus, $\mathcal{N} = \mathbb{R}$ and $K(\eta) = m \log(1 + e^\eta)$. This is m times the cumulant function of the Bernoulli distribution.

3. Taking derivatives, we get

$$K'(\eta) = \frac{me^\eta}{1 + e^\eta} \quad \text{and} \quad K''(\eta) = \frac{me^\eta}{(1 + e^\eta)^2}.$$

Since $e^\eta/(1 + e^\eta) = \theta$ and $1/(1 + e^\eta) = 1 - \theta$, theorem 7.7 gives $\mathbb{E}_\theta(X) = m\theta$ and $\text{Var}_\theta(X) = m\theta(1 - \theta)$, in agreement with table A.1. For a sample of size n the natural statistic is $T_n = \sum_i X_i$ and, by the remark after theorem 7.7, we have $\mathbb{E}_\theta(T_n) = nm\theta$ and $\text{Var}_\theta(T_n) = nm\theta(1 - \theta)$.

□

Solution to exercise 7.3.

1. We can write $x^{\alpha-1} = e^{(\alpha-1)\log x}$ and, collecting all factors into one exponential, for $x > 0$ we find

$$\begin{aligned} f(x; \alpha, \beta) &= \frac{\beta^\alpha}{\Gamma(\alpha)} e^{(\alpha-1)\log x} e^{-\beta x} \\ &= \exp\left((\alpha - 1) \log x + (-\beta) x - (\log \Gamma(\alpha) - \alpha \log \beta)\right). \end{aligned}$$

This is the form of definition 7.8 with $k = 2$, with $h(x) = 1$ for $x > 0$ and $h(x) = 0$ otherwise, natural statistic $(T_1(x), T_2(x)) = (\log x, x)$, natural parameter $(\eta_1(\alpha, \beta), \eta_2(\alpha, \beta)) = (\alpha - 1, -\beta)$ and $B(\alpha, \beta) = \log \Gamma(\alpha) - \alpha \log \beta$.

- As α varies over $(0, \infty)$, the first coordinate $\eta_1 = \alpha - 1$ varies over $(-1, \infty)$ and as β varies over $(0, \infty)$, the second coordinate $\eta_2 = -\beta$ varies over $(-\infty, 0)$, independently of each other. Thus, the set of natural parameter values is the open quadrant $\{(\eta_1, \eta_2) \mid \eta_1 > -1, \eta_2 < 0\}$. This is a non-empty open subset of $\mathbb{R}^2$ and the family has full rank as in definition 7.10.
- Knowing α , we can include the factor $x^{\alpha-1}$ into the function h : for $x > 0$ we have

$$f(x; \beta) = x^{\alpha-1} \exp\left((-\beta)x - (\log \Gamma(\alpha) - \alpha \log \beta)\right),$$

This is a one-parameter exponential family in the sense of definition 7.1, with $h(x) = x^{\alpha-1}$ for $x > 0$, natural statistic $T(x) = x$ and natural parameter $-\beta$.

□

Solution to exercise 7.4.

- Using the formula $K(\eta) = -\log(-\eta)$ on $\mathcal{N} = (-\infty, 0)$ we find $K'(\eta) = -1/\eta$ and $K''(\eta) = 1/\eta^2$. Substituting the natural parameter $\eta = -\theta$ from example 7.4, we can use theorem 7.7 to find $\mathbb{E}_\theta(X) = 1/\theta$ and $\text{Var}_\theta(X) = 1/\theta^2$. This is the mean and variance of the exponential distribution with rate θ .
- For $h(x) = (2\pi\sigma^2)^{-1/2}e^{-x^2/(2\sigma^2)}$ and $T(x) = x$ we have to evaluate

$$Z(\eta) = \int_{-\infty}^{\infty} \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{x^2 - 2\sigma^2\eta x}{2\sigma^2}\right) dx.$$

We can complete the square to find $x^2 - 2\sigma^2\eta x = (x - \sigma^2\eta)^2 - \sigma^4\eta^2$ and thus

$$Z(\eta) = \exp\left(\frac{\sigma^2\eta^2}{2}\right) \int_{-\infty}^{\infty} \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \sigma^2\eta)^2}{2\sigma^2}\right) dx = \exp\left(\frac{\sigma^2\eta^2}{2}\right),$$

since the remaining integrand is the density of $N(\sigma^2\eta, \sigma^2)$ and thus integrates to one. The integral is finite for all η and thus $\mathcal{N} = \mathbb{R}$ and $K(\eta) = \log Z(\eta) = \sigma^2\eta^2/2$. Taking derivatives we find $K'(\eta) = \sigma^2\eta$ and $K''(\eta) = \sigma^2$ and substituting $\eta = \mu/\sigma^2$ in theorem 7.7 we get $\mathbb{E}_\mu(X) = \mu$ and $\text{Var}_\mu(X) = \sigma^2$, as expected. Finally, $K(\eta(\mu)) = \sigma^2(\mu/\sigma^2)^2/2 = \mu^2/(2\sigma^2)$, which is the function $B(\mu)$ from example 7.9.

□

Solution to exercise 7.5.

- By lemma 7.5, applied in the natural parametrisation where $B = K$, the log-likelihood is

$$\ell(\eta) = \sum_{i=1}^n \log h(x_i) + \eta T_n - nK(\eta), \quad T_n = \sum_{i=1}^n T(x_i),$$

and thus the score is $\ell'(\eta) = T_n - nK'(\eta)$. The likelihood equation $\ell'(\eta) = 0$ is therefore $K'(\eta) = T_n/n$, and by theorem 7.7 the left-hand side equals $\mathbb{E}_\eta(T(X_1))$. This is the stated equation.

2. Differentiating once more and using theorem 7.7 again, we find

$$\ell''(\eta) = -nK''(\eta) = -n \operatorname{Var}_\eta(T(X_1)) < 0$$

for all $\eta \in \mathcal{N}$. Thus ℓ is strictly concave on the interval $\mathcal{N}$, and a strictly concave function has at most one critical point, which is then its global maximum. Any root of the likelihood equation is therefore the unique maximiser of ℓ , *i.e.* the unique MLE for η .

3. For the Poisson distribution we have $T(x) = x$ and $K'(\eta) = e^\eta$, and thus the likelihood equation reads $e^\eta = \bar{x}$, with root $\hat{\eta} = \log \bar{x}$ whenever $\bar{x} > 0$. Since $\theta = e^\eta$ is a bijection between $\mathcal{N} = \mathbb{R}$ and $\Theta = (0, \infty)$, theorem 5.6 gives $\hat{\theta} = e^{\hat{\eta}} = \bar{X}$, the MLE of exercise 5.7. For the Bernoulli distribution we have $T(x) = x$ and $K'(\eta) = e^\eta/(1+e^\eta)$, and thus the likelihood equation reads $e^\eta/(1+e^\eta) = \bar{x}$, with root $\hat{\eta} = \log(\bar{x}/(1-\bar{x}))$ whenever $0 < \bar{x} < 1$. Since $\theta = e^\eta/(1+e^\eta)$ is a bijection between $\mathbb{R}$ and $(0, 1)$, theorem 5.6 gives $\hat{\theta} = \bar{X}$, the MLE of example 5.3. The restriction $0 < \bar{x} < 1$ is the condition $0 < s < n$ discussed after that example: when it fails the likelihood equation has no root and the MLE does not exist in the open parameter space.
4. The equation of the first part says that the MLE is the parameter value under which the population mean of the natural statistic T agrees with its sample mean, *i.e.* the MLE is the method of moments estimator based on the moment $\mathbb{E}_\theta(T(X_1))$ rather than on $\mathbb{E}_\theta(X_1)$. Whenever $T(x) = x$, as for the Bernoulli, Poisson and exponential distributions and for the normal mean with known variance, the two moments are the same, and the MLE coincides with the method of moments estimator from definition 4.6; this is the coincidence observed in the remark at the end of the method of moments section of lecture 4. For the gamma distribution of exercise 7.3 the natural statistic is $(\log x, x)$, and thus maximum likelihood matches the means of $\log X_1$ and X_1 , whereas the method of moments in example 4.9 matched the means of X_1 and X_1^2 . Different moments are matched, and thus the two estimators differ.

□

Solution to exercise 7.6.

1. Writing the expectation as an integral (or a sum, for a discrete model), we find

$$\begin{aligned} \mathbb{E}_\eta(e^{sT(X)}) &= \int e^{sT(x)} h(x) e^{\eta T(x) - K(\eta)} dx \\ &= e^{-K(\eta)} \int h(x) e^{(\eta+s)T(x)} dx = e^{K(\eta+s) - K(\eta)}, \end{aligned}$$

where the last step uses the definition of K at the point $\eta + s \in \mathcal{N}$. Taking logarithms gives $\log \mathbb{E}_\eta(e^{sT(X)}) = K(\eta + s) - K(\eta) = M(s)$, as claimed.

2. By the first part we have $e^{M(s)} = \mathbb{E}_\eta(e^{sT(X)})$. Differentiating both sides with respect to s , and taking the derivative inside the expectation on the

right, we get

$$\begin{aligned} M'(s) e^{M(s)} &= \mathbb{E}_\eta(T(X) e^{sT(X)}) \quad \text{and} \\ (M''(s) + M'(s)^2) e^{M(s)} &= \mathbb{E}_\eta(T(X)^2 e^{sT(X)}). \end{aligned}$$

At $s = 0$ we have $M(0) = 0$, and thus $M'(0) = \mathbb{E}_\eta(T(X))$ and $M''(0) = \mathbb{E}_\eta(T(X)^2) - \mathbb{E}_\eta(T(X))^2 = \text{Var}_\eta(T(X))$. On the other hand, from $M(s) = K(\eta + s) - K(\eta)$ we get directly $M'(0) = K'(\eta)$ and $M''(0) = K''(\eta)$. Comparing the two computations gives $\mathbb{E}_\eta(T(X)) = K'(\eta)$ and $\text{Var}_\eta(T(X)) = K''(\eta)$, which is theorem 7.7.

3. For the Poisson distribution we have $T(x) = x$, $\eta = \log \theta$ and $K(\eta) = e^\eta$ by example 7.3 and the computation following definition 7.6, and $\mathcal{N} = \mathbb{R}$, so that $\eta + s \in \mathcal{N}$ for every s . Writing $z = e^s$, the first part gives

$$\mathbb{E}_\theta(z^X) = \mathbb{E}_\eta(e^{sX}) = \exp(e^{\eta+s} - e^\eta) = \exp(\theta z - \theta) = e^{\theta(z-1)},$$

which is the identity used in exercise 5.7.

□

Solution to exercise 7.7.

1. Using the notation $\mu = \theta$ and $\sigma^2 = \theta^2$ from example 7.9, we see that the density of $N(\theta, \theta^2)$ is of the form given in definition 7.8 with $k = 2$, where $h(x) = 1$, the natural statistic is (x, x^2) and the natural parameter is

$$\left(\frac{\mu}{\sigma^2}, -\frac{1}{2\sigma^2} \right) = \left(\frac{1}{\theta}, -\frac{1}{2\theta^2} \right).$$

2. Let $\eta_1 = 1/\theta$ and $\eta_2 = -1/(2\theta^2)$. As θ varies over the interval $(0, \infty)$, the first coordinate η_1 also varies over the interval $(0, \infty)$. Substituting $\theta = 1/\eta_1$ for the second coordinate, we find

$$\eta_2 = -\frac{1}{2\theta^2} = -\frac{\eta_1^2}{2}.$$

Thus, the set of natural parameter values is the curve $\{(\eta_1, -\eta_1^2/2) \mid \eta_1 > 0\}$. Any open subset of $\mathbb{R}^2$ which contains a point on the curve also contains a small disc around this point, and thus points off the curve, obtained by slightly changing η_2 while keeping η_1 fixed. Thus the curve does not contain any non-empty open subsets of $\mathbb{R}^2$. By definition 7.10, the family is not of full rank, as claimed.

□

C.6 Sufficiency

Here we solve the exercises at the end of lecture 8.

Solution to exercise 8.1.

1. If $\sum_i x_i = t$, the event $\{X = x\}$ is contained in the event $\{T = t\}$ and we have

$$\mathbb{P}_\theta(X = x|T = t) = \frac{\mathbb{P}_\theta(X = x)}{\mathbb{P}_\theta(T = t)} = \frac{\theta^t(1 - \theta)^{n-t}}{\binom{n}{t}\theta^t(1 - \theta)^{n-t}} = \frac{1}{\binom{n}{t}}.$$

2. If $\sum_i x_i \neq t$, the events $\{X = x\}$ and $\{T = t\}$ are disjoint and the conditional probability is zero. Thus the result does not depend on θ and T is sufficient for θ by definition 8.1.
3. The joint probability weights are given by

$$f(x; \theta) = \prod_{i=1}^n \theta^{x_i} (1 - \theta)^{1-x_i} = \theta^{\sum_i x_i} (1 - \theta)^{n - \sum_i x_i}$$

for $x \in \{0, 1\}^n$. This function is of the form (8.1) with $g(t; \theta) = \theta^t(1 - \theta)^{n-t}$ and $h(x) = 1$. By theorem 8.2 the statistic T is sufficient, which confirms our direct computation.

□

Solution to exercise 8.2.

1. The density of a single observation is $f(x_i; \theta) = \theta^{-1} \mathbf{1}_{\{0 \leq x_i \leq \theta\}}$. Since the joint density is the product of these densities, and since a product of indicators can be written as the indicator of the intersection of the corresponding events, we find that all n conditions $0 \leq x_i \leq \theta$ are satisfied if and only if the smallest observation is greater than or equal to 0 and the largest observation is less than or equal to θ . Thus we have

$$f(x; \theta) = \theta^{-n} \prod_{i=1}^n \mathbf{1}_{\{0 \leq x_i \leq \theta\}} = \theta^{-n} \mathbf{1}_{\{\max_i x_i \leq \theta\}} \mathbf{1}_{\{\min_i x_i \geq 0\}}.$$

2. Let $g(t; \theta) = \theta^{-n} \mathbf{1}_{\{t \leq \theta\}}$ and $h(x) = \mathbf{1}_{\{\min_i x_i \geq 0\}}$. Then g depends on the data only through $t = \max_i x_i$, the function h does not depend on θ and from the first part of the solution we know that $f(x; \theta) = g(\max_i x_i; \theta) h(x)$ for all x and θ . By theorem 8.2 the statistic $X_{(n)} = \max_i X_i$ is sufficient for θ .
3. The factor $\mathbf{1}_{\{\min_i x_i \geq 0\}}$ depends on the data but does not depend on θ and thus can be included in h . The factor $\mathbf{1}_{\{\max_i x_i \leq \theta\}}$ depends on θ and thus a function involving this factor cannot be in h . This factor must be included in g and, since it only depends on the data through $\max_i x_i$, the maximum is the sufficient statistic resulting from this factorisation. It is a common mistake to write the joint density as θ^{-n} and to conclude that no additional statistic is needed, since the indicator is where the data and the parameter interact.

□

Solution to exercise 8.3.

1. For $x_1, \dots, x_n \geq 0$ the joint density is

$$f(x; \theta) = \prod_{i=1}^n \theta e^{-\theta x_i} = \theta^n \exp\left(-\theta \sum_{i=1}^n x_i\right),$$

and it is zero if any x_i is negative. Thus the factorisation (8.1) holds with $g(t; \theta) = \theta^n e^{-\theta t}$ and $h(x) = \mathbf{1}_{\{\min_i x_i \geq 0\}}$, and $T = \sum_i X_i$ is sufficient by theorem 8.2.

2. For two data vectors x and y with non-negative entries, the first part gives

$$\frac{f(x; \theta)}{f(y; \theta)} = \exp\left(-\theta \left(\sum_i x_i - \sum_i y_i\right)\right).$$

If $\sum_i x_i = \sum_i y_i$, the ratio equals 1 for every θ . If the sums differ by $d \neq 0$, the ratio is $e^{-\theta d}$, which is a strictly monotone function of θ and thus not constant on $(0, \infty)$. The ratio is therefore free of θ if and only if $T(x) = T(y)$ and by theorem 8.6 we find that T is minimal sufficient.

3. Since $\bar{X} = T/n$, the MLE is given by $\hat{\theta} = 1/\bar{X} = n/T$, a function of T . For fixed data, the likelihood function is $L(\theta) = g(T; \theta) h(x) = \theta^n e^{-\theta T} h(x)$, where $h(x)$ is a positive constant not depending on θ . Thus, the maximiser of $L(\theta)$ over θ is the same as the maximiser of $\theta^n e^{-\theta T}$, which depends on the data only through T . Whatever the MLE is, it can be written as a function of the sufficient statistic by the factorisation in theorem 8.2.

□

Solution to exercise 8.4.

1. With μ known, the joint density

$$f(x; \sigma^2) = (2\pi\sigma^2)^{-n/2} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2\right)$$

depends on the data only through $t = \sum_i (x_i - \mu)^2$, so that (8.1) holds with $g(t; \sigma^2) = (2\pi\sigma^2)^{-n/2} \exp(-t/(2\sigma^2))$ and $h(x) = 1$. By theorem 8.2 the statistic $\sum_i (X_i - \mu)^2$ is sufficient for σ^2 .

2. Expanding the squares in the exponent, the joint density is

$$\begin{aligned} f(x; \mu, \sigma^2) &= (2\pi\sigma^2)^{-n/2} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2\right) \\ &= (2\pi\sigma^2)^{-n/2} \exp\left(\frac{\mu}{\sigma^2} \sum_i x_i - \frac{1}{2\sigma^2} \sum_i x_i^2 - \frac{n\mu^2}{2\sigma^2}\right), \end{aligned}$$

which depends on the data only through the pair $T = (\sum_i x_i, \sum_i x_i^2)$, and thus the factorisation (8.1) holds with $h(x) = 1$ and with g equal to the whole expression. By theorem 8.2 the pair T is sufficient for (μ, σ^2) .

3. Using the expanded form of the joint density from the previous part, and writing $a = \sum_i x_i - \sum_i y_i$ and $b = \sum_i x_i^2 - \sum_i y_i^2$, the ratio of the densities of two data vectors is

$$\frac{f(x; \mu, \sigma^2)}{f(y; \mu, \sigma^2)} = \exp\left(\frac{a\mu - b/2}{\sigma^2}\right),$$

since the factors $(2\pi\sigma^2)^{-n/2}$ and $\exp(-n\mu^2/(2\sigma^2))$ cancel. If $a = b = 0$, which is the statement $T(x) = T(y)$, the ratio equals 1 for all parameter values. Conversely, suppose the ratio does not depend on (μ, σ^2) , so that the exponent $(a\mu - b/2)/\sigma^2$ equals a constant c for all $\mu \in \mathbb{R}$ and all $\sigma^2 > 0$. Taking $\sigma^2 = 1$, the expression $a\mu - b/2$ is constant in μ , which forces $a = 0$ and then $b = -2c$. With $a = 0$ the exponent is $-b/(2\sigma^2)$, and this is constant in σ^2 only if $b = 0$. Thus the ratio is free of the parameters if and only if $T(x) = T(y)$, and T is minimal sufficient by theorem 8.6.

4. No. By the previous part T is minimal sufficient, and thus, by definition 8.5, T would have to be a function of $\bar{X}$ if $\bar{X}$ were sufficient. It is not: for $n = 2$ the data sets $x = (1, -1)$ and $y = (2, -2)$ have the same sample mean 0, but $\sum_i x_i^2 = 2$ and $\sum_i y_i^2 = 8$, so that $T(x) \neq T(y)$. Two data sets with the same value of $\bar{X}$ can therefore have non-proportional likelihood functions, and $\bar{X}$ alone loses information about σ^2 .

□

Solution to exercise 8.5.

1. The exponential distribution with rate θ has distribution function $F(y) = 1 - e^{-\theta y}$ for $y \geq 0$, and thus $1 - F(y) = e^{-\theta y}$. Proposition 8.9 gives

$$f_{(1)}(y) = n(e^{-\theta y})^{n-1} \theta e^{-\theta y} = n\theta e^{-n\theta y}, \quad y \geq 0,$$

which is the exponential density with rate $n\theta$. Directly, the minimum exceeds y if and only if every observation does, and by independence

$$\mathbb{P}_\theta(X_{(1)} > y) = \prod_{i=1}^n \mathbb{P}_\theta(X_i > y) = (e^{-\theta y})^n = e^{-n\theta y},$$

which is the probability that an exponential random variable with rate $n\theta$ exceeds y . The two computations agree.

2. From table A.2, an exponential random variable with rate $n\theta$ has mean $1/(n\theta)$ and variance $1/(n\theta)^2$. Thus $\mathbb{E}_\theta(nX_{(1)}) = n/(n\theta) = 1/\theta$, so $nX_{(1)}$ is unbiased for $1/\theta$, and $\text{Var}_\theta(nX_{(1)}) = n^2/(n\theta)^2 = 1/\theta^2$.
3. By lemma 2.2 the sample mean has $\text{Var}_\theta(\bar{X}) = 1/(n\theta^2)$, and thus we find that $\text{Var}_\theta(nX_{(1)})$ equals $n \text{Var}_\theta(\bar{X})$, so that the estimator based on the minimum has n times the variance of the sample mean. Worse, its variance does not decrease at all as n grows. Indeed, by the first part $nX_{(1)}$ has the exponential distribution with rate θ for every n , and thus its sampling distribution does not change with the sample size, and $nX_{(1)}$ is not a consistent estimator for $1/\theta$. It is not a function of $\sum_i X_i$ either: for $n = 2$ the data sets $(1, 3)$ and $(2, 2)$ have the same sum but minima 1 and 2. The

estimator discards the information which the sufficient statistic retains, and lecture 10 shows how conditioning on a sufficient statistic repairs such an estimator.

□

Solution to exercise 8.6.

1. Since X_1 and X_2 are both Bernoulli with success probability θ , we have $\mathbb{E}_\theta(X_1 - X_2) = \theta - \theta = 0$ for all $\theta \in (0, 1)$ and thus $g(T) = X_1 - X_2$ is an unbiased estimator for zero. The variable $g(T)$ equals 1 on the event $\{X_1 = 1, X_2 = 0\}$ and, since the two events are independent, this happens with probability $\theta(1-\theta) > 0$. Thus we have $\mathbb{P}_\theta(g(T) = 0) \leq 1 - \theta(1-\theta) < 1$.
2. The function g satisfies $\mathbb{E}_\theta(g(T)) = 0$ for all θ , but $g(T)$ does not equal zero with probability one, and thus the condition from definition 8.8 is violated: the whole sample is not complete. Nevertheless, the statistic T is sufficient, because once the whole sample is given, there is no variability left. This completes the proof.

□

C.7 Rao–Blackwell and Lehmann–Scheffe

Solution to exercise 10.1.

1. Since U takes only the values 0 and 1, we have $\mathbb{E}_\theta(U) = \mathbb{P}_\theta(X_1 = 0) = e^{-\theta}$ for all θ and thus U is an unbiased estimator for $e^{-\theta}$. The indicator function is a Bernoulli random variable with success probability $e^{-\theta}$, and thus we have $\text{Var}_\theta(U) = e^{-\theta}(1 - e^{-\theta})$.
2. Let $T = X_1 + R$ where $R = \sum_{i=2}^n X_i$. By the given fact, R is Poisson with parameter $(n-1)\theta$ and T is Poisson with parameter $n\theta$, and R is independent of X_1 . For $0 \leq x \leq t$, the event $\{X_1 = x, T = t\}$ is the event $\{X_1 = x, R = t - x\}$ and thus we get

$$\begin{aligned} \mathbb{P}_\theta(X_1 = x | T = t) &= \frac{\mathbb{P}_\theta(X_1 = x) \mathbb{P}_\theta(R = t - x)}{\mathbb{P}_\theta(T = t)} \\ &= \frac{e^{-\theta} \theta^x}{x!} \cdot \frac{e^{-(n-1)\theta} ((n-1)\theta)^{t-x}}{(t-x)!} \cdot \frac{t!}{e^{-n\theta} (n\theta)^t} \\ &= \frac{t!}{x! (t-x)!} \cdot \frac{(n-1)^{t-x}}{n^t} = \binom{t}{x} \left(\frac{1}{n}\right)^x \left(1 - \frac{1}{n}\right)^{t-x}. \end{aligned}$$

As expected, the parameter θ cancels since T is sufficient. Given $T = t$, the first observation can be seen as the number of successes in t independent trials with success probability $1/n$, and in the model each of the t counted events is equally likely to have been any of the n observations.

3. The conditional expectation of an indicator function is a conditional probability. Thus, using the second part of the question, we find

$$V = \mathbb{E}_\theta(U | T) = \mathbb{P}_\theta(X_1 = 0 | T) = \binom{T}{0} \left(1 - \frac{1}{n}\right)^T = \left(1 - \frac{1}{n}\right)^T.$$

4. Using the probability generating function with $s = 1 - 1/n$, we find

$$\mathbb{E}_\theta(V) = \mathbb{E}_\theta\left(\left(1 - \frac{1}{n}\right)^T\right) = \exp\left(n\theta\left(1 - \frac{1}{n} - 1\right)\right) = e^{-\theta}.$$

This agrees with the second statement of the Rao–Blackwell theorem, theorem 10.1.

5. From example 10.6 we know that the sum T is a complete and sufficient statistic, and V is a function of T which is unbiased for $e^{-\theta}$. Thus, by the Lehmann–Scheffe theorem, theorem 10.4, V is the UMVUE for $e^{-\theta}$. The MLE $e^{-\bar{X}} = e^{-T/n}$ is also a function of T , but in exercise 5.7 we have seen that $\mathbb{E}_\theta(e^{-\bar{X}}) = \exp(n\theta(e^{-1/n} - 1)) > e^{-\theta}$, and thus it is biased and cannot be the UMVUE. Nevertheless, for large n , the two estimators are close since $(1 - 1/n)^T = \exp(T \log(1 - 1/n))$ and $\log(1 - 1/n) = -1/n - 1/(2n^2) - \dots$ and thus $V \approx e^{-\bar{X}} e^{-\bar{X}/(2n)}$: the UMVUE is obtained by applying a small downward correction to the MLE in order to compensate for the upward bias.

□

Solution to exercise 10.2.

1. The density $\theta e^{-\theta x}$ is that of the exponential distribution with rate θ , whose mean is $1/\theta$ by table A.2, and thus X_1 is unbiased for $1/\theta$. The joint density of the sample is

$$f(x_1, \dots, x_n; \theta) = \prod_{i=1}^n \theta e^{-\theta x_i} = \theta^n \exp\left(-\theta \sum_{i=1}^n x_i\right)$$

for all $x_i \geq 0$, which depends on the data only through $\sum_i x_i$, and thus T is sufficient by the factorisation theorem, theorem 8.2, with $h \equiv 1$.

2. The random variables X_1 and R are independent, X_1 has density $f(x; \theta) = \theta e^{-\theta x}$ and R has the quoted gamma density with $k = n - 1$, and thus the pair (X_1, R) has joint density $f(x; \theta) f_{n-1}(r)$ for $x, r \geq 0$. The map $(x, r) \mapsto (x, s) = (x, x + r)$ is linear with determinant one, and thus by the change of variable formula, proposition A.19, the pair (X_1, T) has joint density

$$\begin{aligned} f_{X_1, T}(x, s) &= f(x; \theta) f_{n-1}(s - x) = \theta e^{-\theta x} \cdot \frac{\theta^{n-1} (s - x)^{n-2} e^{-\theta(s-x)}}{(n-2)!} \\ &= \frac{\theta^n (s - x)^{n-2} e^{-\theta s}}{(n-2)!} \end{aligned}$$

for $0 < x < s$, and zero otherwise. Dividing by the density $f_n(s) = \theta^n s^{n-1} e^{-\theta s} / (n-1)!$ of T , we obtain the conditional density

$$f(x|s) = \frac{f_{X_1, T}(x, s)}{f_n(s)} = \frac{(n-1)!}{(n-2)!} \cdot \frac{(s-x)^{n-2}}{s^{n-1}} = \frac{(n-1)(s-x)^{n-2}}{s^{n-1}}$$

for $0 < x < s$. Both θ^n and $e^{-\theta s}$ have cancelled, and thus the conditional density does not depend on θ , which is the sufficiency of T seen from definition 8.1. For $n = 2$ the density is $1/s$ on $(0, s)$: given the total of two observations, the first is uniformly distributed over the possible values.

3. Substituting $x = su$, we find

$$\mathbb{E}_\theta(X_1|T=s) = \int_0^s x \frac{(n-1)(s-x)^{n-2}}{s^{n-1}} dx = (n-1)s \int_0^1 u(1-u)^{n-2} du.$$

Integrating by parts, with u differentiated and $(1-u)^{n-2}$ integrated, gives

$$\int_0^1 u(1-u)^{n-2} du = \left[-\frac{u(1-u)^{n-1}}{n-1} \right]_0^1 + \frac{1}{n-1} \int_0^1 (1-u)^{n-1} du = \frac{1}{(n-1)n},$$

and thus $\mathbb{E}_\theta(X_1|T=s) = s/n$. Consequently $\mathbb{E}_\theta(X_1|T) = T/n = \bar{X}$: Rao–Blackwellising the single observation X_1 , which has variance $1/\theta^2$, gives the sample mean, which has variance $1/(n\theta^2)$, and the reduction promised by theorem 10.1 is by a factor of n .

4. Since the X_i are i.i.d. and T is a symmetric function of them, the pair (X_i, T) has the same joint distribution for every i , and thus the conditional expectations $\mathbb{E}_\theta(X_i|T)$ are the same random variable for all i . Summing over i and using the linearity of conditional expectation, we find

$$n \mathbb{E}_\theta(X_1|T) = \sum_{i=1}^n \mathbb{E}_\theta(X_i|T) = \mathbb{E}_\theta(T|T) = T,$$

and thus $\mathbb{E}_\theta(X_1|T) = T/n$. This argument uses only the symmetry of an i.i.d. sample and none of the exponential form, and so it applies to every model in which $\sum_i X_i$ is sufficient.

□

Solution to exercise 10.3.

1. In exercise 2.2 we have found $\mathbb{E}_\theta(M) = n\theta/(n+1)$ and thus

$$\mathbb{E}_\theta(\hat{\theta}) = \frac{n+1}{n} \cdot \frac{n\theta}{n+1} = \theta$$

for all θ . Thus, $\hat{\theta}$ is unbiased.

2. From exercise 8.2 we know that the maximum M is sufficient for θ and by the second statement of proposition 10.5 it is complete. Since $\hat{\theta}$ is a function of M and is unbiased for θ , we can use the Lehmann–Scheffe theorem, theorem 10.4, to conclude that this is the UMVUE for θ .
3. From the solution to exercise 2.2 we know $\text{Var}_\theta(M) = n\theta^2/((n+2)(n+1)^2)$ and thus

$$\text{Var}_\theta(\hat{\theta}) = \left(\frac{n+1}{n}\right)^2 \cdot \frac{n\theta^2}{(n+2)(n+1)^2} = \frac{\theta^2}{n(n+2)}.$$

The variance of a single observation is $\theta^2/12$, as shown in table A.2, and thus, using lemma 2.2, we find

$$\text{Var}_\theta(\tilde{\theta}) = 4 \text{Var}_\theta(\bar{X}) = \frac{4\theta^2}{12n} = \frac{\theta^2}{3n}.$$

The ratio $\text{Var}_\theta(\hat{\theta})/\text{Var}_\theta(\tilde{\theta}) = 3/(n+2)$, which is always less than or equal to one for all $n \geq 1$, equals one only for $n = 1$, where both estimators equal $2X_1$. For large n , the variance of the UMVUE is of order $1/n^2$, compared to the order $1/n$ for $2\bar{X}$, similar to the result in exercise 2.2.

4. From theorem 10.1, the Rao–Blackwell theorem, we know that the conditional expectation $\mathbb{E}_\theta(\tilde{\theta}|M)$ is a random variable, which is a function of M and is unbiased for θ . Since M is complete, the first step of the proof of theorem 10.4 shows that there can only be one such function and thus we can conclude $\mathbb{E}_\theta(\tilde{\theta}|M) = (n+1)M/n$ with probability one. For $n = 1$ we have $M = X_1$ and $\tilde{\theta} = 2X_1 = 2M$, which coincides with the formula $(n+1)M/n = 2M$.

□

Solution to exercise 10.4.

1. Using the definition of the exponential family, we can write the density as

$$f(x; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right) = \frac{e^{-\mu^2/(2\sigma^2)}}{\sqrt{2\pi\sigma^2}} \exp\left(\frac{\mu}{\sigma^2}x - \frac{1}{2\sigma^2}x^2\right),$$

which is of the form given in definition 7.8 with $h(x) = 1$, natural statistic $T(x) = (x, x^2)$ and natural parameter

$$\eta(\mu, \sigma^2) = \left(\frac{\mu}{\sigma^2}, -\frac{1}{2\sigma^2}\right).$$

2. As (μ, σ^2) varies over $\mathbb{R} \times (0, \infty)$, the second component $\eta_2 = -1/(2\sigma^2)$ can take any value in $(-\infty, 0)$. For fixed σ^2 , the first component $\eta_1 = \mu/\sigma^2$ can take any value in $\mathbb{R}$ as μ varies. Thus, the range of η is the open half-plane $\mathbb{R} \times (-\infty, 0)$, which is an open, non-empty subset of $\mathbb{R}^2$ and thus the family has full rank in the sense of definition 7.10.
3. From corollary 8.3 and the first statement of proposition 10.5 we know that the natural statistic $(\sum_i X_i, \sum_i X_i^2)$ of the sample is complete and sufficient. From the pair $(\sum_i X_i, \sum_i X_i^2)$ we can compute $\bar{X} = \sum_i X_i/n$ and $S^2 = (\sum_i X_i^2 - n\bar{X}^2)/(n-1)$, and conversely $\sum_i X_i = n\bar{X}$ and $\sum_i X_i^2 = (n-1)S^2 + n\bar{X}^2$, and thus the map between the two pairs is one-to-one. Since a one-to-one function of a complete sufficient statistic is itself complete and sufficient, the pair $(\bar{X}, S^2)$ is complete and sufficient.

□

Solution to exercise 10.5.

1. Let $Z_i = X_i - \mu$, so that $Z_1, \dots, Z_n$ are i.i.d. $N(0, \sigma^2)$ whatever the value of μ . Since subtracting the same constant from every observation shifts the maximum and the minimum by that constant, we have

$$R = \max_i X_i - \min_i X_i = \max_i Z_i - \min_i Z_i,$$

and the right-hand side is a function of random variables whose joint distribution does not involve μ . Thus the distribution of R is the same under every $\mathbb{P}_\mu$, and R is ancillary for μ in the sense of definition 10.8.

2. In the proof of corollary 10.10 we have seen that, for known σ^2 , the normal distribution is a full-rank exponential family with natural statistic $T(x) = x$, and thus $\bar{X}$ is complete and sufficient for μ by corollary 8.3 and proposition 10.5. Since R is ancillary by the first part, Basu's theorem, theorem 10.9, shows that $\bar{X}$ and R are independent.

3. Now let μ be known and σ^2 unknown, and write $X_i = \mu + \sigma W_i$ with $W_1, \dots, W_n$ i.i.d. standard normal. Then $R = \sigma (\max_i W_i - \min_i W_i)$, and thus for example $\mathbb{E}_{\sigma^2}(R) = \sigma \mathbb{E}(\max_i W_i - \min_i W_i)$ is proportional to σ , and this expectation is strictly positive rather than zero, since W_1 and W_2 are independent and continuous and thus differ with probability one. The distribution of R changes with the parameter, and R is not ancillary for σ^2 . The scaled range R/σ does have a distribution free of σ^2 , but it is not a statistic in the sense of definition 1.3, because computing it requires the unknown σ . Quantities of this kind, whose distribution is known although they involve the parameter, return in lecture 25, where they are used to build confidence intervals.

□

C.8 Fisher Information

Solution to exercise 11.1.

1. The score for the observation is given by $\ell'_1(\theta) = X_1/\theta - (1 - X_1)/(1 - \theta)$. This takes values $1/\theta$ if $X_1 = 1$ (with probability θ) and $-1/(1 - \theta)$ if $X_1 = 0$ (with probability $1 - \theta$). Thus we have

$$\mathbb{E}_\theta(\ell'_1(\theta)) = \theta \cdot \frac{1}{\theta} + (1 - \theta) \cdot \left(-\frac{1}{1 - \theta}\right) = 1 - 1 = 0,$$

as given in lemma 11.1.

2. Combining the two fractions we get

$$\ell'_1(\theta) = \frac{X_1(1 - \theta) - (1 - X_1)\theta}{\theta(1 - \theta)} = \frac{X_1 - \theta}{\theta(1 - \theta)},$$

which is a constant multiple of $X_1 - \theta$. Since $\text{Var}_\theta(X_1) = \theta(1 - \theta)$ we find

$$\text{Var}_\theta(\ell'_1(\theta)) = \frac{\theta(1 - \theta)}{\theta^2(1 - \theta)^2} = \frac{1}{\theta(1 - \theta)}.$$

This is the value of $\mathcal{I}(\theta)$ in example 11.5 and thus, by theorem 11.3, the correct value.

3. From proposition 11.4 we know $\mathcal{I}_n(\theta) = n/(\theta(1 - \theta))$ and thus $1/\mathcal{I}_n(\theta) = \theta(1 - \theta)/n$. By lemma 2.2 the sample mean has $\text{Var}_\theta(\bar{X}) = \text{Var}_\theta(X_1)/n = \theta(1 - \theta)/n$ and thus the variance of $\bar{X}$ equals the reciprocal of the information, as it did for the normal mean in example 11.7.

□

Solution to exercise 11.2.

1. The log-likelihood of a single observation is $\ell_1(\theta) = -2 \log \theta + \log X_1 - X_1/\theta$ and thus the score is given by

$$\ell'_1(\theta) = -\frac{2}{\theta} + \frac{X_1}{\theta^2}.$$

From table A.2 or from the solution to exercise 5.2 we know $\mathbb{E}_\theta(X_1) = 2\theta$ and thus $\mathbb{E}_\theta(\ell'_1(\theta)) = -2/\theta + 2\theta/\theta^2 = 0$.

2. Taking derivatives again, we find $\ell_1''(\theta) = 2/\theta^2 - 2X_1/\theta^3$ and thus

$$\mathcal{I}(\theta) = -\mathbb{E}_\theta(\ell_1''(\theta)) = -\frac{2}{\theta^2} + \frac{2 \cdot 2\theta}{\theta^3} = \frac{2}{\theta^2}.$$

On the other hand, since the score is X_1/θ^2 plus a constant, we can find the variance as

$$\text{Var}_\theta(\ell_1'(\theta)) = \frac{\text{Var}_\theta(X_1)}{\theta^4} = \frac{2\theta^2}{\theta^4} = \frac{2}{\theta^2}.$$

Both approaches give the same result.

3. Since the density can be written as

$$f(x; \theta) = x \exp\left(-\frac{x}{\theta} - 2 \log \theta\right),$$

we have $h(x) = x$, $T(x) = x$, $\eta(\theta) = -1/\theta$ and $B(\theta) = 2 \log \theta$. Thus, the model is an exponential family as in definition 7.1. The natural parameter takes values $\eta < 0$ and the cumulant function is

$$K(\eta) = \log \int_0^\infty x e^{\eta x} dx = \log \frac{1}{\eta^2} = -2 \log(-\eta),$$

where $K'(\eta) = -2/\eta = 2\theta$ is the mean and $K''(\eta) = 2/\eta^2 = 2\theta^2$ is the variance, as predicted by theorem 7.7. Since $\eta'(\theta) = 1/\theta^2$, we find by proposition 11.9 $\mathcal{I}(\theta) = \eta'(\theta)^2 K''(\eta(\theta)) = 2\theta^2/\theta^4 = 2/\theta^2$. This is the same value as we obtained above.

4. From proposition 11.4 we know that the information of the sample is $\mathcal{I}_n(\theta) = 2n/\theta^2$ and thus $1/\mathcal{I}_n(\theta) = \theta^2/(2n) = \text{MSE}(\hat{\theta})$. Since $\hat{\theta}$ is unbiased, its mean squared error equals its variance and thus the variance of the MLE is the reciprocal of the Fisher information for all sample sizes, as it was for the sample mean in exercise 11.1 and example 11.7. From lecture 13 we know that this is the smallest variance an unbiased estimator can have.

□

Solution to exercise 11.3.

1. The log-likelihood of a single observation is $\ell_1(p) = (X_1 - 1) \log(1 - p) + \log p$ and thus the score is given by

$$\ell_1'(p) = \frac{1}{p} - \frac{X_1 - 1}{1 - p}.$$

Using $\mathbb{E}_p(X_1 - 1) = 1/p - 1 = (1 - p)/p$ we find

$$\mathbb{E}_p(\ell_1'(p)) = \frac{1}{p} - \frac{(1 - p)/p}{1 - p} = \frac{1}{p} - \frac{1}{p} = 0.$$

2. Taking derivatives again, we find $\ell_1''(p) = -1/p^2 - (X_1 - 1)/(1 - p)^2$ and taking expectations we get

$$\begin{aligned}\mathcal{I}(p) &= -\mathbb{E}_p(\ell_1''(p)) = \frac{1}{p^2} + \frac{(1-p)/p}{(1-p)^2} \\ &= \frac{1}{p^2} + \frac{1}{p(1-p)} = \frac{(1-p) + p}{p^2(1-p)} = \frac{1}{p^2(1-p)}.\end{aligned}$$

From proposition 11.4 we know that the sample has information $\mathcal{I}_n(p) = n/(p^2(1-p))$.

3. Since the score is $-X_1/(1-p)$ plus a constant, we find

$$\text{Var}_p(\ell_1'(p)) = \frac{\text{Var}_p(X_1)}{(1-p)^2} = \frac{(1-p)/p^2}{(1-p)^2} = \frac{1}{p^2(1-p)}.$$

This agrees with the result from the previous part.

□

Solution to exercise 11.4.

1. Since $K(\eta) = e^\eta$ we have $K''(\eta) = e^\eta$ and using proposition 11.9 we find $\mathcal{I}_\eta(\eta) = e^\eta = \theta$. This is the variance of the natural statistic $T(X) = X$, as expected.
2. Using proposition 11.10 with η (the original parameter) and $\theta = g(\eta) = e^\eta$ (the new parameter), we find $g'(\eta) = e^\eta = \theta$. Thus we have

$$\mathcal{I}_\theta(\theta) = \frac{\mathcal{I}_\eta(\eta)}{g'(\eta)^2} = \frac{\theta}{\theta^2} = \frac{1}{\theta},$$

the value from example 11.6.

3. We have $p_0 = g(\theta) = e^{-\theta}$ with $g'(\theta) = -e^{-\theta}$ and using proposition 11.10 we find

$$\mathcal{I}_{p_0}(p_0) = \frac{1/\theta}{e^{-2\theta}} = \frac{e^{2\theta}}{\theta},$$

or using the notation p_0 we find $\mathcal{I}_{p_0}(p_0) = 1/(p_0^2 \log(1/p_0))$. As $\theta \rightarrow \infty$, the information goes to infinity. This makes sense, since $p_0 = e^{-\theta}$ converges to zero, and the information measures accuracy on the absolute scale of the parameter: a probability which is close to zero is estimated with small absolute error by any reasonable estimator, and the approximate variance $1/\mathcal{I}_n(p_0) = \theta e^{-2\theta}/n$ from lecture 17 is tiny, even if the error relative to the size of p_0 is not small.

□

Solution to exercise 11.5.

1. The map $g(\theta) = 1/\theta$ maps $(0, \infty)$ to $(0, \infty)$ and $g'(\theta) = -1/\theta^2 \neq 0$. From example 11.8 we know that $\mathcal{I}_\theta(\theta) = 1/\theta^2$. Thus, by proposition 11.10, we have

$$\mathcal{I}_\mu(\mu) = \frac{\mathcal{I}_\theta(\theta)}{g'(\theta)^2} = \frac{1/\theta^2}{1/\theta^4} = \theta^2 = \frac{1}{\mu^2}.$$

This is the value given in section 11.6.

2. Using proposition 11.4, in the μ -parametrisation, the information of the sample about μ is given by $n\mathcal{I}_\mu(\mu) = n/\mu^2$. From table A.2 we know that $\text{Var}_\mu(X_1) = 1/\theta^2 = \mu^2$ and thus we have $\text{Var}_\mu(\bar{X}) = \mu^2/n$, by lemma 2.2. The variance of the unbiased estimator $\bar{X}$ for μ equals the reciprocal μ^2/n of the information of the sample, and from lecture 13 we know that no unbiased estimator can be better.

□

Solution to exercise 11.6. The right-hand side of (11.1) with $k = 1$ is an integral over the whole real line, but $f(x; \theta) = 0$ and thus $\partial f / \partial \theta = 0$ for $x < 0$ and for $x > \theta$; since $x = \theta$ is a single point, it does not contribute to the integral; only $0 \leq x \leq \theta$, where $f(x; \theta) = 1/\theta$, is left. We have

$$\int_0^\theta \frac{\partial}{\partial \theta} f(x; \theta) dx = \int_0^\theta \left(-\frac{1}{\theta^2}\right) dx = -\frac{1}{\theta},$$

while the left-hand side can be found as

$$\frac{d}{d\theta} \int_0^\theta f(x; \theta) dx = \frac{d}{d\theta} 1 = 0.$$

The difference between the two sides is $0 - (-1/\theta) = 1/\theta$, which is the value $f(\theta; \theta) = 1/\theta$ of the density at the upper limit of the integral. The derivative of the integral picks up this boundary term (from the moving upper limit) in addition to the derivative of the integrand, but the differentiation under the integral sign ignores this boundary term. □

C.9 The Cramer–Rao Inequality

Solution to exercise 13.1.

1. For $y \geq 0$ we have

$$\mathbb{P}_\theta(Y_1 \leq y) = \mathbb{P}_\theta\left(X_1 \leq \frac{y}{2\theta}\right) = 1 - e^{-y/2},$$

and taking derivatives we find the density $\frac{1}{2}e^{-y/2}$ for $y \geq 0$. This is the exponential density with rate $1/2$, i.e. the $\text{Gamma}(1, 1/2)$ density from table A.2 and $\text{Gamma}(1, 1/2)$ is χ_2^2 by the statement quoted in the question. Thus we have $Y_1 \sim \chi_2^2$.

2. From the first part and definition A.2 we know that each $Y_i = 2\theta X_i$ has the distribution of $Z_{2i-1}^2 + Z_{2i}^2$ for independent, standard normally distributed Z_{2i-1} and Z_{2i} . Since the X_i are independent, we can assume that all $2n$ are independent normal variables. Thus, $2\theta T = \sum_{i=1}^n Y_i$ has the distribution of $Z_1^2 + \dots + Z_{2n}^2$, i.e. χ_{2n}^2 by definition A.2.
3. The density of $T \sim \text{Gamma}(n, \theta)$ is given by $f_T(t) = \theta^n t^{n-1} e^{-\theta t} / \Gamma(n)$ for $t \geq 0$. For $V = 2\theta T$, using proposition A.19 with inverse map $t = v/(2\theta)$ and derivative $1/(2\theta)$, we find

$$f_V(v) = f_T\left(\frac{v}{2\theta}\right) \frac{1}{2\theta} = \frac{\theta^n (v/(2\theta))^{n-1} e^{-v/2}}{\Gamma(n)} \cdot \frac{1}{2\theta} = \frac{(1/2)^n v^{n-1} e^{-v/2}}{\Gamma(n)}$$

for $v \geq 0$. This is the $\text{Gamma}(n, 1/2)$ density again, by table A.2. Since $n = (2n)/2$, this is the χ_{2n}^2 density and it agrees with the result from the previous part.

4. From definition A.2 we know $\mathbb{E}_\theta(2\theta T) = 2n$ and $\text{Var}_\theta(2\theta T) = 4n$, and thus $\mathbb{E}_\theta(T) = n/\theta$ and $\text{Var}_\theta(T) = 4n/(4\theta^2) = n/\theta^2$. For a single observation, the mean is $1/\theta$ and the variance is $1/\theta^2$, as found in table A.2, and thus the sum of n independent copies has mean n/θ and variance n/θ^2 , as shown above.

□

Solution to exercise 13.2.

1. From example 11.5 we know that $\mathcal{I}_n(\theta) = n/(\theta(1-\theta))$. Thus, the bound from theorem 13.2 is $1/\mathcal{I}_n(\theta) = \theta(1-\theta)/n$.
2. The sample mean is an unbiased estimator for θ and using lemma 2.2 we find $\text{Var}_\theta(\bar{X}) = \theta(1-\theta)/n$, since the variance of a Bernoulli distribution is $\theta(1-\theta)$. This bound coincides with the bound from the definition, and thus $\bar{X}$ is efficient by definition 13.4, since the exchange identity (13.1) holds because the Bernoulli distribution is a regular exponential family and $\bar{X}$ has finite variance, as noted after lemma 13.1.

3. The score of the sample in example 11.5 is given by

$$\ell'(\theta) = \frac{T_n}{\theta} - \frac{n - T_n}{1 - \theta} = \frac{T_n - n\theta}{\theta(1 - \theta)} = \frac{n}{\theta(1 - \theta)} (\bar{X} - \theta),$$

and thus we have $\bar{X} - \theta = c(\theta) \ell'(\theta)$ with $c(\theta) = \theta(1-\theta)/n$. Since $g(\theta) = \theta$ has $g'(\theta) = 1$, this is $g'(\theta)/\mathcal{I}_n(\theta)$ as predicted by proposition 13.5.

4. For $g(\theta) = \theta/(1-\theta)$ we have $g'(\theta) = 1/(1-\theta)^2$ and thus using theorem 13.3 we get the bound

$$\frac{g'(\theta)^2}{\mathcal{I}_n(\theta)} = \frac{1}{(1-\theta)^4} \cdot \frac{\theta(1-\theta)}{n} = \frac{\theta}{n(1-\theta)^3}.$$

The Bernoulli distribution is an exponential family with natural statistic $T(x) = x$ and $\mu(\theta) = \theta$. From the discussion after proposition 13.5 we know that only affine functions of the form $a + b\theta$ can be used to construct an efficient estimator. Since the odds are not an affine function of θ , there can be no efficient estimator for the odds.

5. Any estimator $W = W(X_1, \dots, X_n)$ can take at most 2^n different values, one for each possible value of the sample $x \in \{0, 1\}^n$, and the expectation of the estimator is

$$\mathbb{E}_\theta(W) = \sum_{x \in \{0,1\}^n} W(x) \theta^{\sum_i x_i} (1-\theta)^{n-\sum_i x_i},$$

which is a polynomial in θ of degree at most n . A polynomial function is bounded on $(0, 1)$, whereas $\theta/(1-\theta) \rightarrow \infty$ as $\theta \rightarrow 1$ and thus no choice of W makes the estimator $\mathbb{E}_\theta(W)$ equal to the odds for all θ . Thus there is no unbiased estimator for the odds, efficient or otherwise.

□

Solution to exercise 13.3.

1. From example 11.6 we know that $\mathcal{I}_n(\theta) = n/\theta$. Thus, the Cramer–Rao bound for unbiased estimators for θ is θ/n . Since the Poisson distribution has mean and variance θ , lemma 2.2 shows that $\mathbb{E}_\theta(\bar{X}) = \theta$ and $\text{Var}_\theta(\bar{X}) = \theta/n$. Thus, $\bar{X}$ is an unbiased estimator and satisfies the bound: it is an efficient estimator for θ , since the exchange identity (13.1) holds because the Poisson distribution is a regular exponential family and $\bar{X}$ has finite variance, as noted after lemma 13.1.
2. The log-likelihood function is $\ell(\theta) = T_n \log \theta - n\theta - \sum_i \log(X_i!)$ and thus the score function is

$$\ell'(\theta) = \frac{T_n}{\theta} - n = \frac{n}{\theta} (\bar{X} - \theta).$$

Re-arranging this we find $\bar{X} - \theta = (\theta/n) \ell'(\theta)$, which is condition (13.2) with $W = \bar{X}$, $g(\theta) = \theta$ and $c(\theta) = \theta/n$; indeed $g'(\theta)/\mathcal{I}_n(\theta) = 1/(n/\theta) = \theta/n$.

3. From the discussion after proposition 13.5 we know that in a one-parameter exponential family, only affine functions of the mean $\mu(\theta)$ of the natural statistic allow for efficient estimators. In this case we have $\mu(\theta) = \theta$ and $g(\theta) = e^{-\theta}$ is not an affine function of θ . Thus, no efficient estimator for $e^{-\theta}$ exists. There is no contradiction in this result: example 10.6 shows that $(1 - 1/n)^{T_n}$ has the smallest variance amongst all unbiased estimators for $e^{-\theta}$, but this smallest variance is strictly larger than the Cramer–Rao bound. Efficiency is a stronger property than being UMVUE, and in this case the bound cannot be achieved.

□

Solution to exercise 13.4.

1. Example 11.8 gives $\mathcal{I}(\theta) = 1/\theta^2$, and thus $\mathcal{I}_n(\theta) = n/\theta^2$ and the bound is $1/\mathcal{I}_n(\theta) = \theta^2/n$.
2. Since $\tilde{\theta}$ is unbiased, its mean squared error is its variance, $\text{Var}_\theta(\tilde{\theta}) = \theta^2/(n-2)$. The efficiency is

$$\text{eff}_\theta(\tilde{\theta}) = \frac{\theta^2/n}{\theta^2/(n-2)} = \frac{n-2}{n},$$

which is less than one for every $n \geq 3$ and tends to one as $n \rightarrow \infty$: the estimator falls short of the bound for every finite sample size, but only by a factor which disappears in the limit.

3. The exponential distribution with rate θ is an exponential family with $h(x) = 1$, $\eta(\theta) = -\theta$, $T(x) = x$ and $B(\theta) = -\log \theta$, as in example 7.4, and thus $\mu(\theta) = \mathbb{E}_\theta(X_1) = 1/\theta$. By the discussion after proposition 13.5, an efficient estimator exists only for affine functions $a + b/\theta$ of $\mu(\theta)$. The rate $g(\theta) = \theta$ is not of this form, and thus no unbiased estimator for θ is efficient. The same conclusion can be reached directly: condition (13.2)

with $\ell'(\theta) = n/\theta - T$ would read $W = \theta + c(\theta)n/\theta - c(\theta)T$, and since W does not depend on θ , the coefficients $-c(\theta)$ and $\theta + c(\theta)n/\theta$ would have to be constants b and a , giving $\mathbb{E}_\theta(W) = a + bn/\theta$, which cannot equal θ for all θ .

4. The natural parameter $\eta = -\theta$ ranges over the open interval $(-\infty, 0)$, and thus the family is of full rank and T is complete and sufficient by corollary 8.3 and proposition 10.5. Since $\tilde{\theta}$ is a function of T and unbiased for θ , the Lehmann–Scheffe theorem, theorem 10.4, shows that it is the UMVUE for θ . Together with the previous part this is an example of a UMVUE which is not efficient: the bound θ^2/n is not attained by any unbiased estimator, and the best unbiased estimator has variance $\theta^2/(n-2)$.
5. For $g(\theta) = 1/\theta$ we have $g'(\theta) = -1/\theta^2$, and theorem 13.3 gives the bound

$$\frac{g'(\theta)^2}{\mathcal{I}_n(\theta)} = \frac{1/\theta^4}{n/\theta^2} = \frac{1}{n\theta^2} = \frac{\mu^2}{n}.$$

The sample mean is unbiased for μ and has variance $\text{Var}_\theta(X_1)/n = \mu^2/n$ by lemma 2.2, and thus $\bar{X}$ is efficient for μ , since the exchange identity (13.1) holds because the exponential distribution is a regular exponential family and $\bar{X}$ has finite variance, as noted after lemma 13.1. This is consistent with the previous parts: $\mu(\theta) = 1/\theta$ is the mean of the natural statistic, for which proposition 13.5 guarantees an efficient estimator, whereas the rate is a non-affine function of it. Whether an efficient estimator exists depends on which function of the parameter we set out to estimate.

□

Solution to exercise 13.5.

1. Since μ is given and $\tau = \sigma^2$ is the parameter, the log-likelihood function is $\ell(\tau) = -\frac{n}{2} \log(2\pi\tau) - \frac{1}{2\tau} \sum_i (X_i - \mu)^2$ and thus

$$\ell'(\tau) = -\frac{n}{2\tau} + \frac{1}{2\tau^2} \sum_{i=1}^n (X_i - \mu)^2, \quad \ell''(\tau) = \frac{n}{2\tau^2} - \frac{1}{\tau^3} \sum_{i=1}^n (X_i - \mu)^2.$$

Since $\mathbb{E}((X_i - \mu)^2) = \tau$, we can use theorem 11.3 to find

$$\mathcal{I}_n(\tau) = -\mathbb{E}(\ell''(\tau)) = -\frac{n}{2\tau^2} + \frac{n\tau}{\tau^3} = \frac{n}{2\tau^2} = \frac{n}{2\sigma^4}.$$

Thus, the Cramer–Rao bound is $2\sigma^4/n$.

2. The $(X_i - \mu)/\sigma$ are i.i.d. standard normally distributed, i.e. $\sum_i (X_i - \mu)^2/\sigma^2 \sim \chi_n^2$ by definition A.2, with mean n and variance $2n$. Thus, $\hat{\sigma}^2 = \frac{\sigma^2}{n} \sum_i (X_i - \mu)^2/\sigma^2$ has mean σ^2 and variance $\sigma^4 \cdot 2n/n^2 = 2\sigma^4/n$. This is the bound and thus $\hat{\sigma}^2$ is efficient, since the exchange identity (13.1) holds because the normal distribution with known mean is a regular exponential family and $\hat{\sigma}^2$ has finite variance, as noted after lemma 13.1. Alternatively, the score for the first part can be written as $\ell'(\sigma^2) = \frac{n}{2\sigma^4}(\hat{\sigma}^2 - \sigma^2)$, which satisfies condition (13.2) with $c(\sigma^2) = 2\sigma^4/n$.

3. From proposition A.3 we know that $(n-1)S^2/\sigma^2 \sim \chi_{n-1}^2$, with variance $2(n-1)$. Thus we have

$$\text{Var}(S^2) = \frac{\sigma^4}{(n-1)^2} \cdot 2(n-1) = \frac{2\sigma^4}{n-1} > \frac{2\sigma^4}{n}.$$

Thus, the sample variance does not achieve the bound, and the efficiency of the sample variance is $(n-1)/n$. Since the bound $2\sigma^4/n$ still holds for unknown μ , and since S^2 is the UMVUE for σ^2 by example 10.7, this is a second example of a best unbiased estimator which is not efficient: no unbiased estimator for σ^2 can achieve the bound if μ also needs to be estimated, and the loss of one degree of freedom is the price to be paid for not knowing μ .

□

Solution to exercise 13.6.

1. By proposition 11.4 the information of the sample is n times the information of a single observation, in either parametrisation. Proposition 11.10 gives the relation for a single observation, and multiplying both sides by n we find

$$\mathcal{I}_\psi(\psi) = \frac{\mathcal{I}_\theta(\theta)}{g'(\theta)^2} \quad \text{for } \psi = g(\theta).$$

2. Since $\psi = g(\theta)$, the estimator W is an unbiased estimator for the parameter ψ , and theorem 13.2, applied in the ψ -parametrisation, bounds its variance by $1/\mathcal{I}_\psi(\psi)$. Using the first part we get

$$\frac{1}{\mathcal{I}_\psi(\psi)} = \frac{g'(\theta)^2}{\mathcal{I}_\theta(\theta)},$$

which is the bound of theorem 13.3. The two theorems thus give the same bound, and the bound does not depend on how the model is parametrised. This completes the proof.

□

C.10 Multiparameter Fisher Information

Solution to exercise 14.1.

1. The Jacobian matrix of $(\mu, \sigma) \mapsto (\mu, \sigma^2)$ is given by

$$J = \begin{pmatrix} \frac{\partial \mu}{\partial \mu} & \frac{\partial \mu}{\partial \sigma} \\ \frac{\partial \sigma^2}{\partial \mu} & \frac{\partial \sigma^2}{\partial \sigma} \end{pmatrix} = \begin{pmatrix} 1 & 0 \\ 0 & 2\sigma \end{pmatrix}.$$

This matrix is invertible for $\sigma > 0$. Since all three matrices are diagonal, the product $J^\top \mathcal{I}_n(\mu, \sigma^2) J$ is also diagonal and consists of the entries $1 \cdot (n/\sigma^2) \cdot 1 = n/\sigma^2$ and $2\sigma \cdot n/(2\sigma^4) \cdot 2\sigma = 2n/\sigma^2$. Thus, by proposition 14.4, we have $\mathcal{I}_n^\varphi(\mu, \sigma) = \text{diag}(n/\sigma^2, 2n/\sigma^2)$.

2. The log-likelihood function in terms of (μ, σ) is

$$\ell(\mu, \sigma) = -\frac{n}{2} \log(2\pi) - n \log \sigma - \frac{1}{2\sigma^2} \sum_{i=1}^n (X_i - \mu)^2.$$

Taking derivatives we get $\partial\ell/\partial\mu = \sum_i (X_i - \mu)/\sigma^2$ and $\partial\ell/\partial\sigma = -n/\sigma + \sum_i (X_i - \mu)^2/\sigma^3$. Taking derivatives again, we find

$$\begin{aligned} \frac{\partial^2 \ell}{\partial \mu^2} &= -\frac{n}{\sigma^2}, & \frac{\partial^2 \ell}{\partial \mu \partial \sigma} &= -\frac{2}{\sigma^3} \sum_{i=1}^n (X_i - \mu), \\ \frac{\partial^2 \ell}{\partial \sigma^2} &= \frac{n}{\sigma^2} - \frac{3}{\sigma^4} \sum_{i=1}^n (X_i - \mu)^2. \end{aligned}$$

Using $\mathbb{E}(X_i - \mu) = 0$ and $\mathbb{E}((X_i - \mu)^2) = \sigma^2$ we find that the mixed derivative has expectation zero, and the last term has expectation $n/\sigma^2 - 3n/\sigma^2 = -2n/\sigma^2$. Thus, by proposition 14.3, the information matrix is given by $\text{diag}(n/\sigma^2, 2n/\sigma^2)$. This agrees with the result from part (a).

3. The inverse of the matrix from part (a) is given by $\text{diag}(\sigma^2/n, \sigma^2/(2n))$. Thus, by theorem 14.5, any unbiased estimator for σ has variance of at least $\sigma^2/(2n)$. The sample variance S^2 is an unbiased estimator for σ^2 by lemma 2.3, but the square root of the sample variance is not an unbiased estimator for σ : The function $g(x) = -\sqrt{x}$ is strictly convex on $(0, \infty)$ and by Jensen's inequality, lemma A.7, we have $-\sqrt{\mathbb{E}(S^2)} \leq \mathbb{E}(-\sqrt{S^2})$, i.e. $\mathbb{E}(S) \leq \sqrt{\mathbb{E}(S^2)} = \sigma$. The inequality is strict, since S^2 is not constant with probability one (by proposition A.3 it is a multiple of a chi-squared random variable). Thus, $\mathbb{E}(S) < \sigma$, the estimator S is biased for σ and the Cramer–Rao bound, which only concerns unbiased estimators, cannot be compared to $\text{Var}(S)$ directly.

□

Solution to exercise 14.2.

1. We have $\ell_1(\alpha, \beta) = \alpha \log \beta - \log \Gamma(\alpha) + (\alpha - 1) \log X - \beta X$ and

$$\frac{\partial \ell_1}{\partial \alpha} = \log \beta - \psi(\alpha) + \log X \quad \text{and} \quad \frac{\partial \ell_1}{\partial \beta} = \frac{\alpha}{\beta} - X,$$

where $\psi = (\log \Gamma)'$ is the derivative of $\log \Gamma$, which is smooth and thus differentiable on $(0, \infty)$. Taking derivatives again, we get $\partial^2 \ell_1 / \partial \alpha^2 = -\psi'(\alpha)$, $\partial^2 \ell_1 / \partial \alpha \partial \beta = 1/\beta$ and $\partial^2 \ell_1 / \partial \beta^2 = -\alpha/\beta^2$, as given in section 14.5. To verify the first identity of proposition 14.3, we note that the second component of the score has mean $\alpha/\beta - \mathbb{E}(X) = 0$, since $\mathbb{E}(X) = \alpha/\beta$ by table A.2.

2. If $\alpha = 2$ is known, the model consists of one parameter β and by theorem 11.3 the Fisher information is $\mathcal{I}(\beta) = -\mathbb{E}(\partial^2 \ell_1 / \partial \beta^2) = \alpha/\beta^2 = 2/\beta^2$. The lower right corner of $\mathcal{I}(\alpha, \beta)$ is the diagonal element $\mathcal{I}(\alpha, \beta)_{22}$, which by definition is the variance of $\partial \ell_1 / \partial \beta$ with α fixed to the true value, i.e. for known shape. The diagonal element of the inverse answers a different question, asking about the estimation of β when α is unknown. Part (d) shows that this value is larger.

3. The scale is $\theta = g(\beta) = 1/\beta$ where $g'(\beta) = -1/\beta^2$ and by proposition 11.10 we have

$$\mathcal{I}_\theta(\theta) = \frac{\mathcal{I}_\beta(\beta)}{g'(\beta)^2} = \frac{2/\beta^2}{1/\beta^4} = 2\beta^2 = \frac{2}{\theta^2}.$$

This is the same value as in exercise 11.2.

4. We have $a = \psi'(\alpha)$, $b = -1/\beta$ and $c = \alpha/\beta^2$ and thus the determinant of $\mathcal{I}(\alpha, \beta)$ is $ac - b^2 = (\alpha\psi'(\alpha) - 1)/\beta^2$. The lower right element of the inverse is given by

$$(\mathcal{I}(\alpha, \beta)^{-1})_{22} = \frac{a}{ac - b^2} = \frac{\psi'(\alpha)\beta^2}{\alpha\psi'(\alpha) - 1}.$$

For a sample of size n , the information matrix is $n\mathcal{I}(\alpha, \beta)$, the inverse is $\mathcal{I}(\alpha, \beta)^{-1}/n$ and by theorem 14.5 we find the bound $\beta^2\psi'(\alpha)/(n(\alpha\psi'(\alpha) - 1))$ for an unbiased estimator for β as claimed. The bound for known shape is $1/\mathcal{I}_n(\beta) = \beta^2/(n\alpha)$ and the ratio of the two is

$$\frac{\beta^2\psi'(\alpha)/(n(\alpha\psi'(\alpha) - 1))}{\beta^2/(n\alpha)} = \frac{\alpha\psi'(\alpha)}{\alpha\psi'(\alpha) - 1} = \frac{1}{1 - \rho^2},$$

where $\rho^2 = 1/(\alpha\psi'(\alpha))$ is as in (14.2). For $\alpha = 2$ we have $\psi'(2) = \pi^2/6 - 1 \approx 0.6449$ and thus $\alpha\psi'(\alpha) \approx 1.2899$ and the ratio is approximately $1.2899/0.2899 \approx 4.45$: not knowing the shape increases the bound for the rate by a factor of approximately 4.45, as in section 14.5.

□

Solution to exercise 14.3.

1. Since the two samples are independent, the log-likelihood is the sum of the two sample log-likelihoods:

$$\ell(\mu_1, \mu_2) = -\frac{m+n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^m (X_i - \mu_1)^2 - \frac{1}{2\sigma^2} \sum_{j=1}^n (Y_j - \mu_2)^2,$$

with partial derivatives $\partial\ell/\partial\mu_1 = \sum_i (X_i - \mu_1)/\sigma^2$ and $\partial\ell/\partial\mu_2 = \sum_j (Y_j - \mu_2)/\sigma^2$. The second partial derivatives are $\partial^2\ell/\partial\mu_1^2 = -m/\sigma^2$, $\partial^2\ell/\partial\mu_2^2 = -n/\sigma^2$ and $\partial^2\ell/\partial\mu_1\partial\mu_2 = 0$, since each first partial derivative involves only one of the two parameters. None of these second derivatives are random, and thus proposition 14.3 gives $\mathcal{I}(\mu_1, \mu_2) = \text{diag}(m/\sigma^2, n/\sigma^2)$, without any expectations.

2. The difference is $\mu_1 - \mu_2 = a^\top \theta$ with $a = (1, -1)$ and the inverse information matrix is $\text{diag}(\sigma^2/m, \sigma^2/n)$. By theorem 14.5 any unbiased estimator for $\mu_1 - \mu_2$ has variance at least

$$a^\top \mathcal{I}(\mu_1, \mu_2)^{-1} a = \frac{\sigma^2}{m} + \frac{\sigma^2}{n}.$$

The estimator $\bar{X} - \bar{Y}$ is unbiased for $\mu_1 - \mu_2$ by lemma 2.2 and, since the two samples are independent, the variance of this estimator is $\text{Var}(\bar{X}) + \text{Var}(\bar{Y}) = \sigma^2/m + \sigma^2/n$. This is the bound. Thus, $\bar{X} - \bar{Y}$ is an efficient estimator for the difference of the means.

3. If we include σ^2 as an additional parameter, the first derivative is

$$\frac{\partial \ell}{\partial \sigma^2} = -\frac{m+n}{2\sigma^2} + \frac{1}{2\sigma^4} \left(\sum_{i=1}^m (X_i - \mu_1)^2 + \sum_{j=1}^n (Y_j - \mu_2)^2 \right).$$

The mixed derivatives with respect to μ_1 and σ^2 and with respect to μ_2 and σ^2 are $-\sum_i (X_i - \mu_1)/\sigma^4$ and $-\sum_j (Y_j - \mu_2)/\sigma^4$, both with expectation zero, and the second derivative with respect to σ^2 has expectation $(m+n)/(2\sigma^4) - (m+n)\sigma^2/\sigma^6 = -(m+n)/(2\sigma^4)$, as shown in example 14.6. Thus, the information matrix is

$$\mathcal{I}(\mu_1, \mu_2, \sigma^2) = \text{diag}\left(\frac{m}{\sigma^2}, \frac{n}{\sigma^2}, \frac{m+n}{2\sigma^4}\right),$$

still diagonal, and thus σ^2 is orthogonal to both means. The upper left 2×2 block of the inverse is the inverse information matrix from part (b), and the Cramer–Rao bound for $\mu_1 - \mu_2$ is still $\sigma^2/m + \sigma^2/n$, and this bound is achieved by $\bar{X} - \bar{Y}$, whose definition does not depend on σ^2 . For point estimation of the difference, not knowing the variance does not change anything. However, once we want to determine the accuracy of an estimate, the unknown variance starts to matter, since the bound depends on σ^2 and thus needs to be estimated. This is the reason why we go from the normal distribution to the t distribution in lectures 25 and 29.

□

C.11 Consistency of the MLE

Solution to exercise 16.1.

1. For (R1) we have two different rates and thus two different distributions: for example, the means $1/\theta$ are different. For (R2) the support is $[0, \infty)$ for all θ . For (R3) we have $\log f(x; \theta) = \log \theta - \theta x$ and the derivatives w.r.t. θ are $1/\theta - x$, $-1/\theta^2$ and $2/\theta^3$, all of which are continuous at θ . The derivatives of f are of the form (polynomial in x and θ) times $e^{-\theta x}$ and thus are integrable over $[0, \infty)$ uniformly for θ in any compact subinterval of $(0, \infty)$. We can therefore interchange the order of differentiation and integration, and $\mathcal{I}(\theta) = 1/\theta^2$ is continuous and lies strictly between zero and infinity, so that (R3) holds. Finally, for (R4), the third derivative $2/\theta^3$ does not depend on x and for $|\theta - \theta_0| < \delta$ with $\delta = \theta_0/2$ we have $\theta > \theta_0/2$ and thus $2/\theta^3 \leq 16/\theta_0^3$. The constant $M(x) = 16/\theta_0^3$ bounds the third derivative and has finite expectation.
2. The likelihood equation $n/\theta - \sum_i x_i = 0$ has only one solution $\theta = 1/\bar{x}$, since $\sum_i x_i > 0$ with probability one. Since $\ell''(\theta) = -n/\theta^2 < 0$, the log-likelihood function is strictly concave and thus the root is the global maximum of the likelihood function, *i.e.* the MLE as in example 5.2. Since all conditions of theorem 16.3 are satisfied, we find that $1/\bar{X}$ is consistent for θ .
3. By the law of large numbers, theorem A.14, we have $\bar{X} \xrightarrow{P} \mathbb{E}_\theta(X_1) = 1/\theta$. Since the function $g(y) = 1/y$ is continuous at $y = 1/\theta > 0$, the continuous

mapping theorem, theorem A.17, applies and we find $1/\bar{X} \xrightarrow{P} g(1/\theta) = \theta$. This proves the consistency of the MLE.

□

Solution to exercise 16.2.

1. Since $f(x; \theta) = \theta^x(1 - \theta)^{1-x}$ for $x \in \{0, 1\}$, we have

$$\log \frac{f(X_1; \theta)}{f(X_1; \theta_0)} = X_1 \log \frac{\theta}{\theta_0} + (1 - X_1) \log \frac{1 - \theta}{1 - \theta_0},$$

and taking expectations with $\mathbb{E}_{\theta_0}(X_1) = \theta_0$ gives

$$D(\theta) = \theta_0 \log \frac{\theta}{\theta_0} + (1 - \theta_0) \log \frac{1 - \theta}{1 - \theta_0}.$$

2. Using the rule $\log y \leq y - 1$ for logarithms, we find

$$D(\theta) \leq \theta_0 \left(\frac{\theta}{\theta_0} - 1 \right) + (1 - \theta_0) \left(\frac{1 - \theta}{1 - \theta_0} - 1 \right) = (\theta - \theta_0) + (\theta_0 - \theta) = 0.$$

For $D(\theta) = 0$ to hold, we need to have equality in both applications of the inequality, *i.e.* we need $\theta/\theta_0 = 1$ and $(1 - \theta)/(1 - \theta_0) = 1$. Both of these conditions state $\theta = \theta_0$. Thus we have $D(\theta) < 0$ for all $\theta \neq \theta_0$, as stated in theorem 16.2.

□

Solution to exercise 16.3.

1. Since all observations are at most θ , we have $\hat{\theta}_n \leq \theta$ and thus $|\hat{\theta}_n - \theta| > \varepsilon$ is the same as $\hat{\theta}_n < \theta - \varepsilon$. This is the case if and only if all n observations are less than $\theta - \varepsilon$. Using independence we find

$$\mathbb{P}_\theta(|\hat{\theta}_n - \theta| > \varepsilon) = \left(\frac{\theta - \varepsilon}{\theta} \right)^n = \left(1 - \frac{\varepsilon}{\theta} \right)^n \longrightarrow 0 \quad (n \rightarrow \infty),$$

since $0 < 1 - \varepsilon/\theta < 1$. For $\varepsilon \geq \theta$, the probability of this event is zero. Thus $\hat{\theta}_n \xrightarrow{P} \theta$ and the MLE is consistent.

2. Condition (R2) is violated, since the support $(0, \theta)$ depends on θ and (R3) is also violated: for fixed x , the function $\theta \mapsto f(x; \theta)$ has a jump from 0 to $1/\theta$ at $\theta = x$ and thus is not differentiable there. The uniqueness of the MLE fails, since the likelihood equation has no roots: on the region $\theta \geq \max_i x_i$ where L is positive, the derivative $-n\theta^{-n-1}$ never equals zero. Since the proof of theorem 16.3 requires a differentiable likelihood and a root of the derivative to apply Rolle's theorem, neither of these conditions are satisfied. Nevertheless, the statement is true, as can be seen by the direct computation in the first part of the question, or using the mean squared error from exercise 2.2.

□

C.12 Asymptotic Normality of the MLE

Solution to exercise 17.1.

1. Since we have $\mathcal{I}(\theta) = 1/(\theta(1-\theta))$, theorem 17.1 gives that $\sqrt{n}(\bar{X} - \theta) \xrightarrow{d} N(0, \theta(1-\theta))$. Since the mean and variance of the Bernoulli distribution equal θ and $\theta(1-\theta)$, respectively, this is the central limit theorem, theorem A.15, for $\bar{X}$ and thus the two results coincide.
2. From theorem 5.6 we know that the MLE for ψ is given by $\hat{\psi}_n = \log(\bar{X}/(1-\bar{X}))$. Taking derivatives we find the derivative of $g(\theta) = \log \theta - \log(1-\theta)$ to be $g'(\theta) = 1/\theta + 1/(1-\theta) = 1/(\theta(1-\theta))$ and using the delta method, theorem A.18, we find

$$\sqrt{n}(\hat{\psi}_n - \psi) \xrightarrow{d} N\left(0, \frac{\theta(1-\theta)}{\theta^2(1-\theta)^2}\right) = N\left(0, \frac{1}{\theta(1-\theta)}\right).$$

3. Substituting θ by $\hat{\theta}_n$ in the asymptotic variance, we find the standard error $\text{se}(\hat{\psi}_n) = 1/\sqrt{n\hat{\theta}_n(1-\hat{\theta}_n)}$ and thus the approximate 95% interval $\hat{\psi}_n \pm 1.96/\sqrt{n\hat{\theta}_n(1-\hat{\theta}_n)}$. Both of these quantities are undefined if all observations are equal, *i.e.* at the boundary of the domain discussed in section 17.5.

□

Solution to exercise 17.2.

1. Using the result $\mathcal{I}(\theta) = 1/\theta$ from the theorem, we find $\sqrt{n}(\bar{X} - \theta) \xrightarrow{d} N(0, \theta)$. Since the Poisson distribution has mean and variance equal to θ , this is the central limit theorem for $\bar{X}$.
2. For $g(\theta) = e^{-\theta}$ we have $g'(\theta) = -e^{-\theta}$ and using the delta method we find

$$\sqrt{n}(\hat{p}_0 - p_0) \xrightarrow{d} N(0, e^{-2\theta}).$$

In exercise 5.7 we found $\mathbb{E}_\theta(\hat{p}_0) = \exp(n\theta(e^{-1/n} - 1))$ and expanding $e^{-1/n} = 1 - 1/n + 1/(2n^2) - \dots$ we find that the bias is approximately $\theta e^{-\theta}/(2n)$, which is of order $1/n$. Since it is multiplied by $\sqrt{n}$, the bias goes to zero and thus is of a smaller order than the spread $1/\sqrt{n}$ and does not contribute in the limit: the estimator is biased for finite n , but is asymptotically normal with mean p_0 .

□

Solution to exercise 17.3.

1. For $x \geq 0$ and $n > x/\theta$, the event $n(\theta - \hat{\theta}_n) > x$ is the event $\hat{\theta}_n < \theta - x/n$, *i.e.* that all observations are less than $\theta - x/n$. Thus we have

$$\mathbb{P}_\theta(n(\theta - \hat{\theta}_n) > x) = \left(1 - \frac{x}{n\theta}\right)^n \longrightarrow e^{-x/\theta} \quad (n \rightarrow \infty),$$

where we used the standard limit $(1 - a/n)^n \rightarrow e^{-a}$. The right-hand side is the probability that an exponential random variable with mean θ is greater than x . Thus, $n(\theta - \hat{\theta}_n)$ converges in distribution to the corresponding exponential distribution.

2. Since $\sqrt{n}(\hat{\theta}_n - \theta) = -n^{-1/2} \cdot n(\theta - \hat{\theta}_n)$, the second factor converges in distribution to the result from the first part, and the first factor is a constant which converges to zero. Thus, by Slutsky's theorem (theorem A.16), the product converges in distribution to $0 \cdot Y = 0$. This is equivalent to convergence in probability to zero. Thus we cannot have a limit $N(0, \tau^2)$ where $\tau^2 > 0$, since the error of the MLE is of order $1/n$ and not of order $1/\sqrt{n}$. This is a manifestation of the super-efficiency from lecture 13, seen through the lens of the limiting distribution: the variance of the bias-corrected maximum is of order $1/n^2$, smaller than allowed by the Cramer-Rao bound, and thus neither the bound nor theorem 17.1 apply, since the model violates (R2).

□

C.13 The EM Algorithm

Solution to exercise 19.1.

1. Using the formula for the sum of a geometric series, we find

$$\begin{aligned}\mathbb{P}_p(X \geq m) &= \sum_{x=m}^{\infty} (1-p)^{x-1} p = p(1-p)^{m-1} \sum_{j=0}^{\infty} (1-p)^j \\ &= p(1-p)^{m-1} \cdot \frac{1}{p} = (1-p)^{m-1}.\end{aligned}$$

Since the event $\{X \geq m+k\}$ is contained in the event $\{X \geq m\}$, we can use the definition of conditional probability to get

$$\mathbb{P}_p(X \geq m+k | X \geq m) = \frac{(1-p)^{m+k-1}}{(1-p)^{m-1}} = (1-p)^k = \mathbb{P}_p(X \geq k+1).$$

This shows that, given that the first $m-1$ trials were failures, the number $X - (m-1)$ of additional trials required until the first success is still geometrically distributed. The information about the number of failures so far does not change the distribution of the number of additional failures required. This is the discrete analogue of the statement $\mathbb{P}_\theta(X > s+t | X > s) = \mathbb{P}_\theta(X > t)$ for the exponential distribution in exercise 5.6. The shift by one is due to the fact that the support of the geometric distribution starts at 1 instead of at 0.

2. Given $X \geq m$, the conditional probability weights of X are

$$\frac{f(x; p)}{\mathbb{P}_p(X \geq m)} = (1-p)^{x-m} p$$

for $x \geq m$. Substituting $x = m-1+j$ with $j \in \{1, 2, \dots\}$ we get

$$\begin{aligned}\mathbb{E}_p(X | X \geq m) &= \sum_{j=1}^{\infty} (m-1+j)(1-p)^{j-1} p \\ &= (m-1) \sum_{j=1}^{\infty} (1-p)^{j-1} p + \sum_{j=1}^{\infty} j(1-p)^{j-1} p \\ &= (m-1) + \mathbb{E}_p(X),\end{aligned}$$

where the first sum is the total probability 1 and the second sum is the mean of the geometric distribution. Using $\mathbb{E}_p(X) = 1/p$ we get the claim. Alternatively we can get the same result from the first part of the solution: given $X \geq m$, the value $X - (m - 1)$ is geometrically distributed with parameter p and thus has mean $1/p$.

□

Solution to exercise 19.2.

1. The log-likelihood function for the given data is

$$\ell(\pi) = \sum_{i=1}^n \log(\pi f_1(x_i) + (1 - \pi)f_2(x_i)).$$

Taking derivatives we get the likelihood equation

$$\ell'(\pi) = \sum_{i=1}^n \frac{f_1(x_i) - f_2(x_i)}{\pi f_1(x_i) + (1 - \pi)f_2(x_i)} = 0.$$

This equation cannot be solved analytically for general densities.

2. The density of the complete data is $f_c(x, 1; \pi) = \pi f_1(x)$ and $f_c(x, 2; \pi) = (1 - \pi)f_2(x)$. Thus we have

$$\ell_c(\pi) = \sum_{i=1}^n \left[\mathbf{1}_{\{Z_i=1\}} (\log \pi + \log f_1(x_i)) + \mathbf{1}_{\{Z_i=2\}} (\log(1 - \pi) + \log f_2(x_i)) \right].$$

Using Bayes' rule we find the conditional probability of $Z_i = 1$ given $X_i = x_i$ under π_k as $\gamma_i(\pi_k) = \pi_k f_1(x_i) / (\pi_k f_1(x_i) + (1 - \pi_k)f_2(x_i))$. Substituting the indicators using these responsibilities we get

$$Q(\pi | \pi_k) = \log \pi \sum_{i=1}^n \gamma_i(\pi_k) + \log(1 - \pi) \sum_{i=1}^n (1 - \gamma_i(\pi_k)) + \text{const},$$

where the constant term combines the terms $\log f_1(x_i)$ and $\log f_2(x_i)$, which do not depend on π . Setting the derivative $\sum_i \gamma_i / \pi - \sum_i (1 - \gamma_i) / (1 - \pi)$ equal to zero we find $\pi_{k+1} = \frac{1}{n} \sum_i \gamma_i(\pi_k)$ and the second derivative $-\sum_i \gamma_i / \pi^2 - \sum_i (1 - \gamma_i) / (1 - \pi)^2$ is negative for every $\pi \in (0, 1)$, so $Q(\cdot | \pi_k)$ is concave on $(0, 1)$ and the stationary point is the global maximum.

3. Using the notation $f_i = f(x_i; \pi) = \pi f_1(x_i) + (1 - \pi)f_2(x_i)$ for the denominator, we have

$$\begin{aligned} \gamma_i(\pi) - \pi &= \frac{\pi f_1(x_i) - \pi(\pi f_1(x_i) + (1 - \pi)f_2(x_i))}{f_i} \\ &= \frac{\pi(1 - \pi)(f_1(x_i) - f_2(x_i))}{f_i}. \end{aligned}$$

Summing over i and comparing to the first part of the solution we find $\frac{1}{n} \sum_i \gamma_i(\pi) - \pi = \pi(1 - \pi) \ell'(\pi) / n$. A value $\pi^* \in (0, 1)$ is a fixed point of the iteration if and only if the left-hand side is zero at π^* and since $\pi^*(1 - \pi^*) > 0$ this is the case if and only if $\ell'(\pi^*) = 0$. Thus the fixed points of EM are the roots of the likelihood equation as given in proposition 19.3.

□

Solution to exercise 19.3. At $\theta_0 = (1/2, 2, 4)$, the mixture density of observation i is $\frac{1}{2}\varphi_1(x_i - 2) + \frac{1}{2}\varphi_1(x_i - 4)$. Using the standard normal density values $\varphi_1(1) = 0.2420$, $\varphi_1(2) = 0.05399$, $\varphi_1(3) = 0.004432$, $\varphi_1(4) = 0.0001338$ and $\varphi_1(5) = 0.0000015$, we find the five densities to be 0.027062, 0.12320, 0.12320, 0.027062 and 0.0022167. The corresponding logarithms are

$$-3.6096, \quad -2.0939, \quad -2.0939, \quad -3.6096, \quad -6.1118$$

and thus $\ell(\theta_0) = -17.5188$. At $\theta_1 = (0.4001, 0.5445, 5.9709)$ the densities are 0.13763, 0.14389, 0.14939, 0.23922 and 0.14094. The corresponding logarithms are

$$-1.9832, \quad -1.9387, \quad -1.9012, \quad -1.4304, \quad -1.9595$$

and thus $\ell(\theta_1) = -9.2129$. The log-likelihood has increased by 8.31 as required by theorem 19.2. The largest contribution, 4.15, comes from observation $x_5 = 7$: under θ_0 this observation is three standard deviations away from the nearer mean 4 and has tiny density, whereas under θ_1 it is about one standard deviation away from the mean 5.97. The observations $x_1 = 0$ and $x_4 = 6$ contribute 1.63 and 2.18, respectively, for the same reason. The observations $x_2 = 1$ and $x_3 = 5$, which were always within one standard deviation of a mean, contribute little. □

Solution to exercise 19.4.

1. Since we know π , the complete-data log-likelihood is the same as in example 19.4 and the terms $\log \pi$ and $\log(1 - \pi)$ are now constants. The E step determines the responsibilities $\gamma_i = \pi\varphi_\sigma(x_i - \mu_{1,k}) / (\pi\varphi_\sigma(x_i - \mu_{1,k}) + (1 - \pi)\varphi_\sigma(x_i - \mu_{2,k}))$. Note that $Q(\theta|\theta_k)$ depends on $\theta = (\mu_1, \mu_2)$ only via the terms $-\frac{1}{2\sigma^2} \sum_i [\gamma_i(x_i - \mu_1)^2 + (1 - \gamma_i)(x_i - \mu_2)^2]$. The partial derivatives are as shown in the example, and setting them equal to zero gives the two mean updates of (19.2). The update for the weight is no longer needed.
2. Since $\varphi'_\sigma(u) = -(u/\sigma^2)\varphi_\sigma(u)$, we have $\frac{\partial}{\partial \mu_1}\varphi_\sigma(x - \mu_1) = \frac{x - \mu_1}{\sigma^2}\varphi_\sigma(x - \mu_1)$ and thus we can take derivatives of $\ell(\theta) = \sum_i \log(\pi\varphi_\sigma(x_i - \mu_1) + (1 - \pi)\varphi_\sigma(x_i - \mu_2))$ to get

$$\begin{aligned} \frac{\partial \ell}{\partial \mu_1}(\theta) &= \sum_{i=1}^n \frac{\pi\varphi_\sigma(x_i - \mu_1)}{\pi\varphi_\sigma(x_i - \mu_1) + (1 - \pi)\varphi_\sigma(x_i - \mu_2)} \cdot \frac{x_i - \mu_1}{\sigma^2} \\ &= \frac{1}{\sigma^2} \sum_{i=1}^n \gamma_i(\theta) (x_i - \mu_1), \end{aligned}$$

and similarly we find $\frac{\partial \ell}{\partial \mu_2}(\theta) = \frac{1}{\sigma^2} \sum_i (1 - \gamma_i(\theta))(x_i - \mu_2)$. Setting both of these equal to zero gives $\mu_1 = \sum_i \gamma_i(\theta)x_i / \sum_i \gamma_i(\theta)$ and the corresponding equation for μ_2 , stating that θ is mapped to itself by the updates of the first part. Thus, the likelihood equations and the fixed-point equations of the EM algorithm coincide as predicted by proposition 19.3, and we also find the identity of this proposition directly: the derivative of ℓ equals the derivative of $\theta' \mapsto Q(\theta'|\theta)$ at $\theta' = \theta$.

3. As $\pi \rightarrow 1$, the responsibilities converge to $\gamma_i = 1$ for all observations since $\varphi_\sigma(x_i - \mu_{1,k}) > 0$. Thus, the update for μ_1 is now $\mu_{1,k+1} = \bar{x}$ independent of the initial value, and this is the MLE for the mean of a single normal sample in example 5.4, reached in one step. The update for μ_2 is now the undefined ratio $0/0$: a component with no weight has no observations to use to estimate its mean, and the mean of this component is not identifiable as discussed in section 19.5.

□

C.14 Hypothesis Testing

Solution to exercise 20.1.

1. By equation (20.2) the power at θ_1 is $\beta(\theta_1) = \Phi(d - z_{1-\alpha})$, where $d = \sqrt{n}(\theta_1 - \theta_0)/\sigma$. Since Φ is strictly increasing with $\Phi(z_{1-\gamma}) = 1 - \gamma$, we have $\beta(\theta_1) \geq 1 - \gamma$ if and only if $d - z_{1-\alpha} \geq z_{1-\gamma}$, i.e. if and only if

$$\sqrt{n} \geq \frac{(z_{1-\alpha} + z_{1-\gamma}) \sigma}{\theta_1 - \theta_0}.$$

Since $1 - \gamma > \alpha$ we have $z_{1-\gamma} > z_\alpha = -z_{1-\alpha}$, and thus the right-hand side is positive; squaring both sides gives the stated condition.

2. With $z_{0.95} = 1.645$ and $z_{0.9} = 1.282$ the bound is

$$n \geq \left(\frac{(1.645 + 1.282) \cdot 2}{0.5} \right)^2 = 11.708^2 = 137.1,$$

and thus $n = 138$ observations are needed.

3. The bound is proportional to $\sigma^2/(\theta_1 - \theta_0)^2$. Halving the difference $\theta_1 - \theta_0$ multiplies the required sample size by four, to $n = 549$ in the numerical example (the bound becomes 548.3), and doubling σ has the same effect. Detecting a difference half as large costs four times as many observations.

□

Solution to exercise 20.2.

1. Under H_0 the total T is binomially distributed with parameters 10 and $1/2$. Thus, the test size is

$$\beta(1/2) = \mathbb{P}_{1/2}(T \geq c) = \frac{1}{2^{10}} \sum_{k=c}^{10} \binom{10}{k}.$$

2. We can compute the binomial coefficients $\binom{10}{10} = 1$, $\binom{10}{9} = 10$ and $\binom{10}{8} = 45$ to find the sizes $1/1024 = 0.0010$ for $c = 10$, $11/1024 = 0.0107$ for $c = 9$ and $56/1024 = 0.0547$ for $c = 8$. The smallest value c with size less than or equal to 0.05 is $c = 9$ with size 0.0107. Since T can only take the values $0, \dots, 10$, the size can only take the eleven values computed in the first part of the question, and 0.05 is not one of these values (the size jumps from 0.0107 to 0.0547 between $c = 9$ and $c = 8$).

3. The power at p is $\mathbb{P}_p(T \geq 9) = 10p^9(1-p) + p^{10} = p^9(10-9p)$. This gives the values 0.149 for $p = 0.7$, 0.376 for $p = 0.8$ and 0.736 for $p = 0.9$. A coin which shows heads four times in five experiments is detected in fewer than 40% of the experiments; for ten experiments the test can only detect large deviations from fairness with reasonable probability, and a failure to reject the hypothesis is not very meaningful.
4. The p -value is the probability of the event that the total under H_0 is larger than or equal to the observed value, *i.e.* the probability of $\mathbb{P}_{1/2}(T \geq 9) = 11/1024 = 0.0107$. Since this probability is less than or equal to 0.05, we can reject H_0 at a significance level of 0.05. This agrees with the critical region $T \geq 9$ from the second part of the question.

□

Solution to exercise 20.3.

1. The maximum of the distribution is at most m , if and only if all observations are at most m . Since the observations are independent, for $0 \leq m \leq \theta$ we have

$$\mathbb{P}_\theta(M \leq m) = \prod_{i=1}^n \mathbb{P}_\theta(X_i \leq m) = \left(\frac{m}{\theta}\right)^n,$$

This is also the integral of the density from example 8.10.

2. For $c \leq \theta_0$ the size is $\mathbb{P}_{\theta_0}(M > c) = 1 - (c/\theta_0)^n$, as found in the first part. Setting this equal to α gives $c = \theta_0(1 - \alpha)^{1/n}$.
3. For $\theta > \theta_0 \geq c$ the power function is

$$\beta(\theta) = \mathbb{P}_\theta(M > c) = 1 - \left(\frac{c}{\theta}\right)^n = 1 - (1 - \alpha)\left(\frac{\theta_0}{\theta}\right)^n,$$

This function increases from α for $\theta = \theta_0$ to 1 as $\theta \rightarrow \infty$. For $n = 5$, $\theta_0 = 1$ and $\alpha = 0.05$ we have $c = 0.95^{1/5} = 0.990$, and for $\theta = 1.2$ the power is $1 - 0.95/1.2^5 = 1 - 0.95 \cdot 0.402 = 0.618$.

4. Under H_0 all observations are in the interval $[0, \theta_0]$ and thus we have $\mathbb{P}_{\theta_0}(M > \theta_0) = 0$: The second test has size 0 and is a level α test for every α . The power function of this test is $\mathbb{P}_\theta(M > \theta_0) = 1 - (\theta_0/\theta)^n$ for $\theta > \theta_0$. This function equals 0.598 for $\theta = 1.2$ in the numerical example. The difference between the two power functions is $\alpha(\theta_0/\theta)^n$, and this is positive for every alternative, *i.e.* the first test is more powerful everywhere on H_1 , but the first test has type I error probability α which the second test avoids. In the context of section 20.3, the first test is to be preferred: we have agreed on a level α , and a test which does not use the level would be a waste of power. The power gain is at most α , and decreases quickly as θ increases, whereas for the second test a rejection is certain if the observation is above θ_0 , which is impossible under H_0 , and so we can reasonably assume that this is worth the small loss of power. The example shows that the choice of level is a matter for the statistician, and is not a mathematical necessity.

□

Solution to exercise 20.4.

1. The observed value of the standardised statistic is $z = \sqrt{16} (52.3 - 50)/4 = 2.3$ and the one-sided p -value is $p = 1 - \Phi(2.3) = 0.0107$. Since $0.0107 \leq 0.05$ we can reject H_0 at level 0.05, but since $0.0107 > 0.01$ we cannot reject at level 0.01.
2. The two-sided p -value is $2(1 - \Phi(2.3)) = 0.0214$. Again we reject at level 0.05 and we cannot reject at level 0.01.
3. The mean θ is a fixed, unknown number and not a random variable, so we cannot determine a probability for the statement $\theta = 50$. The p -value is a probability about the data, assuming that $\theta = 50$. A correct interpretation is: if the true mean were 50, then a sample mean of at least 52.3 would occur with probability 0.011, and thus the data are significant at level 0.05 but not at level 0.01.
4. Under H_0 the statistic Z is standard normally distributed. Since Φ is continuous and strictly increasing, for $u \in (0, 1)$ we have

$$\begin{aligned}\mathbb{P}_{\theta_0}(1 - \Phi(Z) \leq u) &= \mathbb{P}_{\theta_0}(\Phi(Z) \geq 1 - u) \\ &= \mathbb{P}_{\theta_0}(Z \geq z_{1-u}) = 1 - \Phi(z_{1-u}) = u,\end{aligned}$$

and thus the p -value is uniformly distributed on $(0, 1)$ under H_0 . In particular we have $\mathbb{P}_{\theta_0}(p(X) \leq \alpha) = \alpha$, which is the statement that the test has size α .

□

C.15 The Neyman–Pearson Lemma

Solution to exercise 22.1.

1. The joint probability weights are $f(x; \theta) = \theta^s e^{-n\theta} / \prod_i x_i!$ with $s = \sum_i x_i$. Thus the likelihood ratio is

$$R(x) = \left(\frac{\theta_1}{\theta_0}\right)^s e^{-n(\theta_1 - \theta_0)}.$$

Since $\theta_1/\theta_0 > 1$, this function is strictly increasing as a function of s and the inequality $R(x) > k$ is equivalent to $s > c$ with $c = (\log k + n(\theta_1 - \theta_0))/\log(\theta_1/\theta_0)$. By theorem 22.2 the most powerful test at this level is the one which rejects when $S > c$. The size of this test, $\mathbb{P}_{\theta_0}(S > c)$, can be found using the Poisson distribution with mean $n\theta_0$, which contains the parameters n and θ_0 only. For a given size the threshold c is the same for all $\theta_1 > \theta_0$. By proposition 22.4 the test is uniformly most powerful against $H_1: \theta > \theta_0$.

2. Under H_0 , the sum S is Poisson-distributed with mean 5 and thus we have $\mathbb{P}_{\theta_0}(S > 8) = 1 - 0.9319 = 0.0681 > 0.05$, whereas $\mathbb{P}_{\theta_0}(S > 9) = 1 - 0.9682 = 0.0318 \leq 0.05$. The smallest acceptable threshold for the test is $c = 9$: we can reject $S \geq 10$ and the size of this test is 0.0318.

3. Under $\theta_1 = 2$, the sum is Poisson-distributed with mean 10 and thus the power of the test is $\mathbb{P}_{\theta_1}(S > 9) = 1 - 0.4579 = 0.5421$.
4. The size of the test $\{S > c\}$ is $1 - \mathbb{P}_{\theta_0}(S \leq c)$. This value can only take the values listed for $c = 0, 1, 2, \dots$ and the jumps between these values are from $c = 8$ to $c = 9$, *i.e.* 0.0681 to 0.0318. The value 0.05 is not one of the possible values. The test from part (b) has size 0.0318 and the lemma shows that this is the most powerful test at level 0.0318. The test is also a test of level 0.05, but the lemma does not show that this is the most powerful test among all tests with level 0.05. As the remark in section 22.2 points out, a randomised test could have size exactly 0.05 and higher power.

□

Solution to exercise 22.2.

1. The joint probability weights are $f(x; p) = p^s(1-p)^{n-s}$ with $s = \sum_i x_i$. Under H_0 , these weights equal 2^{-n} for all x . Thus we have

$$R(x) = \frac{p_1^s(1-p_1)^{n-s}}{2^{-n}} = \left(\frac{p_1}{1-p_1}\right)^s (2(1-p_1))^n,$$

as claimed. For $p_1 > 1/2$, we have $p_1/(1-p_1) > 1$ and thus $R(x)$ is strictly increasing in s . The inequality $R(x) > k$ is equivalent to $s > c$ for some c . By theorem 22.2, the most powerful test is to reject when $S > c$, where c is determined by the distribution of S under H_0 , *i.e.* the binomial distribution with parameters n and $1/2$. Since this distribution does not depend on p_1 , the same region of $p_1 > 1/2$ is optimal for all alternatives and by proposition 22.4 the test from exercise 20.2 is uniformly most powerful against $H_1: p > 1/2$.

2. Since $\mathbb{P}_{1/2}(S \geq 14) = 0.0577 > 0.05$ and $\mathbb{P}_{1/2}(S \geq 15) = 0.0207 \leq 0.05$, the largest critical region which does not exceed size 0.05 is $\{S \geq 15\}$. This region has size 0.0207. The power of this test at $p_1 = 0.7$ is $\mathbb{P}_{0.7}(S \geq 15) = 1 - 0.5836 = 0.4164$. For $n = 10$, the corresponding region $\{S \geq 9\}$ from exercise 20.2 has smaller size $11/1024 = 0.0107$ but the power at $p_1 = 0.7$ is $10 \cdot 0.7^9 \cdot 0.3 + 0.7^{10} = 0.1493$. By doubling the sample size, we can use a test with size closer to the nominal value 0.05 and nearly triple the power at this alternative.
3. For $p_1 < 1/2$, we have $p_1/(1-p_1) < 1$ and thus $R(x)$ is strictly decreasing in s . The inequality $R(x) > k$ is now equivalent to $s < c$ and the most powerful test will be a test which rejects for small values of S . For $n = 20$, with the same size 0.0207, the test will reject for $S \leq 5$, by the symmetry of the binomial distribution with parameter $1/2$. Thus, the two one-sided problems will result in tests which reject on the opposite tails of the distribution of S and there can be no single critical region which satisfies both tests.

□

Solution to exercise 22.3.

1. For $x_1, \dots, x_n \geq 0$ we have

$$f(x; \theta_0) = \theta_0^{-n} \mathbf{1}_{\{m \leq \theta_0\}} \quad \text{and} \quad f(x; \theta_1) = \theta_1^{-n} \mathbf{1}_{\{m \leq \theta_1\}},$$

and since $\theta_0 < \theta_1$ there are three cases. If $m \leq \theta_0$, both densities are positive and

$$f(x; \theta_1) = \theta_1^{-n} = (\theta_0/\theta_1)^n \theta_0^{-n} = k f(x; \theta_0).$$

If $\theta_0 < m \leq \theta_1$, we have $f(x; \theta_0) = 0 < f(x; \theta_1)$, and thus $f(x; \theta_1) > k f(x; \theta_0)$. If $m > \theta_1$, both densities vanish and equality holds. Consequently

$$\begin{aligned} \{x \mid f(x; \theta_1) > k f(x; \theta_0)\} &= \{\theta_0 < m \leq \theta_1\}, \\ \{x \mid f(x; \theta_1) = k f(x; \theta_0)\} &= \{m \leq \theta_0\} \cup \{m > \theta_1\}. \end{aligned}$$

2. Since $0 < c < \theta_0$, the size of C is $\mathbb{P}_{\theta_0}(M > c) = 1 - (c/\theta_0)^n = 1 - (1 - \alpha) = \alpha$. The region $C = \{m > c\}$ contains the set $\{\theta_0 < m \leq \theta_1\}$ of part (a), because $c < \theta_0$, and it is contained in the union of the two sets of part (a), which is the whole sample space. Thus C satisfies the condition of theorem 22.2 with $k = (\theta_0/\theta_1)^n$ and by the first statement of the theorem it is most powerful at level α . The power of this test is

$$\mathbb{P}_{\theta_1}(M > c) = 1 - \left(\frac{c}{\theta_1}\right)^n = 1 - (1 - \alpha) \left(\frac{\theta_0}{\theta_1}\right)^n.$$

The threshold $c = \theta_0(1 - \alpha)^{1/n}$ does not depend on θ_1 and the argument from above applies to all $\theta_1 > \theta_0$ with the corresponding k and thus, by proposition 22.4, the test is uniformly most powerful against $H_1: \theta > \theta_0$.

3. The two parts of C'' are disjoint and thus the size of this region is $\mathbb{P}_{\theta_0}(M > \theta_0) + \mathbb{P}_{\theta_0}(M \leq \theta_0 \alpha^{1/n}) = 0 + \alpha = \alpha$. The power of this test is

$$\mathbb{P}_{\theta_1}(M > \theta_0) + \mathbb{P}_{\theta_1}(M \leq \theta_0 \alpha^{1/n}) = 1 - \left(\frac{\theta_0}{\theta_1}\right)^n + \alpha \left(\frac{\theta_0}{\theta_1}\right)^n = 1 - (1 - \alpha) \left(\frac{\theta_0}{\theta_1}\right)^n,$$

This is the same power as for C . Thus both regions are most powerful at level α and, by the second statement of theorem 22.2, they can only differ on the boundary set. And indeed they do: the two regions differ on the sets $C \setminus C'' = \{c < m \leq \theta_0\}$ and $C'' \setminus C = \{m \leq \theta_0 \alpha^{1/n}\}$, both of which are contained in $\{m \leq \theta_0\}$ where $f(x; \theta_1) = k f(x; \theta_0)$. The argument in the uniqueness comment in section 22.2 used the fact that the likelihood ratio has a continuous distribution, but in this case this is not true: on the set $\{M \leq \theta_0\}$, where an event with probability $1 - \alpha$ occurs under H_0 , the ratio $R(X)$ is constant and thus the boundary set has positive probability, despite the model having densities.

□

Solution to exercise 22.4. Since T is sufficient, we can use the factorisation theorem 8.2 to find functions g and h such that

$$f(x; \theta) = g(T(x); \theta) h(x)$$

for all x and all parameter values θ . Now let us choose an x with $f(x; \theta_0) > 0$. Then neither of the two factors in the expression for $f(x; \theta_0) = g(T(x); \theta_0)h(x)$ can be zero and we get

$$R(x) = \frac{f(x; \theta_1)}{f(x; \theta_0)} = \frac{g(T(x); \theta_1) h(x)}{g(T(x); \theta_0) h(x)} = \frac{g(T(x); \theta_1)}{g(T(x); \theta_0)}.$$

Since we can cancel the factor $h(x)$ on both sides, the right-hand side only depends on x via $T(x)$. This completes the proof. $\square$

C.16 Likelihood-Ratio Tests

Here we solve the exercises at the end of lecture 23.

Solution to exercise 23.1.

1. By example 5.4 the log-likelihood

$$\ell(\mu, \sigma^2) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (X_i - \mu)^2$$

is maximised over Θ at $\hat{\mu} = \bar{X}$ and $\hat{\sigma}^2 = Q/n$. Substituting these values, the sum of squares becomes Q and the second term becomes $Q/(2Q/n) = n/2$, and thus the maximised log-likelihood is

$$\ell(\hat{\mu}, \hat{\sigma}^2) = -\frac{n}{2} \log\left(\frac{2\pi Q}{n}\right) - \frac{n}{2}.$$

2. Under H_0 the mean is fixed at μ_0 and only $v = \sigma^2$ is free. Writing $Q_0 = \sum_i (X_i - \mu_0)^2$, the function to maximise is $-\frac{n}{2} \log(2\pi v) - Q_0/(2v)$, which is the function of example 5.4 with Q_0 in place of Q , and thus by the same derivative computation it is maximised at $\hat{\sigma}_0^2 = Q_0/n$, with maximum $-\frac{n}{2} \log(2\pi Q_0/n) - \frac{n}{2}$. The identity (2.1) with μ_0 in place of μ gives $Q_0 = Q + n(\bar{X} - \mu_0)^2$, and thus

$$\hat{\sigma}_0^2 = \frac{Q + n(\bar{X} - \mu_0)^2}{n}.$$

3. By (23.1) and the two previous parts, the terms $-n/2$ cancel and we get

$$W = 2\left(-\frac{n}{2} \log \frac{2\pi Q}{n} + \frac{n}{2} \log \frac{2\pi Q_0}{n}\right) = n \log \frac{Q_0}{Q} = n \log\left(1 + \frac{n(\bar{X} - \mu_0)^2}{Q}\right).$$

Since $Q = (n-1)S^2$, we have $n(\bar{X} - \mu_0)^2/Q = n(\bar{X} - \mu_0)^2/((n-1)S^2) = t^2/(n-1)$, which gives the stated formula. The function $t^2 \mapsto n \log(1 + t^2/(n-1))$ is strictly increasing, and thus the event $W \geq c$ is the event $|t| \geq c'$ for a suitable c' : the generalised likelihood-ratio test rejects for large values of $|t|$.

4. The full parameter $\theta = (\mu, \sigma^2)$ has $k = 2$ coordinates, and H_0 fixes one of them, $\mu = \mu_0$, leaving σ^2 free, so that $r = 1$. The normal model satisfies the regularity conditions, and thus theorem 23.3 gives $W \xrightarrow{d} \chi_1^2$

under H_0 . The approximate test rejects when $W \geq 3.841$, *i.e.* when $t^2 \geq (n-1)(e^{3.841/n} - 1)$. For $n = 10$ this is $t^2 \geq 9(e^{0.3841} - 1) = 9 \cdot 0.4683 = 4.215$, or $|t| \geq 2.053$, whereas the exact test rejects when $|t| \geq 2.262$. The approximate test rejects more often than the exact one, and its true size is $\mathbb{P}(|t_9| \geq 2.053) = 0.070$ instead of the nominal 0.05: the chi-squared approximation is optimistic for small n , because the t_9 distribution has heavier tails than the normal limit which the approximation assumes. For large n the two tests merge, since $n \log(1+t^2/(n-1)) \rightarrow t^2$ by $\log(1+x) \approx x$ and the t_{n-1} distribution tends to the standard normal one; for $n = 100$ the two thresholds are 1.969 and 1.984. Whenever the exact distribution is available, as here, it is the one to use, and lecture 29 does so.

□

Solution to exercise 23.2.

1. Taking logarithms of the probability weights gives

$$\ell(p) = \log \frac{n!}{n_1! \cdots n_m!} + \sum_{j=1}^m n_j \log p_j,$$

where the first term does not depend on p . To maximise $\sum_j n_j \log p_j$ subject to $\sum_j p_j = 1$ we consider the Lagrangian $\sum_j n_j \log p_j - \lambda(\sum_j p_j - 1)$. Setting its partial derivative with respect to p_j to zero gives $n_j/p_j = \lambda$, *i.e.* $p_j = n_j/\lambda$ for every j , and summing over j and using the constraint we find $1 = n/\lambda$, so that $\lambda = n$ and $\hat{p}_j = n_j/n$. This stationary point is the maximum, because $\sum_j n_j \log p_j$ is a concave function of p and the constraint set is convex. If some n_j is zero, the corresponding $\hat{p}_j = 0$ lies on the boundary of the parameter space; the formula still gives the maximum, since any probability assigned to an unobserved outcome only lowers the likelihood.

2. Since H_0 is simple, the restricted maximiser is p^0 , and by (23.1) we get

$$W = 2(\ell(\hat{p}) - \ell(p^0)) = 2\left(\sum_{j=1}^m n_j \log \frac{n_j}{n} - \sum_{j=1}^m n_j \log p_j^0\right) = 2 \sum_{j=1}^m n_j \log \frac{n_j}{n p_j^0},$$

where the combinatorial constant cancels. A term with $n_j = 0$ contributes nothing to either sum, which is the convention $0 \log 0 = 0$.

3. The parameter space $\Theta = \{p \mid p_j > 0, \sum_j p_j = 1\}$ is described by the $m-1$ free coordinates $p_1, \dots, p_{m-1}$, with p_m determined by the constraint, and thus has dimension $k = m-1$. The null hypothesis fixes all of them, so that $r = k = m-1$. The counts are the sums of n i.i.d. indicator vectors, one per repetition, whose distribution is a full-rank exponential family with support not depending on p , so that the regularity conditions hold at every interior point p^0 . Theorem 23.3 therefore gives $W \xrightarrow{d} \chi_{m-1}^2$ under H_0 .
4. Under H_0 : $p_j^0 = 1/6$ the expected counts are $n p_j^0 = 10$ for every face. The terms $n_j \log(n_j/10)$ are

$$\begin{aligned} 13 \log 1.3 &= 3.411, & 8 \log 0.8 &= -1.785, & 11 \log 1.1 &= 1.048, \\ 6 \log 0.6 &= -3.065, & 12 \log 1.2 &= 2.188, & 10 \log 1 &= 0, \end{aligned}$$

with sum 1.797, and thus $w = 2 \cdot 1.797 = 3.59$. Since $3.59 < 11.07 = \chi_{5,0.95}^2$, we do not reject H_0 at level 0.05; the approximate p -value is $\mathbb{P}(\chi_5^2 \geq 3.59) = 0.61$. The observed counts are entirely consistent with a fair die.

5. Write $x_j = d_j/(np_j^0)$, so that $n_j = np_j^0(1 + x_j)$ and $n_j/(np_j^0) = 1 + x_j$. Using $\log(1 + x) = x - x^2/2 + O(x^3)$ we get

$$\begin{aligned} n_j \log \frac{n_j}{np_j^0} &= np_j^0(1 + x_j) \left(x_j - \frac{x_j^2}{2} + O(x_j^3) \right) \\ &= np_j^0 \left(x_j + \frac{x_j^2}{2} + O(x_j^3) \right) = d_j + \frac{d_j^2}{2np_j^0} + O(x_j^3), \end{aligned}$$

since $np_j^0 x_j = d_j$ and $np_j^0 x_j^2 = d_j^2/(np_j^0)$. Summing over j , the linear terms add up to $\sum_j d_j = 0$, and doubling the result gives

$$W \approx \sum_{j=1}^m \frac{d_j^2}{np_j^0} = \sum_{j=1}^m \frac{(n_j - np_j^0)^2}{np_j^0},$$

up to terms of third order in the relative deviations x_j . For the die the deviations from 10 are 3, -2, 1, -4, 2, 0, and Pearson's statistic is $(9 + 4 + 1 + 16 + 4 + 0)/10 = 3.4$, close to the value 3.59 of W ; the two statistics agree to the extent that the relative deviations are small, which for these data they are.

□

Solution to exercise 23.3.

1. By example 4.2 the likelihood is $L(p) = p^S(1 - p)^{n-S}$, so that $\ell(p) = S \log p + (n - S) \log(1 - p)$, and by example 5.3 it is maximised over $(0, 1)$ at $\hat{p} = S/n$ when $0 < S < n$. Since H_0 is simple, the restricted maximiser is p_0 , and (23.1) gives

$$W = 2(\ell(\hat{p}) - \ell(p_0)) = 2 \left(S \log \frac{\hat{p}}{p_0} + (n - S) \log \frac{1 - \hat{p}}{1 - p_0} \right).$$

For $S = 0$ or $S = n$ the supremum of L over $(0, 1)$ is approached at the boundary, where one of the two terms is absent, and the convention $0 \log 0 = 0$ gives the correct value.

2. The parameter is one-dimensional and H_0 fixes it, so that $k = r = 1$; the Bernoulli model satisfies the regularity conditions for $p_0 \in (0, 1)$, as noted in lecture 16. Theorem 23.3 gives $W \xrightarrow{d} \chi_1^2$ under H_0 , and the approximate level α test rejects when $W \geq \chi_{1,1-\alpha}^2$, which for $\alpha = 0.05$ means $W \geq 3.841$.
3. We have $\hat{p} = 32/50 = 0.64$ and $p_0 = 0.5$, and thus

$$w = 2(32 \log 1.28 + 18 \log 0.72) = 2(7.900 - 5.913) = 3.97.$$

Since $3.97 \geq 3.841$, we reject $H_0: p = 1/2$ at level 0.05. The approximate p -value is $\mathbb{P}(\chi_1^2 \geq 3.97) = 2(1 - \Phi(\sqrt{3.97})) = 2(1 - \Phi(1.99)) = 0.047$, just below 0.05, so that the evidence against a fair coin is real but not strong.

4. The map $g(p) = \log(p/(1-p))$ is a strictly increasing bijection from $(0,1)$ onto $\mathbb{R}$ with inverse $p = e^\psi/(1+e^\psi)$, and the likelihood in the new parametrisation is $L_\psi(\psi) = L(g^{-1}(\psi))$. The null hypothesis $\psi = \psi_0$ is the same set of distributions as $p = p_0$ with $p_0 = g^{-1}(\psi_0)$, and thus the numerator of the ratio is $L_\psi(\psi_0) = L(p_0)$ in both cases. Since the supremum $\sup_{p \in (0,1)} L(p)$ is attained at the MLE $\hat{p}$ by the first part, and g is a bijection from $(0,1)$ onto $\mathbb{R}$, the two suprema $\sup_{\psi \in \mathbb{R}} L_\psi(\psi)$ and $\sup_{p \in (0,1)} L(p)$ run over the same set of values, and thus the denominator is $\sup_{\psi \in \mathbb{R}} L_\psi(\psi) = \sup_{p \in (0,1)} L(p) = L(\hat{p})$. Both ratios are therefore equal, and so are the two statistics W . This is the invariance noted in section 23.1; by contrast the standard errors of exercise 17.1 do depend on the parametrisation.

□

C.17 Confidence Intervals

Solution to exercise 25.1.

1. The function $Q = 2\theta T$ is a function of the data and the parameter, and from exercise 13.1 we know that under $\mathbb{P}_\theta$ the distribution of Q is χ_{2n}^2 for all θ , independent of θ . Thus, Q is a pivot. Using $a = \chi_{2n,\alpha/2}^2$ and $b = \chi_{2n,1-\alpha/2}^2$ we find $\mathbb{P}_\theta(a \leq 2\theta T \leq b) = 1 - \alpha$ and since the function Q is increasing in θ , the two inequalities can be solved for θ directly:

$$a \leq 2\theta T \leq b \iff \frac{\chi_{2n,\alpha/2}^2}{2T} \leq \theta \leq \frac{\chi_{2n,1-\alpha/2}^2}{2T}.$$

Thus, the interval

$$[\chi_{2n,\alpha/2}^2/(2T), \chi_{2n,1-\alpha/2}^2/(2T)]$$

is a confidence interval for θ with level $1 - \alpha$ and the probability of this interval covering the correct value is exactly $1 - \alpha$.

2. We have $2T = 50$ and $2n = 20$ and thus the 95% interval is

$$[9.591/50, 34.170/50] = [0.192, 0.683].$$

3. The MLE for the rate is $\hat{\theta} = 1/\bar{x} = 10/25 = 0.4$ and the standard error of the MLE is $0.4/\sqrt{10} = 0.126$. Thus the Wald interval is $0.4 \pm 1.96 \cdot 0.126$, which is the interval $[0.152, 0.648]$. The exact interval is asymmetric about the estimate, extending further to the right than to the left. This is caused by the right skew of the distribution of $\hat{\theta}$ for small exponential samples; the Wald interval, being symmetric, extends to lower values than the exact interval does. At $n = 10$ the normal approximation used in the Wald interval is rough, as noted in section 17.5, and the exact interval should be used instead.

□

Solution to exercise 25.2.

1. From example 8.10 we know that the maximum has density $f_M(t) = nt^{n-1}/\theta^n$ for $0 \leq t \leq \theta$. Thus, for $0 \leq u \leq 1$ we have

$$\mathbb{P}_\theta\left(\frac{M}{\theta} \leq u\right) = \mathbb{P}_\theta(M \leq u\theta) = \int_0^{u\theta} \frac{nt^{n-1}}{\theta^n} dt = \frac{(u\theta)^n}{\theta^n} = u^n.$$

The distribution function of M/θ does not depend on θ and thus M/θ is a pivot.

2. Since $M \leq \theta$ always holds, we have

$$\mathbb{P}_\theta\left(M \leq \theta \leq \frac{M}{\alpha^{1/n}}\right) = \mathbb{P}_\theta\left(\alpha^{1/n} \leq \frac{M}{\theta} \leq 1\right) = 1 - (\alpha^{1/n})^n = 1 - \alpha,$$

where we used the first part of the solution. Thus, $[M, M/\alpha^{1/n}]$ is a confidence interval with level $1 - \alpha$. For $n = 5$ and $\alpha = 0.05$ we have $\alpha^{1/n} = 0.05^{0.2} = 0.549$ and thus the interval is $[3.2, 3.2/0.549] = [3.2, 5.83]$.

3. The condition for coverage is $b^n - a^n = 1 - \alpha$. This condition can be used to find $a = (b^n - 1 + \alpha)^{1/n}$ as a function of b and thus, by choosing the interval $[M/b, M/a]$, we can find the length $M(1/a - 1/b)$ of the interval. Taking derivatives we find $b^{n-1} = a^{n-1} da/db$ and thus

$$\frac{d}{db}\left(\frac{1}{a} - \frac{1}{b}\right) = -\frac{1}{a^2} \frac{da}{db} + \frac{1}{b^2} = -\frac{b^{n-1}}{a^{n+1}} + \frac{1}{b^2} = \frac{1}{b^2}\left(1 - \left(\frac{b}{a}\right)^{n+1}\right) < 0,$$

since $b > a$. The length of the interval decreases as b increases and thus is smallest for the largest possible value $b = 1$, where $a = \alpha^{1/n}$. This is the interval from the previous part of the solution. The reason for this result is that the density nu^{n-1} of the pivot is increasing on $[0, 1]$ and thus the probability $1 - \alpha$ is concentrated on the shortest possible interval which is placed at the upper end of the density.

□

Solution to exercise 25.3.

1. The standard error is $\text{se}(\hat{\theta}_n) = 1/\sqrt{n\mathcal{I}(\bar{X})} = \sqrt{\bar{X}/n}$, and the Wald interval is $\bar{X} \pm z_{1-\alpha/2}\sqrt{\bar{X}/n}$. For the data we have $\bar{x} = 1.5$ and $\sqrt{1.5/20} = 0.274$, and the 95% interval is $1.5 \pm 1.96 \cdot 0.274$, which is $[0.96, 2.04]$.
2. In exercise 17.2 we found $\sqrt{n}(\hat{p}_0 - p_0) \xrightarrow{d} N(0, e^{-2\theta}\theta)$, and replacing θ by $\bar{X}$ in the variance gives the standard error $\text{se}(\hat{p}_0) = e^{-\bar{X}}\sqrt{\bar{X}/n}$, in agreement with the delta-method formula $|g'(\hat{\theta}_n)| \text{se}(\hat{\theta}_n)$ for $g(\theta) = e^{-\theta}$. The approximate interval is $e^{-\bar{X}} \pm z_{1-\alpha/2} e^{-\bar{X}}\sqrt{\bar{X}/n}$. For the data we have $\hat{p}_0 = e^{-1.5} = 0.223$ and $\text{se}(\hat{p}_0) = 0.223 \cdot 0.274 = 0.061$, and the 95% interval is $0.223 \pm 1.96 \cdot 0.061$, which is $[0.103, 0.343]$.
3. Write $[L, U]$ for the interval of the first part. Since g is decreasing, the event $L \leq \theta \leq U$ is the same as the event $g(U) \leq g(\theta) \leq g(L)$, and thus the interval $[e^{-U}, e^{-L}]$ covers p_0 exactly when $[L, U]$ covers θ . Its coverage probability is that of the Wald interval, which converges to $1 - \alpha$ by proposition 25.7. For the data we get $[e^{-2.04}, e^{-0.96}] = [0.130, 0.383]$.

The two intervals differ because g is not linear: the delta-method interval is symmetric about $\hat{p}_0$, while the transformed interval is shifted towards larger values and is asymmetric about $\hat{p}_0 = 0.223$. Both have asymptotic coverage $1 - \alpha$, and for large n the two nearly coincide, since g is close to linear over a short interval. A practical advantage of the transformed interval is that it always lies inside $(0, 1)$, whereas the delta-method interval can extend below zero when $\hat{p}_0$ is small.

□

Solution to exercise 25.4.

1. Under H_0 , the test statistic $\sqrt{n}(\bar{X} - \mu_0)/S$ is the pivot T from example 25.4 and has t_{n-1} distribution, independent of the value of σ^2 . Since the t distribution is symmetric, we have

$$\begin{aligned}\mathbb{P}_{\mu_0, \sigma^2}(|T| > t_{n-1, 1-\alpha/2}) &= \mathbb{P}(T > t_{n-1, 1-\alpha/2}) + \mathbb{P}(T < -t_{n-1, 1-\alpha/2}) \\ &= \frac{\alpha}{2} + \frac{\alpha}{2} = \alpha,\end{aligned}$$

and the test has size α .

2. The acceptance region of the test for μ_0 is $A(\mu_0) = \{x \mid |\bar{x} - \mu_0| \leq t_{n-1, 1-\alpha/2} s/\sqrt{n}\}$. By theorem 25.6, the set of values not rejected,

$$\begin{aligned}C(X) &= \left\{ \mu_0 \mid |\bar{X} - \mu_0| \leq t_{n-1, 1-\alpha/2} \frac{S}{\sqrt{n}} \right\} \\ &= \left[\bar{X} - t_{n-1, 1-\alpha/2} \frac{S}{\sqrt{n}}, \bar{X} + t_{n-1, 1-\alpha/2} \frac{S}{\sqrt{n}} \right],\end{aligned}$$

is a confidence set with level $1 - \alpha$ and it is the t -interval from example 25.4.

3. The 95% t -interval for the given data is $[9.18, 11.42]$. By duality, $H_0: \mu = \mu_0$ is rejected at level 5% if and only if μ_0 is outside this interval. Thus, $H_0: \mu = 9$ is rejected since $9 < 9.18$, and $H_0: \mu = 10$ is not rejected since 10 is inside the interval.
4. The one-sided test does not reject μ_0 if and only if $\sqrt{n}(\bar{X} - \mu_0)/S \leq t_{n-1, 1-\alpha}$, i.e. if $\mu_0 \geq \bar{X} - t_{n-1, 1-\alpha} S/\sqrt{n}$. Thus, the set of values not rejected is the one-sided interval $[\bar{X} - t_{n-1, 1-\alpha} S/\sqrt{n}, \infty)$, whose left-hand boundary is a lower confidence bound for μ with level $1 - \alpha$. Under H_0 , the test statistic is t_{n-1} distributed, and thus the test has size α and the bound has coverage probability exactly $1 - \alpha$, by theorem 25.6.

□

C.18 Bayesian Inference: Priors and Posteriors

Solution to exercise 26.1.

1. The prior mean is $a/(a+b) = 0.3$ and since $a+b = 10$ we have $a = 3$ and $b = 7$. From table A.2 we find that the variance of the Beta(3, 7) distribution is

$$\frac{ab}{(a+b)^2(a+b+1)} = \frac{21}{100 \cdot 11} \approx 0.0191,$$

and thus the prior standard deviation is approximately 0.138.

- Let $\mu = a/(a+b)$ denote the prior mean and $s = a+b$ the equivalent sample size. Then we have $a = \mu s$ and $b = (1-\mu)s$. Using the formula for the variance from the table, we get

$$\frac{ab}{(a+b)^2(a+b+1)} = \frac{\mu(1-\mu)}{s+1}.$$

Given $\mu = 0.3$ and a variance of $0.1^2 = 0.01$ we find that we need $s+1 = 0.21/0.01 = 21$, *i.e.* $s = 20$ and thus we have $a = 6$ and $b = 14$. The equivalent sample size is 20: the smaller spread corresponds to a prior which gives twice the number of observations.

- From example 26.5 we know that the posterior distribution is $\text{Beta}(3+9, 7+11) = \text{Beta}(12, 18)$. The mean of this distribution is $12/30 = 0.4$. In the form (26.3) we can write this as

$$\frac{10}{30} \cdot 0.3 + \frac{20}{30} \cdot 0.45 = 0.1 + 0.3 = 0.4,$$

where $\bar{x} = 9/20 = 0.45$: The data have twice the weight as the prior, and the posterior mean is two thirds of the way between the prior mean and the sample proportion.

□

Solution to exercise 26.2.

- On the interval $(0, 1/2]$ we have $1-\theta \geq 1/2$ and thus $\theta^{-1}(1-\theta)^{-1} \geq \theta^{-1}$ there. Since $\int_0^{1/2} \theta^{-1} d\theta = \infty$, the integral of the prior over $(0, 1)$ is infinite and the prior is improper.
- Consider the kernel $\theta^{\alpha-1}(1-\theta)^{\beta-1}$ for real α and β . Near $\theta = 0$ the factor $(1-\theta)^{\beta-1}$ is bounded between two positive constants and thus the kernel is integrable near 0 if and only if $\theta^{\alpha-1}$ is, which is equivalent to $\alpha > 0$. By symmetry the kernel is integrable near $\theta = 1$ if and only if $\beta > 0$ and thus, since it is bounded away from the endpoints, the integral over $(0, 1)$ is finite if and only if $\alpha > 0$ and $\beta > 0$, in which case it equals $B(\alpha, \beta)$. For the formal posterior we have $\alpha = T$ and $\beta = n - T$ and thus the posterior is proper if and only if $T > 0$ and $n - T > 0$, *i.e.* if and only if $0 < T < n$. In this case it is the $\text{Beta}(T, n - T)$ distribution.
- From table A.2 we find that the posterior mean is $T/(T+n-T) = T/n = \bar{x}$. This is the maximum likelihood estimate from example 5.3. The Haldane prior is the choice of prior where the Bayes estimator and the maximum likelihood estimator coincide. This happens for an equivalent sample size of zero in (26.3).
- For $T = 0$ the likelihood $(1-\theta)^n$ has its maximum at $\theta = 0$ and does not vanish there, while the prior has a non-integrable singularity at $\theta = 0$. Since the data does not contain any successes, the infinite mass of the prior near zero cannot be compensated and thus the integral of $\theta^{-1}(1-\theta)^{n-1}$ is infinite: there is no posterior distribution in this case. In informal language, the improper prior suggests that θ is very close to either 0 or 1 and a data set with no successes cannot exclude the first of these two possibilities.

□

Solution to exercise 26.3.

1. We have

$$\begin{aligned} f(x|\theta) &= \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x_i - \theta)^2}{2\sigma^2}\right) \\ &= (2\pi\sigma^2)^{-n/2} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \theta)^2\right). \end{aligned}$$

Using the identity (2.1) from lecture 2, which holds for all numbers in place of the mean, and using θ , we find $\sum_i (x_i - \theta)^2 = \sum_i (x_i - \bar{x})^2 + n(\bar{x} - \theta)^2$. Since the first term does not depend on θ and since the factor $(2\pi\sigma^2)^{-n/2}$ does not depend on θ , we have

$$f(x|\theta) \propto \exp(-n(\theta - \bar{x})^2/(2\sigma^2)).$$

2. Since the prior density, as a function of θ , is proportional to $\exp(-(\theta - m)^2/(2\tau^2))$ and since we can group the terms in the exponent of the posterior density as powers of θ , we get

$$\begin{aligned} \pi(\theta|x) &\propto \exp\left(-\frac{n(\theta - \bar{x})^2}{2\sigma^2} - \frac{(\theta - m)^2}{2\tau^2}\right) \\ &\propto \exp\left(-\frac{1}{2}\left(\frac{n}{\sigma^2} + \frac{1}{\tau^2}\right)\theta^2 + \left(\frac{n\bar{x}}{\sigma^2} + \frac{m}{\tau^2}\right)\theta\right), \end{aligned}$$

where we have omitted the terms which do not contain θ . Using the abbreviations $1/\tau_n^2 = 1/\tau^2 + n/\sigma^2$ and $m_n = \tau_n^2(m/\tau^2 + n\bar{x}/\sigma^2)$ we can write the exponent as $-\theta^2/(2\tau_n^2) + \theta m_n/\tau_n^2$ and completing the square we find

$$-\frac{\theta^2}{2\tau_n^2} + \frac{\theta m_n}{\tau_n^2} = -\frac{(\theta - m_n)^2}{2\tau_n^2} + \frac{m_n^2}{2\tau_n^2},$$

where we have omitted the last term which does not contain θ . Thus we have $\pi(\theta|x) \propto \exp(-(\theta - m_n)^2/(2\tau_n^2))$ and this is the kernel of the $N(m_n, \tau_n^2)$ density and, since the posterior is a density, this is the corresponding normal distribution. If we multiply out m_n we get the weighted form of the posterior given in proposition 26.6.

3. As $\tau \rightarrow \infty$ with n fixed, the prior precision $1/\tau^2$ goes to zero and thus $\tau_n^2 \rightarrow \sigma^2/n$ and $m_n \rightarrow \bar{x}$: the posterior converges to $N(\bar{x}, \sigma^2/n)$, *i.e.* to the posterior under the flat prior discussed in section 26.4. As $n \rightarrow \infty$ with τ fixed, the weight $(1/\tau^2)/(1/\tau^2 + n/\sigma^2)$ of the prior mean goes to zero and thus $m_n - \bar{x} \rightarrow 0$ and $\tau_n^2 \leq \sigma^2/n$ goes to zero: the posterior is now concentrated around $\bar{x}$ and the prior is washed out.

□

Solution to exercise 26.4.

1. Under $\mathbb{P}_\theta$, the total number T of successes is binomially distributed with $\mathbb{E}_\theta(T) = n\theta$ and $\text{Var}_\theta(T) = n\theta(1 - \theta)$. Since $\hat{\theta}_B$ is an affine function of T , we have

$$\begin{aligned}\mathbb{E}_\theta(\hat{\theta}_B) &= \frac{a + n\theta}{a + b + n} \quad \text{and} \\ \text{bias}(\hat{\theta}_B) &= \frac{a + n\theta - (a + b + n)\theta}{a + b + n} = \frac{a - (a + b)\theta}{a + b + n},\end{aligned}$$

which equals zero if and only if $\theta = a/(a + b)$, the prior mean, and

$$\text{Var}_\theta(\hat{\theta}_B) = \frac{n\theta(1 - \theta)}{(a + b + n)^2}.$$

The bias is of order $1/n$ and the variance is of order $1/n$, both tending to zero as $n \rightarrow \infty$. Thus, by corollary 2.11, the estimator is consistent, although for fixed n it is biased for all $\theta \neq a/(a + b)$.

2. For $a = b = 1$ and $n = 10$ we have $\hat{\theta}_B = (T + 1)/12$ and thus, by theorem 2.7,

$$\text{MSE}(\hat{\theta}_B) = \frac{10\theta(1 - \theta) + (1 - 2\theta)^2}{144}, \quad \text{MSE}(\bar{X}) = \frac{\theta(1 - \theta)}{10}.$$

At $\theta = 1/2$ the bias is zero and we have $\text{MSE}(\hat{\theta}_B) = 2.5/144 \approx 0.0174$ compared to $\text{MSE}(\bar{X}) = 0.025$. Thus, the Bayes estimator is better. At $\theta = 0.1$ we have $\text{MSE}(\hat{\theta}_B) = (0.9 + 0.64)/144 \approx 0.0107$ compared to $\text{MSE}(\bar{X}) = 0.009$ and thus $\bar{X}$ is better. Neither estimator has smaller mean squared error for all θ : the Bayes estimator is better in the middle of the parameter range, where the prior mean is close to the truth, and the reduced variance does not compensate for the bias introduced by the shrinkage towards $1/2$ at the ends of the range.

□

Solution to exercise 26.5. For the first part of the question we can take derivatives. The log-likelihood function for a single observation is $\ell_0(x - \theta)$ and thus the score function is

$$\frac{\partial}{\partial \theta} \log f(x|\theta) = \frac{\partial}{\partial \theta} \ell_0(x - \theta) = -\ell'_0(x - \theta).$$

Since $Z = X - \theta$, with density f_0 , for all values of θ , the Fisher information is

$$\mathcal{I}(\theta) = \mathbb{E}(\ell'_0(X - \theta)^2) = \int \ell'_0(z)^2 f_0(z) dz.$$

This does not depend on θ .

For the second part of the question we have $\log f(x|\sigma) = \ell_0(x/\sigma) - \log \sigma$ and thus the score function is

$$\frac{\partial}{\partial \sigma} (\ell_0(x/\sigma) - \log \sigma) = -\ell'_0(x/\sigma) \frac{x}{\sigma^2} - \frac{1}{\sigma} = -\frac{1 + z\ell'_0(z)}{\sigma}$$

where $z = x/\sigma$. Since $Z = X/\sigma$ has density f_0 for all σ , we find

$$\mathcal{I}(\sigma) = \frac{1}{\sigma^2} \mathbb{E}((1 + Z\ell'_0(Z))^2) = \frac{c}{\sigma^2}, \quad \text{where } c = \int (1 + z\ell'_0(z))^2 f_0(z) dz$$

does not depend on σ .

Finally, the $N(0, \sigma^2)$ model is the scale model with base density the standard normal density and with $\ell_0(z) = -z^2/2 - \log \sqrt{2\pi}$ and thus $\ell'_0(z) = -z$. In this case we have $1 + z\ell'_0(z) = 1 - z^2$ and using the substitutions $\mathbb{E}(Z^2) = 1$ and $\mathbb{E}(Z^4) = 3$ for the standard normal density Z we find

$$c = \mathbb{E}((1 - Z^2)^2) = 1 - 2\mathbb{E}(Z^2) + \mathbb{E}(Z^4) = 2.$$

Thus $\mathcal{I}(\sigma) = 2/\sigma^2$, as claimed. $\square$

C.19 Credible Intervals and Bayesian Testing

Solution to exercise 28.1.

1. Since the posterior density is monotone on each half-line and integrates to one, it converges to zero as $\theta \rightarrow \pm\infty$. The condition $\mathbb{P}(\theta \in H_c|x) = 1 - \alpha \in (0, 1)$ excludes $c \geq \pi(\mu|x)$, where H_c contains at most the point μ and has probability zero, and thus $0 < c < \pi(\mu|x)$. On $(-\infty, \mu]$ the density is continuous and strictly increasing from 0 to $\pi(\mu|x)$ and thus there is exactly one $l < \mu$ with $\pi(l|x) = c$ and $\pi(\theta|x) \geq c$ for $\theta \leq \mu$ if and only if $\theta \geq l$. Similarly, there is exactly one $u > \mu$ with $\pi(u|x) = c$ and the density is at least c on $[\mu, \infty)$ exactly on $[\mu, u]$. Thus we have $H_c = [l, u]$ with $\pi(l|x) = c = \pi(u|x)$.
2. Let $l = \mu - s$ and $u = \mu + t$ with $s, t > 0$. By symmetry we have $\pi(\mu + s|x) = \pi(\mu - s|x) = c = \pi(\mu + t|x)$ and since the density is strictly decreasing on $[\mu, \infty)$ we find $s = t$. The two tails $(-\infty, \mu - t)$ and $(\mu + t, \infty)$ have equal posterior probability by symmetry and together they have probability α , i.e. each tail has probability $\alpha/2$. Thus we have $l = q_{\alpha/2}$ and $u = q_{1-\alpha/2}$, i.e. we have found the equal-tailed interval.
3. The set $\{\theta \mid \pi(\theta|x) \geq c\}$ is for $0 < c < \pi(0|x)$ the interval $[0, u]$ with $\pi(u|x) = c$, obtained by using the same argument as in the first part, but only for the decreasing branch of the density, and using the condition $\mathbb{P}(\theta \leq u|x) = 1 - \alpha$ to get $u = q_{1-\alpha}$. The equal-tailed interval $[q_{\alpha/2}, q_{1-\alpha/2}]$ is then obtained from $[0, q_{1-\alpha}]$ by removing the piece $[0, q_{\alpha/2}]$ and adding the piece $(q_{1-\alpha}, q_{1-\alpha/2}]$, both of which have posterior probability $\alpha/2$. The density is at least $d_1 = \pi(q_{\alpha/2}|x)$ on the first piece and at most $d_2 = \pi(q_{1-\alpha}|x)$ on the second piece, where $d_1 > d_2$ since the density is strictly decreasing. Writing λ_1 and λ_2 for the lengths of the two pieces we find $\alpha/2 \geq d_1\lambda_1$ and $\alpha/2 \leq d_2\lambda_2$ and thus $\lambda_1 \leq \alpha/(2d_1) < \alpha/(2d_2) \leq \lambda_2$: the added piece is longer than the removed piece and the equal-tailed interval is longer than the HPD interval. $\square$

Solution to exercise 28.2.

1. The equal-tailed 95% credible interval is $[0.386, 0.861]$, with posterior probability 2.5% cut off in each tail. The midpoint of this interval is $(0.386 + 0.861)/2 = 0.624$, which is slightly below the posterior mean $9/14 = 0.643$.

2. The posterior density is proportional to $\theta^8(1-\theta)^4$ on $(0, 1)$. Setting the derivative of the logarithm of the posterior density, *i.e.* $8 \log \theta + 4 \log(1-\theta)$, equal to zero, gives the maximum of the density at $8/\theta = 4/(1-\theta)$ and thus the mode is at $\theta = 8/12 = 2/3 \approx 0.667$. Since the mode is above the mean 0.643, the posterior distribution is skewed.
3. By definition 28.2 of the HPD interval, the HPD interval contains the values with the highest posterior density and thus is centred around the mode, so that it lies to the right of the equal-tailed interval. The HPD interval is the shortest interval which has posterior probability 0.95, in this case of length 0.473 against 0.475. As n increases, the posterior distribution approaches a symmetric, normal distribution and the mode and mean move together. In the limit, for a symmetric, unimodal posterior, both intervals coincide; this is a result from exercise 28.1.

□

Solution to exercise 28.3.

1. Since the prior is flat, we have $\text{Beta}(1, 1)$ and using $n = 3$ and $T = 2$ from example 26.5 we find the posterior $\text{Beta}(3, 2)$. Since $B(3, 2) = \Gamma(3)\Gamma(2)/\Gamma(5) = 2/24 = 1/12$, the density of this posterior is $\pi(\theta|x) = 12\theta^2(1-\theta)$ for $0 < \theta < 1$.
2. We can find the posterior probability of H_0 by integrating the density:

$$\mathbb{P}(\theta \leq 1/2|x) = 12 \int_0^{1/2} (\theta^2 - \theta^3) d\theta = 12 \left(\frac{1}{24} - \frac{1}{64} \right) = \frac{5}{16} = 0.3125.$$

Thus we have $\mathbb{P}(H_1|x) = 11/16$ and the posterior odds for H_0 against H_1 are $5/11 \approx 0.45$.

3. The data have multiplied the odds by 5/11. Using Bayes' theorem for the events H_0 and H_1 , we find that the posterior probability of H_j is $\mathbb{P}(H_j)$ times the average of the likelihoods over H_j , with respect to the prior conditioned on H_j , divided by the marginal density of the data. The factor by which the odds change is the ratio of these two averaged likelihoods. In this case, the likelihood is $\theta^2(1-\theta)$ and the conditional priors are uniform on $(0, 1/2)$ and on $(1/2, 1)$, and thus the factor is

$$\frac{2 \int_0^{1/2} \theta^2(1-\theta) d\theta}{2 \int_{1/2}^1 \theta^2(1-\theta) d\theta} = \frac{5/192}{1/12 - 5/192} = \frac{5}{11}.$$

This agrees with the result from the previous part. Note that this is not the ratio of likelihoods at two values, since each hypothesis includes a range of parameter values, and these need to be weighted by the prior before the two hypotheses can be compared. Only for two simple hypotheses do the averages coincide with single likelihood values, as in case (28.3).

4. The bounds l and u satisfy $12 \int_0^l (\theta^2 - \theta^3) d\theta = 0.025$ and $12 \int_0^u (\theta^2 - \theta^3) d\theta = 0.975$, *i.e.* we have

$$4l^3 - 3l^4 = 0.025 \quad \text{and} \quad 4u^3 - 3u^4 = 0.975.$$

These are polynomial equations of degree four and while they can be solved explicitly, the solutions are cumbersome and not very instructive; for a general Beta posterior, the bounds will not be algebraic. The quantiles are best found numerically, for example using the R function `qbeta` to get $l = 0.194$ and $u = 0.932$ and thus the 95% equal-tailed credible interval is $[0.194, 0.932]$.

□

Solution to exercise 28.4.

1. Using the formula from example 28.7, with $n = 25$, $\sigma^2 = 4$, $\theta_0 = 10$, $\theta_1 = 11$ and $\bar{x} = 10.9$, we find that the Bayes factor in favour of H_0 is

$$\frac{f(x|\theta_0)}{f(x|\theta_1)} = \exp\left(\frac{25 \cdot 1 \cdot (21 - 21.8)}{8}\right) = e^{-2.5} = 0.082.$$

The data give evidence in favour of H_1 over H_0 by a factor of $e^{2.5} \approx 12.2$.

2. The prior odds are 3 and thus the posterior odds are $3 \cdot 0.082 = 0.246$ and $\mathbb{P}(H_0|x) = 0.246/1.246 = 0.198$. The data have changed the prior opinion, which was three to one in favour of H_0 , to a posterior opinion of approximately four to one against it.
3. The p -value 0.0122 is the probability, under $\theta = 10$, of obtaining a sample mean of at least 10.9. This value refers to the data, includes outcomes more extreme than the one observed and makes no reference to any alternative. In contrast, the posterior probability 0.198 is a probability for the hypothesis; it uses only the observed data, compares the likelihoods at exactly 10 and 11 and takes into account the prior odds, which were chosen in favour of H_0 . Even if the prior odds had been 1, the posterior probability would have been $0.082/1.082 = 0.076$, or six times the p -value. Both numbers show that the data are evidence against H_0 , but in different ways. Neither number is a rewording of the other.
4. The posterior odds are the prior odds times $e^{-2.5}$, and the posterior odds are greater than 1 if and only if the prior odds are greater than $e^{2.5} \approx 12.2$. Only a prior which assumed H_0 to be more likely than H_1 , by a factor of more than twelve, would result in H_0 being the more likely hypothesis after these data are observed.

□

Solution to exercise 28.5.

1. By (28.1) we have $1 - w = (n/\sigma^2)/(1/\tau^2 + n/\sigma^2) = (n/\sigma^2)\tau_n^2$, and thus $\tau_n^2 = (1 - w)\sigma^2/n$. The posterior mean is $m_n = wm + (1 - w)\bar{X}$, and thus $m_n - \theta = w(m - \theta) + (1 - w)(\bar{X} - \theta)$. Under $\mathbb{P}_\theta$ we have $\bar{X} - \theta \sim N(0, \sigma^2/n)$ by proposition A.3, and an affine function of a normal variable is normal, so that $m_n - \theta \sim N(w(m - \theta), (1 - w)^2\sigma^2/n)$.
2. Write $z = z_{1-\alpha/2}$ and $\delta = \sqrt{n}(m - \theta)/\sigma$ for the standardised distance between the prior mean and the true value. The interval covers θ if and

only if $|m_n - \theta| \leq z\tau_n = z\sqrt{1-w}\sigma/\sqrt{n}$. Subtracting the mean $w(m - \theta)$ and dividing by the standard deviation $(1-w)\sigma/\sqrt{n}$ turns $m_n - \theta$ into a standard normal variable, and the bounds into $(\pm z\sqrt{1-w} - w\delta)/(1-w)$. Thus the coverage probability is

$$\mathbb{P}_\theta(\theta \in \text{interval}) = \Phi\left(\frac{z\sqrt{1-w} - w\delta}{1-w}\right) - \Phi\left(\frac{-z\sqrt{1-w} - w\delta}{1-w}\right).$$

It depends on θ only through δ , and it is symmetric in δ .

3. Here $1/\tau^2 = 1$ and $n/\sigma^2 = 4$, and thus $w = 0.2$, $1-w = 0.8$, $\sqrt{1-w} = 0.894$ and $z\sqrt{1-w} = 1.753$. At $\theta = m = 12$ we have $\delta = 0$ and the coverage probability is $2\Phi(1.753/0.8) - 1 = 2\Phi(2.19) - 1 = 0.971$, above the nominal 0.95. At $\theta = 10$ we have $\delta = 4 \cdot 2/2 = 4$ and $w\delta = 0.8$, and the coverage probability is $\Phi(0.953/0.8) - \Phi(-2.553/0.8) = \Phi(1.19) - \Phi(-3.19) = 0.883 - 0.001 = 0.882$, below the nominal value. As $|\theta - m| \rightarrow \infty$ we have $|\delta| \rightarrow \infty$, both arguments of Φ tend to the same infinity and the coverage probability tends to zero: the interval is pulled towards the prior mean by a fixed fraction w of the distance, and once that pull exceeds the half-width the interval almost never contains the truth.
4. It is not a confidence interval with level 0.95, because definition 25.1 requires coverage probability at least 0.95 for every θ , and the previous part exhibits values of θ where it is smaller. The credible interval trades the uniform guarantee for better behaviour near the prior mean, which is where a well-chosen prior expects the truth to be. As $n \rightarrow \infty$ with θ fixed we have $w \approx \sigma^2/(n\tau^2) \rightarrow 0$, and thus $w\delta \approx \sigma(m-\theta)/(\tau^2\sqrt{n}) \rightarrow 0$ and $\sqrt{1-w} \rightarrow 1$, so that the coverage probability tends to $2\Phi(z) - 1 = 1 - \alpha$ for every θ . This is the large-sample agreement of section 28.3 seen from the frequentist side.

□

C.20 Standard Tests Revisited

Solution to exercise 29.1.

1. We have the sum of observations 154 and thus $\bar{x} = 154/8 = 19.25$. The sum of squared deviations from $\bar{x}$ equals 3.30 and thus we have $s^2 = 3.30/7 = 0.471$ and $s = 0.687$. The observed value of the test statistic is

$$t = \frac{\sqrt{8}(19.25 - 20)}{0.687} = \frac{2.828 \cdot (-0.75)}{0.687} = -3.09.$$

2. From proposition 29.1 we know that the test against $\mu < 20$ rejects if $t < -t_{7,1-\alpha}$. For level 0.05 the critical value is $-t_{7,0.95} = -1.895$ and for level 0.01 the critical value is $-t_{7,0.99} = -2.998$; since -3.09 lies below both of these critical values, the hypothesis H_0 is rejected at both levels. The p -value is given by $\mathbb{P}(T \leq -3.09) = \mathbb{P}(T \geq 3.09)$ for $T \sim t_7$, by symmetry, and since 3.09 is between $t_{7,0.99} = 2.998$ and $t_{7,0.995} = 3.499$, the value is between 0.005 and 0.01.

3. The two-sided test rejects if $|t| > t_{7,1-\alpha/2}$. For level 0.05 we have $|t| = 3.09 > t_{7,0.975} = 2.365$ and thus H_0 is rejected. For level 0.01 we have $3.09 < t_{7,0.995} = 3.499$ and thus H_0 is not rejected. The two-sided p -value is $2\mathbb{P}(T \geq 3.09)$, which is twice the value of the one-sided test, so it is between 0.01 and 0.02.
4. The 99% interval is $\bar{x} \pm t_{7,0.995} s/\sqrt{8} = 19.25 \pm 3.499 \cdot 0.687/2.828 = 19.25 \pm 0.849$. This interval is $[18.40, 20.10]$. The interval contains the value 20 and by theorem 25.6 this is equivalent to the non-rejection of $H_0: \mu = 20$ by the two-sided test with level 0.01 in the previous part. The 95% interval $19.25 \pm 2.365 \cdot 0.243 = [18.68, 19.82]$ does not contain 20 and thus the rejection at level 0.05 is consistent with the results of the other parts.

□

Solution to exercise 29.2.

1. The value of the test statistic (29.2) is $v = 14 \cdot 2.9/1.5 = 27.07$. We can compare this value to quantiles of the χ_{14}^2 distribution: at level 0.05 the critical value is $\chi_{14,0.95}^2 = 23.685$ and since $27.07 > 23.685$ we can reject H_0 , whereas at level 0.01 the critical value is $\chi_{14,0.99}^2 = 29.141$ and since $27.07 < 29.141$ we cannot reject H_0 . The p -value, $\mathbb{P}(W \geq 27.07)$ for $W \sim \chi_{14}^2$, is between 0.01 and 0.025, since v is between $\chi_{14,0.975}^2 = 26.119$ and $\chi_{14,0.99}^2 = 29.141$.
2. The two-sided test with level 0.05 does not reject, if

$$v \in [\chi_{14,0.025}^2, \chi_{14,0.975}^2] = [5.629, 26.119].$$

Since $27.07 > 26.119$, we reject H_0 .

3. Let $c = \chi_{n-1,1-\alpha}^2$. Assuming the true variance is σ^2 , the value $W = (n-1)S^2/\sigma^2$ is χ_{n-1}^2 distributed by proposition A.3 and thus $V = (n-1)S^2/\sigma_0^2 = W\sigma^2/\sigma_0^2$. We have

$$\beta(\sigma^2) = \mathbb{P}_{\sigma^2}(V > c) = \mathbb{P}\left(W > \frac{c\sigma_0^2}{\sigma^2}\right).$$

As σ^2 increases, the threshold $c\sigma_0^2/\sigma^2$ decreases and the probability of W exceeding a smaller number increases, *i.e.* β increases as a function of σ^2 . By definition 20.5, the size of the test for the composite null hypothesis $\sigma^2 \leq \sigma_0^2$ is $\sup_{\sigma^2 \leq \sigma_0^2} \beta(\sigma^2) = \beta(\sigma_0^2) = \mathbb{P}(W > c) = \alpha$, where the bound is achieved at the boundary.

4. The two-sided test with level α rejects large values of V only if $V > \chi_{n-1,1-\alpha/2}^2$. Since this is a larger threshold than the $\chi_{n-1,1-\alpha}^2$ of the one-sided test, the power of the two-sided test against the alternative $\sigma^2 > \sigma_0^2$ is smaller than the power of the one-sided test. An observed value between the two thresholds is rejected by the one-sided test but not by the two-sided test. The two-sided test has power above α against the alternative $\sigma^2 < \sigma_0^2$, where the one-sided test has probability of rejection below α . For the given data, both tests reject the hypothesis at level 0.05, but at level 0.025 the one-sided test, with threshold $\chi_{14,0.975}^2 = 26.119$, still rejects, whereas the two-sided test compares 27.07 to $\chi_{14,0.9875}^2 = 28.42$ and does not reject.

□

Solution to exercise 29.3.

1. We can describe the differences $D_i = X_i - Y_i$ using i.i.d. $N(\mu_D, \sigma_D^2)$, where both parameters are unknown. We test $H_0: \mu_D = 0$ against $H_1: \mu_D > 0$, since we expect the reaction time to be smaller after training. The differences are 4, -1, 6, 3, 6, 5 and the sum of these differences is 23, *i.e.* we have $\bar{d} = 3.833$. The sum of squared deviations from $\bar{d}$ is 34.83, and thus we have $s_D^2 = 34.83/5 = 6.967$ and $s_D = 2.639$. The test statistic from proposition 29.1, for $\mu_0 = 0$, is

$$t = \frac{\sqrt{6} \cdot 3.833}{2.639} = 3.56,$$

where 5 is the number of degrees of freedom. Since $3.56 > t_{5,0.95} = 2.015$ we can reject the hypothesis at level 0.05 and since $3.56 > t_{5,0.99} = 3.365$ we can also reject the hypothesis at level 0.01: the data show a reduction in reaction time.

2. The 95% t -interval from example 25.4 for μ_D is $\bar{d} \pm t_{5,0.975} s_D / \sqrt{6} = 3.833 \pm 2.571 \cdot 2.639 / 2.449 = 3.833 \pm 2.770$, which is $[1.06, 6.60]$.
3. The row means are $\bar{x} = 419/6 = 69.83$ and $\bar{y} = 396/6 = 66.00$, the pooled variance is $s_p^2 = (5 \cdot 59.77 + 5 \cdot 42.80)/10 = 51.28$ and thus we have $s_p = 7.161$. The test statistic (29.3) is

$$t = \frac{69.83 - 66.00}{7.161 \sqrt{1/6 + 1/6}} = \frac{3.833}{4.135} = 0.93,$$

where 10 is the number of degrees of freedom. Since $0.93 < t_{10,0.95} = 1.812$ the colleague does not reject the hypothesis at level 0.05. The two-sample test is wrong, because the two measurements for one subject are not independent (a slow person is likely to be slow after training, too), whereas the two-sample test assumes that the twelve observations are independent. The difference in opinion is caused by the variability of the data. Between subjects the standard deviation of the reaction times is about 7. The two-sample test compares the mean difference 3.83 to this spread. Inside subjects the differences have standard deviation only 2.64. The paired test compares 3.83 to this smaller spread. By taking differences we remove the effect of the subjects from the comparison, and this is the idea of the paired test.

□

Solution to exercise 29.4.

1. For $\mu_0 = 9$ the value of the test statistic (29.1) is $t = 4(10.3 - 9)/2.1 = 2.48$ and, since $|t| > t_{15,0.975} = 2.131$, the two-sided test rejects at level 0.05, as found using duality in exercise 25.4. For $\mu_0 = 11.5$ we have $t = 4(10.3 - 11.5)/2.1 = -2.29$ and $|t| = 2.29 > 2.131$, so H_0 is rejected again. This agrees with the interval, since 11.5 is outside the interval $[9.18, 11.42]$.

2. From theorem 25.6 we know that the two-sided test from proposition 29.3 rejects σ_0^2 at level 0.05 if and only if σ_0^2 is outside the 95% confidence interval $[2.41, 10.56]$ from example 25.5. This is the inverse of the test, so $\sigma_0^2 = 2$ is rejected and $\sigma_0^2 = 9$ is not. Looking at the values directly, for $(n-1)s^2 = 15 \cdot 4.41 = 66.15$ the test statistic corresponding to $\sigma_0^2 = 2$ is $v = 66.15/2 = 33.08$, which is greater than $\chi_{15,0.975}^2 = 27.488$, so H_0 is rejected. For $\sigma_0^2 = 9$ the test statistic $v = 66.15/9 = 7.35$ falls into the interval $[6.262, 27.488]$ and thus H_0 is not rejected.
3. For $\mu_0 = 9.5$ the test statistic is $t = 4(10.3 - 9.5)/2.1 = 1.52$ and, since $1.52 \leq t_{15,0.95} = 1.753$, the one-sided test does not reject at level 0.05. The two-sided interval is the inverse of the two-sided test with level α , *i.e.* the upper tail of the two-sided test has probability $\alpha/2$. If we decided a one-sided test by checking whether μ_0 is less than the lower bound of the interval, we would obtain the one-sided test at level $\alpha/2 = 0.025$, not at level 0.05. Thus, the one-sided test with 0.05 is not the same as the two-sided test. The confidence set corresponding to the one-sided test with 0.05 is, as in exercise 25.4, the one-sided interval $[\bar{x} - t_{15,0.95} s/\sqrt{n}, \infty) = [10.3 - 1.753 \cdot 0.525, \infty) = [9.38, \infty)$, and 9.5 is contained in this interval, as expected.

□

C.21 The Exponential Model: From Start to Finish

Here we solve the exercises at the end of lecture 31.

Solution to exercise 31.1.

1. The log-likelihood function is $\ell(\theta) = n \log \theta - \theta T$, with $T = \sum_i x_i = n\bar{x}$, and the MLE is $\hat{\theta} = n/T$. Using $\hat{\theta}T = n$ and $\theta_0 T = n\theta_0/\hat{\theta}$ we find

$$W = 2(\ell(\hat{\theta}) - \ell(\theta_0)) = 2(n \log \hat{\theta} - n - n \log \theta_0 + n\theta_0/\hat{\theta}) = 2n \left(\log \frac{\hat{\theta}}{\theta_0} - 1 + \frac{\theta_0}{\hat{\theta}} \right).$$

Using $u = \theta_0 \bar{x} = \theta_0/\hat{\theta}$ we have $\log(\hat{\theta}/\theta_0) = -\log u$ and $\theta_0/\hat{\theta} = u$, and thus $W = 2n(u - 1 - \log u)$.

2. Let $g(u) = u - 1 - \log u$ for $u > 0$. Then $g(1) = 0$ and $g'(u) = 1 - 1/u$. Since this value is negative for $u < 1$ and positive for $u > 1$, $u = 1$ must be the unique minimum of g and $g(u) > 0$ for $u \neq 1$. Since $W = 2ng(u)$, we have $W \geq 0$, with equality if and only if $u = 1$, *i.e.* if and only if $\hat{\theta} = \theta_0$.
3. For $\theta_0 = 0.5$ we have $u = 1.25$ and $W = 20(1.25 - 1 - \log 1.25) = 20 \cdot 0.0269 = 0.54$. This value is less than $\chi_{1,0.95}^2 = 3.84$, so $H_0: \theta = 0.5$ is not rejected at level 0.05. For $\theta_0 = 1$ we have $u = 2.5$ and $W = 20(2.5 - 1 - \log 2.5) = 20 \cdot 0.584 = 11.67$. This value is greater than 3.84, so $H_0: \theta = 1$ is rejected: the data are not compatible with a mean lifetime of one year.

This completes the solution.

□

Solution to exercise 31.2.

1. From example 11.8 we know that $\mathcal{I}(\theta) = 1/\theta^2$. Thus, the Jeffreys prior is $\pi_J(\theta) \propto \sqrt{\mathcal{I}(\theta)} = 1/\theta$. Since $\int_0^1 \theta^{-1} d\theta = \infty$, the prior has infinite integral and thus is improper.
2. The posterior density is proportional to $\theta^{-1} \cdot \theta^n e^{-\theta T} = \theta^{n-1} e^{-\theta T}$. This is the $\text{Gamma}(n, T)$ density up to the constant $T^n/\Gamma(n)$. Since $T > 0$, this constant is finite, so the posterior is proper and corresponds to the $\text{Gamma}(n, T)$ distribution.
3. From exercise 11.5 we know that the information about the mean is $\mathcal{I}_\mu(\mu) = 1/\mu^2$. Thus, the Jeffreys prior for μ is $\pi_J(\mu) \propto 1/\mu$. Using the transformation $\pi_J(\theta)$ with $\theta = 1/\mu$ and $|d\theta/d\mu| = 1/\mu^2$ we find the transformed density $\pi_J(1/\mu)/\mu^2 \propto \mu/\mu^2 = 1/\mu$. This is the same prior as $\pi_J(\mu)$ above, as expected from proposition 26.9.
4. For $\theta \sim \text{Gamma}(n, T)$ we have $2T\theta \sim \text{Gamma}(n, 1/2) = \chi_{2n}^2$ and thus the equal-tailed 95% credible interval is $[\chi_{2n, 0.025}^2/(2T), \chi_{2n, 0.975}^2/(2T)]$. These are the same numbers as for the exact confidence interval. The confidence interval covers the fixed value θ in 95% of repeated samples. The credible interval, on the other hand, covers the random variable θ with posterior probability 0.95 given the observed data.

This completes the solution. □

Solution to exercise 31.3.

- (a) The distribution function of one observation is $F(y) = y/\theta$ for $0 \leq y \leq \theta$, with density $f(y) = 1/\theta$ there. By proposition 8.9 the minimum has density

$$f_{(1)}(y) = n(1 - F(y))^{n-1} f(y) = \frac{n}{\theta} \left(1 - \frac{y}{\theta}\right)^{n-1}, \quad 0 \leq y \leq \theta,$$

and zero otherwise. Substituting $u = y/\theta$ we find

$$\mathbb{E}_\theta(X_{(1)}) = \theta \int_0^1 n u (1-u)^{n-1} du = \theta n B(2, n) = \theta n \frac{1! (n-1)!}{(n+1)!} = \frac{\theta}{n+1},$$

using the beta function of table A.2, or equivalently the mean $1/(n+1)$ of the $\text{Beta}(1, n)$ distribution of $X_{(1)}/\theta$.

- (b) By part (a) and the linearity of expectation, $\mathbb{E}_\theta(U) = (n+1)\theta/(n+1) = \theta$ for all θ , and thus U is unbiased. The same substitution gives

$$\begin{aligned} \mathbb{E}_\theta(X_{(1)}^2) &= \theta^2 \int_0^1 n u^2 (1-u)^{n-1} du = \theta^2 n B(3, n) \\ &= \theta^2 \frac{2n(n-1)!}{(n+2)!} = \frac{2\theta^2}{(n+1)(n+2)}, \end{aligned}$$

and thus

$$\begin{aligned} \text{Var}_\theta(X_{(1)}) &= \frac{2\theta^2}{(n+1)(n+2)} - \frac{\theta^2}{(n+1)^2} \\ &= \frac{\theta^2(2(n+1) - (n+2))}{(n+1)^2(n+2)} = \frac{n\theta^2}{(n+1)^2(n+2)} \end{aligned}$$

and $\text{Var}_\theta(U) = (n+1)^2 \text{Var}_\theta(X_{(1)}) = n\theta^2/(n+2)$. Since U is unbiased, $\text{MSE}(U) = \text{Var}_\theta(U) \rightarrow \theta^2$ as $n \rightarrow \infty$, and thus the mean squared error does not tend to zero. This alone does not disprove consistency, since theorem 2.10 is only a sufficient condition; but directly, for $0 < \varepsilon < \theta$ we have

$$\begin{aligned} \mathbb{P}_\theta(|U - \theta| > \varepsilon) &\geq \mathbb{P}_\theta(U < \theta - \varepsilon) = \mathbb{P}_\theta\left(X_{(1)} < \frac{\theta - \varepsilon}{n+1}\right) \\ &= 1 - \left(1 - \frac{\theta - \varepsilon}{(n+1)\theta}\right)^n \rightarrow 1 - e^{-(\theta - \varepsilon)/\theta} > 0, \end{aligned}$$

and thus U does not converge to θ in probability: U is not consistent in the sense of definition 2.8. The reason is that U uses only the smallest observation, which carries little information about the upper endpoint θ .

- (c) Writing the density of one observation as $\theta^{-1}\mathbf{1}_{\{0 \leq x \leq \theta\}}$, the joint density is

$$f(x; \theta) = \theta^{-n} \prod_{i=1}^n \mathbf{1}_{\{0 \leq x_i \leq \theta\}} = \theta^{-n} \mathbf{1}_{\{x_{(n)} \leq \theta\}} \cdot \mathbf{1}_{\{x_{(1)} \geq 0\}},$$

since all x_i are at most θ if and only if the largest is, and all are non-negative if and only if the smallest is. This is of the form (8.1) with $g(x_{(n)}; \theta) = \theta^{-n} \mathbf{1}_{\{x_{(n)} \leq \theta\}}$ and $h(x) = \mathbf{1}_{\{x_{(1)} \geq 0\}}$, and thus $X_{(n)}$ is sufficient for θ by theorem 8.2. The indicator involving θ has to stay in g .

- (d) Conditionally on $X_{(n)} = m$, the other $n-1$ observations are i.i.d. uniform on $(0, m)$, and $X_{(1)}$ is their minimum. By part (a) with $n-1$ in place of n and m in place of θ , the minimum of $n-1$ i.i.d. uniform observations on $(0, m)$ has mean m/n , and thus

$$\mathbb{E}_\theta(X_{(1)} | X_{(n)} = m) = \frac{m}{n}, \quad V = \mathbb{E}_\theta(U | X_{(n)}) = (n+1) \frac{X_{(n)}}{n}.$$

The Rao–Blackwell theorem, theorem 10.1, applies because $X_{(n)}$ is sufficient by part (c) and U is unbiased with finite variance by part (b). It guarantees that V does not depend on θ , which the formula confirms, that V is unbiased for θ , and that $\text{Var}_\theta(V) \leq \text{Var}_\theta(U)$.

- (e) By proposition 10.5 the maximum $X_{(n)}$ is complete, and it is sufficient by part (c). Since V is a function of $X_{(n)}$ and unbiased for θ , the Lehmann–Scheffe theorem, theorem 10.4, shows that V is the UMVUE for θ . Its variance $\theta^2/(n(n+2))$ is smaller than $\text{Var}_\theta(U) = n\theta^2/(n+2)$ by a factor n^2 , and it lies below the Cramer–Rao value θ^2/n . There is no contradiction, because theorem 13.2 does not apply to this model: its conditions, section 11.1, require the support of the density not to depend on θ , and here the support is $(0, \theta)$. The identity $\text{Cov}_\theta(W, \ell'(\theta)) = 1$ on which the proof rests fails, because differentiating $\int W f \, dx$ with respect to θ picks up a term from the moving boundary of the integral. The value θ^2/n is therefore not a bound at all for this model, and estimators based on the maximum are super-efficient in the sense of example 13.6.

□

Solution to exercise 31.4.

- (a) The tray comes from supplier A with probability π , and given the supplier the count is binomial, and thus by the law of total probability

$$f(x; \pi) = \pi f_A(x) + (1 - \pi) f_B(x), \quad x \in \{0, 1, \dots, 10\},$$

a two-component mixture with known components. For n independent trays the observed-data log-likelihood is

$$\begin{aligned} \ell(\pi) &= \sum_{i=1}^n \log(\pi f_A(x_i) + (1 - \pi) f_B(x_i)), \\ \ell'(\pi) &= \sum_{i=1}^n \frac{f_A(x_i) - f_B(x_i)}{\pi f_A(x_i) + (1 - \pi) f_B(x_i)}. \end{aligned}$$

The likelihood equation $\ell'(\pi) = 0$ is, after clearing denominators, a polynomial equation of degree $n - 1$ in π , and the sum inside the logarithm prevents the simplification which made the examples of lecture 5 solvable. The complete data are the pairs (X_i, Z_i) with $f_c(x, 1; \pi) = \pi f_A(x)$ and $f_c(x, 2; \pi) = (1 - \pi) f_B(x)$, and thus

$$\ell_c(\pi) = \sum_{i=1}^n \left[\mathbf{1}_{\{Z_i=1\}} (\log \pi + \log f_A(x_i)) + \mathbf{1}_{\{Z_i=2\}} (\log(1 - \pi) + \log f_B(x_i)) \right],$$

which is linear in the indicators.

- (b) By definition 19.1 the E step computes $Q(\pi | \pi_k) = \mathbb{E}_{\pi_k}(\ell_c(\pi) | X)$. Since ℓ_c is linear in the indicators, and since the pairs are independent, this replaces $\mathbf{1}_{\{Z_i=1\}}$ by its conditional expectation $\mathbb{P}_{\pi_k}(Z_i = 1 | X_i = x_i)$, which by Bayes' rule, proposition A.20, is

$$\gamma_i = \frac{\mathbb{P}_{\pi_k}(Z_i = 1) \mathbb{P}(X_i = x_i | Z_i = 1)}{\mathbb{P}_{\pi_k}(X_i = x_i)} = \frac{\pi_k f_A(x_i)}{\pi_k f_A(x_i) + (1 - \pi_k) f_B(x_i)},$$

and $\mathbf{1}_{\{Z_i=2\}}$ by $1 - \gamma_i$. Thus

$$Q(\pi | \pi_k) = \sum_{i=1}^n \left[\gamma_i \log \pi + (1 - \gamma_i) \log(1 - \pi) \right] + \text{terms not involving } \pi.$$

- (c) Differentiating with the γ_i held fixed gives

$$\frac{\partial Q}{\partial \pi} = \frac{\sum_i \gamma_i}{\pi} - \frac{\sum_i (1 - \gamma_i)}{1 - \pi} = \frac{\sum_i \gamma_i - n\pi}{\pi(1 - \pi)},$$

which vanishes at $\pi_{k+1} = \frac{1}{n} \sum_i \gamma_i$, is positive to the left of this value and negative to the right. Thus π_{k+1} is the maximiser; equivalently, the second derivative $-\sum_i \gamma_i / \pi^2 - \sum_i (1 - \gamma_i) / (1 - \pi)^2$ is negative for every $\pi \in (0, 1)$, so $Q(\cdot | \pi_k)$ is concave on $(0, 1)$ and π_{k+1} is its global maximiser. The update is the complete-data MLE of example 5.3 with the unobserved indicators replaced by responsibilities.

- (d) With $\pi_0 = 1/2$ the weights cancel and $\gamma_i = f_A(x_i)/(f_A(x_i) + f_B(x_i))$. From the table we get, to four decimal places,

$$\gamma = (0.9952, 0.0029, 0.9995, 0.0003, 0.9572) \quad \text{for } x = (2, 7, 1, 8, 3),$$

and thus $\sum_i \gamma_i = 2.9551$ and $\pi_1 = 2.9551/5 = 0.5910$. The responsibilities are all close to 0 or 1. Since the two component distributions barely overlap, each count identifies its supplier almost with certainty, and the trays with 1, 2 and 3 germinated seeds are assigned to supplier A, the trays with 7 and 8 to supplier B. The tray with 3 seeds is the least certain, at 0.957. Further iterations move π_k only slightly, to the fixed point 0.594; the observed-data log-likelihood rises from -10.30 at π_0 to -10.22 at π_1 , as theorem 19.2 requires.

- (e) With p_A unknown, $\log f_A(x_i) = x_i \log p_A + (m - x_i) \log(1 - p_A) + \text{const}$ can no longer be dropped from Q , which now contains the term

$$\sum_i \gamma_i (x_i \log p_A + (m - x_i) \log(1 - p_A)).$$

Setting the derivative of this term with respect to p_A to zero gives

$$\frac{\sum_i \gamma_i x_i}{p_A} - \frac{\sum_i \gamma_i (m - x_i)}{1 - p_A} = 0, \quad i.e. \quad p_{A,k+1} = \frac{\sum_i \gamma_i x_i}{m \sum_i \gamma_i},$$

the responsibility-weighted proportion of germinated seeds among the trays attributed to supplier A, with the responsibilities now computed from $(\pi_k, p_{A,k})$. For the data of part (d) and the responsibilities computed there this gives $p_{A,1} = 5.884/29.55 = 0.199$. The ascent property, theorem 19.2, states that the observed-data log-likelihood never decreases along the EM sequence, $\ell(\pi_{k+1}, p_{A,k+1}) \geq \ell(\pi_k, p_{A,k})$ for all k . It guarantees that the sequence of log-likelihood values converges, and by proposition 19.3 a point at which the iteration comes to rest is a stationary point of ℓ ; it does not guarantee that this point is the global maximum, since a poor starting value can lead to a local maximum, and it says nothing about the speed of convergence.

□

Bibliography

- Hamsa Bastani, Osbert Bastani, Alp Sungu, Haosen Ge, Özge Kabakcı, and Rei Mariman. Generative AI without guardrails can harm learning: Evidence from high school mathematics. *Proceedings of the National Academy of Sciences*, 122(26):e2422633122, 2025. doi: 10.1073/pnas.2422633122.
- W. J. Braun and D. J. Murdoch. *A First Course in Statistical Programming with R*. Cambridge University Press, Cambridge, 3rd edition, 2021.
- G. Casella and R. L. Berger. *Statistical Inference*. Duxbury, Pacific Grove, 2nd edition, 2002.
- D. R. Cox and D. V. Hinkley. *Theoretical Statistics*. Chapman and Hall, London, 1974.
- A. C. Davison. *Statistical Models*. Cambridge University Press, Cambridge, 2003.
- M. H. DeGroot and M. J. Schervish. *Probability and Statistics*. Addison-Wesley, Boston, 4th edition, 2012.
- A. Gelman, J. B. Carlin, H. S. Stern, D. B. Dunson, A. Vehtari, and D. B. Rubin. *Bayesian Data Analysis*. Chapman and Hall/CRC, Boca Raton, 3rd edition, 2013.
- G. H. Givens and J. A. Hoeting. *Computational Statistics*. Wiley, Hoboken, 2nd edition, 2013.
- Mark Hodds, Lara Alcock, and Matthew Inglis. Self-explanation training improves proof comprehension. *Journal for Research in Mathematics Education*, 45(1):62–101, 2014. doi: 10.5951/jresmetheduc.45.1.0062.
- P. M. Lee. *Bayesian Statistics: An Introduction*. Wiley, Chichester, 4th edition, 2012.
- E. L. Lehmann and G. Casella. *Theory of Point Estimation*. Springer, New York, 2nd edition, 1998.
- J. A. Rice. *Mathematical Statistics and Data Analysis*. Duxbury, Belmont, 3rd edition, 2007.
- David Strömberg, Victor Lei, and Yanhui Wu. The generative AI learning penalty: Evidence from Chinese secondary education. Discussion Paper DP21577, Centre for Economic Policy Research, London, 2026.

- G. A. Young and R. L. Smith. *Essentials of Statistical Inference*. Cambridge University Press, Cambridge, 2005.

Index

- alternative hypothesis, 131
- ancillary statistic, 65
- ascent property, 126
- asymptotic normality of the MLE, 110
- asymptotically efficient estimator, 111
- asymptotically unbiased estimator, 7
- Basu's theorem, 65
- Bayes estimator, 176
 - posterior mean, 176
- Bayes factor, 188
- Bayes' theorem, 229
 - for the parameter, 173
- Bayesian test, 187
- bias, 6
- bias-variance decomposition, 8
- Cauchy–Schwarz inequality, 226
- central limit theorem, 228
- change of variable, 229
- Chebyshev's inequality, 225
- chi-squared distribution, 224
- chi-squared test for a variance, 195
- complete data, 124
- complete statistic, 51
 - examples, 64
- complete-data log-likelihood, 124
- completeness, 51
- composite hypothesis, 131
- conditional expectation
 - law of total variance, 227
 - taking out what is known, 227
 - tower property, 227
- confidence bound, 170
- confidence interval, 166
 - by inverting a test, 168
 - one-sided, 170
 - Wald, 169
- confidence level, 166
- confidence set, 166
- conjugate prior, 174
 - exponential family, 174
 - normal-normal, 175
- conservative test, 135
- consistent estimator, 10
 - and asymptotic unbiasedness, 11
 - and mean squared error, 10
 - maximum likelihood estimator, 106
- continuous mapping theorem, 228
- convergence
 - in distribution, 228
 - in probability, 227
- convergence in distribution, 228
- convergence in probability, 227
- covariance, 222
- coverage probability, 166
- Cramer–Rao inequality, 84
 - attainment, 86
 - for a function of the parameter, 85
 - multiparameter, 92
- credibility, 186
- credible interval, 186
 - equal-tailed, 186
 - highest posterior density, 187
- critical value, 132
- critical region, 132
- cumulant function, 41
 - moments of the natural statistic, 41
- delta method, 228
- density, 2
- efficiency, 85
 - asymptotic, 111
- efficient estimator, 85
- EM algorithm, 125
 - ascent property, 126
 - complete-data log-likelihood, 124
 - fixed points, 127
- estimate, 4
- estimator, 4
 - asymptotically efficient, 111

- asymptotically unbiased, 7
- Bayes, 176
- consistent, 10
- efficient, 85
- family of, 4
- maximum likelihood, 25
- method of moments, 20
- unbiased, 6
- uniformly minimum variance unbiased, 64
- expectation
 - linearity, 222
- exponential family, 39
 - completeness, 64
 - conjugate prior, 174
 - Fisher information, 73
 - full rank, 43
 - k -parameter, 42
 - mean and variance, 41
 - one-parameter, 39
 - sample from, 40
 - shape, 39–40
 - sufficient statistic, 49
- F distribution, 225
- F -test, 196
- factorisation theorem, 48
- family of estimators, 4
- Fisher information, 70
 - additivity, 71
 - in an exponential family, 73
 - information identity, 71
 - matrix, 91
 - reparametrisation, 73
- Fisher information matrix, 91
 - information identity, 91
 - properties, 91
- Fisher–Neyman factorisation theorem, 48
- flat prior, 176
- full rank, 43
- generalised likelihood ratio, 152
- generalised likelihood-ratio test, 152
- highest posterior density interval, 187
- HPD interval, *see* highest posterior density interval
- hyperparameter, 173
- hypothesis, 131
 - alternative, 131
 - composite, 131
 - null, 131
 - simple, 131
- identifiability, 22
- improper prior, 176
- index, 307
- inequality
 - Cauchy–Schwarz, 226
 - Chebyshev’s, 225
 - Jensen’s, 226
 - Markov’s, 225
- information identity, 71
- invariance of the MLE, 29
- Jeffreys prior, 176
 - invariance, 176
 - location and scale models, 177
- Jensen’s inequality, 226
- Kullback–Leibler divergence, 105
- label switching, 129
- latent variable, 124
- law of large numbers, 228
- law of total variance, 227
- Lehmann–Scheffe theorem, 64
- level of a test, 133
- likelihood equation, 25
- likelihood function, 18
- likelihood ratio, 146
 - at the true parameter, 104
- likelihood-ratio test, 152
 - Wilks’ theorem, 154
- log-likelihood, 18
- log-partition function, *see* cumulant function
- Markov’s inequality, 225
- maximum likelihood estimator, 25
 - asymptotic normality, 110
 - consistency, 106
 - invariance, 29
 - standard error, 113
- mean squared error, 8
 - bias-variance decomposition, 8
- method of moments, 18
- method of moments estimator, 20
- minimal sufficient statistic, 50
 - criterion, 50

- mixture model, 125
- MLE, *see* maximum likelihood estimator
- moment
 - population, 20
 - sample, 20
- most powerful test, 145
- natural parameter, 39
- natural parameter space, 41
- natural parametrisation, 41
- natural statistic, 39
- nested models, 152
- Neyman–Pearson lemma, 146
- normal distribution
 - conjugate prior, 175
 - independence of sample mean and variance, 66
 - sample mean and sample variance, 224
 - test for the variance, 195
- nuisance parameter, 94
- null hypothesis, 131
- observed data, 124
- one-sided test, 148
- order statistics
 - maximum and minimum, 51
- p -value, 135
- parameter, 2
- parameter space, 2
- parametric model, 2
- pivot, *see* pivotal quantity
- pivotal quantity, 167
- posterior density, 173
 - Bayes’ theorem, 173
- posterior distribution, 173
- posterior mean, 176
- posterior odds, 188
- posterior probability, 187
- power, 132
- power function, 132
- prior density, 173
- prior distribution, 173
 - conjugate, 174
 - flat, 176
 - improper, 176
 - Jeffreys, 176
- prior odds, 188
- probability mass function, *see* probability weights
- probability weights, 2
- random sample, 1
- Rao–Blackwell theorem, 62
 - Rao–Blackwellisation, 63
- Rao–Blackwellisation, 63
- regularity conditions, 103
- sample mean
 - consistency, 10
 - mean and variance, 6
- sample variance, 7
 - normal sample, 224
 - unbiasedness, 7
- sampling distribution, 4
- score function, 20
 - covariance with a statistic, 83
 - mean zero, 70
- score test, 156
- score vector, 90
- significance level, 133
- simple hypothesis, 131
- size of a test, 133
- Slutsky’s theorem, 228
- standard error, 113
 - of the MLE, 113
- statistic, 4
 - ancillary, 65
 - complete, 51
 - minimal sufficient, 50
 - natural, 39
 - sufficient, 47
 - test, 132
- sufficient statistic, 47
 - factorisation theorem, 48
 - in an exponential family, 49
 - minimal, 50
- super-efficiency, 87
- t distribution, 225
- t -interval, 168
- t -test, 194
 - one-sample, 194
 - paired, 196
 - two-sample, 196
- taking out what is known, 227
- test, 132
 - Bayesian, 187

- conservative, 135
- likelihood-ratio, 152
- most powerful, 145
- one-sided, 148
- score, 156
- uniformly most powerful, 148
- Wald, 156
- test function, 132
- test inversion, 168
- test statistic, 132
- tower property, 227
- true parameter value, 2
- type I error, 132
- type II error, 132
- UMVUE, *see* uniformly minimum
variance unbiased estimator
- unbiased estimator, 6
 - asymptotically, 7
- uniformly minimum variance unbiased
estimator, 64
 - Lehmann–Scheffe theorem, 64
- uniformly most powerful test, 148
 - non-existence for two-sided
alternatives, 149
 - one-sided alternative, 148
- variance
 - of a linear function, 222
- Wald interval, 169
- Wald test, 156
- weak law of large numbers, 228
- Wilks' theorem, 154